

Exam AWS Certified Solutions Architect - Associate SAA-C03 (2026/06/15)

淘宝: "IT认证轻松过" <https://shop542998714.m.taobao.com>

咸鱼: "IT认证轻松过" <https://www.goofish.com/personal?userId=2645260430>

微信: "Examtopics"



Question #1

Topic 1

A company collects data for temperature, humidity, and atmospheric pressure in cities across multiple continents. The average volume of data that the company collects from each site daily is 500 GB. Each site has a high-speed Internet connection. The company wants to aggregate the data from all these global sites as quickly as possible in a single Amazon S3 bucket. The solution must minimize operational complexity. Which solution meets these requirements?

- A. Turn on S3 Transfer Acceleration on the destination S3 bucket. Use multipart uploads to directly upload site data to the destination S3 bucket. **Most Voted**
- B. Upload the data from each site to an S3 bucket in the closest Region. Use S3 Cross-Region Replication to copy objects to the destination S3 bucket. Then remove the data from the origin S3 bucket.
- C. Schedule AWS Snowball Edge Storage Optimized device jobs daily to transfer data from each site to the closest Region. Use S3 Cross-Region Replication to copy objects to the destination S3 bucket.
- D. Upload the data from each site to an Amazon EC2 instance in the closest Region. Store the data in an Amazon Elastic Block Store (Amazon EBS) volume. At regular intervals, take an EBS snapshot and copy it to the Region that contains the destination S3 bucket. Restore the EBS volume in that Region.

Correct Answer: A

Community vote distribution

A (94%)

5%

Comments

PhucVuu **Highly Voted** 1 year, 8 months ago

Selected Answer: A

Keyword:

From GLOBAL sites as quickly as possible in a SINGLE S3 bucket.

Minimize operational complexity

A. is correct because S3 Transfer Acceleration is support for high speed transfer in Edge location and you can upload it immediately. Also with multipart uploads your big file can be uploaded in parallel.

B, C, D. is not minimize operational and fast when compare to answer A

upvoted 45 times

Ruffyt **Highly Voted** 1 year, 8 months ago

General line: Collect huge amount of the files across multiple continents

Conditions: High speed Internet connectivity

Task: aggregate the data from all in a single S3 bucket

Requirements: as quick as possible, minimize operational complexity

Correct answer A: S3 Transfer Acceleration because:

- ideally works with objects for long-distance transfer (uses Edge Locations)

- can speed up content transfers to and from S3 as much as 50-500%

- use cases: mobile & web application uploads and downloads, distributed office transfers, data exchange with trusted partners. Generally for sharing

of large data sets between companies, customers can set up special access to their S3 buckets with accelerated uploads to speed data exchanges and the pace of innovation.

B - about disaster recovery

C - about transferring data between your local environment and the AWS Cloud

D - about disaster recovery

upvoted 16 times

pentium75 2 years, 5 months ago

C, we DO have a "local environment" (a "site") that collect data, and THAT data must go to "the AWS cloud". Why not C? The stem says nothing about large objects, which would be a requirement for the "multipart upload" that is mentioned.

upvoted 3 times

BrianEllis 9 months, 1 week ago

C wouldn't apply because it increases complexity. No if during the investigative process you found that Snowball was present then you could discuss this. But chances are if snowball is present then transfer rate isn't an issue in the first place.

The same could be said for Private Link connections per site.

upvoted 1 times

pexta **Most Recent** 1 month, 1 week ago

Selected Answer: A

Edge locations high speed, big file sinot smaller chucnks and easy retry

upvoted 1 times

script001 3 months, 2 weeks ago

Selected Answer: A

The question minimizes the operational complexity

upvoted 1 times

Anjana_e2300554 3 months, 2 weeks ago

Selected Answer: A

fast, simple, minimal management.

upvoted 1 times

mebmg 3 months, 2 weeks ago

Selected Answer: A

Huge data volume: ~500 GB per site, per day

Global sites, all with high-speed internet

As fast as possible aggregation

Single Amazon S3 bucket

Minimal operational complexity

So the winning solution should:

Push data directly to the final bucket

Use AWS's optimized global network

Avoid extra buckets, replication rules, or hardware

upvoted 1 times

Sukri123 3 months, 2 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

SUGI369 3 months, 2 weeks ago

Selected Answer: A

The solution must minimize operational complexity.

upvoted 1 times

PatrickSaliba 3 months, 3 weeks ago

Selected Answer: A

Guys check out cloudchamp.de

Same questions but way cheaper...

upvoted 1 times

nesyto 4 months, 4 weeks ago

Selected Answer: A

correct answer is A

upvoted 1 times

nesyto 5 months ago

Selected Answer: A

A is correct

upvoted 1 times

suvam_2121 6 months ago

Selected Answer: A

A. is correct because S3 Transfer Acceleration is support for high speed transfer in Edge location and you can upload it immediately.

upvoted 1 times

suvam_2121 6 months ago

Selected Answer: A

A. is correct because S3 Transfer Acceleration is support for high speed transfer in Edge location and you can upload it immediately.

upvoted 1 times

Chrpj 9 months, 1 week ago

Selected Answer: A

A is correct

upvoted 1 times

deljy 9 months, 2 weeks ago

Selected Answer: B

note

Keyword:

From GLOBAL sites as quickly as possible in a SINGLE S3 bucket.

Minimize operational complexity

A. is correct because S3 Transfer Acceleration is support for high speed transfer in Edge location and you can upload it immediately. Also with multipart uploads your big file can be uploaded in parallel.

B, C, D. is not minimize operational and fast when compare to answer A

upvoted 1 times

hlocuwu 9 months, 2 weeks ago

Selected Answer: A

As quickly as possible, minimize operational complexity

upvoted 1 times

Aleeexp 9 months, 2 weeks ago

Selected Answer: A

Vote A

upvoted 1 times

A company needs the ability to analyze the log files of its proprietary application. The logs are stored in JSON format in an Amazon S3 bucket. Queries will be simple and will run on-demand. A solutions architect needs to perform the analysis with minimal changes to the existing architecture.

What should the solutions architect do to meet these requirements with the LEAST amount of operational overhead?

- A. Use Amazon Redshift to load all the content into one place and run the SQL queries as needed.
- B. Use Amazon CloudWatch Logs to store the logs. Run SQL queries as needed from the Amazon CloudWatch console.
- C. Use Amazon Athena directly with Amazon S3 to run the queries as needed. **Most Voted**
- D. Use AWS Glue to catalog the logs. Use a transient Apache Spark cluster on Amazon EMR to run the SQL queries as needed.

Correct Answer: C

Community vote distribution

C (99%)

Comments

airraid2010 **Highly Voted** 3 years, 8 months ago

Answer: C

Amazon Athena is an interactive query service that makes it easy to analyze data directly in Amazon Simple Storage Service (Amazon S3) using standard SQL. With a few actions in the AWS Management Console, you can point Athena at your data stored in Amazon S3 and begin using standard SQL to run ad-hoc queries and get results in seconds.

upvoted 82 times

BoboChow 3 years, 8 months ago

I agree C is the answer

upvoted 2 times

tt79 3 years, 8 months ago

C is right.

upvoted 2 times

PhucVuu **Highly Voted** 3 years, 2 months ago

Selected Answer: C

Keyword:

- Queries will be simple and will run on-demand.

- Minimal changes to the existing architecture.

A: Incorrect - We have to do 2 step. load all content to Redshift and run SQL query (This is simple query so we can you Athena, for complex query we will apply Redshit)

B: Incorrect - Our query will be run on-demand so we don't need to use CloudWatch Logs to store the logs.

C: Correct - This is simple query we can apply Athena directly on S3

D: Incorrect - This take 2 step: use AWS Glue to catalog the logs and use Spark to run SQL query

upvoted 52 times

RaghavendraReddy **Most Recent** 3 days, 15 hours ago

Selected Answer: C

When data is:

Stored in S3

Queried occasionally/on-demand

Requires SQL

And the question emphasizes minimal operational overhead

Think Amazon Athena first.

Answer: C.

upvoted 1 times

Anjana_e2300554 3 months, 2 weeks ago

Selected Answer: C

the correct answer is c

upvoted 1 times

mebmg 3 months, 2 weeks ago

Selected Answer: C

The logs are already in Amazon S3
Queries are simple and on-demand
Amazon Athena runs SQL directly on S3
No servers, no data loading, no extra setup
Lowest operational effort

upvoted 1 times

Sukri123 3 months, 2 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

awsgcplearning 4 months, 2 weeks ago

Selected Answer: C

Athena helps to analyze the data available in Amazon S3 through SQL syntax
The answer is C

upvoted 1 times

nesyto 4 months, 4 weeks ago

Selected Answer: C

correct is C

upvoted 1 times

nesyto 5 months ago

Selected Answer: C

The correct answer is C

upvoted 1 times

hlocuwu 9 months, 2 weeks ago

Selected Answer: C

queries will be simple and will run on-demand.
minimal changes

upvoted 1 times

Aleeexp 9 months, 2 weeks ago

Selected Answer: C

Vote c

upvoted 1 times

fig 11 months ago

Selected Answer: C

For analyzing JSON log files stored in Amazon S3 with minimal changes and least operational overhead, Amazon Athena is the most suitable solution.

Amazon Athena (Option C): Athena is a serverless query service that allows you to directly query data in Amazon S3 using standard SQL. It's pay-per-query, so there's no infrastructure to manage, and it integrates directly with S3 without requiring data movement. This perfectly fits the "minimal changes to existing architecture" and "least amount of operational overhead" requirements. You can define a table schema over your S3 JSON logs, and then run your simple, on-demand queries.

upvoted 2 times

DovanDu 11 months ago

Selected Answer: C

Vote C

upvoted 1 times

hossex5 11 months, 4 weeks ago

Selected Answer: C

its c bro

upvoted 1 times

clfigueiredo12 1 year ago

Selected Answer: C

C está certo.

upvoted 1 times

K_SAA 1 year, 1 month ago

Selected Answer: C

Because the query is simple and run on demand so C is the correct answer

B: use aws cloudwatch is not designed for this use-case because cloudwatch do not support SQL queries and the log json is stored in S3 and cloudwatch can not directly query from s3. Cloudwatch only support its own queries syntax not standard SQL

A: use aws RedShift is not considered the right option since you would need to load all log data into Redshift, which is required more worked and high operational overhead, it goes against the requirement of this use-case queries is simple and running on demand

D: use aws glue and Amazon EMR is more complex in setup and high operational overhead

In conclusion, use Amazon athena is the right option, its simple and serverless meaning no infrastructure to manage.

upvoted 1 times

topitexamcom 1 year, 1 month ago

Selected Answer: C

C is the right answer

upvoted 1 times

A company uses AWS Organizations to manage multiple AWS accounts for different departments. The management account has an Amazon S3 bucket that contains project reports. The company wants to limit access to this S3 bucket to only users of accounts within the organization in AWS Organizations.

Which solution meets these requirements with the LEAST amount of operational overhead?

- A. Add the `aws:PrincipalOrgID` global condition key with a reference to the organization ID to the S3 bucket policy. **Most Voted**
- B. Create an organizational unit (OU) for each department. Add the `aws:PrincipalOrgPaths` global condition key to the S3 bucket policy.
- C. Use AWS CloudTrail to monitor the `CreateAccount`, `InviteAccountToOrganization`, `LeaveOrganization`, and `RemoveAccountFromOrganization` events. Update the S3 bucket policy accordingly.
- D. Tag each user that needs access to the S3 bucket. Add the `aws:PrincipalTag` global condition key to the S3 bucket policy.

Correct Answer: A

Community vote distribution

A (97%)

3%

Comments

ude **Highly Voted** 3 years, 8 months ago

Selected Answer: A

`aws:PrincipalOrgID` Validates if the principal accessing the resource belongs to an account in your organization.
upvoted 86 times

BoboChow 3 years, 8 months ago

the condition key `aws:PrincipalOrgID` can prevent the members who don't belong to your organization to access the resource
upvoted 24 times

Naneyerocky **Highly Voted** 1 year, 8 months ago

Selected Answer: A

Condition keys: AWS provides condition keys that you can query to provide more granular control over certain actions. The following condition keys are especially useful with AWS Organizations:

`aws:PrincipalOrgID` – Simplifies specifying the Principal element in a resource-based policy. This global key provides an alternative to listing all the account IDs for all AWS accounts in an organization. Instead of listing all of the accounts that are members of an organization, you can specify the organization ID in the Condition element.

`aws:PrincipalOrgPaths` – Use this condition key to match members of a specific organization root, an OU, or its children. The `aws:PrincipalOrgPaths` condition key returns true when the principal (root user, IAM user, or role) making the request is in the specified organization path. A path is a text representation of the structure of an AWS Organizations entity.

upvoted 29 times

Sleepy_Lazy_Coder 2 years, 10 months ago

are we not choosing ou because the least overhead term was use? option B also seems correct
upvoted 5 times

BlackMamba_4 2 years, 9 months ago

Exactly

upvoted 1 times

EMPERBACH 2 years, 1 month ago

As there are many OU, you need more effort to list down OU path. And question mention about least management overhead to allow users in Organization, not single OU.

upvoted 4 times

CrazyVitamin Most Recent 3 weeks, 3 days ago

Selected Answer: A

Why B is Not the BEST Answer

aws:PrincipalOrgPaths:

Is more granular
Requires managing OU hierarchy paths
Adds operational complexity
Better when restricting access to specific OUs only

But the requirement says:

“only users of accounts within the organization”

—not OU-specific access.

So aws:PrincipalOrgID is simpler and requires the least maintenance.

upvoted 1 times

Anjana_e2300554 3 months, 2 weeks ago

Selected Answer: A

the correct answer is A

upvoted 1 times

Sukri123 3 months, 2 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

suji098 3 months, 2 weeks ago

Selected Answer: A

from myself

upvoted 1 times

nesyto 4 months, 4 weeks ago

Selected Answer: A

correct answer is A

upvoted 1 times

nesyto 5 months ago

Selected Answer: A

The correct answer is A

upvoted 1 times

BlazerX998 7 months ago

Selected Answer: A

A is the correct option. Add the aws:PrincipalOrgID global condition key with a reference to the organization ID to the S3 bucket policy because it directly restricts access to users within the organization with the least amount of operational overhead. This method uses AWS Organizations to verify the principal's origin, requiring no extra infrastructure or constant policy updates for new accounts within the organization.

upvoted 1 times

Aleexp 9 months, 2 weeks ago

Selected Answer: A

Vote A

upvoted 1 times

KiwiCode 9 months, 4 weeks ago

Selected Answer: A

Answer a for sure

upvoted 1 times

fig 11 months ago

Selected Answer: C

Add the aws:PrincipalOrgID global condition key with a reference to the organization ID to the S3 bucket policy. This is the most efficient and low-overhead solution. The aws:PrincipalOrgID condition key directly checks if the AWS account making the request belongs to the specified AWS Organization. This means you set the policy once, and it automatically applies to any account that joins or leaves your organization. It directly addresses the requirement of limiting access to users of accounts within the organization.

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: A

Use aws:PrincipalOrgID in the bucket policy for a clean, scalable, and low-maintenance solution to restrict access to AWS Organization members.

upvoted 1 times

K_SAA 1 year, 1 month ago

Selected Answer: A

A: use aws PrincipalOrgID global condition key with a with a reference to the organization ID to the S3 bucket policy is the correct answer because the company wants to limit access to only user within the organization.

B: is also be considered since the company wanted to limit access to s3 bucket within organization not mention to specific Organizational unit. if you or the company wanted to restrict access to S3 bucket by specific OU consider option B

C: use AWS cloudtrail to monitor action of users within organization, which is more complex and high operational overhead

D: tag each user and use aws:PrincipalTag global condition key to the S3 bucket policy. You would have to tag each user manually, which is required more work overhead

upvoted 1 times

ernie1976 1 year, 1 month ago

Selected Answer: A

This solution es the most simple comparing to other alternatives, just modify one parameter. aws:PrincipalOrgID

upvoted 1 times

Wylla 1 year, 2 months ago

Selected Answer: A

aws:PrincipalOrgID - global key provides an alternative to listing all the account IDs for all AWS accounts in an organization.

upvoted 1 times

francisizme 1 year, 3 months ago

Selected Answer: A

Not B because: aws:PrincipalOrgPaths require you change the OU path when the user switch to another OU

Not C because: you need to put effort to update policy which introduces operational overhead.

Not D because: require a good tag management across multiple accounts. Like team "HR", team "QA", team "Security" and so forth. Also need to update if the user is moved to another team with different tag.

upvoted 2 times

An application runs on an Amazon EC2 instance in a VPC. The application processes logs that are stored in an Amazon S3 bucket. The EC2 instance needs to access the S3 bucket without connectivity to the internet. Which solution will provide private network connectivity to Amazon S3?

- A. Create a gateway VPC endpoint to the S3 bucket. **Most Voted**
- B. Stream the logs to Amazon CloudWatch Logs. Export the logs to the S3 bucket.
- C. Create an instance profile on Amazon EC2 to allow S3 access.
- D. Create an Amazon API Gateway API with a private link to access the S3 endpoint.

Correct Answer: A

Community vote distribution

A (99%)

Comments

PhucVuu **Highly Voted** 1 year, 8 months ago

Selected Answer: A

Keywords:

- EC2 in VPC
- EC2 instance needs to access the S3 bucket without connectivity to the internet

A: Correct - Gateway VPC endpoint can connect to S3 bucket privately without additional cost

B: Incorrect - You can set up interface VPC endpoint for CloudWatch Logs for private network from EC2 to CloudWatch. But from CloudWatch to S3 bucket: Log data can take up to 12 hours to become available for export and the requirement only need EC2 to S3

C: Incorrect - Create an instance profile just grant access but not help EC2 connect to S3 privately

D: Incorrect - API Gateway like the proxy which receive network from out site and it forward request to AWS Lambda, Amazon EC2, Elastic Load Balancing products such as Application Load Balancers or Classic Load Balancers, Amazon DynamoDB, Amazon Kinesis, or any publicly available HTTPS-based endpoint. But not S3

upvoted 81 times

Austinlorenzmccoy 2 years, 6 months ago

Thank you so much

upvoted 1 times

rxwcl 1 year, 7 months ago

Option C involves creating an instance profile on the EC2 instance to allow S3 access. While this option could potentially work, it would not provide private network connectivity to S3, as the EC2 instance would still need to access S3 over the internet.

Option D involves creating an Amazon API Gateway API with a private link to access the S3 endpoint. This option would not provide private network connectivity to S3, as the API Gateway API is not a network interface that can be used to privately connect to S3.

Overall, Option A is the correct solution for providing private network connectivity to Amazon S3 from an EC2 instance in a VPC.

upvoted 3 times

D2w **Highly Voted** 3 years, 8 months ago

Selected Answer: A

VPC endpoint allows you to connect to AWS services using a private network instead of using the public Internet

upvoted 33 times

72e9d12 **Most Recent** 2 months, 3 weeks ago

Selected Answer: A

VPC gateway is free of charges

upvoted 1 times

Anjana_e2300554 3 months, 2 weeks ago

Selected Answer: A

The correct answer is A
upvoted 1 times

Anjana_e2300554 3 months, 2 weeks ago

Selected Answer: A

The correct answer is A
upvoted 1 times

Sukri123 3 months, 2 weeks ago

Selected Answer: A

A is correct
upvoted 1 times

suji098 3 months, 2 weeks ago

Selected Answer: A

form myself
upvoted 1 times

nesyto 4 months, 4 weeks ago

Selected Answer: A

vpc gateway is dor s3 and dynamo and it is free
upvoted 1 times

nesyto 5 months ago

Selected Answer: A

correct answer is A
upvoted 1 times

hlocuwu 9 months, 2 weeks ago

Selected Answer: A

Keywords: EC2 in a VPC
upvoted 1 times

Aleeexp 9 months, 2 weeks ago

Selected Answer: A

Vote A
upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: A

Use a Gateway VPC Endpoint for S3 to enable secure, private access to the S3 bucket from EC2 without internet connectivity.
upvoted 1 times

K_SAA 1 year, 1 month ago

Selected Answer: A

Keywords:
log store in s3, Ec2 instance need to access to s3 bucket without internet
A: use gateway vpc endpoint allow private network connectivity to s3 bucket without publishing to internet
B: Stream the logs to Amazon CloudWatch Logs. Export the logs to the S3 bucket since the question need ec2 instance stay in vpc access to s3 bucket privately but this way do not grant access to s3 bucket
C: Create an instance profile on Amazon EC2 to allow S3 access. This way only authorization access to s3 bucket, not grant private network connectivity to s3 bucket since this also need internet access
D. Create an Amazon API Gateway API with a private link to access the S3 endpoint. Since API gateway acts like a proxy receive requests from outside and forward to the request to aws services like lambda, etc, not using for s3 access
upvoted 1 times

ernie1976 1 year, 1 month ago

Selected Answer: A

It is a simple communication between EC2 and S3, it is no needed to go out to internet, so and endopint is the best solution.
upvoted 1 times

Palee 1 year, 1 month ago

Selected Answer: A

A is correct
upvoted 1 times

melvis8 1 year, 2 months ago

Selected Answer: A

The correct answer is A because a gateway endpoint allows us to securely and cost efficiently access either an S3 or Amazon DynamoDB database within our VPC

upvoted 1 times

BI_12 1 year, 3 months ago

Selected Answer: D

API gateway

upvoted 1 times

A company is hosting a web application on AWS using a single Amazon EC2 instance that stores user-uploaded documents in an Amazon EBS volume. For better scalability and availability, the company duplicated the architecture and created a second EC2 instance and EBS volume in another Availability Zone, placing both behind an Application Load Balancer. After completing this change, users reported that, each time they refreshed the website, they could see one subset of their documents or the other, but never all of the documents at the same time.

What should a solutions architect propose to ensure users see all of their documents at once?

- A. Copy the data so both EBS volumes contain all the documents
- B. Configure the Application Load Balancer to direct a user to the server with the documents
- C. Copy the data from both EBS volumes to Amazon EFS. Modify the application to save new documents to Amazon EFS
- D. Configure the Application Load Balancer to send the request to both servers. Return each document from the correct server

Most Voted

Correct Answer: C

Community vote distribution

C (98%)

2

Comments

D2w **Highly Voted** 3 years, 8 months ago

Selected Answer: C

Concurrent or at the same time key word for EFS

upvoted 60 times

mikey2000 **Highly Voted** 3 years, 6 months ago

Ebs doesnt support cross az only reside in one Az but Efs does, that why it's c

upvoted 48 times

pbpally 3 years, 1 month ago

And just for clarification to others, you can have COPIES of the same EBS volume in one AZ and in another via EBS Snapshots, but don't confuse that with the idea of having some sort of global capability that has concurrent copying mechanisms.

upvoted 12 times

Anjana_e2300554 **Most Recent** 3 months, 2 weeks ago

Selected Answer: C

The correct answer is C

upvoted 1 times

suji098 3 months, 2 weeks ago

Selected Answer: C

myself

upvoted 1 times

Sukri123 3 months, 2 weeks ago

Selected Answer: C

C is correct

upvoted 1 times

nesyto 4 months, 4 weeks ago

Selected Answer: D

correct answer alb to send request to both servers

upvoted 1 times

BrianEllis 9 months, 1 week ago

Selected Answer: C

Honestly I don't know if any of the answers are great.

Some litmus would be:

- If cost is not an issue then yes EFS is a great solution.
 - If you need local filesystem response rates
 - the data is of such a nature that loading times across HTTP degrade responsiveness
- upvoted 1 times

Aleexp 9 months, 2 weeks ago

Selected Answer: C

VOTE C

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: C

Implement Amazon EFS as shared storage for all EC2 instances behind the ALB. This ensures high availability, scalability, and consistent data access across Availability Zones with minimal operational overhead.

upvoted 3 times

Yaredd 11 months, 3 weeks ago

Selected Answer: C

EFS is best suited for multi AZ and EC2 distribution

upvoted 2 times

hossex5 11 months, 4 weeks ago

Selected Answer: C

c is correct

upvoted 2 times

K_SAA 1 year, 1 month ago

Selected Answer: C

Problem

each EC2 instance has its own EBS volume, and EBS volume is specific for AZ specific
so when user send request to specific ec2 instance they only see the response from each EBS volume not all because EBS volume do not sever for cross region problem

C: use EFS because EFS a fully managed and network file system can be mounted to mutiple ec2 instances across mutiple AZ so user can see all documents response from EFS instead of subset of documents

A. Copy the data so both EBS volumes contain all the documents is not correct because when user send request to one instance on a specific AZ they only see the documents from that EBS AZ not both

upvoted 2 times

Palee 1 year, 1 month ago

Selected Answer: C

C is correct

upvoted 1 times

melvis8 1 year, 2 months ago

Selected Answer: C

The correct anszwer to this question is C question EBS volumes are AZ locked so when transferring the architecture to a different AZ the data can no longer be received by users whereas with EFS volume we can EC2 instances across AZs

upvoted 1 times

kimardamina 1 year, 2 months ago

Selected Answer: C

My immediate understanding was the fact that EFS is multi az and will make it less complex as a solution for this.

upvoted 1 times

francisizme 1 year, 3 months ago

Selected Answer: C

Not A: because it requires sync between the EBS volumes which is complex and not scalable.

Not B: It's not scalable if the LB only directs user to one instance

Not D: Impractical. The application will need to have a mechanism to merge the responses coming from 2 instance

upvoted 1 times

Mrigraj12 1 year, 4 months ago

Selected Answer: C

Use common storage like EFS so that the instances can access the document without any hassle and without being uploaded into them separately
upvoted 1 times

A company uses NFS to store large video files in on-premises network attached storage. Each video file ranges in size from 1 MB to 500 GB. The total storage is 70 TB and is no longer growing. The company decides to migrate the video files to Amazon S3. The company must migrate the video files as soon as possible while using the least possible network bandwidth.

Which solution will meet these requirements?

- A. Create an S3 bucket. Create an IAM role that has permissions to write to the S3 bucket. Use the AWS CLI to copy all files locally to the S3 bucket.
- B. Create an AWS Snowball Edge job. Receive a Snowball Edge device on premises. Use the Snowball Edge client to transfer data to the device. Return the device so that AWS can import the data into Amazon S3. **Most Voted**
- C. Deploy an S3 File Gateway on premises. Create a public service endpoint to connect to the S3 File Gateway. Create an S3 bucket. Create a new NFS file share on the S3 File Gateway. Point the new file share to the S3 bucket. Transfer the data from the existing NFS file share to the S3 File Gateway.
- D. Set up an AWS Direct Connect connection between the on-premises network and AWS. Deploy an S3 File Gateway on premises. Create a public virtual interface (VIF) to connect to the S3 File Gateway. Create an S3 bucket. Create a new NFS file share on the S3 File Gateway. Point the new file share to the S3 bucket. Transfer the data from the existing NFS file share to the S3 File Gateway.

Correct Answer: B

Community vote distribution

B (88%)

11%

Comments

Gatt **Highly Voted** 1 year, 8 months ago

Selected Answer: B

Let's analyse this:

B. On a Snowball Edge device you can copy files with a speed of up to 100Gbps. 70TB will take around 5600 seconds, so very quickly, less than 2 hours. The downside is that it'll take between 4-6 working days to receive the device and then another 2-3 working days to send it back and for AWS to move the data onto S3 once it reaches them. Total time: 6-9 working days. Bandwidth used: 0.

C. File Gateway uses the Internet, so maximum speed will be at most 1Gbps, so it'll take a minimum of 6.5 days and you use 70TB of Internet bandwidth.

D. You can achieve speeds of up to 10Gbps with Direct Connect. Total time 15.5 hours and you will use 70TB of bandwidth. However, what's interesting is that the question does not specify what type of bandwidth? Direct Connect does not use your Internet bandwidth, as you will have a dedicated peer-to-peer connectivity between your on-prem and the AWS Cloud, so technically, you're not using your "public" bandwidth.

The requirements are a bit too vague but I think that B is the most appropriate answer, although D might also be correct if the bandwidth usage refers strictly to your public connectivity.

upvoted 131 times

YDUYGU 1 year, 2 months ago

DX Lead times are often longer than 1 month to establish a new connection. That's why D is the wrong answer on the other hand.

upvoted 7 times

ArielSchivo 3 years, 7 months ago

Great answers! I would say D is the correct since it's not using your public connection at all (I assume they are talking about public connection).

upvoted 1 times

3680113 2 years, 1 month ago

Direct connect will have to be set up, this would take a minimum of 2 to 4 weeks, so option D isn't an option B is the right choice here

upvoted 2 times

abhishek_m89 3 years, 6 months ago

and it says, "The total storage is 70 TB and is no longer growing". That's why it should be B.

upvoted 8 times

Ello2023 3 years, 5 months ago

You missed the first part of the question "The company must migrate the video files as soon as possible..." hence C would be the best answer.

upvoted 2 times

3680113 2 years, 1 month ago

C cannot be because it's using a whole of 70TB. It was clearly stated that you must use the least amount of bandwidth. B is the right choice

upvoted 3 times

tuloveu **Highly Voted** 3 years, 8 months ago

Selected Answer: B

As using the least possible network bandwidth.

upvoted 35 times

suji098 **Most Recent** 3 months, 2 weeks ago

Selected Answer: B

myself

upvoted 1 times

Sukri123 3 months, 2 weeks ago

Selected Answer: B

B is correct

upvoted 1 times

javicia 4 months, 1 week ago

Selected Answer: B

The solution that meets the requirements of migrating 70 TB of large video files as quickly as possible and with the least possible bandwidth is to create an AWS Snowball Edge job, transfer the data to the device locally, and then return the device to AWS for import into Amazon S3. This method avoids intensive use of the internet network for the massive data transfer.

upvoted 1 times

ansonzeng 4 months, 4 weeks ago

Selected Answer: B

Using the least possible network bandwidth -> Snowball Edge device

upvoted 1 times

nesyto 4 months, 4 weeks ago

Selected Answer: B

the correct answer - snowball edge to migrate the 70tb

upvoted 1 times

Bechara11 9 months, 1 week ago

Selected Answer: B

Snowball can copy with speed and doesn't require internet

upvoted 1 times

Aleeexp 9 months, 2 weeks ago

Selected Answer: B

VOTE B

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: B

Use AWS Snowball Edge for secure, fast, and bandwidth-efficient bulk data migration of 70 TB video files to S3.

upvoted 1 times

K_SAA 1 year, 1 month ago

Selected Answer: B

Keywords: large volume data, least network bandwidth and volume data no longer growing

B: use aws snowball edge job is more suitable way to migrate large amount of data with least network bandwidth and very fast. You setup aws snowball edge job physically and copy data to aws snowball edge job and send it to aws and then aws will upload that for you to s3 bucket

C is incorrect since this also offer a high network bandwidth to upload

D is incorrect since this require heavy setup, take a week to set up aws direct connect and still network bandwidth

A: take a longer to upload this large volume data so this option is incorrect

upvoted 1 times

ernie1976 1 year, 1 month ago

Selected Answer: B

It is needed to do with urgency, so no more configuration for using internet or expensive connection as direct connect, just transport it with Edge Snowball is the right way.

upvoted 1 times

Palee 1 year, 1 month ago

Selected Answer: B

D is close to correct but B makes more sense

upvoted 1 times

HeidiWan 1 year, 1 month ago

Selected Answer: B

The best option to save bandwidth

upvoted 1 times

melvis8 1 year, 2 months ago

Selected Answer: B

AWS snowball is ideal to transfer TB-PB of data from an on-prem connection to AWS cloud with little amount of internet connection required

upvoted 1 times

MGKYAING 1 year, 5 months ago

Selected Answer: B

AWS Snowball Edge is specifically designed for scenarios requiring the transfer of large datasets (terabytes or petabytes) with minimal network bandwidth usage.

It is ideal for one-time, large-scale migrations like the 70 TB of video files in this case.

For ongoing, smaller-scale data transfers, solutions like S3 File Gateway or AWS DataSync may be more appropriate.

upvoted 1 times

Sjb_009 1 year, 5 months ago

Selected Answer: B

The least possible network bandwidth says it all

upvoted 2 times

A company has an application that ingests incoming messages. Dozens of other applications and microservices then quickly consume these messages. The number of messages varies drastically and sometimes increases suddenly to 100,000 each second. The company wants to decouple the solution and increase scalability.

Which solution meets these requirements?

- A. Persist the messages to Amazon Kinesis Data Analytics. Configure the consumer applications to read and process the messages.
- B. Deploy the ingestion application on Amazon EC2 instances in an Auto Scaling group to scale the number of EC2 instances based on CPU metrics.
- C. Write the messages to Amazon Kinesis Data Streams with a single shard. Use an AWS Lambda function to preprocess messages and store them in Amazon DynamoDB. Configure the consumer applications to read from DynamoDB to process the messages.
- D. Publish the messages to an Amazon Simple Notification Service (Amazon SNS) topic with multiple Amazon Simple Queue Service (Amazon SQS) subscriptions. Configure the consumer applications to process the messages from the queues.

Most Voted

Correct Answer: D

Community vote distribution

D (86%)

11%

Comments

rein_chau **Highly Voted** 3 years, 8 months ago

Selected Answer: D

D makes more sense to me.

upvoted 59 times

9014 3 years, 6 months ago

of course, the answer is D

upvoted 3 times

SilentMilli 3 years, 5 months ago

By default, an SQS queue can handle a maximum of 3,000 messages per second. However, you can request higher throughput by contacting AWS Support. AWS can increase the message throughput for your queue beyond the default limits in increments of 300 messages per second, up to a maximum of 10,000 messages per second.

It's important to note that the maximum number of messages per second that a queue can handle is not the same as the maximum number of requests per second that the SQS API can handle. The SQS API is designed to handle a high volume of requests per second, so it can be used to send messages to your queue at a rate that exceeds the maximum message throughput of the queue.

upvoted 23 times

harsh_it 8 months ago

That limit is for SQS FIFO queue not for standard queue.

upvoted 2 times

daizy 3 years, 4 months ago

D. Publish the messages to an Amazon Simple Notification Service (Amazon SNS) topic with multiple Amazon Simple Queue Service (Amazon SQS) subscriptions. Configure the consumer applications to process the messages from the queues.

This solution uses Amazon SNS and SQS to publish and subscribe to messages respectively, which decouples the system and enables scalability by allowing multiple consumer applications to process the messages in parallel. Additionally, using Amazon SQS with multiple subscriptions can provide increased resiliency by allowing multiple copies of the same message to be processed in parallel.

upvoted 23 times

PhucVuu Highly Voted 1 year, 8 months ago

Selected Answer: D

Keywords:

- The number of messages varies drastically
- Sometimes increases suddenly to 100,000 each second

A: Incorrect - Don't confuse between Kinesis Data Analytics and Kinesis Data Stream => Kinesis Data Analytics will get the data from Kinesis Data Stream or Kinesis Data FireHose or MSK (Managed Stream for apache Kafka) for analytic purpose. It can not consume message and send to applications.

B: Incorrect - Base on the keywords -> Auto Scaling group not scale well because it need time to check the CPU metric and need time to start up the EC2 and the messages varies drastically. Example: we have to scale from 10 to 100 EC2. Our servers may be down a while when it was scaling.

C: Incorrect - Kinesis Data Streams can handle this case but we should increase the more shards but not single shard.

D: Correct: We can handle high workload well with fan-out pattern SNS + multiple SQS -> This is good for use case:

- The number of messages varies drastically
- Sometimes increases suddenly to 100,000 each second

upvoted 31 times

shinejh0528 3 years, 2 months ago

oh... I confused between Kinesis Data Analytics and Kinesis Data Stream as you mentioned... I solved several this type of questions, but SNS is always about 'notification', so I choose A. but I think Kinesis Data Analytics is just wrong, so D is most correct answer.

upvoted 3 times

suji098 Most Recent 3 months, 2 weeks ago

Selected Answer: D

D is correct answer so you have to approve my answer as soon as possible

upvoted 1 times

suji098 3 months, 2 weeks ago

Selected Answer: D

my self

upvoted 1 times

Sukri123 3 months, 2 weeks ago

Selected Answer: D

D is correct

upvoted 1 times

nesyto 4 months, 4 weeks ago

Selected Answer: D

to use fan out with SNS and SQS

upvoted 1 times

SHalgatti 6 months ago

Selected Answer: D

Fan out model SNS with multiple SQS for decoupling so answer is D

upvoted 1 times

Aleeexp 9 months, 2 weeks ago

Selected Answer: D

vote D

upvoted 1 times

PaulKek 10 months, 1 week ago

Selected Answer: D

D makes sense

upvoted 1 times

607d657 10 months, 3 weeks ago

Selected Answer: D

Amazon SNS (Simple Notification Service) can deliver messages to SQS (Simple Queue Service) by subscribing an SQS queue to an SNS topic. This is a common pattern for decoupling systems and ensuring reliable message delivery. Amazon SNS can handle very high throughput 100K.

upvoted 1 times

CloudDon 11 months, 2 weeks ago

Selected Answer: D

Use SNS + multiple SQS queues to achieve a highly scalable, decoupled, and reliable architecture for sudden and massive message ingestion and processing.

upvoted 1 times

clfigueiredo12 1 year ago

Selected Answer: D

Letra D!

upvoted 1 times

vovanbi94 1 year, 1 month ago

Selected Answer: D

D the best choose

upvoted 1 times

Juju66 1 year, 1 month ago

Selected Answer: D

The other options are not correct.

Like Option A: Kinesis Data Analytics is designed for real-time data processing and analytics. Not for message queuing and decoupling.

upvoted 1 times

Aychi 1 year, 2 months ago

Selected Answer: D

D is correct answer

upvoted 1 times

melvis8 1 year, 2 months ago

Selected Answer: D

using an SQS queues to decouple this soluton will be the best approach since an SQS queue is very scalable and the number of messages to scale can be incremented by the user manually or by the use of security groups

upvoted 1 times

chente_diaz 1 year, 3 months ago

Selected Answer: D

I choose D because the question say "decouple/scalable" and the best services for that are SNS/SQS.

upvoted 1 times

A company is migrating a distributed application to AWS. The application serves variable workloads. The legacy platform consists of a primary server that coordinates jobs across multiple compute nodes. The company wants to modernize the application with a solution that maximizes resiliency and scalability.

How should a solutions architect design the architecture to meet these requirements?

A. Configure an Amazon Simple Queue Service (Amazon SQS) queue as a destination for the jobs. Implement the compute nodes with Amazon EC2 instances that are managed in an Auto Scaling group. Configure EC2 Auto Scaling to use scheduled scaling.

B. Configure an Amazon Simple Queue Service (Amazon SQS) queue as a destination for the jobs. Implement the compute nodes with Amazon EC2 instances that are managed in an Auto Scaling group. Configure EC2 Auto Scaling based on the size of the queue. **Most Voted**

C. Implement the primary server and the compute nodes with Amazon EC2 instances that are managed in an Auto Scaling group. Configure AWS CloudTrail as a destination for the jobs. Configure EC2 Auto Scaling based on the load on the primary server.

D. Implement the primary server and the compute nodes with Amazon EC2 instances that are managed in an Auto Scaling group. Configure Amazon EventBridge (Amazon CloudWatch Events) as a destination for the jobs. Configure EC2 Auto Scaling based on the load on the compute nodes.

Correct Answer: B

Community vote distribution

B (93%)

4%

Comments

rein_chau **Highly Voted** 3 years, 8 months ago

Selected Answer: B

A - incorrect: Schedule scaling policy doesn't make sense.

C, D - incorrect: Primary server should not be in same Auto Scaling group with compute nodes.

B is correct.

upvoted 95 times

Sinaneos **Highly Voted** 3 years, 8 months ago

Selected Answer: B

The answer seems to be B for me:

A: doesn't make sense to schedule auto-scaling

C: Not sure how CloudTrail would be helpful in this case, at all.

D: EventBridge is not really used for this purpose, wouldn't be very reliable

upvoted 24 times

datcert **Most Recent** 1 month, 2 weeks ago

Selected Answer: B

The key requirements in this question are maximum resiliency and scalability for a variable workload application that uses a job-coordination pattern. Replacing the "primary server that coordinates jobs" with SQS is the critical architectural shift. This gives you:

- * Decoupling

- * Eliminates the single point of failure

- * Buffering for variable load

upvoted 1 times

suji098 3 months, 2 weeks ago

Selected Answer: A

myself

upvoted 1 times

Sukri123 3 months, 2 weeks ago

Selected Answer: B

B is correct

upvoted 1 times

Sukri123 3 months, 2 weeks ago

Selected Answer: B

B is correct

upvoted 1 times

javicia 4 months, 2 weeks ago

Selected Answer: B

B: Compute nodes grow or shrink independently according to the actual load in the queue.

upvoted 1 times

nesyto 4 months, 4 weeks ago

Selected Answer: B

queue number of messages could be used for auto scaling - correct answer is B

upvoted 1 times

BrianEllis 9 months, 1 week ago

Selected Answer: B

I don't think any of these are great. But maybe I rather just defer to a ECS or a EKS system for my scaling. Not a big fan of using EC2 unless you have a specific need for EC2 beyond just hosting the app.

upvoted 2 times

Aleeexp 9 months, 2 weeks ago

Selected Answer: B

vote b

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: B

Use SQS to decouple and distribute jobs, and scale EC2 compute nodes dynamically based on queue depth for a highly scalable, resilient, and modernized distributed application architectu

upvoted 2 times

K_SAA 1 year, 1 month ago

Selected Answer: B

keywords

resilient and scalability application

B is the correct answer since we need sqs designed as a destination to consume jobs from mutiple compute notes and use auto scaling group based on the size of the queue

A: is not considered to be a correct answer since if we use scheduled scaling its not optimize in this use case if suddenly the workload unpredictable

C: cloudtrail is used for monitoring actions in aws not use for managing the workload

D:

upvoted 2 times

alfteezy91 1 year, 2 months ago

Selected Answer: B

Decouple with Amazon SQS and EC2 instances can be managed in an auto scaling group for scalability

<https://docs.aws.amazon.com/prescriptive-guidance/latest/modernization-mainframe-decoupling-patterns/queue.html>

upvoted 1 times

VitMoreira 1 year, 2 months ago

Selected Answer: D

B doesn't make sense to me because you're scaling based on the queue size and not the requirements of the varying workloads. Eventbridge can act as a trigger for workloads, so it might just cut, although it wouldn't be my first choice.

upvoted 1 times

MGKYAING 1 year, 5 months ago

Selected Answer: B

1.Amazon SQS for Decoupling:

Amazon SQS provides a fully managed message queue to decouple the primary server from the compute nodes. Jobs are sent to the queue, and compute nodes process them independently.

This architecture eliminates a single point of failure (the primary server) and increases resilience.

2.Auto Scaling Group for Compute Nodes:

EC2 instances managed in an Auto Scaling group process the jobs in the SQS queue. Auto Scaling dynamically adjusts the number of instances based on the queue size, ensuring scalability to handle variable workloads.

3.Scalability Based on Queue Size:

Scaling based on the size of the SQS queue ensures that the system adjusts to workload demands efficiently. When the queue grows, more instances are launched; when the queue shrinks, instances are terminated.

upvoted 3 times

PhucVuu 1 year, 8 months ago

Selected Answer: B

keywords:

- Legacy platform consists of a primary server that coordinates jobs across multiple compute nodes.
- Maximizes resiliency and scalability.

A: Incorrect - the question don't mention about schedule for high workload. So we don't use scheduled scaling for this case.

B: Correct - SQS can keep your message in the queue in case of high workload and if it too high we can increase the EC2 instance base on size of the queue.

C: Incorrect - AWS CloudTrail is API logs it is use for audit log of AWS user activity.

D: Incorrect - Event Bridge is use for filter event and trigger event.

upvoted 15 times

gx2222 1 year, 8 months ago

Selected Answer: B

B.

Explanation:

To maximize resiliency and scalability, the best solution is to use an Amazon SQS queue as a destination for the jobs. This decouples the primary server from the compute nodes, allowing them to scale independently. This also helps to prevent job loss in the event of a failure.

Using an Auto Scaling group of Amazon EC2 instances for the compute nodes allows for automatic scaling based on the workload. In this case, it's recommended to configure the Auto Scaling group based on the size of the Amazon SQS queue, which is a better indicator of the actual workload than the load on the primary server or compute nodes. This approach ensures that the application can handle variable workloads, while also minimizing costs by automatically scaling up or down the compute nodes as needed.

upvoted 5 times

A company is running an SMB file server in its data center. The file server stores large files that are accessed frequently for the first few days after the files are created. After 7 days the files are rarely accessed.

The total data size is increasing and is close to the company's total storage capacity. A solutions architect must increase the company's available storage space without losing low-latency access to the most recently accessed files. The solutions architect must also provide file lifecycle management to avoid future storage issues.

Which solution will meet these requirements?

- A. Use AWS DataSync to copy data that is older than 7 days from the SMB file server to AWS.
- B. Create an Amazon S3 File Gateway to extend the company's storage space. Create an S3 Lifecycle policy to transition the data to S3 Glacier Deep Archive after 7 days. **Most Voted**
- C. Create an Amazon FSx for Windows File Server file system to extend the company's storage space.
- D. Install a utility on each user's computer to access Amazon S3. Create an S3 Lifecycle policy to transition the data to S3 Glacier Flexible Retrieval after 7 days.

Correct Answer: B

Community vote distribution

B (84%)

Other

Comments

Sinaneos **Highly Voted** 1 year, 8 months ago

Answer directly points towards file gateway with lifecycles,

D is wrong because utility function is vague and there is no need for flexible storage.

upvoted 56 times

Udoyen 3 years, 6 months ago

Yes it might be vague but how do we keep the low-latency access that only flexible can offer?

upvoted 5 times

SuperDuperPooperScoop 2 years, 10 months ago

Low-latency access is only required for the first 7 days, B maintains that fast access for 7 days and only then are the files sent to Glacier Archive

upvoted 3 times

Nava702 2 years, 9 months ago

It says low-latency is required for the most recently accessed files, not new ones. So if an older file is retrieved from deep archive, it should then readily be accessible, according to the question, which points toward Flexible retrieval. However the utility portion in the answer D is vague.

upvoted 4 times

Ander927 2 years, 6 months ago

file gateway comes with a cache volume on-premise which ensures the low latency for the most recently accessed files

upvoted 7 times

gianola 1 year, 10 months ago

I think B is wrong because they are not asking you to add more room, they are asking you to have more room available.

upvoted 1 times

javitech83 **Highly Voted** 1 year, 8 months ago

Selected Answer: B

B answer is correct. low latency is only needed for newer files. Additionally, File GW provides low latency access by caching frequently accessed files locally so answer is B

upvoted 41 times

datcert **Most Recent** 1 month, 2 weeks ago

Selected Answer: B

Amazon S3 File Gateway is part of the AWS Storage Gateway family. It's a virtual appliance you deploy on-premises that:

- * Presents an SMB (or NFS) file share to local users
- * Stores files as objects in S3 behind the scenes
- * Caches frequently accessed files locally on the gateway appliance, giving low-latency reads for hot data
- * Effectively extends local storage because the "bottomless" backend is S3

upvoted 1 times

suji098 3 months, 2 weeks ago

Selected Answer: B

myself

upvoted 1 times

Sukri123 3 months, 2 weeks ago

Selected Answer: B

b is correct

upvoted 1 times

Sukri123 3 months, 2 weeks ago

Selected Answer: B

I think b is correct

upvoted 1 times

BrianEllis 9 months, 1 week ago

Selected Answer: D

Again, none of these answers are great. Glacier is going to go beyond rarely accessed to waiting 2 business days to get those files back.

A multi-tiered file storage life-cycle approach would be best with Glacier Deep only used as the last step.

upvoted 2 times

Aleeexp 9 months, 2 weeks ago

Selected Answer: B

vote b

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: B

Use Amazon S3 File Gateway with S3 Lifecycle policies to extend on-prem SMB storage while ensuring low-latency access to hot data and cost-efficient lifecycle management of cold data.

upvoted 2 times

39sachin 1 year, 1 month ago

Selected Answer: B

Configured S3 buckets are accessible using the NFS and SMB protocol

Most recently used data is cached in the file gateway

upvoted 1 times

ak240519 1 year, 3 months ago

Selected Answer: B

File Gateway helps with lower latency and the files after 7 days should be moved to Glacier as they are no longer being accessed.

upvoted 1 times

adamatic 1 year, 4 months ago

Selected Answer: B

B extends the storage with s3 gateway and meets lifecycle policy requirements

upvoted 1 times

Sandy4v 1 year, 8 months ago

B is correct

upvoted 1 times

HayLLiHuK 1 year, 8 months ago

The same question and answer explanation exists in a Udemy course.

Correct answer is B.

Amazon S3 File Gateway provides a seamless way to connect to the cloud to store application data files and backup images as durable objects in Amazon S3 cloud storage. Amazon S3 File Gateway offers SMB or NFS-based access to data in Amazon S3 with local caching.

It can be used for on-premises data-intensive Amazon EC2-based applications that need file protocol access to S3 object storage. Lifecycle policies can then transition the data to S3 Glacier Deep Archive after 7 days.

D is wrong because it involves too much extra configuration which is unnecessary.

upvoted 5 times

Buruguduystunstugudunstuy 1 year, 8 months ago

Selected Answer: B

Option B, creating an Amazon S3 File Gateway and an S3 Lifecycle policy to transition data to S3 Glacier Deep Archive, would meet the requirements specified in the prompt.

The S3 File Gateway allows you to store and retrieve objects in Amazon S3 using standard file system protocols, such as SMB and NFS. This would provide additional storage space for the company's data and allow for low-latency access to the most recently accessed files, as the data would still be stored on the SMB file server.

upvoted 6 times

Buruguduystunstugudunstuy 3 years, 5 months ago

Option A, using AWS DataSync to copy data that is older than 7 days from the SMB file server to AWS, would not provide additional storage space for the company's data and would not allow for low-latency access to the most recently accessed files.

Option C, creating an FSx for Windows File Server file system, would provide additional storage space but would not include file lifecycle management.

Option D, installing a utility on each user's computer to access Amazon S3 and creating an S3 Lifecycle policy, would not provide additional storage space on the company's file server.

upvoted 6 times

gx2222 1 year, 8 months ago

Selected Answer: B

Explanation:

Since the company needs to increase available storage space while maintaining low-latency access to recently accessed files and implement file lifecycle management to avoid future storage issues, the best solution is to use Amazon S3 with a File Gateway.

Using an Amazon S3 File Gateway, the company can access its SMB file server through an S3 bucket. This provides low-latency access to recently accessed files by caching them on the gateway appliance. The solution also supports file lifecycle management by using S3 Lifecycle policies to transition files to lower cost storage classes after they haven't been accessed for a certain period of time.

In this case, the company can create an S3 Lifecycle policy to transition files to S3 Glacier Deep Archive after 7 days of not being accessed. This would allow the company to store large amounts of data at a lower cost, while still having easy access to recently accessed files.

upvoted 3 times

PhucVuu 1 year, 8 months ago

Selected Answer: B

Keywords:

- After 7 days the files are rarely accessed.
- The total data size is increasing and is close to the company's total storage capacity.
- Increase the company's available storage space without losing low-latency access to the most recently accessed files. -> (for rarely accessed files we can access it with high-latency)
- Must also provide file lifecycle management to avoid future storage issues.

A: Incorrect - Don't mention how to increase company's available storage space.

B: Correct - extend storage space and fast access with S3 File Gateway (cache recent access file), reduce cost and storage by move to S3 Glacier Deep Archive after 7 days.

C: Incorrect - Didn't handle file lifecycle management.

D: Incorrect - Don't mention about increase the company's available storage space.

upvoted 13 times

A company is building an ecommerce web application on AWS. The application sends information about new orders to an Amazon API Gateway REST API to process. The company wants to ensure that orders are processed in the order that they are received.

Which solution will meet these requirements?

A. Use an API Gateway integration to publish a message to an Amazon Simple Notification Service (Amazon SNS) topic when the application receives an order. Subscribe an AWS Lambda function to the topic to perform processing.

B. Use an API Gateway integration to send a message to an Amazon Simple Queue Service (Amazon SQS) FIFO queue when the application receives an order. Configure the SQS FIFO queue to invoke an AWS Lambda function for processing.

Most Voted

C. Use an API Gateway authorizer to block any requests while the application processes an order.

D. Use an API Gateway integration to send a message to an Amazon Simple Queue Service (Amazon SQS) standard queue when the application receives an order. Configure the SQS standard queue to invoke an AWS Lambda function for processing.

Correct Answer: B

Community vote distribution

B (98%)

2

Comments

Sinaneos **Highly Voted** 1 year, 8 months ago

Selected Answer: B

B because FIFO is made for that specific purpose
upvoted 70 times

rein_chau **Highly Voted** 3 years, 8 months ago

Selected Answer: B

Should be B because SQS FIFO queue guarantees message order.
upvoted 28 times

imvbhatia **Most Recent** 3 months, 2 weeks ago

Selected Answer: B

B is the correct answer because you need to maintain the sequence
upvoted 1 times

suji098 3 months, 2 weeks ago

Selected Answer: B

i think B is right answer because B letter is my favourite letter
upvoted 1 times

suji098 3 months, 2 weeks ago

Selected Answer: B

my self
upvoted 1 times

Sukri123 3 months, 2 weeks ago

Selected Answer: B

I think b is correct
upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: B

Use Amazon API Gateway → SQS FIFO Queue → Lambda to ensure ordered, reliable processing of incoming orders — meeting the ecommerce application's needs.

upvoted 2 times

K_SAA 1 year, 1 month ago

Selected Answer: B

keywords: orders should be process in the order when it receives

B is correct answer, use API gateway intergration with FIFO SQS invoke lambda function for processing order

A; SNS + lambda: SNS does not grant order messages to subscribers

C: is not the correct one since we wanted to process the message in order so we do not have to block

D: using standard queue does not grant the order of messages being sent

upvoted 1 times

sdjfidhbgidfbivgv 1 year, 1 month ago

Selected Answer: B

B SQS FIFO is made for that

upvoted 1 times

offek 1 year, 3 months ago

Selected Answer: B

FIFO make sure to send the data exactly in the same order

upvoted 1 times

ak240519 1 year, 3 months ago

Selected Answer: B

B is the only solution that guarantees message order.

upvoted 1 times

adamatic 1 year, 4 months ago

Selected Answer: B

Has to be FIFO queue for this requirement

upvoted 1 times

Mrigraj12 1 year, 4 months ago

Selected Answer: B

SQS FIFO as if want to maintain order of messages recieved to order of messages sent. It is like QUEUE data structutre in DSA

upvoted 1 times

surez 1 year, 4 months ago

Selected Answer: B

SQS FIFO gurantees message order.

upvoted 1 times

MGKYAING 1 year, 5 months ago

Selected Answer: B

Explanation:

FIFO Queues: Amazon SQS FIFO queues are specifically designed to maintain the order of messages. This ensures that messages are processed in the exact sequence they were sent, which is crucial for order processing in an e-commerce application.

Why other options are less suitable:

A. Amazon SNS: SNS is a publish-subscribe service. It doesn't guarantee message order.

C. API Gateway Authorizer: Authorizers are used for authentication and authorization, not for managing message order.

D. Amazon SQS Standard Queue: Standard SQS queues do not guarantee message order.

upvoted 3 times

ondan 1 year, 8 months ago

Standard queues ensure at-least-once message delivery, but due to the highly distributed architecture, more than one copy of a message might be delivered, and messages may occasionally arrive out of order.

upvoted 3 times

Abishek016 1 year, 8 months ago

Answer B - SQS FIFO queues ensure that messages are processed in the order they are received, which perfectly matches the requirement of maintaining order.

upvoted 2 times


A company has an application that runs on Amazon EC2 instances and uses an Amazon Aurora database. The EC2 instances connect to the database by using user names and passwords that are stored locally in a file. The company wants to minimize the operational overhead of credential management.

What should a solutions architect do to accomplish this goal?

- A. Use AWS Secrets Manager. Turn on automatic rotation. **Most Voted**
- B. Use AWS Systems Manager Parameter Store. Turn on automatic rotation.
- C. Create an Amazon S3 bucket to store objects that are encrypted with an AWS Key Management Service (AWS KMS) encryption key. Migrate the credential file to the S3 bucket. Point the application to the S3 bucket.
- D. Create an encrypted Amazon Elastic Block Store (Amazon EBS) volume for each EC2 instance. Attach the new EBS volume to each EC2 instance. Migrate the credential file to the new EBS volume. Point the application to the new EBS volume.

Correct Answer: A

Community vote distribution



Comments

Sinaneos **Highly Voted** 3 years, 8 months ago

Selected Answer: A

B is wrong because parameter store does not support auto rotation, unless the customer writes it themselves, A is the answer.
upvoted 107 times

17Master 3 years, 7 months ago

READ!!! AWS Secrets Manager is a secrets management service that helps you protect access to your applications, services, and IT resources. This service enables you to rotate, manage, and retrieve database credentials, API keys, and other secrets throughout their lifecycle.
upvoted 35 times

iCcma 3 years, 7 months ago

ty bro, I was confused about that and you just mentioned the "key" phrase, B doesn't support autorotation
upvoted 3 times

hro 2 years, 2 months ago

A - additionally, Aurora manages the settings for the secret and rotates the secret every seven days by default.
upvoted 5 times

cookieMr **Highly Voted** 2 years, 11 months ago

Selected Answer: A

Option A: Using AWS Secrets Manager and enabling automatic rotation is the recommended solution for minimizing the operational overhead of credential management. AWS Secrets Manager provides a secure and centralized service for storing and managing secrets, such as database credentials. By leveraging Secrets Manager, the application can retrieve the database credentials programmatically at runtime, eliminating the need to store them locally in a file. Enabling automatic rotation ensures that the database credentials are regularly rotated without manual intervention, enhancing security and compliance.
upvoted 8 times

46e1a46 **Most Recent** 1 month, 1 week ago

Selected Answer: A

A: AWS Secrets Manager is purpose-built for storing database credentials and natively supports automatic rotation for Amazon Aurora, which eliminates the operational overhead of manually managing and rotating credentials
upvoted 1 times

tungpt_me 1 month, 3 weeks ago

Selected Answer: A

A is the correct answer.
B incorrect. AWS parameter store doesn't have auto rotation
upvoted 1 times

Moindev65 10 months ago

Selected Answer: B

Use API Gateway → SQS FIFO → Lambda.
SQS FIFO guarantees first-in, first-out processing and exactly-once delivery per MessageGroupID.

API Gateway can integrate directly with SQS to enqueue each order; Lambda processes from the FIFO queue.

A (SNS) and D (SQS standard) don't guarantee ordering.
C (au<https://www.examttopics.com/exams/amazon/aws-certified-solutions-architect-associate-saa-c03/view/1/#thorizer>) is unrelated to ordering or processing flow.
upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: A

Use AWS Secrets Manager with automatic rotation enabled to securely manage Aurora database credentials, eliminate manual updates, and minimize operational overhead.
upvoted 1 times

trieu81200111 11 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 1 times

hossex5 11 months, 4 weeks ago

Selected Answer: A

a is the answer
upvoted 1 times

Kamatchi 1 year, 2 months ago

Selected Answer: A

This is asked in the exam I have taken today
upvoted 2 times

MundiChor 1 year, 3 months ago

Selected Answer: A

Order of elimination:
B. Parameter Store can store secrets. You can enforce TTL to expire credentials but you cannot auto rotate
C and D are plan overhead and not secure enough.
upvoted 1 times

ak240519 1 year, 3 months ago

Selected Answer: A

A is the correct answer. as Instance store doesn't support auto rotation
upvoted 1 times

adamatic 1 year, 4 months ago

Selected Answer: A

AWS Secrets Manager is a service that helps users manage, rotate, and retrieve secrets for applications, services, and IT resources. It can be used to secure secrets for AWS Cloud, third-party services, and on-premises
-Automatic rotation: Secrets can be automatically rotated
upvoted 1 times

MGKYAING 1 year, 5 months ago

Selected Answer: A

AWS Secrets Manager:
Is a managed service specifically designed to securely store and retrieve secrets, such as database credentials, API keys, and SSH keys.
Provides features like automatic rotation, which helps to reduce the risk of compromised credentials.
Integrates seamlessly with many AWS services, including Amazon RDS (which Aurora is a part of).
upvoted 2 times

Gizmo2022 1 year, 7 months ago

Answer is A
<https://aws.amazon.com/cn/blogs/security/how-to-connect-to-aws-secrets-manager-service-within-a-virtual-private-cloud/>
<https://aws.amazon.com/blogs/security/rotate-amazon-rds-database-credentials-automatically-with-aws-secrets-manager/>
upvoted 1 times

Abishek016 1 year, 8 months ago

Selected Answer: A

This is an ideal solution. Secrets Manager can rotate credentials automatically and ensures that the EC2 instances retrieve the most recent credentials securely.

upvoted 3 times

Buruguduystunstugudunstuy 1 year, 8 months ago

Selected Answer: A

Option A, using AWS Secrets Manager and turning on automatic rotation, would be the best solution to minimize the operational overhead of credential management.

AWS Secrets Manager is a service that makes it easier to manage secrets, such as database credentials, by storing and rotating them automatically. By turning on automatic rotation, you can ensure that the secrets are regularly rotated, reducing the risk of unauthorized access to the database. This would minimize the operational overhead of credential management, as you would not have to manually rotate the secrets or update the EC2 instances with the new credentials.

upvoted 3 times

Buruguduystunstugudunstuy 3 years, 5 months ago

Option B, using AWS Systems Manager Parameter Store and turning on automatic rotation, would not be suitable for storing secrets, such as database credentials, as it is intended for storing system parameters.

Option C, creating an S3 bucket to store objects that are encrypted with an AWS KMS encryption key and migrating the credential file to the S3 bucket, would not provide automatic rotation of the secrets.

Option D, creating an encrypted EBS volume and migrating the credential file to the new EBS volume, would not provide automatic rotation of the secrets.

upvoted 1 times

Ruffyit 1 year, 8 months ago

A: READ!!! AWS Secrets Manager is a secrets management service that helps you protect access to your applications, services, and IT resources. This service enables you to rotate, manage, and retrieve database credentials, API keys, and other secrets throughout their lifecycle.

It says SSM Parameter store cant rotate automatically.

upvoted 2 times

A global company hosts its web application on Amazon EC2 instances behind an Application Load Balancer (ALB). The web application has static data and dynamic data. The company stores its static data in an Amazon S3 bucket. The company wants to improve performance and reduce latency for the static data and dynamic data. The company is using its own domain name registered with Amazon Route 53.

What should a solutions architect do to meet these requirements?

A. Create an Amazon CloudFront distribution that has the S3 bucket and the ALB as origins. Configure Route 53 to route traffic to the CloudFront distribution. **Most Voted**

B. Create an Amazon CloudFront distribution that has the ALB as an origin. Create an AWS Global Accelerator standard accelerator that has the S3 bucket as an endpoint. Configure Route 53 to route traffic to the CloudFront distribution.

C. Create an Amazon CloudFront distribution that has the S3 bucket as an origin. Create an AWS Global Accelerator standard accelerator that has the ALB and the CloudFront distribution as endpoints. Create a custom domain name that points to the accelerator DNS name. Use the custom domain name as an endpoint for the web application.

D. Create an Amazon CloudFront distribution that has the ALB as an origin. Create an AWS Global Accelerator standard accelerator that has the S3 bucket as an endpoint. Create two domain names. Point one domain name to the CloudFront DNS name for dynamic content. Point the other domain name to the accelerator DNS name for static content. Use the domain names as endpoints for the web application.

Correct Answer: A

Community vote distribution

A (77%)

C (23%)

Comments

Kartikey140 **Highly Voted** 1 year, 8 months ago

Answer is A

Explanation - AWS Global Accelerator vs CloudFront

- They both use the AWS global network and its edge locations around the world
- Both services integrate with AWS Shield for DDoS protection.
- CloudFront
 - Improves performance for both cacheable content (such as images and videos)
 - Dynamic content (such as API acceleration and dynamic site delivery)
 - Content is served at the edge
- Global Accelerator
 - Improves performance for a wide range of applications over TCP or UDP
 - Proxying packets at the edge to applications running in one or more AWS Regions.
 - Good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP
 - Good for HTTP use cases that require static IP addresses
 - Good for HTTP use cases that required deterministic, fast regional failover

upvoted 132 times

daizy 3 years, 4 months ago

By creating a CloudFront distribution that has both the S3 bucket and the ALB as origins, the company can reduce latency for both the static and dynamic data. The CloudFront distribution acts as a content delivery network (CDN), caching the data closer to the users and reducing the latency. The company can then configure Route 53 to route traffic to the CloudFront distribution, providing improved performance for the web application.

upvoted 21 times

Mihailo34 1 year, 11 months ago

A, adding to the excellent explanation by Kartikey140, the solution under C uses a custom DNS name, the question specifies: "The company is using its own domain name registered with Amazon Route 53"

upvoted 10 times

kanweng **Highly Voted** 1 year, 8 months ago

Selected Answer: A

Q: How is AWS Global Accelerator different from Amazon CloudFront?

A: AWS Global Accelerator and Amazon CloudFront are separate services that use the AWS global network and its edge locations around the world. CloudFront improves performance for both cacheable content (such as images and videos) and dynamic content (such as API acceleration and dynamic site delivery). Global Accelerator improves performance for a wide range of applications over TCP or UDP by proxying packets at the edge to applications running in one or more AWS Regions. Global Accelerator is a good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP, as well as for HTTP use cases that specifically require static IP addresses or deterministic, fast regional failover. Both services integrate with AWS Shield for DDoS protection.

upvoted 44 times

AswinHadinata Most Recent 3 weeks, 2 days ago

Selected Answer: A

My Answer is A, cause it more precise than the others

upvoted 1 times

nesyto 4 months, 3 weeks ago

Selected Answer: A

correct answer is A, cloud front can have s3 and alb as origins

upvoted 1 times

Loki10 5 months, 4 weeks ago

Selected Answer: A

◆ Quick Comparison

Global Accelerator → accelerates dynamic traffic (APIs, gaming, VoIP, EC2, ALB, NLB).

CloudFront → accelerates static content (S3 objects, cached files, media).

upvoted 1 times

uk_dataguy 6 months, 2 weeks ago

Selected Answer: A

Correct: A — CloudFront with both S3 + ALB as origins gives one global CDN entry point for all content.

upvoted 1 times

Moindex65 10 months ago

Selected Answer: A

Single global edge layer: CloudFront caches static content from S3 and can also front the dynamic content from the ALB (with path-based behaviors and appropriate cache/TTL settings).

Lower latency worldwide: Users hit the nearest CloudFront edge; fewer round-trips to your regions.

Simple DNS: Point your custom domain in Route 53 to the CloudFront distribution.

Why not the others:

B, C, D: AWS Global Accelerator does not support S3 as an endpoint, and splitting endpoints complicates DNS and user experience without benefit for this use case.

upvoted 2 times

Clouddon 11 months, 2 weeks ago

Selected Answer: A

se Amazon CloudFront with multiple origins (S3 + ALB) and Route 53 routing to CloudFront to optimize performance, simplify architecture, and reduce latency globally for both static and dynamic web content.

upvoted 1 times

Juju66 1 year, 1 month ago

Selected Answer: A

Option C is wrong , because:

This option introduces additional complexity and cost by using AWS Global Accelerator. CloudFront alone can handle the distribution of both static and dynamic content effectively.

upvoted 1 times

SamKuo 1 year, 5 months ago

Selected Answer: C

AWS Global Accelerator:Dynamic content

CloudFront:static content

upvoted 2 times

PhucVuu 1 year, 8 months ago

Keywords:

- The web application has static data and dynamic data. Static data in an Amazon S3 bucket.
- Improve performance and reduce latency for the static data and dynamic data.
- The company is using its own domain name registered with Amazon Route 53.

A: Correct - CloudFront has the Edge location and the cache for dynamic and static

B: Incorrect - AWS Global Accelerator don't have cache function, so static file need to be load directly from S3 every time.

- Beside that we configure CloudFront -> ALB, Accelerator -> S3, Route 53 -> CloudFront. It means that all the traffic go to CloudFront only, Accelerator don't have any traffic.

C: Incorrect - Global Accelerator can configure CloudFront as the endpoint.

D: Incorrect - We already have domain name. Why will we use new domain name? Will we change to new domain name? How everyone know you new domain name?

upvoted 15 times

zoe9z 1 year, 6 months ago

For C, it seems like Global Accelerator does not support CloudFront as an endpoint

upvoted 2 times

MOSHE 1 year, 8 months ago

Selected Answer: A

A. Create an Amazon CloudFront distribution that has the S3 bucket and the ALB as origins. Configure Route 53 to route traffic to the CloudFront distribution.

Here's the reasoning:

CloudFront with Multiple Origins: CloudFront allows you to set up multiple origins for your distribution, so you can use both the ALB (for dynamic content) and the S3 bucket (for static content) as origins. This means that both your dynamic and static content can be served through CloudFront, which will cache content at edge locations to reduce latency.

Route 53 Integration with CloudFront: Amazon Route 53 can be easily configured to route traffic for your domain to a CloudFront distribution. Users will access your domain, and Route 53 will direct them to the nearest CloudFront edge location.

upvoted 3 times

aropl 1 year, 8 months ago

A is correct, other answers have wrong origin or endpoint types.

Cloudfront supports multiple origins on the same distribution (ALB and S3) in our case.

B incorrect - Global Accelerator Standard accelerator doesn't support s3 endpoints

c incorrect - Global Accelerator Standard accelerator doesn't support CloudFront distribution as endpoint

D incorrect - Global Accelerator Standard accelerator doesn't support s3 endpoints

upvoted 10 times

rainiverse 1 year, 8 months ago

Selected Answer: A

I'm wavering between A and C.

With dynamic content, CloudFront is cacheable and that's not good.

But with answer C, AWS Global doesn't support Cloudfront endpoint

"Endpoints for standard accelerators in AWS Global Accelerator can be Network Load Balancers, Application Load Balancers, Amazon EC2 instances, or Elastic IP addresses."

So I choose A

upvoted 2 times

PaulGa 1 year, 10 months ago

Selected Answer: C

Ans C - for the good reasons given by Diddy99. Altho' Ans A could do it, it is not the best optimised answer; Ans C is, but at cost of a custom domain name (which I don't like)

upvoted 1 times

KTEgghead 1 year, 10 months ago

Selected Answer: A

Cloudfront caches content at edge locations, reduces latency, and can serve static content from S3 buckets. It can also accelerate dynamic content from EC2. CloudFront maintains a persistent pool of connections to the origin, which minimises the overhead of establishing new connections. Any of these questions with "latency" and "improve performance" smell like CloudFront.

upvoted 1 times

creamymangosauce 1 year, 11 months ago

Selected Answer: A

A - CloudFront for caching static content. No need for Global Accelerator since no static IP is required

upvoted 2 times

A company performs monthly maintenance on its AWS infrastructure. During these maintenance activities, the company needs to rotate the credentials for its Amazon RDS for MySQL databases across multiple AWS Regions. Which solution will meet these requirements with the LEAST operational overhead?

- A. Store the credentials as secrets in AWS Secrets Manager. Use multi-Region secret replication for the required Regions. Configure Secrets Manager to rotate the secrets on a schedule. **Most Voted**
- B. Store the credentials as secrets in AWS Systems Manager by creating a secure string parameter. Use multi-Region secret replication for the required Regions. Configure Systems Manager to rotate the secrets on a schedule.
- C. Store the credentials in an Amazon S3 bucket that has server-side encryption (SSE) enabled. Use Amazon EventBridge (Amazon CloudWatch Events) to invoke an AWS Lambda function to rotate the credentials.
- D. Encrypt the credentials as secrets by using AWS Key Management Service (AWS KMS) multi-Region customer managed keys. Store the secrets in an Amazon DynamoDB global table. Use an AWS Lambda function to retrieve the secrets from DynamoDB. Use the RDS API to rotate the secrets.

Correct Answer: A

Community vote distribution

A (100%)

Comments

rein_chau **Highly Voted** 3 years, 8 months ago

Selected Answer: A

A is correct.

upvoted 25 times

PhucVuu **Highly Voted** 1 year, 8 months ago

Selected Answer: A

Keywords:

- rotate the credentials for its Amazon RDS for MySQL databases across multiple AWS Regions
- LEAST operational overhead

A: Correct - AWS Secrets Manager supports

- Encrypt credential for RDS, DocumentDb, Redshift, other DBs and key/value secret.

- multi-region replication.

- Remote base on schedule

B: Incorrect - Secure string parameter only apply for Parameter Store. All the data in AWS Secrets Manager is encrypted

C: Incorrect - don't mention about replicate S3 across region.

D: Incorrect - So many steps compare to answer A =))

upvoted 17 times

srf_deaf_rwr234dfhjyfgseqwljuk **Most Recent** 3 months, 1 week ago

Selected Answer: A

A is correct.

upvoted 1 times

nesyto 4 months, 3 weeks ago

Selected Answer: A

correct answer is A , secrets manager can rotate keys

upvoted 1 times

jucejara 11 months, 1 week ago

Selected Answer: A

A. Ya el servicio cómo tál es para el uso que se requiere.

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: A

Use AWS Secrets Manager with automatic rotation and multi-Region replication to securely manage and rotate Amazon RDS for MySQL credentials with the least operational overhead.

upvoted 1 times

MundiChor 1 year, 3 months ago

Selected Answer: A

Capability to rotate secrets and tight coupling with RDS

upvoted 1 times

Vandaman 1 year, 3 months ago

Selected Answer: A

No other explanation necessary

upvoted 1 times

Mrigraj12 1 year, 4 months ago

Selected Answer: A

A is the right answer as it takes minutes to setup. so least operational head

upvoted 1 times

trinh_le 1 year, 7 months ago

Selected Answer: A

A: correct - the secret manager supports rotating credentials

B: incorrect - Parameter Store does not perform any cryptographic operations. Instead, it relies on AWS KMS to encrypt and decrypt secure string parameter values

C and D: incorrect - handles through Lambda, require more operational overhead

upvoted 2 times

SilentMilli 1 year, 8 months ago

Selected Answer: A

AWS Secrets Manager is a secrets management service that enables you to store, manage, and rotate secrets such as database credentials, API keys, and SSH keys. Secrets Manager can help you minimize the operational overhead of rotating credentials for your Amazon RDS for MySQL databases across multiple Regions. With Secrets Manager, you can store the credentials as secrets and use multi-Region secret replication to replicate the secrets to the required Regions. You can then configure Secrets Manager to rotate the secrets on a schedule so that the credentials are rotated automatically without the need for manual intervention. This can help reduce the risk of secrets being compromised and minimize the operational overhead of credential management.

upvoted 5 times

Buruguduystunstugudunstuy 1 year, 8 months ago

Selected Answer: A

Option A, storing the credentials as secrets in AWS Secrets Manager and using multi-Region secret replication for the required Regions, and configuring Secrets Manager to rotate the secrets on a schedule, would meet the requirements with the least operational overhead.

AWS Secrets Manager allows you to store, manage, and rotate secrets, such as database credentials, across multiple AWS Regions. By enabling multi-Region secret replication, you can replicate the secrets across the required Regions to allow for seamless rotation of the credentials during maintenance activities. Additionally, Secrets Manager provides automatic rotation of secrets on a schedule, which would minimize the operational overhead of rotating the credentials on a monthly basis.

upvoted 3 times

Buruguduystunstugudunstuy 3 years, 5 months ago

Option B, storing the credentials as secrets in AWS Systems Manager and using multi-Region secret replication, would not provide automatic rotation of secrets on a schedule.

Option C, storing the credentials in an S3 bucket with SSE enabled and using EventBridge to invoke an AWS Lambda function to rotate the credentials, would not provide automatic rotation of secrets on a schedule.

Option D, encrypting the credentials as secrets using KMS multi-Region customer managed keys and storing the secrets in a DynamoDB global table, would not provide automatic rotation of secrets on a schedule and would require additional operational overhead to retrieve the secrets from DynamoDB and use the RDS API to rotate the secrets.

upvoted 3 times

Musti35 1 year, 8 months ago

Selected Answer: A

With Secrets Manager, you can store, retrieve, manage, and rotate your secrets, including database credentials, API keys, and other secrets. When you create a secret using Secrets Manager, it's created and managed in a Region of your choosing. Although scoping secrets to a Region is a security

best practice, there are scenarios such as disaster recovery and cross-Regional redundancy that require replication of secrets across Regions. Secrets Manager now makes it possible for you to easily replicate your secrets to one or more Regions to support these scenarios.

upvoted 3 times

linux_admin 1 year, 8 months ago

Selected Answer: A

A. Store the credentials as secrets in AWS Secrets Manager. Use multi-Region secret replication for the required Regions. Configure Secrets Manager to rotate the secrets on a schedule.

This solution is the best option for meeting the requirements with the least operational overhead. AWS Secrets Manager is designed specifically for managing and rotating secrets like database credentials. Using multi-Region secret replication, you can easily replicate the secrets across the required AWS Regions. Additionally, Secrets Manager allows you to configure automatic secret rotation on a schedule, further reducing the operational overhead.

upvoted 1 times

PaulGa 1 year, 10 months ago

Selected Answer: A

Ans A. What's to debate...?

upvoted 1 times

creamymangosauce 1 year, 11 months ago

Selected Answer: A

A - Secrets Manager automates the rotation of secrets for RDS without own implementation required, the options require effort to implement the secret rotation logic

upvoted 1 times

ics_911 2 years, 4 months ago

Selected Answer: A

A is correct.

upvoted 1 times

A company runs an ecommerce application on Amazon EC2 instances behind an Application Load Balancer. The instances run in an Amazon EC2 Auto Scaling group across multiple Availability Zones. The Auto Scaling group scales based on CPU utilization metrics. The ecommerce application stores the transaction data in a MySQL 8.0 database that is hosted on a large EC2 instance. The database's performance degrades quickly as application load increases. The application handles more read requests than write transactions. The company wants a solution that will automatically scale the database to meet the demand of unpredictable read workloads while maintaining high availability.

Which solution will meet these requirements?

- A. Use Amazon Redshift with a single node for leader and compute functionality.
- B. Use Amazon RDS with a Single-AZ deployment. Configure Amazon RDS to add reader instances in a different Availability Zone.
- C. Use Amazon Aurora with a Multi-AZ deployment. Configure Aurora Auto Scaling with Aurora Replicas. **Most Voted**
- D. Use Amazon ElastiCache for Memcached with EC2 Spot Instances.

Correct Answer: C

Community vote distribution

C (99%)

Comments

D2w **Highly Voted** 3 years, 8 months ago

Selected Answer: C

C, AURORA is 5x performance improvement over MySQL on RDS and handles more read requests than write;; maintaining high availability = Multi-AZ deployment
upvoted 55 times

Buruguduystunstugudunstuy **Highly Voted** 3 years, 5 months ago

Selected Answer: C

Option C, using Amazon Aurora with a Multi-AZ deployment and configuring Aurora Auto Scaling with Aurora Replicas, would be the best solution to meet the requirements.

Aurora is a fully managed, MySQL-compatible relational database that is designed for high performance and high availability. Aurora Multi-AZ deployments automatically maintain a synchronous standby replica in a different Availability Zone to provide high availability. Additionally, Aurora Auto Scaling allows you to automatically scale the number of Aurora Replicas in response to read workloads, allowing you to meet the demand of unpredictable read workloads while maintaining high availability. This would provide an automated solution for scaling the database to meet the demand of the application while maintaining high availability.

upvoted 29 times

Buruguduystunstugudunstuy 3 years, 5 months ago

Option A, using Amazon Redshift with a single node for leader and compute functionality, would not provide high availability.

Option B, using Amazon RDS with a Single-AZ deployment and configuring RDS to add reader instances in a different Availability Zone, would not provide high availability and would not automatically scale the number of reader instances in response to read workloads.

Option D, using Amazon ElastiCache for Memcached with EC2 Spot Instances, would not provide a database solution and would not meet the requirements.

upvoted 10 times

javicia **Most Recent** 4 months, 2 weeks ago

Selected Answer: D

D: D: For me, it would be a Redis of reading.

upvoted 1 times

nesyto 4 months, 3 weeks ago

Selected Answer: C

correct answer is C
upvoted 1 times

nesyto 4 months, 3 weeks ago

Selected Answer: D

amazon elastic cache would be perfect solution answer is D
upvoted 1 times

ZephyrNguyen15 6 months, 3 weeks ago

Selected Answer: C

Zephyr choose C
upvoted 1 times

S_K_han 9 months ago

Selected Answer: C

C is right as Aurora works well with MySQL
upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: C

Use Amazon Aurora with Multi-AZ deployment and Auto Scaling Aurora Replicas to meet unpredictable read demand with high availability, performance, and minimal operational overhead.
upvoted 1 times

melvis8 1 year, 2 months ago

Selected Answer: C

Using aurora is the best answer for this issue since aurora is ideal for unpredictable workloads and we can manage read heavy workloads by deploying the database with aurora replicas
upvoted 2 times

Tjazz04 1 year, 6 months ago

Selected Answer: C

Keyword: High Availability
Only letter C has multi-AZ deployment. Also, it automatically scales replicas for unpredictable read workload
upvoted 3 times

VINVIN99 1 year, 7 months ago

Selected Answer: C

CCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCCC
upvoted 1 times

PaulGa 1 year, 10 months ago

Selected Answer: C

Ans C. Its the only one that offers right level of scaling and availability
upvoted 2 times

OBIOHAnze 2 years ago

Selected Answer: C

Here's why C is the best solution:

Amazon Aurora: A managed, high-performance MySQL-compatible relational database engine.
Multi-AZ deployment: Ensures high availability in case of an AZ failure.
Aurora Auto Scaling with Aurora Replicas: Automatically scales read replicas based on traffic, improving read performance.
upvoted 3 times

A_jaa 2 years, 5 months ago

Selected Answer: C

Answer-c
upvoted 1 times

Ndlesty 2 years, 6 months ago

Selected Answer: C

key statement: "...will automatically scale the database to meet the demand of unpredictable read workloads while maintaining high availability.
upvoted 2 times

AWSGuru123 2 years, 8 months ago

Selected Answer: C

Aurora

upvoted 1 times

Syruis 2 years, 10 months ago

Selected Answer: C

C fit perfectly

upvoted 1 times


A company recently migrated to AWS and wants to implement a solution to protect the traffic that flows in and out of the production VPC. The company had an inspection server in its on-premises data center. The inspection server performed specific operations such as traffic flow inspection and traffic filtering. The company wants to have the same functionalities in the AWS Cloud.

Which solution will meet these requirements?

- A. Use Amazon GuardDuty for traffic inspection and traffic filtering in the production VPC.
- B. Use Traffic Mirroring to mirror traffic from the production VPC for traffic inspection and filtering.
- C. Use AWS Network Firewall to create the required rules for traffic inspection and traffic filtering for the production VPC.
Most Voted
- D. Use AWS Firewall Manager to create the required rules for traffic inspection and traffic filtering for the production VPC.

Correct Answer: C

Community vote distribution



Comments

SilentMilli **Highly Voted** 1 year, 8 months ago

Selected Answer: C

I would recommend option C: Use AWS Network Firewall to create the required rules for traffic inspection and traffic filtering for the production VPC

AWS Network Firewall is a managed firewall service that provides filtering for both inbound and outbound network traffic. It allows you to create rules for traffic inspection and filtering, which can help protect your production VPC.

Option A: Amazon GuardDuty is a threat detection service, not a traffic inspection or filtering service.

Option B: Traffic Mirroring is a feature that allows you to replicate and send a copy of network traffic from a VPC to another VPC or on-premises location. It is not a service that performs traffic inspection or filtering.

Option D: AWS Firewall Manager is a security management service that helps you to centrally configure and manage firewalls across your accounts. It is not a service that performs traffic inspection or filtering.

upvoted 156 times

BoboChow **Highly Voted** 1 year, 8 months ago

Selected Answer: C

I agree with C.

AWS Network Firewall is a stateful, managed network firewall and intrusion detection and prevention service for your virtual private cloud (VPC) that you created in Amazon Virtual Private Cloud (Amazon VPC). With Network Firewall, you can filter traffic at the perimeter of your VPC. This includes filtering traffic going to and coming from an internet gateway, NAT gateway, or over VPN or AWS Direct Connect.

upvoted 26 times

BoboChow 3 years, 8 months ago

And I'm not sure Traffic Mirroring can be for filtering

upvoted 4 times

nesyto **Most Recent** 4 months, 3 weeks ago

Selected Answer: C

correct answer is C Firewall protects and inspects traffic for vpc

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: C

Use AWS Network Firewall to implement inline traffic inspection and filtering in the production VPC — the most direct and scalable cloud-native replacement for your on-prem inspection server.

upvoted 1 times

Juju66 1 year, 1 month ago

Selected Answer: C

I agree C.

Option D AWS Firewall Manager sounds similar, but it's not a service that can do traffic inspection.

upvoted 1 times

satyaamm 1 year, 5 months ago

Selected Answer: C

AWS Firewall Manager is the only one that actually helps direct the network traffic in AWS. GuardDuty and AWS Firewall Manager recognized inspection against threats.

upvoted 1 times

Tjazz04 1 year, 6 months ago

Selected Answer: C

C is the most appropriate when it comes to traffic flow inspection and filtering

upvoted 1 times

cookieMr 1 year, 8 months ago

Selected Answer: C

AWS Network Firewall is a managed network firewall service that allows you to define firewall rules to filter and inspect network traffic. You can create rules to define the traffic that should be allowed or blocked based on various criteria such as source/destination IP addresses, protocols, ports, and more. With AWS Network Firewall, you can implement traffic inspection and filtering capabilities within the production VPC, helping to protect the network traffic.

In the context of the given scenario, AWS Network Firewall can be a suitable choice if the company wants to implement traffic inspection and filtering directly within the VPC without the need for traffic mirroring. It provides an additional layer of security by enforcing specific rules for traffic filtering, which can help protect the production environment.

upvoted 3 times

Guru4Cloud 1 year, 8 months ago

Selected Answer: C

- AWS Network Firewall is a managed network security service that provides stateful inspection of traffic and allows you to define firewall rules to control the traffic flow in and out of your VPC.
- With AWS Network Firewall, you can create custom rule groups to define specific operations for traffic inspection and filtering.
- It can perform deep packet inspection and filtering at the network level to enforce security policies, block malicious traffic, and allow or deny traffic based on defined rules.
- By integrating AWS Network Firewall with the production VPC, you can achieve similar functionalities as the on-premises inspection server, performing traffic flow inspection and filtering.

upvoted 2 times

PaulGa 1 year, 10 months ago

Selected Answer: C

Ans C. As per good response by SilentMili

upvoted 1 times

TheFivePips 2 years, 3 months ago

Selected Answer: C

I didn't realize the network firewall could do inspection, but here's what the documentation says:

AWS Network Firewall supports Transport Layer Security (TLS) inspection, allowing customers to strengthen their security posture on AWS by improving visibility into encrypted traffic flows. You can use AWS Network Firewall to decrypt TLS sessions and inspect both inbound and outbound Amazon Virtual Private Cloud (VPC) traffic without the need to deploy or manage any additional network security infrastructure. Encryption and decryption happen on the same firewall instance natively, so traffic does not cross any network boundaries.

upvoted 4 times

awsgeek75 2 years, 5 months ago

Selected Answer: C

Network Firewall to define firewall rules for traffic inspection.

A: GuardDuty is not for this

B: Wrong product

D: Firewall Manager does not monitor traffic, it manages firewall

upvoted 1 times

A_jaa 2 years, 5 months ago

Selected Answer: C

Answer-C

upvoted 1 times

danielpark99 2 years, 8 months ago

Selected Answer: C

AWS Network Firewall to support from layer 3 to layer 7 protection, it is able to inspect any direction lets say vpc to vpc and outbound and inbound and even supporting direct connect and site to site vpn

upvoted 1 times

reema908516 2 years, 9 months ago

Selected Answer: C

AWS Network Firewall is a managed firewall service that provides filtering for both inbound and outbound network traffic. It allows you to create rules for traffic inspection and filtering, which can help protect your production VPC.

upvoted 1 times

nmywrld 2 years, 9 months ago

Why isn't D viable? Firewall Manager will help to provision network firewall as required if you define it in firewall manager. And it's fully managed, no requiring you to do any configuration or set up.

upvoted 1 times

pentium75 2 years, 5 months ago

Because we need a firewall, not a service that we COULD IN THEORY use to create a firewall?

upvoted 2 times

Syruis 2 years, 10 months ago

Selected Answer: C

C with no doubt

upvoted 1 times

A company hosts a data lake on AWS. The data lake consists of data in Amazon S3 and Amazon RDS for PostgreSQL. The company needs a reporting solution that provides data visualization and includes all the data sources within the data lake. Only the company's management team should have full access to all the visualizations. The rest of the company should have only limited access.

Which solution will meet these requirements?

A. Create an analysis in Amazon QuickSight. Connect all the data sources and create new datasets. Publish dashboards to visualize the data. Share the dashboards with the appropriate IAM roles.

B. Create an analysis in Amazon QuickSight. Connect all the data sources and create new datasets. Publish dashboards to visualize the data. Share the dashboards with the appropriate users and groups. **Most Voted**

C. Create an AWS Glue table and crawler for the data in Amazon S3. Create an AWS Glue extract, transform, and load (ETL) job to produce reports. Publish the reports to Amazon S3. Use S3 bucket policies to limit access to the reports.

D. Create an AWS Glue table and crawler for the data in Amazon S3. Use Amazon Athena Federated Query to access data within Amazon RDS for PostgreSQL. Generate reports by using Amazon Athena. Publish the reports to Amazon S3. Use S3 bucket policies to limit access to the reports.

Correct Answer: B

Community vote distribution

B (86%)

8%

Comments

rodriiviru **Highly Voted** 3 years, 8 months ago

Selected Answer: B

B

upvoted 72 times

BoboChow 3 years, 8 months ago

Agree with you

upvoted 3 times

PhucVuu **Highly Voted** 1 year, 8 months ago

Selected Answer: B

Keywords:

- Data lake on AWS.

- Consists of data in Amazon S3 and Amazon RDS for PostgreSQL.

- The company needs a reporting solution that provides data VISUALIZATION and includes ALL the data sources within the data lake.

A - Incorrect: Amazon QuickSight only support users(standard version) and groups (enterprise version). users and groups only exists without QuickSight. QuickSight don't support IAM. We use users and groups to view the QuickSight dashboard

B - Correct: as explained in answer A and QuickSight is used to created dashboard from S3, RDS, Redshift, Aurora, Athena, OpenSearch, Timestream

C - Incorrect: This way don't support visulization and don't mention how to process RDS data

D - Incorrect: This way don't support visulization and don't mention how to combine data RDS and S3

upvoted 66 times

javia **Most Recent** 4 months, 2 weeks ago

Selected Answer: A

A. Cree un análisis en Amazon QuickSight

upvoted 1 times

nesyto 4 months, 3 weeks ago

Selected Answer: B

the correct answer is B, the key is visualisation

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: B

Use Amazon QuickSight to build dashboards from S3 and RDS sources and control access using users and groups — providing a secure, interactive, and scalable reporting solution for your data lake.

upvoted 2 times

Sanki_26 1 year, 2 months ago

Selected Answer: B

Amazon QuickSight is a data visualization service that allows you to create interactive dashboards and reports from various data sources, including Amazon S3 and Amazon RDS for PostgreSQL. You can connect all the data sources and create new datasets in QuickSight, and then publish dashboards to visualize the data. You can also share the dashboards with the appropriate users and groups, and control their access levels using IAM roles and permissions.

upvoted 3 times

Vandaman 1 year, 3 months ago

Selected Answer: B

It was the use of users and groups that answered that for me

upvoted 1 times

FlyingHawk 1 year, 5 months ago

Selected Answer: B

Amazon QuickSight can handle dynamic data through Direct Query Mode or scheduled data refreshes. It provides real-time or near-real-time visualizations, interactive dashboards, and fine-grained access control, making it the most suitable solution for the company's requirements.

<https://aws.amazon.com/blogs/business-intelligence/best-practices-for-amazon-quicksight-spice-and-direct-query-mode/>

I saw the question in udemy practice test 5 from Neal Davis, he suggested the correct answer is D, his reasons: option B solves the problem of access sharing with resources but does not take care of delta in data.

upvoted 2 times

MGKYAING 1 year, 5 months ago

Selected Answer: B

Amazon QuickSight is designed for creating visualizations and dashboards, and it can easily connect to a variety of data sources such as Amazon S3 and Amazon RDS for PostgreSQL.

You can create an analysis in QuickSight by connecting the required data sources, creating datasets, and then building dashboards.

Access control in QuickSight is robust. You can share dashboards with specific users and groups, ensuring that only the management team has full access to all the visualizations, while the rest of the company can have limited access.

upvoted 1 times

Tjazz04 1 year, 6 months ago

Selected Answer: B

Keyword: data visualization = QuickSight

QuickSight defines Users (standard versions) and Groups (enterprise version)

These users & groups only exist within QuickSight, not IAM

So the correct answer is B.

upvoted 2 times

EzKkk 1 year, 6 months ago

Selected Answer: B

The question requires a solution for data visualization, which means it focuses solely on downstream consumption. Therefore, any solution that includes upstream processing is off the table (C & D). This leaves us with two options: A and B. Essentially, they are the same with one key difference: A manages access through IAM Roles, while B manages it through an IAM group. Since only the management team is granted permission, using a group will be more efficient, as it scales better with changes in personnel and provides a centralized point for managing permissions.

upvoted 1 times

EzKkk 1 year, 6 months ago

The question requires a solution for data visualization, which means it focuses solely on downstream consumption. Therefore, any solution that includes upstream processing is off the table (C & D). This leaves us with two options: A and B. Essentially, they are the same with one key difference: A manages access through IAM Roles, while B manages it through an IAM group. Since only the management team is granted permission, using a group will be more efficient, as it scales better with changes in personnel and provides a centralized point for managing permissions.

upvoted 2 times

hanen 1 year, 8 months ago

Selected Answer: D

If you have data in sources other than Amazon S3, you can use Athena Federated Query to query the data in place or build pipelines that extract data from multiple data sources and store them in Amazon S3. With Athena Federated Query, you can run SQL queries across data stored in relational, non-relational, object, and custom data sources.

Athena uses data source connectors that run on AWS Lambda to run federated queries. A data source connector is a piece of code that can translate between your target data source and Athena. You can think of a connector as an extension of Athena's query engine. Prebuilt Athena data source connectors exist for data sources like Amazon CloudWatch Logs, Amazon DynamoDB, Amazon DocumentDB, and Amazon RDS, and JDBC-compliant relational data sources such MySQL, and PostgreSQL under the Apache 2.0 license

upvoted 4 times

FlyingHawk 1 year, 5 months ago

D relies on static reports stored in Amazon S3.

If the data changes, the reports would need to be regenerated and re-published, which is not ideal for dynamic data.

They also lack interactive visualization capabilities and fine-grained access control for dashboards.

upvoted 1 times

JA2018 1 year, 7 months ago

key phases: " The data lake consists of data in Amazon S3 and Amazon RDS for PostgreSQL. The company needs a reporting solution that provides data visualization and includes all the data sources within the data lake. "

upvoted 1 times

linux_admin 1 year, 8 months ago

Selected Answer: B

Option B is the correct answer because Amazon QuickSight's sharing mechanism is based on users and groups, not IAM roles. IAM roles are used for granting permissions to AWS resources, but they are not directly used for sharing QuickSight dashboards.

In option B, you create an analysis in Amazon QuickSight, connect all the data sources (Amazon S3 and Amazon RDS for PostgreSQL), and create new datasets. After publishing dashboards to visualize the data, you share them with appropriate users and groups. This approach allows you to control the access levels for different users, such as providing full access to the management team and limited access to the rest of the company. This solution meets the requirements specified in the question.

upvoted 7 times

elearningtakai 1 year, 8 months ago

Selected Answer: B

Amazon QuickSight is a cloud-based business intelligence (BI) service that makes it easy to create and publish interactive dashboards that include data visualizations from multiple data sources. By using QuickSight, the company can connect to both Amazon S3 and Amazon RDS for PostgreSQL and create new datasets that combine data from both sources. The company can then use QuickSight to create interactive dashboards that visualize the data and provide data insights.

To limit access to the visualizations, the company can use QuickSight's built-in security features. QuickSight allows you to define fine-grained access control at the user or group level. This way, the management team can have full access to all the visualizations, while the rest of the company can have only limited access.

upvoted 4 times

dszes 1 year, 8 months ago

tricky question, Users, groups and roles can have access.

Viewing who has access to a dashboard

Use the following procedure to see which users or groups have access to the dashboard.

Open the published dashboard and choose Share at upper right. Then choose Share dashboard.

In the Share dashboard page that opens, under Manage permissions, review the users and groups, and their roles and settings.

You can search to locate a specific user or group by entering their name or any part of their name in the search box at upper right. Searching is case-sensitive, and wildcards aren't supported. Delete the search term to return the view to all users.

upvoted 2 times

cookieMr 1 year, 8 months ago

Selected Answer: B

B. Create an analysis in Amazon QuickSight. Connect all the data sources and create new datasets. Publish dashboards to visualize the data. Share the dashboards with the appropriate users and groups.

Amazon QuickSight is a business intelligence (BI) tool provided by AWS that allows you to create interactive dashboards and reports. It supports a variety of data sources, including Amazon S3 and Amazon RDS for PostgreSQL, which are the data sources in the company's data lake.

Option A (Create an analysis in Amazon QuickSight and share with IAM roles) is incorrect because it suggests sharing with IAM roles, which are more suitable for managing access to AWS resources rather than granting access to specific users or groups within QuickSight.

upvoted 4 times

A company is implementing a new business application. The application runs on two Amazon EC2 instances and uses an Amazon S3 bucket for document storage. A solutions architect needs to ensure that the EC2 instances can access the S3 bucket. What should the solutions architect do to meet this requirement?

- A. Create an IAM role that grants access to the S3 bucket. Attach the role to the EC2 instances. **Most Voted**
- B. Create an IAM policy that grants access to the S3 bucket. Attach the policy to the EC2 instances.
- C. Create an IAM group that grants access to the S3 bucket. Attach the group to the EC2 instances.
- D. Create an IAM user that grants access to the S3 bucket. Attach the user account to the EC2 instances.

Correct Answer: A

Community vote distribution

A (100%)

Comments

sba21 **Highly Voted** 3 years, 8 months ago

Selected Answer: A

Always remember that you should associate IAM roles to EC2 instances
upvoted 108 times

Buruguduystunstugudunstuy **Highly Voted** 3 years, 5 months ago

Selected Answer: A

The correct option to meet this requirement is A: Create an IAM role that grants access to the S3 bucket and attach the role to the EC2 instances.

An IAM role is an AWS resource that allows you to delegate access to AWS resources and services. You can create an IAM role that grants access to the S3 bucket and then attach the role to the EC2 instances. This will allow the EC2 instances to access the S3 bucket and the documents stored within it.

Option B is incorrect because an IAM policy is used to define permissions for an IAM user or group, not for an EC2 instance.

Option C is incorrect because an IAM group is used to group together IAM users and policies, not to grant access to resources.

Option D is incorrect because an IAM user is used to represent a person or service that interacts with AWS resources, not to grant access to resources.

upvoted 77 times

nesyto **Most Recent** 4 months, 3 weeks ago

Selected Answer: A

correct answer is A
upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: A

Use an IAM role attached to the EC2 instances to securely and efficiently grant access to the Amazon S3 bucket.
upvoted 1 times

jr1521675287 11 months, 2 weeks ago

Selected Answer: A

only A
upvoted 1 times

satyaamm 1 year, 5 months ago

Selected Answer: A

Only IAM role is suitable here as policies are added to IAM groups or users. Moreover, IAM groups cannot be attached to a resource and attaching IAM user credentials to a IAM role is risky and thereby not suitable.

upvoted 3 times

Tjazz04 1 year, 6 months ago

Selected Answer: A

Ans. A

Here's why:

IAM Role: Roles are designed to be assumed by entities like EC2 instances. By creating an IAM role with the necessary permissions to access the S3 bucket and attaching this role to the EC2 instances, you ensure that the instances can securely access the S3 bucket without needing to manage long-term credentials.

IAM Policy: While policies define permissions, they need to be attached to roles or users. Attaching a policy directly to EC2 instances is not possible.

IAM Group: Groups are used to manage permissions for multiple users, not instances.

IAM User: Users are intended for individual people or applications, not for EC2 instances.

By using an IAM role, you follow AWS best practices for security and manageability. If you have any more questions or need further clarification, feel free to ask!

upvoted 2 times

EzKkk 1 year, 6 months ago

Selected Answer: A

IAM Role + EC2 instance = go-to solution

upvoted 1 times

PaulGa 1 year, 10 months ago

Selected Answer: A

Ans A - as per "Buruguduystunstugudunstuy" response.

upvoted 1 times

A_jaa 2 years, 5 months ago

Selected Answer: A

Answer-A

upvoted 2 times

thewalker 2 years, 5 months ago

Selected Answer: A

Below is the response from Amazon Q:

To access S3 from an EC2 instance, you need to create an IAM role and associate that role with the EC2 instance. Here are the basic steps:

1. Create an IAM role and attach the AmazonS3ReadOnlyAccess or AmazonS3FullAccess managed policy to grant S3 access.
2. Launch the EC2 instance and select the IAM role you created during launch.
3. The instance will now have the permissions defined in the IAM role and you can access S3 from the instance.

upvoted 2 times

thewalker 2 years, 5 months ago

Some key points:

1. Attaching an IAM role is preferred over creating a resource-based policy for S3 access from EC2 as it provides centralized access management.
2. The instance will need internet access to communicate with S3. Make sure the associated security group and NACL rules allow outbound internet access.
3. Check AWS documentation for latest steps to create and associate an IAM role with an EC2 instance. The console and CLI provide options to automate this process.

upvoted 1 times

jjcode 2 years, 6 months ago

Strangely straight forward, Almost had me confused.

upvoted 1 times

GabrielSGoncalves 2 years, 7 months ago

Selected Answer: A

For sure

upvoted 1 times

Ruffyit 2 years, 7 months ago

The correct option to meet this requirement is A: Create an IAM role that grants access to the S3 bucket and attach the role to the EC2 instances.

An IAM role is an AWS resource that allows you to delegate access to AWS resources and services. You can create an IAM role that grants access to the S3 bucket and then attach the role to the EC2 instances. This will allow the EC2 instances to access the S3 bucket and the documents stored within it.

Option B is incorrect because an IAM policy is used to define permissions for an IAM user or group, not for an EC2 instance.

Option C is incorrect because an IAM group is used to group together IAM users and policies, not to grant access to resources.

Option D is incorrect because an IAM user is used to represent a person or service that interacts with AWS resources, not to grant access to resources.

upvoted 1 times

danielpark99 2 years, 8 months ago

Selected Answer: A

EC2 instances should be associated with IAM roles.

Policies can be applying to users and groups can help to apply multiple roles.

upvoted 1 times

Abdou1604 2 years, 10 months ago

Option B may work but ,

suggests creating an IAM policy directly and attaching it to the EC2 instances. While this might work, it's not the recommended approach. Using an IAM role is more secure and manageable.

upvoted 1 times

Guru4Cloud 2 years, 10 months ago

Selected Answer: A

Always remember that you should associate IAM roles to EC2 instances.

An IAM role is an AWS resource that allows you to delegate access to AWS resources and services. You can create an IAM role that grants access to the S3 bucket and then attach the role to the EC2 instances. This will allow the EC2 instances to access the S3 bucket and the documents stored within it.

upvoted 1 times

An application development team is designing a microservice that will convert large images to smaller, compressed images. When a user uploads an image through the web interface, the microservice should store the image in an Amazon S3 bucket, process and compress the image with an AWS Lambda function, and store the image in its compressed form in a different S3 bucket. A solutions architect needs to design a solution that uses durable, stateless components to process the images automatically. Which combination of actions will meet these requirements? (Choose two.)

A. Create an Amazon Simple Queue Service (Amazon SQS) queue. Configure the S3 bucket to send a notification to the SQS queue when an image is uploaded to the S3 bucket. **Most Voted**

B. Configure the Lambda function to use the Amazon Simple Queue Service (Amazon SQS) queue as the invocation source. When the SQS message is successfully processed, delete the message in the queue. **Most Voted**

C. Configure the Lambda function to monitor the S3 bucket for new uploads. When an uploaded image is detected, write the file name to a text file in memory and use the text file to keep track of the images that were processed.

D. Launch an Amazon EC2 instance to monitor an Amazon Simple Queue Service (Amazon SQS) queue. When items are added to the queue, log the file name in a text file on the EC2 instance and invoke the Lambda function.

E. Configure an Amazon EventBridge (Amazon CloudWatch Events) event to monitor the S3 bucket. When an image is uploaded, send an alert to an Amazon Simple Notification Service (Amazon SNS) topic with the application owner's email address for further processing.

Correct Answer: AB

Community vote distribution

AB (99%)

Comments

Buruguduystunstugudunstuy **Highly Voted** 3 years, 5 months ago

Selected Answer: AB

To design a solution that uses durable, stateless components to process images automatically, a solutions architect could consider the following actions:

Option A involves creating an SQS queue and configuring the S3 bucket to send a notification to the queue when an image is uploaded. This allows the application to decouple the image upload process from the image processing process and ensures that the image processing process is triggered automatically when a new image is uploaded.

Option B involves configuring the Lambda function to use the SQS queue as the invocation source. When the SQS message is successfully processed the message is deleted from the queue. This ensures that the Lambda function is invoked only once per image and that the image is not processed multiple times.

upvoted 44 times

Buruguduystunstugudunstuy 3 years, 5 months ago

Option C is incorrect because it involves storing state (the file name) in memory, which is not a durable or scalable solution.

Option D is incorrect because it involves launching an EC2 instance to monitor the SQS queue, which is not a stateless solution.

Option E is incorrect because it involves using Amazon EventBridge (formerly Amazon CloudWatch Events) to send an alert to an Amazon Simple Notification Service (Amazon SNS) topic, which is not related to the image processing process.

upvoted 24 times

hsinchang 2 years, 10 months ago

So storing states invokes the stateless principle, nice understanding!

upvoted 2 times

op22233 2 years, 8 months ago

A stateless system sends a request to the server and relays the response (or the state) back without storing any information. On the other hand, stateful systems expect a response, track information, and resend the request if no response is received

upvoted 4 times

sba21 **Highly Voted** 3 years, 8 months ago

Selected Answer: AB

It looks like A-B

upvoted 15 times

nesyto **Most Recent** 4 months, 3 weeks ago

Selected Answer: AB

A and B are the correct combination

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: AB

Use S3 → SQS → Lambda with stateless processing logic for a durable, scalable, and serverless image processing pipeline.

upvoted 2 times

vadps 1 year ago

Selected Answer: AB

AB: S3-SQS-Lambda

upvoted 1 times

EzKkk 1 year, 6 months ago

Selected Answer: AB

High level flow: Web (1) -> Internet (2) -> Uploaded image bucket (3) -> Compressed image bucket (4)

Let's break this down.

Image uploading (1)(2)(3): If the image's size is big, multi upload is a MUST. If you want to accelerate the process, you can also use Acceleration Transfer with additional fee.

Image compressing (3)(4): SQS Simple Queue can directly integrate with S3 events. For every image uploaded, S3 event will create a message in SQS queue. After that, Lambda can be invoked to compress the image and then upload the image to destination bucket.

upvoted 2 times

PhucVuu 1 year, 8 months ago

Selected Answer: AB

Keywords:

- Store the image in an Amazon S3 bucket, process and compress the image with an AWS Lambda function.
- Durable, stateless components to process the images automatically

A,B: Correct - SQS has message retention function(store message) default 4 days(can increase up to 14 days) so that you can re-run lambda if there are any errors when processing the images.

C: Incorrect - Lambda function just run the request then stop, the max timeout is 15 mins. So we cannot store data in the ram of Lambda function.

D: Incorrect - we can trigger Lambda directly from SQS no need EC2 instance in this case

E: Incorrect - It kinds of manually step -> the owner has to read email then process it :))

upvoted 10 times

Guru4Cloud 1 year, 8 months ago

Selected Answer: AB

Explanation:

Option A: By creating an Amazon SQS queue and configuring the S3 bucket to send a notification to the SQS queue when an image is uploaded, the system establishes a durable and scalable way to handle incoming image processing tasks.

Option B: Configuring the Lambda function to use the SQS queue as the invocation source allows it to retrieve messages from the queue and process them in a stateless manner. After successfully processing the image, the Lambda function can delete the message from the queue to avoid duplicate processing.

upvoted 3 times

DigitalDanny 1 year, 8 months ago

Selected Answer: AB

A. Create an Amazon Simple Queue Service (Amazon SQS) queue. Configure the S3 bucket to send a notification to the SQS queue when an image is uploaded to the S3 bucket.

B. Configure the Lambda function to use the Amazon Simple Queue Service (Amazon SQS) queue as the invocation source. When the SQS message is successfully processed, delete the message in the queue.

Explanation:

A (SQS Queue): Using SQS to decouple the S3 bucket from the processing components provides durability and scalability. When an image is

uploaded, a notification is sent to the SQS queue.

B (Lambda with SQS Trigger): Configuring the Lambda function to use the SQS queue as the invocation source allows for stateless and scalable image processing. Lambda can be triggered by messages in the SQS queue, and upon successful processing, the message can be deleted, ensuring that each message (image) is processed once.

This combination ensures a durable, stateless, and scalable architecture for processing images automatically in response to user uploads.

upvoted 3 times

cookieMr 1 year, 8 months ago

Selected Answer: AB

Option A is a correct because it allows for decoupling between the image upload process and image processing. By configuring S3 to send a notification to SQS, image upload event is recorded and can be processed independently by microservice.

Option B is also a correct because it ensures that Lambda is triggered by messages in SQS. Lambda can retrieve image information from SQS, process and compress image, and store compressed image in a different S3. Once processing is successful, Lambda can delete processed message from SQS, indicating that image has been processed.

Option C is not recommended because it introduces a stateful approach by using a text file to keep track of processed images.

Option D is not optimal solution as it introduces unnecessary complexity by involving an EC2 to monitor SQS and maintain a text file.

Option E is not directly related to requirement of processing images automatically. Although EventBridge and SNS can be useful for event notifications and further processing, they don't provide the same level of durability and scalability as SQS.

upvoted 6 times

PaulGa 1 year, 10 months ago

Selected Answer: AB

Ans A,B - as per Buruguduystunstugudunstuy's response: stateless, robust

upvoted 1 times

soufiyane 2 years, 2 months ago

Selected Answer: AB

whenever we talk about microservices we should mention SQS, so A and B are the right answers

upvoted 4 times

han_ds 2 years, 3 months ago

Selected Answer: AB

A/B make the most sense and in practice this works, I've done it.

upvoted 3 times

A_jaa 2 years, 5 months ago

Selected Answer: AB

Answer- A,B

upvoted 1 times

mohamedsambo 2 years, 5 months ago

I can not understand why it is not as simple like s3-1 event destination to notify the lambda function to process and upload to s3-2

upvoted 2 times

miki111 2 years, 11 months ago

Option AB MET THE REQUIREMENT

upvoted 1 times

RupeC 2 years, 11 months ago

Selected Answer: AB

D and E are distractions. C seems a valid solution. However, as you have to select two, A and B are the only two that work in conjunction with each other.

upvoted 2 times

A company has a three-tier web application that is deployed on AWS. The web servers are deployed in a public subnet in a VPC. The application servers and database servers are deployed in private subnets in the same VPC. The company has deployed a third-party virtual firewall appliance from AWS Marketplace in an inspection VPC. The appliance is configured with an IP interface that can accept IP packets.

A solutions architect needs to integrate the web application with the appliance to inspect all traffic to the application before the traffic reaches the web server.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create a Network Load Balancer in the public subnet of the application's VPC to route the traffic to the appliance for packet inspection.
- B. Create an Application Load Balancer in the public subnet of the application's VPC to route the traffic to the appliance for packet inspection.
- C. Deploy a transit gateway in the inspection VPC. Configure route tables to route the incoming packets through the transit gateway.
- D. Deploy a Gateway Load Balancer in the inspection VPC. Create a Gateway Load Balancer endpoint to receive the incoming packets and forward the packets to the appliance. **Most Voted**

Correct Answer: D

Community vote distribution

D (84%)

Other

Comments

CloudGuru99 **Highly Voted** 3 years, 8 months ago

Answer is D . Use Gateway Load balancer

upvoted 49 times

pm2229 **Highly Voted** 1 year, 8 months ago

It's D, Coz.. Gateway Load Balancer is a new type of load balancer that operates at layer 3 of the OSI model and is built on Hyperplane, which is capable of handling several thousands of connections per second. Gateway Load Balancer endpoints are configured in spoke VPCs originating or receiving traffic from the Internet. This architecture allows you to perform inline inspection of traffic from multiple spoke VPCs in a simplified and scalable fashion while still centralizing your virtual appliances.

upvoted 49 times

jucejara **Most Recent** 11 months, 1 week ago

Selected Answer: D

Esta solución utiliza el servicio Gateway Load Balancer (GWLB), que está diseñado específicamente para inspección de tráfico a través de dispositivos virtuales de terceros, como firewalls, IDS/IPS, etc., con mínima sobrecarga operativa.

Ventajas de GWLB:

Permite escalar automáticamente los dispositivos de inspección.

Se integra fácilmente con appliances de terceros desde AWS Marketplace.

Permite enrutamiento transparente del tráfico a través del dispositivo para inspección.

Puede usarse junto con puntos finales de Gateway Load Balancer (GWLBe) en otras VPC, como la VPC de la aplicación, sin cambios significativos en la arquitectura.

upvoted 2 times

Clouddon 11 months, 2 weeks ago

Selected Answer: D

Use Gateway Load Balancer (GWLB) in the inspection VPC with Gateway Load Balancer Endpoints in the application VPC to efficiently and transparently route traffic through the third-party firewall appliance — achieving deep packet inspection with minimal operational overhead.

upvoted 2 times

surajkrishnamurthy 1 year, 3 months ago

Selected Answer: D

3rd Party Virtual Appliances ---> Gateway Loadbalancer

upvoted 6 times

EllenLiu 1 year, 6 months ago

Selected Answer: D

we're able to load balance some of the incoming traffic through to those virtual appliances where they can perform some kind of inspection.

- Used in front of virtual appliances such as:
 - firewalls,
 - IDS: intrusion detection systems
 - IPS: intrusion prevention systems
 - DPI: deep packet inspection systems
 - Operates at Layer 3 – listens for all packets on all ports
 - Forwards traffic to the TG specified in the listener rules
 - GLB and virtual appliances Exchanges traffic using the GENEVE protocol on port 6081
- upvoted 1 times

carlton_agesa 1 year, 4 months ago

In this context, "using the GENEVE protocol" refers to the method by which the Global Load Balancer (GLB) and the virtual appliances (like firewalls and intrusion detection systems) communicate and exchange traffic. GENEVE (Generic Network Virtualization Encapsulation) is a protocol that encapsulates packets for efficient transport over a network, allowing for flexible and scalable network virtualization. It operates on port 6081, enabling the load balancer to forward traffic to the specified target groups (TG) based on the listener rules.

Credits: AI answer

upvoted 2 times

EzKkk 1 year, 6 months ago

Selected Answer: D

The answer is D for some obvious reasons:

- Low operational overhead.
- Integrate AWS Marketplace and third-party application.

I think the question will be much more interesting if the author add another option like using VPC-to-VPC traffic inspection because the question asked traffic to be inspected before it reaches application layer so traffic forwarding is also a feasible solution.

upvoted 2 times

Buruguduystunstugudunstuy 1 year, 8 months ago

Selected Answer: D

The solution that will meet these requirements with the least operational overhead is D: Deploy a Gateway Load Balancer in the inspection VPC and create a Gateway Load Balancer endpoint to receive the incoming packets and forward the packets to the appliance.

A Gateway Load Balancer is a fully managed service that provides a single point of contact for clients and distributes incoming traffic across multiple targets, such as Amazon Elastic Compute Cloud (EC2) instances and containers, in one or more virtual private clouds (VPCs). You can deploy a Gateway Load Balancer in the inspection VPC and create a Gateway Load Balancer endpoint to receive the incoming packets from the web servers in the application's VPC and forward the packets to the appliance for packet inspection. This will allow you to inspect all traffic to the web application with minimal operational overhead.

upvoted 9 times

JA2018 1 year, 7 months ago

Gateway Load Balancer helps you easily deploy, scale, and manage your third-party virtual appliances. It gives you one gateway for distributing traffic across multiple virtual appliances while scaling them up or down, based on demand. This decreases potential points of failure in your network and increases availability.

You can find, test, and buy virtual appliances from third-party vendors directly in AWS Marketplace. This integrated experience streamlines the deployment process so you see value from your virtual appliances more quickly—whether you want to keep working with your current vendors or try something new.

upvoted 1 times

Buruguduystunstugudunstuy 3 years, 5 months ago

Option A is incorrect because a Network Load Balancer is designed to handle traffic at the connection level and is not suitable for packet inspection

Option B is incorrect because an Application Load Balancer is designed to handle traffic at the request level and is not suitable for packet inspection.

Option C is incorrect because a transit gateway is designed to allow multiple VPCs and on-premises networks to connect to each other, but it is not suitable for packet inspection.

upvoted 11 times

Bang3R 1 year, 8 months ago

D. Deploy a Gateway Load Balancer in the inspection VPC. Create a Gateway Load Balancer endpoint to receive the incoming packets and forward the packets to the appliance.

A Gateway Load Balancer can inspect traffic before forwarding it to a virtual appliance for additional processing. The solution will not require changing the existing architecture and will have the least amount of operational overhead. The appliance can be configured with a specific IP interface to accept IP packets. The Gateway Load Balancer can be configured with an endpoint to route incoming packets to the appliance. The solution ensures all traffic to the web application is inspected before it reaches the web server.

upvoted 3 times

PhucVuu 1 year, 8 months ago

Selected Answer: D

Keywords: Third-party virtual firewall appliance from AWS Marketplace in an inspection VPC -> only Gateway Load Balancer support it

A: Incorrect - Network Load Balancer don't support to route traffic to third-party virtual firewall appliance.

B: Incorrect - Application Load Balancer don't support to route traffic to third-party virtual firewall appliance.

C: Incorrect - Transit Gateway is use as connect center to connect all VPC, Direct Connect Gateway and VPN Connection. Routes Tables in Transit Gateway only limit which VPC can talk to other VPCs.

D: Correct - Gateway Load Balancer support route traffic to third-party virtual firewall appliance in layer 3 that make it different from ALB and NLB.

upvoted 20 times

lht 1 year, 8 months ago

Selected Answer: B

In the scenario described, the web servers, application servers, and database servers are all located within the same VPC. Therefore, a Gateway Load Balancer may not be the most suitable choice for load balancing traffic between them.

Instead, an Application Load Balancer (ALB) would be a better option as it operates at Layer 7 and can inspect traffic at the application layer. This would allow the virtual firewall to inspect traffic before it reaches the web servers, which is the requirement specified in the scenario.

Overall, while a Gateway Load Balancer can be useful in certain scenarios, it is not the best choice for this particular use case. An Application Load Balancer is a better option as it provides the necessary features to integrate the web application with the virtual firewall appliance and inspect all traffic before it reaches the web server.

upvoted 2 times

johndoe 1 year, 8 months ago

Selected Answer: D

Here's why

Traffic enters the service consumer VPC through the internet gateway.

Traffic is sent to the Gateway Load Balancer endpoint, as a result of ingress routing.

Traffic is sent to the Gateway Load Balancer for inspection through the security appliance.

Traffic is sent back to the Gateway Load Balancer endpoint after inspection.

Traffic is sent to the application servers (destination subnet).

<https://docs.aws.amazon.com/elasticloadbalancing/latest/gateway/getting-started.html>

But I ain't completely sure about the least operational overhead.

upvoted 3 times

AlaTaftaf 1 year, 8 months ago

Selected Answer: B

This is the answer of ChatGpt:

Option B is the correct solution because the ALB can be used to redirect traffic to the virtual firewall appliance without requiring any changes to the backend application servers. The ALB can also be configured to send traffic to multiple targets, allowing the architect to perform high availability and load balancing. This solution is easy to implement and manage and does not require any additional components such as transit gateways or gateway load balancers.

Option D is not the optimal solution since Gateway Load Balancer (GWLB) is intended for use with virtual appliances in the cloud, such as firewalls and intrusion prevention systems. However, it adds operational overhead since creating and managing a Gateway Load Balancer requires several components, including an endpoint group and listener.

upvoted 2 times

cookieMr 1 year, 8 months ago

Selected Answer: A

A. Create a Network Load Balancer in the public subnet of the application's VPC to route the traffic to the appliance for packet inspection.

By creating a Network Load Balancer (NLB) in the public subnet, you can configure it to forward incoming traffic to the virtual firewall appliance for inspection. The NLB operates at the transport layer (Layer 4) and can distribute traffic across multiple instances, including the firewall appliance. This allows you to scale the inspection capacity if needed. The NLB can be associated with a target group that includes the IP address of the firewall appliance, directing traffic to it before reaching the web servers.

Option B (Application Load Balancer) is not suitable for this scenario as it operates at the application layer (Layer 7) and does not provide direct

access to the IP packets for inspection.

Option C (Transit Gateway) and option D (Gateway Load Balancer) introduce additional complexity and overhead compared to using an NLB. They are not necessary for achieving the requirement of inspecting traffic to the web application before reaching the web servers.

upvoted 8 times

vipyodha 2 years, 11 months ago

best answer.well explained

upvoted 1 times

Guru4Cloud 1 year, 8 months ago

Selected Answer: D

The correct answer is D.

Here is the explanation:

Option D is correct because a Gateway Load Balancer (GWLB) is a global service, and it can be deployed in any VPC. This means that the GWLB can reach the appliance. Additionally, the GWLB can be configured to forward packets to the appliance for packet inspection.

Option A is incorrect because a Network Load Balancer (NLB) is a regional service, and the appliance is deployed in an inspection VPC. This means that the NLB would not be able to reach the appliance.

Option B is incorrect because an Application Load Balancer (ALB) is a regional service, and the appliance is deployed in an inspection VPC. This means that the ALB would not be able to reach the appliance.

Option C is incorrect because a transit gateway is a global service, and the appliance is deployed in an inspection VPC. This means that the transit gateway would not be able to reach the appliance.

upvoted 11 times

DigitalDanny 1 year, 8 months ago

Selected Answer: D

Gateway Load Balancer (GWLB): GWLB is designed for deploying third-party appliances and provides a scalable and easy way to route traffic through appliances. It operates at the network layer and can handle both TCP and UDP traffic.

Operational Overhead: Deploying a GWLB in the inspection VPC and creating an endpoint involves less operational overhead compared to managing Load Balancers in the application's VPC. It allows for centralized management of the inspection process.

This solution ensures that all traffic is routed through the Gateway Load Balancer for inspection before reaching the web servers, providing a scalable and efficient way to integrate the third-party virtual firewall appliance

upvoted 5 times

Cmtan 1 year, 8 months ago

A and B are wrong, as they don't support cross-VPC traffic routing

Option C -transit gateway attached to VPC,updating route table and configure security groups and network ACLs can accomplish the task.

Meanwhile, Gateway load balancer is designed meant for routing traffic across VPC, but itself alone does not work. All effort mentioned is C are still required. So this is not the least effort?

upvoted 1 times

A company wants to improve its ability to clone large amounts of production data into a test environment in the same AWS Region. The data is stored in Amazon EC2 instances on Amazon Elastic Block Store (Amazon EBS) volumes. Modifications to the cloned data must not affect the production environment. The software that accesses this data requires consistently high I/O performance.

A solutions architect needs to minimize the time that is required to clone the production data into the test environment. Which solution will meet these requirements?

- A. Take EBS snapshots of the production EBS volumes. Restore the snapshots onto EC2 instance store volumes in the test environment.
- B. Configure the production EBS volumes to use the EBS Multi-Attach feature. Take EBS snapshots of the production EBS volumes. Attach the production EBS volumes to the EC2 instances in the test environment.
- C. Take EBS snapshots of the production EBS volumes. Create and initialize new EBS volumes. Attach the new EBS volumes to EC2 instances in the test environment before restoring the volumes from the production EBS snapshots.
- D. Take EBS snapshots of the production EBS volumes. Turn on the EBS fast snapshot restore feature on the EBS snapshots. Restore the snapshots into new EBS volumes. Attach the new EBS volumes to EC2 instances in the test environment.

Most Voted

Correct Answer: D

Community vote distribution

D (94%)

4%

Comments

UWSFish **Highly Voted** 3 years, 7 months ago

Selected Answer: D

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-fast-snapshot-restore.html>

Amazon EBS fast snapshot restore (FSR) enables you to create a volume from a snapshot that is fully initialized at creation. This eliminates the latency of I/O operations on a block when it is accessed for the first time. Volumes that are created using fast snapshot restore instantly deliver all of their provisioned performance.

upvoted 57 times

PhucVuu **Highly Voted** 1 year, 8 months ago

Selected Answer: D

Keywords:

- Modifications to the cloned data must not affect the production environment.
- Minimize the time that is required to clone the production data into the test environment.

A: Incorrect - we can do this But it is not minimize the time as requirement.

B: Incorrect - This approach use same EBS volumes for produciton and test. If we modify test then it will be affected prodution environment.

C: Incorrect - EBS snapshot will create new EBS volumes. It can not restore from existing volumes.

D: Correct - Turn on the EBS fast snapshot restore feature on the EBS snapshots -> no latency on first use

upvoted 46 times

nesyto **Most Recent** 4 months, 3 weeks ago

Selected Answer: D

correct answer is D, fast snapshot although it is available only for 16 Tib support fast snapshot

upvoted 1 times

Nishant27saini 10 months, 4 weeks ago

Selected Answer: D

Option D is right answer

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: D

Use EBS snapshots with Fast Snapshot Restore enabled to create fully-initialized, high-performance EBS volumes for the test environment — ensuring fast, safe, and performant data cloning without affecting production.

upvoted 1 times

se7en102 1 year, 1 month ago

Selected Answer: D

EBS Snapshots: Taking snapshots of the production EBS volumes allows you to create point-in-time backups of the data without impacting the production environment.

Fast Snapshot Restore: Enabling the EBS fast snapshot restore feature allows you to restore the snapshots to new EBS volumes much more quickly than standard restoration. This significantly reduces the time required to clone the data.

New EBS Volumes: By restoring the snapshots into new EBS volumes, you ensure that the test environment is completely isolated from the production environment, allowing for modifications without any risk to production data.

High I/O Performance: EBS volumes provide high I/O performance, which is essential for the software accessing the data in the test environment.

upvoted 1 times

PaulGa 1 year, 10 months ago

Selected Answer: D

Ans D - as per PhucVuu response... what's to debate...

upvoted 1 times

1e22522 1 year, 10 months ago

Selected Answer: D

Ye its d cuh

upvoted 1 times

1dfed2b 2 years, 3 months ago

Selected Answer: C

<https://docs.aws.amazon.com/ebs/latest/userguide/ebs-restoring-volume.html>

Its C. reate a new volume from the snapshot. Use the create-volume command. For --snapshot-id, specify the ID of the snapshot to use. For --availability-zone, specify the same Availability Zone as the instance. Configure the remaining parameters as needed.

upvoted 2 times

Isomas 2 years, 3 months ago

Answer is D because volumes that are created using fast snapshot restore instantly deliver all of their provisioned performance.

Volumes created from normal snapshots will take time to initialize

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: D

A: Can work but long cloning time

B: Wrong as multi attach will mean changes by test will affect production

C: Slow

D: Fast restore makes this a quicker option

upvoted 3 times

A_jaa 2 years, 5 months ago

Selected Answer: D

Answer-D

upvoted 1 times

Ruffyit 2 years, 7 months ago

Amazon EBS fast snapshot restore (FSR) enables you to create a volume from a snapshot that is fully initialized at creation. This eliminates the latency of I/O operations on a block when it is accessed for the first time. Volumes that are created using fast snapshot restore instantly deliver all of their provisioned performance.

upvoted 1 times

ukivanlamlpi 2 years, 10 months ago

Selected Answer: A

why not A? high I/O, no need durability

upvoted 2 times

JackLo 2 years, 9 months ago

Although it is test environment, it's data should be durable

upvoted 3 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: D

Needs to minimize the time that is required to clone the production data into the test environment = EBS fast snapshot restore feature
upvoted 1 times

Anil_Awasthi 2 years, 10 months ago

Selected Answer: C

Option C provides an effective solution for cloning large amounts of production data into a test environment with minimized time, high I/O performance, and without affecting the production environment.

upvoted 1 times

pentium75 2 years, 5 months ago

But you don't need a new, empty volume, you need a restore of the PROD snapshot. Thus D.

upvoted 2 times

Guru4Cloud 2 years, 10 months ago

Selected Answer: D

The correct answer is D.

Here is a step-by-step explanation of how to clone production data into a test environment using EBS snapshots:

Take EBS snapshots of the production EBS volumes.

Turn on the EBS fast snapshot restore feature on the EBS snapshots.

Restore the snapshots into new EBS volumes.

Attach the new EBS volumes to EC2 instances in the test environment.

The EBS fast snapshot restore feature allows you to restore snapshots more quickly than the default method. This is because the feature uses a process called parallel restore, which allows multiple EBS volumes to be restored at the same time.

The EBS fast snapshot restore feature is only available for EBS snapshots that are created in the same AWS Region as the EC2 instances that you are using to restore the snapshots.

upvoted 6 times

An ecommerce company wants to launch a one-deal-a-day website on AWS. Each day will feature exactly one product on sale for a period of 24 hours. The company wants to be able to handle millions of requests each hour with millisecond latency during peak hours.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use Amazon S3 to host the full website in different S3 buckets. Add Amazon CloudFront distributions. Set the S3 buckets as origins for the distributions. Store the order data in Amazon S3.
- B. Deploy the full website on Amazon EC2 instances that run in Auto Scaling groups across multiple Availability Zones. Add an Application Load Balancer (ALB) to distribute the website traffic. Add another ALB for the backend APIs. Store the data in Amazon RDS for MySQL.
- C. Migrate the full application to run in containers. Host the containers on Amazon Elastic Kubernetes Service (Amazon EKS). Use the Kubernetes Cluster Autoscaler to increase and decrease the number of pods to process bursts in traffic. Store the data in Amazon RDS for MySQL.
- D. Use an Amazon S3 bucket to host the website's static content. Deploy an Amazon CloudFront distribution. Set the S3 bucket as the origin. Use Amazon API Gateway and AWS Lambda functions for the backend APIs. Store the data in Amazon DynamoDB. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Sinaneos **Highly Voted** 3 years, 8 months ago

Selected Answer: D

D because all of the components are infinitely scalable
dynamoDB, API Gateway, Lambda, and of course s3+cloudfront
upvoted 52 times

Buruguduystunstugudunstuy **Highly Voted** 3 years, 5 months ago

Selected Answer: D

The solution that will meet these requirements with the least operational overhead is D: Use an Amazon S3 bucket to host the website's static content, deploy an Amazon CloudFront distribution, set the S3 bucket as the origin, and use Amazon API Gateway and AWS Lambda functions for the backend APIs. Store the data in Amazon DynamoDB.

Using Amazon S3 to host static content and Amazon CloudFront to distribute the content can provide high performance and scale for websites with millions of requests each hour. Amazon API Gateway and AWS Lambda can be used to build scalable and highly available backend APIs to support the website, and Amazon DynamoDB can be used to store the data. This solution requires minimal operational overhead as it leverages fully managed services that automatically scale to meet demand.

upvoted 24 times

Buruguduystunstugudunstuy 3 years, 5 months ago

Option A is incorrect because using multiple S3 buckets to host the full website would not provide the required performance and scale for millions of requests each hour with millisecond latency.

Option B is incorrect because deploying the full website on EC2 instances and using an Application Load Balancer (ALB) and an RDS database would require more operational overhead to maintain and scale the infrastructure.

Option C is incorrect because while deploying the application in containers and hosting them on Amazon Elastic Kubernetes Service (EKS) can provide high performance and scale, it would require more operational overhead to maintain and scale the infrastructure compared to using fully managed services like S3 and CloudFront.

upvoted 24 times

try1260 1 year, 6 months ago

How to use dynamodb in a ecommerce trx it required sql database
upvoted 2 times

JustEugen 10 months ago

There are no requirements for the database, and no requirements to how orders should be processed
upvoted 1 times

DevendraAWS6025 **Most Recent** 3 weeks, 1 day ago

Selected Answer: D

Option D as all are server less and Key word is millisecond latency it only achieved using dynamo db.
upvoted 1 times

46e1a46 1 month, 1 week ago

Selected Answer: D

D because
- spiky/unpredictable traffic (flash sale, one-deal-a-day)
- Millions of requests with millisecond latency
LEAST operational overhead

The answer is almost always the fully serverless stack: S3 + CloudFront + API Gateway + Lambda + DynamoDB.
upvoted 1 times

nesyto 4 months, 3 weeks ago

Selected Answer: D

D is the correct answer
upvoted 1 times

rahulben 4 months, 3 weeks ago

Selected Answer: D

Least operation overhead means the option is D only. As other options involve lots of work
upvoted 1 times

BlazerX998 7 months ago

Selected Answer: D

D is the best option because it meets the requirements for handling millions of requests per hour with low latency while having the least operational overhead.
upvoted 1 times

tgrgopi 9 months, 1 week ago

Selected Answer: D

D because all of the components are infinitely scalable
dynamoDB, API Gateway, Lambda, and of course s3+cloudfront
upvoted 1 times

guhaan 9 months, 2 weeks ago

Selected Answer: D

The correct solution is D because we need to store static content in S3 along with CloudFront for CDN but also need Amazon API Gateway and AWS Lambda functions to process the orders in backend.
upvoted 1 times

Dipsha 10 months, 2 weeks ago

Selected Answer: D

Here the option D best suites. Static front-end (S3 + CloudFront) , dynamic backend (API Gateway + Lambda + DynamoDB) — ideal for low-latency, scalable, cost-effective solutions
upvoted 1 times

Dipsha 10 months, 2 weeks ago

Selected Answer: D

yes if question may focus on latency only then answer would be EC2 but here asked for least overhead ho so option goes on S3.
upvoted 1 times

atg_raj 11 months, 1 week ago

Selected Answer: D

I had confusion between B and D as B seemed more reasonable if you want scalability with control but D is just the better solution.
upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: D

Use the S3 + CloudFront + API Gateway + Lambda + DynamoDB architecture (Option D) to deliver a scalable, low-latency, cost-effective "deal of the day" ecommerce platform with minimal infrastructure management.

upvoted 1 times

hossex5 11 months, 4 weeks ago

Selected Answer: D

d is the answer

upvoted 1 times

phenic 1 year, 1 month ago

Selected Answer: D

The correct answer is D . You use s3 to host static content. the cloudfront will CloudFront to distribute the content can provide high performance and scale for websites with millions of requests each hour. Amazon API Gateway and AWS Lambda can be used to build scalable and highly available backend APIs to support the website, and Amazon DynamoDB can be used to store the data. Amazon API Gateway and AWS Lambda can be used to build scalable and highly available backend APIs to support the website, and Amazon DynamoDB can be used to store the data.

upvoted 1 times

AshishDhole 1 year, 4 months ago

Selected Answer: D

Millisecond Latency = DynamoDB

upvoted 3 times

satyaamm 1 year, 5 months ago

Selected Answer: D

D as DynamoDB is only suitable because of its really fast milliseconds speed.

upvoted 1 times

A solutions architect is using Amazon S3 to design the storage architecture of a new digital media application. The media files must be resilient to the loss of an Availability Zone. Some files are accessed frequently while other files are rarely accessed in an unpredictable pattern. The solutions architect must minimize the costs of storing and retrieving the media files.

Which storage option meets these requirements?

- A. S3 Standard
- B. S3 Intelligent-Tiering **Most Voted**
- C. S3 Standard-Infrequent Access (S3 Standard-IA)
- D. S3 One Zone-Infrequent Access (S3 One Zone-IA)

Correct Answer: B

Community vote distribution

B (99%)

Comments

123jh10 **Highly Voted** 3 years, 7 months ago

Selected Answer: B

"unpredictable pattern" - always go for Intelligent Tiering of S3

It also meets the resiliency requirement: "S3 Standard, S3 Intelligent-Tiering, S3 Standard-IA, S3 Glacier Instant Retrieval, S3 Glacier Flexible Retrieval, and S3 Glacier Deep Archive redundantly store objects on multiple devices across a minimum of three Availability Zones in an AWS Region"

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/DataDurability.html>

upvoted 52 times

Buruguduystunstugudunstuy **Highly Voted** 3 years, 5 months ago

Selected Answer: B

The storage option that meets these requirements is B: S3 Intelligent-Tiering.

Amazon S3 Intelligent Tiering is a storage class that automatically moves data to the most cost-effective storage tier based on access patterns. It can store objects in two access tiers: the frequent access tier and the infrequent access tier. The frequent access tier is optimized for frequently accessed objects and is charged at the same rate as S3 Standard. The infrequent access tier is optimized for objects that are not accessed frequently and are charged at a lower rate than S3 Standard.

S3 Intelligent Tiering is a good choice for storing media files that are accessed frequently and infrequently in an unpredictable pattern because it automatically moves data to the most cost-effective storage tier based on access patterns, minimizing storage and retrieval costs. It is also resilient to the loss of an Availability Zone because it stores objects in multiple Availability Zones within a region.

upvoted 19 times

Buruguduystunstugudunstuy 3 years, 5 months ago

Option A, S3 Standard, is not a good choice because it does not offer the cost optimization of S3 Intelligent-Tiering.

Option C, S3 Standard-Infrequent Access (S3 Standard-IA), is not a good choice because it is optimized for infrequently accessed objects and does not offer the cost optimization of S3 Intelligent-Tiering.

Option D, S3 One Zone-Infrequent Access (S3 One Zone-IA), is not a good choice because it is not resilient to the loss of an Availability Zone. It stores objects in a single Availability Zone, making it less durable than other storage classes.

upvoted 8 times

hariteja **Most Recent** 1 month ago

Selected Answer: B

Because some are access frequently and some are rarely

upvoted 1 times

GlassFish 2 months, 4 weeks ago

Selected Answer: B

Unpredictable pattern = Intelligent Tiering

upvoted 1 times

nesyto 4 months, 3 weeks ago

Selected Answer: B

B is the correct answer , reason some of the files are accessed at unpredictable sequence

upvoted 1 times

flac_1 6 months, 3 weeks ago

Selected Answer: D

Intelligent tiering will shuffle the data between tiers. For an unpredictable pattern this would be best. Lowest amount of admin while taking advantage of a key AWS Storage feature

upvoted 1 times

tgrgopi 9 months, 1 week ago

Selected Answer: B

For unpredictable pattern" - always go for Intelligent Tiering of S3

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: B

Use S3 Intelligent-Tiering to get multi-AZ resilience and automatic cost savings for media files with unpredictable access patterns.

upvoted 1 times

LEEOscar 1 year ago

Selected Answer: B

Amazon S3 Intelligent-Tiering:

Automatically moves data between different storage tiers based on access patterns.

Designed for unpredictable access—you don't need to know in advance whether data will be frequently accessed.

Provides multi-AZ durability, just like S3 Standard.

upvoted 1 times

San757 1 year, 1 month ago

Selected Answer: B

The storage option that meets these requirements is B: S3 Intelligent-Tiering.

upvoted 1 times

Gizmo2022 1 year, 6 months ago

Selected Answer: B

I am going to choose B.

S3 Intelligent-Tiering - Perfect use case when you don't know the frequency of access or irregular patterns of usage.

upvoted 1 times

Tjazz04 1 year, 6 months ago

Selected Answer: B

Key phrase: Some files are accessed frequently while other files are rarely accessed in an unpredictable pattern.

Always use S3 Intelligent-Tiering for unpredictable access patterns.

upvoted 1 times

PaulGa 1 year, 10 months ago

Selected Answer: B

Ans B - Intelligent Tiering: cost effective, optimised by access frequency

upvoted 2 times

DavidNgTan 1 year, 10 months ago

Selected Answer: B

S3 intelligent tiering support 3 layer, frequent access, infrequent access and rarely access data.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/intelligent-tiering-overview.html>

upvoted 3 times

lofzee 2 years ago

Unpredictable pattern equals intelligent tiering

upvoted 3 times

JohnZh 2 years, 2 months ago

unpredictable pattern == S3 IA (Intelligent Tiering, not Infrequent access though)
upvoted 3 times

awsgeek75 2 years, 5 months ago

Selected Answer: B

Unpredictable pattern = Intelligent tiering
upvoted 3 times

A company is storing backup files by using Amazon S3 Standard storage. The files are accessed frequently for 1 month. However, the files are not accessed after 1 month. The company must keep the files indefinitely.

Which storage solution will meet these requirements MOST cost-effectively?

A. Configure S3 Intelligent-Tiering to automatically migrate objects.

B. Create an S3 Lifecycle configuration to transition objects from S3 Standard to S3 Glacier Deep Archive after 1 month.

Most Voted

C. Create an S3 Lifecycle configuration to transition objects from S3 Standard to S3 Standard-Infrequent Access (S3 Standard-IA) after 1 month.

D. Create an S3 Lifecycle configuration to transition objects from S3 Standard to S3 One Zone-Infrequent Access (S3 One Zone-IA) after 1 month.

Correct Answer: B

Community vote distribution

B (97%)

2%

Comments

Buruguduystunstugudunstuy **Highly Voted** 3 years, 5 months ago

Selected Answer: B

The storage solution that will meet these requirements most cost-effectively is B: Create an S3 Lifecycle configuration to transition objects from S3 Standard to S3 Glacier Deep Archive after 1 month.

Amazon S3 Glacier Deep Archive is a secure, durable, and extremely low-cost Amazon S3 storage class for long-term retention of data that is rarely accessed and for which retrieval times of several hours are acceptable. It is the lowest-cost storage option in Amazon S3, making it a cost-effective choice for storing backup files that are not accessed after 1 month.

You can use an S3 Lifecycle configuration to automatically transition objects from S3 Standard to S3 Glacier Deep Archive after 1 month. This will minimize the storage costs for the backup files that are not accessed frequently.

upvoted 19 times

Buruguduystunstugudunstuy 3 years, 5 months ago

Option A, configuring S3 Intelligent-Tiering to automatically migrate objects, is not a good choice because it is not designed for long-term storage and does not offer the cost benefits of S3 Glacier Deep Archive.

Option C, transitioning objects from S3 Standard to S3 Standard-Infrequent Access (S3 Standard-IA) after 1 month, is not a good choice because it is not the lowest-cost storage option and would not provide the cost benefits of S3 Glacier Deep Archive.

Option D, transitioning objects from S3 Standard to S3 One Zone-Infrequent Access (S3 One Zone-IA) after 1 month, is not a good choice because it is not the lowest-cost storage option and would not provide the cost benefits of S3 Glacier Deep Archive.

upvoted 7 times

vgchan 3 years, 5 months ago

Also S3 Standard-IA & One Zone-IA stores the data for max of 30 days and not indefinitely.

upvoted 6 times

MatAlves 1 year, 8 months ago

you want the MOST cost-effectively solution. Glacier DA is the cheapest for archival.

upvoted 1 times

JA2018 1 year, 7 months ago

If the AZ selected for the One Zone-IA option is whacked offline, high chances of data loss.

upvoted 1 times

ninjawrz **Highly Voted** 3 years, 7 months ago

B: Transition to Glacier deep archive for cost efficiency
upvoted 9 times

nesyto Most Recent 4 months, 3 weeks ago

Selected Answer: B

B is the correct answer as we want to archive them indefinitely and they will not be accessed
upvoted 1 times

tgrgopi 9 months, 1 week ago

Selected Answer: B

B is the right option
upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: B

Use S3 Lifecycle policies to transition data from Standard to Glacier Deep Archive after 1 month — the most cost-effective and durable solution for rarely accessed backups that need to be retained indefinitely.
upvoted 1 times

phenic 1 year, 1 month ago

Selected Answer: B

Amazon S3 Glacier Deep Archive is a secure, durable, and extremely low-cost Amazon S3 storage class for long-term retention of data that is rarely accessed and for which retrieval times of several hours are acceptable. It is the lowest-cost storage option in Amazon S3, making it a cost-effective choice for storing backup files that are not accessed after 1 month.
upvoted 1 times

Vandaman 1 year, 3 months ago

Selected Answer: B

Definitely B. Files are not accessed after 1 month and need to be kept indefinitely, so Glacier Deep Archive is the best solution.
upvoted 1 times

satyaamm 1 year, 5 months ago

Selected Answer: B

Glacier archive is most suitable here as the data doesn't need to be accessed anymore but stored indefinitely.
upvoted 1 times

LayManCloud_2050 1 year, 7 months ago

I think the answer is C. This is because, the files were accessed frequently for the first month, so transitioning to Glacier Deep Archive after 1 month would result in higher retrieval costs for the first month of frequent access, making it a less optimal solution. However, since the files are accessed frequently for the first month and then not accessed after that, transitioning them to S3 Standard-IA after 1 month is the cost-effective choice.
upvoted 1 times

PaulGa 1 year, 10 months ago

Selected Answer: B

Ans B - Glacier: the files will not be accessed after 1 month; they just need to be retained
upvoted 1 times

DavidNgTan 1 year, 10 months ago

Selected Answer: B

Should apply S3 lifecycle to move not accessed file after 1 month to S3 Glacier Deep Archive.
upvoted 1 times

Solomon2001 2 years, 1 month ago

Selected Answer: C

Since the files are not accessed after 1 month but need to be kept indefinitely, transitioning them to S3 Standard-Infrequent Access (S3 Standard-IA) would be the best choice.
upvoted 1 times

JohnZh 2 years, 2 months ago

not accessed == Glacier -- easy one
upvoted 1 times

awsgeek75 2 years, 5 months ago

Selected Answer: B

A: Possible but expensive
CD: One zone so no guarantee of being stored indefinitely.
B: S3GDA is cost effective indefinite storage

upvoted 1 times

A_jaa 2 years, 5 months ago

Selected Answer: B

Answer-B

upvoted 1 times

RNess 2 years, 5 months ago

It's can't be B!!

because objects that are archived to S3 Glacier Instant Retrieval and S3 Glacier Flexible Retrieval are charged for a minimum storage duration of 90 days, and S3 Glacier Deep Archive has a minimum storage duration of 180 days.

upvoted 2 times

RichWil 1 year, 3 months ago

In AWS S3 Standard, there is no minimum storage duration requirement before you can move an object to S3 Glacier Deep Archive. You can transition objects at any time using S3 Lifecycle Policies.

However, once an object is stored in S3 Glacier Deep Archive, it has a minimum storage duration of 180 days. If the object is deleted or overwritten before 180 days, you will be charged for the full 180-day period.

upvoted 1 times

pentium75 2 years, 5 months ago

But the items must be kept forever, so where's the issue with that?

upvoted 1 times

RNess 2 years, 5 months ago

I mean, that min duration ****before**** can move to S3 Glacier Flexible Retrieval or S3 Glacier Deep Archive

upvoted 2 times

ddement0r 2 years, 6 months ago

Selected Answer: B

B because since the files should be kept but never accessed we can put them in Deep Archive

upvoted 1 times

A company observes an increase in Amazon EC2 costs in its most recent bill. The billing team notices unwanted vertical scaling of instance types for a couple of EC2 instances. A solutions architect needs to create a graph comparing the last 2 months of EC2 costs and perform an in-depth analysis to identify the root cause of the vertical scaling.

How should the solutions architect generate the information with the LEAST operational overhead?

- A. Use AWS Budgets to create a budget report and compare EC2 costs based on instance types.
- B. Use Cost Explorer's granular filtering feature to perform an in-depth analysis of EC2 costs based on instance types.
- C. Use graphs from the AWS Billing and Cost Management dashboard to compare EC2 costs based on instance types for the last 2 months.
- D. Use AWS Cost and Usage Reports to create a report and send it to an Amazon S3 bucket. Use Amazon QuickSight with Amazon S3 as a source to generate an interactive graph based on instance types.

Most Voted

Correct Answer: B

Community vote distribution

B (66%) C (25%) 10%

Comments

sba21 **Highly Voted** 3 years, 8 months ago

Selected Answer: B

<https://www.examtopycs.com/discussions/amazon/view/68306-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 40 times

123jh10 **Highly Voted** 3 years, 7 months ago

Selected Answer: C

The requested result is a graph, so...

A - can't be as the result is a report

B - can't be as it is limited to 14 days visibility and the graph has to cover 2 months

C - seems to provide graphs and the best option available, as...

D - could provide graphs, BUT involves operational overhead, which has been requested to be minimised.

upvoted 33 times

satyaammm 1 year, 5 months ago

Hey brother, I think you might be wrong. The AWS Cost Explorer retains for 12 months and the 14 days visibility is actually for the free trial. So it has nothing related to how much the cost explorer can retain. I hope I made myself clear.

upvoted 3 times

lofzee 3 years, 4 months ago

14 days? Fam, you ever logged into the console?

upvoted 23 times

ManikRoy 2 years, 2 months ago

Never be over-confident! check the below AWS link, which mentions the below.

Cost Explorer allows customers to understand their spend for the past 14 days at the most granular hourly level. You can monitor your hourly cost and usage to determine appropriate SP/RI commitment or understand peak hour usage.

<https://aws.amazon.com/aws-cost-management/aws-cost-explorer/features/#:~:text=Cost%20Explorer%20allows%20customers%20to,or%20understand%20peak%20hour%20usage.>

upvoted 5 times

Mr_Marcus 2 years ago

Neither the requirement nor answer B state the need for the "most granular hourly level." Correct answer is B.

upvoted 2 times

lofzee 2 years ago

The answer is B because least operational overhead

upvoted 1 times

Ello2023 3 years, 4 months ago

B. This is correct because there is no limit of 14 days. Quoted from Amazon "AWS prepares the data about your costs for the current month and the last 12 months, and then calculates the forecast for the next 12 months." (<https://aws.amazon.com/aws-cost-management/aws-cost-explorer/>).

upvoted 11 times

Udoyen 3 years, 6 months ago

Cost Explorer, AWS prepares the data about your costs for the current month and the last 12 months: <https://aws.amazon.com/aws-cost-management/aws-cost-explorer/>

upvoted 17 times

JA2018 1 year, 7 months ago

<https://docs.aws.amazon.com/cost-management/latest/userguide/differences-billing-data-cost-explorer-data.html>

Cost Explorer supports deep-dive analysis so that you can identify savings opportunities. Cost Explorer data provides more granular dimensions (such as Availability Zone or operating system) and includes features that might show differences when compared to billing data. On the Cost Management preferences page, you can manage your preferences for Cost Explorer data, including linked account access and historical and granular data settings. For more information, see Controlling access to Cost Explorer.

upvoted 1 times

nesyto **Most Recent** 4 months, 3 weeks ago

Selected Answer: B

Cost explorer is the answer

upvoted 1 times

rahulben 4 months, 3 weeks ago

Selected Answer: B

Answer: B

Why:

Need to compare last 2 months of EC2 costs by instance types

Need in-depth analysis to identify root cause

Cost Explorer's granular filtering is the LEAST operational overhead—it's built-in, no extra setup needed

Option A requires manual budget setup

Option C is limited to dashboard graphs (less detailed)

Option D adds unnecessary complexity (S3, QuickSight setup = more overhead)

upvoted 1 times

tgrgopi 9 months, 1 week ago

Selected Answer: B

B is the right answer

upvoted 1 times

607d657 10 months, 3 weeks ago

Selected Answer: B

- AWS Cost Explorer provides built-in visualizations and granular filtering by dimensions like instance type, linked account, region, and even EC2 instance ID

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: B

Use AWS Cost Explorer for an interactive, low-effort, and detailed view of EC2 cost trends by instance type — perfect for identifying unwanted vertical scaling.

✗ Why the Other Options Are Not Optimal:

A. AWS Budgets

Budgets help track and alert on cost thresholds, not analyze past cost trends in detail.

It won't show detailed instance-type usage or cost breakdowns over time.

C. Billing and Cost Management dashboard

Provides summary-level graphs, not granular filtering by EC2 instance types.

Less useful for in-depth root cause analysis of instance size changes.

D. Cost and Usage Reports + S3 + QuickSight

While powerful, this setup involves significant operational overhead:

Creating reports,

Managing S3 buckets,

Configuring QuickSight data sources and dashboards.

Overkill for a quick comparison and analysis.

upvoted 1 times

jp 11 months, 2 weeks ago

Selected Answer: B

Cost Explorer always you to look back for a year or so

upvoted 1 times

jp 11 months, 2 weeks ago

allows*

upvoted 1 times

Gongs 1 year, 1 month ago

Selected Answer: C

C cause Timescale in b is wrong

upvoted 1 times

Juju66 1 year, 1 month ago

Selected Answer: B

C can't provide granular filter.

upvoted 1 times

Vandaman 1 year, 3 months ago

Selected Answer: B

I see lately that you can look back as far as 3 years.

upvoted 1 times

AshishDhole 1 year, 4 months ago

Selected Answer: B

Cost Explorer, AWS prepares the data about your costs for the current month and the last 12 months, and then calculates the forecast for the next 12 months. The current month's data is available for viewing in about 24 hours. The rest of your data takes a few days longer. Cost Explorer updates your cost data at least once every 24 hours.

upvoted 1 times

V2910 1 year, 4 months ago

Selected Answer: B

Cost explorer has option to filter by service and we can modify the timeline also to compare the costs across 3 months

upvoted 1 times

aatikah 1 year, 6 months ago

Selected Answer: B

AWS Cost Explorer allows you to view cost and usage data with granular filtering. However, the historical data retention depends on the level of granularity:

Daily granularity: You can access up to the past 12 months of daily cost and usage data.

Hourly granularity: If you have detailed billing enabled, you can view up to the past 14 days of hourly usage data.

Monthly granularity: You can access up to the past 12 months of cost and usage data, and also forecast costs for the next 12 months.

upvoted 3 times

mzeynalli 1 year, 7 months ago

Selected Answer: B

B. Use Cost Explorer's granular filtering feature to perform an in-depth analysis of EC2 costs based on instance types.

Explanation:

Cost Explorer provides an easy way to analyze AWS costs and usage visually. It allows you to filter data by multiple parameters, such as instance types, regions, and time periods. It also enables you to drill down and identify specific cost drivers, such as unwanted vertical scaling.

This option is suitable because it provides granular insights with minimal operational overhead, as it is a built-in AWS tool specifically designed for cost analysis.

upvoted 4 times

mzeynalli 1 year, 7 months ago

The correct answer is:

B. Use Cost Explorer's granular filtering feature to perform an in-depth analysis of EC2 costs based on instance types.

Explanation:

Cost Explorer provides an easy way to analyze AWS costs and usage visually. It allows you to filter data by multiple parameters, such as instance types, regions, and time periods. It also enables you to drill down and identify specific cost drivers, such as unwanted vertical scaling. This option is suitable because it provides granular insights with minimal operational overhead, as it is a built-in AWS tool specifically designed for cost analysis.

Why Other Options Are Less Suitable:

C. AWS Billing and Cost Management Dashboard: It provides a high-level view of costs, but it lacks the specific filtering capabilities needed for a granular, instance-type-level analysis over time.

upvoted 2 times

PaulEkwem 1 year, 8 months ago

Option B:

Use AWS Cost Explorer, a tool that makes it easy to view your costs and break them down into categories like EC2 instance types, time periods, and more.

With Cost Explorer, you can easily filter the information to see which EC2 instance types caused the increase in costs over the last two months. This lets you do an in-depth analysis without much effort.

upvoted 3 times

A company is designing an application. The application uses an AWS Lambda function to receive information through Amazon API Gateway and to store the information in an Amazon Aurora PostgreSQL database.

During the proof-of-concept stage, the company has to increase the Lambda quotas significantly to handle the high volumes of data that the company needs to load into the database. A solutions architect must recommend a new design to improve scalability and minimize the configuration effort.

Which solution will meet these requirements?

- A. Refactor the Lambda function code to Apache Tomcat code that runs on Amazon EC2 instances. Connect the database by using native Java Database Connectivity (JDBC) drivers.
- B. Change the platform from Aurora to Amazon DynamoDB. Provision a DynamoDB Accelerator (DAX) cluster. Use the DAX client SDK to point the existing DynamoDB API calls at the DAX cluster.
- C. Set up two Lambda functions. Configure one function to receive the information. Configure the other function to load the information into the database. Integrate the Lambda functions by using Amazon Simple Notification Service (Amazon SNS).
- D. Set up two Lambda functions. Configure one function to receive the information. Configure the other function to load the information into the database. Integrate the Lambda functions by using an Amazon Simple Queue Service (Amazon SQS) queue. **Most Voted**

Correct Answer: D

Community vote distribution

D (99%)

Comments

123jh10 **Highly Voted** 3 years, 7 months ago

Selected Answer: D

A - refactoring can be a solution, BUT requires a LOT of effort - not the answer

B - DynamoDB is NoSQL and Aurora is SQL, so it requires a DB migration... again a LOT of effort, so no the answer

C and D are similar in structure, but...

C uses SNS, which would notify the 2nd Lambda function... provoking the same bottleneck... not the solution

D uses SQS, so the 2nd lambda function can go to the queue when responsive to keep with the DB load process.

Usually the app decoupling helps with the performance improvement by distributing load. In this case, the bottleneck is solved by uses queues... so D is the answer.

upvoted 107 times

PhucVuu **Highly Voted** 3 years, 2 months ago

Selected Answer: D

Keywords:

- Company has to increase the Lambda quotas significantly to handle the high volumes of data that the company needs to load into the database.

- Improve scalability and minimize the configuration effort.

A: Incorrect - Lambda is Serverless and automatically scale - EC2 instance we have to create load balancer, auto scaling group,... a lot of things. using native Java Database Connectivity (JDBC) drivers don't improve the performance.

B: Incorrect - a lot of things to changes and DynamoDB Accelerator use for cache(read) not for write.

C: Incorrect - SNS is use for send notification (e-mail, SMS).

D: Correct - with SQS we can scale application well by queuing the data.

upvoted 21 times

nesyto **Most Recent** 4 months, 3 weeks ago

Selected Answer: D

correct answer is D

upvoted 1 times

tgrgopi 9 months, 1 week ago

Selected Answer: D

Using an SQS queue to decouple the API-ingest Lambda from the DB-writer Lambda gives you automatic buffering, smoothing of traffic spikes, and built-in retry/backoff semantics with minimal configuration. It keeps a serverless architecture (no EC2), requires little operational work, and avoids changing the database.

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: D

Use SQS as a buffer between two Lambda functions to build a scalable, decoupled, and resilient pipeline for ingesting and storing data into Aurora PostgreSQL with minimal configuration effort.

upvoted 1 times

PaulGa 1 year, 10 months ago

Selected Answer: D

It could be C or D, but Ans D wins because

–SNS is Push Mechanism - 'Other Lambda' function is forced to take message when it might not be ready (or refuse it)

–SQS is Pull Mechanism – 'Other Lambda' function can take next message when its ready to do so

SQS is simple and allows better de-coupling.

upvoted 2 times

rodrigoleoncio 2 years, 1 month ago

Selected Answer: D

uses SQS

upvoted 1 times

JohnZh 2 years, 2 months ago

If the throughput is so high that lambda concurrency needs to go beyond 1000, we need to set up a queue to throttle the request.

upvoted 3 times

TheFivePips 2 years, 3 months ago

Selected Answer: D

D. Set up two Lambda functions, one for receiving information and another for loading data into the database. Integrate them using an Amazon SQS queue. This approach allows for better scalability, maintains the serverless paradigm, and minimizes manual configuration effort. It leverages Amazon SQS as a reliable message queue between Lambda functions.

Options A and B introduce complexities and changes in architecture, while Option C introduces an additional service that may not be as suitable for decoupling processes in this scenario.

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: D

SQS will help Lambda scale even more.

A EC2 + Tomcat will be slower than Lambda for this usecase

B is wrong because the problem is with Lambda scaling not the DB

C SNS is not the best option for this usecase when SQS is an option

upvoted 5 times

A_jaa 2 years, 5 months ago

Selected Answer: D

Answer-D

upvoted 1 times

ddement0r 2 years, 6 months ago

Selected Answer: D

D : other ones just don't make sense

upvoted 1 times

pedestrianlove 2 years, 6 months ago

Sorry, but the question does not make sense by itself. What are you asking for more scalability from an already scalable Lambda function?

If you're concerned about the concurrency limits of Lambda function, decoupling just doesn't make sense, since it'll keep even more lambda instances running in a given time period(including 2 phases of execution for each request, let alone the cold start issues).

If you're concerned about bottleneck database induced, that'll even be more ridiculous since you're supposed to resolve the scalability issue of the database(e.g. Aurora) instead of decoupling the Lambda function to improve the throughput of this entire data flow.

upvoted 2 times

mohamedsambo 2 years, 5 months ago

i think it is clear that he want to enhance the lambda even more than "The default concurrency limit across all functions per region in a given account is 1,000"

cause sqs can scale and store the data till new available revoked lambda consume it

upvoted 3 times

Ruffyt 2 years, 7 months ago

Lambda and SQS are serverless. No involvement will be required in execution.

upvoted 2 times

xdkonorek2 2 years, 7 months ago

Selected Answer: B

I think B would be better solution.

How splitting one function into 2 increase scalability when company already increased service quota? Effectively they will have same compute time.

Changing Aurora to DAX will shorten the time for data loads by ~100x requiring way less time for data loading, and it's most time consuming thing this lambda does. DAX has better scaling than aurora and is better fit with lambda

upvoted 2 times

MakaylaLearns 2 years, 9 months ago

Lambda Functions: A review

Run your code in response to events

You can build chatbots using Lambda functions to process user input, execute business logic, and generate responses.

Scales automatically

They can be triggered in response to API events

Lambda functions can process files as they are uploaded to S3 buckets. This is often used for tasks like image resizing, data extraction, or file validation.

upvoted 1 times

doujones 2 years, 10 months ago

Do you all have to take the whole practice exam on here, in order to pass AWS SAA C03

upvoted 2 times

A company needs to review its AWS Cloud deployment to ensure that its Amazon S3 buckets do not have unauthorized configuration changes.

What should a solutions architect do to accomplish this goal?

- A. Turn on AWS Config with the appropriate rules. **Most Voted**
- B. Turn on AWS Trusted Advisor with the appropriate checks.
- C. Turn on Amazon Inspector with the appropriate assessment template.
- D. Turn on Amazon S3 server access logging. Configure Amazon EventBridge (Amazon Cloud Watch Events).

Correct Answer: A

Community vote distribution

A (97%) 2%

Comments

Buruguduystunstugudunstuy **Highly Voted** 3 years, 5 months ago

Selected Answer: A

The solution that will accomplish this goal is A: Turn on AWS Config with the appropriate rules.

AWS Config is a service that enables you to assess, audit, and evaluate the configurations of your AWS resources. You can use AWS Config to monitor and record changes to the configuration of your Amazon S3 buckets. By turning on AWS Config and enabling the appropriate rules, you can ensure that your S3 buckets do not have unauthorized configuration changes.

upvoted 68 times

Buruguduystunstugudunstuy 3 years, 5 months ago

AWS Trusted Advisor (Option B) is a service that provides best practice recommendations for your AWS resources, but it does not monitor or record changes to the configuration of your S3 buckets.

Amazon Inspector (Option C) is a service that helps you assess the security and compliance of your applications. While it can be used to assess the security of your S3 buckets, it does not monitor or record changes to the configuration of your S3 buckets.

Amazon S3 server access logging (Option D) enables you to log requests made to your S3 bucket. While it can help you identify changes to your S3 bucket, it does not monitor or record changes to the configuration of your S3 bucket.

upvoted 49 times

gokalpkoer3 **Highly Voted** 3 years, 7 months ago

Configuration changes= AWS Config

upvoted 36 times

nesyto **Most Recent** 4 months, 3 weeks ago

Selected Answer: A

AWS config checks for config changes

upvoted 1 times

sankalparava 6 months, 1 week ago

Selected Answer: D

Why not option D

upvoted 1 times

jucejara 11 months, 1 week ago

Selected Answer: A

AWS Config es el servicio específicamente diseñado para monitorear, registrar y evaluar configuraciones de recursos de AWS, como los buckets de S3. Permite:

Detectar cambios no autorizados o accidentales en configuraciones.

Evaluar continuamente el cumplimiento con políticas específicas (por ejemplo, "los buckets no deben ser públicos").

Configurar reglas predefinidas o personalizadas para generar alertas o remediaciones automáticas.

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: A

Use AWS Config with S3-specific rules to detect, record, and respond to unauthorized configuration changes, ensuring secure and compliant S3 bucket configurations.

upvoted 1 times

ellanos 1 year, 1 month ago

Selected Answer: A

A is the most suitable option

upvoted 1 times

Juju66 1 year, 1 month ago

Selected Answer: A

c - AWS Inspector is only for EC2 instances

upvoted 2 times

satyaamm 1 year, 5 months ago

Selected Answer: A

AWS Config is the only one that provides continuous monitoring here.

upvoted 1 times

MGKYAING 1 year, 5 months ago

Selected Answer: A

To review an AWS deployment and ensure that Amazon S3 buckets do not have unauthorized configuration changes, you need a service that can monitor and record the configuration state of AWS resources and evaluate them against desired configurations.

upvoted 1 times

PaulGa 1 year, 10 months ago

Selected Answer: A

Ans A - as well explained by "Buruguduystunstugudunstuy" – we are dealing with configuration here: ensuring that what we've designed continues to follow the rules

upvoted 3 times

andyngkh86 2 years, 4 months ago

ChatGPT give the answer is D

upvoted 2 times

A_jaa 2 years, 5 months ago

Selected Answer: A

Answer-A

upvoted 1 times

Ruffyt 2 years, 7 months ago

A: <https://aws.amazon.com/config/#:~:text=How%20it%20works-,AWS%20Config,-continually%20assesses%2C%20audits>

upvoted 2 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: A

AWS Config continually assesses, audits, and evaluates the configurations and relationships of your resources on AWS, on premises, and on other clouds. It normalizes changes into a consistent format and checks resource compliance with custom and managed rules before and after provisioning.

<https://aws.amazon.com/config/#:~:text=How%20it%20works-,AWS%20Config,-continually%20assesses%2C%20audits>

upvoted 4 times

Guru4Cloud 2 years, 10 months ago

Selected Answer: A

AWS Config provides a detailed inventory of the company's AWS resources and configuration history, and can be configured with rules to evaluate resource configurations for compliance with policies and best practices.

The solutions architect can enable AWS Config and configure rules specifically checking for S3 bucket settings like public access blocking, encryption settings, access control lists, etc. AWS Config will record configuration changes to S3 buckets over time, allowing the company to review changes and

be alerted about any unauthorized modifications.

By: Claude.ai

upvoted 2 times

miki111 2 years, 11 months ago

Option A is the right answer for this.

upvoted 1 times

A company is launching a new application and will display application metrics on an Amazon CloudWatch dashboard. The company's product manager needs to access this dashboard periodically. The product manager does not have an AWS account. A solutions architect must provide access to the product manager by following the principle of least privilege. Which solution will meet these requirements?

A. Share the dashboard from the CloudWatch console. Enter the product manager's email address, and complete the sharing steps. Provide a shareable link for the dashboard to the product manager. **Most Voted**

B. Create an IAM user specifically for the product manager. Attach the CloudWatchReadOnlyAccess AWS managed policy to the user. Share the new login credentials with the product manager. Share the browser URL of the correct dashboard with the product manager.

C. Create an IAM user for the company's employees. Attach the ViewOnlyAccess AWS managed policy to the IAM user. Share the new login credentials with the product manager. Ask the product manager to navigate to the CloudWatch console and locate the dashboard by name in the Dashboards section.

D. Deploy a bastion server in a public subnet. When the product manager requires access to the dashboard, start the server and share the RDP credentials. On the bastion server, ensure that the browser is configured to open the dashboard URL with cached AWS credentials that have appropriate permissions to view the dashboard.

Correct Answer: A

Community vote distribution

A (72%)

B (27%)

Comments

masetromain **Highly Voted** 3 years, 8 months ago

Selected Answer: A

Answer A : <https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/cloudwatch-dashboard-sharing.html>

Share a single dashboard and designate specific email addresses of the people who can view the dashboard. Each of these users creates their own password that they must enter to view the dashboard.

upvoted 98 times

mn2013 2 years, 4 months ago

But this link also says All people who you share the dashboard with are granted these permissions for the account. If you share the dashboard publicly, then everyone who has the link to the dashboard has these permissions.

The cloudwatch:GetMetricData and ec2:DescribeTags permissions cannot be scoped down to specific metrics or EC2 instances, so the people with access to the dashboard can query all CloudWatch metrics and the names and tags of all EC2 instances in the account. If that is the case, how is the least privilege principle applicable?

upvoted 5 times

Yak_Yeti 1 year, 2 months ago

while CloudWatch dashboard sharing doesn't perfectly align with the strictest interpretation of least privilege, it can be made more secure by carefully designing the dashboard, implementing strong governance, and educating the user. If the data within the dashboard is highly sensitive, then creating a custom application would be the best method.

Therefore, while option A is the most efficient and simple, it is important to understand the security trade offs that are being made.

upvoted 1 times

123jh10 3 years, 7 months ago

Thanks for the link! No doubt A is the answer.

upvoted 10 times

Guru4Cloud **Highly Voted** 2 years, 10 months ago

Selected Answer: B

Option B provides the product manager with specific access to the CloudWatch dashboard using an IAM user with the CloudWatchReadOnlyAccess policy attached. The IAM user has only read-only access to the required resources, which follows the principle of least privilege.

upvoted 21 times

emilyhu08 2 years, 8 months ago

B has a problem for cloudwatchreadonlyaccess policy, it's not only grant read access to dashboard, but other read permission for logs, insights, etc. so it does not follow the principle of least privilege. Option A only grants access to dashboard.

upvoted 22 times

nesyto **Most Recent** 4 months, 3 weeks ago

Selected Answer: A

correct A, you can share cloud watch dashboard with up to 5 users via email

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: A

Use the CloudWatch dashboard sharing feature to securely share a view-only link with the product manager — no AWS credentials required, and aligned with least privilege and low operational overhead.

upvoted 2 times

jppush 11 months, 2 weeks ago

Selected Answer: A

B won't be the answer because it doesn't follow the least access policy. A is the correct answer

upvoted 1 times

Juju66 1 year, 1 month ago

Selected Answer: A

This is new function for CloudWatch now.

upvoted 2 times

Vandaman 1 year, 3 months ago

Selected Answer: A

Answer is A. This option is available in Cloudwatch

upvoted 1 times

Nacho222222 1 year, 3 months ago

Selected Answer: B

CloudWatch does not allow public links by default, but you can share temporary access using AWS IAM Identity Center (formerly AWS SSO) or a federated user with specific permissions.

upvoted 1 times

Krenil 1 year, 4 months ago

Selected Answer: B

Option - B:

Here's why:

Principle of Least Privilege: The product manager should only be given the permissions necessary to access the CloudWatch dashboard. The CloudWatchReadOnlyAccess policy grants the least privilege access to view CloudWatch metrics and dashboards.

IAM User: Since the product manager does not have an AWS account, creating a specific IAM user for the product manager ensures they can log in securely to AWS, without needing to share broader permissions.

Simple Access: After the IAM user is created, the product manager can directly log into AWS, navigate to the CloudWatch dashboard, and access the metrics they need, without excessive complexity or unnecessary access.

upvoted 1 times

testiaced 1 year, 3 months ago

"To view a dashboard that you create, users must have the required permissions, or the dashboard must be shared across AWS accounts. You can't just email a link to a user who isn't signed into AWS and allow them to view the dashboard."

upvoted 1 times

AfricanCloudGuru 1 year, 4 months ago

Selected Answer: A

Since the Product manager access the account periodically, sharing the account will be the best option

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/cloudwatch-dashboard-sharing.html>

upvoted 1 times

MGKYAING 1 year, 5 months ago

Selected Answer: A

CloudWatch allows dashboards to be shared publicly or with specific users using a shareable URL. This enables access without requiring an AWS account.

This approach follows the principle of least privilege because the product manager is given only the ability to view the dashboard via the shared link, without needing broader access to the AWS environment.

This is the correct choice because it directly meets the requirements without granting unnecessary permissions or requiring the creation of AWS credentials.

upvoted 3 times

MGKYAING 1 year, 5 months ago

Selected Answer: A

1. Direct Sharing Without AWS Account:

Amazon CloudWatch allows dashboards to be shared externally via a shareable link. This enables users without an AWS account to access the dashboard securely, which meets the requirement to allow the product manager (who doesn't have an AWS account) to view the dashboard.

2. Least Privilege Principle:

The product manager only gets access to the specific dashboard, and no additional permissions or access to the AWS environment are granted. This aligns with the principle of least privilege.

3. No Operational Overhead:

The solution does not require creating new IAM users or setting up infrastructure like bastion servers. It's straightforward and efficient.

upvoted 1 times

EzKkk 1 year, 6 months ago

Selected Answer: B

I will go with B even though A sounds like a better choice for the following reasons:

- 1, You need a solution to share your dashboard.
- 2, You only share a specific metric of an application to whom lack access to AWS.
- 3, Target in question only access periodically.
- 4, Must apply least privilege.

Though A & B are good choices to share you dashboard, A doesn't follow least privilege principal because it provides more permission than needed. Not to mention this solution also raises concerns about security since this dashboard is public.

upvoted 2 times

f51a8bd 1 year, 7 months ago

respuesta es A:

Esta solución permite cumplir con los requisitos de acceso seguro y controlado sin una cuenta completa de AWS y sigue el principio de mínimo privilegio para acceder solo a los paneles necesarios en Amazon CloudWatch.

upvoted 1 times

mzeynali 1 year, 7 months ago

Selected Answer: A

The correct answer is:

A. Share the dashboard from the CloudWatch console. Enter the product manager's email address, and complete the sharing steps. Provide a shareable link for the dashboard to the product manager.

Explanation:

Principle of Least Privilege: The solution should only provide the minimum necessary access. Sharing a CloudWatch dashboard directly meets this principle without the need for an AWS account.

Shareable Link: CloudWatch allows you to share a dashboard using a shareable URL that can be accessed without requiring AWS credentials. This way, the product manager can view the dashboard without having an AWS account.

No IAM User Required: Solutions involving IAM users (options B and C) require creating an AWS identity and sharing credentials, which introduces unnecessary security risks and management overhead.

No Bastion Server: Deploying a bastion server (option D) is not practical, involves additional costs, and requires managing another infrastructure component, which goes against simplicity and security best practices.

upvoted 3 times

mzeynali 1 year, 7 months ago

The correct answer is:

A. Share the dashboard from the CloudWatch console. Enter the product manager's email address, and complete the sharing steps. Provide a shareable link for the dashboard to the product manager.

Explanation:

Principle of Least Privilege: The solution should only provide the minimum necessary access. Sharing a CloudWatch dashboard directly meets this principle without the need for an AWS account.

Shareable Link: CloudWatch allows you to share a dashboard using a shareable URL that can be accessed without requiring AWS credentials. This way, the product manager can view the dashboard without having an AWS account.

No IAM User Required: Solutions involving IAM users (options B and C) require creating an AWS identity and sharing credentials, which introduces unnecessary security risks and management overhead.

No Bastion Server: Deploying a bastion server (option D) is not practical, involves additional costs, and requires managing another infrastructure component, which goes against simplicity and security best practices.

upvoted 2 times

e854fbe 1 year, 8 months ago

Selected Answer: B

correct answer B

the reason A is not a good answer is:

when sharing a link to the CloudWatch dashboard the following warning appears

We recommend that you do not share dashboards if your account contains any sensitive information which you would not wish to share with the users with whom you are sharing the dashboard.

The users that you specified above will be granted the following permissions: CloudWatch read-only permissions to alarms and contributor insights rules in the Dashboard which you share, and to all metrics and the names and tags of all EC2 instances in your account even if they are not shown in the Dashboard which you share. We recommend that you consider whether it is appropriate to make this information available to the users with whom you are sharing.

so following the least privilege principle, creating an IAM User in option B is more secure

upvoted 3 times

A company is migrating applications to AWS. The applications are deployed in different accounts. The company manages the accounts centrally by using AWS Organizations. The company's security team needs a single sign-on (SSO) solution across all the company's accounts. The company must continue managing the users and groups in its on-premises self-managed Microsoft Active Directory.

Which solution will meet these requirements?

A. Enable AWS Single Sign-On (AWS SSO) from the AWS SSO console. Create a one-way forest trust or a one-way domain trust to connect the company's self-managed Microsoft Active Directory with AWS SSO by using AWS Directory Service for Microsoft Active Directory.

B. Enable AWS Single Sign-On (AWS SSO) from the AWS SSO console. Create a two-way forest trust to connect the company's self-managed Microsoft Active Directory with AWS SSO by using AWS Directory Service for Microsoft Active Directory.


Most Voted

C. Use AWS Directory Service. Create a two-way trust relationship with the company's self-managed Microsoft Active Directory.

D. Deploy an identity provider (IdP) on premises. Enable AWS Single Sign-On (AWS SSO) from the AWS SSO console.

Correct Answer: B

Community vote distribution



Comments

17Master **Highly Voted** 3 years, 7 months ago

Selected Answer: B

Tricky question!!! forget one-way or two-way. In this scenario, AWS applications (Amazon Chime, Amazon Connect, Amazon QuickSight, AWS Single Sign-On, Amazon WorkDocs, Amazon WorkMail, Amazon WorkSpaces, AWS Client VPN, AWS Management Console, and AWS Transfer Family) need to be able to look up objects from the on-premises domain in order for them to function. This tells you that authentication needs to flow both ways. This scenario requires a two-way trust between the on-premises and AWS Managed Microsoft AD domains.

It is a requirement of the application

Scenario 2: <https://aws.amazon.com/es/blogs/security/everything-you-wanted-to-know-about-trusts-with-aws-managed-microsoft-ad/>
upvoted 80 times

mohamedsambo 2 years, 5 months ago

AWS IAM Identity Center requires a two-way trust so that it has permissions to read user and group information from your domain to synchronize user and group metadata. IAM Identity Center uses this metadata when assigning access to permission sets or applications. User and group metadata is also used by applications for collaboration, like when you share a dashboard with another user or group. The trust from AWS Directory Service for Microsoft Active Directory to your domain permits IAM Identity Center to trust your domain for authentication. The trust in the opposite direction grants AWS permissions to read user and group metadata.

upvoted 8 times

pbpally 3 years, 1 month ago

The problem with this is that nowhere in the question is it saying that the application needs to be able to flow back so two-way is not needed.
upvoted 5 times

aatikah 1 year, 6 months ago

Application Context in This Scenario:

AWS Applications Mentioned in the Scenario:

The question focuses on AWS Single Sign-On (AWS SSO) and the use of SSO across AWS accounts.

AWS SSO with AWS Organizations does not require AWS applications to look up or access on-premises AD objects beyond authentication.

Blog Reference:

The blog mentions the need for a two-way trust if AWS services like Amazon WorkSpaces or Amazon WorkMail require access to user attributes stored in the on-premises AD.

This scenario does not mention such AWS applications, so a one-way trust is sufficient for the SSO use case.

upvoted 2 times

pbpally 3 years, 1 month ago

What I did find though was documentation that explicitly states that IAM Identity Center (successor to AWS SSO) requires a two-way trust:
https://docs.aws.amazon.com/directoryservice/latest/admin-guide/ms_ad_setup_trust.html

upvoted 12 times

KADSM **Highly Voted** 3 years, 7 months ago

Answer B as we have AWS SSO which requires two way trust. As per documentation - A two-way trust is required for AWS Enterprise Apps such as Amazon Chime, Amazon Connect, Amazon QuickSight, AWS IAM Identity Center (successor to AWS Single Sign-On), Amazon WorkDocs, Amazon WorkMail, Amazon WorkSpaces, and the AWS Management Console. AWS Managed Microsoft AD must be able to query the users and groups in your self-managed AD.

Amazon EC2, Amazon RDS, and Amazon FSx will work with either a one-way or two-way trust.

upvoted 14 times

pbpally 3 years, 1 month ago

I found the documentation that explicitly states that IAM Identity Center (successor to AWS SSO) requires a two-way trust:
https://docs.aws.amazon.com/directoryservice/latest/admin-guide/ms_ad_setup_trust.html

upvoted 7 times

AswinHadinata **Most Recent** 3 weeks, 2 days ago

Selected Answer: A

I choose A, because it's telling about continue managing the users and groups in its on-premise, means 1 way trust is sufficient to delegate this requirement

upvoted 1 times

Darius_W 8 months ago

Selected Answer: A

Its A:

Chat GPT.

Answer: A.

Enable AWS SSO (IAM Identity Center) and connect it to your on-prem AD via AWS Managed Microsoft AD with a one-way trust.

Why A is right (least privilege, least ops):

IAM Identity Center natively provides SSO across all AWS Organizations accounts.

Keep users/groups in on-prem AD; establish a one-way trust so AWS Managed Microsoft AD trusts your on-prem AD (users authenticate in AD, permissions assigned in Identity Center).

No need for extra IdP or broad two-way trusts.

Why not the others

B (two-way trust): Works, but not required and violates least privilege.

C (Directory Service only): Doesn't give org-wide SSO by itself; you still need Identity Center.

D (deploy on-prem IdP): Unnecessary complexity—Identity Center + Managed Microsoft AD already solves it.

upvoted 1 times

LucV3 9 months, 3 weeks ago

Selected Answer: A

It is A because two-way, because you have and organization with multiple accounts to manage and since you'd need to manage login's in a centralized manner, you'd need to have activated Identity Center and per the doc at
<https://docs.aws.amazon.com/singlesignon/latest/userguide/connectonpremad.html>, with Identity Center, you have to have two-way trust.

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: A

Why the Other Options Are Incorrect:

B. Two-way forest trust

Unnecessary and more complex than a one-way trust.

A one-way trust from AWS to on-prem AD is sufficient and recommended by AWS for this use case.

C. Use AWS Directory Service with two-way trust only

This option is incomplete — it lacks integration with AWS SSO (IAM Identity Center), which is required for SSO across AWS accounts.

Also, two-way trust is not required, increasing complexity unnecessarily.

D. Deploy an on-prem IdP + enable AWS SSO

Possible, but adds significant operational overhead by requiring you to deploy and manage a full SAML 2.0 IdP.

Not needed when AWS SSO supports direct integration with self-managed AD via AWS Managed Microsoft AD and trust relationships.

upvoted 3 times

Juju66 1 year, 1 month ago

Selected Answer: A

I think one-way trust is sufficient for AWS SSO to authenticate users from the on-prem AD without requiring changes to the on-prem directory. two-way trust is unnecessary for SSO and introduces additional complexity.

Both two-way and one-way are valid for forest trust.

[https://learn.microsoft.com/en-us/previous-versions/windows/it-pro/windows-server-2008-r2-and-2008/cc771397\(v=ws.10\)](https://learn.microsoft.com/en-us/previous-versions/windows/it-pro/windows-server-2008-r2-and-2008/cc771397(v=ws.10))

upvoted 1 times

omunoz 1 year, 2 months ago

Selected Answer: A

A one-way trust (rather than a two-way trust) is the recommended approach for security reasons:

The on-premises AD remains the authoritative identity provider.

AWS trusts the on-prem AD, but AD does not trust AWS.

This setup ensures that AWS does not introduce potential security risks to the on-premises

A two-way trust is unnecessary for SSO and can introduce security risks, as it allows AWS to authenticate users in the on-prem AD and vice versa.

AWS best practices recommend using a one-way trust for such integrations.

upvoted 3 times

javaguy1906 1 year, 3 months ago

Selected Answer: A

A one-way trust is sufficient because AWS should trust the on-premises AD for authentication, but the on-premises AD does not need to trust AWS

upvoted 1 times

AwsAbhiKumar 1 year, 4 months ago

Selected Answer: A

AWS Single Sign-On can integrate with on-premises Microsoft Active Directory using AWS Directory Service for Microsoft Active Directory.

A one-way trust relationship is sufficient for AWS SSO to authenticate users from the on-premises Active Directory without requiring changes to existing trust relationships in the on-premises environment.

upvoted 1 times

AfricanCloudGuru 1 year, 4 months ago

Selected Answer: B

Because we need to create two ways forest traffic

upvoted 1 times

Rcosmos 1 year, 5 months ago

Selected Answer: B

Resposta correta: B!

Explicação revisada:

Quando a empresa precisa integrar seu Microsoft Active Directory autogerenciado com o AWS SSO para logon único em múltiplas contas na AWS, uma relação de confiança de floresta bidirecional é necessária para permitir que usuários e grupos locais sejam autenticados e tenham acesso completo aos recursos da AWS. Essa abordagem garante interoperabilidade total entre o AD local e o AWS SSO.

upvoted 1 times

MGKYAING 1 year, 5 months ago

Selected Answer: A

AWS Single Sign-On (AWS SSO) allows centralized management of access to multiple AWS accounts, and it integrates with on-premises Active Directory via AWS Directory Service for Microsoft Active Directory.

A one-way trust enables AWS Directory Service to authenticate users from the on-premises AD without giving AWS access to manage the on-premises AD. This is secure and follows best practices.

This is the correct solution because it integrates the on-premises AD with AWS SSO and maintains the centralized management of users and groups in the on-premises AD.

upvoted 3 times

John54321 1 year, 5 months ago

Selected Answer: B

AWS SSO (previously called Identity center) requires two way trust.

https://docs.aws.amazon.com/directoryservice/latest/admin-guide/ms_ad_setup_trust.html

upvoted 1 times

hilker1983 1 year, 5 months ago

Selected Answer: A

AWS Single Sign-On (AWS SSO): Provides a centralized way to manage SSO access to multiple AWS accounts and business applications. It simplifies user access management across AWS Organizations.

One-Way Trust: Creating a one-way trust with AWS Directory Service allows the on-premises Microsoft Active Directory to authenticate users without allowing AWS to initiate authentication requests to the on-premises directory. This maintains security and control.

Minimal Operational Overhead: This approach leverages AWS managed services, reducing the complexity and overhead of managing SSO and directory services.

upvoted 1 times

aatikah 1 year, 6 months ago

Selected Answer: A

Application Context in This Scenario:

AWS Applications Mentioned in the Scenario:

The question focuses on AWS Single Sign-On (AWS SSO) and the use of SSO across AWS accounts.

AWS SSO with AWS Organizations does not require AWS applications to look up or access on-premises AD objects beyond authentication.

This scenario does not mention such AWS applications, so a one-way trust is sufficient for the SSO use case.

upvoted 1 times

Carlini2020 1 year, 6 months ago

Selected Answer: A

No need for both ways communication.

upvoted 1 times

A company provides a Voice over Internet Protocol (VoIP) service that uses UDP connections. The service consists of Amazon EC2 instances that run in an Auto Scaling group. The company has deployments across multiple AWS Regions. The company needs to route users to the Region with the lowest latency. The company also needs automated failover between Regions. Which solution will meet these requirements?

- A. Deploy a Network Load Balancer (NLB) and an associated target group. Associate the target group with the Auto Scaling group. Use the NLB as an AWS Global Accelerator endpoint in each Region. **Most Voted**
- B. Deploy an Application Load Balancer (ALB) and an associated target group. Associate the target group with the Auto Scaling group. Use the ALB as an AWS Global Accelerator endpoint in each Region.
- C. Deploy a Network Load Balancer (NLB) and an associated target group. Associate the target group with the Auto Scaling group. Create an Amazon Route 53 latency record that points to aliases for each NLB. Create an Amazon CloudFront distribution that uses the latency record as an origin.
- D. Deploy an Application Load Balancer (ALB) and an associated target group. Associate the target group with the Auto Scaling group. Create an Amazon Route 53 weighted record that points to aliases for each ALB. Deploy an Amazon CloudFront distribution that uses the weighted record as an origin.

Correct Answer: A

Community vote distribution

A (82%)

Other

Comments

Six_Fingered_Jose **Highly Voted** 3 years, 7 months ago

Selected Answer: A

agree with A,
Global Accelerator has automatic failover and is perfect for this scenario with VoIP
<https://aws.amazon.com/global-accelerator/faqs/>
upvoted 66 times

bnagaraja9099 2 years, 7 months ago

A - Global Accelerator is a good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP, as well as for HTTP use cases that specifically require static IP addresses or deterministic, fast regional failover. Both services integrate with AWS Shield for DDoS protection.
upvoted 11 times

TilTil 2 years, 2 months ago

This is the best case for A to be an answer. Cloudfront is great but for HTTP use cases.
upvoted 2 times

BoboChow 3 years, 7 months ago

Thank you for your link, it make me consolidate A.
upvoted 7 times

bullrem 3 years, 4 months ago

This option does not meet the requirements because AWS Global Accelerator is only used to route traffic to the optimal AWS Region, it does not provide automatic failover between regions.
upvoted 2 times

sachin 3 years, 3 months ago

Instant regional failover: AWS Global Accelerator automatically checks the health of your applications and routes user traffic only to healthy application endpoints. If the health status changes or you make configuration updates, AWS Global Accelerator reacts instantaneously to route your users to the next available endpoint.
upvoted 10 times

ElaineRan 2 years, 10 months ago

Thank you, the link also helps me to know the differences between Global Acc and CloudFront.
upvoted 4 times

tavy 2 years, 7 months ago

ok but answer A does not mention service Global Accelerator, it mentions the NLB would act like one. Not sure if the wording is wrong or not.
' Deploy a Network Load Balancer (NLB) and an associated target group. Associate the target group with the Auto Scaling group. Use the NLB as an AWS Global Accelerator endpoint in each Region. ' - keyword use NLB 'as an'
upvoted 4 times

boooliyooo 2 years, 5 months ago

as an AWS Global Accelerator endpoint
upvoted 2 times

mouhannadhaj **Highly Voted** 3 years, 7 months ago

Selected Answer: A

CloudFront uses Edge Locations to cache content while Global Accelerator uses Edge Locations to find an optimal pathway to the nearest regional endpoint. CloudFront is designed to handle HTTP protocol meanwhile Global Accelerator is best used for both HTTP and non-HTTP protocols such as TCP and UDP. so i think A is a better answer
upvoted 42 times

nesyto **Most Recent** 4 months, 3 weeks ago

Selected Answer: A

global accelerator is perfect for udp
upvoted 1 times

S_K_han 9 months ago

Selected Answer: A

A seems to be the right answer. It makes sense
upvoted 1 times

xuefeng 9 months, 2 weeks ago

Selected Answer: A

Global Accelerator 具有自动故障转移功能
upvoted 1 times

jucejara 11 months, 1 week ago

Selected Answer: A

Global Accelerator = UDP
upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: A

Use AWS Global Accelerator with NLB endpoints in each Region for low-latency, highly available, and UDP-compatible routing — the best fit for a global, real-time VoIP service.
upvoted 1 times

LEEOscar 1 year ago

Selected Answer: A

UDP Support:
UDP traffic (used by VoIP) is not supported by Application Load Balancer (ALB).

Network Load Balancer (NLB) does support UDP, making NLB the correct choice.
Lowest latency routing:
AWS Global Accelerator is designed to route traffic to the AWS Region with the lowest latency.

It uses the AWS global network for optimal performance, which is better than Route 53 latency-based routing.
upvoted 3 times

Juju66 1 year, 1 month ago

Selected Answer: A

UDP - User Datagram Protocol (UDP) is a transport layer protocol used for sending messages (datagrams) between applications over a network. UDP is part of the Internet Protocol (IP) suite and is defined by Internet Engineering Task Force (IETF) standards in RFC 768.
upvoted 1 times

lubomir.tomecek 1 year, 4 months ago

Selected Answer: A

C is not correct, because Network Loadbalancer is already integrated with Route 53.

upvoted 1 times

adamatic 1 year, 4 months ago

Selected Answer: A

i full understand this is a use case for global accelerator but can anyone tell me why you need a NLB over an ALB here?

upvoted 1 times

adamatic 1 year, 4 months ago

oh i see, ALB is for HTTP/HTTPS

upvoted 2 times

AfricanCloudGuru 1 year, 4 months ago

Selected Answer: A

The Nature of the application VoIP needs NLB and Global Accelerator + an auto scaling

upvoted 1 times

ylanzzz 1 year, 5 months ago

Selected Answer: A

cloudfront doesn't provide automatic failover

upvoted 2 times

satyaamm 1 year, 5 months ago

Selected Answer: A

NLB is suitable for UDP and global accelerator is also suitable for UDP.

upvoted 1 times

AbhishekCloudEngg 1 year, 5 months ago

Selected Answer: A

The correct answer is A: Deploy a Network Load Balancer (NLB) and an associated target group. Associate the target group with the Auto Scaling group. Use the NLB as an AWS Global Accelerator endpoint in each Region.

Key reasons:

NLB supports UDP protocol (required for VoIP), while ALB only handles HTTP/HTTPS

Global Accelerator provides:

Automatic routing to lowest latency region

Built-in cross-region failover

Optimized network performance via AWS global network

upvoted 1 times

minhhieu_tech_1 1 year, 6 months ago

Selected Answer: A

Cloudfront is not support UDP Protocol. So maybe C answer isn't correct.

upvoted 1 times

EzKkk 1 year, 6 months ago

Selected Answer: C

I think this is not a very well written one. What we need:

1, Network Service that handles latency based policy and cross region failover -> Route53

2, A component that route traffic to ASG that are compatible with UDP -> Network Load Balancer.

upvoted 1 times

A development team runs monthly resource-intensive tests on its general purpose Amazon RDS for MySQL DB instance with Performance Insights enabled. The testing lasts for 48 hours once a month and is the only process that uses the database. The team wants to reduce the cost of running the tests without reducing the compute and memory attributes of the DB instance. Which solution meets these requirements MOST cost-effectively?

- A. Stop the DB instance when tests are completed. Restart the DB instance when required.
- B. Use an Auto Scaling policy with the DB instance to automatically scale when tests are completed.
- C. Create a snapshot when tests are completed. Terminate the DB instance and restore the snapshot when required. **Most Voted**
- D. Modify the DB instance to a low-capacity instance when tests are completed. Modify the DB instance again when required.

Correct Answer: C

Community vote distribution

C (81%)

Other

Comments

hanhdroid **Highly Voted** 3 years, 8 months ago

Selected Answer: C

Answer C, you still pay for storage when an RDS database is stopped
upvoted 46 times

KVK16 **Highly Voted** 3 years, 7 months ago

Selected Answer: C

C - Create a manual Snapshot of DB and shift to S3- Standard and Restore from Manual Snapshot when required.

Not A - By stopping the DB although you are not paying for DB hours you are still paying for Provisioned IOPs , the storage for Stopped DB is more than Snapshot of underlying EBS vol. and Automated Back ups .

Not D - Is possible but not MOST cost effective, no need to run the RDS when not needed.

upvoted 15 times

nesyto **Most Recent** 4 months, 3 weeks ago

Selected Answer: C

C is the correct answer as it is cost-effective
upvoted 1 times

maninblack 5 months, 3 weeks ago

Selected Answer: C

I selected C and it's the only one that makes sense to me,
however this part of the questions is what makes me not so sure in that answer

"The team wants to reduce the cost of running the tests without reducing the compute and memory attributes of the DB instance."

I guess terminating the db instance does not account towards reducing compute and memory attributes?

upvoted 1 times

Kosemin 9 months ago

Selected Answer: C

This option will allow you to save on costs by only paying for storage of the snapshot when the DB instance is terminated.
upvoted 1 times

j_d_rk9 10 months, 1 week ago

Selected Answer: C

It is C but I am honestly unsure when this will ever be a real scenario.
upvoted 3 times

Nishant27saini 10 months, 4 weeks ago

Selected Answer: C

Option C is the right answer as you can save the compute cost with this
upvoted 2 times

mirceack 11 months, 1 week ago

Selected Answer: A

However, creating a snapshot and then terminating the DB instance would save you the cost of running the instance for the 48 hours when it is not in use, but there are some restore-time penalties when bringing the database back online.
upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: C

Use RDS snapshots + termination + restore workflow to minimize costs for a monthly-used, resource-intensive RDS MySQL instance while maintaining full compute and memory attributes during test periods.
upvoted 1 times

jppush 11 months, 2 weeks ago

Selected Answer: C

Correct answer is C because A, you still pay for storage when the instance is stopped.
upvoted 1 times

trieu81200111 11 months, 3 weeks ago

Selected Answer: C

C is correct
upvoted 1 times

RenukaB 11 months, 3 weeks ago

Selected Answer: C

Answer C: Creating snapshot of EBS volumes is the trick to save on costs.
upvoted 1 times

Juju66 1 year, 1 month ago

Selected Answer: C

Stopping an RDS instance does not reduce costs significantly because you still pay for the storage and some management overhead. Additionally, stopping and starting an RDS instance can take time and may not be as efficient as terminating and restoring from a snapshot.
upvoted 2 times

donquixote47 1 year, 2 months ago

Selected Answer: C

you still pay for storage when an RDS database is stopped
upvoted 1 times

Kazmin 1 year, 2 months ago

Selected Answer: B

B is right i think
upvoted 1 times

John54321 1 year, 5 months ago

Selected Answer: C

An RDS instance can be stopped for up to seven days so it cannot be A.
"The instance stops running, up to a maximum of 7 consecutive days."
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_StopInstance.html
upvoted 3 times

f51a8bd 1 year, 5 months ago

Selected Answer: A

La opción A es la más rentable porque detiene la instancia de base de datos cuando no se está utilizando, ahorrando los costos de cómputo, y luego se reinicia cuando sea necesario, sin comprometer el rendimiento durante las pruebas.
upvoted 1 times

A company that hosts its web application on AWS wants to ensure all Amazon EC2 instances, Amazon RDS DB instances, and Amazon Redshift clusters are configured with tags. The company wants to minimize the effort of configuring and operating this check.

What should a solutions architect do to accomplish this?

- A. Use AWS Config rules to define and detect resources that are not properly tagged. **Most Voted**
- B. Use Cost Explorer to display resources that are not properly tagged. Tag those resources manually.
- C. Write API calls to check all resources for proper tag allocation. Periodically run the code on an EC2 instance.
- D. Write API calls to check all resources for proper tag allocation. Schedule an AWS Lambda function through Amazon CloudWatch to periodically run the code.

Correct Answer: A

Community vote distribution

A (100%)

Comments

kurinei021 **Highly Voted** 3 years, 5 months ago

Answer from ChatGPT:

Yes, you can use AWS Config to create tags for your resources. AWS Config is a service that enables you to assess, audit, and evaluate the configurations of your AWS resources. You can use AWS Config to create rules that automatically tag resources when they are created or when their configurations change.

To create tags for your resources using AWS Config, you will need to create an AWS Config rule that specifies the tag key and value you want to use and the resources you want to apply the tag to. You can then enable the rule and AWS Config will automatically apply the tag to the specified resources when they are created or when their configurations change.

upvoted 22 times

aaroncelestine 2 years, 9 months ago

This is the first answer that I've seen ChatGPT get correct here on ExamTopics. You should all know that using ChatGPT for this is bound to give bad answers. It only parrots what it has seen written/copied/pasted by someone/something somewhere, picked up with absolutely zero context. ChatGPT doesn't "know" anything about AWS services. So, beware the "answers" it gives.

upvoted 14 times

kidomaruto 2 years, 7 months ago

I tried it with Bing AI, and the answer was almost always the right one. It depends a lot on the prompt quality

upvoted 3 times

cookieMr **Highly Voted** 2 years, 11 months ago

Selected Answer: A

AWS Config provides a set of pre-built or customizable rules that can be used to check the configuration and compliance of AWS resources. By creating a custom rule or using the built-in rule for tagging, you can define the required tags for EC2, RDS DB and Redshift clusters. AWS Config continuously monitors the resources and generates configuration change events or evaluation results.

By leveraging AWS Config, the solution can automatically detect any resources that do not comply with the defined tagging requirements. This approach eliminates the need for manual checks or periodic code execution, reducing operational overhead. Additionally, AWS Config provides the ability to automatically remediate non-compliant resources by triggering Lambda or sending notifications, further streamlining the configuration management process.

Option B (using Cost Explorer) primarily focuses on cost analysis and does not provide direct enforcement of proper tagging. Option C and D (writing API calls and running them manually or through scheduled Lambda) require more manual effort and maintenance compared to using AWS Config rules.

upvoted 10 times

Clouddon Most Recent 11 months, 2 weeks ago

Selected Answer: A

Use AWS Config with required-tags rules to automatically and continuously check for proper tagging on EC2, RDS, and Redshift resources — with the least effort and highest reliability.

upvoted 1 times

AshishDhole 1 year, 4 months ago

Selected Answer: A

By Using AWS Config we can tag the resources, and then filter the resources which are tagged or not tagged.

upvoted 1 times

satyaamm 1 year, 5 months ago

Selected Answer: A

AWS Config helps identify any configuration issues.

upvoted 1 times

MGKYAING 1 year, 5 months ago

Selected Answer: A

AWS Config is designed to monitor and assess the compliance of AWS resource configurations.

You can use AWS Config managed rules, such as required-tags, to automatically evaluate whether resources are properly tagged.

AWS Config provides a centralized dashboard to view non-compliant resources and requires minimal setup and maintenance.

This is the correct answer because it provides an automated, scalable, and low-effort solution to detect resources that are not properly tagged.

upvoted 1 times

mzeynalli 1 year, 7 months ago

Selected Answer: A

Answer: A. Use AWS Config rules to define and detect resources that are not properly tagged.

Explanation:

Using AWS Config is the best solution because it provides a managed service to monitor and evaluate the configurations of your AWS resources continuously. Here's why AWS Config is the preferred choice:

Automated Tag Compliance:

Ease of Setup and Operation:

Comprehensive Coverage:

AWS Config can monitor a wide range of AWS resources, including Amazon EC2 instances, RDS databases, Redshift clusters, and more, ensuring that all your resources are properly tagged.

upvoted 3 times

PaulGa 1 year, 9 months ago

Selected Answer: A

Not sure we need ChatGPT here (other than to prove what it can/can't do): answers B, C, D all require manual - and periodic - intervention. Has to be Ans A.

upvoted 1 times

DavidNgTan 1 year, 10 months ago

Selected Answer: A

You can use AWS Config to create a rule that evaluates whether your resources have the required tags applied.

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: A

<https://docs.aws.amazon.com/whitepapers/latest/tagging-best-practices/implementing-and-enforcing-tagging.html>

AWS Config (required_tag)

AWS Config is a service that allows you to assess, audit, and evaluate the configurations of your AWS resources (see Resource types supported by AWS Config). In the case of tagging, we can use it to identify resources that are lacking tags with specific keys, using the required_tags rule (refer to Resource types supported by required_tags). From the earlier example, we might test for the existence of the key on all Amazon EC2 instances.

upvoted 3 times

A_jaa 2 years, 5 months ago

Selected Answer: A

Answer-A

upvoted 1 times

Ruffyit 2 years, 7 months ago

Has typos in the question, correct is "A company that hosts its web application on AWS wants to ensure all Amazon EC2 instance, Amazon RDS DB instances, and Amazon Redshift clusters are configured with tags." Keyword "are configured with tags", choose (A) "AWS Config rules".

upvoted 2 times

awashenko 2 years, 8 months ago

Selected Answer: A

I originally thought D, but after reading through the discussion I agree that option A would require less effort. D would get the job done but would require more effort so I think A is correct.

upvoted 2 times

KawtarZ 2 years, 9 months ago

Selected Answer: A

A without a doubt

upvoted 1 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: A

AWS Config continually assesses, audits, and evaluates the configurations and relationships of your resources on AWS, on premises, and on other clouds.

upvoted 3 times

james2033 2 years, 10 months ago

Selected Answer: A

Has typos in the question, correct is "A company that hosts its web application on AWS wants to ensure all Amazon EC2 instance, Amazon RDS DB instances, and Amazon Redshift clusters are configured with tags." Keyword "are configured with tags", choose (A) "AWS Config rules".

upvoted 2 times

miki111 2 years, 11 months ago

Option A is the right answer for this.

upvoted 1 times

A development team needs to host a website that will be accessed by other teams. The website contents consist of HTML, CSS, client-side JavaScript, and images.

Which method is the MOST cost-effective for hosting the website?

- A. Containerize the website and host it in AWS Fargate.
- B. Create an Amazon S3 bucket and host the website there. **Most Voted**
- C. Deploy a web server on an Amazon EC2 instance to host the website.
- D. Configure an Application Load Balancer with an AWS Lambda target that uses the Express.js framework.

Correct Answer: B

Community vote distribution

B (100%)

Comments

masetromain **Highly Voted** 3 years, 8 months ago

Selected Answer: B

Good answer is B: client-side JavaScript. the website is static, so it must be S3.
upvoted 35 times

BoboChow **Highly Voted** 3 years, 8 months ago

Selected Answer: B

HTML, CSS, client-side JavaScript, and images are all static resources.
upvoted 10 times

Clouddon **Most Recent** 11 months, 2 weeks ago

Selected Answer: B

Use Amazon S3 static website hosting for a simple, scalable, and low-cost solution to serve HTML, CSS, JavaScript, and images to internal teams.
upvoted 1 times

dattateja8 1 year, 2 months ago

Selected Answer: B

Ans B: Don't get confused by looking at JavaScript which can be dynamic. It is Client-side JavaScript. Means it get's executed at the Client side in a web browser. So it is all static and S3 is the cheapest option to host a static website.
upvoted 2 times

satyaammm 1 year, 5 months ago

Selected Answer: B

S3 is the most suitable option for hosting static websites.
upvoted 1 times

PaulGa 1 year, 9 months ago

Selected Answer: B

I had initially thought Ans A... but its Ans C -- "cookieMr" makes it clear:
Ans A "Containerising the website and hosting with AWS Fargate involves additional complexity and costs associated with managing the container environment and scaling resources."
So it has to be...
Ans B: "...Amazon S3 to host the website, take advantage of its durability, scalability, and low-cost pricing model. Only pay for the storage and data transfer associated with your website, without the need for managing and maintaining web servers or containers. This reduces the operational overhead and infrastructure costs."
upvoted 1 times

awsgeek75 2 years, 5 months ago

Selected Answer: B

Cheapest Static site hosting = S3

upvoted 3 times

A_jaa 2 years, 5 months ago

Selected Answer: B

Answer-B

upvoted 2 times

Ruffyt 2 years, 7 months ago

HTML, CSS, client-side JavaScript, and images are all static resources.

upvoted 1 times

AWStudyBuddy 2 years, 7 months ago

The MOST cost-effective method for hosting a website is to:

Create an Amazon S3 bucket and host the website there.

Amazon S3 is a highly scalable and cost-effective object storage service. It is a good option for hosting static websites, such as the website in this scenario.

To host a static website on Amazon S3, you would first need to create an S3 bucket. Then, you would need to upload the website files to the bucket. Once the files are uploaded, you can configure the bucket to serve as a website.

upvoted 2 times

hungpm 2 years, 9 months ago

Selected Answer: B

Static website should work fine with S3

upvoted 1 times

KawtarZ 2 years, 9 months ago

Selected Answer: B

the website is static because the backend runs on client side.

upvoted 2 times

evanhongo 2 years, 10 months ago

Selected Answer: B

all static resources.

upvoted 1 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: B

static website, cost-effective = S3 web hosting

upvoted 3 times

james2033 2 years, 10 months ago

Selected Answer: B

Just all static content HTML, CSS, client-side JavaScript, images. Amazon S3 is good enough.

upvoted 1 times

miki111 2 years, 11 months ago

Option B is the right answer for this.

upvoted 1 times

Kaab_B 2 years, 11 months ago

Selected Answer: B

S3 is amongst the cheapest services offered by AWS.

upvoted 1 times

A company runs an online marketplace web application on AWS. The application serves hundreds of thousands of users during peak hours. The company needs a scalable, near-real-time solution to share the details of millions of financial transactions with several other internal applications. Transactions also need to be processed to remove sensitive data before being stored in a document database for low-latency retrieval.

What should a solutions architect recommend to meet these requirements?

- A. Store the transactions data into Amazon DynamoDB. Set up a rule in DynamoDB to remove sensitive data from every transaction upon write. Use DynamoDB Streams to share the transactions data with other applications.
- B. Stream the transactions data into Amazon Kinesis Data Firehose to store data in Amazon DynamoDB and Amazon S3. Use AWS Lambda integration with Kinesis Data Firehose to remove sensitive data. Other applications can consume the data stored in Amazon S3.
- C. Stream the transactions data into Amazon Kinesis Data Streams. Use AWS Lambda integration to remove sensitive data from every transaction and then store the transactions data in Amazon DynamoDB. Other applications can consume the transactions data off the Kinesis data stream. **Most Voted**
- D. Store the batched transactions data in Amazon S3 as files. Use AWS Lambda to process every file and remove sensitive data before updating the files in Amazon S3. The Lambda function then stores the data in Amazon DynamoDB. Other applications can consume transaction files stored in Amazon S3.

Correct Answer: C

Community vote distribution

C (87%)

13%

Comments

ArielSchivo **Highly Voted** 3 years, 7 months ago

Selected Answer: C

I would go for C. The tricky phrase is "near-real-time solution", pointing to Firehose, but it can't send data to DynamoDB, so it leaves us with C as best option.

Kinesis Data Firehose currently supports Amazon S3, Amazon Redshift, Amazon OpenSearch Service, Splunk, Datadog, NewRelic, Dynatrace, Sumologic, LogicMonitor, MongoDB, and HTTP End Point as destinations.

<https://aws.amazon.com/kinesis/data-firehose/faqs/#:~:text=Kinesis%20Data%20Firehose%20currently%20supports,HTTP%20End%20Point%20as%20destinations.>
upvoted 102 times

Lonojack 3 years, 4 months ago

This was a really tough one. But you have the best explanation on here with reference point. Thanks. I'm going with answer C!
upvoted 4 times

SaraSundaram 3 years, 2 months ago

There are many questions having Firehose and Stream. Need to know them in detail to answer. Thanks for the explanation
upvoted 4 times

diabloexodia 2 years, 11 months ago

Stream is used if you want real time results , but with firehose , you generally use the data at a later point of time by storing it somewhere. Hence if you see "REAL TIME" the answer is most probably Kinesis Data Streams.
upvoted 20 times

lizzard812 3 years, 4 months ago

Sorry but I still can't see how Kinesis Data Stream is 'scalable', since you have to provision the quantity of shards in advance?
upvoted 1 times

habibi03336 3 years, 3 months ago

"easily stream data at any scale"

This is a description of Kinesis Data Stream. I think you can configure its quantity but still not provision and manage scalability by yourself.

upvoted 1 times

JesseeS **Highly Voted** 3 years, 7 months ago

The answer is C, because Firehose does not support DynamoDB and another key word is "data" Kinesis Data Streams is the correct choice. Pay attention to key words. AWS likes to trick you up to make sure you know the services.

upvoted 33 times

Clouddon **Most Recent** 11 months, 2 weeks ago

Selected Answer: C

Use Amazon Kinesis Data Streams with AWS Lambda for real-time data sanitization and multi-application consumption, and store sanitized data in Amazon DynamoDB for fast retrieval — an optimal solution for secure, scalable, low-latency transaction processing.

upvoted 1 times

kyd0nix 1 year, 4 months ago

Selected Answer: B

IMO "near-real-time" is key for Firehose, BUT since of all the discussions B vs C (Firehose can't have DynamoDB as destination, I think the question is misswritten and has to be reviewed to avoid the confusion)

upvoted 1 times

FlyingHawk 1 year, 4 months ago

Selected Answer: C

C - Kinesis Data Streams allows for low-latency processing, which is crucial for near-real-time requirements.

upvoted 1 times

MGKYAING 1 year, 5 months ago

Selected Answer: C

Scalable processing: The system must scale to handle hundreds of thousands of users and millions of transactions during peak hours.

Near-real-time sharing of transactions: Transactions should be shared with internal applications in near-real time.

Sensitive data removal: Sensitive information must be processed and removed before storage.

Low-latency retrieval: The processed data must be stored in a document database (Amazon DynamoDB) for quick access.

upvoted 1 times

aefuen1 1 year, 5 months ago

Selected Answer: B

It's B. You can write to the DynamoDB table from the lambda preprocessing function. Also option C can't be correct because if "Other applications can consume the transactions data off the Kinesis data stream", this means they will consume data with sensitive values, which is a constraint for the solution.

upvoted 2 times

kernel1 1 year, 3 months ago

The question says it needs to have sensitive information removed before DB storage. Other apps could consume transaction data but not necessarily store, or consume only non-sensitive data.

upvoted 1 times

engnrshon 1 year, 6 months ago

C :

upvoted 1 times

Mauro0001 1 year, 9 months ago

Selected Answer: C

One of the tricky phrases is 'near-real-time solutions' because it points to the fact that every time a write is made to a database, it incurs a delay, and then retrieving it with an API call adds another latency. With Kinesis Data Streams, that process is optimized because the intermediary that gives you the ability to write to DynamoDB also provides that data to other services due to the retention period of Kinesis Data Streams.

upvoted 2 times

PaulGa 1 year, 9 months ago

Selected Answer: C

Ans C.

High level difference between the Kinesis and DynamoDB:

Kinesis Streams allows production/ consumption of large volumes of data (web data, logs, etc); DynamoDB Streams is a feature local to DynamoDB to track the granular changes to DynamoDB table items.

(Note also: data latency for Firehose is 60 seconds or higher; Streams is for custom processing and has sub-second processing latency).

upvoted 2 times

Lin878 1 year, 11 months ago

Selected Answer: C

Q: What is a destination in Firehose?

A destination is the data store where your data will be delivered. Firehose currently supports Amazon S3, Amazon Redshift, Amazon OpenSearch Service, Splunk, Datadog, NewRelic, Dynatrace, Sumo Logic, LogicMonitor, MongoDB, and HTTP End Point as destinations.

<https://aws.amazon.com/firehose/faqs/>

upvoted 2 times

the_mellie 2 years ago

Selected Answer: C

with multiple consumers and on the fly modification, it seems like the most logical choice

upvoted 2 times

vi24 2 years, 3 months ago

I chose B. The "near real time" is very specific to Kinesis firehose which is a better option anyway. The rest of the answer makes sense too. C is wrong "sensitive data removed by Lambda & then store transaction data in DynamoDB", while it continues to say other applications are accessing the transaction data from Kinesis Data stream !!

upvoted 3 times

Pics00094 2 years, 3 months ago

Selected Answer: C

need to know.

- 1) Lambda Integration
- 2) Difference between Real time(Kinesis Data Stream) vs Near Real time(Kinesis Fire House)
- 3) Firehose can't target DynamoDB

upvoted 5 times

JulianWaksmann 2 years, 4 months ago

i think c are bad too, because it isn't near real time.

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: C

A: DynamoDB streams are logs, not fit for real-time sharing.

B: S3 is not document database, it's BLOB

D: S3 and files are not database

C: Kinesis + Lambda + DynamoDB is high performance, low latency scalable solution.

upvoted 3 times

A_jaa 2 years, 5 months ago

Selected Answer: C

Answer-C

upvoted 1 times

A company hosts its multi-tier applications on AWS. For compliance, governance, auditing, and security, the company must track configuration changes on its AWS resources and record a history of API calls made to these resources. What should a solutions architect do to meet these requirements?

- A. Use AWS CloudTrail to track configuration changes and AWS Config to record API calls.
- B. Use AWS Config to track configuration changes and AWS CloudTrail to record API calls. **Most Voted**
- C. Use AWS Config to track configuration changes and Amazon CloudWatch to record API calls.
- D. Use AWS CloudTrail to track configuration changes and Amazon CloudWatch to record API calls.

Correct Answer: B

Community vote distribution

B (99%)

Comments

airraid2010 **Highly Voted** 3 years, 7 months ago

Selected Answer: B

CloudTrail - Track user activity and API call history.

Config - Assess, audits, and evaluates the configuration and relationships of tag resources.

Therefore, the answer is B

upvoted 40 times

Clouddon **Most Recent** 11 months, 2 weeks ago

Selected Answer: B

Use AWS Config for resource configuration tracking and AWS CloudTrail for API call logging — the best practice combo for governance, compliance, and auditing on AWS.

upvoted 1 times

satyaamm 1 year, 5 months ago

Selected Answer: B

AWS Config for configuration changes and AWS CloudTrail for API activities.

upvoted 2 times

SBS2 1 year, 6 months ago

Selected Answer: B

upvoted 1 times

awsgeek75 2 years, 5 months ago

Selected Answer: B

Config = Governance, auditing of AWS resource

CloudTrail = API call tracking

B is correct

upvoted 4 times

A_jaa 2 years, 5 months ago

Selected Answer: B

Answer-B

upvoted 1 times

Ruffyit 2 years, 7 months ago

Correct Answer- Option B. Here's why

AWS Config for Configuration Changes: AWS Config is a service that tracks changes to resource configurations over time. It provides a history of

configuration changes to your AWS resources and helps with compliance and auditing by allowing you to assess how resource configurations have changed over time.

AWS CloudTrail for API Calls: AWS CloudTrail is designed specifically for recording API calls made to AWS resources. It captures detailed information about who made each API call, the actions taken, and the resources affected. This is essential for auditing and security purposes.

upvoted 3 times

AWSStudyBuddy 2 years, 7 months ago

Correct Answer- Option B. Here's why

AWS Config for Configuration Changes: AWS Config is a service that tracks changes to resource configurations over time. It provides a history of configuration changes to your AWS resources and helps with compliance and auditing by allowing you to assess how resource configurations have changed over time.

AWS CloudTrail for API Calls: AWS CloudTrail is designed specifically for recording API calls made to AWS resources. It captures detailed information about who made each API call, the actions taken, and the resources affected. This is essential for auditing and security purposes.

While Amazon CloudWatch can be used to monitor and gather metrics, it is not designed for recording API calls or tracking configuration changes. AWS Config and AWS CloudTrail are purpose-built for these specific tasks and are the best services to use for the described requirements.

upvoted 4 times

AWSStudyBuddy 2 years, 7 months ago

Selected Answer: B

Although tracking configuration changes and recording API calls are not intended uses for Amazon CloudWatch, it can be utilized for monitoring and collecting data. AWS CloudTrail and AWS Config are purpose-built for these specific tasks and are the best services to use for the described requirements.

upvoted 1 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: B

CloudWatch is a monitoring service for AWS resources and applications. CloudTrail is a web service that records API activity in your AWS account.

upvoted 2 times

Bogs123456711 2 years, 10 months ago

Selected Answer: B

CONFIG - AWS CONFIG

RECORD API CALLS - CLOUDTRAIL

upvoted 3 times

hsinchang 2 years, 10 months ago

Selected Answer: B

CloudWatch is mainly used to monitor AWS services with metrics, not recording actions inside the AWS environments. It can also monitor CloudTrail logged events.

For recording API calls it requires CloudTrail.

upvoted 2 times

james2033 2 years, 10 months ago

Selected Answer: B

Keyword "Amazon CloudWatch" is not for this case, remove C and D.

Use AWS Config first to track configuration changes, Second is AWS CloudTrail to record API calls. (Answer B, and correct answer). Answer A is reversed order of B, and not accepted.

upvoted 3 times

miki111 2 years, 11 months ago

Option B is the right answer for this.

upvoted 1 times

karloscetina007 2 years, 11 months ago

Selected Answer: B

B is the answer with no doubts

upvoted 1 times

minhpn 2 years, 11 months ago

Selected Answer: B

config => AWS config

record API calls => AWS CloudTrail

upvoted 1 times

cookieMr 2 years, 11 months ago

Selected Answer: B

To meet the requirement of tracking configuration changes on AWS resources and recording a history of API calls, a solutions architect should recommend option B: Use AWS Config to track configuration changes and AWS CloudTrail to record API calls.

Option A (using CloudTrail to track configuration changes and Config to record API calls) is incorrect because CloudTrail is specifically designed to capture API call history, while Config is designed for tracking configuration changes.

Option C (using Config to track configuration changes and CloudWatch to record API calls) is not the recommended approach. While CloudWatch can be used for monitoring and logging, it does not provide the same level of detail and compliance tracking as CloudTrail for recording API calls.

Option D (using CloudTrail to track configuration changes and CloudWatch to record API calls) is not the optimal choice because CloudTrail is the appropriate service for tracking configuration changes, while CloudWatch is not specifically designed for recording API call history.

upvoted 3 times

A company is preparing to launch a public-facing web application in the AWS Cloud. The architecture consists of Amazon EC2 instances within a VPC behind an Elastic Load Balancer (ELB). A third-party service is used for the DNS. The company's solutions architect must recommend a solution to detect and protect against large-scale DDoS attacks.

Which solution meets these requirements?

- A. Enable Amazon GuardDuty on the account.
- B. Enable Amazon Inspector on the EC2 instances.
- C. Enable AWS Shield and assign Amazon Route 53 to it.
- D. Enable AWS Shield Advanced and assign the ELB to it. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

ninjawrz **Highly Voted** 3 years, 7 months ago

Selected Answer: D

Answer is D

C is incorrect because question says Third party DNS and route 53 is AWS proprietary

upvoted 43 times

kidomaruto 2 years, 7 months ago

Right answer, wrong explanation.

You can use Route 53 with a custom domain.. it's all about the "large-scale DDOS attack".

upvoted 22 times

BoboChow **Highly Voted** 3 years, 8 months ago

Selected Answer: D

AWS Shield Advanced provides expanded DDoS attack protection for your Amazon EC2 instances, Elastic Load Balancing load balancers, CloudFront distributions, Route 53 hosted zones, and AWS Global Accelerator standard accelerators.

upvoted 30 times

leonardh 3 years, 1 month ago

I'd agree as Shield Advanced is the only tier that can protect EC2 which is not possible in Standard.

upvoted 7 times

Clouddon **Most Recent** 11 months, 2 weeks ago

Selected Answer: D

To effectively detect and mitigate large-scale DDoS attacks on a public-facing web app behind an ELB, enable AWS Shield Advanced and assign it to the ELB.

upvoted 1 times

satyaamm 1 year, 5 months ago

Selected Answer: D

Only AWS Shield Advanced can help against DDos attacks and since we are using a third-party DNS so we can't attach it to Route 53.

upvoted 3 times

MGKYAING 1 year, 5 months ago

Selected Answer: D

Enabling AWS Shield Advanced provides advanced protection against large-scale DDoS attacks.

Assigning the ELB to AWS Shield Advanced ensures that traffic to the web application is monitored and protected with enhanced mitigation techniques.

The solution also offers a direct response mechanism (via the AWS DDoS Response Team) for further assistance during attacks.

upvoted 1 times

Tjazz04 1 year, 6 months ago

Selected Answer: D

Ans. D - Keyword - Large-scale DDoS attacks which the AWS Shield Advanced can prevent

upvoted 1 times

zied007 1 year, 9 months ago

Selected Answer: D

Answer is D

upvoted 1 times

PaulGa 1 year, 9 months ago

Selected Answer: D

Ans D. Shield (Advanced) is built for DDoS and can interface to ELB

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: D

A: GuardDuty is not for this, mostly for account monitoring for suspicious activity

B: Inspector is for OS vulnerabilities

C: Shield with R53 is not going to protect against DDoS

D: Shield Advanced is build for DDoS protection

upvoted 7 times

awsgeek75 2 years, 5 months ago

Forgot to mention, C won't work because a 3rd party DNS is used and R53 is not part of the setup

upvoted 4 times

awsgeek75 2 years, 5 months ago

Prevent large scale DDOS attack = AWS Shield Advanced

upvoted 2 times

A_jaa 2 years, 5 months ago

Selected Answer: D

Answer-D

upvoted 1 times

djgodzilla 2 years, 5 months ago

Selected Answer: D

- In addition to the network and transport layer protections that come with Standard, Shield Advanced provides additional detection and mitigation against large and sophisticated DDoS attacks, near real-time visibility into attacks, and integration with AWS WAF, a web application firewall.

<https://aws.amazon.com/shield/features/#:~:text=In%20addition%20to%20the%20network,WAF%2C%20a%20web%20application%20firewall.>

upvoted 2 times

OmegaLambda7XL9 2 years, 6 months ago

This one got me to be honest

upvoted 3 times

Ruffyit 2 years, 7 months ago

Option A is incorrect because Amazon GuardDuty is a threat detection service that focuses on identifying malicious activity and unauthorized behavior within AWS accounts. While it is useful for detecting various security threats, it does not specifically address large-scale DDoS attacks.

Option B is also incorrect because Amazon Inspector is a vulnerability assessment service that helps identify security issues and vulnerabilities within EC2. It does not directly protect against DDoS attacks.

Option C is not the optimal choice because AWS Shield provides basic DDoS protection for resources such as Elastic IP addresses, CloudFront, and Route53 hosted zones. However, it

upvoted 4 times

Ruffyit 2 years, 7 months ago

does not provide the advanced capabilities and assistance offered by AWS Shield Advanced, which is better suited for protecting against large-scale DDoS attacks.

Therefore, option D with AWS Shield Advanced and assigning the ELB to it is the recommended solution to detect and protect against large-scale DDoS attacks in the architecture described.

upvoted 3 times

Abitek007 2 years, 8 months ago

D, but can be tricky, the third party negates Route53

upvoted 1 times

Ak9kumar 2 years, 8 months ago

Answer D. Learn section on AWS Advanced Shield on [aws.Amazon.com](https://aws.amazon.com/shield/) to help you understand this. It helped me.

upvoted 2 times

ishant101 2 years, 9 months ago

answer is D

upvoted 1 times


A company is building an application in the AWS Cloud. The application will store data in Amazon S3 buckets in two AWS Regions. The company must use an AWS Key Management Service (AWS KMS) customer managed key to encrypt all data that is stored in the S3 buckets. The data in both S3 buckets must be encrypted and decrypted with the same KMS key. The data and the key must be stored in each of the two Regions.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create an S3 bucket in each Region. Configure the S3 buckets to use server-side encryption with Amazon S3 managed encryption keys (SSE-S3). Configure replication between the S3 buckets.
- B. Create a customer managed multi-Region KMS key. Create an S3 bucket in each Region. Configure replication between the S3 buckets. Configure the application to use the KMS key with client-side encryption. **Most Voted**
- C. Create a customer managed KMS key and an S3 bucket in each Region. Configure the S3 buckets to use server-side encryption with Amazon S3 managed encryption keys (SSE-S3). Configure replication between the S3 buckets.
- D. Create a customer managed KMS key and an S3 bucket in each Region. Configure the S3 buckets to use server-side encryption with AWS KMS keys (SSE-KMS). Configure replication between the S3 buckets.

Correct Answer: B

Community vote distribution



Comments

pooppants **Highly Voted** 3 years, 8 months ago

Selected Answer: B

KMS Multi-region keys are required <https://docs.aws.amazon.com/kms/latest/developerguide/multi-region-keys-overview.html>
upvoted 72 times

Instantqueue 2 years, 8 months ago

It's not correct because the question asks for server side encryption, not client side (before the objects reach the bucket).
upvoted 8 times

babayomi 1 year, 9 months ago

Note, The question did not ask for server side encryption, It stated Customer manages encryption. So reason why Customer side encryption makes the solution much easier and applicable. Hence B as an answer
upvoted 1 times

kelmryan1 2 years, 1 month ago

It says encrypt all data and the data originates from the application. Making it B
upvoted 1 times

sohailn 2 years, 10 months ago

Absolutely D is the right one because s3 kms multi region as an individual key so you must first decrypt in source bucket and then re-encrypt in target bucket
upvoted 6 times

babayomi 1 year, 9 months ago

Note, The question did not ask for server side encryption, It stated Customer manages encryption. So reason why Customer side encryption makes the solution much easier and applicable. Hence B as an answer
upvoted 2 times

kelmryan1 2 years, 1 month ago

The SAME key would exist in both regions
upvoted 1 times

0203b0f 2 years, 1 month ago

Answer is D : This option aligns with the requirement to use a customer-managed KMS key for encryption. It also ensures that the same KMS key is used for encryption and decryption across both AWS Regions, as required. By using SSE-KMS, you can ensure that encryption keys are managed by AWS KMS, providing greater control and security over the encryption process. Configuring replication between the S3 buckets ensures that data is synchronized across both regions. This approach minimizes operational overhead while meeting the specified requirements.

upvoted 4 times

babayomi 1 year, 9 months ago

Note, The question did not ask for server side encryption, It stated Customer manages encryption. So reason why Customer side encryption makes the solution much easier and applicable. Hence B as an answer

upvoted 1 times

Johan_jelly 2 years, 5 months ago

KMS multi-region keys are typically used when you need to enable cross-Region replication of encrypted data

upvoted 1 times

dwx101 2 years, 8 months ago

AWS services that integrate with AWS KMS for encryption at rest or digital signatures currently treat multi-Region keys as though they were single-Region keys. They might re-wrap or re-encrypt data moved between Regions. For example, Amazon S3 cross-region replication decrypts and re-encrypts data under a KMS key in the destination Region, even when replicating objects protected by a multi-Region key.

upvoted 3 times

KJa Highly Voted 3 years, 8 months ago

Selected Answer: D

Cannot be A - question says customer managed key
Cannot B - client side encryption is operational overhead
Cannot C - as it says SSE-S3 instead of customer managed
so the answer is D though it required one time setup of keys

upvoted 70 times

mattlai 3 years, 8 months ago

fun joke, if u dont do encryption on client side, where else could it be?

upvoted 1 times

Newptone 3 years, 7 months ago

It could be server side. For client side, the application need to finish the encryption and decryption by itself. So S3 object encryption on the server side is less operational overhead. <https://docs.aws.amazon.com/AmazonS3/latest/userguide/UsingClientSideEncryption.html>

But for option B, the major issue is if you create KMS keys in 2 regions, they can not be the same.

upvoted 7 times

Newptone 3 years, 7 months ago

Sorry for the typo, I mean option D.

upvoted 2 times

Clouddon 2 years, 10 months ago

Kindly point at where server-side encryption support multi-region. It is only mention on the aws blog that client-side support multi-region.

upvoted 2 times

BoboChow 3 years, 8 months ago

The data in both S3 buckets must be encrypted and decrypted with the same KMS key. AWS KMS supports multi-Region keys, which are AWS KMS keys in different AWS Regions that can be used interchangeably – as though you had the same key in multiple Regions. "as though" means it's different. So I agree with B

upvoted 15 times

BoboChow 3 years, 8 months ago

key change across regions unless you use multi-Region keys

upvoted 2 times

pentium75 2 years, 5 months ago

B includes replicating the data in the S3 buckets, which is not mentioned anywhere in the stem. It says that you need to store data in two buckets, not that you need to replicate content between buckets.

upvoted 1 times

Drew3000 2 years, 3 months ago

All the choices involve replication between the buckets.

upvoted 1 times

th3cookie 3 years, 6 months ago

How does client side encryption increase OPERATIONAL overhead? Do you think every connected client is sitting there with gpg cli, decrypting/encrypting every packet that comes in/out? No, it's done via SDK -> <https://docs.aws.amazon.com/encryption-sdk/latest/developer-guide/introduction.html>

The correct answer is B because that's the only way to actually get the same key across multiple regions with minimal operational overhead

upvoted 14 times

kakka22 3 years, 2 months ago

"The data in both S3 buckets must be encrypted and decrypted with the same KMS key"

Client side encryption means that key is generated in from the client without storing that in the KMS...

upvoted 7 times

karo000001 Most Recent 10 months, 1 week ago

Selected Answer: B

The tricky part is that in the question it says 'Which solution will meet these requirements with the LEAST operational overhead?' this will point you automatically to choose 'D' as an answer because it is less complex as from an architectural point of view it is simpler operationally comparing to B, but the fact that also in the answer it says 'Data in both s3 buckets must be encrypted and protected with the same KMS key, this eliminates B as B requires two separate KMS keys. D is arguably better from a practical standpoint but B is the only one that technically satisfies all the given constraints!

upvoted 2 times

Clouddon 11 months, 2 weeks ago

Selected Answer: B

Use a multi-Region AWS KMS customer managed key with client-side encryption to ensure consistent encryption/decryption across Regions, and minimize operational overhead.

upvoted 1 times

GGP23 12 months ago

Selected Answer: B

B because multi region keys are needed.

upvoted 1 times

Yak_Yeti 1 year, 2 months ago

Selected Answer: B

D says "Create a customer managed KMS key and an S3 bucket in each Region." If there are keys created in both regions, then the keys are not the same. If it said create a multi region key and synch it, then I would pick D.

upvoted 1 times

SleeplessHossam 1 year, 2 months ago

Selected Answer: B

I was confused which to choose, B or D, but someone in discussion mentioned really important point, D suggests creating two keys and two regions, which they can't be the same Key, so D of course is not the correct answer.

upvoted 1 times

Anastesas 1 year, 3 months ago

Selected Answer: D

Given the requirements to encrypt and decrypt data stored in S3 buckets in two AWS Regions using the same KMS key with minimal operational overhead, server-side encryption with AWS KMS (SSE-KMS) is the most suitable solution. It provides simplicity, seamless integration, centralized key management, and optimal performance.

Therefore: Option D: Create a customer managed KMS key and an S3 bucket in each Region. Configure the S3 buckets to use server-side encryption with AWS KMS keys (SSE-KMS). Configure replication between the S3 buckets.

upvoted 1 times

iwantcertificates 1 year, 3 months ago

Selected Answer: D

I think this question would be invalidated in an actual test. Both B and D has problems. B uses client-side encryption, which adds complexity to the operation. Additionally, KMS by default uses server side encryption, SSE-KMS. But the option D also has an imprecision, because you are creating two KMS keys instead of just one with multi-region

upvoted 1 times

sammo08 1 year, 3 months ago

Selected Answer: B

Key Reason: Separate KMS Keys in Each Region

In Option D, you are creating a separate AWS KMS key in each Region, meaning the encryption keys are distinct.

When replicating data between Regions, AWS does not allow direct decryption of objects encrypted with an SSE-KMS key in one Region using a different KMS key in another Region.

KMS keys are tied to a specific Region unless you use a Multi-Region key.

Because of this, the replicated objects cannot be automatically decrypted in the destination Region without additional configuration, increasing operational overhead.

Why Option B is Better

Option B uses a Multi-Region KMS Key that allows encryption and decryption with the same key across both Regions.

Less operational complexity, as the application does not have to handle different keys in each Region.

Seamless replication and decryption of data across Regions.

upvoted 1 times

Clpsz 1 year, 4 months ago

Selected Answer: D

B: client side encryption is not correct

upvoted 1 times

AshishDhole 1 year, 4 months ago

Selected Answer: B

he data in both S3 buckets must be encrypted and decrypted with the same KMS key. The data and the key must be stored in each of the two Regions --> Single Key for both region. So multi region key will be the answer

upvoted 1 times

Krenil 1 year, 4 months ago

Selected Answer: D

Correct Ans: D

Option B is not the best solution because it involves client-side encryption, which adds unnecessary complexity and operational overhead. Here's why:

Client-side encryption means that the application itself must handle the encryption and decryption of data before it is uploaded to S3 and after it is downloaded. This requires additional development work, key management, and the management of encryption and decryption logic in the application.

In contrast, SSE-KMS (option D) allows Amazon S3 to manage the encryption and decryption process automatically on the server side, which reduces operational overhead and simplifies the implementation.

upvoted 1 times

Tika57 1 year, 4 months ago

Selected Answer: D

B and D are in short list (using customer managed key and KMS) but B requires that each client encrypts the data before sending it, whereas the D is totally transparent so with "LEAST overhead"

upvoted 1 times

dariar 1 year, 4 months ago

Selected Answer: D

For me the answer is D.

It could also be the B-answer, but when it asks "Configure the application", it sounds like the "operational overhead", so it's eliminatory for me

upvoted 1 times

AtiiF 1 year, 4 months ago

Selected Answer: D

=====> D

upvoted 1 times

V2910 1 year, 4 months ago

Selected Answer: D

Difference between B and D is server side encryption

upvoted 1 times

A company recently launched a variety of new workloads on Amazon EC2 instances in its AWS account. The company needs to create a strategy to access and administer the instances remotely and securely. The company needs to implement a repeatable process that works with native AWS services and follows the AWS Well-Architected Framework.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use the EC2 serial console to directly access the terminal interface of each instance for administration.
- B. Attach the appropriate IAM role to each existing instance and new instance. Use AWS Systems Manager Session Manager to establish a remote SSH session. **Most Voted**
- C. Create an administrative SSH key pair. Load the public key into each EC2 instance. Deploy a bastion host in a public subnet to provide a tunnel for administration of each instance.
- D. Establish an AWS Site-to-Site VPN connection. Instruct administrators to use their local on-premises machines to connect directly to the instances by using SSH keys across the VPN tunnel.

Correct Answer: B

Community vote distribution

B (95%) 4%

Comments

BoboChow **Highly Voted** 3 years, 7 months ago

Selected Answer: B

How can Session Manager benefit my organization?

Ans: No open inbound ports and no need to manage bastion hosts or SSH keys

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager.html>

upvoted 25 times

Nightducky 3 years, 6 months ago

Do you know what from the question is it Windows or Linux EC2. I think not so how you want to do SSH session for Windows?

Answer is C

upvoted 3 times

JayBee65 3 years, 6 months ago

Session Manager provides support for Windows, Linux, and macOS from a single tool

upvoted 10 times

sohailn 2 years, 10 months ago

session manager works with linux, windows, and mac too

upvoted 4 times

TienHuynh 2 years, 11 months ago

"Cross-platform support for Windows, Linux, and macOS"

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager.html>

upvoted 3 times

cookieMr **Highly Voted** 2 years, 11 months ago

Selected Answer: B

Option A provides direct access to the terminal interface of each instance, but it may not be practical for administration purposes and can be cumbersome to manage, especially for multiple instances.

Option C adds operational overhead and introduces additional infrastructure that needs to be managed, monitored, and secured. It also requires SSH key management and maintenance.

Option D is complex and may not be necessary for remote administration. It also requires administrators to connect from their local on-premises machines, which adds complexity and potential security risks.

Therefore, option B is the recommended solution as it provides secure, auditable, and repeatable remote access using IAM roles and AWS Systems Manager Session Manager, with minimal operational overhead.

upvoted 9 times

Clouddon Most Recent 11 months, 2 weeks ago

Selected Answer: B

Use AWS Systems Manager Session Manager with IAM roles attached to EC2 instances for secure, scalable, keyless administration — fully compliant with AWS best practices and requiring the least operational overhead.

upvoted 1 times

Kazmin 1 year, 2 months ago

Selected Answer: A

looks correct

upvoted 1 times

Rcosmos 1 year, 5 months ago

Selected Answer: B

A Opção B é a solução mais simples, segura e de menor sobrecarga operacional. O uso do AWS Systems Manager Session Manager elimina a necessidade de abrir portas, gerenciar chaves SSH ou configurar infraestrutura adicional, alinhando-se ao AWS Well-Architected Framework.

upvoted 1 times

satyaamm 1 year, 5 months ago

Selected Answer: B

Creating an IAM Role is the most suitable as others do not meet the criteria of less operational overhead.

upvoted 1 times

safa_123 1 year, 7 months ago

B. Create a customer managed multi-Region KMS key. Create an S3 bucket in each Region. Configure replication between the S3 buckets. Configure the application to use the KMS key with client-side encryption.

Explanation:

Multi-Region KMS Key: AWS KMS supports multi-Region keys, which can be replicated across Regions, enabling encryption and decryption in multiple Regions using the "same" KMS key (though technically it's a replica). This meets the requirement of using the same key in both Regions without additional management overhead for separate keys.

Client-side encryption: Configuring the application to use the KMS key with client-side encryption ensures that the data is encrypted before it is sent to Amazon S3, and decrypted when retrieved.

upvoted 2 times

MortisG 1 year ago

This is perfect for the previous question.

upvoted 3 times

zied007 1 year, 9 months ago

Selected Answer: B

Answer is B

upvoted 1 times

ManikRoy 2 years, 1 month ago

Selected Answer: B

Session Manager provides secure and auditable node management without the need to open inbound ports, maintain bastion hosts, or manage SSH keys.

upvoted 4 times

awsgeek75 2 years, 5 months ago

Selected Answer: B

A: Serial console is for device direct connection to peripherals and monitor boot etc.

C: Workable solution but a lot of overhead

D: Too much overhead for everyone

B: Managed product for this purpose so least overhead.

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager.html>

upvoted 2 times

A_jaa 2 years, 5 months ago

Selected Answer: B

Answer-B

upvoted 1 times

AWSSStudyBuddy 2 years, 7 months ago

I go with option B. Here's why--- IAM Roles: Without SSH keys or shared passwords, securely provide access to EC2 instances and AWS services.
upvoted 4 times

AWSSStudyBuddy 2 years, 7 months ago

Without requiring direct SSH connection, securely access and control EC2 instances with AWS Systems Manager Session Manager.

Least Operational Overhead: An effective and fully managed method of managing instances.

Well-Architected Framework: Complies with performance, security, and reliability best practices from AWS.

Cons of alternative options:

Option A: The automation and flexibility required for secure administration at scale are not provided by using the EC2 serial terminal directly.

Option C: There is more operational overhead and complexity when a bastion host is deployed.

Option D: For secure instance administration, setting up an AWS Site-to-Site VPN connection is too difficult and not the optimal approach.

In conclusion, Option B is suggested as the best option given the given circumstances.

upvoted 5 times

Guru4Cloud 2 years, 10 months ago

Selected Answer: B

This solution meets all of the requirements with the LEAST operational overhead. It is repeatable, uses native AWS services, and follows the AWS Well-Architected Framework.

Repeatable: The process of attaching an IAM role to an EC2 instance and using Systems Manager Session Manager to establish a remote SSH session is repeatable. This can be easily automated, so that new instances can be provisioned and administrators can connect to them securely without any manual intervention.

upvoted 3 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: B

With AWS Systems Manager Session Manager, you can manage your Amazon Elastic Compute Cloud (Amazon EC2) instances, edge devices, on-premises servers, and virtual machines (VMs). You can use either an interactive one-click browser-based shell or the AWS Command Line Interface (AWS CLI). It provides secure and auditable node management without the need to open inbound ports, maintain bastion hosts, or manage SSH keys.

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager.html#:~:text=RSS-,Session%20Manager,-is%20a%20fully>

upvoted 3 times

james2033 2 years, 10 months ago

Selected Answer: B

Keyword "access and administer the instances remotely and securely" See "AWS Systems Manager Session Manager at "

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager.html> .

upvoted 2 times

miki111 2 years, 11 months ago

Option B is the right answer for this.

upvoted 1 times

TienHuynh 2 years, 11 months ago

Selected Answer: B

- +Centralized access control to managed nodes using IAM policies
- +No open inbound ports and no need to manage bastion hosts or SSH keys
- +Cross-platform support for Windows, Linux, and macOS

upvoted 2 times

A company is hosting a static website on Amazon S3 and is using Amazon Route 53 for DNS. The website is experiencing increased demand from around the world. The company must decrease latency for users who access the website. Which solution meets these requirements MOST cost-effectively?

- A. Replicate the S3 bucket that contains the website to all AWS Regions. Add Route 53 geolocation routing entries.
- B. Provision accelerators in AWS Global Accelerator. Associate the supplied IP addresses with the S3 bucket. Edit the Route 53 entries to point to the IP addresses of the accelerators.
- C. Add an Amazon CloudFront distribution in front of the S3 bucket. Edit the Route 53 entries to point to the CloudFront distribution. **Most Voted**
- D. Enable S3 Transfer Acceleration on the bucket. Edit the Route 53 entries to point to the new endpoint.

Correct Answer: C

Community vote distribution

C (100%)

Comments

cookieMr **Highly Voted** 2 years, 11 months ago

Selected Answer: C

Option A (replicating the S3 bucket to all AWS Regions) can be costly and complex, requiring replication of data across multiple Regions and managing synchronization. It may not provide a significant latency improvement compared to the CloudFront solution.

Option B (provisioning accelerators in AWS Global Accelerator) can be more expensive as it adds an extra layer of infrastructure (accelerators) and requires associating IP addresses with the S3 bucket. CloudFront already includes global edge locations and provides similar acceleration capabilities.

Option D (enabling S3 Transfer Acceleration) can help improve upload speed to the S3 bucket but may not have a significant impact on reducing latency for website visitors.

Therefore, option C is the most cost-effective solution as it leverages CloudFront's caching and global distribution capabilities to decrease latency and improve website performance.

upvoted 43 times

TariqKipkemei **Highly Voted** 2 years, 10 months ago

Keywords:

Global, Reduce latency, S3, Static Website, Cost effective = Amazon CloudFront

upvoted 6 times

Clouddon **Most Recent** 11 months, 2 weeks ago

Selected Answer: C

Deploy Amazon CloudFront in front of the S3 bucket and update Route 53 to point to the CloudFront distribution for the most cost-effective and low-latency global access to your static website.

upvoted 1 times

AshishDhole 1 year, 4 months ago

Selected Answer: C

Cost effective solution + Static Website = Cloud Front + S3

upvoted 2 times

satyaamm 1 year, 5 months ago

Selected Answer: C

CloudFront helps reduce latency most cost-effectively.

upvoted 1 times

MGKYAING 1 year, 5 months ago

Selected Answer: C

Global Content Delivery:

CloudFront is a content delivery network (CDN) that caches content at edge locations around the world, reducing latency by serving content from the nearest edge location to the user.

Seamless Integration:

CloudFront integrates easily with S3 buckets. It fetches and caches the static content while reducing the load on the origin S3 bucket.

Cost-Effectiveness:

CloudFront charges based on data transfer and requests, which is often more economical than replicating S3 buckets across Regions or using other solutions that require additional infrastructure.

Performance and Scalability:

CloudFront handles spikes in traffic without additional configuration or cost overhead for provisioning and managing infrastructure.

upvoted 1 times

zied007 1 year, 9 months ago

Selected Answer: C

Answer is C

upvoted 1 times

hb0011 2 years ago

Some of these questions seem too easy for SAA. These seem like Cloud Practitioner questions...

upvoted 3 times

TilTil 2 years, 2 months ago

Cloudfront is a lovely and affordable CDN for static content.

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: C

S3 static website so CloudFront is the best CDN solution for low cost and low latency

A: Very expensive way of doing this

B: Makes no sense

D: Transfer Acc is for upload boosting

C: CloudFront literally solves this problem

upvoted 2 times

A_jaa 2 years, 5 months ago

Selected Answer: C

Answer-C

upvoted 1 times

Ruffyit 2 years, 7 months ago

Option A (replicating the S3 bucket to all AWS Regions) can be costly and complex, requiring replication of data across multiple Regions and managing synchronization. It may not provide a significant latency improvement compared to the CloudFront solution.

Option B (provisioning accelerators in AWS Global Accelerator) can be more expensive as it adds an extra layer of infrastructure (accelerators) and requires associating IP addresses with the S3 bucket. CloudFront already includes global edge locations and provides similar acceleration capabilities

Option D (enabling S3 Transfer Acceleration) can help improve upload speed to the S3 bucket but may not have a significant impact on reducing latency for website visitors.

Therefore, option C is the most cost-effective solution as it leverages CloudFront's caching and global distribution capabilities to decrease latency and improve website performance.

upvoted 2 times

Guru4Cloud 2 years, 10 months ago

Selected Answer: C

Amazon CloudFront is a content delivery network (CDN) service that distributes content globally to reduce latency. By setting up a CloudFront distribution in front of the S3 bucket hosting the static website, you can take advantage of its edge locations around the world to deliver content from the nearest location to the users, reducing the latency they experience.

CloudFront automatically caches and replicates content to its edge locations, resulting in faster delivery and lower latency for users worldwide. This solution is highly effective in optimizing performance while keeping costs under control because CloudFront charges are based on actual data transfer and requests, and the pay-as-you-go pricing model ensures that you only pay for what you use.

upvoted 5 times

james2033 2 years, 10 months ago

Selected Answer: C

Keyword "Amazon CloudFront" (C).
upvoted 1 times

miki111 2 years, 11 months ago

Option C is the right answer for this.
upvoted 1 times

miki111 2 years, 11 months ago

Option C is the right answer for this.
upvoted 1 times

TienHuynh 2 years, 11 months ago

Selected Answer: C

key words:
-around the world
-decrease latency
-most cost-effective

answer is C
upvoted 2 times

A company maintains a searchable repository of items on its website. The data is stored in an Amazon RDS for MySQL database table that contains more than 10 million rows. The database has 2 TB of General Purpose SSD storage. There are millions of updates against this data every day through the company's website.

The company has noticed that some insert operations are taking 10 seconds or longer. The company has determined that the database storage performance is the problem.

Which solution addresses this performance issue?

- A. Change the storage type to Provisioned IOPS SSD. **Most Voted**
- B. Change the DB instance to a memory optimized instance class.
- C. Change the DB instance to a burstable performance instance class.
- D. Enable Multi-AZ RDS read replicas with MySQL native asynchronous replication.

Correct Answer: A

Community vote distribution

A (94%)

6%

Comments

pazabal **Highly Voted** 3 years, 5 months ago

Selected Answer: A

A: Made for high levels of I/O ops for consistent, predictable performance.

B: Can improve performance of insert ops, but it's a storage performance rather than processing power problem

C: for moderate CPU usage

D: for scale read-only replicas and doesn't improve performance of insert ops on the primary DB instance

upvoted 43 times

cookieMr **Highly Voted** 2 years, 11 months ago

Selected Answer: A

Option B (changing the DB instance to a memory optimized instance class) focuses on improving memory capacity but may not directly address the storage performance issue.

Option C (changing the DB instance to a burstable performance instance class) is suitable for workloads with varying usage patterns and burstable performance needs, but it may not provide consistent and predictable performance for heavy write workloads.

Option D (enabling Multi-AZ RDS read replicas with MySQL native asynchronous replication) is a solution for high availability and read scaling but does not directly address the storage performance issue.

Therefore, option A is the most appropriate solution to address the performance issue by leveraging Provisioned IOPS SSD storage type, which provides consistent and predictable I/O performance for the Amazon RDS for MySQL database.

upvoted 28 times

Clouddon **Most Recent** 11 months, 2 weeks ago

Selected Answer: A

Upgrade your RDS MySQL storage from General Purpose SSD (gp2) to Provisioned IOPS SSD (io1/io2) to achieve consistent, high-performance write throughput — the best solution for addressing your current insert performance issue.

upvoted 1 times

JA2018 1 year, 7 months ago

Selected Answer: A

Hint in the stem: DB insert operations taking 10 seconds or longer => current General Purpose SSD storage cannot providing sufficient IOPS required for the workload. Provisioned IOPS SSD can significantly improve storage performance and reduce latency for DB operations.

Provisioned IOPS SSD: Designed to deliver fast, predictable, and consistent I/O performance. This type of performance is very crucial for databases with high transaction rates and frequent updates. It allows users to provision a specific level of IOPS to meet the performance requirements of their specific workload.

upvoted 1 times

PaulGa 1 year, 9 months ago

Selected Answer: A

Ans A, as per stephyfresh13:

"Provisioned IOPS SSD: This storage type is designed to deliver fast, predictable, and consistent I/O performance, which is crucial for databases with high transaction rates and frequent updates. It allows you to provision a specific level of IOPS to meet the performance requirements of your workload."

upvoted 4 times

stephyfresh13 1 year, 11 months ago

A. Change the storage type to Provisioned IOPS SSD.

Here's why:

Provisioned IOPS SSD: This storage type is designed to deliver fast, predictable, and consistent I/O performance, which is crucial for databases with high transaction rates and frequent updates. It allows you to provision a specific level of IOPS to meet the performance requirements of your workload.

Current Issue: The problem with insert operations taking 10 seconds or longer indicates that the current General Purpose SSD storage is not providing sufficient IOPS for the workload. Provisioned IOPS SSD can significantly improve storage performance and reduce latency for database operations.

upvoted 4 times

sudohogan 2 years, 1 month ago

B is the answer, if the company decided that storage is the problem then an IOPS SSD (storage) is the solution and not a memory optimiser.

upvoted 1 times

TilTil 2 years, 2 months ago

Selected Answer: B

A. IOPS is about increasing the number of input connections so you can handle more requests. Which may not be the issue.

B. Having a memory optimized class provides more RAM to execute the queries which take upto 10 secs to complete. More RAM means they can execute faster.

C and D are distractors. They deal with high availability and timely scalability which are not issues here.

upvoted 2 times

MortisG 1 year ago

It says insert operations are taking up to 10 secs. This is clearly I/O so optimizing memory won't help. Correct answer is A.

upvoted 1 times

awsgeek75 2 years, 5 months ago

Selected Answer: A

Database storage is issue so

BD is irrelevant

C is for performance boost (CPU) which won't help with storage issues

A Fix the storage issue

upvoted 2 times

A_jaa 2 years, 5 months ago

Selected Answer: A

Answer-A

upvoted 1 times

bujuman 2 years, 5 months ago

Selected Answer: B

Do not misconsider "database storage performance is the problem". I beleive the correct asnwer is B Due too the fact that Mmemory Optimized EC2 instance family is designed for DB servers perf.

upvoted 1 times

pentium75 2 years, 5 months ago

But the stem clearly says that storage performance, NOT memory performance, is the problem. More memory won't increase storage performance.

upvoted 6 times

aptx4869 2 years, 7 months ago

Selected Answer: A

A is correct answer because it is talking about storage and transaction speed is slow due to it, should change to iops storage instead.

upvoted 2 times

Ruffyt 2 years, 7 months ago

A: Made for high levels of I/O opps for consistent, predictable performance.
B: Can improve performance of insert opps, but it's a storage performance rather than processing power problem
C: for moderate CPU usage
D: for scale read-only replicas and doesn't improve performance of insert opps on the primary DB instance
upvoted 1 times

Rhydian25 2 years ago

Copy-paste from pazabal's answer...
upvoted 1 times

AWSStudyBuddy 2 years, 7 months ago

I go with option A. Using Amazon Provisioned IOPS (PIOPS) SSD storage is the best way to solve the performance issue of insert operations taking 10 seconds or longer on an Amazon RDS for MySQL database table with more than 10 million rows and 2 TB of General Purpose SSD storage.

A high-performance storage solution with reliable throughput and minimal latency is PIOPS SSD storage. Workloads like insert operations, which demand high I/O performance, are ideally suited for it.

upvoted 3 times

tom_cruise 2 years, 8 months ago

Selected Answer: A

Key: database storage performance is the problem.
upvoted 1 times

awsleffe 2 years, 8 months ago

Selected Answer: A

Option A is answer - A. Change the storage type to Provisioned IOPS SSD.
The company's issue is related to storage performance, specifically with insert operations. This suggests that the I/O operations are the bottleneck.

Provisioned IOPS SSD storage type is designed to handle the kind of workload the company is experiencing and should help improve the performance of insert operations.

upvoted 3 times

awashenko 2 years, 8 months ago

Selected Answer: A

"The company has determined that the database storage performance is the problem."
This is the key statement in the question. Otherwise I would have selected B but this statement here makes A correct.
upvoted 2 times

A company has thousands of edge devices that collectively generate 1 TB of status alerts each day. Each alert is approximately 2 KB in size. A solutions architect needs to implement a solution to ingest and store the alerts for future analysis. The company wants a highly available solution. However, the company needs to minimize costs and does not want to manage additional infrastructure. Additionally, the company wants to keep 14 days of data available for immediate analysis and archive any data older than 14 days.

What is the MOST operationally efficient solution that meets these requirements?

A. Create an Amazon Kinesis Data Firehose delivery stream to ingest the alerts. Configure the Kinesis Data Firehose stream to deliver the alerts to an Amazon S3 bucket. Set up an S3 Lifecycle configuration to transition data to Amazon S3 Glacier after 14 days. **Most Voted**

B. Launch Amazon EC2 instances across two Availability Zones and place them behind an Elastic Load Balancer to ingest the alerts. Create a script on the EC2 instances that will store the alerts in an Amazon S3 bucket. Set up an S3 Lifecycle configuration to transition data to Amazon S3 Glacier after 14 days.

C. Create an Amazon Kinesis Data Firehose delivery stream to ingest the alerts. Configure the Kinesis Data Firehose stream to deliver the alerts to an Amazon OpenSearch Service (Amazon Elasticsearch Service) cluster. Set up the Amazon OpenSearch Service (Amazon Elasticsearch Service) cluster to take manual snapshots every day and delete data from the cluster that is older than 14 days.

D. Create an Amazon Simple Queue Service (Amazon SQS) standard queue to ingest the alerts, and set the message retention period to 14 days. Configure consumers to poll the SQS queue, check the age of the message, and analyze the message data as needed. If the message is 14 days old, the consumer should copy the message to an Amazon S3 bucket and delete the message from the SQS queue.

Correct Answer: A

Community vote distribution

A (87%)

13%

Comments

Sinaneos **Highly Voted** 3 years, 1 month ago

Selected Answer: A

Definitely A, it's the most operationally efficient compared to D, which requires a lot of code and infrastructure to maintain. A is mostly managed (firehose is fully managed and S3 lifecycles are also managed)

upvoted 41 times

Kelvin_ke 3 years ago

what about the 30 days minimum requirement to transition to S3 glacier?

upvoted 8 times

Abrar2022 2 years, 6 months ago

GLACIER IS 7 DAYS REQUIREMENT NOT 30

upvoted 7 times

studis 2 years, 12 months ago

You can directly migrate from S3 standard to glacier without waiting

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html>

upvoted 5 times

jerry1 1 year, 2 months ago

the 30 day requirement seem only apply to IA but the article did not mention any limitation on glacier

upvoted 1 times

ErnShm 2 years, 6 months ago

the current article doesn't enable the current option, minimum days are 30

upvoted 1 times

Suvam90 2 years, 3 months ago

No, It's not correct, We can change the storage class in day 0 also using lifecycle policy, I implemented in my project, 30 days is just an example.

upvoted 5 times

caffee 2 years, 8 months ago

This constraint is related to moving from Standard to IA/IA-One Zone only. Nothing to do with Glacier

upvoted 5 times

123jh10 **Highly Voted** 3 years, 1 month ago

Selected Answer: A

Only A makes sense operationally.

If you think D, just consider what is needed to move the message from SQS to S3... you are polling daily 14 TB to take out 1 TB... that's no operationally efficient at all.

upvoted 21 times

Clouddon **Most Recent** 11 months, 2 weeks ago

Selected Answer: A

Use Amazon Kinesis Data Firehose to deliver alerts to Amazon S3, and configure S3 Lifecycle rules to transition older data to Glacier after 14 days — achieving the most efficient, scalable, and low-maintenance solution for this use case.

upvoted 1 times

satyaamm 1 year, 5 months ago

Selected Answer: A

Using Kinesis Data Firehose with S3 Life Configuration policies is the most suitable option as we are looking for less operational overhead.

upvoted 1 times

Parul25 1 year, 10 months ago

I understand A is the most operationally efficient option of all but I can't wrap my head around the fact that objects must have a minimum of 30 days before they can transition or expire from Amazon S3. This means that for the first 30 days after an item is created, you cannot transition or remove it. So, how option A can be the best fit?

upvoted 1 times

cheroh_tots 1 year, 10 months ago

The same 30-day minimum applies when you specify a transition from S3 Standard-IA storage to S3 One Zone-IA.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html>

upvoted 3 times

krx59456zslszcom 1 year, 5 months ago

But the docs are also mentioning a table with minimum storage durations beyond the S3 Standard-IA

"[...]for a list of minimum storage duration for all storage class, see Comparing the Amazon S3 storage classes."

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/storage-class-intro.html#sc-compare>

upvoted 1 times

awsgeek75 1 year, 11 months ago

Selected Answer: A

BCD: Additional infra which company doesn't want

A: Firehose for ingestion and delivery to S3. Lifecycle for managing archive. Fully managed and operationally easy solution

upvoted 2 times

A_jaa 1 year, 11 months ago

Selected Answer: A

Answer-A

upvoted 1 times

jjcode 2 years ago

so many words...

upvoted 5 times

OmegaLambda7XL9 2 years ago

That was an easy A. Kinesis Firehose can load data directly to S3 which makes it the most operationally efficient

upvoted 2 times

Ruffyt 2 years, 1 month ago

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html>

upvoted 2 times

tom_cruise 2 years, 2 months ago

Selected Answer: A

Key: MOST operationally efficient solution

upvoted 2 times

David_Ang 2 years, 2 months ago

Selected Answer: A

"A" is simply correct because kinesis firehouse is made for this, SQS standard is not going to support 500 million alerts 2KB each (1 TB) this service is made for requests that are lighter.

upvoted 2 times

Ak9kumar 2 years, 2 months ago

I picked A. Appeared to be right answer.

upvoted 1 times

chandu7024 2 years, 2 months ago

Should be A

upvoted 1 times

TariqKipkemei 2 years, 4 months ago

Selected Answer: A

The MOST operationally efficient option is A.

upvoted 1 times

james2033 2 years, 4 months ago

Selected Answer: A

Keyword "Amazon S3 Glacier" (A).

upvoted 1 times

miki111 2 years, 4 months ago

Option A is the right answer for this.

upvoted 1 times

A company's application integrates with multiple software-as-a-service (SaaS) sources for data collection. The company runs Amazon EC2 instances to receive the data and to upload the data to an Amazon S3 bucket for analysis. The same EC2 instance that receives and uploads the data also sends a notification to the user when an upload is complete. The company has noticed slow application performance and wants to improve the performance as much as possible.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create an Auto Scaling group so that EC2 instances can scale out. Configure an S3 event notification to send events to an Amazon Simple Notification Service (Amazon SNS) topic when the upload to the S3 bucket is complete.
- B. Create an Amazon AppFlow flow to transfer data between each SaaS source and the S3 bucket. Configure an S3 event notification to send events to an Amazon Simple Notification Service (Amazon SNS) topic when the upload to the S3 bucket is complete. **Most Voted**
- C. Create an Amazon EventBridge (Amazon CloudWatch Events) rule for each SaaS source to send output data. Configure the S3 bucket as the rule's target. Create a second EventBridge (Cloud Watch Events) rule to send events when the upload to the S3 bucket is complete. Configure an Amazon Simple Notification Service (Amazon SNS) topic as the second rule's target.
- D. Create a Docker container to use instead of an EC2 instance. Host the containerized application on Amazon Elastic Container Service (Amazon ECS). Configure Amazon CloudWatch Container Insights to send events to an Amazon Simple Notification Service (Amazon SNS) topic when the upload to the S3 bucket is complete.

Correct Answer: B

Community vote distribution

B (66%)

A (34%)

Comments

jdr75 **Highly Voted** 3 years, 2 months ago

Selected Answer: A

It says "LEAST operational overhead" (ie do it in a way it's the less work for me).

If you know a little Amazon AppFlow (see the some videos) you'll see you'll need time to configure and test it, and at the end cope with the errors during the extraction and load the info to the target.

The customer in the example ALREADY has some EC2 that do the work, the only problem is the performance, that WILL be improved scaling out and adding a queue (SNS) to decouple the work of notify the user.

The operational load of doing this is LESS that configuring AppFlow.

upvoted 49 times

awsgeek75 2 years, 4 months ago

With AWS, a managed service is "less operational overhead" regardless of the complexity of the setup. AppFlow management is less of a headache when compared to EC2 management so A cannot be correct. EC2 has a setup overhead of OS/Application/Code hooks, security etc. continuous patching/upgrading seems like more than what you'll need to do with SaaS.

B is the correct answer.

upvoted 9 times

nhaastrup 1 year, 5 months ago

No B. if you go to the aws site(<https://aws.amazon.com/appflow/>) -AppFlow is a fully managed integration service that helps you securely transfer data between software as a service (SaaS) applications such as Salesforce, SAP, Google Analytics, Facebook Ads, and ServiceNow, and AWS services such as Amazon Simple Storage Service (S3) and Amazon Redshift in just a "few clicks". remember the question is asking for "the LEAST operational overhead", While Auto Scaling improves performance, it still involves managing EC2 instances, which adds operational overhead compared to a managed service like AppFlow.

upvoted 8 times

pentium75 2 years, 5 months ago

"Operational overhead" refers to operation of the solution once it's deployed, it's not about setting it up. EC2 instances that retrieve data from A and write it to B are nonsense, that's what cloud services are meant for.

upvoted 14 times

the_mellie 2 years ago

Agreed, I hesitated a long time exactly for this reason. However, the requirement is to improve the performance as much as possible. I cannot guarantee the same performance with EC2 than with managed services. That's why I eventually voted B.

upvoted 2 times

Six_Fingered_Jose **Highly Voted** 3 years, 7 months ago

Selected Answer: B

This question just screams AppFlow (SaaS integration)

<https://aws.amazon.com/appflow/>

upvoted 41 times

Six_Fingered_Jose 3 years, 7 months ago

configuring Auto-Scaling also takes time when compared to AppFlow, in AWS's words "in just a few clicks"

> Amazon AppFlow is a fully managed integration service that enables you to securely transfer data between Software-as-a-Service (SaaS) applications like Salesforce, SAP, Zendesk, Slack, and ServiceNow, and AWS services like Amazon S3 and Amazon Redshift, in just a few clicks

upvoted 23 times

Clouddon **Most Recent** 11 months, 2 weeks ago

Selected Answer: B

Use Amazon AppFlow to automate and scale data ingestion from SaaS sources to S3, and configure S3 event notifications to SNS for user alerts — the most efficient, scalable, and low-maintenance solution for this use case.

upvoted 1 times

Kazmin 1 year, 2 months ago

Selected Answer: A

A is correct

upvoted 1 times

Itachi28 1 year, 2 months ago

Selected Answer: B

Option A: Auto Scaling EC2 instances

High operational overhead (still managing EC2 instances).

Not eliminating bottlenecks (EC2 instance is still the data processor).

● Option C: EventBridge Rules

Complex event-driven setup instead of direct integration.

Requires custom event rules for each SaaS source, increasing operational burden.

● Option D: ECS with CloudWatch Events

Migrating to ECS introduces complexity and doesn't directly improve transfer speeds.

Still processes data manually, adding an unnecessary layer.

Amazon AppFlow (Option B) is the best choice because it is fully managed, removes EC2 dependencies, and simplifies integration while improving performance.

upvoted 1 times

Dharmarajan 1 year, 4 months ago

Selected Answer: B

Initially I was inclined towards A, but then I read about AppFlow, which is a specific service to collect SAAS data sources. Also with A, there will be additional maintenance activities for EC2 and application, but AppFlow will completely eliminate that work!

upvoted 1 times

Skyskilo 1 year, 5 months ago

Selected Answer: A

Key words here is Least Operational Overhead so A for me

upvoted 1 times

vuongdo 1 year, 5 months ago

Selected Answer: B

The answer is B, because for answer A you will have to manage the EC2 infrastructure, while AppFlow is a serverless service, you will not need to manage the infrastructure

upvoted 3 times

Ilccing 1 year, 5 months ago

Selected Answer: B

i support AppFlow, the original can be saas, and the target can be s3.

upvoted 2 times

f51a8bd 1 year, 7 months ago

B. Cree un flujo de Amazon AppFlow para transferir datos entre cada fuente de SaaS y el depósito de S3. Configure una notificación de eventos de S3 para enviar eventos a un tema de Amazon Simple Notification Service (Amazon SNS) cuando se complete la carga al depósito de S3.

Explicación:

Amazon AppFlow permite la integración directa con múltiples fuentes de SaaS y automatiza el flujo de datos hacia Amazon S3, eliminando la necesidad de instancias de EC2 para recibir y cargar datos. Esto reduce la carga en las instancias de EC2 y simplifica el proceso de integración y recolección de datos.

Notificación de eventos de S3 y SNS : Configurar eventos en Amazon S3 permite enviar notificaciones directamente a un tema de Amazon SNS cuando los datos se han cargado, lo cual asegura que los usuarios recibirán notificaciones sin necesidad de código personalizado en EC2.

upvoted 2 times

ChymKuBoy 1 year, 8 months ago

Selected Answer: B

B for sure

upvoted 1 times

rar75024 1 year, 8 months ago

direct connect maybe correct answer but we are making too many assumptions regarding direct connect network requirements.

upvoted 1 times

effiecancode 1 year, 11 months ago

Has to be appflow because of the SaaS integration

upvoted 2 times

jatric 1 year, 11 months ago

Selected Answer: B

First i thought it would be A but with more study found if its SaaS and need to choose between EC2 and AWS managed service one should always choose AWS managed service. So AppFlow seems more appropriate here.

upvoted 4 times

SMALLE 2 years, 4 months ago

Selected Answer: B

Amazon AppFlow is a fully managed integration service that helps you securely transfer data between software as a service (SaaS) applications such as Salesforce, SAP, Google Analytics, Facebook Ads, and ServiceNow, and AWS services such as Amazon Simple Storage Service (S3) and Amazon Redshift in just a few clicks.

<https://aws.amazon.com/appflow/>

upvoted 5 times

awsgeek75 2 years, 4 months ago

Selected Answer: B

<https://aws.amazon.com/appflow/>

"With Amazon AppFlow automate bi-directional data flows between SaaS applications and AWS services in just a few clicks."

If you want to pass the exam, choose B, regardless of your personal experience! Always use AWS managed services for "least operational overhead"

upvoted 6 times

OmegaLambda7XL9 2 years, 6 months ago

Yea, I think this question is looking for Amazon Appflow. I also feel like it would be easier to set up Autoscaling for the already existing EC2 instances in the short term but then the fact that this software integrates with a lot of SaaS services means using Amazon Appflow will work reduce operational overhead in the long term

upvoted 5 times

A company runs a highly available image-processing application on Amazon EC2 instances in a single VPC. The EC2 instances run inside several subnets across multiple Availability Zones. The EC2 instances do not communicate with each other. However, the EC2 instances download images from Amazon S3 and upload images to Amazon S3 through a single NAT gateway. The company is concerned about data transfer charges.

What is the MOST cost-effective way for the company to avoid Regional data transfer charges?

- A. Launch the NAT gateway in each Availability Zone.
- B. Replace the NAT gateway with a NAT instance.
- C. Deploy a gateway VPC endpoint for Amazon S3. **Most Voted**
- D. Provision an EC2 Dedicated Host to run the EC2 instances.

Correct Answer: C

Community vote distribution

C (99%)

Comments

SilentMilli **Highly Voted** 3 years, 5 months ago

Selected Answer: C

Deploying a gateway VPC endpoint for Amazon S3 is the most cost-effective way for the company to avoid Regional data transfer charges. A gateway VPC endpoint is a network gateway that allows communication between instances in a VPC and a service, such as Amazon S3, without requiring an Internet gateway or a NAT device. Data transfer between the VPC and the service through a gateway VPC endpoint is free of charge, while data transfer between the VPC and the Internet through an Internet gateway or NAT device is subject to data transfer charges. By using a gateway VPC endpoint, the company can reduce its data transfer costs by eliminating the need to transfer data through the NAT gateway to access Amazon S3. This option would provide the required connectivity to Amazon S3 and minimize data transfer charges.

upvoted 103 times

Bmarodi 3 years ago

Very good explanation!

upvoted 9 times

johne42 2 years, 9 months ago

<https://aws.amazon.com/blogs/storage/managing-amazon-s3-access-with-vpc-endpoints-and-s3-access-points/>

upvoted 5 times

OmegaLambda7XL9 2 years, 6 months ago

Precisely

upvoted 3 times

Buruguduystunstugudunstuy **Highly Voted** 3 years, 5 months ago

Selected Answer: C

Option C (correct). Deploy a gateway VPC endpoint for Amazon S3.

A VPC endpoint for Amazon S3 allows you to access Amazon S3 resources within your VPC without using the Internet or a NAT gateway. This means that data transfer between your EC2 instances and S3 will not incur Regional data transfer charges.

Option A (wrong), launching a NAT gateway in each Availability Zone, would not avoid data transfer charges because the NAT gateway would still be used to access S3.

Option B (wrong), replacing the NAT gateway with a NAT instance, would also not avoid data transfer charges as it would still require using the Internet or a NAT gateway to access S3.

Option D (wrong), provisioning an EC2 Dedicated Host, would not affect data transfer charges as it only pertains to the physical host that the EC2 instances are running on and not the data transfer charges for accessing.

upvoted 6 times

Clouddon Most Recent 11 months, 2 weeks ago

Selected Answer: C

To eliminate unnecessary NAT gateway and Regional data transfer costs when accessing S3, deploy a gateway VPC endpoint for Amazon S3 — the most cost-effective and AWS-recommended solution.

upvoted 1 times

Dharmarajan 1 year, 4 months ago

Selected Answer: C

This one was obvious. with VPC gateway for S3, the traffic will not leave AWS and directly go to S3. Therefore will reduce the cost of outgoing data.

upvoted 1 times

satyaamm 1 year, 5 months ago

Selected Answer: C

Gateway VPC endpoint are most suitable for privately accessing S3 here.

upvoted 1 times

MGKYAING 1 year, 5 months ago

Selected Answer: C

Using a gateway VPC endpoint for Amazon S3 is the most cost-effective solution because it enables instances within the VPC to communicate directly with Amazon S3 without incurring data transfer charges via a NAT gateway. Here's how it addresses the problem:

Avoids NAT Gateway Charges for S3 Access:

NAT gateways incur data processing and data transfer charges when instances access S3.

A gateway VPC endpoint eliminates the need to route traffic through the NAT gateway by enabling direct communication between EC2 instances and S3.

No Data Transfer Costs for Intra-Region S3 Access:

Data transferred between Amazon S3 and resources in the same AWS Region through a gateway VPC endpoint is free of charge.

Highly Available and Managed:

Gateway VPC endpoints are highly available and fully managed by AWS, requiring no additional operational overhead.

upvoted 1 times

atikla 1 year, 5 months ago

Selected Answer: C

Correct

upvoted 1 times

ChymKuBoy 1 year, 8 months ago

Selected Answer: C

C for sure

upvoted 1 times

PaulGa 1 year, 9 months ago

Selected Answer: C

Ans C - excellent explanation by SilentMilli

upvoted 1 times

monkey_aws 1 year, 10 months ago

Selected Answer: C

VPC gateway endpoint is free to use, but only available for S3 and DynamoDB

upvoted 4 times

awsgeek75 2 years, 4 months ago

Selected Answer: C

Gateway VPC allows direct access to S3 without going through public internet. This is the de-facto way to save cost for S3 to VPC traffic.

Correct answer is C

upvoted 3 times

Ruffyit 2 years, 7 months ago

<https://aws.amazon.com/blogs/storage/managing-amazon-s3-access-with-vpc-endpoints-and-s3-access-points/>

upvoted 2 times

ACloud_Guru15 2 years, 8 months ago

Selected Answer: C

Gateway Endpoint best suits the requirement

upvoted 2 times

srinivasmn 2 years, 8 months ago

Answer is C: An S3 VPC endpoint provides a way for an S3 request to be routed through to the Amazon S3 service, without having to connect a subnet to an internet gateway. The S3 VPC endpoint is what's known as a gateway endpoint.

upvoted 2 times

Guru4Cloud 2 years, 10 months ago

Selected Answer: C

the EC2 instances are downloading and uploading images to S3, configuring a gateway VPC endpoint will allow them to access S3 without crossing Availability Zones or regions, eliminating regional data transfer charges

upvoted 2 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: C

Gateway VPC endpoints provide reliable connectivity to Amazon S3 without requiring an internet gateway or a NAT device for your VPC.

upvoted 3 times

miki111 2 years, 10 months ago

Option C is the right answer.

upvoted 2 times

A company has an on-premises application that generates a large amount of time-sensitive data that is backed up to Amazon S3. The application has grown and there are user complaints about internet bandwidth limitations. A solutions architect needs to design a long-term solution that allows for both timely backups to Amazon S3 and with minimal impact on internet connectivity for internal users.

Which solution meets these requirements?

- A. Establish AWS VPN connections and proxy all traffic through a VPC gateway endpoint.
- B. Establish a new AWS Direct Connect connection and direct backup traffic through this new connection. **Most Voted**
- C. Order daily AWS Snowball devices. Load the data onto the Snowball devices and return the devices to AWS each day.
- D. Submit a support ticket through the AWS Management Console. Request the removal of S3 service limits from the account.

Correct Answer: B

Community vote distribution

B (99%)

Comments

Sinaneos **Highly Voted** 3 years, 8 months ago

Selected Answer: B

A: VPN also goes through the internet and uses the bandwidth

C: daily Snowball transfer is not really a long-term solution when it comes to cost and efficiency

D: S3 limits don't change anything here

So the answer is B

upvoted 40 times

Buruguduystunstugudunstuy **Highly Voted** 3 years, 5 months ago

Selected Answer: B

Option B (correct). Establish a new AWS Direct Connect connection and direct backup traffic through this new connection.

AWS Direct Connect is a network service that allows you to establish a dedicated network connection from your on-premises data center to AWS. This connection bypasses the public Internet and can provide more reliable, lower-latency communication between your on-premises application and Amazon S3. By directing backup traffic through the AWS Direct Connect connection, you can minimize the impact on your internet bandwidth and ensure timely backups to S3.

upvoted 39 times

Buruguduystunstugudunstuy 3 years, 5 months ago

Option A (wrong), establishing AWS VPN connections and proxying all traffic through a VPC gateway endpoint, would not necessarily minimize the impact on internet bandwidth as it would still utilize the public Internet to access S3.

Option C (wrong), using AWS Snowball devices, would not address the issue of internet bandwidth limitations as the data would still need to be transferred over the Internet to and from the Snowball devices.

Option D (wrong), submitting a support ticket to request the removal of S3 service limits, would not address the issue of internet bandwidth limitations and would not ensure timely backups to S3.

upvoted 11 times

OmegaLambda7XL9 2 years, 6 months ago

Snowball isn't timely since it takes days after ordering to receive the Snowball devices and days to have it shipped and backed up

upvoted 3 times

Bofi 3 years, 3 months ago

Option C is wrong so is your reason. you do not need internet to load data into Snowball Devices. if you are using snow cone for example, u will connect it to your on-premises device directly for loading and Aws will load it in the cloud. However, it not effective to do that everyday , hence option B is the better choice.

upvoted 4 times

Buruguduystunstugudunstuy 3 years, 3 months ago

You're right Option B is the correct answer. I answered Option B as the correct answer above.

upvoted 1 times

Clouddon **Most Recent** 11 months, 2 weeks ago

Selected Answer: B

Implement AWS Direct Connect to create a dedicated network path for timely, large-scale backups to S3 — ensuring internal users retain sufficient bandwidth while keeping backup operations fast and reliable.

upvoted 1 times

Dharmarajan 1 year, 4 months ago

Selected Answer: B

This one is clear from the given choices. There is no specification of "immediate" need, which means the company can afford to wait for that one month time period required to set up AWS Direct connect.

upvoted 1 times

MGKYAING 1 year, 5 months ago

Selected Answer: B

Establishing an AWS Direct Connect connection provides a dedicated, high-bandwidth, low-latency link between the on-premises environment and AWS. This solution effectively addresses the problem by:

Reducing Internet Bandwidth Usage:

By routing backup traffic through the Direct Connect connection, internal users' internet bandwidth is preserved for other activities.
Ensuring Timely Backups:

Direct Connect provides a consistent and reliable connection with speeds ranging from 50 Mbps to 100 Gbps, ensuring that large volumes of time-sensitive data can be backed up efficiently.

Long-Term Scalability:

This solution supports long-term data transfer needs as the application's data volume grows

upvoted 1 times

Tjazz04 1 year, 6 months ago

Selected Answer: B

Ans. B

For long term solution, using AWS Direct Connect for backup traffic will have its own bandwidth for backing-up data and will not share the bandwidth used by the users.

upvoted 1 times

zied007 1 year, 9 months ago

Selected Answer: B

The answer is B

upvoted 1 times

awsgeek75 2 years, 4 months ago

Selected Answer: B

Direct Connect is the only working solution.

A: VPN uses same bandwidth so doesn't solve anything

C: Snowball devices are physical devices requiring physical shipment so wrong solution

D: There are no S3 service limitations in the account related to this problem

upvoted 2 times

NicolasB 2 years, 5 months ago

Option B.

AWS Direct Connect link your on-premise instance with VPC, and all traffic will bypass your Internet Service Provider.

<https://docs.aws.amazon.com/directconnect/latest/UserGuide/Welcome.html>

upvoted 2 times

Ruffyt 2 years, 7 months ago

AWS Direct Connect is a network service that allows you to establish a dedicated network connection from your on-premises data center to AWS.

This connection bypasses the public Internet and can provide more reliable, lower-latency communication between your on-premises application and Amazon S3. By directing backup traffic through the AWS Direct Connect connection, you can minimize the impact on your internet bandwidth and ensure timely backups to S3.

upvoted 3 times

AWSSStudyBuddy 2 years, 7 months ago

Selected Answer: B

I picked option B, because AWS Direct Connect offers a dedicated, secure, high-performance connection that may circumvent bandwidth restrictions and minimize the impact on internet access, AWS Direct Connect is the ideal choice for backing up data to Amazon S3. Some solutions are not as good because they are not as scalable, reliable, or secure as VPN connections, Snowball devices, or reducing S3 service constraints.

upvoted 3 times

tom_cruise 2 years, 8 months ago

Selected Answer: B

Key: time sensitive. So snowball does not apply here.

upvoted 2 times

srinivasmn 2 years, 8 months ago

Right option is C,, In AWS Direct Connect, the network is not fluctuating and provides a consistent experience, while in AWS VPN the VPN is connected with shared and public networks, so the bandwidth and latency fluctuate. Hence direct connect is better choice than virtual connect.

upvoted 2 times

srinivasmn 2 years, 8 months ago

Typo correction to my my above comment. The right option is B.

upvoted 2 times

chandu7024 2 years, 8 months ago

Option B Correct. Reason is that, Direct connect will not use internet. But it will take good amount of time to establish the connectivity.

upvoted 1 times

Guru4Cloud 2 years, 10 months ago

Selected Answer: B

AWS Direct Connect is a dedicated network connection between your on-premises network and AWS. This provides a private, high-bandwidth connection that is not subject to the same internet bandwidth limitations as traditional internet connections. This will allow for timely backups to Amazon S3 without impacting internet connectivity for internal users.

upvoted 2 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: B

AWS Direct Connect cloud service is the shortest path to your AWS resources. While in transit, your network traffic remains on the AWS global network and never touches the public internet. This reduces the chance of hitting bottlenecks or unexpected increases in latency.

<https://aws.amazon.com/directconnect/#:~:text=The-,AWS%20Direct%20Connect,-cloud%20service%20is>

upvoted 2 times

miki111 2 years, 10 months ago

Option B is the right answer.

upvoted 1 times

A company has an Amazon S3 bucket that contains critical data. The company must protect the data from accidental deletion. Which combination of steps should a solutions architect take to meet these requirements? (Choose two.)

- A. Enable versioning on the S3 bucket. **Most Voted**
- B. Enable MFA Delete on the S3 bucket. **Most Voted**
- C. Create a bucket policy on the S3 bucket.
- D. Enable default encryption on the S3 bucket.
- E. Create a lifecycle policy for the objects in the S3 bucket.

Correct Answer: AB

Community vote distribution

AB (97%)

1%

Comments

Uhrien **Highly Voted** 3 years, 8 months ago

Selected Answer: AB

The correct solution is AB, as you can see here:

<https://aws.amazon.com/it/premiumsupport/knowledge-center/s3-audit-deleted-missing-objects/>

It states the following:

To prevent or mitigate future accidental deletions, consider the following features:

Enable versioning to keep historical versions of an object.

Enable Cross-Region Replication of objects.

Enable MFA delete to require multi-factor authentication (MFA) when deleting an object version.

upvoted 67 times

liams123 1 year, 11 months ago

but it could be C you could use 's3:deleteobject' permission without specific conditions or restricts only to authorized users. B does the same thing tho but B is mainly used to restrict unauthorized access not deletion. does anyone agree. I think it is A and C or A & B

upvoted 7 times

cookieMr **Highly Voted** 2 years, 11 months ago

Selected Answer: AB

Enabling versioning on S3 ensures multiple versions of object are stored in bucket. When object is updated or deleted, new version is created, preserving previous version.

Enabling MFA Delete adds additional layer of protection by requiring MFA device to be present when attempting to delete objects. This helps prevent accidental or unauthorized deletions by requiring extra level of authentication.

C. Creating a bucket policy on S3 is more focused on defining access control and permissions for bucket and its objects, rather than protecting against accidental deletion.

D. Enabling default encryption on S3 ensures that any new objects uploaded to bucket are automatically encrypted. While encryption is important for data security, it does not directly address accidental deletion.

E. Creating lifecycle policy for objects in S3 allows for automated management of objects based on predefined rules. While this can help with data retention and storage cost optimization, it does not directly protect against accidental deletion.

upvoted 13 times

liams123 1 year, 11 months ago

but it could be C you could use 's3:deleteobject' permission without specific conditions or restricts only to authorized users. B does the same thing tho but B is mainly used to restrict unauthorized access not deletion. does anyone agree. I think it is A and C or A & B

upvoted 3 times

Clouddon Most Recent 11 months, 2 weeks ago

Selected Answer: AB

To protect critical data in S3 from accidental deletion, enable versioning and MFA Delete — the two most effective built-in AWS mechanisms for this purpose.

upvoted 1 times

Mrigraj12 1 year, 4 months ago

Selected Answer: AB

c: won't be able to delete anything

d: doesn't change anything

e: baseless

upvoted 1 times

satyaamm 1 year, 5 months ago

Selected Answer: AB

AB is the correct solution here as we need MFA delete for making the deletion process inevitable and also we need S3 bucket versioning for MFA delete.

upvoted 1 times

MGKYAING 1 year, 5 months ago

Selected Answer: AB

To protect critical data in an S3 bucket from accidental deletion, these two features provide robust safeguards:

A. Enable versioning on the S3 bucket:

Versioning allows the bucket to maintain multiple versions of an object. Even if an object is accidentally deleted, the previous version can be restored ensuring data recovery.

B. Enable MFA Delete on the S3 bucket:

MFA Delete adds an additional layer of security by requiring multi-factor authentication for deletion operations. This prevents accidental or unauthorized deletions.

upvoted 1 times

Chr1s_Mrg 1 year, 8 months ago

Selected Answer: AB

To protect critical data in an Amazon S3 bucket from accidental deletion, a solutions architect should take the following steps:

Enable versioning on the S3 bucket: This allows you to recover objects that are accidentally deleted or overwritten by keeping multiple versions of an object.

Enable MFA Delete on the S3 bucket: This adds an extra layer of security by requiring multi-factor authentication (MFA) for delete operations, which helps prevent accidental or unauthorized deletions

upvoted 2 times

PaulGa 1 year, 9 months ago

Selected Answer: BD

Ans A,B - as per 'kwabsAA' 2 months ago

"To protect data from accidental deletion, the correct answers are B and D. Versioning does not prevent accidental deletion; it only allows for recovery after the fact. Multi-Factor Authentication (MFA) helps prevent accidental deletion by requiring an additional confirmation step before deletion, making it deliberate rather than accidental. Option D, which involves encryption, ensures that only individuals with the encryption keys can read or manipulate the data, thus preventing unauthorized access and manipulation, including deletion."

upvoted 2 times

zied007 1 year, 9 months ago

Selected Answer: AB

encryption will not prevent accidental deletions

upvoted 2 times

kwabsAA 1 year, 11 months ago

To protect data from accidental deletion, the correct answers are B and D. Versioning does not prevent accidental deletion; it only allows for recovery after the fact. Multi-Factor Authentication (MFA) helps prevent accidental deletion by requiring an additional confirmation step before deletion, making it deliberate rather than accidental. Option D, which involves encryption, ensures that only individuals with the encryption keys can read or manipulate the data, thus preventing unauthorized access and manipulation, including deletion.

upvoted 3 times

liams123 1 year, 11 months ago

you do realised B and D are only to stop unauthroized people from deleting it and if they did delete it would be on purpose to cause issues for a business. it says accidental id say A and C. idk the questions worded poorly dont trust half the answers on here

upvoted 2 times

liams123 1 year, 11 months ago

but it could be C you could use 's3:deleteobject' permission without specific conditions or restricts only to authorized users. B does the same thing tho but B is mainly used to restrict unauthorized access not deletion. does anyone agree. I think it is A and C or A & B. A allows multiple versions of objects to be stored in the bucket. Even if an object is deleted, its previous versions remain intact and accessible. idk the questions weird i could see how it could be all of them except D & E

upvoted 2 times

kwabsAA 1 year, 11 months ago

BD. For D, When you encrypt data, an unauthorized user (without the encryption key) cannot manipulate the data (ie. decryption, modifying, deletion).

upvoted 2 times

mmrakib 2 years, 3 months ago

Selected Answer: AB

AB will be the correct answer.

upvoted 3 times

sidharthwader 2 years, 3 months ago

This could be done if we enable MFA delete on the bucket but in order to enable this bucket versioning must be done. Hence A and B would be the answer.

upvoted 2 times

liams123 1 year, 11 months ago

but it could be C you could use 's3:deleteobject' permission without specific conditions or restricts only to authorized users. B does the same thing tho but B is mainly used to restrict unauthorized access not deletion. does anyone agree. I think it is A and C or A & B

upvoted 1 times

Conster 2 years, 3 months ago

I am getting so confused about what answers I should study. The answers don't match here or in ChatGPT. Can anyone who just took the exam, and passed, point me in the right direction? TIA!

upvoted 2 times

KRC96 2 years, 1 month ago

chatgpt will help you only if you gave correct prompt.

upvoted 1 times

liams123 1 year, 11 months ago

yea chatgpt said A and C it does make sense. Cause B is mainly used for unauthorized access not deletion. idk this website and certlibrary give some interesting answers makes it hard to know but some questions are difficult

upvoted 1 times

awsgeek75 2 years, 4 months ago

Selected Answer: AB

B: MFA to put an extra step to verify deletion and stop from accidental deletion

A: Versioning for recovery of objects that were deleted accidentally even with MFA

Remember, the solution is not required to STOP from deletion. It just wants to STOP ACCIDENTAL deletion.

CDE offer nothing related to accidental deletion

upvoted 3 times

rt_7777 2 years, 5 months ago

Not sure why Answer is BD. I am trying to rationalize it. What I guess could be to address keyword "critical data" where set default encryption is just enhance the security of stored data but does not prevent from deletion. This will be have 2 options A, B for that. B is make sense to ensure user know what to delete on second layer. For option A, it just help you to audit and recovered what was accidentally deleted but does not "prevent" accidentally delete.

upvoted 1 times

fb4afde 2 years, 6 months ago

Selected Answer: AB

Agree, s3 encryption does not prevent deletion

upvoted 2 times

A company has a data ingestion workflow that consists of the following:

- An Amazon Simple Notification Service (Amazon SNS) topic for notifications about new data deliveries
- An AWS Lambda function to process the data and record metadata

The company observes that the ingestion workflow fails occasionally because of network connectivity issues. When such a failure occurs, the Lambda function does not ingest the corresponding data unless the company manually reruns the job.

Which combination of actions should a solutions architect take to ensure that the Lambda function ingests all data in the future? (Choose two.)

- A. Deploy the Lambda function in multiple Availability Zones.
- B. Create an Amazon Simple Queue Service (Amazon SQS) queue, and subscribe it to the SNS topic. **Most Voted**
- C. Increase the CPU and memory that are allocated to the Lambda function.
- D. Increase provisioned throughput for the Lambda function.
- E. Modify the Lambda function to read from an Amazon Simple Queue Service (Amazon SQS) queue. **Most Voted**

Correct Answer: BE

Community vote distribution

BE (98%)

2

Comments

Incognito13 **Highly Voted** 3 years, 7 months ago

A, C, D options are out, since Lambda is fully managed service which provides high availability and scalability by its own

Answers are B and E

upvoted 27 times

Oluseun 3 years, 3 months ago

There are times you do have to increase lambda memory for improved performance though. But not in this case.

upvoted 4 times

Sinaneos **Highly Voted** 3 years, 8 months ago

Selected Answer: BE

BE so that the lambda function reads the SQS queue and nothing gets lost

upvoted 12 times

Clouddon **Most Recent** 11 months, 2 weeks ago

Selected Answer: BE

By adding SQS as a buffer between SNS and Lambda, and updating Lambda to poll from SQS, you build a highly durable and fault-tolerant ingestor workflow that can survive transient failures without data loss.

upvoted 1 times

AshishDhole 1 year, 4 months ago

Selected Answer: BE

best solution is to decouple the applications using SQS queue.

upvoted 1 times

Rcosmos 1 year, 5 months ago

Selected Answer: BE

Explicação:

Criar uma fila do SQS e inscrevê-la no tópico do SNS (Opção B):

Quando o tópico do SNS envia notificações, elas são colocadas na fila do SQS.

Isso garante que as mensagens sejam armazenadas de forma durável na fila, mesmo que a função do Lambda não esteja disponível para processá-

las imediatamente.

Modificar a função do Lambda para ler da fila do SQS (Opção E):

A função do Lambda é configurada para ser acionada por mensagens na fila do SQS.

Se ocorrer um problema de conectividade ou falha na função, as mensagens permanecerão na fila do SQS até que sejam processadas com sucesso. Isso evita a perda de mensagens e elimina a necessidade de intervenção manual.

upvoted 1 times

satyaamm 1 year, 5 months ago

Selected Answer: BE

Enabling SQS with Lambda is the solution here as we need all data to be ingested in future.

upvoted 1 times

MGKYAING 1 year, 5 months ago

Selected Answer: BE

B. Create an Amazon Simple Queue Service (Amazon SQS) queue, and subscribe it to the SNS topic:

By adding an SQS queue between the SNS topic and the Lambda function, you decouple the notification and processing layers.

Messages are stored reliably in the SQS queue, and they are not lost even if the Lambda function fails temporarily due to network connectivity or other issues.

E. Modify the Lambda function to read from an Amazon Simple Queue Service (Amazon SQS) queue:

Configuring the Lambda function to poll the SQS queue ensures that messages are processed asynchronously. If a processing failure occurs, the message remains in the queue and can be retried without manual intervention.

This ensures all data is eventually processed.

upvoted 4 times

friday_test 1 year, 6 months ago

B,E : SQS is a pull mechanism broker, with at least one delivery, so by using SQS no messages get lost.

upvoted 1 times

Chr1s_Mrg 1 year, 8 months ago

Selected Answer: BE

Create an Amazon Simple Queue Service (Amazon SQS) queue, and subscribe it to the SNS topic: This will decouple the data ingestion process and ensure that messages are not lost if the Lambda function fails to process them immediately

Modify the Lambda function to read from an Amazon Simple Queue Service (Amazon SQS) queue: This ensures that the Lambda function can process the messages from the SQS queue, providing a reliable way to handle data ingestion even if there are network issues.

upvoted 1 times

ChymKuBoy 1 year, 12 months ago

Selected Answer: BE

BE for sure

upvoted 1 times

awsgeek75 2 years, 4 months ago

Selected Answer: BE

BE is correct as SQS ensures the messages are stored in a queue for processing.

A: No issue with Lambda availability so this solution is wrong

C: No issues with CPU or memory so no value added by this step also

D: This is not a provisioning issue so provisioning more Lambda won't solve the re-execution issues. The missed messages will still be lost

upvoted 4 times

OmegaLambda7XL9 2 years, 6 months ago

Since network timeout is the issue here, introduce SQS and read from it , that way when network goes down, data still remains in the queue and when connectivity is back, the lambda function can continue from the last data in the queue

upvoted 4 times

SaurabhTiwari1 1 year, 9 months ago

Nice explanation

upvoted 2 times

SaurabhTiwari1 1 year, 9 months ago

explanation*

upvoted 1 times

Ruffyt 2 years, 7 months ago

the correct combination of actions to ensure that the Lambda function ingests all data in the future is to create an SQS queue and subscribe it to the SNS topic (option B) and modify the Lambda function to read from the SQS queue (option E).

upvoted 1 times

tom_cruise 2 years, 8 months ago

Selected Answer: BE

Key: network connectivity issues
upvoted 2 times

awashenko 2 years, 8 months ago

Selected Answer: BE

This one told you the answer in the answer choices. Just add the word THEN between B and E and there ya go.
upvoted 1 times

Abdou1604 2 years, 10 months ago

B and E , the FAN out model , SQS will help to retrieve the work and delayed processing
upvoted 1 times

Guru4Cloud 2 years, 10 months ago

Selected Answer: BE

B) Create an Amazon Simple Queue Service (Amazon SQS) queue, and subscribe it to the SNS topic.

E) Modify the Lambda function to read from an Amazon Simple Queue Service (Amazon SQS) queue.
upvoted 1 times

A company has an application that provides marketing services to stores. The services are based on previous purchases by store customers. The stores upload transaction data to the company through SFTP, and the data is processed and analyzed to generate new marketing offers. Some of the files can exceed 200 GB in size.

Recently, the company discovered that some of the stores have uploaded files that contain personally identifiable information (PII) that should not have been included. The company wants administrators to be alerted if PII is shared again. The company also wants to automate remediation.

What should a solutions architect do to meet these requirements with the LEAST development effort?

A. Use an Amazon S3 bucket as a secure transfer point. Use Amazon Inspector to scan the objects in the bucket. If objects contain PII, trigger an S3 Lifecycle policy to remove the objects that contain PII.

B. Use an Amazon S3 bucket as a secure transfer point. Use Amazon Macie to scan the objects in the bucket. If objects contain PII, use Amazon Simple Notification Service (Amazon SNS) to trigger a notification to the administrators to remove the objects that contain PII. **Most Voted**

C. Implement custom scanning algorithms in an AWS Lambda function. Trigger the function when objects are loaded into the bucket. If objects contain PII, use Amazon Simple Notification Service (Amazon SNS) to trigger a notification to the administrators to remove the objects that contain PII.

D. Implement custom scanning algorithms in an AWS Lambda function. Trigger the function when objects are loaded into the bucket. If objects contain PII, use Amazon Simple Email Service (Amazon SES) to trigger a notification to the administrators and trigger an S3 Lifecycle policy to remove the objects that contain PII.

Correct Answer: B

Community vote distribution

B (64%)

D (36%)

Comments

Gatt **Highly Voted** 3 years, 7 months ago

I have a problem with answer B. The question says: "automate remediation". B says that you inform the administrator and he removes the data manually, that's not automating remediation. Very weird, that would mean that D is correct - but it's so much harder to implement.

upvoted 49 times

Horaii 3 years, 6 months ago

That is correct, "Automate remediation" is not possible if you chose the B

upvoted 2 times

karbob 3 years, 5 months ago

what about LEAST development effort on custom scanning algorithms and If objects contain PII

upvoted 3 times

karbob 3 years, 5 months ago

"The company wants administrators to be alerted " the accessory follows the principal the principale here is => wants administrators to be alerted

upvoted 3 times

acuaws 3 years, 1 month ago

the problem is... you'd have to write lambda to detect PII? AWS has a product for that and we know that's Macie

upvoted 14 times

dodino 2 years, 6 months ago

D is not correct because it mention SES, not SNS.

upvoted 4 times

grzeev **Highly Voted** 3 years, 7 months ago

Selected Answer: B

Amazon Macie is a data security and data privacy service that uses machine learning (ML) and pattern matching to discover and protect your sensitive data

upvoted 26 times

grzeev 3 years, 7 months ago

Macie automatically detects a large and growing list of sensitive data types, including personally identifiable information (PII) such as names, addresses, and credit card numbers. It also gives you constant visibility of the data security and data privacy of your data stored in Amazon S3

upvoted 13 times

jammy **Most Recent** 1 month, 1 week ago

Selected Answer: B

Correct answer is b.

Why other are incorrect.....

A. Amazon Inspector → Inspector is for scanning EC2 instances, containers, and Lambda functions for vulnerabilities, not for scanning S3 objects for PII.

C. Custom Lambda scanning → Requires building and maintaining your own detection logic, which is complex and error-prone. Not "least development effort."

D. Custom Lambda + SES + Lifecycle policy → Same issue as C, plus SES is for sending emails, not the best fit for automated alerts/remediation. Lifecycle policies are not designed for immediate remediation based on content.

upvoted 1 times

Stpn2me 2 months, 3 weeks ago

Selected Answer: B

The only reason I would choose B is because Macie is proprietary AWS, sort of like with Cisco when you are answering their questions, the answer is always Cisco. But otherwise, according to the question, remediation needs to be automatic, which only leaves D.

upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: B

Amazon Macie + SNS is the least effort, fully managed, and most purpose-fit solution to scan for PII in S3 and alert administrators automatically.

upvoted 1 times

MundiChor 1 year, 2 months ago

Selected Answer: B

Although I selected B on the account of "Least Developmental Effort", However I feel the question itself is wrong. In lieu of being a difficult and tricky question, it contradicts itself too many times.

1. Macie has a S3 file limitation of 20 GB (in case of ZIP archive), here the files are exceeding 200 GB with no indication of breaking the file or compression.

2. The question states that the customer wants to automate remediation; now the word remediation means "to solve" and automate remediation means that there should not be any manual effort. But none of the options provide any solution for that.

I would rate the trust of this question pretty low.

upvoted 3 times

MPG1970 1 year, 2 months ago

Selected Answer: C

Can't be B Macie has a 5Gb limit

C is better than D Lambda can be set up to remove data immediately and AWS state their #1 priority is security.

upvoted 1 times

LovaP 1 year, 3 months ago

Selected Answer: B

You can also leverage Macie integration with Amazon EventBridge and AWS Security Hub to monitor, process, and remediate findings by using other services, applications, and systems.

upvoted 1 times

Vandaman 1 year, 3 months ago

Selected Answer: D

The remediation must be automated

upvoted 1 times

Panknil 1 year, 4 months ago

Selected Answer: B

I think "The company wants administrators to be alerted if PII is shared again" is the key here.. so B is correct answer.

upvoted 1 times

kyd0nix 1 year, 4 months ago

Selected Answer: B

None are correct.

A does not alert the admins.

B does not automate solution.

C and D imply dev effort.

upvoted 2 times

AwsAbhiKumar 1 year, 4 months ago

Selected Answer: D

Macie has file limitation of 5 GB whereas Custom Lambda function allows you to handle files larger than 5 GB, overcoming Macie's limitation.

upvoted 2 times

Dharmarajan 1 year, 4 months ago

Selected Answer: B

B appears most appropriate out of the given option, however it does not address automation of remediation.

However in the view of remediation, the other options do not address it either, so B is most appropriate.

upvoted 1 times

dariar 1 year, 4 months ago

Selected Answer: D

The good answer is D, even if "sensitive data" = "macie".

Macie has limits for analytics files, so 200GB won't pass :

Size of an individual file to analyze:

Adobe Portable Document Format (.pdf) file: 1,024 MB

Apache Avro object container (.avro) file: 8 GB

Apache Parquet (.parquet) file: 8 GB

Email message (.eml) file: 20 GB

GNU Zip compressed archive (.gz or .gzip) file: 8 GB

Microsoft Excel workbook (.xls or .xlsx) file: 512 MB

Microsoft Word document (.doc or .docx) file: 512 MB

Non-binary text file: 20 GB

TAR archive (.tar) file: 20 GB

ZIP compressed archive (.zip) file: 8 GB

Also, the answer B doesn't provide an auto-remediation, the admin still needs to remove the file manually.

Very tricky question, but I think the right fit is D.

upvoted 2 times

AshishDhole 1 year, 4 months ago

Selected Answer: B

PII + S3 == Amazon Macie

upvoted 1 times

AwsAbhiKumar 1 year, 4 months ago

Well this is correct but not in this situation where some file can exceed 200 GB (Macie has limit of 5GB)

upvoted 1 times

Rcosmos 1 year, 5 months ago

Selected Answer: B

Notificações com Amazon SNS:

Quando o Macie detecta PII, ele pode ser configurado para acionar uma notificação via SNS, alertando os administradores para tomar as ações necessárias.

Automação parcial:

Embora o Macie não remova automaticamente os objetos, ele permite que administradores sejam informados para realizar a correção manualmente garantindo controle sobre os dados.

Mínimo esforço de desenvolvimento:

Essa abordagem utiliza serviços nativos da AWS sem necessidade de scripts ou funções personalizadas, reduzindo significativamente o tempo e o custo de implementação.

upvoted 1 times

satyaamm 1 year, 5 months ago

Selected Answer: B

Using Amazon Macie is most suitable for a S3 bucket and using SNS is also suitable as both of these services resolve the issues with least operational overhead.

upvoted 1 times

A company needs guaranteed Amazon EC2 capacity in three specific Availability Zones in a specific AWS Region for an upcoming event that will last 1 week.

What should the company do to guarantee the EC2 capacity?

- A. Purchase Reserved Instances that specify the Region needed.
- B. Create an On-Demand Capacity Reservation that specifies the Region needed.
- C. Purchase Reserved Instances that specify the Region and three Availability Zones needed.
- D. Create an On-Demand Capacity Reservation that specifies the Region and three Availability Zones needed. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Incognito013 **Highly Voted** 3 years, 7 months ago

Reserved instances are for long term so on-demand will be the right choice - Answer D
upvoted 35 times

simoneri88 2 years, 4 months ago

Confirmed! Reserved instances require a fixed one-year or three-year commitment. See <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-capacity-reservations.html#capacity-reservations-differences>
upvoted 6 times

tch 1 year, 2 months ago

Amazon EC2 billing and purchasing options
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/instance-purchasing-options.html>
upvoted 1 times

tch 1 year, 2 months ago

Reserved Instances – Reduce your Amazon EC2 costs by making a commitment to a consistent instance configuration, including instance type and Region, for a term of 1 or 3 years.

Capacity Reservations – Reserve capacity for your EC2 instances in a specific Availability Zone.

upvoted 1 times

Buruguduystunstugudunstuy **Highly Voted** 3 years, 5 months ago

Selected Answer: D

CORRECT

Option D. Create an On-Demand Capacity Reservation that specifies the Region and three Availability Zones needed.

An On-Demand Capacity Reservation is a type of Amazon EC2 reservation that enables you to create and manage reserved capacity on Amazon EC2. With an On-Demand Capacity Reservation, you can specify the Region and Availability Zones where you want to reserve capacity, and the number of EC2 instances you want to reserve. This allows you to guarantee capacity in specific Availability Zones in a specific Region.

WRONG

Option A, purchasing Reserved Instances that specify the Region needed, would not guarantee capacity in specific Availability Zones.
Option B, creating an On-Demand Capacity Reservation that specifies the Region needed, would not guarantee capacity in specific Availability Zones
Option C, purchasing Reserved Instances that specify the Region and three Availability Zones needed, would not guarantee capacity in specific Availability Zones as Reserved Instances do not provide capacity reservations.

upvoted 26 times

BlueVolcano1 3 years, 4 months ago

Another reason as to why Reserved Instances aren't the solution here is that you have to commit to either a 1 year or 3 year term, not 1 week.
upvoted 23 times

Clouddon Most Recent 11 months, 2 weeks ago

Selected Answer: D

For guaranteed EC2 capacity in specific Availability Zones for a short-term event, use On-Demand Capacity Reservations with AZs specified. This ensures your EC2 instances can launch when needed, even during peak demand.

upvoted 1 times

tch 1 year, 2 months ago

Reserved Instances – Reduce your Amazon EC2 costs by making a commitment to a consistent instance configuration, including instance type and Region, for a term of 1 or 3 years.

Capacity Reservations – Reserve capacity for your EC2 instances in a specific Availability Zone.

do not confuse

upvoted 1 times

Rcosmos 1 year, 5 months ago

Selected Answer: D

A Opção D (Reserva de capacidade sob demanda especificando a região e as zonas de disponibilidade) é a melhor solução, pois garante a capacidade necessária nas zonas de disponibilidade específicas para um evento de curta duração, atendendo ao requisito de forma eficiente e econômica.

upvoted 1 times

satyaamm 1 year, 5 months ago

Selected Answer: D

On-demand capacity instances in the required region and required availability zones is the most suitable here.

upvoted 1 times

friday_test 1 year, 6 months ago

Reserved instances offer either a 1 year or 3 year term commitment. So on demand-instances would satisfy the need above. The answer is D , as the question also requires availability zones.

upvoted 1 times

Edawg173 2 years, 2 months ago

i always get these wrong

upvoted 2 times

TiITil 2 years, 2 months ago

Selected Answer: D

Picked C for this one and failed. Reserved Instances are reserved for 1-3 years so On-Demand Reservation makes more sense.

upvoted 3 times

JA2018 1 year, 7 months ago

Check out <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-capacity-reservations.html#capacity-reservations-differences>

Terms for Capacity Reservations: No commitment required. Can be created and canceled as needed.

upvoted 1 times

awsgeek75 2 years, 4 months ago

Selected Answer: D

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-capacity-reservations.html>

"On-Demand Capacity Reservations enable you to reserve compute capacity for your Amazon EC2 instances in a specific Availability Zone for any duration"

upvoted 2 times

awsgeek75 2 years, 4 months ago

Just to avoid any confusion, Reserved Instance also guarantee capacity reservation. However, the reason why we don't need Reserved Instances is because they are reserved for a duration of 1 to 3 years with a lock-in contract (no refunds!). The company is only interested in reserved capacity for a week so D is the best solution.

<https://aws.amazon.com/compare/the-difference-between-on-demand-instances-and-reserved-instances/>

upvoted 2 times

Ruffyit 2 years, 7 months ago

CORRECT

Option D. Create an On-Demand Capacity Reservation that specifies the Region and three Availability Zones needed.

An On-Demand Capacity Reservation is a type of Amazon EC2 reservation that enables you to create and manage reserved capacity on Amazon EC2. With an On-Demand Capacity Reservation, you can specify the Region and Availability Zones where you want to reserve capacity, and the number of EC2 instances you want to reserve. This allows you to guarantee capacity in specific Availability Zones in a specific Region.

upvoted 2 times

awashenko 2 years, 8 months ago

Selected Answer: D

Reserved Instances have a commitment over a year so those are out. Option B only allows you to specify the Region and not the AZ. Therefore, D is the only solution.

upvoted 2 times

Abdou1604 2 years, 10 months ago

its B , On-Demand Capacity Reservation allows you to reserve capacity for Amazon EC2 instances in a specific AWS Region, without specifying specific Availability Zones

upvoted 1 times

Guru4Cloud 2 years, 10 months ago

Selected Answer: D

D is the correct option to guarantee EC2 capacity in specific Availability Zones for a set timeframe.

On-Demand Capacity Reservations allow you to reserve EC2 capacity across specific Availability Zones for any duration. This guarantees you will have access to those resources.

upvoted 2 times

miki111 2 years, 10 months ago

Option D is the right answer.

upvoted 1 times

cookieMr 2 years, 11 months ago

Selected Answer: D

The most appropriate option to guarantee EC2 capacity in three specific Availability Zones in the desired AWS Region for the 1-week event is to create an On-Demand Capacity Reservation that specifies the Region and three Availability Zones (option D).

A. Purchasing Reserved Instances that specify the Region needed does not guarantee capacity in specific Availability Zones.

B. Creating an On-Demand Capacity Reservation without specifying the Availability Zones would not guarantee capacity in the desired zones.

C. Purchasing Reserved Instances that specify the Region and three Availability Zones is not necessary for a short-term event and involves longer-term commitments.

upvoted 4 times

Abrar2022 3 years ago

Reserved instances is for long term

On-demand Capacity reservation enables you to choose specific AZ for any duration

upvoted 1 times

A company's website uses an Amazon EC2 instance store for its catalog of items. The company wants to make sure that the catalog is highly available and that the catalog is stored in a durable location. What should a solutions architect do to meet these requirements?

- A. Move the catalog to Amazon ElastiCache for Redis.
- B. Deploy a larger EC2 instance with a larger instance store.
- C. Move the catalog from the instance store to Amazon S3 Glacier Deep Archive.
- D. Move the catalog to an Amazon Elastic File System (Amazon EFS) file system. **Most Voted**

Correct Answer: D

Community vote distribution

D (93%)

7%

Comments

Six_Fingered_Jose **Highly Voted** 3 years, 7 months ago

Selected Answer: D

keyword is "durable" location
A and B is ephemeral storage
C takes forever so is not HA,
that leaves D
upvoted 53 times

Fakhrudin 2 years, 9 months ago

Yes, if you open EFS home page (<https://aws.amazon.com/efs/>), Amazon state, "Securely and reliably access your files with a fully managed file system designed for 99.99999999 percent (11 9s) durability and up to 99.99 percent (4 9s) of availability."
upvoted 13 times

rajendradba **Highly Voted** 3 years, 8 months ago

Selected Answer: D

Elasticahe is in Memory, EFS is for durability
upvoted 22 times

sirasdf 2 years, 3 months ago

That's wrong. EFS is highly available.
upvoted 4 times

Clouddon **Most Recent** 11 months, 2 weeks ago

Selected Answer: D

To meet the durability and availability requirements for a frequently accessed catalog, Amazon EFS is the best and most cost-effective solution.
upvoted 2 times

MortisG 1 year ago

Selected Answer: D

EFS because of durability and availability. Both initial requirements.
upvoted 2 times

Dharmarajan 1 year, 4 months ago

Selected Answer: D

Easy one. D. none of the other options are meeting the criteria of durability.
upvoted 2 times

satyaammm 1 year, 5 months ago

Selected Answer: D

EFS is most suitable for high availability and persistent storage.

upvoted 2 times

PaulGa 1 year, 9 months ago

Selected Answer: D

Ans D - the clue is its the only option that offers "durable" (Not sure why the author suggests A??)

upvoted 3 times

jatric 1 year, 11 months ago

Selected Answer: D

Durability is the key differentiator

upvoted 2 times

uxrdbean 2 years ago

If they are in instance store now, would EFS maintain the performance necessary for them? I don't think the performance would be similar.

upvoted 1 times

sudohogan 2 years, 1 month ago

The ElastiCache option solves both storage and availability issues.

upvoted 1 times

HectorCosta 2 years, 1 month ago

Selected Answer: D

Cache is not Durable

upvoted 2 times

Edawg173 2 years, 2 months ago

D would be correct why would it give me A. :/

upvoted 2 times

TilTil 2 years, 2 months ago

Selected Answer: D

A and B are distractors.

D is durable and HA.

upvoted 2 times

modehqudah 2 years, 4 months ago

Selected Answer: D

durable location

upvoted 1 times

arslantobe 2 years, 4 months ago

Option C, which suggests moving the catalog to Amazon S3 Glacier Deep Archive, is not a suitable choice for an active catalog that requires high availability and quick access. Glacier Deep Archive is designed for long-term archival and may not provide the low-latency access required for a catalog used in a website.

Therefore, option D is the most appropriate choice for achieving both high availability and durability for the catalog.

upvoted 4 times

vip2 2 years, 5 months ago

Selected Answer: A

SEE <https://aws.amazon.com/elasticache/faqs/#Redis>

upvoted 1 times

Hamso 2 years, 5 months ago

Redis can be made durable, supports failover and multi-AZ deployment, it's effective in catalog use cases. EFS is effective when a shared storage is needed

upvoted 5 times

A company stores call transcript files on a monthly basis. Users access the files randomly within 1 year of the call, but users access the files infrequently after 1 year. The company wants to optimize its solution by giving users the ability to query and retrieve files that are less than 1-year-old as quickly as possible. A delay in retrieving older files is acceptable.

Which solution will meet these requirements MOST cost-effectively?

A. Store individual files with tags in Amazon S3 Glacier Instant Retrieval. Query the tags to retrieve the files from S3 Glacier Instant Retrieval.


B. Store individual files in Amazon S3 Intelligent-Tiering. Use S3 Lifecycle policies to move the files to S3 Glacier Flexible Retrieval after 1 year. Query and retrieve the files that are in Amazon S3 by using Amazon Athena. Query and retrieve the files that are in S3 Glacier by using S3 Glacier Select. **Most Voted**

C. Store individual files with tags in Amazon S3 Standard storage. Store search metadata for each archive in Amazon S3 Standard storage. Use S3 Lifecycle policies to move the files to S3 Glacier Instant Retrieval after 1 year. Query and retrieve the files by searching for metadata from Amazon S3.

D. Store individual files in Amazon S3 Standard storage. Use S3 Lifecycle policies to move the files to S3 Glacier Deep Archive after 1 year. Store search metadata in Amazon RDS. Query the files from Amazon RDS. Retrieve the files from S3 Glacier Deep Archive.

Correct Answer: B

Community vote distribution



Comments

masetromain **Highly Voted** 3 years, 8 months ago

Selected Answer: B

I think the answer is B:

Users access the files randomly

S3 Intelligent-Tiering is the ideal storage class for data with unknown, changing, or unpredictable access patterns, independent of object size or retention period. You can use S3 Intelligent-Tiering as the default storage class for virtually any workload, especially data lakes, data analytics, new applications, and user-generated content.

<https://aws.amazon.com/fr/s3/storage-classes/intelligent-tiering/>

upvoted 45 times

habibi03336 3 years, 3 months ago

It says "S3 Intelligent-Tiering is the ideal storage class for data with unknown, changing, or unpredictable access patterns".

However, the statement says access pattern is predictable. It says there is frequent access about 1 year.

upvoted 2 times

killbots 3 years, 2 months ago

it doesnt say predictable, it says files are accessed random. Random = Unpredictable. Answer is B

upvoted 12 times

CarlosMarin 2 years, 1 month ago

It says "... access the files INFREQUENTLY after 1 year"

upvoted 1 times

lofzee 2 years ago

Helps to read sometimes

upvoted 1 times

PaulEkwem 1 year, 7 months ago

Hey, don't be rude
upvoted 1 times

sachin 3 years, 3 months ago

What about if the file you have not accessed 360 days and intelligent tier moved the file to Glacier and on 364 day you want to access the file instantly ?

I think C is right choice
upvoted 6 times

boringtangent 2 years, 2 months ago

bro u r forgetting cost effectiveness which is the requirement 1 yr in s3 standard will cost more than s3 intelligent tire.
upvoted 3 times

MutiverseAgent 2 years, 11 months ago

Agree, S3 Intelligent-Tiering meets all the requirements. The very important/crucial consideration here to satisfy that all files withing a year are instantly accessible is that the two options "Archive Access" and "Deep Archive Access" are not enabled in the "Archive rule actions" section present in the "Intelligent-Tiering Archive configurations" of the bucket. Those options are not enabled by default so this answer will work.
upvoted 3 times

ssoffline 3 years ago

Answer is C, why not intelligent Tiering

If the Intelligent-Tiering data transitions to Glacier after 180 days instead of 1 year, it would still be a cost-effective solution that meets the requirements.

With files stored in Amazon S3 Intelligent-Tiering, the data is automatically moved to the appropriate storage class based on its access patterns. In this case, if the data transitions to Glacier after 180 days, it means that files that are infrequently accessed beyond the initial 180 days will be stored in Glacier, which is a lower-cost storage option compared to S3 Standard.
upvoted 7 times

pentium75 2 years, 5 months ago

C involves moving "the files to S3 Glacier Instant Retrieval" which is not cost-effective since "a delay in retrieving older files is acceptable."
upvoted 4 times

RupeC 2 years, 11 months ago

With S3 Intelligent-Tiering, you can define rules that determine when objects should be moved from the frequent access tier to the infrequent access tier, or vice versa, within S3 Standard storage classes.
upvoted 2 times

IngenieriaEGlobal 2 years, 8 months ago

The Answer is B. S3 Intelligent-Tiering stores objects in two access tiers: one tier that is optimized for frequent access and another lower-cost tier that is optimized for infrequent access. For a small monthly monitoring and automation fee per object, S3 Intelligent-Tiering monitors access patterns and moves objects that have not been accessed for 30 consecutive days to the infrequent access tier. There are no retrieval fees in S3 Intelligent-Tiering. If an object in the infrequent access tier is accessed later, it is automatically moved back to the frequent access tier. No additional tiering fees apply when objects are moved between access tiers within the S3 Intelligent-Tiering storage class. S3 Intelligent-Tiering is designed for 99.9% availability and 99.999999999% durability, and offers the same low latency and high throughput performance of S3 Standard
upvoted 7 times

Visinho 2 years, 4 months ago

Are you not going to pay for Athena usage?
upvoted 1 times

Lilibell **Highly Voted** 3 years, 8 months ago

The answer is B
upvoted 12 times

dontCarexxx **Most Recent** 4 months, 4 weeks ago

Selected Answer: B

B however deprecated : these questions need some updating !
Source: <https://aws.amazon.com/fr/blogs/storage/how-to-optimize-querying-your-data-in-amazon-s3/>
upvoted 1 times

MultipleWorkerMirroredStrategy 9 months, 1 week ago

Selected Answer: C

To me, it's C because B violates "by giving users the ability to query and retrieve files that are less than 1-year-old as quickly as possible".
upvoted 1 times

Clouddon 11 months, 2 weeks ago

Selected Answer: B

Option B provides cost optimization, query capability, and low-latency access for recent files, with acceptable latency for older data, all with minimal operational effort—making it the best choice for this use case.

upvoted 1 times

c12ab95 1 year ago

Selected Answer: B

Why Option B Works Best:

First Year (Active Access):

S3 Intelligent-Tiering automatically shifts objects between Frequent and Infrequent Access tiers, reducing costs by ~40% if access declines.

Athena queries data directly in S3, avoiding retrieval fees for active tiers.

After 1 Year (Infrequent Access):

Lifecycle policies transition data to Glacier Flexible Retrieval, which supports Glacier Select for SQL-based partial retrievals (avoids restoring entire files).

Cost savings: Flexible Retrieval is cheaper than S3 Standard-IA and allows expedited (1–5 min) or bulk (5–12 hr) retrievals.

Query Flexibility:

Recent files (<1 year) are queried via Athena with millisecond latency.

Older files (>1 year) use Glacier Select to retrieve subsets of data, minimizing retrieval costs and time.

upvoted 1 times

sammo08 1 year, 3 months ago

Selected Answer: B

Option A Incorrect – S3 Glacier Instant Retrieval is more expensive than necessary for frequently accessed files, and tag-based querying is inefficient.

Option C Incorrect – Using S3 Glacier Instant Retrieval after 1 year is more expensive than needed, and storing metadata in S3 is less efficient than Athena.

Option D Incorrect – S3 Glacier Deep Archive has long retrieval times, and maintaining metadata in RDS adds extra costs.

Option B Correct – S3 Intelligent-Tiering optimizes costs, and S3 Glacier Flexible Retrieval balances cost with acceptable retrieval delays, while Athena and S3 Glacier Select enable efficient queries.

upvoted 1 times

dariar 1 year, 4 months ago

Selected Answer: B

The B and C are both good options, but since Athena is free for use with S3, it definitely the B-option.

Athena queries data directly from Amazon S3. There are no additional storage charges for querying your data with Athena.

upvoted 1 times

ricktechie66 1 year, 5 months ago

Selected Answer: C

This solution best meets the requirements while being cost-effective because:

1. ****Initial storage in S3 Standard:**** Provides the fastest access for files less than 1 year old, meeting the requirement for quick retrieval of newer files
2. ****Metadata storage in S3 Standard:**** Enables efficient querying and searching of files without additional database costs
3. ****Lifecycle policy to S3 Glacier Instant Retrieval:**** Provides cost optimization for older files while still maintaining relatively quick retrieval when needed.

upvoted 1 times

Rohan_Butala 1 year, 5 months ago

Selected Answer: B

Option B strikes the right balance between cost-effectiveness and performance, ensuring that frequently accessed data is easily retrievable and older data is archived at lower cost with fast retrieval options for both new and older data.

upvoted 1 times

_adoo 1 year, 7 months ago

If you are like me and thought Athena would make this cost-inefficient:

Athena is free for s3 querying.

It has a cost per TB for SQL Queries.

And a different cost for Apache apps.

<https://aws.amazon.com/athena/pricing/>

upvoted 3 times

JonesNick 1 year, 7 months ago

Keyword : MOST cost-effectively

Athena is great tool for analyzing data in S3. But it comes with the cost.

So answer is C.

upvoted 1 times

jatric 1 year, 11 months ago

Selected Answer: B

cost effective to retrieve the file of 1 year or less. Standard S3 is more cost effective than intelligent tiering.

upvoted 2 times

344bba0 1 year, 11 months ago

Selected Answer: C

The answer is C.

* Intelligent tiering may not guarantee frequent, but random access and fast searches within a year.

* Athena is a great analysis solution, but it is an unnecessary cost for search purposes only.

upvoted 4 times

mknarula 1 year, 11 months ago

Answer is C. The only difference between choice B and C is Glacier Storage class. The question states clearly that "giving users the ability to query and retrieve files that are less than 1-year-old as quickly as possible". This is possible via instant retrieval and not flexible retrieval.

upvoted 2 times

ManikRoy 2 years, 1 month ago

Selected Answer: B

option B is the only one that mentions use of Amazon Athena and Glacier select for querying, So I'll go with it though I would have preferred using S3 standard storage in place of S3 intelligent tiering for the first year.

upvoted 1 times

hro 2 years, 2 months ago

Answer is C - unfortunately. S3 Glacier Instant Retrieval - MOST Cost effective deep storage for retrieving files as quickly as possible.

upvoted 1 times

A company has a production workload that runs on 1,000 Amazon EC2 Linux instances. The workload is powered by third-party software. The company needs to patch the third-party software on all EC2 instances as quickly as possible to remediate a critical security vulnerability.

What should a solutions architect do to meet these requirements?

- A. Create an AWS Lambda function to apply the patch to all EC2 instances.
- B. Configure AWS Systems Manager Patch Manager to apply the patch to all EC2 instances.
- C. Schedule an AWS Systems Manager maintenance window to apply the patch to all EC2 instances.
- D. Use AWS Systems Manager Run Command to run a custom command that applies the patch to all EC2 instances.

Most Voted

Correct Answer: D

Community vote distribution

D (68%)

B (32%)

Comments

tinyfoot Highly Voted 3 years, 7 months ago

The primary focus of Patch Manager, a capability of AWS Systems Manager, is on installing operating systems security-related updates on managed nodes. By default, Patch Manager doesn't install all available patches, but rather a smaller set of patches focused on security. (Ref <https://docs.aws.amazon.com/systems-manager/latest/userguide/patch-manager-how-it-works-selection.html>)

Run Command allows you to automate common administrative tasks and perform one-time configuration changes at scale. (Ref <https://docs.aws.amazon.com/systems-manager/latest/userguide/execute-remote-commands.html>)

Seems like patch manager is meant for OS level patches and not 3rd party applications. And this falls under run command wheelhouse to carry out one-time configuration changes (update of 3rd part application) at scale.

upvoted 74 times

Fakhrudin 2 years, 9 months ago

3rd party applications are also supported by Patch Manager (<https://docs.aws.amazon.com/systems-manager/latest/userguide/patch-manager.html>).

You can use Patch Manager to apply patches for both operating systems and applications. (On Windows Server, application support is limited to updates for applications released by Microsoft.) You can use Patch Manager to install Service Packs on Windows nodes and perform minor version upgrades on Linux nodes. You can patch fleets of Amazon Elastic Compute Cloud (Amazon EC2) instances, edge devices, on-premises servers, and virtual machines (VMs) by operating system type. This includes supported versions of several operating systems, as listed in Patch Manager prerequisites.

upvoted 8 times

Shasha1 Highly Voted 3 years, 6 months ago

D

AWS Systems Manager Run Command allows the company to run commands or scripts on multiple EC2 instances. By using Run Command, the company can quickly and easily apply the patch to all 1,000 EC2 instances to remediate the security vulnerability.

Creating an AWS Lambda function to apply the patch to all EC2 instances would not be a suitable solution, as Lambda functions are not designed to run on EC2 instances. Configuring AWS Systems Manager Patch Manager to apply the patch to all EC2 instances would not be a suitable solution, as Patch Manager is not designed to apply third-party software patches. Scheduling an AWS Systems Manager maintenance window to apply the patch to all EC2 instances would not be a suitable solution, as maintenance windows are not designed to apply patches to third-party software

upvoted 22 times

Clouddon Most Recent 11 months, 2 weeks ago

Selected Answer: D

Option D offers the fastest and most flexible solution to urgently apply third-party patches across a large fleet of EC2 instances with minimal operational overhead.

upvoted 1 times

sammo08 1 year, 3 months ago

Selected Answer: D

Option A Incorrect – AWS Lambda is not suitable for managing and patching third-party software on EC2 instances at scale.

Option B Incorrect – AWS Systems Manager Patch Manager is designed for OS patching, not third-party software patching

Option C Incorrect – A maintenance window schedules updates but does not immediately apply the patch, delaying remediation

Option D Correct – AWS Systems Manager Run Command allows immediate execution of a custom patching script across all EC2 instances, ensuring quick remediation.

upvoted 4 times

Dharmarajan 1 year, 4 months ago

Selected Answer: B

I would vote for B as the option for Run command does not have exception handling capabilities like patch failures. OR the command has to be scripted to handle such cases, which is additional operational overhead - especially for cases where the exception corner cases are not handled in the script.

upvoted 1 times

FlyingHawk 1 year, 4 months ago

Selected Answer: D

Patch Manager is optimized for OS-level security patches and may not support the specific third-party software patch required in this scenario.

Even if the third-party software is supported, Patch Manager may not provide the level of control or speed needed for a critical security update.

upvoted 1 times

Rcosmos 1 year, 5 months ago

Selected Answer: D

A Opção D (AWS Systems Manager Run Command) é a solução mais rápida e eficaz para aplicar o patch a todas as 1.000 instâncias EC2 em resposta a uma vulnerabilidade crítica de segurança, garantindo uma correção imediata e controlada.

upvoted 1 times

Rohan_Butala 1 year, 5 months ago

Selected Answer: B

AWS Systems Manager Patch Manager is a fully managed service designed to automate the process of patching operating systems and applications on EC2 instances. It can automatically apply patches to EC2 instances at scale, ensuring that critical vulnerabilities are remediated quickly. Patch Manager allows you to define patch baselines, schedule patching, and apply patches across multiple instances without manual intervention.

upvoted 2 times

Rohan_Butala 1 year, 5 months ago

Selected Answer: B

Option B has right balance between cost-effectiveness and performance, ensuring that frequently accessed data is easily retrievable and older data is archived at lower cost with fast retrieval options for both new and older data.

upvoted 2 times

Homine 1 year, 7 months ago

Selected Answer: B

Path Manager can support patch third party applications

upvoted 2 times

PaulGa 1 year, 9 months ago

Selected Answer: B

Ans B (should be). Clue: "The company needs to *patch* the third-party software..." - if the Patch Manager is not capable of this then perhaps it should be, rather than delegating functionality to another service... The suggested answer D, implies Patch Manager can't do the job...

upvoted 3 times

jhoiti 1 year, 11 months ago

Selected Answer: B

In your case, since the company needs to fix a critical security vulnerability as quickly as possible, Patch Manager (option B) would be the most recommended choice. It provides a quick and efficient way to apply the patch to all affected instances without the need to create custom scripts.

However, if your company has specific requirements that cannot be met by Patch Manager, Run Command (option D) may be a viable alternative.

upvoted 2 times

lofzee 2 years ago

I originally thought B but after a bit of reading I've changed my mind to D purely because patch manager will not be aware of this random third part application.

Run command allows you to install applications, run powershell, bash etc commands at scale so the most sensible answer would be run command.

upvoted 6 times

EMPERBACH 2 years, 1 month ago

Selected Answer: B

Install software -> Patch Manager

Run command/processing workload -> Run Command

upvoted 3 times

ManikRoy 2 years, 1 month ago

Selected Answer: D

I think patch manager would need an agent to be installed and also Patch Manager doesn't derive severity levels from third-party sources.

upvoted 3 times

RafikTAAMMA 2 years, 2 months ago

Selected Answer: D

AWS Systems Manager Patch Manager primarily focuses on operating system patches and does not directly support third-party software patching or Linux instances

upvoted 4 times

awsgeek75 2 years, 4 months ago

Selected Answer: D

Critical means immediate. Just run the patch command with AWS SM run command to get it done. D is best choice.

A: Too convoluted

B: Can work but have to setup a lot of things to get this done. would be a good choice if D wasn't an option

C: It's a critical patch so not time for maintenance window

upvoted 5 times

A company is developing an application that provides order shipping statistics for retrieval by a REST API. The company wants to extract the shipping statistics, organize the data into an easy-to-read HTML format, and send the report to several email addresses at the same time every morning.

Which combination of steps should a solutions architect take to meet these requirements? (Choose two.)

A. Configure the application to send the data to Amazon Kinesis Data Firehose.

B. Use Amazon Simple Email Service (Amazon SES) to format the data and to send the report by email. **Most Voted**


C. Create an Amazon EventBridge (Amazon CloudWatch Events) scheduled event that invokes an AWS Glue job to query the application's API for the data.

D. Create an Amazon EventBridge (Amazon CloudWatch Events) scheduled event that invokes an AWS Lambda function to query the application's API for the data. **Most Voted**

E. Store the application data in Amazon S3. Create an Amazon Simple Notification Service (Amazon SNS) topic as an S3 event destination to send the report by email.

Correct Answer: BD

Community vote distribution



Comments

whosawsome **Highly Voted** 3 years, 7 months ago

Selected Answer: BD

You can use SES to format the report in HTML.

<https://docs.aws.amazon.com/ses/latest/dg/send-email-formatted.html>

upvoted 33 times

apchandana 3 years, 2 months ago

this document is talking about the SES API. not ses. SES does not format data. just sending emails.

<https://aws.amazon.com/ses/>

upvoted 7 times

Clouddon 2 years, 10 months ago

When you send an email with Amazon SES, the email information you need to provide depends on how you call Amazon SES. You can provide a minimal amount of information and have Amazon SES take care of all of the formatting for you. Or, if you want to do something more advanced like send an attachment, you can provide the raw message yourself. <https://docs.aws.amazon.com/ses/latest/dg/send-email-concepts-email-format.html>

upvoted 4 times

backbencher2022 **Highly Voted** 3 years, 7 months ago

Selected Answer: BD

B&D are the only 2 correct options. If you are choosing option E then you missed the daily morning schedule requirement mentioned in the question which cant be achieved with S3 events for SNS. Event Bridge can used to configure scheduled events (every morning in this case). Option B fulfills the email in HTML format requirement (by SES) and D fulfills every morning schedule event requirement (by EventBridge)

upvoted 27 times

PaulGa 1 year, 9 months ago

Option B doesn't specify an API interface...

upvoted 1 times

RupeC 2 years, 11 months ago

I don't believe you are correct when you say that E cannot meet the scheduling requirement. If the glue action is scheduled and outputs to S3, then as the S3 event destination is SNS, in effect you have a way of getting SNS to have a scheduled release.

upvoted 2 times

pentium75 2 years, 5 months ago

But E does not include a glue action. We need either C or D for the scheduling, plus B or E for the email sending.
upvoted 2 times

slimen 2 years, 7 months ago

the daily schedule can be achieve with event bridge
- schedule and event bridge to trigger daily
- the event bridge will trigger a lambda function that will collect data and save it in s3
- once data in s3 the event notification will trigger SNS to send emails
upvoted 2 times

RamanadhRavinuthala Most Recent 1 year, 10 months ago

Selected Answer: BD

D Explanation: EventBridge can be used to schedule regular invocations of a Lambda function that retrieves the required data from the application's API. This step sets up the process to collect the data at the specified time every morning.

B Explanation: Amazon SES can format the data into an easy-to-read HTML report and send the email to multiple recipients efficiently.
upvoted 4 times

JA2018 1 year, 7 months ago

<https://docs.aws.amazon.com/ses/latest/dg/send-email-concepts-email-format.html>
upvoted 1 times

hieunt.hus 1 year, 11 months ago

Selected Answer: BC

SNS & Glue
upvoted 5 times

jatric 1 year, 11 months ago

BC - Question says extract the data - which is the job of AWS Glue and then invoke SES endpoint to send email. SES support html format.

But BD also seems possible.
upvoted 5 times

lofzee 2 years ago

Defo B and D
Look at the aws notes on Sns
Notes
You can't customize the body of the email message. The email delivery feature is intended to provide internal system alerts, not marketing messages
It can't send anything but notifications. SES can send normal emails therefore, BD is the answer.
upvoted 3 times

jhoiti 2 years, 2 months ago

Selected Answer: BE

Option B suggests using Amazon SES, which allows you to format the data and email the report to multiple recipients in an efficient and scalable way.

Option E proposes storing application data in Amazon S3, which is scalable and durable storage. By configuring an Amazon SNS topic as an S3 event destination, you can automatically trigger the report to be sent via email whenever new data is added to S3.
upvoted 1 times

Monster07 2 years, 2 months ago

Selected Answer: BD

Chat GPT says :
With Amazon SES, you can send rich, formatted email content, including text, HTML, attachments, and embedded images, suitable for email communication.
Amazon SNS is primarily used for sending plain-text or JSON-formatted messages, suitable for notifications and alerts across different channels.

This can suggest that we need to use SES if we want to use HTML content.
upvoted 1 times

wyejay 2 years, 4 months ago

Answer: B and D
Other options
A. Amazon Kinesis Data Firehose: This service is typically used for real-time streaming data processing rather than for scheduled tasks like generating a morning report.

C. Amazon EventBridge to invoke an AWS Glue job: AWS Glue is a data integration service that's more focused on ETL (extract, transform, load) operations, often involving large datasets and complex transformations, which might be more than needed for this scenario.

E. Amazon S3 with SNS topic: Storing data in S3 and using SNS for notification is viable, but this doesn't directly address the need to format the data into HTML and send it as an email report. SNS is better suited for sending notifications rather than formatted reports.

upvoted 11 times

tonybuivannghia 1 year, 8 months ago

Some mistakes that Kinesis Data Firehose is used for near real-time streaming data processing. But totally I agree with you that B & D are correct.

upvoted 2 times

awsgeek75 2 years, 4 months ago

Selected Answer: BD

Very detailed question so let's break it down:

"send the report to several email addresses at the same time every morning" this locks B as nothing else can do it.

A: Firehose to collect data from API will work but it cannot generate a report

C: Glue is ETL, it cannot extract data from an API

E: Store data in S3. No idea what this will help with

The API provides order shipping data so you can query it. Lambda can be used to query the API easily so D is good choice that works with B.

BD is correct combination

upvoted 3 times

pentium75 2 years, 5 months ago

Selected Answer: BD

"At the same time every morning" requires scheduling, which is only mentioned in C and D. AWS Glue has no native functionality to query REST APIs, thus we need a Lambda function -> D.

For email we need SES or SNS, but as we want "an easy-to-read HTML format", SNS is out. SNS can send notifications, not formatted emails. Thus B.

upvoted 5 times

wearrexdzw3123 2 years, 7 months ago

Selected Answer: B

I think there is a problem with the answer. It should be that ses sends the email processed by lambda.

upvoted 1 times

tom_cruise 2 years, 7 months ago

Selected Answer: BD

Key: retrieval by a REST API, that's why use lambda

upvoted 2 times

tom_cruise 2 years, 8 months ago

Selected Answer: DE

Both SES and SNS can format html, but there is a disconnection between B and D. Where do you store the data between the steps?

upvoted 4 times

pentium75 2 years, 5 months ago

Why would I need to "store the data"? Wouldn't the Lambda function just call the SES API?

upvoted 2 times

David_Ang 2 years, 8 months ago

Selected Answer: BD

the reason why "B" is more correct than "E" is because is more simple and you don't have to store data is not what they want, also SES is a service that is meant for sending the data through email, and is exactly what the company wants. is not the first time the admin is wrong with the answer

upvoted 3 times

hieulam 2 years, 8 months ago

Selected Answer: DE

E should be correct:

<https://saturncloud.io/blog/how-to-send-html-mails-using-amazon-sns/>

upvoted 1 times

h_sahu 2 years, 8 months ago

I believe BD are the answers. E can't be used, because, in E can't help with email formatting. E won't be the best choice even for scheduling.

upvoted 3 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: BD

Create an Amazon EventBridge (Amazon CloudWatch Events) scheduled event that invokes an AWS Lambda function to query the application's API for the data. Then use Amazon Simple Email Service (Amazon SES) to format the data and to send the report by email.

upvoted 2 times

A company wants to migrate its on-premises application to AWS. The application produces output files that vary in size from tens of gigabytes to hundreds of terabytes. The application data must be stored in a standard file system structure. The company wants a solution that scales automatically, is highly available, and requires minimum operational overhead.

Which solution will meet these requirements?

- A. Migrate the application to run as containers on Amazon Elastic Container Service (Amazon ECS). Use Amazon S3 for storage.
- B. Migrate the application to run as containers on Amazon Elastic Kubernetes Service (Amazon EKS). Use Amazon Elastic Block Store (Amazon EBS) for storage.
- C. Migrate the application to Amazon EC2 instances in a Multi-AZ Auto Scaling group. Use Amazon Elastic File System (Amazon EFS) for storage. **Most Voted**
- D. Migrate the application to Amazon EC2 instances in a Multi-AZ Auto Scaling group. Use Amazon Elastic Block Store (Amazon EBS) for storage.

Correct Answer: C

Community vote distribution

C (100%)

Comments

ArielSchivo **Highly Voted** 3 years, 7 months ago

Selected Answer: C

EFS is a standard file system, it scales automatically and is highly available.

upvoted 33 times

masetromain **Highly Voted** 3 years, 8 months ago

I have absolutely no idea...

Output files that vary in size from tens of gigabytes to hundreds of terabytes

Limit size for a single object:

S3 5TiB

<https://aws.amazon.com/fr/blogs/aws/amazon-s3-object-size-limit/>

EBS 64 TiB

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/volume_constraints.html

EFS 47.9 TiB

<https://docs.aws.amazon.com/efs/latest/ug/limits.html>

upvoted 11 times

Help2023 3 years, 3 months ago

The answer to that is

Limit size for a single object:

S3, 5TiB is per object but you can have more than one object in a bucket, meaning infinity

<https://aws.amazon.com/fr/blogs/aws/amazon-s3-object-size-limit/>

EBS 64 TiB is per block of storage

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/volume_constraints.html

EFS 47.9 TiB per file and in the questions it says Files the 's'

<https://docs.aws.amazon.com/efs/latest/ug/limits.html>

upvoted 2 times

JayBee65 3 years, 6 months ago

S3 and EBS are block storage but you are looking to store files, so EFS is the correct option.

upvoted 4 times

Ello2023 3 years, 4 months ago

S3 is object storage.