

By creating an Amazon CloudFront distribution with origins in each AWS Region where the application is deployed, you can leverage CloudFront's global edge network to route traffic to the closest available Region. CloudFront will automatically route the traffic based on the client's location and the health of the origins using CloudFront health checks.

Option A (creating Amazon Route 53 health checks with an active-active failover configuration) is not suitable for this scenario as it is primarily used for failover between different endpoints within the same Region, rather than routing traffic to different Regions.

upvoted 2 times

pentium75 2 years, 5 months ago

Option A does not speak of Route 53 failover routing policies.

upvoted 1 times

Axeashes 2 years, 12 months ago

Selected Answer: A

<https://aws.amazon.com/blogs/compute/building-a-multi-region-serverless-application-with-amazon-api-gateway-and-aws-lambda/>

upvoted 4 times

Gooniegoogoo 2 years, 11 months ago

that is from 2017.. i wonder if it is still relevant..

upvoted 1 times

DrWatson 3 years ago

Selected Answer: A

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/dns-failover.html>

upvoted 2 times

antropaws 3 years ago

Selected Answer: A

I understand that you can use Route 53 to provide regional failover.

upvoted 1 times

alexandercamachop 3 years ago

Selected Answer: A

To route traffic to multiple AWS Regions and provide regional failover capabilities for a stateless web application running on AWS Lambda functions invoked by Amazon API Gateway, you can use Amazon Route 53 with an active-active failover configuration.

By creating Amazon Route 53 health checks for each Region and configuring an active-active failover configuration, Route 53 can monitor the health of the endpoints in each Region and route traffic to healthy endpoints. In the event of a failure in one Region, Route 53 automatically routes traffic to the healthy endpoints in other Regions.

This setup ensures high availability and failover capabilities for your web application across multiple AWS Regions.

upvoted 4 times


A company has two VPCs named Management and Production. The Management VPC uses VPNs through a customer gateway to connect to a single device in the data center. The Production VPC uses a virtual private gateway with two attached AWS Direct Connect connections. The Management and Production VPCs both use a single VPC peering connection to allow communication between the applications.

What should a solutions architect do to mitigate any single point of failure in this architecture?

- A. Add a set of VPNs between the Management and Production VPCs.
- B. Add a second virtual private gateway and attach it to the Management VPC.
- C. Add a second set of VPNs to the Management VPC from a second customer gateway device. **Most Voted**
- D. Add a second VPC peering connection between the Management VPC and the Production VPC.

Correct Answer: C

Community vote distribution



Comments

Guru4Cloud **Highly Voted** 1 year, 9 months ago

Selected Answer: C

C is the correct option to mitigate the single point of failure.

The Management VPC currently has a single VPN connection through one customer gateway device. This is a single point of failure.

Adding a second set of VPN connections from the Management VPC to a second customer gateway device provides redundancy and eliminates this single point of failure.

upvoted 7 times

Guru4Cloud 1 year, 9 months ago

As @Abrar2022 explains

(production) VPN 1-----> cgw 1
(management) VPN 2-----> cgw

upvoted 4 times

Abrar2022 **Highly Voted** 2 years ago

(production) VPN 1-----> cgw 1
(management) VPN 2-----> cgw 2

upvoted 6 times

Abrar2022 2 years ago

ANSWER IS C

upvoted 2 times

Dantecito **Most Recent** 1 year, 4 months ago

Selected Answer: D

but what if the VPC peering connection fails? how The Management and Production vpc communicate?

upvoted 1 times

bsbs1234 1 year, 8 months ago

C,

(production) --PrivateGateway----->Direct Connect Gateway 1 ---> cgw 1 ---> DataCenter

(production) -- PrivateGateway -----> Direct Connect Gateway 2 --->cgw 2 --> DataCenter

(Management) --> VPN ----> (Direct Connect Gateway 1?) --->cgw1 ----> dataCenter----> device in dataCenter

upvoted 2 times

omoakin 2 years ago

I agree to C

upvoted 2 times

cloudenthusiast 2 years ago

Selected Answer: C

option D is not a valid solution for mitigating single points of failure in the architecture. I apologize for the confusion caused by the incorrect information.

To mitigate single points of failure in the architecture, you can consider implementing option C: adding a second set of VPNs to the Management VPC from a second customer gateway device. This will introduce redundancy at the VPN connection level for the Management VPC, ensuring that if one customer gateway or VPN connection fails, the other connection can still provide connectivity to the data center.

upvoted 5 times

Efren 2 years ago

Selected Answer: C

Redundant VPN connections: Instead of relying on a single device in the data center, the Management VPC should have redundant VPN connections established through multiple customer gateways. This will ensure high availability and fault tolerance in case one of the VPN connections or customer gateways fails.

upvoted 5 times

nosense 2 years ago

Selected Answer: C

<https://www.examttopics.com/discussions/amazon/view/53908-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 3 times

A company runs its application on an Oracle database. The company plans to quickly migrate to AWS because of limited resources for the database, backup administration, and data center maintenance. The application uses third-party database features that require privileged access.

Which solution will help the company migrate the database to AWS MOST cost-effectively?

- A. Migrate the database to Amazon RDS for Oracle. Replace third-party features with cloud services.
- B. Migrate the database to Amazon RDS Custom for Oracle. Customize the database settings to support third-party features.
Most Voted
- C. Migrate the database to an Amazon EC2 Amazon Machine Image (AMI) for Oracle. Customize the database settings to support third-party features.
- D. Migrate the database to Amazon RDS for PostgreSQL by rewriting the application code to remove dependency on Oracle APEX.

Correct Answer: B

Community vote distribution

B (97%)

3%

Comments

Abrar2022 **Highly Voted** 2 years, 6 months ago

Selected Answer: B

RDS Custom since it's related to 3rd vendor
RDS Custom since it's related to 3rd vendor
RDS Custom since it's related to 3rd vendor
upvoted 7 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: B

Key constraints: Limited resources for DB admin and cost. 3rd party db features with privileged access.
A: Won't work due to 3rd party features
C: AMI with Oracle may work but again overhead of backed, maintenance etc
D: Too much overhead in rewrite
B: Actually supports Oracle 3rd party features
Caution: If this is only about APEX as suggested in option D, then A is also a possible answer:
<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Appendix.Oracle.Options.APEX.html>
upvoted 3 times

awsgeek75 1 year, 11 months ago

Action ignore the last line of my previous comment, A is not a valid option in any case as it suggest replacing 3rd party features with cloud services which is not possible without more details.
upvoted 2 times

pentium75 1 year, 11 months ago

Selected Answer: B

"Amazon RDS Custom is a managed database service for applications that require customization of the underlying operating system and database environment. Benefits of RDS automation with the access needed for legacy, packaged, and custom applications."

That should allow the "privileged access".
upvoted 2 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: B

Migrate the database to Amazon RDS Custom for Oracle. Customize the database settings to support third-party features.

upvoted 4 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: B

Custom database features = Amazon RDS Custom for Oracle

upvoted 4 times

antropaws 2 years, 6 months ago

Selected Answer: B

Most likely B.

upvoted 2 times

omoakin 2 years, 6 months ago

CCCCCCCCCCCCCCCCCCCC

upvoted 1 times

aqmdla2002 2 years, 6 months ago

Selected Answer: B

<https://aws.amazon.com/about-aws/whats-new/2021/10/amazon-rds-custom-oracle/>

upvoted 3 times

hiroohiroo 2 years, 6 months ago

Selected Answer: B

https://docs.aws.amazon.com/ja_jp/AmazonRDS/latest/UserGuide/Oracle.Resources.html

upvoted 2 times

karbob 2 years, 6 months ago

Amazon RDS Custom for Oracle, which is not an actual service. !!!!

upvoted 2 times

nosense 2 years, 6 months ago

Option C is also a valid solution, but it is not as cost-effective as option B.

Option C requires the company to manage its own database infrastructure, which can be expensive and time-consuming. Additionally, the company will need to purchase and maintain Oracle licenses.

upvoted 3 times

y0 2 years, 6 months ago

RDS Custom enables the capability to access the underlying database and OS so as to configure additional settings to support 3rd party. This feature is applicable only for Oracle and Postgresql

upvoted 2 times

y0 2 years, 6 months ago

Sorry Oracle and sql server (not postgresql)

upvoted 2 times

omoakin 2 years, 6 months ago

I will say C cos of this

"application uses third-party "

upvoted 1 times

cloudenthusiast 2 years, 6 months ago

Selected Answer: C

Should not it be since for Ec2, the company will have full control over the database and this is the reason that they are moving to AWS in the first place "The company plans to quickly migrate to AWS because of limited resources for the database, backup administration, and data center maintenance?"

upvoted 1 times

pentium75 1 year, 11 months ago

"Amazon RDS Custom (B) is a managed database service for applications that require customization of the underlying operating system and database environment. Benefits of RDS automation with the access needed for legacy, packaged, and custom applications."

upvoted 2 times

Efren 2 years, 6 months ago

Selected Answer: B

RDS Custom when is something related to 3rd vendor, for me

upvoted 2 times

nosense 2 years, 6 months ago

not sure, but b probably

upvoted 3 times


A company has a three-tier web application that is in a single server. The company wants to migrate the application to the AWS Cloud. The company also wants the application to align with the AWS Well-Architected Framework and to be consistent with AWS recommended best practices for security, scalability, and resiliency.

Which combination of solutions will meet these requirements? (Choose three.)

- A. Create a VPC across two Availability Zones with the application's existing architecture. Host the application with existing architecture on an Amazon EC2 instance in a private subnet in each Availability Zone with EC2 Auto Scaling groups. Secure the EC2 instance with security groups and network access control lists (network ACLs).
- B. Set up security groups and network access control lists (network ACLs) to control access to the database layer. Set up a single Amazon RDS database in a private subnet.
- C. Create a VPC across two Availability Zones. Refactor the application to host the web tier, application tier, and database tier. Host each tier on its own private subnet with Auto Scaling groups for the web tier and application tier. **Most Voted**
- D. Use a single Amazon RDS database. Allow database access only from the application tier security group.
- E. Use Elastic Load Balancers in front of the web tier. Control access by using security groups containing references to each layer's security groups. **Most Voted**
- F. Use an Amazon RDS database Multi-AZ cluster deployment in private subnets. Allow database access only from application tier security groups. **Most Voted**

Correct Answer: CEF

Community vote distribution



Comments

awsgeek75 **Highly Voted** 1 year, 11 months ago

Selected Answer: CEF

The wording on this question makes things ambiguous for C. But, remember well-architected so:

A: Not ideal as it is suggesting using existing architecture but with autoscaling EC2. Doesn't leave room for improvement on scaling or reliability on each tier.

B: Single RDS, not well-architected

D: Again, single RDS

E,F are good options and C is only remaining good one.

upvoted 9 times

awsgeek75 1 year, 11 months ago

C is badly worded IMHO because of this part " Refactor the application to host the web tier, application tier, and database tier." The database tier just makes it confusing when you don't read E and F.

upvoted 3 times

jjcode **Highly Voted** 1 year, 10 months ago

i would flag this on the test and do it last.

upvoted 7 times

4919c76 **Most Recent** 1 year, 1 month ago

Selected Answer: EF

The web tier can only be in a public subnet so C is wrong

upvoted 1 times

ChhatwaniB 1 year, 1 month ago

in Option C , it saying all layer should be in it's own private subnet . Isn't the web layer should be in public subnet ?

upvoted 1 times

ChhatwaniB 1 year, 1 month ago

in Option C , it saying all layer should be in it's own private subnet . Isn't the web layer should be in public subnet ?

upvoted 1 times

Burrito69 1 year, 8 months ago

remove singles and remove network ACLs

upvoted 3 times

argl1995 2 years, 5 months ago

option A cannot be the answer as Security group is at instance level whereas a NACL is at the subnet level. Having said that option C is the right one as the VPC cannot span across the regions and here it is mentioned two AZs for which I am guessing it is a default VPC which is created in each region with a subnet in each AZ.

upvoted 2 times

argl1995 2 years, 5 months ago

So, CEF is the right answer

upvoted 2 times

Gooniegoogoo 2 years, 5 months ago

How can you create a VPC across 2 AZ? i only see EF here.. if they mean 2 separate VPC then that is different but a VPC cannot span two AZ..

upvoted 1 times

lemur88 2 years, 3 months ago

A VPC most definitely can span across 2 AZ. You may be thinking of subnets.

upvoted 3 times

marufexplorer 2 years, 5 months ago

I also agree with CEF but chatGPT answer is ACE. A and C is the similar
Another Logic F is not True because in the question not mentioned about DB

upvoted 1 times

awsgeek75 1 year, 10 months ago

ChatGPT is a language parser. It is not an AWS solution architect!

upvoted 3 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: CEF

CEF is best

upvoted 2 times

antropaws 2 years, 6 months ago

Selected Answer: CEF

It's clearly CEF.

upvoted 2 times

Abrar2022 2 years, 6 months ago

Selected Answer: CEF

C-scalable and resilient

E-high availability of the application

F-Multi-AZ configuration provides high availability

upvoted 6 times

omoakin 2 years, 6 months ago

B- to control access to database

C-scalable and resilient

E-high availability of the application

upvoted 1 times

lucdt4 2 years, 6 months ago

Selected Answer: CEF

CEF

A: application's existing architecture is wrong (single AZ)

B: single AZ

D: Single AZ

upvoted 2 times

cloudenthusiast 2 years, 6 months ago

C.

This solution follows the recommended architecture pattern of separating the web, application, and database tiers into different subnets. It provides better security, scalability, and fault tolerance.

E.By using Elastic Load Balancers (ELBs), you can distribute traffic to multiple instances of the web tier, increasing scalability and availability.

Controlling access through security groups allows for fine-grained control and ensures only authorized traffic reaches each layer.

F.

Deploying an Amazon RDS database in a Multi-AZ configuration provides high availability and automatic failover. Placing the database in private subnets enhances security. Allowing database access only from the application tier security groups limits exposure and follows the principle of least privilege.

upvoted 5 times

mwwt2022 1 year, 11 months ago

good explanation

upvoted 2 times

nosense 2 years, 6 months ago

Selected Answer: CEF

Only this valid for best practices and well architected

upvoted 5 times


A company is migrating its applications and databases to the AWS Cloud. The company will use Amazon Elastic Container Service (Amazon ECS), AWS Direct Connect, and Amazon RDS.

Which activities will be managed by the company's operational team? (Choose three.)

- A. Management of the Amazon RDS infrastructure layer, operating system, and platforms
- B. Creation of an Amazon RDS DB instance and configuring the scheduled maintenance window **Most Voted**
- C. Configuration of additional software components on Amazon ECS for monitoring, patch management, log management, and host intrusion detection **Most Voted**
- D. Installation of patches for all minor and major database versions for Amazon RDS
- E. Ensure the physical security of the Amazon RDS infrastructure in the data center
- F. Encryption of the data that moves in transit through Direct Connect **Most Voted**

Correct Answer: BCF

Community vote distribution



Comments

pentium75 **Highly Voted** 1 year, 11 months ago

Selected Answer: BCF

ADE = AWS responsibility

upvoted 10 times

awsgeek75 **Highly Voted** 1 year, 10 months ago

Selected Answer: BCF

Just to clarify on F. Direct Connect is an ISP and AWS offering, I consider it as a physical connection just like you get from your ISP at home. There is not security on it until you build security on the connection. AWS provides Direct Connect but it does not provide encryption level security on data movement through it by default. It's the customer's responsibility.

upvoted 9 times

Guru4Cloud **Most Recent** 2 years, 3 months ago

Selected Answer: BCF

B: Creating an RDS instance and configuring the maintenance window is done by the customer.

C: Adding monitoring, logging, etc on ECS is managed by the customer.

F: Encrypting Direct Connect traffic is handled by the customer.

upvoted 5 times

james2033 2 years, 4 months ago

Selected Answer: BCF

In question has 3 keyword "Amazon ECS", "AWS Direct Connect", "Amazon RDS". With per Amazon services, choose 1 according answer. Has 6 items, need pick 3 items.

ECS --> choose C.

Direct Connect --> choose F.

RDS --> Exclude A (by keyword "infrastructure layer"), Choose B. Exclusive D (by keyword "patches for all minor and major database versions for Amazon RDS"). Exclusive E (by keyword "Ensure the physical security of the Amazon RDS"). Easy question.

upvoted 5 times

kapit 2 years, 5 months ago

BC & F (no automatic encryption with direct connect

upvoted 2 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: BF

Amazon ECS is a fully managed service, the ops team only focus on building their applications, not the environment.

Only option B and F makes sense.

upvoted 1 times

pentium75 1 year, 11 months ago

Plus C (we were asked for three). Configuration (!) of components for monitoring, log management etc.; those services exist from AWS but you need to configure them (which logs do you want to store for how long etc.).

upvoted 2 times

antropaws 2 years, 6 months ago

Selected Answer: BCF

100% BCF.

upvoted 2 times

lucdt4 2 years, 6 months ago

Selected Answer: BCF

BCF

B: Mentioned RDS

C: Mentioned ECS

F: Mentioned Direct connect

upvoted 5 times

hiroohiroo 2 years, 6 months ago

Selected Answer: BCF

Yes BCF

upvoted 2 times

omoakin 2 years, 6 months ago

I agree BCF

upvoted 2 times

nosense 2 years, 6 months ago

Selected Answer: BCF

Bcf for me

upvoted 3 times

A company runs a Java-based job on an Amazon EC2 instance. The job runs every hour and takes 10 seconds to run. The job runs on a scheduled interval and consumes 1 GB of memory. The CPU utilization of the instance is low except for short surges during which the job uses the maximum CPU available. The company wants to optimize the costs to run the job.

Which solution will meet these requirements?

- A. Use AWS App2Container (A2C) to containerize the job. Run the job as an Amazon Elastic Container Service (Amazon ECS) task on AWS Fargate with 0.5 virtual CPU (vCPU) and 1 GB of memory.
- B. Copy the code into an AWS Lambda function that has 1 GB of memory. Create an Amazon EventBridge scheduled rule to run the code each hour. **Most Voted**
- C. Use AWS App2Container (A2C) to containerize the job. Install the container in the existing Amazon Machine Image (AMI). Ensure that the schedule stops the container when the task finishes.
- D. Configure the existing schedule to stop the EC2 instance at the completion of the job and restart the EC2 instance when the next job starts.

Correct Answer: B

Community vote distribution

B (100%)

Comments

omarshaban **Highly Voted** 2 years, 4 months ago
THIS WAS IN MY EXAM
upvoted 12 times

awsgeek75 **Highly Voted** 2 years, 4 months ago
Selected Answer: B
Never done it myself but apparently you can run Java in Lambda all the way to latest version
<https://docs.aws.amazon.com/lambda/latest/dg/lambda-java.html>
upvoted 6 times

Danilus **Most Recent** 1 year, 7 months ago
Selected Answer: B
key -The company wants to optimize the costs to run the job.
the answer is B because lambda provides the most cost-effective,easy to manage and scalable solution.it minimizes the need to manage infrastructure because you only pay for the compute time used during job execution
upvoted 2 times

noircesar25 2 years, 3 months ago
can someone explain what makes A wrong, im aware that C hasnt covered all the requirements but A seems good with fargate serverless and autoscaling functionalities, plus AWS app2container is for .NET and JAVA
upvoted 4 times

bujuman 2 years, 2 months ago
Statement: A company runs a Java-based job on an Amazon EC2 instance
Requirement: The company wants to optimize the costs to run the job
Regarding option A: App2Container is more likely for migrating legacy application to container based application.
Which is not the main purpose of this use case.
We are asked to reduce cost on a application that is already running under EC2 instance.
So option B has the high weight, cause lambda could perfectly do the job with the minimal cost
upvoted 4 times

JA2018 1 year, 6 months ago
you only need to pay for the compute resources consumed as and when each Lambda function is actually running (for the running duration)

upvoted 1 times

Murtadhaceit 2 years, 6 months ago

Selected Answer: B

This question is intended for Lambda. Just searched for Lambda with Event bridge. I

upvoted 3 times

potomac 2 years, 7 months ago

Selected Answer: B

Lambda allows you to allocate memory for your functions in increments of 1 MB, ranging from a minimum of 128 MB to a maximum of 10,240 MB (10 GB).

upvoted 2 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: B

Remember - AWS Lambda function can go up to 10 GB of memory, instead of free tier only allow 512MB.

upvoted 6 times

james2033 2 years, 10 months ago

Selected Answer: B

"AWS Batch jobs as EventBridge targets" at <https://docs.aws.amazon.com/batch/latest/userguide/batch-cwe-target.html>

AWS Batch + Amazon EventBridge <https://docs.aws.amazon.com/batch/latest/userguide/batch-cwe-target.html> .

AWS Lambda just for a point of time per period. Choose B.

upvoted 2 times

TariqKipkemei 2 years, 12 months ago

Selected Answer: B

10 seconds to run, optimize the costs, consumes 1 GB of memory = AWS Lambda function.

upvoted 3 times

alexandercamachop 3 years ago

Selected Answer: B

AWS Lambda automatically scales resources to handle the workload, so you don't have to worry about managing the underlying infrastructure. It provisions the necessary compute resources based on the configured memory size (1 GB in this case) and executes the job in a serverless environment.

By using Amazon EventBridge, you can create a scheduled rule to trigger the Lambda function every hour, ensuring that the job runs on the desired interval.

upvoted 2 times

Yadav_Sanjay 3 years ago

Selected Answer: B

B - Within 10 sec and 1 GB Memory (Lambda Memory 128MB to 10GB)

upvoted 3 times

Yadav_Sanjay 3 years ago

<https://docs.aws.amazon.com/lambda/latest/operatorguide/computing-power.html>

upvoted 2 times

Efren 3 years ago

Selected Answer: B

Agreed, B Lambda

upvoted 3 times

A company wants to implement a backup strategy for Amazon EC2 data and multiple Amazon S3 buckets. Because of regulatory requirements, the company must retain backup files for a specific time period. The company must not alter the files for the duration of the retention period.

Which solution will meet these requirements?

- A. Use AWS Backup to create a backup vault that has a vault lock in governance mode. Create the required backup plan.
- B. Use Amazon Data Lifecycle Manager to create the required automated snapshot policy.
- C. Use Amazon S3 File Gateway to create the backup. Configure the appropriate S3 Lifecycle management.
- D. Use AWS Backup to create a backup vault that has a vault lock in compliance mode. Create the required backup plan.

Most Voted

Correct Answer: D

Community vote distribution

D (100%)

Comments

Efren **Highly Voted** 3 years ago

D, Governance is like the goverment, they can do things you cannot , like delete files or backups :D Compliance, nobody can!
upvoted 39 times

cmbt 2 years, 11 months ago

Finally I understood!
upvoted 4 times

joshnort 2 years, 11 months ago

Great analogy
upvoted 8 times

nosense **Highly Voted** 3 years ago

Selected Answer: D

D bcs in governance we can delete backup
upvoted 5 times

Danilus **Most Recent** 1 year, 7 months ago

Selected Answer: D

key-The company must not alter the files for the duration of the retention period
the answer is D because backup with vault lock in compliance mode make sure that the files when are locked they cannot be deleted or modified for the specified retention period it makes the files immutable
why is not A? because the governance mode... with this mode backups allows flexibility to alter or delete backups
upvoted 2 times

f2e2419 2 years, 5 months ago

Selected Answer: D

D. Use AWS Backup to create a backup vault that has a vault lock in compliance mode. Create the required backup plan
upvoted 3 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: D

D. Use AWS Backup to create a backup vault that has a vault lock in compliance mode. Create the required backup plan
upvoted 3 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: D

Use AWS Backup to create a backup vault that has a vault lock in compliance mode. Create the required backup plan
upvoted 3 times

ccat91 2 years, 10 months ago

Selected Answer: D

Compliance mode
upvoted 2 times

TariqKipkemei 2 years, 12 months ago

Selected Answer: D

Must not alter the files for the duration of the retention period = Compliance Mode
upvoted 2 times

antropaws 3 years ago

Selected Answer: D

D for sure.
upvoted 2 times

dydzah 3 years ago

Selected Answer: D

<https://docs.aws.amazon.com/aws-backup/latest/devguide/vault-lock.html>
upvoted 3 times

cloudenthusiast 3 years ago

Selected Answer: D

compliance mode
upvoted 4 times

A company has resources across multiple AWS Regions and accounts. A newly hired solutions architect discovers a previous employee did not provide details about the resources inventory. The solutions architect needs to build and map the relationship details of the various workloads across all accounts.

Which solution will meet these requirements in the MOST operationally efficient way?

- A. Use AWS Systems Manager Inventory to generate a map view from the detailed view report.
- B. Use AWS Step Functions to collect workload details. Build architecture diagrams of the workloads manually.
- C. Use Workload Discovery on AWS to generate architecture diagrams of the workloads. **Most Voted**
- D. Use AWS X-Ray to view the workload details. Build architecture diagrams with relationships.

Correct Answer: C

Community vote distribution

C (97%)

3%

Comments

Danilus 1 year, 7 months ago

Key - Build and map the relationship details of the various workloads.

Key - MOST operationally efficient way.

The answer is C because Workload Discovery is a tool that allows you to find automatically the workloads on the cloud, giving you a visual representation of the infrastructure and the relationship between the resources through different accounts.

upvoted 2 times

zinabu 2 years, 2 months ago

workload discovery=architecture diagram

upvoted 3 times

osmk 2 years, 4 months ago

<https://docs.aws.amazon.com/solutions/latest/workload-discovery-on-aws/solution-overview.html> Workload Discovery on AWS is a visualization tool that automatically generates architecture diagrams of your workload on AWS. You can use this solution to build, customize, and share detailed workload visualizations based on live data from AWS

upvoted 3 times

awsgeek75 2 years, 5 months ago

Selected Answer: C

A: Systems Manager Inventory -> Metadata

B: Not possible (correct me if I'm wrong)

D: X-Ray is for application debugging

C: Workload Discovery is purpose built tool for this type of usage

upvoted 4 times

NayeraB 2 years, 3 months ago

Even if B is possible, it has "manually" in it which we won't do because we're lazy in this question

upvoted 3 times

potomac 2 years, 7 months ago

Selected Answer: C

Workload Discovery on AWS (formerly called AWS Perspective) is a tool to visualize AWS Cloud workloads. Use Workload Discovery on AWS to build, customize, and share detailed architecture diagrams of your workloads based on live data from AWS.

upvoted 3 times

TariqKipkemei 2 years, 7 months ago

Selected Answer: C

use Workload Discovery on AWS

upvoted 3 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: C

Workload Discovery is purpose-built to automatically generate visual mappings of architectures across accounts and Regions. This makes it the most operationally efficient way to meet the requirements.

upvoted 4 times

MrAWSAssociate 2 years, 11 months ago

Selected Answer: C

Option A: AWS SSM offers "Software inventory": Collect software catalog and configuration for your instances.

Option C: Workload Discovery on AWS: is a tool for maintaining an inventory of the AWS resources across your accounts and various Regions and mapping relationships between them, and displaying them in a web UI.

upvoted 4 times

DrWatson 3 years ago

Selected Answer: A

<https://aws.amazon.com/blogs/mt/visualizing-resources-with-workload-discovery-on-aws/>

upvoted 1 times

pentium75 2 years, 5 months ago

That is C

upvoted 2 times

Abrar2022 3 years ago

Selected Answer: C

AWS Workload Discovery - create diagram, map and visualise AWS resources across AWS accounts and Regions

upvoted 3 times

Abrar2022 3 years ago

Workload Discovery on AWS can map AWS resources across AWS accounts and Regions and visualize them in a UI provided on the website.

upvoted 2 times

hiroohiroo 3 years ago

Selected Answer: C

https://aws.amazon.com/jp/builders-flash/202209/workload-discovery-on-aws/?awsf.filter-name=*all

upvoted 3 times

omoakin 3 years ago

Only C makes sense

upvoted 3 times

cloudenthusiast 3 years ago

Selected Answer: C

Workload Discovery on AWS is a service that helps visualize and understand the architecture of your workloads across multiple AWS accounts and Regions. It automatically discovers and maps the relationships between resources, providing an accurate representation of the architecture.

upvoted 3 times

Efren 3 years ago

Not sure here tbh

To efficiently build and map the relationship details of various workloads across multiple AWS Regions and accounts, you can use the AWS Systems Manager Inventory feature in combination with AWS Resource Groups. Here's a solution that can help you achieve this:

AWS Systems Manager Inventory:

upvoted 1 times

nosense 3 years ago

Selected Answer: C

only c mapping relationships

upvoted 2 times

A company uses AWS Organizations. The company wants to operate some of its AWS accounts with different budgets. The company wants to receive alerts and automatically prevent provisioning of additional resources on AWS accounts when the allocated budget threshold is met during a specific period.

Which combination of solutions will meet these requirements? (Choose three.)

A. Use AWS Budgets to create a budget. Set the budget amount under the Cost and Usage Reports section of the required AWS accounts.

B. Use AWS Budgets to create a budget. Set the budget amount under the Billing dashboards of the required AWS accounts.

Most Voted

C. Create an IAM user for AWS Budgets to run budget actions with the required permissions.

D. Create an IAM role for AWS Budgets to run budget actions with the required permissions.

Most Voted

E. Add an alert to notify the company when each account meets its budget threshold. Add a budget action that selects the IAM identity created with the appropriate config rule to prevent provisioning of additional resources.

F. Add an alert to notify the company when each account meets its budget threshold. Add a budget action that selects the IAM identity created with the appropriate service control policy (SCP) to prevent provisioning of additional resources.

Most Voted

Correct Answer: BDF

Community vote distribution

BDF (71%)

ADF (15%)

12%

Comments

vesen22 **Highly Voted** 3 years ago

Selected Answer: BDF

I don't see why adf has the most voted when almost everyone has chosen bdf, smh
https://acloudguru.com/videos/acg-fundamentals/how-to-set-up-an-aws-billing-and-budget-alert?utm_source=google&utm_medium=paid-search&utm_campaign=cloud-transformation&utm_term=ssi-global-acg-core-dsa&utm_content=free-trial&gclid=Cj0KCQjwmtGjBhDhARIsAEqfDEcDfXdlul2NxxgSMxKraclITZimW0tDBRpsJPpx8IS9T4NndKhbUqPlaAlzhEALw_wcB
 upvoted 14 times

omarshaban **Highly Voted** 2 years, 4 months ago

IN MY EXAM

upvoted 11 times

FlyingHawk **Most Recent** 1 year, 6 months ago

Selected Answer: B

AWS budget is under the Budget and Planning section in Billing and Cost Management cosole
 upvoted 1 times

FlyingHawk 1 year, 6 months ago

<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/view-billing-dashboard.html>

upvoted 1 times

TariqKipkemei 2 years, 7 months ago

Selected Answer: DF

Its 11/Nov/2023. Options D&F are definitely required.
 As for the budget, right from the aws console, the only place to set this up is:
 AWS Billing>Cost Management>Budgets.

upvoted 5 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: BDF

How to create a budget:

Billing console > budget > create budget!

upvoted 6 times

Chris22usa 2 years, 11 months ago

ACF:

Option B is incorrect because the budget amount should be set under the Cost and Usage Reports section, not the Billing dashboards.

upvoted 1 times

pentium75 2 years, 5 months ago

"Create an AWS Budget: Go to the AWS Billing Dashboard"

<https://awslabs.github.io/scale-out-computing-on-aws/workshops/TKO-Scale-Out-Computing/modules/071-budgets/>

upvoted 1 times

Abrar2022 3 years ago

Selected Answer: BDF

How to create a budget:

Billing console > budget > create budget!

upvoted 3 times

udo2020 3 years ago

It is BDF because there is actually a Billing Dashboard available.

upvoted 8 times

hiroohiroo 3 years ago

Selected Answer: BDF

https://docs.aws.amazon.com/ja_jp/awsaccountbilling/latest/aboutv2/view-billing-dashboard.html

upvoted 6 times

y0 3 years ago

BDF - Budgets can be set from the billing dashboard in AWS console

upvoted 3 times

cloudenthusiast 3 years ago

Selected Answer: ADF

Currently, AWS does not have a specific feature called "AWS Billing Dashboards."

upvoted 6 times

Efren 3 years ago

if im not wrong, those are correct

upvoted 3 times


A company runs applications on Amazon EC2 instances in one AWS Region. The company wants to back up the EC2 instances to a second Region. The company also wants to provision EC2 resources in the second Region and manage the EC2 instances centrally from one AWS account.

Which solution will meet these requirements MOST cost-effectively?

- A. Create a disaster recovery (DR) plan that has a similar number of EC2 instances in the second Region. Configure data replication.
- B. Create point-in-time Amazon Elastic Block Store (Amazon EBS) snapshots of the EC2 instances. Copy the snapshots to the second Region periodically.
- C. Create a backup plan by using AWS Backup. Configure cross-Region backup to the second Region for the EC2 instances.
Most Voted
- D. Deploy a similar number of EC2 instances in the second Region. Use AWS DataSync to transfer the data from the source Region to the second Region.

Correct Answer: C

Community vote distribution



Comments

cloudehustiaat **Highly Voted** 3 years ago

Selected Answer: C

Using AWS Backup, you can create backup plans that automate the backup process for your EC2 instances. By configuring cross-Region backup, you can ensure that backups are replicated to the second Region, providing a disaster recovery capability. This solution is cost-effective as it leverages AWS Backup's built-in features and eliminates the need for manual snapshot management or deploying and managing additional EC2 instances in the second Region.

upvoted 8 times

Guru4Cloud **Highly Voted** 2 years, 9 months ago

Selected Answer: C

C is the most cost-effective solution that meets all the requirements.

AWS Backup provides automated backups across Regions for EC2 instances. This handles the backup requirement.

AWS Backup is more cost-effective for cross-Region EC2 backups than using EBS snapshots manually or DataSync.

upvoted 5 times

wizcloudifa 2 years, 1 month ago

Well Datasync is to transfer data from on-prem to AWS rather than AWS R1 to R2, and it is used for a one-time migration rather than repetitive data migration activities like backups hence its out (on so many levels)

upvoted 2 times

elmyth **Most Recent** 1 year, 7 months ago

Selected Answer: B

We have implemented solution similar to B, it's really small and simple script, so we don't have much manual work here, and it's really the most cost-effective.

upvoted 3 times

Certibanksman 1 year, 9 months ago

Selected Answer: B

Option B (EBS snapshots with cross-Region copy) is the most cost-effective solution for backing up EC2 instances to a second Region while allowing for centralized management and easy recovery when needed.

upvoted 3 times

Daminaij 1 year, 9 months ago

The answer is c

upvoted 1 times

bogobob 2 years, 6 months ago

Selected Answer: D

How does AWS Backup address that "The company also wants to provision EC2 resources in the second Region"?

upvoted 3 times

pentium75 2 years, 5 months ago

How do A, B or D address that? They want to "provision EC2 resources", nobody says that this should be copies of the existing servers. And if it should be copies of the existing servers, wouldn't we need the same (not "a similar") number of servers? We have no idea how many applications or how many servers they have.

upvoted 3 times

NSA_Poker 2 years ago

With AWS Backup you can 'back up the EC2 instances to a second Region' by implementing cross-Region backup & 'provision EC2 resources in the second Region' by restoring the backup using the AWS Backup console.

<https://docs.aws.amazon.com/aws-backup/latest/devguide/restore-resource.html>

upvoted 1 times

TariqKipkemei 2 years, 12 months ago

Selected Answer: C

AWS backup

upvoted 2 times

omoakin 3 years ago

CCCCC

. Create a backup plan by using AWS Backup. Configure cross-Region backup to the second Region for the EC2 instances.

upvoted 1 times

Blingy 3 years ago

CCCCCCC

upvoted 2 times

Efren 3 years ago

C, i would say same, always AWS Backup

upvoted 1 times


A company that uses AWS is building an application to transfer data to a product manufacturer. The company has its own identity provider (IdP). The company wants the IdP to authenticate application users while the users use the application to transfer data. The company must use Applicability Statement 2 (AS2) protocol.

Which solution will meet these requirements?

- A. Use AWS DataSync to transfer the data. Create an AWS Lambda function for IdP authentication.
- B. Use Amazon AppFlow flows to transfer the data. Create an Amazon Elastic Container Service (Amazon ECS) task for IdP authentication.
- C. Use AWS Transfer Family to transfer the data. Create an AWS Lambda function for IdP authentication. **Most Voted**
- D. Use AWS Storage Gateway to transfer the data. Create an Amazon Cognito identity pool for IdP authentication.

Correct Answer: C

Community vote distribution



Comments

TariqKipkemei **Highly Voted** 2 years, 12 months ago

Selected Answer: C

Option C stands out stronger because AWS Transfer Family securely scales your recurring business-to-business file transfers to AWS Storage services using SFTP, FTPS, FTP, and AS2 protocols.

And AWS Lambda can be used to authenticate users with the company's IdP.

upvoted 11 times

baba365 2 years, 11 months ago

Ans : C

To authenticate your users, you can use your existing identity provider with AWS Transfer Family. You integrate your identity provider using an AWS Lambda function, which authenticates and authorizes your users for access to Amazon S3 or Amazon Elastic File System (Amazon EFS).

<https://docs.aws.amazon.com/transfer/latest/userguide/custom-identity-provider-users.html>

upvoted 6 times

Danilus **Most Recent** 1 year, 7 months ago

Selected Answer: C

key-The company must use Applicability Statement 2 (AS2) protocol

The only service in AWS that accepts this protocol is AWS Transfer Family, which provides a managed solution for secure file transfers using AS2, along with essential features such as encryption, digital signatures, and message disposition notifications (MDN)

upvoted 2 times

zinabu 2 years, 2 months ago

aws transfer family for data transfer and lamda function for idp authentication

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: C

<https://aws.amazon.com/about-aws/whats-new/2022/07/aws-transfer-family-support-applicability-statement-2-as2/>

upvoted 4 times

potomac 2 years, 7 months ago

Selected Answer: C

To authenticate your users, you can use your existing identity provider with AWS Transfer Family. You integrate your identity provider using an AWS Lambda function, which authenticates and authorizes your users for access to Amazon S3 or Amazon Elastic File System (Amazon EFS).

upvoted 3 times

potomac 2 years, 7 months ago

Selected Answer: C

Applicability Statement 2 (AS2) is a business-to-business (B2B) messaging protocol used to exchange Electronic Data Interchange (EDI) documents. With AWS Transfer Family's AS2 capabilities, you can securely exchange AS2 messages at scale while maintaining compliance and interoperability with your trading partners.

upvoted 4 times

thanhnv142 2 years, 7 months ago

D is ok

upvoted 1 times

hsinchang 2 years, 10 months ago

its own IdP -> Lambda

upvoted 3 times

dydzah 3 years ago

Selected Answer: C

<https://docs.aws.amazon.com/transfer/latest/userguide/custom-identity-provider-users.html>

upvoted 2 times

examtopictempacc 3 years ago

Selected Answer: C

C is correct. AWS Transfer Family supports the AS2 protocol, which is required by the company. Also, AWS Lambda can be used to authenticate users with the company's IdP, which meets the company's requirement.

upvoted 3 times

EA100 3 years ago

Answer - D

AS2 is a widely used protocol for secure and reliable data transfer. In this scenario, the company wants to transfer data using the AS2 protocol and authenticate application users using their own identity provider (IdP). AWS Storage Gateway provides a hybrid cloud storage solution that enables data transfer between on-premises environments and AWS.

By using AWS Storage Gateway, you can set up a gateway that supports the AS2 protocol for data transfer. Additionally, you can configure authentication using an Amazon Cognito identity pool. Amazon Cognito provides a comprehensive authentication and user management service that integrates with various identity providers, including your own IdP.

Therefore, Option D is the correct solution as it leverages AWS Storage Gateway for AS2 data transfer and allows authentication using an Amazon Cognito identity pool integrated with the company's IdP.

upvoted 2 times

deechean 2 years, 9 months ago

AWS Transfer Family also support AS2

upvoted 2 times

hiroohiroo 3 years ago

Selected Answer: C

<https://repost.aws/articles/ARo2ihKKThT2Cue5j6yVUGsQ/articles/ARo2ihKKThT2Cue5j6yVUGsQ/aws-transfer-family-announces-support-for-sending-as2-messages-over-https?>

upvoted 1 times

omoakin 3 years ago

C is correct

upvoted 1 times

omoakin 3 years ago

This is a new Qtn n AS2 is newly supported by AWS Transfer family.....good timing to know ur stuffs.

upvoted 2 times

nosense 3 years ago

Option D looks the better option because it is more secure, scalable, cost-effective, and easy to use than option C.

upvoted 1 times

cloudenthusiast 3 years ago

Selected Answer: D

AWS Storage Gateway supports the AS2 protocol for transferring data. By using AWS Storage Gateway, the company can integrate its own IdP authentication by creating an Amazon Cognito identity pool. Amazon Cognito provides user authentication and authorization capabilities, allowing the company to authenticate application users using its own IdP.

AWS Transfer Family does not currently support the AS2 protocol. AS2 is a specific protocol used for secure and reliable data transfer, often used in business-to-business (B2B) scenarios. In this case, option C, which suggests using AWS Transfer Family, would not meet the requirement of using the AS2 protocol.

upvoted 3 times

omoakin 3 years ago

AWS Transfer Family now supports the Applicability Statement 2 (AS2) protocol, complementing existing protocol support for SFTP, FTPS, and FTP

upvoted 2 times

y0 3 years ago

This is not a case for storage gateway which is more used for a hybrid like environment. Here, to transfer data, we can think of Datasync or Transfer family and considering AS2 protocol, transfer family looks good

upvoted 3 times

Efren 3 years ago

ChatGP

To meet the requirements of using an identity provider (IdP) for user authentication and the AS2 protocol for data transfer, you can implement the following solution:

AWS Transfer Family: Use AWS Transfer Family, specifically AWS Transfer for SFTP or FTPS, to handle the data transfer using the AS2 protocol. AWS Transfer for SFTP and FTPS provide fully managed, highly available SFTP and FTPS servers in the AWS Cloud.

Not sure about Lambda tho

upvoted 3 times

Efren 3 years ago

Maybe yes

The Lambda authorizer authenticates the token with the third-party identity provider.

upvoted 2 times

cloudenthusiast 3 years ago

Also from ChatGPT

AWS Transfer Family supports multiple protocols, including AS2, and can be used for data transfer. By utilizing AWS Transfer Family, the company can integrate its own IdP authentication by creating an AWS Lambda function.

Both options D and C are valid solutions for the given requirements. The choice between them would depend on additional factors such as specific preferences, existing infrastructure, and overall architectural considerations.

upvoted 2 times

A solutions architect is designing a RESTAPI in Amazon API Gateway for a cash payback service. The application requires 1 GB of memory and 2 GB of storage for its computation resources. The application will require that the data is in a relational format.

Which additional combination of AWS services will meet these requirements with the LEAST administrative effort? (Choose two.)

- A. Amazon EC2
- B. AWS Lambda **Most Voted**
- C. Amazon RDS **Most Voted**
- D. Amazon DynamoDB
- E. Amazon Elastic Kubernetes Services (Amazon EKS)

Correct Answer: BC

Community vote distribution

BC (89%)

11%

Comments

cloudenthusiast **Highly Voted** 3 years ago

Selected Answer: BC

"The application will require that the data is in a relational format" so DynamoDB is out. RDS is the choice. Lambda is serverless.
upvoted 17 times

DougZ 11 months, 1 week ago

"LEAST administrative effort"

upvoted 1 times

smurfing2k17 **Most Recent** 2 years, 1 month ago

Why cant it be AC? We don't know the time of job runs right?

upvoted 2 times

JA2018 1 year, 6 months ago

requirement/ key not stated in stem

upvoted 1 times

TariqKipkemei 2 years, 12 months ago

Selected Answer: BC

AWS Lambda and Amazon RDS

upvoted 3 times

handsonlabsaws 3 years ago

Selected Answer: AC

"2 GB of storage for its COMPUTATION resources" the maximum for Lambda is 512MB.

upvoted 3 times

PLN6302 2 years, 9 months ago

Lambda now supports upto 10GB of memory

upvoted 8 times

Kp88 2 years, 10 months ago

I thought the same but seems like you can go all the way to 10gb. 512mb is the free tier

<https://docs.aws.amazon.com/lambda/latest/dg/configuration-function-common.html#configuration-ephemeral-storage>

upvoted 3 times

r3mo 2 years, 12 months ago

At first I was thinking the same. But the computation memory for the lambda function is 1gb not 2gb. Hence. if you go to basic settings when you create the lambda function you can select a in the memori settings the 1024 MB (1Gb) and that solve the problem.

upvoted 2 times

Efren 3 years ago

Selected Answer: BC

Relational Data RDS and computing for Lambda

upvoted 4 times

nosense 3 years ago

bc for me

upvoted 3 times

A company uses AWS Organizations to run workloads within multiple AWS accounts. A tagging policy adds department tags to AWS resources when the company creates tags.

An accounting team needs to determine spending on Amazon EC2 consumption. The accounting team must determine which departments are responsible for the costs regardless of AWS account. The accounting team has access to AWS Cost Explorer for all AWS accounts within the organization and needs to access all reports from Cost Explorer.

Which solution meets these requirements in the MOST operationally efficient way?

- A. From the Organizations management account billing console, activate a user-defined cost allocation tag named department. Create one cost report in Cost Explorer grouping by tag name, and filter by EC2. **Most Voted**
- B. From the Organizations management account billing console, activate an AWS-defined cost allocation tag named department. Create one cost report in Cost Explorer grouping by tag name, and filter by EC2.
- C. From the Organizations member account billing console, activate a user-defined cost allocation tag named department. Create one cost report in Cost Explorer grouping by the tag name, and filter by EC2.
- D. From the Organizations member account billing console, activate an AWS-defined cost allocation tag named department. Create one cost report in Cost Explorer grouping by tag name, and filter by EC2.

Correct Answer: A

Community vote distribution

A (100%)

Comments

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: A

By activating a user-defined cost allocation tag named "department" and creating a cost report in Cost Explorer that groups by the tag name and filters by EC2, the accounting team will be able to track and attribute costs to specific departments across all AWS accounts within the organization. This approach allows for consistent cost allocation and reporting regardless of the AWS account structure.

upvoted 9 times

luisgu **Highly Voted** 2 years, 6 months ago

Selected Answer: A

<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/custom-tags.html>

upvoted 6 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: A

Management not user.

upvoted 5 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: A

From the Organizations management account billing console, activate a user-defined cost allocation tag named department. Create one cost report in Cost Explorer grouping by tag name, and filter by EC2.

upvoted 5 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: A

From the Organizations management account billing console, activate a user-defined cost allocation tag named department. Create one cost report in Cost Explorer grouping by tag name, and filter by EC2.

upvoted 3 times

hiroohiroo 2 years, 6 months ago

Selected Answer: A

https://docs.aws.amazon.com/ja_jp/awsaccountbilling/latest/aboutv2/activating-tags.html

upvoted 4 times

nosense 2 years, 6 months ago

Selected Answer: A

a for me

upvoted 3 times

A company wants to securely exchange data between its software as a service (SaaS) application Salesforce account and Amazon S3. The company must encrypt the data at rest by using AWS Key Management Service (AWS KMS) customer managed keys (CMKs). The company must also encrypt the data in transit. The company has enabled API access for the Salesforce account.

- A. Create AWS Lambda functions to transfer the data securely from Salesforce to Amazon S3.
- B. Create an AWS Step Functions workflow. Define the task to transfer the data securely from Salesforce to Amazon S3.
- C. Create Amazon AppFlow flows to transfer the data securely from Salesforce to Amazon S3. **Most Voted**
- D. Create a custom connector for Salesforce to transfer the data securely from Salesforce to Amazon S3.

Correct Answer: C

Community vote distribution

C (100%)

Comments

cloudenthusiast **Highly Voted** 3 years ago

Selected Answer: C

Amazon AppFlow is a fully managed integration service that allows you to securely transfer data between different SaaS applications and AWS services. It provides built-in encryption options and supports encryption in transit using SSL/TLS protocols. With AppFlow, you can configure the data transfer flow from Salesforce to Amazon S3, ensuring data encryption at rest by utilizing AWS KMS CMKs.

upvoted 15 times

Guru4Cloud **Highly Voted** 2 years, 9 months ago

Selected Answer: C

- ° Amazon AppFlow can securely transfer data between Salesforce and Amazon S3.
- ° AppFlow supports encrypting data at rest in S3 using KMS CMKs.
- ° AppFlow supports encrypting data in transit using HTTPS/TLS.
- ° AppFlow provides built-in support and templates for Salesforce and S3, requiring less custom configuration than solutions like Lambda, Step Functions, or custom connectors.
- ° So Amazon AppFlow is the easiest way to meet all the requirements of securely transferring data between Salesforce and S3 with encryption at rest and in transit.

upvoted 10 times

1e22522 **Most Recent** 1 year, 10 months ago

Selected Answer: C

i do like myself some appflow flow

upvoted 2 times

zinabu 2 years, 2 months ago

SAAS=aws appflow

upvoted 4 times

cvoiceip 2 years, 5 months ago

Ans : C

Salesforce -----> Amazon AppFlow -----> S3

upvoted 3 times

hsinchang 2 years, 10 months ago

securely transfer data between Software-as-a-Service (SaaS) applications and AWS -> AppFlow

upvoted 3 times

TariqKipkemei 2 years, 12 months ago

Selected Answer: C

With Amazon AppFlow automate bi-directional data flows between SaaS applications and AWS services in just a few clicks

upvoted 2 times

DrWatson 3 years ago

Selected Answer: C

<https://docs.aws.amazon.com/appflow/latest/userguide/what-is-appflow.html>

upvoted 3 times

Abrar2022 3 years ago

All you need to know is that AWS AppFlow securely transfers data between different SaaS applications and AWS services

upvoted 3 times

hiroohiroo 3 years ago

Selected Answer: C

<https://docs.aws.amazon.com/appflow/latest/userguide/salesforce.html>

upvoted 4 times

Efren 3 years ago

Selected Answer: C

SaaS with another service, AppFlow

upvoted 2 times

A company is developing a mobile gaming app in a single AWS Region. The app runs on multiple Amazon EC2 instances in an Auto Scaling group. The company stores the app data in Amazon DynamoDB. The app communicates by using TCP traffic and UDP traffic between the users and the servers. The application will be used globally. The company wants to ensure the lowest possible latency for all users.

Which solution will meet these requirements?

A. Use AWS Global Accelerator to create an accelerator. Create an Application Load Balancer (ALB) behind an accelerator endpoint that uses Global Accelerator integration and listening on the TCP and UDP ports. Update the Auto Scaling group to register instances on the ALB.

B. Use AWS Global Accelerator to create an accelerator. Create a Network Load Balancer (NLB) behind an accelerator endpoint that uses Global Accelerator integration and listening on the TCP and UDP ports. Update the Auto Scaling group to register instances on the NLB. **Most Voted**

C. Create an Amazon CloudFront content delivery network (CDN) endpoint. Create a Network Load Balancer (NLB) behind the endpoint and listening on the TCP and UDP ports. Update the Auto Scaling group to register instances on the NLB. Update CloudFront to use the NLB as the origin.

D. Create an Amazon CloudFront content delivery network (CDN) endpoint. Create an Application Load Balancer (ALB) behind the endpoint and listening on the TCP and UDP ports. Update the Auto Scaling group to register instances on the ALB. Update CloudFront to use the ALB as the origin.

Correct Answer: B

Community vote distribution

B (100%)

Comments

zinabu **Highly Voted** 1 year, 8 months ago

TCP/UDP/IP based communication with server =NLB
for global low latency communication if IP/udp/tcp based = aws global accelerator
upvoted 5 times

sandordini **Most Recent** 1 year, 7 months ago

Selected Answer: B

Mobile gaming, UDP > AWS Global Accelerator + NLB
upvoted 3 times

Mikado211 1 year, 11 months ago

Selected Answer: B

UDP == NLB
NLB can't be used with Cloudfront, so we have to play with AWS Global accelerator
upvoted 4 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: B

Use AWS Global Accelerator to create an accelerator. Create a Network Load Balancer (NLB) behind an accelerator endpoint that uses Global Accelerator integration and listening on the TCP and UDP ports. Update the Auto Scaling group to register instances on the NLB
upvoted 3 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: B

TCP and UDP = global accelerator and Network Load Balancer

upvoted 3 times

antropaws 2 years, 6 months ago

Selected Answer: B

Clearly B.

upvoted 2 times

eddie5049 2 years, 6 months ago

Selected Answer: B

NLB + Accelerator

upvoted 4 times

hiroohiroo 2 years, 6 months ago

Selected Answer: B

AWS Global Accelerator+NLB

upvoted 4 times

Efren 2 years, 6 months ago

Selected Answer: B

UDP, Global Accelerator plus NLB

upvoted 2 times

nosense 2 years, 6 months ago

Selected Answer: B

AWS Global Accelerator is a better solution for the mobile gaming app than CloudFront

upvoted 4 times

A company has an application that processes customer orders. The company hosts the application on an Amazon EC2 instance that saves the orders to an Amazon Aurora database. Occasionally when traffic is high the workload does not process orders fast enough.

What should a solutions architect do to write the orders reliably to the database as quickly as possible?

- A. Increase the instance size of the EC2 instance when traffic is high. Write orders to Amazon Simple Notification Service (Amazon SNS). Subscribe the database endpoint to the SNS topic.
- B. Write orders to an Amazon Simple Queue Service (Amazon SQS) queue. Use EC2 instances in an Auto Scaling group behind an Application Load Balancer to read from the SQS queue and process orders into the database. **Most Voted**
- C. Write orders to Amazon Simple Notification Service (Amazon SNS). Subscribe the database endpoint to the SNS topic. Use EC2 instances in an Auto Scaling group behind an Application Load Balancer to read from the SNS topic.
- D. Write orders to an Amazon Simple Queue Service (Amazon SQS) queue when the EC2 instance reaches CPU threshold limits. Use scheduled scaling of EC2 instances in an Auto Scaling group behind an Application Load Balancer to read from the SQS queue and process orders into the database.

Correct Answer: B

Community vote distribution

B (100%)

Comments

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: B

By decoupling the write operation from the processing operation using SQS, you ensure that the orders are reliably stored in the queue, regardless of the processing capacity of the EC2 instances. This allows the processing to be performed at a scalable rate based on the available EC2 instances, improving the overall reliability and speed of order processing.

upvoted 12 times

omarshaban **Highly Voted** 1 year, 11 months ago

IN MY EXAM

upvoted 10 times

Guru4Cloud **Most Recent** 2 years, 3 months ago

Selected Answer: B

Decoupling the order processing from the application using Amazon SQS and leveraging Auto Scaling to handle the processing of orders based on the workload in the SQS queue is indeed the most efficient and scalable approach. This architecture addresses both reliability and performance concerns during traffic spikes.

upvoted 4 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: B

Write orders to an Amazon Simple Queue Service (Amazon SQS) queue. Use EC2 instances in an Auto Scaling group behind an Application Load Balancer to read from the SQS queue and process orders into the database.

upvoted 3 times

antropaws 2 years, 6 months ago

Selected Answer: B

100% B.

upvoted 3 times

omoakin 2 years, 6 months ago

BBBBBBBBBB

upvoted 3 times

An IoT company is releasing a mattress that has sensors to collect data about a user's sleep. The sensors will send data to an Amazon S3 bucket. The sensors collect approximately 2 MB of data every night for each mattress. The company must process and summarize the data for each mattress. The results need to be available as soon as possible. Data processing will require 1 GB of memory and will finish within 30 seconds.

Which solution will meet these requirements MOST cost-effectively?

- A. Use AWS Glue with a Scala job
- B. Use Amazon EMR with an Apache Spark script
- C. Use AWS Lambda with a Python script **Most Voted**
- D. Use AWS Glue with a PySpark job

Correct Answer: C

Community vote distribution

C (100%)

Comments

cloudehustiaist **Highly Voted** 2 years, 6 months ago

Selected Answer: C

AWS Lambda charges you based on the number of invocations and the execution time of your function. Since the data processing job is relatively small (2 MB of data), Lambda is a cost-effective choice. You only pay for the actual usage without the need to provision and maintain infrastructure.
upvoted 7 times

joechen2023 2 years, 5 months ago

but the question states "Data processing will require 1 GB of memory and will finish within 30 seconds." so it can't be C as Lambda support maximum 512M
upvoted 2 times

DougZ 11 months, 1 week ago

As it only needs 30 seconds, Lambda (max 15 mins) should be a good candidate
upvoted 1 times

nilandd44gg 2 years, 4 months ago

C is valid.
Lambda quotas:
Memory - 128 MB to 10,240 MB, in 1-MB increments.

Note: Lambda allocates CPU power in proportion to the amount of memory configured. You can increase or decrease the memory and CPU power allocated to your function using the Memory (MB) setting. At 1,769 MB, a function has the equivalent of one vCPU.

Function timeout 900 seconds (15 minutes)

4 KB, for all environment variables associated with the function, in aggregate
<https://docs.aws.amazon.com/lambda/latest/dg/gettingstarted-limits.html>
upvoted 3 times

BillaRanga 1 year, 10 months ago

Lambda can support upto 10 GB, But 512M is under free tier
upvoted 3 times

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: C

The data processing is lightweight, only requiring 1 GB memory and finishing in under 30 seconds. Lambda is designed for short, transient workloads like this.

Lambda scales automatically, invoking the function as needed when new data arrives. No servers to manage. Lambda has a very low cost. You only pay for the compute time used to run the function, billed in 100ms increments. Much cheaper than provisioning EMR or Glue. Processing can begin as soon as new data hits the S3 bucket by triggering the Lambda function. Provides low latency.

upvoted 5 times

Chiquitabandita Most Recent 2 years, 1 month ago

I understand C is a common answer "throw Lambda" seems to be a common theme for questions that need processing under 15 minutes for the test. but in reality, can the other solutions be viable options as well?

upvoted 4 times

Mikado211 2 years ago

That's the point here, technically all the options are good and will work, but since we are on a small amount of data Lambda will be the cheapest one, usually Glue or EMR will be kept for a big amount of data.

Here is a topic where people did a comparison in comments : https://www.reddit.com/r/aws/comments/9umxv1/aws_glue_vs_lambda_costbenefit/

upvoted 4 times

TariqKipkemei 2 years, 1 month ago

Selected Answer: C

"processing will require 1 GB of memory and will finish within 30 seconds", perfect for AWS Lambda.

upvoted 5 times

antropaws 2 years, 6 months ago

Selected Answer: C

I reckon C, but I would consider other well founded options.

upvoted 3 times

nosense 2 years, 6 months ago

Selected Answer: C

c anyway the MOST cost-effectively

upvoted 3 times

A company hosts an online shopping application that stores all orders in an Amazon RDS for PostgreSQL Single-AZ DB instance. Management wants to eliminate single points of failure and has asked a solutions architect to recommend an approach to minimize database downtime without requiring any changes to the application code.

Which solution meets these requirements?

- A. Convert the existing database instance to a Multi-AZ deployment by modifying the database instance and specifying the Multi-AZ option. **Most Voted**
- B. Create a new RDS Multi-AZ deployment. Take a snapshot of the current RDS instance and restore the new Multi-AZ deployment with the snapshot.
- C. Create a read-only replica of the PostgreSQL database in another Availability Zone. Use Amazon Route 53 weighted record sets to distribute requests across the databases.
- D. Place the RDS for PostgreSQL database in an Amazon EC2 Auto Scaling group with a minimum group size of two. Use Amazon Route 53 weighted record sets to distribute requests across instances.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: A

A. Convert the existing database instance to a Multi-AZ deployment by modifying the database instance and specifying the Multi-AZ option
upvoted 6 times

Abrar2022 **Highly Voted** 2 years, 6 months ago

"minimize database downtime" so why create a new DB just modify the existing one so no time is wasted.
upvoted 5 times

awsgeek75 **Most Recent** 1 year, 10 months ago

Selected Answer: A

A is correct but reason need to be clarified:
<https://aws.amazon.com/blogs/database/best-practices-for-converting-a-single-az-amazon-rds-instance-to-a-multi-az-instance/>

The instance doesn't automatically convert to Multi-AZ immediately. By default it will convert at next maintenance window but you can convert it immediately. Compared to B this is much better. CD are too many changes overall so unsuitable.
upvoted 3 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: A

Eliminate single points of failure = Multi-AZ deployment
upvoted 5 times

antropaws 2 years, 6 months ago

Selected Answer: A

A) <https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.MultiAZSingleStandby.html#Concepts.MultiAZ.Migrating>
upvoted 2 times

cloudenthusiast 2 years, 6 months ago

Selected Answer: A

Compared to other solutions that involve creating new instances, restoring snapshots, or setting up replication manually, converting to a Multi-AZ deployment is a simpler and more streamlined approach with lower overhead.

Overall, option A offers a cost-effective and efficient way to minimize database downtime without requiring significant changes or additional complexities.

upvoted 3 times

Efren 2 years, 6 months ago

A for HA, but also read replica can convert itself to master if the master is down... so not sure if C?

upvoted 1 times

Efren 2 years, 6 months ago

Sorry, the Route 53 doesnt make sense to sent requests to RR , what if is a write?

upvoted 2 times

nosense 2 years, 6 months ago

Selected Answer: A

i guess aa

upvoted 3 times


A company is developing an application to support customer demands. The company wants to deploy the application on multiple Amazon EC2 Nitro-based instances within the same Availability Zone. The company also wants to give the application the ability to write to multiple block storage volumes in multiple EC2 Nitro-based instances simultaneously to achieve higher application availability.

Which solution will meet these requirements?

- A. Use General Purpose SSD (gp3) EBS volumes with Amazon Elastic Block Store (Amazon EBS) Multi-Attach
- B. Use Throughput Optimized HDD (st1) EBS volumes with Amazon Elastic Block Store (Amazon EBS) Multi-Attach
- C. Use Provisioned IOPS SSD (io2) EBS volumes with Amazon Elastic Block Store (Amazon EBS) Multi-Attach **Most Voted**
- D. Use General Purpose SSD (gp2) EBS volumes with Amazon Elastic Block Store (Amazon EBS) Multi-Attach

Correct Answer: C

Community vote distribution



Comments

potomac **Highly Voted** 2 years, 1 month ago

Selected Answer: C

Multi-Attach is supported exclusively on Provisioned IOPS SSD (io1 and io2) volumes.
upvoted 13 times

awsgeek75 **Highly Voted** 1 year, 10 months ago

Selected Answer: C

hdd < gp2 < gp3 < io2
upvoted 7 times

In2718 **Most Recent** 11 months ago

Selected Answer: A

A follow-up as to why this is a terrible question - the application could either have throughput or IOPS as its weakest link. In a real-world meeting with the designer team, I would tell them that until you provide real numbers under load in UAT you are going to get gp3 as a best all-around solution.
upvoted 1 times

In2718 11 months ago

Selected Answer: A

This is a terribly written question - there is no benchmark specifications for either IOPS or throughput. Because of this if you guess anything other than gp3 you would be guessing incorrectly. In real-life without such specifications, you would choose gp3. Terrible question.
upvoted 1 times

master9 1 year, 11 months ago

Selected Answer: C

AWS IO2 does support Multi-Attach. Multi-Attach allows you to share access to an EBS data volume between up to 16 Nitro-based EC2 instances within the same Availability Zone. Each attached instance has full read and write permission to the shared volume. This feature is intended to make it easier to achieve higher application availability for customers that want to deploy applications that manage storage consistency from multiple writers in shared storage infrastructure. However, please note that Multi-Attach on io2 is available in certain regions only.
upvoted 7 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: C

C. Use Provisioned IOPS SSD (io2) EBS volumes with Amazon Elastic Block Store (Amazon EBS) Multi-Attach
upvoted 5 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: C

Multi-Attach is supported exclusively on Provisioned IOPS SSD (io1 and io2) volumes.

[https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html#:~:text=Multi%2DAttach%20is%20supported%20exclusively%20on%20Provisioned%20IOPS%20SSD%20\(io1%20and%20io2\)%20volumes](https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html#:~:text=Multi%2DAttach%20is%20supported%20exclusively%20on%20Provisioned%20IOPS%20SSD%20(io1%20and%20io2)%20volumes)
upvoted 4 times

Axeashes 2 years, 5 months ago

Multi-Attach is supported exclusively on Provisioned IOPS SSD (io1 and io2) volumes.

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html>
upvoted 4 times

Uzi_m 2 years, 6 months ago

The correct answer is A.

Currently, Multi Attach EBS feature is supported by gp3 volumes also.

Multi-Attach is supported for certain EBS volume types, including io1, io2, gp3, st1, and sc1 volumes.

upvoted 2 times

Kp88 2 years, 4 months ago

No , Read this --> <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html#considerations>

upvoted 4 times

AshishRocks 2 years, 6 months ago

Answer should be D

upvoted 1 times

Kp88 2 years, 4 months ago

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html#considerations>

upvoted 2 times

AshishRocks 2 years, 6 months ago

By ChatGPT - Create General Purpose SSD (gp2) volumes: Provision multiple gp2 volumes with the required capacity for your application.

upvoted 1 times

Kp88 2 years, 4 months ago

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html#considerations>

upvoted 1 times

AshishRocks 2 years, 6 months ago

Multi-Attach does not support Provisioned IOPS SSD (io2) volumes. Multi-Attach is currently available only for General Purpose SSD (gp2), Throughput Optimized HDD (st1), and Cold HDD (sc1) EBS volumes.

upvoted 1 times

Abrar2022 2 years, 6 months ago

Multi-Attach is supported exclusively on Provisioned IOPS SSD (io1 or io2) volumes.

upvoted 2 times

elmogy 2 years, 6 months ago

Selected Answer: C

only io1/io2 supports Multi-Attach

upvoted 4 times

Uzi_m 2 years, 6 months ago

Multi-Attach is supported for certain EBS volume types, including io1, io2, gp3, st1, and sc1 volumes.

upvoted 1 times

Kp88 2 years, 4 months ago

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html#considerations>

upvoted 2 times

examtopictempacc 2 years, 6 months ago

Selected Answer: C

only io1/io2 supports Multi-Attach

upvoted 4 times

Vlad 2 years, 6 months ago

Selected Answer: A

Option D suggests using General Purpose SSD (gp2) EBS volumes with Amazon EBS Multi-Attach. While gp2 volumes support multi-attach, gp3 volumes offer a more cost-effective solution with enhanced performance characteristics.

upvoted 1 times

Vlad 2 years, 6 months ago

I'm sorry :

Multi-Attach enabled volumes can be attached to up to 16 instances built on the Nitro System that are in the same Availability Zone. Multi-Attach is supported exclusively on Provisioned IOPS SSD (io1 or io2) volumes.

upvoted 3 times

Vlad 2 years, 6 months ago

The answer is C:

upvoted 2 times

EA100 2 years, 6 months ago

Answer - C

C. Use Provisioned IOPS SSD (io2) EBS volumes with Amazon Elastic Block Store (Amazon EBS) Multi-Attach.

While both option C and option D can support Amazon EBS Multi-Attach, using Provisioned IOPS SSD (io2) EBS volumes provides higher performance and lower latency compared to General Purpose SSD (gp2) volumes. This makes io2 volumes better suited for demanding and mission-critical applications where performance is crucial.

If the goal is to achieve higher application availability and ensure optimal performance, using Provisioned IOPS SSD (io2) EBS volumes with Multi-Attach will provide the best results.

upvoted 3 times

A company designed a stateless two-tier application that uses Amazon EC2 in a single Availability Zone and an Amazon RDS Multi-AZ DB instance. New company management wants to ensure the application is highly available.

What should a solutions architect do to meet this requirement?

- A. Configure the application to use Multi-AZ EC2 Auto Scaling and create an Application Load Balancer **Most Voted**
- B. Configure the application to take snapshots of the EC2 instances and send them to a different AWS Region
- C. Configure the application to use Amazon Route 53 latency-based routing to feed requests to the application
- D. Configure Amazon Route 53 rules to handle incoming requests and create a Multi-AZ Application Load Balancer

Correct Answer: A

Community vote distribution

A (100%)

Comments

nosense **Highly Voted** 2 years ago

Selected Answer: A

it's A

upvoted 6 times

cloudehustiaist **Highly Voted** 2 years ago

Selected Answer: A

By combining Multi-AZ EC2 Auto Scaling and an Application Load Balancer, you achieve high availability for the EC2 instances hosting your stateless two-tier application.

upvoted 5 times

Guru4Cloud **Most Recent** 1 year, 9 months ago

Selected Answer: A

A. Configure the application to use Multi-AZ EC2 Auto Scaling and create an Application Load Balancer

upvoted 2 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: A

Highly available = Multi-AZ EC2 Auto Scaling and Application Load Balancer.

upvoted 3 times

antropaws 2 years ago

Selected Answer: A

Most likely A.

upvoted 2 times

A company uses AWS Organizations. A member account has purchased a Compute Savings Plan. Because of changes in the workloads inside the member account, the account no longer receives the full benefit of the Compute Savings Plan commitment. The company uses less than 50% of its purchased compute power.

- A. Turn on discount sharing from the Billing Preferences section of the account console in the member account that purchased the Compute Savings Plan.
- B. Turn on discount sharing from the Billing Preferences section of the account console in the company's Organizations management account. **Most Voted**
- C. Migrate additional compute workloads from another AWS account to the account that has the Compute Savings Plan.
- D. Sell the excess Savings Plan commitment in the Reserved Instance Marketplace.

Correct Answer: B

Community vote distribution

B (74%)

D (26%)

Comments

norris81 **Highly Voted** 3 years ago

Selected Answer: B

<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/ri-turn-off.html>

Sign in to the AWS Management Console and open the AWS Billing console at <https://console.aws.amazon.com/billing/>

Note

Ensure you're logged in to the management account of your AWS Organizations.

upvoted 11 times

NayeraB **Highly Voted** 2 years, 3 months ago

Option E, Take it out of the salary of the guy who made the decision to purchase an entire compute plan without studying the company's needs.

upvoted 7 times

jingen11 **Most Recent** 1 year, 7 months ago

Selected Answer: B

I dont think there is a way for us to sell excess Savings plan. Only Selling Reserved instances is possible in marketplace

upvoted 2 times

jingen11 1 year, 7 months ago

<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/ri-turn-off.html#ri-turn-on-process>

upvoted 2 times

Stranko 2 years, 3 months ago

Selected Answer: D

I'd go with D, due to "The company uses less than 50% of its purchased compute power". Like, why are you sharing it between other accounts of the company, if the company itself doesn't need it? If you provisioned too much you can sell the overprovisioned capacity on the market. I'd understand B if it was about the account using about 50% of the plan and other accounts running similar workloads, but no such thing is stated.

upvoted 2 times

mr123dd 2 years, 5 months ago

Selected Answer: D

in the question, it does not clarify then number of accounts the company has, if they only has one account, I think it is D,

upvoted 1 times

Mujahid_1 2 years, 5 months ago

what are you guys doing
this section is for discussion not for copy paste
upvoted 4 times

pentium75 2 years, 5 months ago

Selected Answer: B

B, it's a generic Compute Savings Plan that can be used for compute workloads in the other accounts.

A doesn't work, discount sharing must be enabled for all accounts (at least for those that provide and share the discounts).

C is not possible, there's a reason why the workloads are in different accounts.

D would be a last resort if there wouldn't be any other workloads in the own organization, but here are.

upvoted 4 times

baba365 2 years, 8 months ago

So what exactly is the question?
upvoted 7 times

awsgeek75 2 years, 4 months ago

It's an English test on complete the sentence...
upvoted 4 times

pentium75 2 years, 5 months ago

What to do
upvoted 2 times

michalf84 2 years, 8 months ago

Selected Answer: D

I saw similar question in older exam one can sell on the market unused capacity
upvoted 1 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: B

B. Turn on discount sharing from the Billing Preferences section of the account console in the company's Organizations management account
upvoted 3 times

Lx016 2 years, 9 months ago

Bro, no need to copy paste the answer that is already written. Need an explanation, I see that you just copy pasting the potential answers without any explanation in each discussion.
upvoted 26 times

live_reply_developers 2 years, 11 months ago

Selected Answer: D

"For example, you might want to sell Reserved Instances after moving instances to a new AWS Region, changing to a new instance type, ending projects before the term expiration, when your business needs change, or if you have unneeded capacity."

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ri-market-general.html>

upvoted 2 times

pentium75 2 years, 5 months ago

D would make sense if the company wouldn't have other accounts with workloads. Or if it would be EC2 Savings Plans that would not match the instance types in other accounts. But it's a generic Compute Savings Plan that surely can be used in another account. Thus B.
upvoted 2 times

awsgeek75 2 years, 4 months ago

I am also confused between B and D as the last part of the question "The company uses less than 50% of its purchased compute power." could imply that the whole company (not just this member account) only uses 50% of the computer power. If they said the member account only uses 50% then it would be clear cut B.
upvoted 2 times

TariqKipkemei 2 years, 11 months ago

Selected Answer: B

answer is B.

<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/ri-turn-off.html#:~:text=choose%20Save%20Turning%20on%20shared%20reserved%20instances%20and%20Savings%20Plans%20discounts,-You%20can%20use>

upvoted 2 times

Felix_br 3 years ago

Selected Answer: D

The company uses less than 50% of its purchased compute power.

For this reason i believe D is the best solution : <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ri-market-general.html>

upvoted 3 times

Abrar2022 3 years ago

The company Organization's management account can turn on/off shared reserved instances.

upvoted 2 times

cloudenthusiast 3 years ago

Selected Answer: B

To summarize, option C (Migrate additional compute workloads from another AWS account to the account that has the Compute Savings Plan) is a valid solution to address the underutilization of the Compute Savings Plan. However, it involves workload migration and may require careful planning and coordination. Consider the feasibility and impact of migrating workloads before implementing this solution.

upvoted 3 times

EA100 3 years ago

Answer - C

If a member account within AWS Organizations has purchased a Compute Savings Plan

upvoted 1 times

EA100 3 years ago

Asnwer - C

upvoted 1 times

A company is developing a microservices application that will provide a search catalog for customers. The company must use REST APIs to present the frontend of the application to users. The REST APIs must access the backend services that the company hosts in containers in private VPC subnets.

Which solution will meet these requirements?

- A. Design a WebSocket API by using Amazon API Gateway. Host the application in Amazon Elastic Container Service (Amazon ECS) in a private subnet. Create a private VPC link for API Gateway to access Amazon ECS.
- B. Design a REST API by using Amazon API Gateway. Host the application in Amazon Elastic Container Service (Amazon ECS) in a private subnet. Create a private VPC link for API Gateway to access Amazon ECS. **Most Voted**
- C. Design a WebSocket API by using Amazon API Gateway. Host the application in Amazon Elastic Container Service (Amazon ECS) in a private subnet. Create a security group for API Gateway to access Amazon ECS.
- D. Design a REST API by using Amazon API Gateway. Host the application in Amazon Elastic Container Service (Amazon ECS) in a private subnet. Create a security group for API Gateway to access Amazon ECS.

Correct Answer: B

Community vote distribution

B (100%)

Comments

cloudenthusiast **Highly Voted** 3 years ago

Selected Answer: B

REST API with Amazon API Gateway: REST APIs are the appropriate choice for providing the frontend of the microservices application. Amazon API Gateway allows you to design, deploy, and manage REST APIs at scale.

Amazon ECS in a Private Subnet: Hosting the application in Amazon ECS in a private subnet ensures that the containers are securely deployed within the VPC and not directly exposed to the public internet.

Private VPC Link: To enable the REST API in API Gateway to access the backend services hosted in Amazon ECS, you can create a private VPC link. This establishes a private network connection between the API Gateway and ECS containers, allowing secure communication without traversing the public internet.

upvoted 14 times

MNotABot **Highly Voted** 2 years, 11 months ago

Question itself says: "The company must use REST APIs", hence WebSocket APIs are not applicable and such options are eliminated straight away.

upvoted 9 times

MatAlves **Most Recent** 1 year, 9 months ago

Selected Answer: B

"VPC links enable you to create private integrations that connect your HTTP API routes to private resources in a VPC, such as Application Load Balancers or Amazon ECS container-based applications."

upvoted 3 times

freedafeng 1 year, 10 months ago

I think the connection should be from the application to the ECS in the private VPC, instead of from the API Gateway to the ECS in the private VPC. API Gateway only needs to connect to the application...

upvoted 1 times

awsgeek75 2 years, 4 months ago

Selected Answer: B

AC are wrong as they are not REST API

D, you don't make SG for API Gateway to EC2, you have to make a VPC Link. More details at <https://docs.aws.amazon.com/apigateway/latest/developerguide/http-api-vpc-links.html>

upvoted 5 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: B

To allow the REST APIs to securely access the backend, a private VPC link should be created from API Gateway to the ECS containers. A private VPC link provides private connectivity between API Gateway and the VPC without using public IP addresses or requiring an internet gateway/NAT

upvoted 5 times

Axeashes 2 years, 11 months ago

Selected Answer: B

<https://docs.aws.amazon.com/apigateway/latest/developerguide/http-api-private-integration.html>

upvoted 2 times

TariqKipkemei 2 years, 11 months ago

Selected Answer: B

A VPC link is a resource in Amazon API Gateway that allows for connecting API routes to private resources inside a VPC.

upvoted 3 times

samehpalass 2 years, 11 months ago

B is the right choice

upvoted 2 times

Yadav_Sanjay 2 years, 11 months ago

Why Not D

upvoted 3 times

potomac 2 years, 7 months ago

A security group acts as a firewall for associated EC2 instances, controlling both inbound and outbound traffic at the instance level.

upvoted 2 times

nosense 3 years ago

Selected Answer: B

b is right, bcs vpc link provided security connection

upvoted 4 times

A company stores raw collected data in an Amazon S3 bucket. The data is used for several types of analytics on behalf of the company's customers. The type of analytics requested determines the access pattern on the S3 objects.

The company cannot predict or control the access pattern. The company wants to reduce its S3 costs.

Which solution will meet these requirements?

- A. Use S3 replication to transition infrequently accessed objects to S3 Standard-Infrequent Access (S3 Standard-IA)
- B. Use S3 Lifecycle rules to transition objects from S3 Standard to Standard-Infrequent Access (S3 Standard-IA)
- C. Use S3 Lifecycle rules to transition objects from S3 Standard to S3 Intelligent-Tiering **Most Voted**
- D. Use S3 Inventory to identify and transition objects that have not been accessed from S3 Standard to S3 Intelligent-Tiering

Correct Answer: C

Community vote distribution

C (100%)

Comments

nosense **Highly Voted** 2 years ago

Selected Answer: C

S3 Inventory can't to move files to another class
upvoted 8 times

Murtadhaceit **Highly Voted** 1 year, 6 months ago

Selected Answer: C

Unpredictable access pattern = Intelligent-Tiering.
upvoted 5 times

allanm **Most Recent** 10 months ago

Selected Answer: C

Tricky answers - While S3 inventory can identify all objects within the bucket, it cannot be used to transition S3 classes.
upvoted 1 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

C. Use S3 Lifecycle rules to transition objects from S3 Standard to S3 Intelligent-Tiering
upvoted 4 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: C

Cannot predict access pattern = S3 Intelligent-Tiering.
upvoted 4 times

Efren 2 years ago

Selected Answer: C

Not known patterns, Intelligent Tier
upvoted 4 times

A company has applications hosted on Amazon EC2 instances with IPv6 addresses. The applications must initiate communications with other external applications using the internet. However the company's security policy states that any external service cannot initiate a connection to the EC2 instances.

What should a solutions architect recommend to resolve this issue?

- A. Create a NAT gateway and make it the destination of the subnet's route table
- B. Create an internet gateway and make it the destination of the subnet's route table
- C. Create a virtual private gateway and make it the destination of the subnet's route table
- D. Create an egress-only internet gateway and make it the destination of the subnet's route table **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

wRhIH **Highly Voted** 2 years, 11 months ago

For exam,
egress-only internet gateway: IPv6
NAT gateway: IPv4
upvoted 58 times

b82faaf 2 years, 6 months ago

This is very helpful, thanks.
upvoted 4 times

MatAlves 1 year, 9 months ago

Good stuff.

"An egress-only internet gateway is for use with IPv6 traffic only. To enable outbound-only internet communication over IPv4, use a NAT gateway instead."

<https://docs.aws.amazon.com/vpc/latest/userguide/egress-only-internet-gateway.html>
upvoted 2 times

RDM10 2 years, 8 months ago

thanks a lot
upvoted 4 times

cloudenthusiast **Highly Voted** 3 years ago

Selected Answer: D

An egress-only internet gateway (EIGW) is specifically designed for IPv6-only VPCs and provides outbound IPv6 internet access while blocking inbound IPv6 traffic. It satisfies the requirement of preventing external services from initiating connections to the EC2 instances while allowing the instances to initiate outbound communications.

upvoted 9 times

cloudenthusiast 3 years ago

Since the company's security policy explicitly states that external services cannot initiate connections to the EC2 instances, using a NAT gateway (option A) would not be suitable. A NAT gateway allows outbound connections from private subnets to the internet, but it does not restrict inbound connections from external sources.

upvoted 5 times

pentium75 2 years, 5 months ago

"A NAT gateway ... does not restrict inbound connections from external sources." Actually it does, but only for IPv4.

upvoted 2 times

[Removed] 3 years ago

Enable outbound IPv6 traffic using an egress-only internet gateway

<https://docs.aws.amazon.com/vpc/latest/userguide/egress-only-internet-gateway.html>

upvoted 3 times

MatAlves **Most Recent** 1 year, 9 months ago

Selected Answer: D

"An egress-only internet gateway is for use with IPv6 traffic only. To enable outbound-only internet communication over IPv4, use a NAT gateway instead."

<https://docs.aws.amazon.com/vpc/latest/userguide/egress-only-internet-gateway.html>

upvoted 2 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: D

D. Create an egress-only internet gateway and make it the destination of the subnet's route table

upvoted 2 times

TariqKipkemei 2 years, 11 months ago

Selected Answer: D

Outbound traffic only = Create an egress-only internet gateway and make it the destination of the subnet's route table

upvoted 2 times

radev 3 years ago

Selected Answer: D

Egress-Only internet Gateway

upvoted 4 times

A company is creating an application that runs on containers in a VPC. The application stores and accesses data in an Amazon S3 bucket. During the development phase, the application will store and access 1 TB of data in Amazon S3 each day. The company wants to minimize costs and wants to prevent traffic from traversing the internet whenever possible.

Which solution will meet these requirements?

- A. Enable S3 Intelligent-Tiering for the S3 bucket
- B. Enable S3 Transfer Acceleration for the S3 bucket
- C. Create a gateway VPC endpoint for Amazon S3. Associate this endpoint with all route tables in the VPC **Most Voted**
- D. Create an interface endpoint for Amazon S3 in the VPC. Associate this endpoint with all route tables in the VPC

Correct Answer: C

Community vote distribution

C (100%)

Comments

litos168 **Highly Voted** 2 years, 4 months ago

Amazon S3 supports both gateway endpoints and interface endpoints. With a gateway endpoint, you can access Amazon S3 from your VPC, without requiring an internet gateway or NAT device for your VPC, and with no additional cost. However, gateway endpoints do not allow access from on-premises networks, from peered VPCs in other AWS Regions, or through a transit gateway. For those scenarios, you must use an interface endpoint, which is available for an additional cost.

upvoted 14 times

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: C

Gateway VPC Endpoint: A gateway VPC endpoint enables private connectivity between a VPC and Amazon S3. It allows direct access to Amazon S3 without the need for internet gateways, NAT devices, VPN connections, or AWS Direct Connect.

Minimize Internet Traffic: By creating a gateway VPC endpoint for Amazon S3 and associating it with all route tables in the VPC, the traffic between the VPC and Amazon S3 will be kept within the AWS network. This helps in minimizing data transfer costs and prevents the need for traffic to traverse the internet.

Cost-Effective: With a gateway VPC endpoint, the data transfer between the application running in the VPC and the S3 bucket stays within the AWS network, reducing the need for data transfer across the internet. This can result in cost savings, especially when dealing with large amounts of data.

upvoted 8 times

cloudenthusiast 2 years, 6 months ago

Option B (Enable S3 Transfer Acceleration for the S3 bucket) is a feature that uses the CloudFront global network to accelerate data transfers to and from Amazon S3. While it can improve data transfer speed, it still involves traffic traversing the internet and doesn't directly address the goal of minimizing costs and preventing internet traffic whenever possible.

upvoted 3 times

awsgeek75 **Most Recent** 1 year, 10 months ago

Selected Answer: C

<https://aws.amazon.com/blogs/architecture/choosing-your-vpc-endpoint-strategy-for-amazon-s3/>

A: Storage cost is not described as an issue here

B: Tx Accelerator is for external (global user) traffic acceleration

D: Interface endpoint is on-prem to S3

C: gateway VPC is specifically for S3 to AWS resources

upvoted 4 times

dkw2342 1 year, 9 months ago

Interface endpoints are not exclusively for on-prem to S3.

The only reason why option D is wrong is because "Associate this endpoint with all route tables in the VPC" makes no sense.
upvoted 1 times

bsbs1234 2 years, 2 months ago

I think both C&D will work.

But D will have extra cost. So C is correct.

upvoted 3 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: C

C. Create a gateway VPC endpoint for Amazon S3. Associate this endpoint with all route tables in the VPC

upvoted 2 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: C

Prevent traffic from traversing the internet = Gateway VPC endpoint for S3.

upvoted 4 times

Anmol_1010 2 years, 6 months ago

Key word transversing to internet

upvoted 2 times

Efren 2 years, 6 months ago

Selected Answer: C

Gateway endpoint for S3

upvoted 3 times

nosense 2 years, 6 months ago

Selected Answer: C

vpc endpoint for s3

upvoted 5 times

A company has a mobile chat application with a data store based in Amazon DynamoDB. Users would like new messages to be read with as little latency as possible. A solutions architect needs to design an optimal solution that requires minimal application changes.

Which method should the solutions architect select?

A. Configure Amazon DynamoDB Accelerator (DAX) for the new messages table. Update the code to use the DAX endpoint.

Most Voted

B. Add DynamoDB read replicas to handle the increased read load. Update the application to point to the read endpoint for the read replicas.

C. Double the number of read capacity units for the new messages table in DynamoDB. Continue to use the existing DynamoDB endpoint.

D. Add an Amazon ElastiCache for Redis cache to the application stack. Update the application to point to the Redis cache endpoint instead of DynamoDB.

Correct Answer: A

Community vote distribution

A (95%)

5%

Comments

pentium75 1 year, 11 months ago

Selected Answer: A

B and C do not reduce latency. D would reduce latency but require significant application changes.

upvoted 2 times

wizcloudifa 1 year, 7 months ago

D would not reduce latency, all the messages are new, they won't be stored in cache and they are unique messages(dynamic content) so ElastiCache would be pointless here, it's more preferable for static content/frequently accessed content.... A makes perfect sense

upvoted 3 times

Cyberkayu 1 year, 11 months ago

Selected Answer: C

0 code change @C

ABD. In memory cache, read replica, elasticache. Chat application and content is dynamic, cache will still pull data from prod database

upvoted 1 times

pentium75 1 year, 11 months ago

C has 0 code changes but doesn't address the issue.

upvoted 7 times

danielmakita 2 years, 1 month ago

Would go for A.

Minimal application changes != No application changes

upvoted 2 times

thanhnv142 2 years, 1 month ago

"requires minimal application changes" - Do not choose A because it requires updates of codes.

upvoted 1 times

thanhnv142 2 years, 1 month ago

C is correct
A, B and D all require code changes to the app.
upvoted 1 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: A

A. Configure Amazon DynamoDB Accelerator (DAX) for the new messages table. Update the code to use the DAX endpoint.
upvoted 2 times

haoAWS 2 years, 5 months ago

Selected Answer: A

Read replica does improve the read speed, but it cannot improve the latency because there is always latency between replicas. So A works and B not work.
upvoted 2 times

mattcl 2 years, 5 months ago

C , "requires minimal application changes"
upvoted 1 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: A

little latency = Amazon DynamoDB Accelerator (DAX) .
upvoted 4 times

DrWatson 2 years, 6 months ago

Selected Answer: A

I go with A <https://aws.amazon.com/blogs/mobile/building-a-full-stack-chat-application-with-aws-and-nextjs/> but I have some doubts about this <https://aws.amazon.com/blogs/database/how-to-build-a-chat-application-with-amazon-elasticache-for-redis/>
upvoted 3 times

cloudenthusiast 2 years, 6 months ago

Selected Answer: A

Amazon DynamoDB Accelerator (DAX): DAX is an in-memory cache for DynamoDB that provides low-latency access to frequently accessed data. By configuring DAX for the new messages table, read requests for the table will be served from the DAX cache, significantly reducing the latency.

Minimal Application Changes: With DAX, the application code can be updated to use the DAX endpoint instead of the standard DynamoDB endpoint. This change is relatively minimal and does not require extensive modifications to the application's data access logic.

Low Latency: DAX caches frequently accessed data in memory, allowing subsequent read requests for the same data to be served with minimal latency. This ensures that new messages can be read by users with minimal delay.
upvoted 4 times

cloudenthusiast 2 years, 6 months ago

Option B (Add DynamoDB read replicas) involves creating read replicas to handle the increased read load, but it may not directly address the requirement of minimizing latency for new message reads.
upvoted 1 times

Efren 2 years, 6 months ago

Tricky one, in doubt also with B, read replicas.
upvoted 1 times

awsgeek75 1 year, 11 months ago

Yes it's tricky but least code changes is the tie breaker. DAX has zero code changes.
upvoted 2 times

nosense 2 years, 6 months ago

Selected Answer: A

a is valid
upvoted 3 times

A company hosts a website on Amazon EC2 instances behind an Application Load Balancer (ALB). The website serves static content. Website traffic is increasing, and the company is concerned about a potential increase in cost.

- A. Create an Amazon CloudFront distribution to cache static files at edge locations **Most Voted**
- B. Create an Amazon ElastiCache cluster. Connect the ALB to the ElastiCache cluster to serve cached files
- C. Create an AWS WAF web ACL and associate it with the ALB. Add a rule to the web ACL to cache static files
- D. Create a second ALB in an alternative AWS Region. Route user traffic to the closest Region to minimize data transfer costs

Correct Answer: A

Community vote distribution

A (100%)

Comments

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: A

Amazon CloudFront: CloudFront is a content delivery network (CDN) service that caches content at edge locations worldwide. By creating a CloudFront distribution, static content from the website can be cached at edge locations, reducing the load on the EC2 instances and improving the overall performance.

Caching Static Files: Since the website serves static content, caching these files at CloudFront edge locations can significantly reduce the number of requests forwarded to the EC2 instances. This helps to lower the overall cost by offloading traffic from the instances and reducing the data transfer costs.

upvoted 6 times

awsgeek75 **Highly Voted** 1 year, 11 months ago

Selected Answer: A

The problem with this question is that no sane AWS architect will choose any of these options and go for S3 caching. But given the choices, A is the only one which will solve the problem within reasonable cost.

upvoted 5 times

Guru4Cloud **Most Recent** 2 years, 3 months ago

Selected Answer: A

A. Create an Amazon CloudFront distribution to cache static files at edge locations

upvoted 3 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: A

Serves static content = Amazon CloudFront distribution.

upvoted 3 times

nosense 2 years, 6 months ago

Selected Answer: A

a for me

upvoted 3 times

A company has multiple VPCs across AWS Regions to support and run workloads that are isolated from workloads in other Regions. Because of a recent application launch requirement, the company's VPCs must communicate with all other VPCs across all Regions.

Which solution will meet these requirements with the LEAST amount of administrative effort?

- A. Use VPC peering to manage VPC communication in a single Region. Use VPC peering across Regions to manage VPC communications.
- B. Use AWS Direct Connect gateways across all Regions to connect VPCs across regions and manage VPC communications.
- C. Use AWS Transit Gateway to manage VPC communication in a single Region and Transit Gateway peering across Regions to manage VPC communications. **Most Voted**
- D. Use AWS PrivateLink across all Regions to connect VPCs across Regions and manage VPC communications

Correct Answer: C

Community vote distribution

C (100%)

Comments

Felix_br **Highly Voted** 2 years, 6 months ago

The correct answer is: C. Use AWS Transit Gateway to manage VPC communication in a single Region and Transit Gateway peering across Regions to manage VPC communications.

AWS Transit Gateway is a network hub that you can use to connect your VPCs and on-premises networks. It provides a single point of control for managing your network traffic, and it can help you to reduce the number of connections that you need to manage.

Transit Gateway peering allows you to connect two Transit Gateways in different Regions. This can help you to create a global network that spans multiple Regions.

To use Transit Gateway to manage VPC communication in a single Region, you would create a Transit Gateway in each Region. You would then attach your VPCs to the Transit Gateway.

To use Transit Gateway peering to manage VPC communication across Regions, you would create a Transit Gateway peering connection between the Transit Gateways in each Region.

upvoted 25 times

TariqKipkemei 2 years, 5 months ago

thank you for this comprehensive explanation

upvoted 3 times

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: C

AWS Transit Gateway: Transit Gateway is a highly scalable service that simplifies network connectivity between VPCs and on-premises networks. By using a Transit Gateway in a single Region, you can centralize VPC communication management and reduce administrative effort.

Transit Gateway Peering: Transit Gateway supports peering connections across AWS Regions, allowing you to establish connectivity between VPCs in different Regions without the need for complex VPC peering configurations. This simplifies the management of VPC communications across Regions

upvoted 8 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: C

C is like a managed solution for A. A can work but with a lot of overhead (CIDR blocks uniqueness requirement). B and D are not the right products

upvoted 2 times

potomac 2 years, 1 month ago

Selected Answer: C

multiple regions + multiple VPCs --> Transit Gateway
upvoted 3 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: C

Definitely C.
Very well explained by @Felix_br
upvoted 2 times

omoakin 2 years, 6 months ago

Ccccccccccccccccccc
if you have services in multiple Regions, a Transit Gateway will allow you to access those services with a simpler network configuration.
upvoted 3 times

A company is designing a containerized application that will use Amazon Elastic Container Service (Amazon ECS). The application needs to access a shared file system that is highly durable and can recover data to another AWS Region with a recovery point objective (RPO) of 8 hours. The file system needs to provide a mount target in each Availability Zone within a Region.


A solutions architect wants to use AWS Backup to manage the replication to another Region.

Which solution will meet these requirements?

- A. Amazon FSx for Windows File Server with a Multi-AZ deployment
- B. Amazon FSx for NetApp ONTAP with a Multi-AZ deployment
- C. Amazon Elastic File System (Amazon EFS) with the Standard storage class **Most Voted**
- D. Amazon FSx for OpenZFS

Correct Answer: C

Community vote distribution



Comments

elmogy **Highly Voted** 2 years, 6 months ago

Selected Answer: C

<https://aws.amazon.com/efs/faq/>

Q: What is Amazon EFS Replication?

EFS Replication can replicate your file system data to another Region or within the same Region without requiring additional infrastructure or a custom process. Amazon EFS Replication automatically and transparently replicates your data to a second file system in a Region or AZ of your choice. You can use the Amazon EFS console, AWS CLI, and APIs to activate replication on an existing file system. EFS Replication is continual and provides a recovery point objective (RPO) and a recovery time objective (RTO) of minutes, helping you meet your compliance and business continuity goals.

upvoted 13 times

wizcloudifa **Most Recent** 1 year, 7 months ago

Selected Answer: C

The key thing to notice here in the question is "with a recovery point objective (RPO) of 8 hours", as it is 8 hours of time it recovery can be easily managed by EFS, no need to go for costlier and not built for this use-case(share file system) options like NetApp ONTAP(proprietary data cluster), OpenZFS(not a built in filesystem in AWS) or FSx for windows(file system for windows compatible workloads)

upvoted 3 times

awsgeek75 1 year, 11 months ago

Selected Answer: C

A: ECS is not Windows File Server so won't work

B: ONTAP is proprietary data cluster completely unrelated to this question

D: OpenZFS needs a Linux kind of host for access. Not a built-in filesystem in AWS by default

upvoted 2 times

pentium75 1 year, 11 months ago

Selected Answer: C

"The file system <https://www.examtopycs.com/exams/amazon/aws-certified-solutions-architect-associate-saa-c03/view/48/#> needs to provide a mount target in each (!) Availability Zone within a Region", most regions have three AZs, but FSx Multi-AZ provides only nodes "spread across two AZs". While "or Amazon EFS file systems that use Regional storage classes [such as Standard], you can create a mount target in each Availability Zone in an AWS Region."

upvoted 3 times

pentium75 1 year, 11 months ago

Huh, comment has been scrambled a bit. Anyway

FSx Multi-AZ: Mount targets in two AZs
EFS Standard: Can create mount target in each AZ
upvoted 3 times

Goutham4981 2 years ago

Selected Answer: C

In the absence of this information, we can only make an assumption based on the provided requirements. The requirement for a shared file system that can recover data to another AWS Region with a recovery point objective (RPO) of 8 hours, and the need for a mount target in each Availability Zone within a Region, are all natively supported by Amazon EFS with the Standard storage class.

While Amazon FSx for NetApp ONTAP does provide shared file systems and supports both Windows and Linux, it does not natively support replication to another region through AWS Backup.
upvoted 3 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: C

C. Amazon Elastic File System (Amazon EFS) with the Standard storage class
upvoted 2 times

cd93 2 years, 3 months ago

Selected Answer: B

B or C, but since question didn't mention operating system type, I guess we should go with B because it is more versatile (EFS supports Linux only), although ECS containers do support windows instances...
upvoted 2 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: C

Both option B and C will support this requirement.

<https://aws.amazon.com/efs/faq/#:~:text=What%20is%20Amazon%20EFS%20Replication%3F>

<https://aws.amazon.com/fsx/netapp-ontap/faqs/#:~:text=How%20do%20I%20configure%20cross%2Dregion%20replication%20for%20the%20data%20in%20my%20file%20system%3F>
upvoted 2 times

omoakin 2 years, 6 months ago

BBBBBBBBBBBBBBBB
upvoted 1 times

Anmol_1010 2 years, 6 months ago

Efs, can be mounted only in 1 region
So the answer is B
upvoted 3 times

Rob1L 2 years, 6 months ago

Selected Answer: C

C: EFS
upvoted 3 times

y0 2 years, 6 months ago

Selected Answer: C

AWS Backup can manage replication of EFS to another region as mentioned below
<https://docs.aws.amazon.com/efs/latest/ug/awsbackup.html>
upvoted 2 times

norris81 2 years, 6 months ago

<https://aws.amazon.com/efs/faq/>

During a disaster or fault within an AZ affecting all copies of your data, you might experience loss of data that has not been replicated using Amazon EFS Replication. EFS Replication is designed to meet a recovery point objective (RPO) and recovery time objective (RTO) of minutes. You can use AWS Backup to store additional copies of your file system data and restore them to a new file system in an AZ or Region of your choice. Amazon EFS file system backup data created and managed by AWS Backup is replicated to three AZs and is designed for 99.99999999% (11 nines) durability.
upvoted 2 times

nosense 2 years, 6 months ago

Amazon EFS is a scalable and durable elastic file system that can be used with Amazon ECS. However, it does not support replication to another AWS Region.

upvoted 1 times

fakrap 2 years, 6 months ago

To use EFS replication in a Region that is disabled by default, you must first opt in to the Region, so it does support.

upvoted 2 times

elmogy 2 years, 6 months ago

it does support replication to another AWS Region

check the same link you are replying to :/

<https://aws.amazon.com/efs/faq/>

Q: What is Amazon EFS Replication?

EFS Replication can replicate your file system data to another Region or within the same Region without requiring additional infrastructure or a custom process. Amazon EFS Replication automatically and transparently replicates your data to a second file system in a Region or AZ of your choice. You can use the Amazon EFS console, AWS CLI, and APIs to activate replication on an existing file system. EFS Replication is continual and provides a recovery point objective (RPO) and a recovery time objective (RTO) of minutes, helping you meet your compliance and business continuity goals.

upvoted 2 times

nosense 2 years, 6 months ago

Selected Answer: B

shared file system that is highly durable and can recover data

upvoted 2 times

Efren 2 years, 6 months ago

Why not EFS?

upvoted 2 times

A company is expecting rapid growth in the near future. A solutions architect needs to configure existing users and grant permissions to new users on AWS. The solutions architect has decided to create IAM groups. The solutions architect will add the new users to IAM groups based on department.

Which additional action is the MOST secure way to grant permissions to the new users?

- A. Apply service control policies (SCPs) to manage access permissions
- B. Create IAM roles that have least privilege permission. Attach the roles to the IAM groups
- C. Create an IAM policy that grants least privilege permission. Attach the policy to the IAM groups **Most Voted**
- D. Create IAM roles. Associate the roles with a permissions boundary that defines the maximum permissions

Correct Answer: C

Community vote distribution

C (95%)

5%

Comments

Rob1L **Highly Voted** 3 years ago

Selected Answer: C

Option B is incorrect because IAM roles are not directly attached to IAM groups.

upvoted 10 times

RoroJ 3 years ago

IAM Roles can be attached to IAM Groups:

https://docs.aws.amazon.com/directoryservice/latest/admin-guide/assign_role.html

upvoted 4 times

antropaws 3 years ago

Read your own link: You can assign an existing IAM role to an AWS Directory Service user or group. Not to IAM groups.

upvoted 11 times

Efren **Highly Voted** 3 years ago

Selected Answer: C

Agreed with C

https://docs.aws.amazon.com/IAM/latest/UserGuide/id_groups_manage_attach-policy.html

Attaching a policy to an IAM user group

upvoted 7 times

MatAlves **Most Recent** 1 year, 9 months ago

Selected Answer: C

"Manage access in AWS by creating policies and attaching them to IAM identities (users, groups of users, or roles) or AWS resources."

"An IAM role is an identity within your AWS account that has specific permissions. It's similar to an IAM user, but isn't associated with a specific person."

"IAM roles do not have any permanent credentials associated with them and are instead assumed by IAM users, AWS services, or applications that need temporary security credentials to access AWS resources"

upvoted 2 times

MatAlves 1 year, 9 months ago

https://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies.html

<https://docs.aws.amazon.com/IAM/latest/UserGuide/id.html>

<https://blog.awsfundamentals.com/aws-iam-roles-terms-concepts-and-examples>

upvoted 2 times

zinabu 2 years, 2 months ago

create role=for resource like EC2 and lambda

create a Policy =for groups or user access policy for the resources like S3 bucket

upvoted 4 times

pentium75 2 years, 5 months ago

Selected Answer: C

Not A or D because this is not about restricting maximum permissions, it is about securely granting permissions

Not B because IAM roles are not attached to IAM groups.

C because IAM policies are attached to IAM groups.

upvoted 5 times

potomac 2 years, 7 months ago

Selected Answer: C

A is wrong

SCPs are mainly used along with AWS Organizations organizational units (OUs). SCPs do not replace IAM Policies such that they do not provide actual permissions. To perform an action, you would still need to grant appropriate IAM Policy permissions.

upvoted 3 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: C

Create an IAM policy that grants least privilege permission. Attach the policy to the IAM groups

upvoted 2 times

TariqKipkemei 2 years, 11 months ago

Selected Answer: C

An IAM policy is an object in AWS that, when associated with an identity or resource, defines their permissions. Permissions in the policies determine whether a request is allowed or denied. You manage access in AWS by creating policies and attaching them to IAM identities (users, groups of users, or roles) or AWS resources.

So, option B will also work.

But Since I can only choose one, C would be it.

upvoted 3 times

MrAWSAssociate 2 years, 11 months ago

Selected Answer: C

You can attach up to 10 IAM policy for a 'user group'.

upvoted 2 times

antropaws 3 years ago

Selected Answer: C

C is the correct one.

upvoted 2 times

nosense 3 years ago

Selected Answer: B

should be b

upvoted 2 times

imazsyed 3 years ago

it should be C

upvoted 4 times

nosense 3 years ago

Option C is not as secure as option B because IAM policies are attached to individual users and cannot be used to manage permissions for groups of users.

upvoted 2 times

omoakin 3 years ago

IAM Roles manage who has access to your AWS resources, whereas IAM policies control their permissions. A Role with no Policy attached to it won't have to access any AWS resources. A Policy that is not attached to an IAM role is effectively unused.

upvoted 5 times

Clouddon 2 years, 9 months ago

https://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies.html

upvoted 2 times

pentium75 2 years, 5 months ago

IAM roles are not attached to IAM groups.

IAM policies are attached to IAM roles, IAM groups or IAM users. IAM roles are used by services.

upvoted 2 times

A group requires permissions to list an Amazon S3 bucket and delete objects from that bucket. An administrator has created the following IAM policy to provide access to the bucket and applied that policy to the group. The group is not able to delete objects in the bucket. The company follows least-privilege access rules.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Action": [
        "s3:ListBucket",
        "s3:DeleteObject"
      ],
      "Resource": [
        "arn:aws:s3:::bucket-name"
      ],
      "Effect": "Allow"
    }
  ]
}
```

Which statement should a solutions architect add to the policy to correct bucket access?

A.

```
"Action": [
  "s3:*Object"
],
"Resource": [
  "arn:aws:s3:::bucket-name/*"
],
"Effect": "Allow"
```

B.

```
"Action": [
  "s3:*"
],
"Resource": [
  "arn:aws:s3:::bucket-name/*"
],
"Effect": "Allow"
```

C.

```
"Action": [
  "s3:DeleteObject"
],
"Resource": [
  "arn:aws:s3:::bucket-name*"
],
"Effect": "Allow"
```

D.

```
"Action": [
  "s3:DeleteObject"
],
"Resource": [
  "arn:aws:s3:::bucket-name/*"
],
"Effect": "Allow"
```

Most Voted

Correct Answer: D

Community vote distribution

Comments

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: D

option B action is S3:*. this means all actions. The company follows least-privilege access rules. Hence option D
upvoted 6 times

nosense **Highly Voted** 2 years, 6 months ago

Selected Answer: D

d work
upvoted 5 times

Efren 2 years, 6 months ago

Agreed
upvoted 2 times

TariqKipkemei **Most Recent** 2 years, 5 months ago

Selected Answer: D

D is the answer
upvoted 4 times

AncaZalog 2 years, 5 months ago

what's the difference between B and D? on B the statements are just placed in another order
upvoted 2 times

TariqKipkemei 2 years, 5 months ago

option B action is S3:*. this means all actions. The company follows least-privilege access rules. Hence option D
upvoted 2 times

serepetru 2 years, 6 months ago

What is the difference between C and D?
upvoted 2 times

Ta_Les 2 years, 6 months ago

the "/" at the end of the last line on D
upvoted 6 times

sheilawu 1 year, 6 months ago

so annoying when you oversee this"/"
upvoted 1 times

Rob1L 2 years, 6 months ago

Selected Answer: D

D for sure
upvoted 2 times

A law firm needs to share information with the public. The information includes hundreds of files that must be publicly readable. Modifications or deletions of the files by anyone before a designated future date are prohibited.

Which solution will meet these requirements in the MOST secure way?

- A. Upload all files to an Amazon S3 bucket that is configured for static website hosting. Grant read-only IAM permissions to any AWS principals that access the S3 bucket until the designated date.
- B. Create a new Amazon S3 bucket with S3 Versioning enabled. Use S3 Object Lock with a retention period in accordance with the designated date. Configure the S3 bucket for static website hosting. Set an S3 bucket policy to allow read-only access to the objects. **Most Voted**
- C. Create a new Amazon S3 bucket with S3 Versioning enabled. Configure an event trigger to run an AWS Lambda function in case of object modification or deletion. Configure the Lambda function to replace the objects with the original versions from a private S3 bucket.
- D. Upload all files to an Amazon S3 bucket that is configured for static website hosting. Select the folder that contains the files. Use S3 Object Lock with a retention period in accordance with the designated date. Grant read-only IAM permissions to any AWS principals that access the S3 bucket.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Lin878 **Highly Voted** 1 year, 11 months ago

Selected Answer: B

Object Lock works only in buckets that have S3 Versioning enabled.
<https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock.html>
upvoted 6 times

dydzah **Highly Voted** 3 years ago

Selected Answer: B

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock.html>
upvoted 5 times

eduvaldi **Most Recent** 1 year ago

Selected Answer: B

Versioning and Object Lock is necessary
upvoted 1 times

MatAlves 1 year, 9 months ago

Selected Answer: B

Versioning Enabled + Object Lock = B
upvoted 2 times

potomac 2 years, 7 months ago

Selected Answer: B

S3 bucket policy
upvoted 4 times

thanhnv142 2 years, 7 months ago

B is correct.
A doesnot have S3 object lock, but deletion is prohibited, which implies object lock

C does not have S3 as static web, but have to share the s3 with the public
D mentions files - but S3 manages objects, not file

upvoted 2 times

hydro143 2 years, 8 months ago

D?

Its like B, but also with read-only access limitations for anyone with IAM permissions. Also versioning in B doesn't help with anything.

upvoted 2 times

ManikRoy 2 years, 1 month ago

Enabling versioning is a pre-requisite for object lock.

upvoted 2 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: B

Create a new Amazon S3 bucket with S3 Versioning enabled. Use S3 Object Lock with a retention period in accordance with the designated date. Configure the S3 bucket for static website hosting. Set an S3 bucket policy to allow read-only access to the objects.

upvoted 3 times

TariqKipkemei 2 years, 11 months ago

Selected Answer: B

Create a new Amazon S3 bucket with S3 Versioning enabled. Use S3 Object Lock with a retention period in accordance with the designated date. Configure the S3 bucket for static website hosting. Set an S3 bucket policy to allow read-only access to the objects.

upvoted 4 times

antropaws 3 years ago

Selected Answer: B

Clearly B.

upvoted 3 times

nosense 3 years ago

Selected Answer: B

Option A allows the files to be modified or deleted by anyone with read-only IAM permissions. Option C allows the files to be modified or deleted by anyone who can trigger the AWS Lambda function.

Option D allows the files to be modified or deleted by anyone with read-only IAM permissions to the S3 bucket

upvoted 5 times

wizcloudifa 2 years, 1 month ago

no it doesnt, did you not notice this part: "S3 Object Lock with a retention period in accordance with the designated date", this part avoids deletion/modification of files

upvoted 2 times

A company is making a prototype of the infrastructure for its new website by manually provisioning the necessary infrastructure. This infrastructure includes an Auto Scaling group, an Application Load Balancer and an Amazon RDS database. After the configuration has been thoroughly validated, the company wants the capability to immediately deploy the infrastructure for development and production use in two Availability Zones in an automated fashion.

What should a solutions architect recommend to meet these requirements?

- A. Use AWS Systems Manager to replicate and provision the prototype infrastructure in two Availability Zones
- B. Define the infrastructure as a template by using the prototype infrastructure as a guide. Deploy the infrastructure with AWS CloudFormation. **Most Voted**
- C. Use AWS Config to record the inventory of resources that are used in the prototype infrastructure. Use AWS Config to deploy the prototype infrastructure into two Availability Zones.
- D. Use AWS Elastic Beanstalk and configure it to use an automated reference to the prototype infrastructure to automatically deploy new environments in two Availability Zones.

Correct Answer: B

Community vote distribution

B (81%)

A (19%)

Comments

Guru4Cloud **Highly Voted** 2 years, 9 months ago

Just Think Infrastructure as Code == Cloud Formation
upvoted 9 times

MatAlves **Highly Voted** 1 year, 9 months ago

Selected Answer: A

The difference between CloudFormation and Beanstalk might be trick, but just for the exam think:

Cloudformation -> Infra as Code
Beanstalk -> deploy and manage applications
upvoted 6 times

LeonSauveterre 1 year, 6 months ago

That's right. And I think you've voted for A by mistake. I think you wanna vote for B ;)
upvoted 1 times

awsgeek75 **Most Recent** 2 years, 5 months ago

Selected Answer: B

A: Wrong product
C: Wrong product
D: EBS can only handle EC2 so RDS won't be replicated automatically
B: CloudFormation = IaC
upvoted 4 times

capino 2 years, 9 months ago

Selected Answer: B

Just Think Infrastructure as Code == Cloud Formation
upvoted 5 times

haoAWS 2 years, 11 months ago

Why D is not correct?
upvoted 2 times

Kiki_Pass 2 years, 10 months ago

I guess it's because Beanstalk is PaaS (platform as a service) while CloudFormation is IaC (infrastructure as code). The question emphasis more on infrastructure

upvoted 3 times

wRhIH 2 years, 11 months ago

I guess "TEMPLATE" leads to CloudFormation

upvoted 3 times

TariqKipkemei 2 years, 11 months ago

Selected Answer: B

Infrastructure as code = AWS CloudFormation

upvoted 5 times

antropaws 3 years ago

Selected Answer: B

Clearly B.

upvoted 3 times

Felix_br 3 years ago

Selected Answer: B

AWS CloudFormation is a service that allows you to define and provision infrastructure as code. This means that you can create a template that describes the resources you want to create, and then use CloudFormation to deploy those resources in an automated fashion.

In this case, the solutions architect should define the infrastructure as a template by using the prototype infrastructure as a guide. The template should include resources for an Auto Scaling group, an Application Load Balancer, and an Amazon RDS database. Once the template is created, the solutions architect can use CloudFormation to deploy the infrastructure in two Availability Zones.

upvoted 4 times

omoakin 3 years ago

B

Define the infrastructure as a template by using the prototype infrastructure as a guide. Deploy the infrastructure with AWS CloudFormation

upvoted 2 times

nosense 3 years ago

Selected Answer: B

b obvious

upvoted 5 times

A business application is hosted on Amazon EC2 and uses Amazon S3 for encrypted object storage. The chief information security officer has directed that no application traffic between the two services should traverse the public internet.

Which capability should the solutions architect use to meet the compliance requirements?

A. AWS Key Management Service (AWS KMS)

B. VPC endpoint **Most Voted**

C. Private subnet

D. Virtual private gateway

Correct Answer: B

Community vote distribution

B (100%)

Comments

cloudenthusiast **Highly Voted** 2 years ago

Selected Answer: B

A VPC endpoint enables you to privately access AWS services without requiring internet gateways, NAT gateways, VPN connections, or AWS Direct Connect connections. It allows you to connect your VPC directly to supported AWS services, such as Amazon S3, over a private connection within the AWS network.

By creating a VPC endpoint for Amazon S3, the traffic between your EC2 instances and S3 will stay within the AWS network and won't traverse the public internet. This provides a more secure and compliant solution, as the data transfer remains within the private network boundaries.

upvoted 10 times

TariqKipkemei **Highly Voted** 1 year, 11 months ago

Selected Answer: B

Prevent traffic from traversing the internet = VPC endpoint for S3.

upvoted 5 times

Tendu **Most Recent** 1 year, 3 months ago

Selected Answer: B

Answer B -VPC Endpoint

upvoted 1 times

antropaws 2 years ago

Selected Answer: B

B until proven contrary.

upvoted 3 times

handsonlabsaws 2 years ago

Selected Answer: B

B for sure

upvoted 3 times

Blingy 2 years ago

BBBBBBBBB

upvoted 2 times

A company hosts a three-tier web application in the AWS Cloud. A Multi-AZ Amazon RDS for MySQL server forms the database layer. Amazon ElastiCache forms the cache layer. The company wants a caching strategy that adds or updates data in the cache when a customer adds an item to the database. The data in the cache must always match the data in the database.

Which solution will meet these requirements?

- A. Implement the lazy loading caching strategy
- B. Implement the write-through caching strategy **Most Voted**
- C. Implement the adding TTL caching strategy
- D. Implement the AWS AppConfig caching strategy

Correct Answer: B

Community vote distribution

B (100%)

Comments

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: B

In the write-through caching strategy, when a customer adds or updates an item in the database, the application first writes the data to the database and then updates the cache with the same data. This ensures that the cache is always synchronized with the database, as every write operation triggers an update to the cache.

upvoted 31 times

cloudenthusiast 2 years, 6 months ago

Lazy loading caching strategy (option A) typically involves populating the cache only when data is requested, and it does not guarantee that the data in the cache always matches the data in the database.

Adding TTL (Time-to-Live) caching strategy (option C) involves setting an expiration time for cached data. It is useful for scenarios where the data can be considered valid for a specific period, but it does not guarantee that the data in the cache is always in sync with the database.

AWS AppConfig caching strategy (option D) is a service that helps you deploy and manage application configurations. It is not specifically designed for caching data synchronization between a database and cache layer.

upvoted 38 times

Kp88 2 years, 4 months ago

Great explanation, thanks

upvoted 3 times

zinabu **Most Recent** 1 year, 8 months ago

write-through caching strategy

upvoted 2 times

awsgeek75 1 year, 10 months ago

Selected Answer: B

More helpful reading for why B is the answer:

<https://docs.aws.amazon.com/AmazonElastiCache/latest/mem-ug/Strategies.html#Strategies.WriteThrough>

upvoted 4 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: B

B. Implement the write-through caching strategy

upvoted 4 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: B

The answer is definitely B.

I couldn't provide any more details than what has been shared by @cloudenthusiast.

upvoted 2 times

nosense 2 years, 6 months ago

Selected Answer: B

write-through caching strategy updates the cache at the same time as the database

upvoted 3 times

A company wants to migrate 100 GB of historical data from an on-premises location to an Amazon S3 bucket. The company has a 100 megabits per second (Mbps) internet connection on premises. The company needs to encrypt the data in transit to the S3 bucket. The company will store new data directly in Amazon S3.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use the s3 sync command in the AWS CLI to move the data directly to an S3 bucket
- B. Use AWS DataSync to migrate the data from the on-premises location to an S3 bucket **Most Voted**
- C. Use AWS Snowball to move the data to an S3 bucket
- D. Set up an IPsec VPN from the on-premises location to AWS. Use the s3 cp command in the AWS CLI to move the data directly to an S3 bucket

Correct Answer: B

Community vote distribution

B (62%)

A (38%)

Comments

cloudehustia **Highly Voted** 3 years ago

Selected Answer: B

AWS DataSync is a fully managed data transfer service that simplifies and automates the process of moving data between on-premises storage and Amazon S3. It provides secure and efficient data transfer with built-in encryption, ensuring that the data is encrypted in transit.

By using AWS DataSync, the company can easily migrate the 100 GB of historical data from their on-premises location to an S3 bucket. DataSync will handle the encryption of data in transit and ensure secure transfer.

upvoted 13 times

luiscc **Highly Voted** 3 years ago

Selected Answer: B

Using DataSync, the company can easily migrate the 100 GB of historical data to an S3 bucket. DataSync will handle the encryption of data in transit, so the company does not need to set up a VPN or worry about managing encryption keys.

Option A, using the s3 sync command in the AWS CLI to move the data directly to an S3 bucket, would require more operational overhead as the company would need to manage the encryption of data in transit themselves. Option D, setting up an IPsec VPN from the on-premises location to AWS, would also require more operational overhead and would be overkill for this scenario. Option C, using AWS Snowball, could work but would require more time and resources to order and set up the physical device.

upvoted 8 times

JustEugen **Most Recent** 8 months, 3 weeks ago

Selected Answer: B

I think the question is unfair and has a significant grey area. I have personally selected B to use DataSync over S3, despite DataSync having a larger operational overhead.

I selected DataSync because S3 sync uses just HTTPS on the application layer, which is primarily not the best option for migrating the data on premises, while DataSync uses Transport Layer and enforces the usage of TLS 1.2.

Additionally to the security question, we have 100 GB of data and only 100 Mbps, which would require 2 - 2.5 hours to complete a job, and in case of failure, we would have to manually handle errors and retries, while DataSync will do that for us

upvoted 1 times

In2718 9 months, 3 weeks ago

Selected Answer: A

Chat says A -

-The AWS CLI uses HTTPS (TLS) by default, so data is encrypted in transit without extra setup.

-100 GB over a 100 Mbps link is only a few hours, so no need for Snowball or VPN.

-s3 sync is simple and low-overhead for a one-time migration and can be reused if you need to re-run to catch any changes.

-DataSync works too but requires deploying/maintaining an agent—unnecessary overhead for 100 GB.

-A VPN adds setup/ops complexity with no benefit here.

upvoted 1 times

In2718 10 months, 4 weeks ago

Selected Answer: A

In reading through the threads, option A seems to be the better choice based on comments regarding bandwidth and transfer times, S3 sync is the least overhead. Not a well written question, but have to consider reading into what they are trying to convey.

upvoted 1 times

Dz1990 1 year ago

Selected Answer: B

s3 sync Disadvantages:

Manual process = You handle errors, retries, progress monitoring

Basic tool = No optimization for large transfers

Network issues = Manual restart if connection drops

No built-in throttling = Could saturate bandwidth

Higher operational overhead = More hands-on management needed

upvoted 1 times

Manimgh 1 year, 1 month ago

Selected Answer: B

encrypted = DataSync

upvoted 1 times

bora4motion 1 year, 1 month ago

Selected Answer: B

I will not use cli (A) to send 100GB to s3 from web. This means I will have to stay by computer? or use a different ec2 to login into it and initiate cli from there? hell no!!

B Datasync is the elegant way.

upvoted 1 times

dollykidurian 1 year, 2 months ago

My counter to all those, picking choice A, you will have to manually re-run the command if on-premise internet goes down however with datasync the service auto resume. (less operational overhead and less monitoring)

upvoted 1 times

zdi561 1 year, 4 months ago

Selected Answer: B

You can also copy or synchronize data from on-premises to AWS manually with the AWS CLI, but doing so requires a fair bit of system tuning see <https://aws.amazon.com/blogs/storage/migrating-and-managing-large-datasets-on-amazon-s3/>

upvoted 1 times

FlyingHawk 1 year, 6 months ago

Selected Answer: A

For A, it is more straightforward approach, with 100Mbps = 12.5MB/s, 100 GB = 100 * 1024 MB = 102,400 MB, 102400 / 12.5 = 8192s = 2.3 hours, so if the network is stable, it should be less operational overhead than B which requires the installation of agent and network configuration, so compared to S3 sync, DataSync requires more upfront configuration but provides more robust, automated transfer capabilities. In an AWS certification exam, they typically expect you to choose the solution that provides the most straightforward, efficient approach with the least operational overhead while meeting all specified requirements.

For this scenario, they would likely expect you to choose A.

upvoted 1 times

FlyingHawk 1 year, 6 months ago

In the context of AWS exam, AWS may expect we choose B as AWS DataSync is a managed service designed specifically for data migration tasks, A is more ad-hoc approach suit for small data-set. 100 GB will take 2.22 hours, although it is acceptable.

upvoted 2 times

LeonSauveterre 1 year, 6 months ago

Selected Answer: A

Option A meets all the requirements with the LEAST operational overhead (while option B will introduce unnecessary complexity for a one-time transfer). The data will be encrypted in transit (via HTTPS), and a 100 Mbps connection is sufficient to transfer data of 100 GB.

If the task involves periodic transfers or error recovery, DataSync (option B) could be considered as an alternative.

upvoted 2 times

1e22522 1 year, 10 months ago

Selected Answer: B

Why would u u se the CLI

upvoted 2 times

bujuman 2 years, 2 months ago

Selected Answer: A

Assertions:

- needs to encrypt the data in transit to the S3 bucket.
- The company will store new data directly in Amazon S3.

Requirements:

- with the LEAST operational overhead

Even Though options A and B could do the job, option A requires VM maintenance because it is not a once-off migration (The company will store new data directly in Amazon S3)

NB: According to me, we must stuck to the question and avoid to interpret

upvoted 3 times

bujuman 2 years, 2 months ago

Erratum:

Assertions:

- needs to encrypt the data in transit to the S3 bucket.
- The company will store new data directly in Amazon S3.

Requirements:

- with the LEAST operational overhead

Even Though options A and B could do the job, option B requires VM maintenance because it is not a once-off migration (The company will store new data directly in Amazon S3)

NB: According to me, we must stuck to the question and avoid to interpret

upvoted 1 times

pentium75 2 years, 5 months ago

Selected Answer: A

A - one single command, uses encryption automatically

B - Must install, configure and eventually decommission DataSync

C - Overkill

D - No need for VPN

upvoted 5 times

awsgeek75 2 years, 4 months ago

I agree, A is a million times simpler than B in terms of operational setup. AWS CLI is just one install on a server on client side and one command (literally) to sync the data.

upvoted 3 times

1rob 2 years, 5 months ago

Selected Answer: A

By default, all data transmitted from the client computer running the AWS CLI and AWS service endpoints is encrypted by sending everything through a HTTPS/TLS connection. You don't need to do anything to enable the use of HTTPS/TLS. It is always enabled unless you explicitly disable it for an individual command by using the --no-verify-ssl command line option.

This is simpler compared to datasync, which will cost operational overhead to configure.

upvoted 2 times

potomac 2 years, 7 months ago

Selected Answer: B

storage data (including metadata) is encrypted in transit, but how it's encrypted throughout the transfer depends on your source and destination locations.

upvoted 2 times

A company containerized a Windows job that runs on .NET 6 Framework under a Windows container. The company wants to run this job in the AWS Cloud. The job runs every 10 minutes. The job's runtime varies between 1 minute and 3 minutes.

Which solution will meet these requirements MOST cost-effectively?

- A. Create an AWS Lambda function based on the container image of the job. Configure Amazon EventBridge to invoke the function every 10 minutes.
- B. Use AWS Batch to create a job that uses AWS Fargate resources. Configure the job scheduling to run every 10 minutes.
- C. Use Amazon Elastic Container Service (Amazon ECS) on AWS Fargate to run the job. Create a scheduled task based on the container image of the job to run every 10 minutes. **Most Voted**
- D. Use Amazon Elastic Container Service (Amazon ECS) on AWS Fargate to run the job. Create a standalone task based on the container image of the job. Use Windows task scheduler to run the job every 10 minutes.

Correct Answer: C

Community vote distribution

C (51%)

B (29%)

A (20%)

Comments

baba365 **Highly Voted** 2 years, 8 months ago

Lambda supports only Linux-based container images.

<https://docs.aws.amazon.com/lambda/latest/dg/images-create.html>

upvoted 16 times

awsgeek75 2 years, 5 months ago

Not really. Lambda supports .Net 6 directly: <https://aws.amazon.com/blogs/compute/introducing-the-net-6-runtime-for-aws-lambda/>

upvoted 7 times

AmrFawzy93 **Highly Voted** 3 years ago

Selected Answer: C

By using Amazon ECS on AWS Fargate, you can run the job in a containerized environment while benefiting from the serverless nature of Fargate, where you only pay for the resources used during the job's execution. Creating a scheduled task based on the container image of the job ensures that it runs every 10 minutes, meeting the required schedule. This solution provides flexibility, scalability, and cost-effectiveness.

upvoted 8 times

bora4motion 1 year, 1 month ago

why not lambda?

upvoted 1 times

zdi561 **Most Recent** 1 year, 4 months ago

Selected Answer: A

Lambda is the best for cost, the job is kicked off every 10min and it only takes several min. Set 10min timeout, there is no cost for idling. Batch and ECS cost more for sure.

upvoted 1 times

zdi561 1 year, 4 months ago

Selected Answer: A

Lambda can be in .net container, and it is cheaper than any other options.

upvoted 1 times

LeonSauveterre 1 year, 6 months ago

Selected Answer: C

For B: Even though AWS Batch is free, I think it's more suitable for large-scale or relatively complex workloads, not lightweight periodic jobs. So B is not cost-effective (The cost depends on the compute resources you provision to execute the jobs, and maybe "cost-effective" somewhat involves less operational overhead because you have spent your money)

If you require queueing for task prioritization or processing a backlog of tasks, or work on data processing over datasets where Batch excels, AWS Batch would be better!

upvoted 1 times

muhammadahmer36 1 year, 10 months ago

A A A A A

upvoted 2 times

zinabu 2 years, 2 months ago

selected answer : A

AWS Lambda now supports .NET 6 as both a managed runtime and a container base image

upvoted 5 times

xBUGx 2 years, 2 months ago

Selected Answer: A

<https://aws.amazon.com/about-aws/whats-new/2022/02/aws-lambda-adds-support-net6/>

upvoted 3 times

awsgeek75 2 years, 5 months ago

The question is weirdly phrased for .Net based containers. "A company containerized a Windows job that runs on .NET 6 Framework under a Windows container." This could mean that the job requires .Net 6 Framework OR it could mean the job requires Windows and .Net Framework 6. If the job is just based on .Net 6 then Lambda can run it. I am just a bit cautious about language because other parameters fall under Lambda. Question may have been wrongly quoted here.

upvoted 4 times

pentium75 2 years, 5 months ago

Selected Answer: C

I guess this is an old question from before August 2023, when AWS Batch did not support Windows containers, while ECS already did since September 2021. Thus it would be C, though now B does also work. Since both Batch and ECS are free, we'd pay only for the Fargate resources (which are identical in both cases), now B and C would be correct.

A doesn't work because Lambda still does not support Windows containers.

D doesn't make sense because the container would have to run 24/7

upvoted 7 times

ftaws 2 years, 5 months ago

I think that Batch with Fargate is more cheaper than ECS.

upvoted 1 times

pentium75 2 years, 5 months ago

Both Batch and ECS are free.

<https://aws.amazon.com/de/ecs/pricing/>

<https://aws.amazon.com/de/batch/pricing/>

upvoted 2 times

kt7 2 years, 6 months ago

Selected Answer: B

Batch supports fargate now

upvoted 6 times

ccmc 2 years, 7 months ago

Selected Answer: B

aws batch supports fargate

upvoted 3 times

deechean 2 years, 9 months ago

Selected Answer: C

C works. For A, the lambda support container image, but the container image much implement the Lambda Runtime API.

upvoted 1 times

markoniz 2 years, 9 months ago

Absolutely agree with this one ... Lambda do not support Windows container, on the other hand ECS is adequate solution

upvoted 3 times

Hades2231 2 years, 9 months ago

Selected Answer: B

As they support Batch on Fargate now (Aug 2023), the correct answer should be B?
upvoted 3 times

RDM10 2 years, 8 months ago

that's exactly my question too.

In one of the discussions, they same lambda is for jobs for 15 min. But for other question, they same batch is the best. I do not understand why we cant use batch?

upvoted 1 times

Smart 2 years, 9 months ago

Selected Answer: A

<https://docs.aws.amazon.com/lambda/latest/dg/csharp-image.html#csharp-image-clients>
upvoted 1 times

pentium75 2 years, 5 months ago

But it's clearly "a Windows job". Lambda does not support Windows containers. (.NET 6 could also run under Linux, but they'd need to modify the container in any case.)

upvoted 2 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: C

C is the most cost-effective solution for running a short-lived Windows container job on a schedule.

Using Amazon ECS scheduled tasks on Fargate eliminates the need to provision EC2 resources. You pay only for the duration the task runs.

Scheduled tasks handle scheduling the jobs and scaling resources automatically. This is lower cost than managing your own scaling via Lambda or Batch.

ECS also supports Windows containers natively unlike Lambda (option A).

Option D still requires provisioning and paying for full time EC2 resources to run a task scheduler even when tasks are not running.

upvoted 2 times

A company wants to move from many standalone AWS accounts to a consolidated, multi-account architecture. The company plans to create many new AWS accounts for different business units. The company needs to authenticate access to these AWS accounts by using a centralized corporate directory service.

Which combination of actions should a solutions architect recommend to meet these requirements? (Choose two.)

A. Create a new organization in AWS Organizations with all features turned on. Create the new AWS accounts in the organization. **Most Voted**

B. Set up an Amazon Cognito identity pool. Configure AWS IAM Identity Center (AWS Single Sign-On) to accept Amazon Cognito authentication.

C. Configure a service control policy (SCP) to manage the AWS accounts. Add AWS IAM Identity Center (AWS Single Sign-On) to AWS Directory Service.

D. Create a new organization in AWS Organizations. Configure the organization's authentication mechanism to use AWS Directory Service directly.

E. Set up AWS IAM Identity Center (AWS Single Sign-On) in the organization. Configure IAM Identity Center, and integrate it with the company's corporate directory service. **Most Voted**

Correct Answer: AE

Community vote distribution

AE (100%)

Comments

cloudenthusiast **Highly Voted** 2 years ago

Selected Answer: AE

A. By creating a new organization in AWS Organizations, you can establish a consolidated multi-account architecture. This allows you to create and manage multiple AWS accounts for different business units under a single organization.

E. Setting up AWS IAM Identity Center (AWS Single Sign-On) within the organization enables you to integrate it with the company's corporate directory service. This integration allows for centralized authentication, where users can sign in using their corporate credentials and access the AWS accounts within the organization.

Together, these actions create a centralized, multi-account architecture that leverages AWS Organizations for account management and AWS IAM Identity Center (AWS Single Sign-On) for authentication and access control.

upvoted 11 times

Guru4Cloud **Most Recent** 1 year, 9 months ago

Selected Answer: AE

A) Using AWS Organizations allows centralized management of multiple AWS accounts in a single organization. New accounts can easily be created within the organization.

E) Integrating AWS IAM Identity Center (AWS SSO) with the company's corporate directory enables federated single sign-on. Users can log in once to access accounts and resources across AWS.

Together, Organizations and IAM Identity Center provide consolidated management and authentication for multiple accounts using existing corporate credentials.

upvoted 3 times

samehpalass 1 year, 11 months ago

Selected Answer: AE

A:AWS Organization

E:Authentication because option C (SCP) for Authorization

upvoted 4 times

baba365 1 year, 11 months ago

Ans: CD

'centralized corporate directory service' with new accounts in AWS Organizations
upvoted 2 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: AE

Create a new organization in AWS Organizations with all features turned on. Create the new AWS accounts in the organization. Set up AWS IAM Identity Center (AWS Single Sign-On) in the organization. Configure IAM Identity Center, and integrate it with the company's corporate directory service.

AWS IAM Identity Center (successor to AWS Single Sign-On) helps you securely create or connect your workforce identities and manage their access centrally across AWS accounts and applications.

[https://aws.amazon.com/iam/identity-center/#:~:text=AWS%20IAM%20Identity%20Center%20\(successor%20to%20AWS%20Single%20Sign%20On\)%20helps%20you%20securely%20create%20or%20connect%20your%20workforce%20identities%20and%20manage%20their%20access%20centrally%20across%20AWS%20accounts%20and%20applications.](https://aws.amazon.com/iam/identity-center/#:~:text=AWS%20IAM%20Identity%20Center%20(successor%20to%20AWS%20Single%20Sign%20On)%20helps%20you%20securely%20create%20or%20connect%20your%20workforce%20identities%20and%20manage%20their%20access%20centrally%20across%20AWS%20accounts%20and%20applications.)

upvoted 2 times

nosense 2 years ago

ae is right

upvoted 2 times

A company is looking for a solution that can store video archives in AWS from old news footage. The company needs to minimize costs and will rarely need to restore these files. When the files are needed, they must be available in a maximum of five minutes.

What is the MOST cost-effective solution?

- A. Store the video archives in Amazon S3 Glacier and use Expedited retrievals. **Most Voted**
- B. Store the video archives in Amazon S3 Glacier and use Standard retrievals.
- C. Store the video archives in Amazon S3 Standard-Infrequent Access (S3 Standard-IA).
- D. Store the video archives in Amazon S3 One Zone-Infrequent Access (S3 One Zone-IA).

Correct Answer: A

Community vote distribution

A (97%)

3%

Comments

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: A

By choosing Expedited retrievals in Amazon S3 Glacier, you can reduce the retrieval time to minutes, making it suitable for scenarios where quick access is required. Expedited retrievals come with a higher cost per retrieval compared to standard retrievals but provide faster access to your archived data.

upvoted 13 times

nosense **Highly Voted** 2 years, 6 months ago

glacier expedited retrieval times of typically 1-5 minutes.

upvoted 6 times

wsdasdasdqwda 2 years, 1 month ago

Fully agree. Check here for evidences: <https://aws.amazon.com/s3/storage-classes/glacier/#:~:text=S3%20Glacier%20Flexible%20Retrieval%20provides,amounts%20of%20data%20typically%20in>

upvoted 2 times

Ravan **Most Recent** 1 year, 9 months ago

Selected Answer: A

The most cost-effective solution that also meets the requirement of having the files available within a maximum of five minutes when needed is:

A. Store the video archives in Amazon S3 Glacier and use Expedited retrievals.

Amazon S3 Glacier is designed for long-term storage of data archives, providing a highly durable and secure solution at a low cost. With Expedited retrievals, data can be retrieved within a few minutes, which meets the requirement of having the files available within five minutes when needed. This option provides the balance between cost-effectiveness and retrieval speed, making it the best choice for the company's needs.

upvoted 2 times

pentium75 1 year, 11 months ago

Selected Answer: A

Occasional cost for retrieval from Glacier is nothing compared to the huge storage cost savings compared to C. Still meets the five minute requirement.

upvoted 2 times

master9 1 year, 11 months ago

Selected Answer: C

The retrieval price will play an important role here. I selected the "C" option because in "Glacier and use Expedited retrievals" its around \$0.004 per GB/month and for STD-IA \$0.0125 per GB/month
<https://www.cloudforecast.io/blog/aws-s3-pricing-and-optimization-guide/>

upvoted 1 times

pentium75 1 year, 11 months ago

But they "will rarely need to restore their files", thus the low cost for occasional expedited retrievals will be nothing compared to the huge storage cost savings.

upvoted 2 times

ngo01214 2 years, 2 months ago

s3 expedited can only be applied on glacier flexible retrieval storage class and s3 intelligent tiering archive access tier. so the answer should be C

upvoted 3 times

pentium75 1 year, 11 months ago

A mentions "G3 Glacier" which has been renamed to "S3 Glacier Flexible Retrieval" and meets the requirements.

upvoted 2 times

Smart 2 years, 3 months ago

Selected Answer: A

I am going with option A, but it is a poorly written question. "For all but the largest archives (more than 250 MB), data accessed by using Expedited retrievals is typically made available within 1–5 minutes. "

upvoted 2 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: A

Answer - A

Fast availability: Although retrieval times for objects stored in Amazon S3 Glacier typically range from minutes to hours, you can use the Expedited retrievals option to expedite access to your archives. By using Expedited retrievals, the files can be made available in a maximum of five minutes when needed. However, Expedited retrievals do incur higher costs compared to standard retrievals.

upvoted 2 times

hsinchang 2 years, 4 months ago

Selected Answer: A

Expedited retrievals are designed for urgent requests and can provide access to data in as little as 1-5 minutes for most archive objects. Standard retrievals typically finish within 3-5 hours for objects stored in the S3 Glacier Flexible Retrieval storage class or S3 Intelligent-Tiering Archive Access tier. These retrievals typically finish within 12 hours for objects stored in the S3 Glacier Deep Archive storage class or S3 Intelligent-Tiering Deep Archive Access tier. So A.

upvoted 3 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: A

Expedited retrievals allow you to quickly access your data that's stored in the S3 Glacier Flexible Retrieval storage class or the S3 Intelligent-Tiering Archive Access tier when occasional urgent requests for restoring archives are required. Data accessed by using Expedited retrievals is typically made available within 1–5 minutes.

upvoted 2 times

MrAWSAssociate 2 years, 5 months ago

Selected Answer: A

A for sure!

upvoted 2 times

Doyin8807 2 years, 6 months ago

C because A is not the most cost effective

upvoted 1 times

luiscc 2 years, 6 months ago

Selected Answer: A

Expedited retrieval typically takes 1-5 minutes to retrieve data, making it suitable for the company's requirement of having the files available in a maximum of five minutes.

upvoted 5 times

Efren 2 years, 6 months ago

Selected Answer: A

Glacier expedite

upvoted 3 times

EA100 2 years, 6 months ago

Answer - A

Fast availability: Although retrieval times for objects stored in Amazon S3 Glacier typically range from minutes to hours, you can use the Expedited retrievals option to expedite access to your archives. By using Expedited retrievals, the files can be made available in a maximum of five minutes when needed. However, Expedited retrievals do incur higher costs compared to standard retrievals.

upvoted 2 times

A company is building a three-tier application on AWS. The presentation tier will serve a static website. The logic tier is a containerized application. This application will store data in a relational database. The company wants to simplify deployment and to reduce operational costs.

Which solution will meet these requirements?

- A. Use Amazon S3 to host static content. Use Amazon Elastic Container Service (Amazon ECS) with AWS Fargate for compute power. Use a managed Amazon RDS cluster for the database. **Most Voted**
- B. Use Amazon CloudFront to host static content. Use Amazon Elastic Container Service (Amazon ECS) with Amazon EC2 for compute power. Use a managed Amazon RDS cluster for the database.
- C. Use Amazon S3 to host static content. Use Amazon Elastic Kubernetes Service (Amazon EKS) with AWS Fargate for compute power. Use a managed Amazon RDS cluster for the database.
- D. Use Amazon EC2 Reserved Instances to host static content. Use Amazon Elastic Kubernetes Service (Amazon EKS) with Amazon EC2 for compute power. Use a managed Amazon RDS cluster for the database.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Yadav_Sanjay **Highly Voted** 2 years, 6 months ago

Selected Answer: A

ECS is slightly cheaper than EKS
upvoted 12 times

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: A

Amazon S3 is a highly scalable and cost-effective storage service that can be used to host static website content. It provides durability, high availability, and low latency access to the static files.

Amazon ECS with AWS Fargate eliminates the need to manage the underlying infrastructure. It allows you to run containerized applications without provisioning or managing EC2 instances. This reduces operational overhead and provides scalability.

By using a managed Amazon RDS cluster for the database, you can offload the management tasks such as backups, patching, and monitoring to AWS. This reduces the operational burden and ensures high availability and durability of the database.

upvoted 5 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: A

B: CloudFront = Extra cost for something they don't want (CDN)

C: Kubernetes is more operationally complex than ECS containers on Fargate.

D: EC2 expensive

A: S3 is cheap for static content. ECS with Fargate is easiest implantation. Managed RDS is very low op overhead

upvoted 5 times

wsdasdasdqwdaw 2 years, 1 month ago

Why not B ?

upvoted 1 times

wsdasdasdqwdaw 2 years, 1 month ago

Aaa I got it. With CF we are adding additional cost => A.

upvoted 2 times

cyber_bedouin 2 years ago

A is better because ECS Fargate = "containerized application"

upvoted 2 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: A

Use Amazon S3 to host static content. Use Amazon Elastic Container Service (Amazon ECS) with AWS Fargate for compute power. Use a managed Amazon RDS cluster for the database.

upvoted 3 times

jaydesai8 2 years, 5 months ago

Selected Answer: A

S3= hosting static contents

Ecs = Little cheaper than EKS

RDS = Database

upvoted 4 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: A

Use Amazon S3 to host static content. Use Amazon Elastic Container Service (Amazon ECS) with AWS Fargate for compute power. Use a managed Amazon RDS cluster for the database

upvoted 3 times

A company seeks a storage solution for its application. The solution must be highly available and scalable. The solution also must function as a file system be mountable by multiple Linux instances in AWS and on premises through native protocols, and have no minimum size requirements. The company has set up a Site-to-Site VPN for access from its on-premises network to its VPC.

Which storage solution meets these requirements?

- A. Amazon FSx Multi-AZ deployments
- B. Amazon Elastic Block Store (Amazon EBS) Multi-Attach volumes
- C. Amazon Elastic File System (Amazon EFS) with multiple mount targets **Most Voted**
- D. Amazon Elastic File System (Amazon EFS) with a single mount target and multiple access points

Correct Answer: C

Community vote distribution

C (100%)

Comments

Felix_br **Highly Voted** 2 years, 6 months ago

Selected Answer: C

The other options are incorrect for the following reasons:

- A. Amazon FSx Multi-AZ deployments Amazon FSx is a managed file system service that provides access to file systems that are hosted on Amazon EC2 instances. Amazon FSx does not support native protocols, such as NFS.
- B. Amazon Elastic Block Store (Amazon EBS) Multi-Attach volumes Amazon EBS is a block storage service that provides durable, block-level storage volumes for use with Amazon EC2 instances. Amazon EBS Multi-Attach volumes can be attached to multiple EC2 instances at the same time, but they cannot be mounted by multiple Linux instances through native protocols, such as NFS.
- D. Amazon Elastic File System (Amazon EFS) with a single mount target and multiple access points A single mount target can only be used to mount the file system on a single EC2 instance. Multiple access points are used to provide access to the file system from different VPCs.

upvoted 14 times

unbendable 2 years, 1 month ago

Amazon FSx ONTAP supports clients mounting it with NFS. <https://docs.aws.amazon.com/fsx/latest/ONTAPGuide/attach-linux-client.html>. Though A is not clear about which FSx product is used

upvoted 2 times

dkw2342 1 year, 9 months ago

"A single mount target can only be used to mount the file system on a single EC2 instance. Multiple access points are used to provide access to the file system from different VPCs."

This is clearly wrong. You can have exactly one EFS mount target per subnet (AZ), and of course this mount target can be used by many clients (EC2 instances, containers etc.) - see diagram here for example: <https://docs.aws.amazon.com/efs/latest/ug/accessing-fs.html>

In my opinion, C and D are equally valid answers.

upvoted 2 times

Saliigen 1 year, 5 months ago

D is not valid answer because "the solution must be highly available" then it must be multi-AZ: every AZ requires a mount target

upvoted 1 times

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: C

Amazon EFS is a fully managed file system service that provides scalable, shared storage for Amazon EC2 instances. It supports the Network File System version 4 (NFSv4) protocol, which is a native protocol for Linux-based systems. EFS is designed to be highly available, durable, and scalable.

upvoted 9 times

awsgeek75 Most Recent 1 year, 11 months ago

Selected Answer: C

A: FSx is a File Server, not a mountable file system

B: EBS can't be mounted on on-prem devices

D: Access points are not same as mount points

C: EFS support multi mount targets and on-prem devices: <https://docs.aws.amazon.com/efs/latest/ug/mounting-fs-mount-helper-direct.html>

upvoted 3 times

iwannabeawsgod 2 years, 1 month ago

EFS POSIX LINUX

upvoted 3 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: C

C. Amazon Elastic File System (Amazon EFS) with multiple mount targets

upvoted 3 times

boubie44 2 years, 6 months ago

i don't understand why not D?

upvoted 1 times

lucdt4 2 years, 6 months ago

the requirement is mountable by multiple Linux

-> C (multiple mount targets)

upvoted 3 times

A 4-year-old media company is using the AWS Organizations all features feature set to organize its AWS accounts. According to the company's finance team, the billing information on the member accounts must not be accessible to anyone, including the root user of the member accounts.

Which solution will meet these requirements?

- A. Add all finance team users to an IAM group. Attach an AWS managed policy named Billing to the group.
- B. Attach an identity-based policy to deny access to the billing information to all users, including the root user.
- C. Create a service control policy (SCP) to deny access to the billing information. Attach the SCP to the root organizational unit (OU). **Most Voted**
- D. Convert from the Organizations all features feature set to the Organizations consolidated billing feature set.

Correct Answer: C

Community vote distribution

C (100%)

Comments

cloudehustia **Highly Voted** 2 years, 6 months ago

Selected Answer: C

Service Control Policies (SCP): SCPs are an integral part of AWS Organizations and allow you to set fine-grained permissions on the organizational units (OUs) within your AWS Organization. SCPs provide central control over the maximum permissions that can be granted to member accounts, including the root user.

Denying Access to Billing Information: By creating an SCP and attaching it to the root OU, you can explicitly deny access to billing information for all accounts within the organization. SCPs can be used to restrict access to various AWS services and actions, including billing-related services.

Granular Control: SCPs enable you to define specific permissions and restrictions at the organizational unit level. By denying access to billing information at the root OU, you can ensure that no member accounts, including root users, have access to the billing information.

upvoted 8 times

Kiki_Pass **Highly Voted** 2 years, 4 months ago

but SCP do not apply to the management account (full admin power)?

upvoted 5 times

TwinSpark 1 year, 7 months ago

i can understand this information coming from the famous course in udemy. I thought same, but after some research i now think it is a wrong information.

"SCPs affect all users and roles in attached accounts, including the root user. The only exceptions are those described in Tasks and entities not restricted by SCPs."

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scps.html#:~:text=SCPs%20affect%20all%20users%20and,affected%20any%20service%20linked%20role.

upvoted 3 times

potomac **Most Recent** 2 years, 1 month ago

Selected Answer: C

SCP is for authorization

upvoted 3 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: C

C. Create a service control policy (SCP) to deny access to the billing information. Attach the SCP to the root organizational unit (OU)

upvoted 3 times

PRASAD180 2 years, 5 months ago

C Crt 100%

upvoted 2 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: C

Service control policy are a type of organization policy that you can use to manage permissions in your organization. SCPs offer central control over the maximum available permissions for all accounts in your organization. SCPs help you to ensure your accounts stay within your organization's access control guidelines. SCPs are available only in an organization that has all features enabled.

upvoted 4 times

Abrar2022 2 years, 6 months ago

By denying access to billing information at the root OU, you can ensure that no member accounts, including root users, have access to the billing information.

upvoted 2 times

nosense 2 years, 6 months ago

Selected Answer: C

c for me

upvoted 2 times

An ecommerce company runs an application in the AWS Cloud that is integrated with an on-premises warehouse solution. The company uses Amazon Simple Notification Service (Amazon SNS) to send order messages to an on-premises HTTPS endpoint so the warehouse application can process the orders. The local data center team has detected that some of the order messages were not received.

A solutions architect needs to retain messages that are not delivered and analyze the messages for up to 14 days.

Which solution will meet these requirements with the LEAST development effort?

- A. Configure an Amazon SNS dead letter queue that has an Amazon Kinesis Data Stream target with a retention period of 14 days.
- B. Add an Amazon Simple Queue Service (Amazon SQS) queue with a retention period of 14 days between the application and Amazon SNS.
- C. Configure an Amazon SNS dead letter queue that has an Amazon Simple Queue Service (Amazon SQS) target with a retention period of 14 days. **Most Voted**
- D. Configure an Amazon SNS dead letter queue that has an Amazon DynamoDB target with a TTL attribute set for a retention period of 14 days.

Correct Answer: C

Community vote distribution

C (80%)

B (20%)

Comments

pentium75 **Highly Voted** 1 year, 11 months ago

Selected Answer: C

"Configuring an Amazon SNS dead-letter queue for a subscription ...

A dead-letter queue is an Amazon SQS queue that an Amazon SNS subscription can target for messages that can't be delivered to subscribers successfully", this is exactly what C says. <https://docs.aws.amazon.com/sns/latest/dg/sns-configure-dead-letter-queue.html>

B, an SQS queue "between the application and Amazon SNS" would change the application logic. SQS cannot push messages to the "on-premises https endpoint", rather the destination would have to retrieve messages from the queue. Besides, option B would eventually deliver the messages that failed on the first attempt, which is NOT what is asked for. The goal is to retain undeliverable messages for analysis (NOT to deliver them), and this is typically achieved with a dead letter queue.

upvoted 11 times

Rob1L **Highly Voted** 2 years, 6 months ago

Selected Answer: C

The message retention period in Amazon SQS can be set between 1 minute and 14 days (the default is 4 days). Therefore, you can configure your SQS DLQ to retain undelivered SNS messages for 14 days. This will enable you to analyze undelivered messages with the least development effort.

upvoted 5 times

osmk **Most Recent** 1 year, 11 months ago

A dead-letter queue is an Amazon SQS queue that an Amazon SNS subscription can target for messages that can't be delivered to subscribers successfully. <https://docs.aws.amazon.com/sns/latest/dg/sns-dead-letter-queues.html>

upvoted 3 times

awsgeek75 1 year, 11 months ago

Selected Answer: C

LEAST development effort!

A: Custom dead letter queue using Kinesis Data Stream (laughable solution!) so lots of coding

B: Change app logic to put SQS between SNS and the app. Also too much coding

D: Same as A, too much code change

C: SNS dead letter queue is by default a SQS queue so no coding required

upvoted 5 times

Mikado211 1 year, 12 months ago

Selected Answer: C

Problem here SNS dead letter queue is a SQS queue, so technically speaking both B and C are right. But I suppose that they want us to speak about SNS dead letter queue, that nobody do... meh, frustrating.

upvoted 4 times

Mikado211 1 year, 12 months ago

Aaaah ok.

So with B == you place the SQS queue between the application and the SNS topic

with C == you place the SQS queue as a DLQ for the SNS topic

Of course it's C !

upvoted 6 times

aws94 2 years ago

Selected Answer: C

<https://docs.aws.amazon.com/sns/latest/dg/sns-configure-dead-letter-queue.html>

upvoted 3 times

daniel1 2 years, 1 month ago

Selected Answer: C

GPT4 to the rescue:

The most appropriate solution would be to configure an Amazon SNS dead letter queue with an Amazon Simple Queue Service (Amazon SQS) target with a retention period of 14 days (Option C). This setup would ensure that any undelivered messages are retained in the SQS queue for up to 14 days for analysis, with minimal development effort required.

upvoted 2 times

ealpuce 2 years ago

ChatGPT is not a reliable source.

upvoted 9 times

Wayne23Fang 2 years, 1 month ago

Selected Answer: B

I like (B) since it is put SQS before SNS so we could prepare for retention. (C) dead letter Queue is kind of "rescue" effort. Also (C) should mention reprocessing dead letter.

upvoted 1 times

pentium75 1 year, 11 months ago

"Reprocessing dead letters" is not desired here. They want to "retain messages that are not delivered and analyze the messages for up to 14 days", which is what C does.

upvoted 1 times

thanhnv142 2 years, 1 month ago

C is correct. It used a combination of SNS and SQS so it better than B.

upvoted 2 times

iwannabeawsgod 2 years, 1 month ago

Selected Answer: C

C is the answer

upvoted 2 times

Devsin2000 2 years, 2 months ago

B is correct Answer. SQS Retain messages in queues for up to 14 days

C is incorrect because there is nothing called Amazon SNS dead letter queue

upvoted 2 times

RDM10 2 years, 2 months ago

<https://docs.aws.amazon.com/sns/latest/dg/sns-configure-dead-letter-queue.html>

upvoted 6 times

pentium75 1 year, 11 months ago

C "Configure an Amazon SNS dead letter queue"

AWS "Configuring an Amazon SNS dead-letter queue"

<https://docs.aws.amazon.com/sns/latest/dg/sns-configure-dead-letter-queue.html>

upvoted 2 times

lemur88 2 years, 3 months ago

Selected Answer: C

<https://docs.aws.amazon.com/sns/latest/dg/sns-dead-letter-queues.html>

upvoted 2 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: C

C. Configure an Amazon SNS dead letter queue that has an Amazon Simple Queue Service (Amazon SQS) target with a retention period of 14 days. By using an Amazon SQS queue as the target for the dead letter queue, you ensure that the undelivered messages are reliably stored in a queue for up to 14 days. Amazon SQS allows you to specify a retention period for messages, which meets the retention requirement without additional development effort.

upvoted 2 times

mtmayer 2 years, 3 months ago

Selected Answer: B

Dead Letter is a SQS feature not SNS.

A dead-letter queue is an Amazon SQS queue that an Amazon SNS subscription can target for messages that can't be delivered to subscribers successfully. Messages that can't be delivered due to client errors or server errors are held in the dead-letter queue for further analysis or reprocessing. For more information, see [Configuring an Amazon SNS dead-letter queue for a subscription and Amazon SNS message delivery retries](https://docs.aws.amazon.com/sns/latest/dg/sns-dead-letter-queues.html)

upvoted 3 times

pentium75 1 year, 11 months ago

"See [Configuring an Amazon SNS \(!\) dead-letter queue](#)", exactly, thus C.

upvoted 2 times

xyb 2 years, 4 months ago

Selected Answer: B

In SNS, DLQs store the messages that failed to be delivered to subscribed endpoints. For more information, see [Amazon SNS Dead-Letter Queues](#).

In SQS, DLQs store the messages that failed to be processed by your consumer application. This failure mode can happen when producers and consumers fail to interpret aspects of the protocol that they use to communicate. In that case, the consumer receives the message from the queue, but fails to process it, as the message doesn't have the structure or content that the consumer expects. The consumer can't delete the message from the queue either. After exhausting the receive count in the redrive policy, SQS can sideline the message to the DLQ. For more information, see [Amazon SQS Dead-Letter Queues](#).

<https://aws.amazon.com/blogs/compute/designing-durable-serverless-apps-with-dlqs-for-amazon-sns-amazon-sqs-aws-lambda/>

upvoted 2 times

pentium75 1 year, 11 months ago

"Configuring an Amazon SNS dead-letter queue for a subscription

A dead-letter queue is an Amazon SQS queue that an Amazon SNS subscription can target for messages that can't be delivered to subscribers successfully. "

upvoted 2 times

TariqKipkemei 2 years, 5 months ago

C is best to handle this requirement. Although good to note that dead-letter queue is an SQS queue.

"A dead-letter queue is an Amazon SQS queue that an Amazon SNS subscription can target for messages that can't be delivered to subscribers successfully. Messages that can't be delivered due to client errors or server errors are held in the dead-letter queue for further analysis or reprocessing."

<https://docs.aws.amazon.com/sns/latest/dg/sns-dead-letter-queues.html#:~:text=A%20dead%20letter%20queue%20is%20an%20Amazon%20SQS%20queue>

upvoted 2 times

Felix_br 2 years, 6 months ago

C - Amazon SNS dead letter queues are used to handle messages that are not delivered to their intended recipients. When a message is sent to an Amazon SNS topic, it is first delivered to the topic's subscribers. If a message is not delivered to any of the subscribers, it is sent to the topic's dead letter queue.

Amazon SQS is a fully managed message queuing service that enables you to decouple and scale microservices, distributed systems, and serverless applications. Amazon SQS queues can be configured to have a retention period, which is the amount of time that messages will be kept in the queue before they are deleted.

To meet the requirements of the company, you can configure an Amazon SNS dead letter queue that has an Amazon SQS target with a retention period of 14 days. This will ensure that any messages that are not delivered to the on-premises warehouse application will be stored in the Amazon SQS queue for up to 14 days. The company can then analyze the messages in the Amazon SQS queue to determine why they were not delivered.

upvoted 3 times

A gaming company uses Amazon DynamoDB to store user information such as geographic location, player data, and leaderboards. The company needs to configure continuous backups to an Amazon S3 bucket with a minimal amount of coding. The backups must not affect availability of the application and must not affect the read capacity units (RCUs) that are defined for the table.

Which solution meets these requirements?

- A. Use an Amazon EMR cluster. Create an Apache Hive job to back up the data to Amazon S3.
- B. Export the data directly from DynamoDB to Amazon S3 with continuous backups. Turn on point-in-time recovery for the table. **Most Voted**
- C. Configure Amazon DynamoDB Streams. Create an AWS Lambda function to consume the stream and export the data to an Amazon S3 bucket.
- D. Create an AWS Lambda function to export the data from the database tables to Amazon S3 on a regular basis. Turn on point-in-time recovery for the table.

Correct Answer: B

Community vote distribution

B (87%)

13%

Comments

elmogy **Highly Voted** 2 years, 6 months ago

Selected Answer: B

Continuous backups is a native feature of DynamoDB, it works at any scale without having to manage servers or clusters and allows you to export data across AWS Regions and accounts to any point-in-time in the last 35 days at a per-second granularity. Plus, it doesn't affect the read capacity or the availability of your production tables.

<https://aws.amazon.com/blogs/aws/new-export-amazon-dynamodb-table-data-to-data-lake-amazon-s3/>

upvoted 12 times

awsgeek75 **Highly Voted** 1 year, 11 months ago

Selected Answer: B

A: Impacts RCU
C: Requires coding of Lambda to read from stream to S3
D: More coding in Lambda
B: AWS Managed solution with no coding

upvoted 5 times

In2718 **Most Recent** 10 months, 4 weeks ago

Selected Answer: C

The correct answer should be C. The question states... "configure continuous backups to an Amazon S3 bucket with a minimal amount of coding". It says "minimal" amount of coding, and "continuous" backup. Answer B is not the correct answer.

upvoted 1 times

potomac 2 years, 1 month ago

Selected Answer: B

DynamoDB export to S3 is a fully managed solution for exporting DynamoDB data to an Amazon S3 bucket at scale.

upvoted 4 times

baba365 2 years, 2 months ago

A DynamoDB stream is an ordered flow of information about changes to items in a DynamoDB table... for C.U.D events (Create, Update, Delete) and its logs are retained for only 24hrs .

upvoted 4 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: B

Export the data directly from DynamoDB to Amazon S3 with continuous backups. Turn on point-in-time recovery for the table.
upvoted 3 times

ukivanlamlpi 2 years, 4 months ago

Selected Answer: C

continuous backup, no impact to availability ==> DynamoDB stream
B. export is one off, no continuous and demand on read capacity
upvoted 4 times

hsinchang 2 years, 4 months ago

minimal amount of coding rules out Lambda
upvoted 4 times

Chris22usa 2 years, 5 months ago

ChatGPT answer is C and it indicates continuous backup process uses DynamoDB stream actually
upvoted 2 times

pentium75 1 year, 11 months ago

ChatGPT is usually wrong on these topics.
upvoted 2 times

Gajendr 1 year, 11 months ago

Wrong.
"DynamoDB full exports are charged based on the size of the DynamoDB table (table data and local secondary indexes) at the point in time for which the export is done. DynamoDB incremental exports are charged based on the size of data processed from your continuous backups for the time period being exported."
<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/S3DataExport.HowItWorks.html>
upvoted 1 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: B

Using DynamoDB table export, you can export data from an Amazon DynamoDB table from any time within your point-in-time recovery window to an Amazon S3 bucket. Exporting a table does not consume read capacity on the table, and has no impact on table performance and availability.
upvoted 1 times

norris81 2 years, 6 months ago

Selected Answer: B

<https://repost.aws/knowledge-center/back-up-dynamodb-s3>
<https://aws.amazon.com/blogs/aws/new-amazon-dynamodb-continuous-backups-and-point-in-time-recovery-pitr/>
<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Streams.Lambda.html>

There is no edit
upvoted 2 times

cloudenthusiast 2 years, 6 months ago

Selected Answer: B

Continuous Backups: DynamoDB provides a feature called continuous backups, which automatically backs up your table data. Enabling continuous backups ensures that your table data is continuously backed up without the need for additional coding or manual interventions.

Export to Amazon S3: With continuous backups enabled, DynamoDB can directly export the backups to an Amazon S3 bucket. This eliminates the need for custom coding to export the data.

Minimal Coding: Option B requires the least amount of coding effort as continuous backups and the export to Amazon S3 functionality are built-in features of DynamoDB.

No Impact on Availability and RCUs: Enabling continuous backups and exporting data to Amazon S3 does not affect the availability of your application or the read capacity units (RCUs) defined for the table. These operations happen in the background and do not impact the table's performance or consume additional RCUs.

upvoted 4 times

Efren 2 years, 6 months ago

Selected Answer: B

DynamoDB Export to S3 feature
Using this feature, you can export data from an Amazon DynamoDB table anytime within your point-in-time recovery window to an Amazon S3 bucket.
upvoted 3 times

Efren 2 years, 6 months ago

B also for me

upvoted 3 times

norris81 2 years, 6 months ago

<https://repost.aws/knowledge-center/back-up-dynamodb-s3>

<https://aws.amazon.com/blogs/aws/new-amazon-dynamodb-continuous-backups-and-point-in-time-recovery-pitr/>

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Streams.Lambda.html>

upvoted 1 times

Efren 2 years, 6 months ago

you could mention what is the best answer from you :)

upvoted 2 times

A solutions architect is designing an asynchronous application to process credit card data validation requests for a bank. The application must be secure and be able to process each request at least once.

Which solution will meet these requirements MOST cost-effectively?

A. Use AWS Lambda event source mapping. Set Amazon Simple Queue Service (Amazon SQS) standard queues as the event source. Use AWS Key Management Service (SSE-KMS) for encryption. Add the kms:Decrypt permission for the Lambda execution role. **Most Voted**


B. Use AWS Lambda event source mapping. Use Amazon Simple Queue Service (Amazon SQS) FIFO queues as the event source. Use SQS managed encryption keys (SSE-SQS) for encryption. Add the encryption key invocation permission for the Lambda function.

C. Use the AWS Lambda event source mapping. Set Amazon Simple Queue Service (Amazon SQS) FIFO queues as the event source. Use AWS KMS keys (SSE-KMS). Add the kms:Decrypt permission for the Lambda execution role.

D. Use the AWS Lambda event source mapping. Set Amazon Simple Queue Service (Amazon SQS) standard queues as the event source. Use AWS KMS keys (SSE-KMS) for encryption. Add the encryption key invocation permission for the Lambda function.

Correct Answer: A

Community vote distribution



Comments

elmogy **Highly Voted** 3 years ago

Selected Answer: A

SQS FIFO is slightly more expensive than standard queue
<https://calculator.aws/#/addService/SQS>

I would still go with the standard because of the keyword "at least once" because FIFO process "exactly once". That leaves us with A and D, I believe that lambda function only needs to decrypt so I would choose A
 upvoted 12 times

pentium75 **Highly Voted** 2 years, 5 months ago

Selected Answer: A

"Process each request at least once" = Standard queue, rules out B and C which use more expensive FIFO queue

Permissions are added to Lambda execution roles, not Lambda functions, thus D is out.
 upvoted 11 times

In2718 **Most Recent** 9 months, 3 weeks ago

Selected Answer: B

In real life the answer is B --- you would never run credit card transmissions "at least once". Every real-world credit transaction scenario uses FIFO. Also, you could use the keys stated in B. Ridiculous question, should not have been used. This is being stated just so you know the answer for the real world.
 upvoted 2 times

emakid 1 year, 11 months ago

Selected Answer: A

Use AWS Lambda event source mapping. Set Amazon Simple Queue Service (Amazon SQS) standard queues as the event source. Use AWS Key Management Service (SSE-KMS) for encryption. Add the kms permission for the Lambda execution role.
 upvoted 2 times

JackyCCK 2 years, 3 months ago

D is not FIFO either
upvoted 1 times

EdenWang 2 years, 6 months ago

Selected Answer: B

With the SSE-SQS encryption type, you do not need to create, manage, or pay for SQS-managed encryption keys.
upvoted 1 times

pentium75 2 years, 5 months ago

And what the hell is "encryption key invocation permission for the Lambda function"?
upvoted 5 times

wsdassdasdqwda 2 years, 7 months ago

Initially though it is B, but it is said that the messages should be processed at least once, not the same order, and Standard SQS is "almost" FIFO, which changed my opinion and I would go with A as correct.
upvoted 4 times

BrijMohan08 2 years, 9 months ago

Selected Answer: A

<https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/standard-queues.html>
upvoted 5 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: B

Using SQS FIFO queues ensures each message is processed at least once in order. SSE-SQS provides encryption that is handled entirely by SQS without needing decrypt permissions.

Standard SQS queues (Options A and D) do not guarantee order.

Using KMS keys (Options C and D) requires providing the Lambda role with decrypt permissions, adding complexity.

SQS FIFO queues with SSE-SQS encryption provide orderly, secure, server-side message processing that Lambda can consume without needing to manage decryption. This is the most efficient and cost-effective approach.
upvoted 8 times

pentium75 2 years, 5 months ago

But permissions are added to Lambda execution roles, not functions
upvoted 5 times

Clouddon 2 years, 8 months ago

Amazon SQS offers standard as the default queue type. Standard queues support a nearly unlimited number of API calls per second, per API action (SendMessage, ReceiveMessage, or DeleteMessage). Standard queues support at-least-once message delivery. However, occasionally (because of the highly distributed architecture that allows nearly unlimited throughput), more than one copy of a message might be delivered out of order. Standard queues provide best-effort ordering which ensures that messages are generally delivered in the same order as they're sent. Whereas, FIFO (First-In-First-Out) queues have all the capabilities of the standard queues, but are designed to enhance messaging between applications when the order of operations and events is critical, or where duplicates can't be tolerated. (is correct)
upvoted 5 times

hsinchang 2 years, 10 months ago

Least Privilege Policy leads to A over D.
upvoted 2 times

TariqKipkemei 2 years, 11 months ago

Selected Answer: B

Considering this is credit card validation process, there needs to be a strict 'process exactly once' policy offered by the SQS FIFO, and also SQS already supports server-side encryption with customer-provided encryption keys using the AWS Key Management Service (SSE-KMS) or using SQS-owned encryption keys (SSE-SQS). Both encryption options greatly reduce the operational burden and complexity involved in protecting data. Additionally, with the SSE-SQS encryption type, you do not need to create, manage, or pay for SQS-managed encryption keys. Therefore option B stands out for me.
upvoted 1 times

TariqKipkemei 2 years, 7 months ago

I retract my answer and change it to A, there is a requirement to process each request 'at least once'. Only standard queues can deliver messages at least once.
There is also a requirement for the most 'cost-effective' option. Standard queues are the cheaper option.

<https://aws.amazon.com/sqs/pricing/#:~:text=SQS%20requests%20priced%3F>
upvoted 3 times

darren_song 2 years, 11 months ago

Selected Answer: A

https://docs.aws.amazon.com/zh_tw/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-least-privilege-policy.html

upvoted 2 times

Abrar2022 3 years ago

Selected Answer: A

at least once and cost effective suggests SQS standard

upvoted 2 times

Felix_br 3 years ago

Selected Answer: B

Solution B is the most cost-effective solution to meet the requirements of the application.

Amazon Simple Queue Service (SQS) FIFO queues are a good choice for this application because they guarantee that messages are processed in the order in which they are received. This is important for credit card data validation because it ensures that fraudulent transactions are not processed before legitimate transactions.

SQS managed encryption keys (SSE-SQS) are a good choice for encrypting the messages in the SQS queue because they are free to use. AWS Key Management Service (KMS) keys (SSE-KMS) are also a good choice for encrypting the messages, but they do incur a cost.

upvoted 2 times

pentium75 2 years, 5 months ago

"They guarantee that messages are processed in the order in which they are received. This is important" but not asked for!

upvoted 3 times

omoakin 3 years ago

AAAAAAA

upvoted 1 times

Yadav_Sanjay 3 years ago

Selected Answer: A

should be A. Key word - at least once and cost effective suggests SQS standard

upvoted 3 times

Efren 3 years ago

It has to be default, no FIFO. It doesn't say just one, it says at least once, so that is default queue that is cheaper than FIFO. Between the default options, not sure to be honest

upvoted 4 times

jayce5 3 years ago

No, when it comes to "credit card data validation," it should be FIFO. If you use the standard approach, there is a chance that people who come after will get processed before those who come first.

upvoted 1 times

pentium75 2 years, 5 months ago

Question clearly says "process each request at least once" which is the description of a standard queue. Your opinion how these transactions should be processed does not matter if it contradicts the requirements given.

Besides, it is about "credit card data validation", NOT payments. Nothing happens if they check twice if your credit card is valid.

upvoted 2 times

A company has multiple AWS accounts for development work. Some staff consistently use oversized Amazon EC2 instances, which causes the company to exceed the yearly budget for the development accounts. The company wants to centrally restrict the creation of AWS resources in these accounts.

Which solution will meet these requirements with the LEAST development effort?

- A. Develop AWS Systems Manager templates that use an approved EC2 creation process. Use the approved Systems Manager templates to provision EC2 instances.
- B. Use AWS Organizations to organize the accounts into organizational units (OUs). Define and attach a service control policy (SCP) to control the usage of EC2 instance types. **Most Voted**
- C. Configure an Amazon EventBridge rule that invokes an AWS Lambda function when an EC2 instance is created. Stop disallowed EC2 instance types.
- D. Set up AWS Service Catalog products for the staff to create the allowed EC2 instance types. Ensure that staff can deploy EC2 instances only by using the Service Catalog products.

Correct Answer: B

Community vote distribution

B (93%)

7%

Comments

alexandercamachop **Highly Voted** 2 years, 6 months ago

Selected Answer: B

Anytime you see Multiple AWS Accounts, and needs to consolidate is AWS Organization. Also anytime we need to restrict anything in an organization, it is SCP policies.

upvoted 9 times

omarshaban **Highly Voted** 1 year, 11 months ago

IN MY EXAM

upvoted 6 times

Cyberkayu **Most Recent** 1 year, 11 months ago

Selected Answer: B

B. Multiple AWS account, consolidate under one AWS Organization, top down policy (SCP) to all member account to restrict EC2 Type.

upvoted 3 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: B

Use AWS Organizations to organize the accounts into organizational units (OUs). Define and attach a service control policy (SCP) to control the usage of EC2 instance types.

upvoted 3 times

Ale1973 2 years, 4 months ago

Selected Answer: D

I have a question regarding this answer, what do they mean by "development effort"?:

If they mean the work it takes to implement the solution (using develop as implement), option B achieves the constraint with little administrative overhead (there is less to do to configure this option).

If by "development effort", they mean less effort for the development team, when development team try to deploy instances and gets errors because they are not allowed, this generates overhead. In this case the best option is D.

What did you think?

upvoted 2 times

pentium75 1 year, 11 months ago

"Development effort" = Develop the solution that the question asks for. We don't care about the developers whose permissions we want to restrict.
upvoted 4 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: B

Use AWS Organizations to organize the accounts into organizational units (OUs). Define and attach a service control policy (SCP) to control the usage of EC2 instance types
upvoted 3 times

Blingy 2 years, 6 months ago

BBBBBBBBB

upvoted 2 times

elmogy 2 years, 6 months ago

Selected Answer: B

I would choose B

The other options would require some level of programming or custom resource creation:

A. Developing Systems Manager templates requires development effort

C. Configuring EventBridge rules and Lambda functions requires development effort

D. Creating Service Catalog products requires development effort to define the allowed EC2 configurations.

Option B - Using Organizations service control policies - requires no custom development. It involves:

Organizing accounts into OUs

Creating an SCP that defines allowed/disallowed EC2 instance types

Attaching the SCP to the appropriate OUs

This is a native AWS service with a simple UI for defining and managing policies. No coding or resource creation is needed.

So option B, using Organizations service control policies, will meet the requirements with the least development effort.

upvoted 4 times

cloudenthusiast 2 years, 6 months ago

Selected Answer: B

AWS Organizations: AWS Organizations is a service that helps you centrally manage multiple AWS accounts. It enables you to group accounts into organizational units (OUs) and apply policies across those accounts.

Service Control Policies (SCPs): SCPs in AWS Organizations allow you to define fine-grained permissions and restrictions at the account or OU level. By attaching an SCP to the development accounts, you can control the creation and usage of EC2 instance types.

Least Development Effort: Option B requires minimal development effort as it leverages the built-in features of AWS Organizations and SCPs. You can define the SCP to restrict the use of oversized EC2 instance types and apply it to the appropriate OUs or accounts.

upvoted 5 times

Efren 2 years, 6 months ago

B for me as well

upvoted 2 times

A company wants to use artificial intelligence (AI) to determine the quality of its customer service calls. The company currently manages calls in four different languages, including English. The company will offer new languages in the future. The company does not have the resources to regularly maintain machine learning (ML) models.

The company needs to create written sentiment analysis reports from the customer service call recordings. The customer service call recording text must be translated into English.

Which combination of steps will meet these requirements? (Choose three.)

- A. Use Amazon Comprehend to translate the audio recordings into English.
- B. Use Amazon Lex to create the written sentiment analysis reports.
- C. Use Amazon Polly to convert the audio recordings into text.
- D. Use Amazon Transcribe to convert the audio recordings in any language into text. **Most Voted**
- E. Use Amazon Translate to translate text in any language to English. **Most Voted**
- F. Use Amazon Comprehend to create the sentiment analysis reports. **Most Voted**

Correct Answer: DEF

Community vote distribution

DEF (100%)

Comments

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: DEF

Amazon Transcribe will convert the audio recordings into text, Amazon Translate will translate the text into English, and Amazon Comprehend will perform sentiment analysis on the translated text to generate sentiment analysis reports.

upvoted 7 times

awsgeek75 **Highly Voted** 1 year, 11 months ago

Selected Answer: DEF

A: Comprehend cannot translate
 B: Lex is like a chatbot so not useful
 C: Polly converts text to audio (polly the parrot!) so this is wrong
 D: Can convert audio to text
 E: Can translate
 F: Can do sentiment analysis reports

upvoted 6 times

wsdasdasdqwda **Most Recent** 2 years, 1 month ago

It is: DEF

upvoted 2 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: DEF

D. Use Amazon Transcribe to convert the audio recordings in any language into text.
 E. Use Amazon Translate to translate text in any language to English.
 F. Use Amazon Comprehend to create the sentiment analysis reports.

upvoted 3 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: DEF

Amazon Transcribe to convert speech to text. Amazon Translate to translate text to english. Amazon Comprehend to perform sentiment analysis on translated text.

upvoted 2 times

HareshPrajapati 2 years, 6 months ago

afree with DEF

upvoted 2 times

Blingy 2 years, 6 months ago

I'd go with DEF too

upvoted 2 times

elmogy 2 years, 6 months ago

Selected Answer: DEF

agree with DEF

upvoted 3 times

Efren 2 years, 6 months ago

agreed as well, weird

upvoted 2 times

Guru4Cloud 2 years, 3 months ago

@efren - It is not weird - This need to know the services for it

upvoted 3 times

A company uses Amazon EC2 instances to host its internal systems. As part of a deployment operation, an administrator tries to use the AWS CLI to terminate an EC2 instance. However, the administrator receives a 403 (Access Denied) error message.

The administrator is using an IAM role that has the following IAM policy attached:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": ["ec2:TerminateInstances"],
      "Resource": ["*"]
    },
    {
      "Effect": "Deny",
      "Action": ["ec2:TerminateInstances"],
      "Condition": {
        "NotIpAddress": {
          "aws:SourceIp": [
            "192.0.2.0/24",
            "203.0.113.0/24"
          ]
        }
      }
    },
    {
      "Effect": "Deny",
      "Action": ["ec2:TerminateInstances"],
      "Resource": ["*"]
    }
  ]
}
```

What is the cause of the unsuccessful request?

- A. The EC2 instance has a resource-based policy with a Deny statement.
- B. The principal has not been specified in the policy statement.
- C. The "Action" field does not grant the actions that are required to terminate the EC2 instance.
- D. The request to terminate the EC2 instance does not originate from the CIDR blocks 192.0.2.0/24 or 203.0.113.0/24.

Most Voted

Correct Answer: D

Community vote distribution

D (100%)

Comments

TariqKipkemei **Highly Voted** 2 years, 11 months ago

Selected Answer: D

the command is coming from a source IP which is not in the allowed range.

upvoted 5 times

chasingsummer **Most Recent** 2 years, 5 months ago

Selected Answer: D

I ran a Policy Simulator and indeed, D is right answer.

Here is the JSON policy:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "ec2:TerminateInstances",
      "Resource": "*"
    },
    {
      "Effect": "Deny",
      "Action": "ec2:TerminateInstances",
      "Condition": {
        "NotIpAddress": {
          "aws:SourceIp": [
            "192.0.2.0/24",
            "203.0.113.0/24"
          ]
        }
      },
      "Resource": "*"
    }
  ]
}
```

upvoted 2 times

chasingsummer 2 years, 5 months ago

The condition operator is "NotIpAddress" so I am not sure about D as right answer.

upvoted 2 times

awsgeek75 2 years, 4 months ago

Deny when IP address is not in (NotIpAddress). AWS has a weird way of stating Deny and it almost sound like double negative meaning positive. But read this doc for more clarity:

https://docs.aws.amazon.com/IAM/latest/UserGuide/reference_policies_examples_aws_deny-ip.html

It has the exact same example! Good luck!

upvoted 2 times

LeonSauveterre 1 year, 6 months ago

This is exactly what I need, thanks!

upvoted 1 times

awsgeek75 2 years, 5 months ago

Selected Answer: D

If you want to read more about this, see how it works: https://docs.aws.amazon.com/IAM/latest/UserGuide/reference_policies_examples_aws_deny-ip.html

Same policy as in this question with almost same use case.

D is correct answer.

upvoted 4 times

elmogy 3 years ago

Selected Answer: D

"aws:SourceIP" indicates the IP address that is trying to perform the action.

upvoted 2 times

nosense 3 years ago

Selected Answer: D

d for sure

upvoted 3 times

A company is conducting an internal audit. The company wants to ensure that the data in an Amazon S3 bucket that is associated with the company's AWS Lake Formation data lake does not contain sensitive customer or employee data. The company wants to discover personally identifiable information (PII) or financial information, including passport numbers and credit card numbers.

Which solution will meet these requirements?

- A. Configure AWS Audit Manager on the account. Select the Payment Card Industry Data Security Standards (PCI DSS) for auditing.
- B. Configure Amazon S3 Inventory on the S3 bucket. Configure Amazon Athena to query the inventory.
- C. Configure Amazon Macie to run a data discovery job that uses managed identifiers for the required data types. **Most Voted**
- D. Use Amazon S3 Select to run a report across the S3 bucket.

Correct Answer: C

Community vote distribution

C (100%)

Comments

elmogy **Highly Voted** 2 years, 6 months ago

Selected Answer: C

agree with C

upvoted 5 times

Blingy **Highly Voted** 2 years, 6 months ago

Macie = Sensitive PII

upvoted 5 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: C

PII or sensitive data = Macie

upvoted 4 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: C

Configure Amazon Macie to run a data discovery job that uses managed identifiers for the required data types.

upvoted 3 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: C

Amazon Macie is a data security service that uses machine learning (ML) and pattern matching to discover and help protect your sensitive data.

upvoted 3 times

cloudenthusiast 2 years, 6 months ago

Selected Answer: C

Amazon Macie is a service that helps discover, classify, and protect sensitive data stored in AWS. It uses machine learning algorithms and managed identifiers to detect various types of sensitive information, including personally identifiable information (PII) and financial information. By configuring Amazon Macie to run a data discovery job with the appropriate managed identifiers for the required data types (such as passport numbers and credit card numbers), the company can identify and classify any sensitive data present in the S3 bucket.

upvoted 5 times

A company uses on-premises servers to host its applications. The company is running out of storage capacity. The applications use both block storage and NFS storage. The company needs a high-performing solution that supports local caching without re-architecting its existing applications.

Which combination of actions should a solutions architect take to meet these requirements? (Choose two.)

- A. Mount Amazon S3 as a file system to the on-premises servers.
- B. Deploy an AWS Storage Gateway file gateway to replace NFS storage. **Most Voted**
- C. Deploy AWS Snowball Edge to provision NFS mounts to on-premises servers.
- D. Deploy an AWS Storage Gateway volume gateway to replace the block storage. **Most Voted**
- E. Deploy Amazon Elastic File System (Amazon EFS) volumes and mount them to on-premises servers.

Correct Answer: BD

Community vote distribution

BD (100%)

Comments

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: BD

By combining the deployment of an AWS Storage Gateway file gateway and an AWS Storage Gateway volume gateway, the company can address both its block storage and NFS storage needs, while leveraging local caching capabilities for improved performance.

upvoted 7 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: BD

A: Not possible
C: Snowball edge is snowball with computing. It's not a NAS!
E: Technically yes but requires VPN or Direct Connect so re-architecture
B & D both use Storage Gateway which can be used as NFS and Block storage
<https://aws.amazon.com/storagegateway/>

upvoted 4 times

ftaws 1 year, 11 months ago

Use the Storage Gateway -> It means that use S3 for storage ?

upvoted 3 times

thanhnv142 2 years, 1 month ago

DE

B is not correct because NFS is a file system while storage gw is a storage. To replace a file system, need another file system which is EFS.

upvoted 3 times

Tekk97 2 years ago

That's what I thought. but I think B is work too.

upvoted 2 times

wizcloudifa 1 year, 7 months ago

if you focus on the wording of option B its "Storage gateway File gateway" not volume gateway hence it is perfect replacement for NFS files.

upvoted 1 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: BD

Deploy an AWS Storage Gateway file gateway to replace NFS storage
Deploy an AWS Storage Gateway volume gateway to replace the block storage

upvoted 3 times

elmogy 2 years, 6 months ago

Selected Answer: BD

local caching is a key feature of AWS Storage Gateway solution

<https://aws.amazon.com/storagegateway/features/>

<https://aws.amazon.com/blogs/storage/aws-storage-gateway-increases-cache-4x-and-enhances-bandwidth-throttling/#:~:text=AWS%20Storage%20Gateway%20increases%20cache%204x%20and%20enhances,for%20Volume%20Gateway%20customers%20.%205%20Conclusion%20>

upvoted 4 times

Piccalo 2 years, 6 months ago

Selected Answer: BD

B and D is the correct answer

upvoted 2 times

A company has a service that reads and writes large amounts of data from an Amazon S3 bucket in the same AWS Region. The service is deployed on Amazon EC2 instances within the private subnet of a VPC. The service communicates with Amazon S3 over a NAT gateway in the public subnet. However, the company wants a solution that will reduce the data output costs.

Which solution will meet these requirements MOST cost-effectively?

- A. Provision a dedicated EC2 NAT instance in the public subnet. Configure the route table for the private subnet to use the elastic network interface of this instance as the destination for all S3 traffic.
- B. Provision a dedicated EC2 NAT instance in the private subnet. Configure the route table for the public subnet to use the elastic network interface of this instance as the destination for all S3 traffic.
- C. Provision a VPC gateway endpoint. Configure the route table for the private subnet to use the gateway endpoint as the route for all S3 traffic. **Most Voted**
- D. Provision a second NAT gateway. Configure the route table for the private subnet to use this NAT gateway as the destination for all S3 traffic.

Correct Answer: C

Community vote distribution

C (100%)

Comments

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: C

A VPC gateway endpoint allows you to privately access Amazon S3 from within your VPC without using a NAT gateway or NAT instance. By provisioning a VPC gateway endpoint for S3, the service in the private subnet can directly communicate with S3 without incurring data transfer costs for traffic going through a NAT gateway.

upvoted 11 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: C

As a rule of thumb, EC2 <-> S3 in your workload should always try to use a VPC gateway unless there is an explicit restriction (account etc.) which disallows it.

upvoted 4 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: C

Using a VPC endpoint for S3 allows the EC2 instances to access S3 directly over the Amazon network without traversing the internet. This significantly reduces data output charges.

upvoted 3 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: C

use VPC gateway endpoint to route traffic internally and save on costs.

upvoted 2 times

elmogy 2 years, 6 months ago

Selected Answer: C

private subnet needs to communicate with S3 --> VPC endpoint right away

upvoted 3 times

A company uses Amazon S3 to store high-resolution pictures in an S3 bucket. To minimize application changes, the company stores the pictures as the latest version of an S3 object. The company needs to retain only the two most recent versions of the pictures.

The company wants to reduce costs. The company has identified the S3 bucket as a large expense.

Which solution will reduce the S3 costs with the LEAST operational overhead?

- A. Use S3 Lifecycle to delete expired object versions and retain the two most recent versions. **Most Voted**
- B. Use an AWS Lambda function to check for older versions and delete all but the two most recent versions.
- C. Use S3 Batch Operations to delete noncurrent object versions and retain only the two most recent versions.
- D. Deactivate versioning on the S3 bucket and retain the two most recent versions.

Correct Answer: A

Community vote distribution

A (100%)

Comments

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: A

S3 Lifecycle policies allow you to define rules that automatically transition or expire objects based on their age or other criteria. By configuring an S3 Lifecycle policy to delete expired object versions and retain only the two most recent versions, you can effectively manage the storage costs while maintaining the desired retention policy. This solution is highly automated and requires minimal operational overhead as the lifecycle management is handled by S3 itself.

upvoted 6 times

VellaDevil **Highly Voted** 2 years, 5 months ago

Selected Answer: A

A --> "you can also provide a maximum number of noncurrent versions to retain."
<https://docs.aws.amazon.com/AmazonS3/latest/userguide/intro-lifecycle-rules.html>

upvoted 5 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: A

B: Too much work with Lambda

C: Possible but requires lot of work

D: Oxymoron statement... i.e. how do you remove version and retain version at same time without additional overhead? Custom solution may be more work.

A: S3 Lifecycle is designed to retain object and version with set criteria

upvoted 4 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: A

Use S3 Lifecycle to delete expired object versions and retain the two most recent versions.

upvoted 3 times

TariqKipkemei 2 years, 5 months ago

Selected Answer: A

S3 Lifecycle to the rescue...whooooosh

upvoted 3 times

antropaws 2 years, 6 months ago

Selected Answer: A

A is correct.

upvoted 3 times

Konb 2 years, 6 months ago

Selected Answer: A

Agree with LONGMEN

upvoted 4 times

A company needs to minimize the cost of its 1 Gbps AWS Direct Connect connection. The company's average connection utilization is less than 10%. A solutions architect must recommend a solution that will reduce the cost without compromising security.

Which solution will meet these requirements?

- A. Set up a new 1 Gbps Direct Connect connection. Share the connection with another AWS account.
- B. Set up a new 200 Mbps Direct Connect connection in the AWS Management Console.
- C. Contact an AWS Direct Connect Partner to order a 1 Gbps connection. Share the connection with another AWS account.
- D. Contact an AWS Direct Connect Partner to order a 200 Mbps hosted connection for an existing AWS account. **Most Voted**

Correct Answer: D

Community vote distribution

D (83%)

B (17%)

Comments

Abrar2022 **Highly Voted** 2 years, 6 months ago

Selected Answer: D

Hosted Connection 50 Mbps, 100 Mbps, 200 Mbps,
Dedicated Connection 1 Gbps, 10 Gbps, and 100 Gbps
upvoted 11 times

norris81 **Highly Voted** 2 years, 6 months ago

Selected Answer: D

D

For Dedicated Connections, 1 Gbps, 10 Gbps, and 100 Gbps ports are available. For Hosted Connections, connection speeds of 50 Mbps, 100 Mbps, 200 Mbps, 300 Mbps, 400 Mbps, 500 Mbps, 1 Gbps, 2 Gbps, 5 Gbps and 10 Gbps may be ordered from approved AWS Direct Connect Partners. See AWS Direct Connect Partners for more information.

upvoted 5 times

Ravan **Most Recent** 1 year, 9 months ago

Selected Answer: D

No, you cannot directly adjust the speed of an existing Direct Connect connection through the AWS Management Console.

To adjust the speed of an existing Direct Connect connection, you typically need to contact your Direct Connect service provider. They can assist you in modifying the speed of your connection based on your requirements. Depending on the provider, this process may involve submitting a request or contacting their support team to initiate the necessary changes. Keep in mind that adjusting the speed of your Direct Connect connection may also involve contractual and billing considerations.

upvoted 4 times

awsgeek75 1 year, 11 months ago

Selected Answer: D

A: Not secure as sharing with another account
B: I don't think this possible as you need ISP to setup Direct Connect
C: Less secure due to sharing
D: Direct connect partners can provide hosted solutions for existing accounts so correct answer
upvoted 5 times

awsgeek75 1 year, 11 months ago

For B I'm wrong above, it's because you cannot order 200MB connection through management console.

upvoted 3 times

pentium75 1 year, 11 months ago

Selected Answer: D

< 1 Gbps = Hosted (through partner)

upvoted 3 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: B

If you already have an existing AWS Direct Connect connection configured at 1 Gbps, and you wish to reduce the connection bandwidth to 200 Mbp to minimize costs, you should indeed contact your AWS Direct Connect Partner and request to lower the connection speed to 200 Mbps.

upvoted 4 times

Guru4Cloud 2 years, 3 months ago

I meant D.. DDDDDDDDDD

upvoted 6 times

omoakin 2 years, 6 months ago

BBBBBBBBBBBBBB

upvoted 2 times

elmogy 2 years, 6 months ago

Selected Answer: D

company need to setup a cheaper connection (200 M) but B is incorrect because you can only order port speeds of 1, 10, or 100 Gbps for more flexibility you can go with hosted connection, You can order port speeds between 50 Mbps and 10 Gbps.

<https://docs.aws.amazon.com/whitepapers/latest/aws-vpc-connectivity-options/aws-direct-connect.html>

upvoted 5 times

cloudenthusiast 2 years, 6 months ago

Selected Answer: B

By opting for a lower capacity 200 Mbps connection instead of the 1 Gbps connection, the company can significantly reduce costs. This solution ensures a dedicated and secure connection while aligning with the company's low utilization, resulting in cost savings.

upvoted 4 times

pentium75 1 year, 11 months ago

But 200M cannot be ordered through Management Console, only partners.

upvoted 3 times

nosense 2 years, 6 months ago

Selected Answer: D

A hosted connection is a lower-cost option that is offered by AWS Direct Connect Partners

upvoted 5 times

Efren 2 years, 6 months ago

Also, there are not 200 MBps direct connection speed.

upvoted 1 times

nosense 2 years, 6 months ago

Hosted Connection 50 Mbps, 100 Mbps, 200 Mbps,
Dedicated Connection 1 Gbps, 10 Gbps, and 100 Gbps

B would require the company to purchase additional hardware or software

upvoted 3 times

A company has multiple Windows file servers on premises. The company wants to migrate and consolidate its files into an Amazon FSx for Windows File Server file system. File permissions must be preserved to ensure that access rights do not change.

Which solutions will meet these requirements? (Choose two.)

A. Deploy AWS DataSync agents on premises. Schedule DataSync tasks to transfer the data to the FSx for Windows File Server file system. **Most Voted**

B. Copy the shares on each file server into Amazon S3 buckets by using the AWS CLI. Schedule AWS DataSync tasks to transfer the data to the FSx for Windows File Server file system.

C. Remove the drives from each file server. Ship the drives to AWS for import into Amazon S3. Schedule AWS DataSync tasks to transfer the data to the FSx for Windows File Server file system.

D. Order an AWS Snowcone device. Connect the device to the on-premises network. Launch AWS DataSync agents on the device. Schedule DataSync tasks to transfer the data to the FSx for Windows File Server file system. **Most Voted**

E. Order an AWS Snowball Edge Storage Optimized device. Connect the device to the on-premises network. Copy data to the device by using the AWS CLI. Ship the device back to AWS for import into Amazon S3. Schedule AWS DataSync tasks to transfer the data to the FSx for Windows File Server file system.

Correct Answer: AD

Community vote distribution

AD (96%)

4%

Comments

pentium75 **Highly Voted** 1 year, 11 months ago

Selected Answer: AD

B, C and E would copy the files to S3 first where permissions would be lost
upvoted 10 times

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: AD

A This option involves deploying DataSync agents on your on-premises file servers and using DataSync to transfer the data directly to the FSx for Windows File Server. DataSync ensures that file permissions are preserved during the migration process.

D

This option involves using an AWS Snowcone device, a portable data transfer device. You would connect the Snowcone device to your on-premises network, launch DataSync agents on the device, and schedule DataSync tasks to transfer the data to FSx for Windows File Server. DataSync handles the migration process while preserving file permissions.

upvoted 8 times

Guru4Cloud **Most Recent** 2 years, 3 months ago

Selected Answer: BD

Why not - BD?

upvoted 1 times

Guru4Cloud 2 years, 3 months ago

° This option uses S3 as an intermediary, ensuring that file permissions are preserved during the initial data copy. DataSync can then transfer the data from S3 to FSx while maintaining the permissions.

° This option uses a Snowcone device with DataSync agents to replicate the on-premises permission structure directly to FSx. This approach is suitable for maintaining file permissions during migration.

upvoted 2 times

pentium75 1 year, 11 months ago

B copies the data to S3 first where file permissions would be lost.

upvoted 3 times

elmogy 2 years, 6 months ago

Selected Answer: AD

the key is file permissions are preserved during the migration process. only datasync supports that

upvoted 5 times

coolkidsclubvip 2 years, 3 months ago

Bro, all 5 answers mentioned DataSync.....

upvoted 5 times

Devsin2000 2 years, 2 months ago

Yes but AD have only DataSync, whereas others have others have AWS CLI used.

upvoted 3 times

nosense 2 years, 6 months ago

Selected Answer: AD

Option B would require copy the data to Amazon S3 before transferring it to Amazon FSx for Windows File Server

Option C would require the company to remove the drives from each file server and ship them to AWS

upvoted 2 times

barracouto 2 years, 3 months ago

Also, S3 doesn't retain permissions because it isn't a file system.

upvoted 4 times

A company wants to ingest customer payment data into the company's data lake in Amazon S3. The company receives payment data every minute on average. The company wants to analyze the payment data in real time. Then the company wants to ingest the data into the data lake.

Which solution will meet these requirements with the MOST operational efficiency?

- A. Use Amazon Kinesis Data Streams to ingest data. Use AWS Lambda to analyze the data in real time.
- B. Use AWS Glue to ingest data. Use Amazon Kinesis Data Analytics to analyze the data in real time.
- C. Use Amazon Kinesis Data Firehose to ingest data. Use Amazon Kinesis Data Analytics to analyze the data in real time.
- D. Use Amazon API Gateway to ingest data. Use AWS Lambda to analyze the data in real time.

Most Voted

Correct Answer: C

Community vote distribution

C (88%)

12%

Comments

Axeashes **Highly Voted** 2 years, 5 months ago

Kinesis Data Firehose is near real time (min. 60 sec). - The question is focusing on real time processing/analysis + efficiency -> Kinesis Data Stream is real time ingestion.

<https://www.amazonaws.cn/en/kinesis/data-firehose/#:~:text=Near%20real%2Dtime,is%20sent%20to%20the%20service.>

upvoted 12 times

Axeashes 2 years, 5 months ago

Unless the intention is real time analytics not real time ingestion !

upvoted 6 times

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: C

By leveraging the combination of Amazon Kinesis Data Firehose and Amazon Kinesis Data Analytics, you can efficiently ingest and analyze the payment data in real time without the need for manual processing or additional infrastructure management. This solution provides a streamlined and scalable approach to handle continuous data ingestion and analysis requirements.

upvoted 11 times

In2718 **Most Recent** 9 months, 3 weeks ago

Selected Answer: A

Wow, going against the grain... Chat AI says answer A.

A. Use Amazon Kinesis Data Streams to ingest data. Use AWS Lambda to analyze the data in real time.

Why this is the most operationally efficient among the choices:

Serverless + native streaming: Kinesis Data Streams (use on-demand mode) + Lambda requires no capacity planning and scales automatically. Lambda can write the analyzed results straight to Amazon S3.

Option C is invalid: Kinesis Data Analytics does not read from Firehose; it reads from Kinesis Data Streams or Kafka. So "Firehose to ingest + KDA to analyze" won't wire up directly.

Option B: AWS Glue is primarily ETL/batch (Glue Streaming still needs a stream source; it doesn't ingest by itself).

Option D: API Gateway + Lambda works for REST events but adds cost/limits and isn't purpose-built for streaming ingestion.

upvoted 3 times

zdi561 1 year, 4 months ago

Selected Answer: C

The question says that the data is analyzed first then ingest into s3 for data lake
upvoted 1 times

wizcloudifa 1 year, 7 months ago

Selected Answer: C

Kinesis Firehouse = ingesting
Kinesis Datastreams = storing
Kinesis analytics = doing analysis
upvoted 6 times

awsgeek75 1 year, 11 months ago

Selected Answer: C

Data is stored on S3 so real-time data analytics can be done with Kinesis Data Analytics which rules out Lambda solutions (A and D) as they are more operationally complex.
B is not useful it is more of ETL.

Firehose is actually to distribute data but given that company is already receiving data somehow so Firehose can basically distribute it to S3 with minimum latency. I have to admit this was confusing. I would have used Kinesis Streams to store on S3 and Data Analytics but combination is confusing!

upvoted 3 times

Mr_Marcus 1 year, 6 months ago

"Data is stored on S3..."

Nope. Re-read the first sentence. S3 is the destination, not the source.
The task is to ingest, analyze in real time, and store in S3.

upvoted 2 times

1rob 2 years ago

Selected Answer: C

"payment data every minute on average" is a good-to-go- for firehose.
Also firehose is more operational efficient compared to Data Streams.

upvoted 3 times

lucasbg 2 years ago

Selected Answer: A

I think this is A. The purpose of Firehose is to ingest and deliver to a data store, not to an analytics service. And in fact you can use lambda for real time analysis, such I find A more aligned.

upvoted 2 times

pentium75 1 year, 11 months ago

But developing and maintaining a custom Lambda function "to analyze the data in real time" is surely not as 'operationally efficient' as using an existing service such as Kinesis Data Analytics.

upvoted 4 times

DDongi 2 years, 1 month ago

Firehose has a 60 sec delay so real time analytics should be without real time data isn't that problematic? Why would you have then real time analytics then in the first place?

upvoted 2 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: C

Kinesis Data Streams focuses on ingesting and storing data streams while Kinesis Data Firehose focuses on delivering data streams to select destinations, as the motive of the question is to do analytics, the answer should be C.

upvoted 4 times

hsinchang 2 years, 4 months ago

Selected Answer: C

Kinesis Data Streams focuses on ingesting and storing data streams while Kinesis Data Firehose focuses on delivering data streams to select destinations, as the motive of the question is to do analytics, the answer should be C.

upvoted 2 times

james2033 2 years, 4 months ago

Selected Answer: C

Quote "Connect with 30+ fully integrated AWS services and streaming destinations such as Amazon Simple Storage Service (S3)" at <https://aws.amazon.com/kinesis/data-firehose/> . Amazon Kinesis Data Analytics <https://aws.amazon.com/kinesis/data-analytics/>

upvoted 2 times

TariqKipkemei 2 years, 4 months ago

Selected Answer: C

Use Kinesis Firehose to capture and deliver the data to Kinesis Analytics to perform analytics.

upvoted 2 times

Anmol_1010 2 years, 6 months ago

Did anyone took tge exam recently,
How many questiona were there

upvoted 2 times

omoakin 2 years, 6 months ago

Can we understand why admin's answers are mostly wrong? Or is this done on purpose?

upvoted 3 times

nosense 2 years, 6 months ago

Selected Answer: C

Amazon Kinesis Data Firehose the most optimal variant

upvoted 4 times

kailu 2 years, 6 months ago

Shouldn't C be more appropriate?

upvoted 5 times

MostofMichelle 2 years, 6 months ago

You're right. I believe the answers are wrong on purpose, so good thing votes can be made on answers and discussions are allowed.

upvoted 2 times

A company runs a website that uses a content management system (CMS) on Amazon EC2. The CMS runs on a single EC2 instance and uses an Amazon Aurora MySQL Multi-AZ DB instance for the data tier. Website images are stored on an Amazon Elastic Block Store (Amazon EBS) volume that is mounted inside the EC2 instance.

Which combination of actions should a solutions architect take to improve the performance and resilience of the website? (Choose two.)

- A. Move the website images into an Amazon S3 bucket that is mounted on every EC2 instance
- B. Share the website images by using an NFS share from the primary EC2 instance. Mount this share on the other EC2 instances.
- C. Move the website images onto an Amazon Elastic File System (Amazon EFS) file system that is mounted on every EC2 instance. **Most Voted**
- D. Create an Amazon Machine Image (AMI) from the existing EC2 instance. Use the AMI to provision new instances behind an Application Load Balancer as part of an Auto Scaling group. Configure the Auto Scaling group to maintain a minimum of two instances. Configure an accelerator in AWS Global Accelerator for the website
- E. Create an Amazon Machine Image (AMI) from the existing EC2 instance. Use the AMI to provision new instances behind an Application Load Balancer as part of an Auto Scaling group. Configure the Auto Scaling group to maintain a minimum of two instances. Configure an Amazon CloudFront distribution for the website. **Most Voted**

Correct Answer: CE

Community vote distribution

CE (75%)

AE (25%)

Comments

cloudenthusiast **Highly Voted** 3 years ago

Selected Answer: CE

By combining the use of Amazon EFS for shared file storage and Amazon CloudFront for content delivery, you can achieve improved performance and resilience for the website.

upvoted 14 times

Abrar2022 **Highly Voted** 3 years ago

You don't mount S3

upvoted 7 times

wizcloudifa **Most Recent** 2 years, 1 month ago

Selected Answer: CE

First of all you should understand, a website using CMS is a dynamic one not static, so A is out, B is more complicated than C, so C, and between global accelerator and cloudfront, Cloudfront suits better as there is no legacy protocols data(UDP, etc) that needs to be accessed, hence E

upvoted 5 times

foha2012 2 years, 5 months ago

I choose AE. Although I don't know if s3 can be mounted on ec2 ?? Maybe wrong wording. Efs is a better choice but its not a natural selection for strong images.

upvoted 2 times

allanm 10 months ago

Read the question - it asks for a combination of 2 solutions. S3 is out because the architecture is AMI/EC2 driven and doesn't state that it needs S3 to drive the storage layer, rather EFS so A is out

upvoted 1 times

awsgeek75 2 years, 5 months ago

I made the same mistake but mounting S3 on EC2 is a painful operation so EFS makes more sense (C).
Option E takes care of caching static images on CDN so that problem is solved along with resilience etc.

upvoted 4 times

pentium75 2 years, 5 months ago

Selected Answer: CE

Not A because you can't mount an S3 bucket on an EC2 instance. You could use a file gateway and share an S3 bucket via NFS and mount that on EC2, but that is not mentioned here and would also not make sense.

upvoted 4 times

seldiora 1 year, 8 months ago

it is possible to mount s3 to EC2, just quite difficult: <https://aws.amazon.com/blogs/storage/mounting-amazon-s3-to-an-amazon-ec2-instance-using-a-private-connection-to-s3-file-gateway/>

upvoted 2 times

potomac 2 years, 7 months ago

Selected Answer: CE

You can mount EFS file systems to multiple Amazon EC2 instances remotely and securely without having to log in to the instances by using the AWS Systems Manager Run Command.

upvoted 3 times

wsdasdasdqwdaw 2 years, 7 months ago

A is out of the game for sure. Mount S3 to EC2 ... madness. The question is CE or DE, but it is CE because of AWS Global Accelerator is match with NLB, not ALB as it is stated in option D, thus CE as many of all here noted.

upvoted 5 times

thanhnv142 2 years, 7 months ago

A and E is correct. We have a cloud front + S3 combo

upvoted 1 times

wsdasdasdqwdaw 2 years, 7 months ago

S3 can't be mounted on EC2 it is not A for sure!

upvoted 4 times

NickGordon 2 years, 7 months ago

<https://aws.amazon.com/blogs/storage/mounting-amazon-s3-to-an-amazon-ec2-instance-using-a-private-connection-to-s3-file-gateway/>

upvoted 3 times

pentium75 2 years, 5 months ago

"Using a ... S3 file gateway"

upvoted 2 times

thanhnv142 2 years, 7 months ago

A and E.

C is not correct because You dont mount a new EFS onto existing EC2. If you do that, you have to migrate all existing data in EBS into EFS. Then remove all the EBS. Should never do this.

upvoted 1 times

pentium75 2 years, 5 months ago

I can't follow. EFS provides NFS mount points, how can you not mount those onto existing EC2?

upvoted 2 times

franbarberan 2 years, 8 months ago

Selected Answer: CE

<https://bluexp.netapp.com/blog/ebs-efs-amazons3-best-cloud-storage-system>

upvoted 2 times

Smart 2 years, 9 months ago

Selected Answer: CE

Not A - S3 cannot be mounted (up until few months ago). Exam does not test for the updates in last 6 months.

upvoted 4 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: AE

You have summarized the reasons why options A and E are the best choices very well.

Migrating static website assets like images to Amazon S3 enables high scalability, durability and shared access across instances. This improves performance.

Using Auto Scaling with load balancing provides elasticity and resilience. Adding a CloudFront distribution further boosts performance through caching and content delivery.

upvoted 2 times

pentium75 2 years, 5 months ago

You can't directly mount an S3 bucket on EC2.

upvoted 2 times

Ale1973 2 years, 10 months ago

Selected Answer: AE

Both options AE and CE would work, but I choose AE, because, on my opinion, S3 is best suited for performance and resilience.

upvoted 3 times

pentium75 2 years, 5 months ago

You can't directly mount an S3 bucket on EC2

upvoted 2 times

MicketyMouse 2 years, 10 months ago

Selected Answer: CE

EFS, unlike EBS, can be mounted across multiple EC2 instances and hence C over A.

upvoted 2 times

TariqKipkemei 2 years, 11 months ago

Selected Answer: AE

Technically both options AE and CE would work. But S3 is best suited for unstructured data, and the key benefit of mounting S3 on EC2 is that it provides a cost-effective alternative of using object storage for applications dealing with large files, as compared to expensive file or block storage. At the same time it provides more performant, scalable and highly available storage for these applications.

Even though there is no mention of 'cost efficient' in this question, in the real world cost is the no.1 factor.

In the exam I believe both options would be a pass.

<https://aws.amazon.com/blogs/storage/mounting-amazon-s3-to-an-amazon-ec2-instance-using-a-private-connection-to-s3-file-gateway/>

upvoted 4 times

pentium75 2 years, 5 months ago

You can't directly mount an S3 bucket on EC2, only through file gateway

upvoted 1 times

AshutoshSingh1923 2 years, 11 months ago

Selected Answer: CE

Option C provides moving the website images onto an Amazon EFS file system that is mounted on every EC2 instance. Amazon EFS provides a scalable and fully managed file storage solution that can be accessed concurrently from multiple EC2 instances. This ensures that the website images can be accessed efficiently and consistently by all instances, improving performance

In Option E The Auto Scaling group maintains a minimum of two instances, ensuring resilience by automatically replacing any unhealthy instances. Additionally, configuring an Amazon CloudFront distribution for the website further improves performance by caching content at edge locations closer to the end-users, reducing latency and improving content delivery.

Hence combining these actions, the website's performance is improved through efficient image storage and content delivery

upvoted 3 times

Vadbro7 2 years, 11 months ago

Which answer is correct? the most voted ones or the Suggested answers?

upvoted 1 times

9be0170 1 year, 11 months ago

the most voted always

upvoted 2 times


A company runs an infrastructure monitoring service. The company is building a new feature that will enable the service to monitor data in customer AWS accounts. The new feature will call AWS APIs in customer accounts to describe Amazon EC2 instances and read Amazon CloudWatch metrics.

What should the company do to obtain access to customer accounts in the MOST secure way?

- A. Ensure that the customers create an IAM role in their account with read-only EC2 and CloudWatch permissions and a trust policy to the company's account. **Most Voted**
- B. Create a serverless API that implements a token vending machine to provide temporary AWS credentials for a role with read-only EC2 and CloudWatch permissions.
- C. Ensure that the customers create an IAM user in their account with read-only EC2 and CloudWatch permissions. Encrypt and store customer access and secret keys in a secrets management system.
- D. Ensure that the customers create an Amazon Cognito user in their account to use an IAM role with read-only EC2 and CloudWatch permissions. Encrypt and store the Amazon Cognito user and password in a secrets management system.

Correct Answer: A

Community vote distribution



Comments

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: A

By having customers create an IAM role with the necessary permissions in their own accounts, the company can use AWS Identity and Access Management (IAM) to establish cross-account access. The trust policy allows the company's AWS account to assume the customer's IAM role temporarily, granting access to the specified resources (EC2 instances and CloudWatch metrics) within the customer's account. This approach follows the principle of least privilege, as the company only requests the necessary permissions and does not require long-term access keys or user credentials from the customers.

upvoted 16 times

Piccalo **Highly Voted** 2 years, 6 months ago

Selected Answer: A

A. Roles give temporary credentials

upvoted 9 times

Efren 2 years, 6 months ago

Agreed . Role is the keyword

upvoted 3 times

awsgeek75 **Most Recent** 1 year, 10 months ago

Selected Answer: A

B: Sharing credentials, even temporary, is insecure

C: Access and secret keys. That won't work and sharing secrets outside of account is not secure for this use case

A: Keyword "trust policy"

D: Again, sharing username and pwd and sharing in any way is not secure

upvoted 2 times

pentium75 1 year, 11 months ago

Selected Answer: A

Not B (would be about access to the company's account, not the customers' accounts)

Not C (storing credentials in a custom system is a big nono)

Not D (Cognito has nothing to do here and "user and password" is terrible)

upvoted 3 times

1rob 2 years ago

Selected Answer: D

The company's infrastructure monitoring service needs to call aws API's in the MOST secure way. So you have to focus on restricting access to the APIs and there is where cognito comes in to play.

upvoted 2 times

pentium75 1 year, 11 months ago

What is unsecure with A?

upvoted 2 times

1rob 1 year, 11 months ago

The company runs an infrastructure monitoring service. Nowhere is stated that this service lives in an aws account. So A and C I wouldn't choose. B is a bit too vague. So I end up with D.

upvoted 1 times

awsgeek75 1 year, 11 months ago

Are you suggesting to restrict CloudWatch API with Cognito roles?

upvoted 2 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: A

A is the most secure approach for accessing customer accounts.

Having customers create a cross-account IAM role with the appropriate permissions, and configuring the trust policy to allow the monitoring service principal account access, implements secure delegation and least privilege access.

upvoted 2 times

A company needs to connect several VPCs in the us-east-1 Region that span hundreds of AWS accounts. The company's networking team has its own AWS account to manage the cloud network.

What is the MOST operationally efficient solution to connect the VPCs?

- A. Set up VPC peering connections between each VPC. Update each associated subnet's route table
- B. Configure a NAT gateway and an internet gateway in each VPC to connect each VPC through the internet
- C. Create an AWS Transit Gateway in the networking team's AWS account. Configure static routes from each VPC. **Most Voted**
- D. Deploy VPN gateways in each VPC. Create a transit VPC in the networking team's AWS account to connect to each VPC.

Correct Answer: C

Community vote distribution

C (100%)

Comments

hsinchang **Highly Voted** 2 years, 4 months ago

Selected Answer: C

The main difference between AWS Transit Gateway and VPC peering is that AWS Transit Gateway is designed to connect multiple VPCs together in a hub-and-spoke model, while VPC peering is designed to connect two VPCs together in a peer-to-peer model.

As we have several VPCs here, the answer should be C.

upvoted 20 times

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: C

AWS Transit Gateway is a highly scalable and centralized hub for connecting multiple VPCs, on-premises networks, and remote networks. It simplifies network connectivity by providing a single entry point and reducing the number of connections required. In this scenario, deploying an AWS Transit Gateway in the networking team's AWS account allows for efficient management and control over the network connectivity across multiple VPCs.

upvoted 7 times

awsgeek75 **Most Recent** 1 year, 10 months ago

Selected Answer: C

A: This option is suggesting hundreds of peering connection for EACH VPC. Nope!

B: NAT gateway is for network translation not VPC interconnectivity so this is wrong

C: Transit GW + static routes will connect all VPCs <https://aws.amazon.com/transit-gateway/>

D: VPN gateway is for on-prem to VPN for a VPC. There is no on-prem here so this is wrong

upvoted 3 times

TariqKipkemei 2 years, 1 month ago

Selected Answer: C

Connect, Monitor and Manage Multiple VPCs in one place = AWS Transit Gateway

upvoted 3 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: C

C is the most operationally efficient solution for connecting a large number of VPCs across accounts.

Using AWS Transit Gateway allows all the VPCs to connect to a central hub without needing to create a mesh of VPC peering connections between each VPC pair.

This significantly reduces the operational overhead of managing the network topology as new VPCs are added or changed.

The networking team can centrally manage the Transit Gateway routing and share it across accounts using Resource Access Manager.

upvoted 3 times

MirKhobaeb 2 years, 6 months ago

Answer is C

upvoted 2 times

MirKhobaeb 2 years, 6 months ago

A transit gateway is a network transit hub that you can use to interconnect your virtual private clouds (VPCs) and on-premises networks. As your cloud infrastructure expands globally, inter-Region peering connects transit gateways together using the AWS Global Infrastructure. Your data is automatically encrypted and never travels over the public internet.

upvoted 3 times

nosense 2 years, 6 months ago

Selected Answer: C

I voted for c

upvoted 3 times

nosense 2 years, 6 months ago

An AWS Transit Gateway is a highly scalable and secure way to connect VPCs in multiple AWS accounts. It is a central hub that routes traffic between VPCs, on-premises networks, and AWS services.

upvoted 4 times

A company has Amazon EC2 instances that run nightly batch jobs to process data. The EC2 instances run in an Auto Scaling group that uses On-Demand billing. If a job fails on one instance, another instance will reprocess the job. The batch jobs run between 12:00 AM and 06:00 AM local time every day.

Which solution will provide EC2 instances to meet these requirements MOST cost-effectively?

- A. Purchase a 1-year Savings Plan for Amazon EC2 that covers the instance family of the Auto Scaling group that the batch job uses.
- B. Purchase a 1-year Reserved Instance for the specific instance type and operating system of the instances in the Auto Scaling group that the batch job uses.
- C. Create a new launch template for the Auto Scaling group. Set the instances to Spot Instances. Set a policy to scale out based on CPU usage. **Most Voted**
- D. Create a new launch template for the Auto Scaling group. Increase the instance size. Set a policy to scale out based on CPU usage.

Correct Answer: C

Community vote distribution

C (100%)

Comments

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: C

Purchasing a 1-year Savings Plan (option A) or a 1-year Reserved Instance (option B) may provide cost savings, but they are more suitable for long-running, steady-state workloads. Since your batch jobs run for a specific period each day, using Spot Instances with the ability to scale out based on CPU usage is a more cost-effective choice.

upvoted 15 times

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: C

C is the most cost-effective solution in this scenario.

Using Spot Instances allows EC2 capacity to be purchased at significant discounts compared to On-Demand prices. The auto scaling group can scale out to add Spot Instances when needed for the batch jobs.

If Spot Instances become unavailable, regular On-Demand Instances will be launched instead to maintain capacity. The potential for interruptions is acceptable since failed jobs can be re-run.

upvoted 7 times

sandordini **Most Recent** 1 year, 7 months ago

Selected Answer: C

Stateless, most cost-effective >> Spot

upvoted 3 times

awsgeek75 1 year, 11 months ago

Selected Answer: C

You don't need any scaling really as the job runs on another EC2 instance if it fails on first one. A. B. D are all more expensive than C due to spot instance being cheaper than reserved instances.

upvoted 3 times

TariqKipkemei 2 years, 4 months ago

Selected Answer: C

Spot Instances to the rescue....whoosh

upvoted 2 times

wRhIH 2 years, 5 months ago

" If a job fails on one instance, another instance will reprocess the job". This ensures Spot Instances are enough for this case

upvoted 5 times

Abrar2022 2 years, 6 months ago

Selected Answer: C

Since your batch jobs run for a specific period each day, using Spot Instances with the ability to scale out based on CPU usage is a more cost-effective choice.

upvoted 2 times

Blingy 2 years, 6 months ago

C FOR ME COS OF SPOT INSTACES

upvoted 3 times

udo2020 2 years, 6 months ago

First I think it is B but because of cost saving I think it should be C spot instances.

upvoted 2 times

nosense 2 years, 6 months ago

Selected Answer: C

c for me

upvoted 2 times


A social media company is building a feature for its website. The feature will give users the ability to upload photos. The company expects significant increases in demand during large events and must ensure that the website can handle the upload traffic from users.

Which solution meets these requirements with the MOST scalability?

- A. Upload files from the user's browser to the application servers. Transfer the files to an Amazon S3 bucket.
- B. Provision an AWS Storage Gateway file gateway. Upload files directly from the user's browser to the file gateway.
- C. Generate Amazon S3 presigned URLs in the application. Upload files directly from the user's browser into an S3 bucket.
Most Voted
- D. Provision an Amazon Elastic File System (Amazon EFS) file system. Upload files directly from the user's browser to the file system.

Correct Answer: C

Community vote distribution



Comments

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: C

This approach allows users to upload files directly to S3 without passing through the application servers, reducing the load on the application and improving scalability. It leverages the client-side capabilities to handle the file uploads and offloads the processing to S3.

upvoted 16 times

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: C

C is the best solution to meet the scalability requirements.

Generating S3 presigned URLs allows users to upload directly to S3 instead of application servers. This removes the application servers as a bottleneck for upload traffic.

S3 can scale to handle very high volumes of uploads with no limits on storage or throughput. Using presigned URLs leverages this scalability.

upvoted 5 times

hro **Most Recent** 1 year, 8 months ago

C - You may use presigned URLs to allow someone to upload an object to your Amazon S3 bucket. Using a presigned URL will allow an upload without requiring another party to have AWS security credentials or permissions.

upvoted 3 times

awsgeek75 1 year, 11 months ago

Selected Answer: C

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/using-presigned-url.html>

"You can also use presigned URLs to allow someone to upload a specific object to your Amazon S3 bucket. This allows an upload without requiring another party to have AWS security credentials or permissions. "

upvoted 2 times

Goutham4981 2 years ago

Selected Answer: A

S3 presigned url is used for sharing objects from an s3 bucket and not for uploading to an s3 bucket

upvoted 3 times

Murtadhaceit 2 years ago

No. It allows to download and upload.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/using-presigned-url.html>

upvoted 4 times

TariqKipkemei 2 years, 4 months ago

Selected Answer: C

You may use presigned URLs to allow someone to upload an object to your Amazon S3 bucket. Using a presigned URL will allow an upload without requiring another party to have AWS security credentials or permissions.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/PresignedUrlUploadObject.html>

upvoted 2 times

baba365 2 years, 5 months ago

Hello Moderator. This question and answer should be rephrased because:

1. S3 pre-signed URLs are used to share objects FROM S3 buckets
2. How scalable are pre-signed URLs when they are time constrained?

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/ShareObjectPreSignedURL.html>

upvoted 2 times

pentium75 1 year, 11 months ago

Both is wrong

Presigned URLs can be used for upload

The solution is scalable because you can issue thousands of pre-signed URLs, and thousands of users can upload images at the same time.

User wants to upload picture -> server generates presigned URL and sends it to the app -> app uploads file

upvoted 3 times

nosense 2 years, 6 months ago

Selected Answer: C

the most scalable because it allows users to upload files directly to Amazon S3,

upvoted 4 times

A company has a web application for travel ticketing. The application is based on a database that runs in a single data center in North America. The company wants to expand the application to serve a global user base. The company needs to deploy the application to multiple AWS Regions. Average latency must be less than 1 second on updates to the reservation database.

The company wants to have separate deployments of its web platform across multiple Regions. However, the company must maintain a single primary reservation database that is globally consistent.

Which solution should a solutions architect recommend to meet these requirements?

A. Convert the application to use Amazon DynamoDB. Use a global table for the center reservation table. Use the correct Regional endpoint in each Regional deployment. **Most Voted**

B. Migrate the database to an Amazon Aurora MySQL database. Deploy Aurora Read Replicas in each Region. Use the correct Regional endpoint in each Regional deployment for access to the database.

C. Migrate the database to an Amazon RDS for MySQL database. Deploy MySQL read replicas in each Region. Use the correct Regional endpoint in each Regional deployment for access to the database.

D. Migrate the application to an Amazon Aurora Serverless database. Deploy instances of the database to each Region. Use the correct Regional endpoint in each Regional deployment to access the database. Use AWS Lambda functions to process event streams in each Region to synchronize the databases.

Correct Answer: A

Community vote distribution

A (52%)

B (48%)

Comments

cloudehustia **Highly Voted** 3 years ago

Selected Answer: A

Using DynamoDB's global tables feature, you can achieve a globally consistent reservation database with low latency on updates, making it suitable for serving a global user base. The automatic replication provided by DynamoDB eliminates the need for manual synchronization between Regions.
upvoted 24 times

FlyingHawk 1 year, 6 months ago

DynamoDB Global Tables rely on eventual consistency for replication between Regions. This means that updates in one Region may take milliseconds to seconds to propagate to other Regions.

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ReadConsistency.html>

The requirement specifies global consistency, which DynamoDB cannot guarantee.

upvoted 1 times

upliftinghut **Highly Voted** 2 years, 4 months ago

Selected Answer: B

Aurora: less than 1 second: <https://aws.amazon.com/rds/aurora/global-database/>

DynamoDB: from 0.5 to 2.5 second: https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/V2globaltables_HowItWorks.html

upvoted 6 times

TheLaPlanta 2 years, 2 months ago

B doesn't say Aurora Global

upvoted 13 times

MatAlves 1 year, 9 months ago

DynamoDB doesn't meet the <1s req though.

upvoted 2 times

JustEugen Most Recent 8 months, 3 weeks ago

Selected Answer: A

I think the correct answer is A, as A and B are missing details.

If you follow the answer, neither A nor B is a solution:

- In A, it is not specified if the Aurora Global Database will be used
- In B, it uses Dynamo DB global table with eventual consistency mode, which won't guarantee that data will be written, cause it can be rejected in case of conflict

But if we think about details:

- For A, use DynamoDB Multi-Region strong consistency (MRSC)
- For B, use the Aurora Global Database

The best option will be B, as it involves only a single writer, which aligns with our request for global consistency. With DynamoDB MRSC, we will have a significant overhead of handling possible write conflicts

upvoted 1 times

JustEugen 8 months, 3 weeks ago

I am sorry, not A, the correct is B

upvoted 1 times

In2718 9 months, 3 weeks ago

Selected Answer: B

Single primary, globally consistent writes: Aurora Global Database uses one writer Region and read-only secondary Regions, avoiding multi-writer conflicts.

< 1s average update latency cross-Region: Storage-level replication typically delivers sub-second lag from the primary to secondary Regions—ideal for a global web tier reading locally while writing to the primary.

Separate web stacks per Region: Point each Regional deployment to its local Aurora reader endpoint for reads; route writes to the primary Region.

upvoted 1 times

Charlesvg 11 months, 2 weeks ago

Selected Answer: A

Option B provides only READ replicas globally. Which does not satisfies the UPDATES in less that 1 seconds for the application deployed on every region

upvoted 1 times

eduvaldi 1 year ago

Selected Answer: B

just for latency Its B

upvoted 1 times

tch 1 year, 3 months ago

Selected Answer: A

answer should be A.

we can consider if B change to "Amazon Aurora Global Database"...

upvoted 1 times

tch 1 year, 3 months ago

Global tables - multi-Region replication for DynamoDB

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/GlobalTables.htm> Amazon DynamoDB global tables are a fully managed, multi-Region, and multi-active database option that delivers fast and localized read and write performance for massively scaled global applications.

upvoted 1 times

Dantecito 1 year, 4 months ago

Selected Answer: A

A is the only option that allows you to write from any reagon. I understand the < 1 second latency, but DynamoDB global tables use a last-writer-wins reconciliation between concurrent updates.

B and C only supports read scaling, not low-latency global writes. It will be much much much much than 1 second to write on the databse.

D I stopped reading at lambda so NO!.

upvoted 1 times

zdi561 1 year, 4 months ago

Selected Answer: A

Dynamodb is consistent and read/write in all node, Aurora db , only the master is writable. For reservation , writing is criticle.

upvoted 1 times

khatingarun 1 year, 4 months ago

Selected Answer: A

A is the correct answer
upvoted 1 times

zdi561 1 year, 4 months ago

Selected Answer: B

dynamodb is not globally consistent. DynamoDB doesn't support strongly consistent reads across Regions.
https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/V2globaltables_HowItWorks.html
upvoted 2 times

hilker1983 1 year, 5 months ago

Selected Answer: B

correct answer bbbb
upvoted 1 times

EllenLiu 1 year, 5 months ago

Selected Answer: A

A: dynamoDB global table preview strong consistent
<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/multi-region-strong-consistency-gt.html>
upvoted 1 times

Penjerla 1 year, 6 months ago

Selected Answer: A

Aurora replicas have to be in the same region as primary
<https://docs.aws.amazon.com/prescriptive-guidance/latest/aurora-replication-options/compare-solutions.html>
upvoted 2 times

FlyingHawk 1 year, 6 months ago

Selected Answer: B

Since DynamoDB Global Tables allow writes in all Regions, conflicts can occur if the same item is updated simultaneously in multiple Regions. DynamoDB uses a "last writer wins" conflict resolution strategy, which may lead to data integrity issues:
https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/V2globaltables_HowItWorks.html#V2globaltables_HowItWorks.conflict-resolution
If you require less than 1-second latency for write updates across Regions, DynamoDB may not consistently meet this requirement due to the asynchronous nature of its replication.
For this scenario, DynamoDB Global Tables do not fully meet the requirements due to their reliance on eventual consistency. If strong consistency is a must, go with Amazon Aurora Global Database.
upvoted 2 times

MatAlves 1 year, 8 months ago

Selected Answer: B

The question asks "Average latency must be less than 1 second on updates to the reservation database."

A is incorrect:

"Changes to a DynamoDB global tables are replicated asynchronously, with typical latency of between 0.5 - 2.5 seconds between AWS Regions in the same geographic area."

B is the answer:

"All Aurora Replicas return the same data for query results with minimal replica lag. This lag is usually much less than 100 milliseconds after the primary instance has written an update."

upvoted 2 times

MatAlves 1 year, 9 months ago

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/Aurora.Replication.html>

<https://community.aws/content/2drxEND7MtTOb2bWs2J0NICGewP/ddb-globaltables-lag?lang=en>

upvoted 3 times

SVDK 2 years, 4 months ago

Selected Answer: A

How can you update your database in the different regions with read replicas? You need to be able to read and write to the database from the different regions.
upvoted 3 times

A company has migrated multiple Microsoft Windows Server workloads to Amazon EC2 instances that run in the us-west-1 Region. The company manually backs up the workloads to create an image as needed.

In the event of a natural disaster in the us-west-1 Region, the company wants to recover workloads quickly in the us-west-2 Region. The company wants no more than 24 hours of data loss on the EC2 instances. The company also wants to automate any backups of the EC2 instances.

Which solutions will meet these requirements with the LEAST administrative effort? (Choose two.)

A. Create an Amazon EC2-backed Amazon Machine Image (AMI) lifecycle policy to create a backup based on tags. Schedule the backup to run twice daily. Copy the image on demand.

B. Create an Amazon EC2-backed Amazon Machine Image (AMI) lifecycle policy to create a backup based on tags. Schedule the backup to run twice daily. Configure the copy to the us-west-2 Region. **Most Voted**

C. Create backup vaults in us-west-1 and in us-west-2 by using AWS Backup. Create a backup plan for the EC2 instances based on tag values. Create an AWS Lambda function to run as a scheduled job to copy the backup data to us-west-2.

D. Create a backup vault by using AWS Backup. Use AWS Backup to create a backup plan for the EC2 instances based on tag values. Define the destination for the copy as us-west-2. Specify the backup schedule to run twice daily. **Most Voted**

E. Create a backup vault by using AWS Backup. Use AWS Backup to create a backup plan for the EC2 instances based on tag values. Specify the backup schedule to run twice daily. Copy on demand to us-west-2.

Correct Answer: BD

Community vote distribution

BD (100%)

Comments

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: BD

Option B suggests using an EC2-backed Amazon Machine Image (AMI) lifecycle policy to automate the backup process. By configuring the policy to run twice daily and specifying the copy to the us-west-2 Region, the company can ensure regular backups are created and copied to the alternate region.

Option D proposes using AWS Backup, which provides a centralized backup management solution. By creating a backup vault and backup plan based on tag values, the company can automate the backup process for the EC2 instances. The backup schedule can be set to run twice daily, and the destination for the copy can be defined as the us-west-2 Region.

upvoted 10 times

cloudenthusiast 2 years, 6 months ago

Both options automate the backup process and include copying the backups to the us-west-2 Region, ensuring data resilience in the event of a disaster. These solutions minimize administrative effort by leveraging automated backup and copy mechanisms provided by AWS services.

upvoted 7 times

awsgeek75 **Highly Voted** 1 year, 10 months ago

Selected Answer: BD

LEAST admin overhead:

A: On demand so wrong

C: Lambda is overhead

E: On-demand is wrong

BD is the only choice. Although D seems to cover for B also, happy to be corrected.

upvoted 8 times

pmlabs Most Recent 2 years, 2 months ago

B D seems to meet the requirements fully
upvoted 2 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: BD

B and D are the options that meet the requirements with the least administrative effort.

B uses EC2 image lifecycle policies to automatically create AMIs of the instances twice daily and copy them to the us-west-2 region. This automates regional backups.

D leverages AWS Backup to define a backup plan that runs twice daily and copies backups to us-west-2. AWS Backup automates EC2 instance backups.

Together, these options provide automated, regional EC2 backup capabilities with minimal administrative overhead.

upvoted 2 times

TariqKipkemei 2 years, 4 months ago

Selected Answer: BD

options B and D will provide least administrative effort.

upvoted 2 times

antropaws 2 years, 6 months ago

Selected Answer: BD

I also vote B and D.

upvoted 2 times

nosense 2 years, 6 months ago

Selected Answer: BD

solutions are both automated and require no manual intervention to create or copy backups

upvoted 5 times

A company operates a two-tier application for image processing. The application uses two Availability Zones, each with one public subnet and one private subnet. An Application Load Balancer (ALB) for the web tier uses the public subnets. Amazon EC2 instances for the application tier use the private subnets.


Users report that the application is running more slowly than expected. A security audit of the web server log files shows that the application is receiving millions of illegitimate requests from a small number of IP addresses. A solutions architect needs to resolve the immediate performance problem while the company investigates a more permanent solution.

What should the solutions architect recommend to meet this requirement?

- A. Modify the inbound security group for the web tier. Add a deny rule for the IP addresses that are consuming resources.
- B. Modify the network ACL for the web tier subnets. Add an inbound deny rule for the IP addresses that are consuming resources. **Most Voted**
- C. Modify the inbound security group for the application tier. Add a deny rule for the IP addresses that are consuming resources.
- D. Modify the network ACL for the application tier subnets. Add an inbound deny rule for the IP addresses that are consuming resources.

Correct Answer: B

Community vote distribution



Comments

lucdt4 **Highly Voted** 2 years, 6 months ago

Selected Answer: B

A wrong because security group can't deny (only allow)
upvoted 29 times

cloudenthusiast **Highly Voted** 2 years, 6 months ago

Selected Answer: B

In this scenario, the security audit reveals that the application is receiving millions of illegitimate requests from a small number of IP addresses. To address this issue, it is recommended to modify the network ACL (Access Control List) for the web tier subnets.

By adding an inbound deny rule specifically targeting the IP addresses that are consuming resources, the network ACL can block the illegitimate traffic at the subnet level before it reaches the web servers. This will help alleviate the excessive load on the web tier and improve the application's performance.

upvoted 10 times

awsgeek75 **Most Recent** 1 year, 10 months ago

Selected Answer: B

A: Wrong as SG cannot deny. By default everything is deny in SG and you allow stuff
CD: App tier is not under attack so these are irrelevant options
B: Correct as NACL is exactly for this access control list to define rules for CIDR or IP addresses
upvoted 3 times

Besisco 1 year, 4 months ago

agree, even if an app tier was under attack, the block rule should be applied at a web tier ACL
upvoted 1 times

TariqKipkemei 2 years, 1 month ago

Selected Answer: B

Modify the network ACL for the web tier subnets. Add an inbound deny rule for the IP addresses that are consuming resources.

upvoted 3 times

potomac 2 years, 1 month ago

Selected Answer: B

A is wrong

Security groups act at the network interface level, not the subnet level, and they support Allow rules only.

upvoted 3 times

Devsin2000 2 years, 2 months ago

Selected Answer: A

The security Group can be applied to an ALB at web tier.

upvoted 1 times

Goutham4981 2 years ago

Security group can't deny.

upvoted 3 times

OSHOAIB 1 year, 11 months ago

Security group rules are always permissive; you can't create rules that deny access.

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/security-group-rules.html>

upvoted 3 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: B

Since the bad requests are targeting the web tier, adding ACL deny rules for those IP addresses on the web subnets will block the traffic before it reaches the instances.

Security group changes (Options A and C) would not be effective since the requests are not even reaching those resources.

Modifying the application tier ACL (Option D) would not stop the bad traffic from hitting the web tier.

upvoted 3 times

fakrap 2 years, 6 months ago

Selected Answer: B

A is wrong because you cannot put any deny in security group

upvoted 3 times

Rob1L 2 years, 6 months ago

Selected Answer: B

You cannot Deny on SG, so it's B

upvoted 6 times

nosense 2 years, 6 months ago

Selected Answer: A

Option B is not as effective as option A

upvoted 5 times

cloudenthusiast 2 years, 6 months ago

A and C out due to the fact that SG does not have deny on allow rules.

upvoted 3 times

y0 2 years, 6 months ago

Security group only have allow rules

upvoted 3 times

nosense 2 years, 6 months ago

yeah, my mistake. B should be

upvoted 2 times

A global marketing company has applications that run in the ap-southeast-2 Region and the eu-west-1 Region. Applications that run in a VPC in eu-west-1 need to communicate securely with databases that run in a VPC in ap-southeast-2.

Which network design will meet these requirements?

- A. Create a VPC peering connection between the eu-west-1 VPC and the ap-southeast-2 VPC. Create an inbound rule in the eu-west-1 application security group that allows traffic from the database server IP addresses in the ap-southeast-2 security group.
- B. Configure a VPC peering connection between the ap-southeast-2 VPC and the eu-west-1 VPC. Update the subnet route tables. Create an inbound rule in the ap-southeast-2 database security group that references the security group ID of the application servers in eu-west-1.
- C. Configure a VPC peering connection between the ap-southeast-2 VPC and the eu-west-1 VPC. Update the subnet route tables. Create an inbound rule in the ap-southeast-2 database security group that allows traffic from the eu-west-1 application server IP addresses. **Most Voted**
- D. Create a transit gateway with a peering attachment between the eu-west-1 VPC and the ap-southeast-2 VPC. After the transit gateways are properly peered and routing is configured, create an inbound rule in the database security group that references the security group ID of the application servers in eu-west-1.

Correct Answer: C

Community vote distribution

C (90%)

10%

Comments

VellaDevil **Highly Voted** 2 years, 11 months ago

Selected Answer: C

Answer: C --> "You cannot reference the security group of a peer VPC that's in a different Region. Instead, use the CIDR block of the peer VPC."
<https://docs.aws.amazon.com/vpc/latest/peering/vpc-peering-security-groups.html>

upvoted 52 times

MatAlves 1 year, 9 months ago

Wow, big thanks!

upvoted 2 times

hsinchang 2 years, 10 months ago

Thanks for this clarification!

upvoted 4 times

Axeashes **Highly Voted** 2 years, 12 months ago

Selected Answer: C

"You cannot reference the security group of a peer VPC that's in a different Region. Instead, use the CIDR block of the peer VPC."
<https://docs.aws.amazon.com/vpc/latest/peering/vpc-peering-security-groups.html>

upvoted 13 times

In2718 **Most Recent** 9 months, 3 weeks ago

Selected Answer: C

Improved answer - for eu-west-1 to ap-southeast-2 over peering, use subnet/VPC CIDRs (ideally via a prefix list) rather than individual IPs.

upvoted 1 times

potomac 2 years, 7 months ago

Selected Answer: C

After establishing the VPC peering connection, the subnet route tables need to be updated in both VPCs to route traffic to the other VPC's CIDR blocks through the peering connection.

upvoted 3 times

Bennyboy789 2 years, 9 months ago

Selected Answer: C

VPC Peering Connection: This allows communication between instances in different VPCs as if they are on the same network. It's a straightforward approach to connect the two VPCs.

Subnet Route Tables: After establishing the VPC peering connection, the subnet route tables need to be updated in both VPCs to route traffic to the other VPC's CIDR blocks through the peering connection.

Inbound Rule in Database Security Group: By creating an inbound rule in the ap-southeast-2 database security group that allows traffic from the eu-west-1 application server IP addresses, you ensure that only the specified application servers from the eu-west-1 VPC can access the database server in the ap-southeast-2 VPC.

upvoted 4 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: B

B) Configure VPC peering between ap-southeast-2 and eu-west-1 VPCs. Update routes. Allow traffic in ap-southeast-2 database SG from eu-west-1 application server SG.

This option establishes the correct network connectivity for the applications in eu-west-1 to reach the databases in ap-southeast-2:

VPC peering connects the two VPCs across regions - [https://docs.aws.amazon.com/vpc/latest/peering/what-is-vpc-peering.html#:~:text=You%20can%20create%20a%20VPC,%2DRegion%20VPC%20peering%20connection\).](https://docs.aws.amazon.com/vpc/latest/peering/what-is-vpc-peering.html#:~:text=You%20can%20create%20a%20VPC,%2DRegion%20VPC%20peering%20connection).)

Updating route tables enables routing between the VPCs

Security group rule allowing traffic from eu-west-1 application server SG to ap-southeast-2 database SG secures connectivity

upvoted 2 times

awsgeek75 2 years, 4 months ago

No, you cannot use a SG reference from another region so last part "Create an inbound rule in the ap-southeast-2 database security group that references the security group ID of the application servers in eu-west-1" cannot be setup. This is why B is wrong.

upvoted 3 times

Guru4Cloud 2 years, 9 months ago

Options A, C, D have flaws:

Option A peer direction is wrong

Option C opens databases to application server IP addresses rather than SG

Option D uses transit gateway which is unnecessary for just two VPCs

upvoted 3 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: C

Selected C but B can also work

upvoted 1 times

TariqKipkemei 2 years, 10 months ago

I just tried from the the console, You can specify the name or ID of another security group in the same region. To specify a security group in another AWS account (EC2-Classic only), prefix it with the account ID and a forward slash, for example: 111122223333/OtherSecurityGroup.

You can Specify a single IP address, or an IP address range in CIDR notation in the same/other region.

In the exam both option B and C would be a pass. In the real world both option will work.

upvoted 3 times

TariqKipkemei 2 years, 7 months ago

Correction, You cannot reference the security group of a peer VPC that's in a different Region. Instead, use the CIDR block of the peer VPC.

The C is the only option here.

<https://docs.aws.amazon.com/vpc/latest/peering/vpc-peering-security-groups.html#:~:text=You%20cannot-,reference,-the%20security%20group>

upvoted 4 times

awsgeek75 2 years, 4 months ago

This is why B is wrong. You can never access cross region security group id

upvoted 1 times

Chris22usa 2 years, 11 months ago

I realize D is right as ChatGpt indicates. Because here is not a problem just one application in a VPC connection to another in different region. Actually there many applications in different VPCs in a region which need to connect any other application crossing in other region. So two transit gateway need to installed in two regions for multiple to multiple VPCs connections.

upvoted 2 times

Iragmt 2 years, 11 months ago

However, there was also a part of "create an inbound rule in the database security group that references the security group ID of the application servers in eu-west-1"

therefore, still C because we cannot reference SG ID of diff VPC, we should use the CIDR block

upvoted 1 times

Chris22usa 2 years, 11 months ago

post it on ChaptGpt and it give me answer D. what heck with this?

upvoted 2 times

haoAWS 2 years, 11 months ago

Selected Answer: C

B is wrong because It is in a different region, so reference to the security group ID will not work. A is wrong because you need to update the route table. The answer should be C.

upvoted 1 times

mattcl 2 years, 11 months ago

is B. what happens if application server IP addresses changes (Option C). You must change manually the IP in the security group again.

upvoted 1 times

antropaws 2 years, 11 months ago

Selected Answer: C

I thought B, but I vote C after checking Axeashes response.

upvoted 1 times

HelioNeto 3 years ago

Selected Answer: C

I think the answer is C because the security groups are in different VPCs. When the question wants to allow traffic from app vpc to database vpc i think using peering connection you will be able to add the security groups rules using private ip addresses of app servers. I don't think the database VPC will identify the security group id of another VPC.

upvoted 2 times

REzirezi 3 years ago

D You cannot create a VPC peering connection between VPCs in different regions.

upvoted 3 times

fakrap 3 years ago

You can peer any two VPCs in different Regions, as long as they have distinct, non-overlapping CIDR blocks. This ensures that all of the private IP addresses are unique, and it allows all of the resources in the VPCs to address each other without the need for any form of network address translation (NAT).

upvoted 2 times

nosense 3 years ago

Selected Answer: B

b for me. bcs correct inbound rule, and not overhead

upvoted 2 times

cloudenthusiast 3 years ago

Selected Answer: B

Option B suggests configuring a VPC peering connection between the ap-southeast-2 VPC and the eu-west-1 VPC. By establishing this peering connection, the VPCs can communicate with each other over their private IP addresses.

Additionally, updating the subnet route tables is necessary to ensure that the traffic destined for the remote VPC is correctly routed through the VPC peering connection.

To secure the communication, an inbound rule is created in the ap-southeast-2 database security group. This rule references the security group ID of the application servers in the eu-west-1 VPC, allowing traffic only from those instances. This approach ensures that only the authorized application servers can access the databases in the ap-southeast-2 VPC.

upvoted 5 times

A company is developing software that uses a PostgreSQL database schema. The company needs to configure multiple development environments and databases for the company's developers. On average, each development environment is used for half of the 8-hour workday.

Which solution will meet these requirements MOST cost-effectively?

- A. Configure each development environment with its own Amazon Aurora PostgreSQL database
- B. Configure each development environment with its own Amazon RDS for PostgreSQL Single-AZ DB instances
- C. Configure each development environment with its own Amazon Aurora On-Demand PostgreSQL-Compatible database
Most Voted
- D. Configure each development environment with its own Amazon S3 bucket by using Amazon S3 Object Select

Correct Answer: C

Community vote distribution

C (67%)

B (33%)

Comments

cloudehustiaft **Highly Voted** 3 years ago

Selected Answer: C

Option C suggests using Amazon Aurora On-Demand PostgreSQL-Compatible databases for each development environment. This option provides the benefits of Amazon Aurora, which is a high-performance and scalable database engine, while allowing you to pay for usage on an on-demand basis. Amazon Aurora On-Demand instances are typically more cost-effective for individual development environments compared to the provisioned capacity options.

upvoted 12 times

cloudehustiaft 3 years ago

Option B suggests using Amazon RDS for PostgreSQL Single-AZ DB instances for each development environment. While Amazon RDS is a reliable and cost-effective option, it may have slightly higher costs compared to Amazon Aurora On-Demand instances.

upvoted 6 times

Iragmt 2 years, 11 months ago

I'm thinking that it should be B, since question does not mention any requirement only cost effective and this is just an development environment I guess we can leverage the use of RDS free tier also

upvoted 3 times

Stranko **Highly Voted** 2 years, 3 months ago

Selected Answer: C

Guys, when you use the pricing calculator the cost between option B and C is really close. I doubt anyone wants to test on your knowledge of exact pricings in your region. I think that "On Demand" being explicitly specified in option C and not being specified in option B is the main difference here the exam wants to test. In that case I'd assume that option B means a constantly running instance and not "On Demand" which would make the choice pretty obvious. Again, I don't think AWS exam will test you on knowing that a single AZ is cheaper by 0,005 cents than Aurora :D

upvoted 12 times

tch **Most Recent** 1 year, 3 months ago

Selected Answer: C

There is no upfront commitment with Aurora. You pay an hourly charge for each instance that you launch, and when you're finished with an Aurora DB instance, you can delete it. You do not need to overprovision storage as a safety margin, and you only pay for the storage you actually consume from AWS official link

upvoted 1 times

zdi561 1 year, 4 months ago

Selected Answer: B

B RDS single az is the cheapest. Aurora needs at least two azs.

upvoted 1 times

tonybuivanngia 1 year, 7 months ago

Selected Answer: C

I choose C because Aurora on-Demand is Aurora Serverless:

The Aurora Serverless is cost effective. Scale out fine-grained increments to provide just the right number of database resources and pay only for capacity consumed.

upvoted 2 times

a7md0 1 year, 11 months ago

Selected Answer: B

Single-AZ DB instances cheaper

upvoted 1 times

trinh_le 2 years, 1 month ago

Selected Answer: B

Single AZ more cost effective

upvoted 2 times

chasingsummer 2 years, 4 months ago

Selected Answer: B

1 instance(s) x 0.245 USD hourly x (4 / 24 hours in a day) x 730 hours in a month = 29.8083 USD ----> Amazon RDS PostgreSQL instances cost (monthly)

1 instance(s) x 0.26 USD hourly x (4 / 24 hours in a day) x 730 hours in a month = 31.6333 USD ----> Amazon Aurora PostgreSQL-Compatible DB instances cost (monthly)

upvoted 2 times

upliftinghut 2 years, 4 months ago

Selected Answer: C

C is correct because B is cheaper but they don't mention to stop the DB when not in use

upvoted 3 times

awsgeek75 2 years, 5 months ago

Selected Answer: C

On-Demand is cheaper than Aurora or RDS because of low weekly usage

upvoted 2 times

pentium75 2 years, 5 months ago

Selected Answer: C

We have environments that are used on average 4 hours per workday = 20 hours per week. So with option C (Aurora on-demand aka serverless) we pay for 20 hours per week. With option B (RDS) we pay for 168 hours per week (the answer does not mention anything about automating shutdown etc.).

So even if Aurora Serverless is slightly more expensive than RDS, C is cheaper because we pay only 20 (not 168) hours per week.

upvoted 4 times

Salilgen 1 year, 5 months ago

Question doesn't talk about serverless. Moreover, you can select on-demand pricing model for RDS for PostgreSQL Single-AZ too.

IMO answer is B

upvoted 1 times

Mikado211 2 years, 6 months ago

Selected Answer: B

Aurora on demand is (a little) more expensive than Aurora

Aurora is more expensive than RDS single instance

So cost effectiveness == RDS.

(B)

upvoted 1 times

pentium75 2 years, 5 months ago

But if you use the database only 20 hours per week (5 x 4), wouldn't you pay way less with Aurora serverless than with RDS?

upvoted 3 times

Murtadhaceit 2 years, 6 months ago

Selected Answer: B

AWS Services Calculator is showing B cheaper by less than a dollar for the same settings for both. I used "db.r6g.large" for RDS (Single-AZ) and Aurora and put 4 hours/day.

upvoted 7 times

Stranko 2 years, 3 months ago

I used the calculator, single AZ is cheaper for the exact same usage duration, if you pick On-Demand option for it too. In Aurora case (option C) you have "On Demand" explicitly specified, so if it has to be specified then I suppose that B option is about a constantly running instance. If B had an "On Demand" added, I'd vote B too.

upvoted 2 times

JoseVincent68 2 years, 6 months ago

Selected Answer: B

Amazon RDS Single AZ is cheaper than Aurora Multi-AZ

upvoted 1 times

Wayne23Fang 2 years, 7 months ago

Selected Answer: B

Aurora instances will cost you ~20% more than RDS MySQL Given the running hours the same.
Also Aurora is HA.

upvoted 1 times

baba365 2 years, 8 months ago

... just trying to trick you. Aurora on demand is Aurora Serverless.

upvoted 6 times

Anmol_1010 2 years, 7 months ago

that is good piece of infoamtion

upvoted 1 times

deechean 2 years, 9 months ago

Selected Answer: C

Aurora allows you to pay for the hours used. 4 hour every day, you only need 1/6 cost of 24 hours per day. You can check the Aurora pricing calculator.

upvoted 4 times

A company uses AWS Organizations with resources tagged by account. The company also uses AWS Backup to back up its AWS infrastructure resources. The company needs to back up all AWS resources.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use AWS Config to identify all untagged resources. Tag the identified resources programmatically. Use tags in the backup plan. **Most Voted**
- B. Use AWS Config to identify all resources that are not running. Add those resources to the backup vault.
- C. Require all AWS account owners to review their resources to identify the resources that need to be backed up.
- D. Use Amazon Inspector to identify all noncompliant resources.

Correct Answer: A

Community vote distribution

A (100%)

Comments

cloudenthusiast **Highly Voted** 3 years ago

Selected Answer: A

This solution allows you to leverage AWS Config to identify any untagged resources within your AWS Organizations accounts. Once identified, you can programmatically apply the necessary tags to indicate the backup requirements for each resource. By using tags in the backup plan configuration, you can ensure that only the tagged resources are included in the backup process, reducing operational overhead and ensuring all necessary resources are backed up.

upvoted 5 times

Gape4 **Most Recent** 1 year, 11 months ago

Selected Answer: A

I will go for A. C and D doesn't make sense. B- resources not running? No

upvoted 2 times

TariqKipkemei 2 years, 7 months ago

Selected Answer: A

Use AWS config to deploy the tag rule and remediate resources that are not compliant.

upvoted 4 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: A

This option has the least operational overhead:

AWS Config continuously evaluates resource configurations and can identify untagged resources
Resources can be programmatically tagged via the AWS SDK based on Config data
Backup plans can use tag criteria to automatically back up newly tagged resources
No manual review or resource discovery needed

upvoted 4 times

Bill1000 3 years ago

Selected Answer: A

Vote A

upvoted 3 times

nosense 3 years ago

Selected Answer: A

a valid for me

upvoted 4 times

A social media company wants to allow its users to upload images in an application that is hosted in the AWS Cloud. The company needs a solution that automatically resizes the images so that the images can be displayed on multiple device types. The application experiences unpredictable traffic patterns throughout the day. The company is seeking a highly available solution that maximizes scalability.

What should a solutions architect do to meet these requirements?

- A. Create a static website hosted in Amazon S3 that invokes AWS Lambda functions to resize the images and store the images in an Amazon S3 bucket. **Most Voted**
- B. Create a static website hosted in Amazon CloudFront that invokes AWS Step Functions to resize the images and store the images in an Amazon RDS database.
- C. Create a dynamic website hosted on a web server that runs on an Amazon EC2 instance. Configure a process that runs on the EC2 instance to resize the images and store the images in an Amazon S3 bucket.
- D. Create a dynamic website hosted on an automatically scaling Amazon Elastic Container Service (Amazon ECS) cluster that creates a resize job in Amazon Simple Queue Service (Amazon SQS). Set up an image-resizing program that runs on an Amazon EC2 instance to process the resize jobs.

Correct Answer: A

Community vote distribution

A (100%)

Comments

cloudenthusiast **Highly Voted** 3 years ago

Selected Answer: A

By using Amazon S3 and AWS Lambda together, you can create a serverless architecture that provides highly scalable and available image resizing capabilities. Here's how the solution would work:

Set up an Amazon S3 bucket to store the original images uploaded by users.

Configure an event trigger on the S3 bucket to invoke an AWS Lambda function whenever a new image is uploaded.

The Lambda function can be designed to retrieve the uploaded image, perform the necessary resizing operations based on device requirements, and store the resized images back in the S3 bucket or a different bucket designated for resized images.

Configure the Amazon S3 bucket to make the resized images publicly accessible for serving to users.

upvoted 19 times

surajkrishnamurthy **Most Recent** 1 year, 4 months ago

Selected Answer: A

Is Static Website Read only ?

Eventhough I have selected Answer as A

upvoted 1 times

LeonSauveterre 1 year, 6 months ago

Selected Answer: A

OK. A is better, although I chose D originally.

I guess it shouldn't be D simply because it overkills for this scenario and introduces higher operational overhead (even though option D can handle more complex workloads than Lambda if needed).

upvoted 1 times

cnureddy 2 years ago

How can end user upload an image to S3 bucket with static hosting. I believe it should be dynamic website (Answer D)

upvoted 2 times

mr123dd 2 years, 5 months ago

image = static = S3 or cloudfront
but image is unstructured data so you dont store it in a relational database like RDS
and Step Function is not for processing
So A

upvoted 4 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: A

This meets all the key requirements:

S3 static website provides high availability and auto scaling to handle unpredictable traffic
Lambda functions invoked from the S3 site can resize images on the fly
Storing images in S3 buckets provides durability, scalability and high throughput
Serverless approach with S3 and Lambda maximizes scalability and availability

upvoted 2 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: A

Scalability = S3, Lamda
automatically resize images = Lambda

upvoted 3 times

A company is running a microservices application on Amazon EC2 instances. The company wants to migrate the application to an Amazon Elastic Kubernetes Service (Amazon EKS) cluster for scalability. The company must configure the Amazon EKS control plane with endpoint private access set to true and endpoint public access set to false to maintain security compliance. The company must also put the data plane in private subnets. However, the company has received error notifications because the node cannot join the cluster.

Which solution will allow the node to join the cluster?

- A. Grant the required permission in AWS Identity and Access Management (IAM) to the AmazonEKSNodeRole IAM role.
- B. Create interface VPC endpoints to allow nodes to access the control plane. **Most Voted**
- C. Recreate nodes in the public subnet. Restrict security groups for EC2 nodes.
- D. Allow outbound traffic in the security group of the nodes.

Correct Answer: B

Community vote distribution

B (60%)

A (40%)

Comments

cloudenthusiast **Highly Voted** 3 years ago

Selected Answer: B

By creating interface VPC endpoints, you can enable the necessary communication between the Amazon EKS control plane and the nodes in private subnets. This solution ensures that the control plane maintains endpoint private access (set to true) and endpoint public access (set to false) for security compliance.

upvoted 23 times

y0 **Highly Voted** 3 years ago

Selected Answer: A

Check this : <https://docs.aws.amazon.com/eks/latest/userguide/create-node-role.html>

Also, EKS does not require VPC endpoints. This is not the right use case for EKS

upvoted 21 times

h0ng97_spare_002 2 years, 2 months ago

<https://docs.aws.amazon.com/eks/latest/userguide/create-node-role.html#:~:text=Before,launched>

"Before you can launch nodes and register them into a cluster, you must create an IAM role for those nodes to use when they are launched."

upvoted 4 times

JA2018 1 year, 6 months ago

which implies AmazonEKSNodeRole IAM role had already been configured..... that leaves answer B as the only viable choice

upvoted 1 times

TwinSpark 2 years, 1 month ago

correct i was going for B, but A looks better.

<https://docs.aws.amazon.com/eks/latest/userguide/cluster-endpoint.html>

"When you enable endpoint private access for your cluster, Amazon EKS creates a Route 53 private hosted zone on your behalf and associates it with your cluster's VPC. This private hosted zone is managed by Amazon EKS, and it doesn't appear in your account's Route 53 resources. "

upvoted 1 times

bora4motion **Most Recent** 1 year, 1 month ago

Selected Answer: B

to me it looks like a network problem, not a "rights" problem.

upvoted 1 times

Dantecito 1 year, 4 months ago

Selected Answer: B

B. "The company has received error notifications because the node cannot join the cluster" it's a network connectivity issue, not permissions.
upvoted 1 times

4729e6c 1 year, 4 months ago

Selected Answer: B

Since the company has configured the Amazon EKS control plane with private access enabled (true) and public access disabled (false), the EKS API server is only accessible through private network routes. The worker nodes, which are in private subnets, must be able to communicate with the EKS control plane for proper functionality.

When the nodes cannot join the cluster, the most common cause is that they are unable to reach the EKS API server. In this case, since the control plane is private-only, AWS recommends using interface VPC endpoints (AWS PrivateLink) to enable private communication between the nodes and the control plane.

upvoted 1 times

Salilgen 1 year, 5 months ago

Selected Answer: A

IMO answer is A.

In private cluster you need interface VPC endpoints to necessary connect to some AWS services

(<https://docs.aws.amazon.com/eks/latest/userguide/private-clusters.html>) but EKS node always connect to control plane by EKS owned ENI.

See this: it is very clear.

<https://keetmalin.medium.com/eks-cluster-network-architecture-for-worker-nodes-635e067c8c2a>

upvoted 2 times

LeonSauveterre 1 year, 6 months ago

Selected Answer: B

The question is trying to tell us:

1. "private access = true" and "public access = false". So the control plane endpoint is private and only accessible from within the VPC.
2. Nodes (data plane) are in private subnets.

Option A is NECESSARY but lack of permissions would generally cause authorization errors, not connectivity errors (that lead to failure of joining the cluster). Only after you got this right, you would likely receive errors about joining. So apparently we have already configured auth correctly, meaning A is not the answer.

Option C exposes the nodes to the internet. Wrong.

Option D is important for nodes to communicate with AWS services (because of allowing outbound traffic), but it's not sufficient if the required VPC interface endpoints are not even there.

upvoted 1 times

a7md0 1 year, 11 months ago

Selected Answer: A

AmazonEKSNoderole IAM role

<https://docs.aws.amazon.com/eks/latest/userguide/create-node-role.html>

upvoted 2 times

emakid 1 year, 11 months ago

Selected Answer: B

When Amazon EKS nodes cannot join the cluster, especially when the control plane is set to private access only, the issue typically revolves around networking and connectivity. When the EKS control plane is configured with private access only, the nodes must communicate with the control plane over private IP addresses. Creating VPC endpoints (specifically, `com.amazonaws.<region>.eks`) allows traffic between the EKS nodes and the control plane to be routed privately within the VPC, which resolves the connectivity issue.

upvoted 3 times

Gape4 1 year, 11 months ago

Selected Answer: B

I think is B.

upvoted 2 times

MandAsh 1 year, 12 months ago

Selected Answer: B

Error they have mentioned is at network level. They are not saying authorisation is failed rather node is unable to connect to cluster aka connectivity issue. So answer it must be B

upvoted 2 times

Rocconno 2 years ago

Selected Answer: B

<https://docs.aws.amazon.com/eks/latest/userguide/private-clusters.html>

"Any self-managed nodes must be deployed to subnets that have the VPC interface endpoints that you require. If you create a managed node group the VPC interface endpoint security group must allow the CIDR for the subnets, or you must add the created node security group to the VPC interface endpoint security group."

upvoted 2 times

stalk98 2 years, 1 month ago

I Think is A

upvoted 1 times

trinh_le 2 years, 1 month ago

Selected Answer: B

B is good to go

upvoted 3 times

JackyCCK 2 years, 2 months ago

S3/DynamoDB - VPC endpoint, other service should use interface endpoint so B is incorrect

upvoted 2 times

bujuman 2 years, 2 months ago

Selected Answer: B

Because of these two assertions:

- Amazon EKS control plane with endpoint private access set to true and endpoint public access set to false to maintain security compliance.

(The company must also put the data plane in private subnets.

The best answer is related to Networking, Private Subnets (EKS Ctr Plane is strictly private and Data Plane stick under private subnets) and not related to EKS autodeployment that sure need an IAM policy. So according to me, answer B is the best answer.

upvoted 3 times

potomac 2 years, 7 months ago

Selected Answer: A

Before can launch nodes and register nodes into a EKS cluster, must create an IAM role for those nodes to use when they are launched.

upvoted 2 times

A company is migrating an on-premises application to AWS. The company wants to use Amazon Redshift as a solution.

Which use cases are suitable for Amazon Redshift in this scenario? (Choose three.)

A. Supporting data APIs to access data with traditional, containerized, and event-driven applications

B. Supporting client-side and server-side encryption **Most Voted**

C. Building analytics workloads during specified hours and when the application is not active **Most Voted**


D. Caching data to reduce the pressure on the backend database

E. Scaling globally to support petabytes of data and tens of millions of requests per minute **Most Voted**

F. Creating a secondary replica of the cluster by using the AWS Management Console

Correct Answer: BCE

Community vote distribution



Comments

elmogy **Highly Voted** 3 years ago

Selected Answer: BCE

Amazon Redshift is a data warehouse solution, so it is suitable for:

- Supporting encryption (client-side and server-side)
- Handling analytics workloads, especially during off-peak hours when the application is less active
- Scaling to large amounts of data and high query volumes for analytics purposes

The following options are incorrect because:

- A) Data APIs are not typically used with Redshift. It is more for running SQL queries and analytics.
- D) Redshift is not typically used for caching data. It is for analytics and data warehouse purposes.
- F) Redshift clusters do not create replicas in the management console. They are standalone clusters. you could create DR cluster from snapshot and restore to another region (automated or manual) but I do not think this what is meant in this option.

upvoted 21 times

pentium75 2 years, 5 months ago

"Data APIs are not typically used with Redshift" -> "With the Data API, you can programmatically access data in your Amazon Redshift cluster from different AWS services such as AWS Lambda, Amazon SageMaker notebooks, AWS Cloud9, and also your on-premises applications using the AWS SDK. This allows you to build cloud-native, containerized, serverless, web-based, and event-driven applications on the AWS Cloud."

upvoted 2 times

Rob1L **Highly Voted** 3 years ago

Selected Answer: BCE

B. Supporting client-side and server-side encryption: Amazon Redshift supports both client-side and server-side encryption for improved data security.

C. Building analytics workloads during specified hours and when the application is not active: Amazon Redshift is optimized for running complex analytic queries against very large datasets, making it a good choice for this use case.

E. Scaling globally to support petabytes of data and tens of millions of requests per minute: Amazon Redshift is designed to handle petabytes of data, and to deliver fast query and I/O performance for virtually any size dataset.

upvoted 6 times

Dz1990 **Most Recent** 1 year ago

Selected Answer: ACE

Going with ACE here!

The key insight about option B: encryption is a feature, not a use case. The question asks what you do with Redshift, not what it has. It's like asking "Why buy a Tesla?" and answering "it has seatbelts" vs "to reduce carbon footprint." See the difference?

For option E - yeah, "millions of requests" sounds intense for a data warehouse, but Redshift legitimately scales to petabytes and handles massive

analytical workloads. It's not about transactional speed; it's about analytical scale. With concurrency scaling, it can handle way more than people think.

ACE = actual business use cases: modern API access, off-peak analytics, and massive data scaling. BCE = mixing features with use cases, which is a classic AWS exam trap!

upvoted 1 times

DanRuggles 3 weeks, 2 days ago

Amazon Redshift fully supports the Redshift Data API to access data from traditional, containerized, and event-driven applications. The Redshift Data API provides a simple, secure HTTP endpoint (via AWS SDK) that lets you run SQL statements on Redshift without:

- Managing database drivers (JDBC/ODBC)
- Handling persistent connections
- Dealing with complex network configurations

upvoted 1 times

Dantecito 1 year, 4 months ago

Selected Answer: ABC

Chatgpt says ABC.

D. Not for caching.

E. It is not optimized for high-volume like millions of requests per minute.

F. Does not natively support creating a secondary replica of the cluster.

upvoted 1 times

LeonSauveterre 1 year, 6 months ago

Selected Answer: BCE

YOU MUST KNOW: Amazon Redshift is a cloud-based data warehouse service designed for analytics and reporting on large datasets, optimized for complex queries and analytical workloads rather than transactional or real-time applications.

After you memorized that, option A is incorrect. If you want to work on operational workloads, then Aurora, Amazon RDS, DynamoDB, and such, are better choices.

Option D: Caching is not a core use case for Redshift. You can choose DAX (Amazon DynamoDB Accelerator) or ElastiCache to mitigate database pressure.

Option E: I choose that. BUT, Redshift is not designed for *high-concurrency* transactional workloads like "10s of mils of reqs/min", although it CAN handle massive datasets. Well... It seems partially true to me compared to other definitely wrong options.

upvoted 2 times

rpmaws 1 year, 9 months ago

Selected Answer: ACE

B is not correct, how it can do encryption at client side ?

upvoted 2 times

3bdf1cc 1 year, 11 months ago

Found this related to A -- but specific to Redshift Serverless - but should qualify as a Redshift use case

The Data API enables you to seamlessly access data from Redshift Serverless with all types of traditional, cloud-native, and containerized serverless web service-based applications and event-driven applications.

<https://www.amazonaws.cn/en/blog-selection/use-the-amazon-redshift-data-api-to-interact-with-amazon-redshift-serverless/>

upvoted 1 times

NSA_Poker 2 years ago

Selected Answer: BCE

The following are obviously incorrect:

(D) Redshift is not as suitable as ElastiCache for caching.

(F) A secondary replica of the cluster is not supported.

The debate is between BCE & ACE or simplified, between A & C.

(A) is incorrect bc there is a difference btw Amazon Redshift Data API & API Gateway. API Gateway supports containerized and serverless workloads, as well as web applications. Amazon Redshift Data API is a built in API to access Redshift data with web services-based applications, including AWS Lambda, Amazon SageMaker notebooks, and AWS Cloud9.

<https://aws.amazon.com/blogs/big-data/build-a-serverless-analytics-application-with-amazon-redshift-and-amazon-api-gateway/>

(B) is correct. You have the following options of protecting data at rest in Amazon Redshift. Use server-side encryption OR use client-side encryption <https://docs.aws.amazon.com/redshift/latest/mgmt/security-encryption.html>

upvoted 2 times

JackyCCK 2 years, 2 months ago

Redshift is OLAP(online analytical processing) so D is wrong, "when the application is not active"

upvoted 2 times

JackyCCK 2 years, 2 months ago

*C is wrong

upvoted 1 times

awsgeek75 2 years, 5 months ago

Selected Answer: ACE

A, C, E are for data and Redshift is data warehouse.
B is too generic of a choice
D caching is not the main purpose of Redshift
F replication is not main use of Redshift

CE are easy

Between AB, I chose A because Redshift supports data API and client-side encryption is not Redshift specific

upvoted 3 times

1rob 2 years, 5 months ago

Selected Answer: ABD

A: source <https://aws.amazon.com/blogs/big-data/using-the-amazon-redshift-data-api-to-interact-with-amazon-redshift-clusters/>
B: source: <https://docs.aws.amazon.com/redshift/latest/mgmt/security-encryption.html>
C: not sure, but you can configure scheduled queries, but the remark " and when the application is not active " , that is not relevant.
D: source https://docs.aws.amazon.com/redshift/latest/dg/c_challenges_achieving_high_performance_queries.html
E: Scaling globally is not supported; redshift is only a regional service.
F: only read replica is supported. So not a secondary replica of the cluster.

upvoted 2 times

pentium75 2 years, 5 months ago

Selected Answer: ABD

A: <https://aws.amazon.com/de/blogs/big-data/get-started-with-the-amazon-redshift-data-api/>
B: <https://docs.aws.amazon.com/redshift/latest/mgmt/security-encryption.html>
D: https://docs.aws.amazon.com/redshift/latest/dg/c_challenges_achieving_high_performance_queries.html#result-caching

Not C: Redshift is a Data Warehouse; you can use that for analytics, but it is not directly related to an "application"

Not E: "Petabytes of data" yes, but "tens of millions of requests per minute" is not a typical feature of Redshift

Nor F: Replicas are not a Redshift feature

upvoted 2 times

TariqKipkemei 2 years, 7 months ago

Selected Answer: ACE

Technically both options A and B apply, this is from the links below:

A. You can access your Amazon Redshift database using the built-in Amazon Redshift Data API.

<https://docs.aws.amazon.com/redshift/latest/mgmt/data-api.html#:~:text=in%20Amazon%20Redshift-,Data%20API,-.%20Using%20this%20API>

B. You can encrypt data client-side and upload the encrypted data to Amazon Redshift. In this case, you manage the encryption process, the encryption keys, and related tools.

<https://docs.aws.amazon.com/redshift/latest/mgmt/security-encryption.html#:~:text=Use-,client%2Dside,-encryption%20%E2%80%93%20You%20can>

upvoted 2 times

potomac 2 years, 7 months ago

Selected Answer: ABC

Amazon Redshift provides a Data API that you can use to painlessly access data from Amazon Redshift with all types of traditional, cloud-native, and containerized, serverless web services-based and event-driven applications.

Amazon Redshift supports up to 500 concurrent queries per cluster, which may be expanded by adding more nodes to the cluster.

upvoted 3 times

potomac 2 years, 7 months ago

change to ABD

To reduce query runtime and improve system performance, Amazon Redshift caches the results of certain types of queries in memory on the leader node. When a user submits a query, Amazon Redshift checks the results cache for a valid, cached copy of the query results. If a match is found in the result cache, Amazon Redshift uses the cached results and doesn't run the query. Result caching is transparent to the user.

upvoted 1 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: BCE

The key use cases for Amazon Redshift that fit this scenario are:

B) Redshift supports both client-side and server-side encryption to protect sensitive data.

C) Redshift is well suited for running batch analytics workloads during off-peak times without affecting OLTP systems.

E) Redshift can scale to massive datasets and concurrent users to support large analytics workloads.

upvoted 3 times

cd93 2 years, 9 months ago

Selected Answer: BCD

Why E lol? It's a data warehouse! it has no need to support millions of requests, it is not mentioned anywhere (<https://aws.amazon.com/redshift/features>)

In fact Redshift editor supports max 500 connections and workgroup support max 2000 connections at once, see it's quota page
Redshift has a cache layer, D is correct

upvoted 3 times

mrsoa 2 years, 10 months ago

Selected Answer: BCE

BCE, For B this is why

<https://docs.aws.amazon.com/redshift/latest/mgmt/security-encryption.html>

upvoted 2 times

A company provides an API interface to customers so the customers can retrieve their financial information. The company expects a larger number of requests during peak usage times of the year.

The company requires the API to respond consistently with low latency to ensure customer satisfaction. The company needs to provide a compute host for the API.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use an Application Load Balancer and Amazon Elastic Container Service (Amazon ECS).
- B. Use Amazon API Gateway and AWS Lambda functions with provisioned concurrency. **Most Voted**
- C. Use an Application Load Balancer and an Amazon Elastic Kubernetes Service (Amazon EKS) cluster.
- D. Use Amazon API Gateway and AWS Lambda functions with reserved concurrency.

Correct Answer: B

Community vote distribution

B (79%)

A (21%)

Comments

cloudenthusiast **Highly Voted** 3 years ago

Selected Answer: B

In the context of the given scenario, where the company wants low latency and consistent performance for their API during peak usage times, it would be more suitable to use provisioned concurrency. By allocating a specific number of concurrent executions, the company can ensure that there are enough function instances available to handle the expected load and minimize the impact of cold starts. This will result in lower latency and improved performance for the API.

upvoted 13 times

Bennyboy789 **Highly Voted** 2 years, 9 months ago

Selected Answer: B

Provisioned - minimizing cold starts and providing low latency.

upvoted 8 times

LeonSauveterre **Most Recent** 1 year, 6 months ago

Selected Answer: B

Concurrency of lambda function, provisioned vs reserved:

Reserved Concurrency:

1. Desired instances can start without interference from other functions
2. Guarantees available execution capacity for critical functions

Provisioned Concurrency:

1. Also guarantees consistent performance
2. [!!] Allows you to pre-warm a specific number of function instances to reduce cold start latency

upvoted 3 times

awsgeek75 2 years, 5 months ago

Selected Answer: A

<https://docs.aws.amazon.com/lambda/latest/dg/lambda-concurrency.html#reserved-and-provisioned>

Consistency decreases if you exceed your provisioned instance. Let's say you have 1000 (default) provisioned instances and the load is 1500. The new 500 will have to wait until the first 1000 concurrent calls finish. This is solved by increasing the provisioned concurrency to 1500.

upvoted 2 times

1rob 2 years, 6 months ago

Selected Answer: A

So I have my doubts here. The question also states ;"The company needs to provide a compute host for the API." Imho this implies to have some sort of physical host which has to be provided by the customer. Translating this further to aws this would mean an EC2 instance. And then when I would go for ECS in stead of EKS.
Please share your opinion.

upvoted 7 times

pdragon1981 2 years, 5 months ago

Exactly, initially I was thinking on B but if company must provide a host I would say that only option A is feasible

upvoted 3 times

pdragon1981 2 years, 5 months ago

Sorry I understand bad the text, correct answer is B, as for my understanding now the host is the device that the customer needs to connect with API Gateway, below explains well the logic

<https://aws.amazon.com/api-gateway/>

upvoted 4 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: B

This option provides the least operational overhead:

API Gateway handles the API requests and integration with Lambda

Lambda automatically scales compute without managing servers

Provisioned concurrency ensures consistent low latency by keeping functions initialized

No need to manage containers or orchestration platforms as with ECS/EKS

upvoted 3 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: B

The company requires the API to respond consistently with low latency to ensure customer satisfaction especially during high peak periods, there is no mention of cost efficient. Hence provisioned concurrency is the best option.

Provisioned concurrency is the number of pre-initialized execution environments you want to allocate to your function. These execution environments are prepared to respond immediately to incoming function requests. Configuring provisioned concurrency incurs charges to your AWS account.

<https://docs.aws.amazon.com/lambda/latest/dg/provisioned-concurrency.html#:~:text=for%20a%20function.,Provisioned%20concurrency,-%E2%80%93Provisioned%20concurrency%20is>

upvoted 2 times

MirKhobaeb 3 years ago

Selected Answer: B

AWS Lambda provides a highly scalable and distributed infrastructure that automatically manages the underlying compute resources. It automatically scales your API based on the incoming request load, allowing it to respond consistently with low latency, even during peak times. AWS Lambda takes care of infrastructure provisioning, scaling, and resource management, allowing you to focus on writing the code for your API logic.

upvoted 4 times


A company wants to send all AWS Systems Manager Session Manager logs to an Amazon S3 bucket for archival purposes.

Which solution will meet this requirement with the MOST operational efficiency?

- A. Enable S3 logging in the Systems Manager console. Choose an S3 bucket to send the session data to. **Most Voted**
- B. Install the Amazon CloudWatch agent. Push all logs to a CloudWatch log group. Export the logs to an S3 bucket from the group for archival purposes.
- C. Create a Systems Manager document to upload all server logs to a central S3 bucket. Use Amazon EventBridge to run the Systems Manager document against all servers that are in the account daily.
- D. Install an Amazon CloudWatch agent. Push all logs to a CloudWatch log group. Create a CloudWatch logs subscription that pushes any incoming log events to an Amazon Kinesis Data Firehose delivery stream. Set Amazon S3 as the destination.

Correct Answer: A

Community vote distribution



Comments

master9 **Highly Voted** 2 years, 5 months ago

Selected Answer: A

send logs to Amazon S3 from AWS Systems Manager Session Manager. Here are the steps to do so:

Enable S3 Logging: Open the AWS Systems Manager console. In the navigation pane, choose Session Manager. Choose the Preferences tab, and then choose Edit. Select the check box next to Enable under S3 logging.

Create an S3 Bucket: To store the Session Manager logs, create an S3 bucket to hold the audit logs from the Session Manager interactive shell usage.

Configure IAM Role: AWS Systems Manager Agent (SSM Agent) uses the same AWS Identity and Access Management (IAM) role to activate itself and upload logs to Amazon S3. You can use either an IAM instance profile that's attached to an Amazon Elastic Compute Cloud (Amazon EC2) instance or the IAM role that's configured for the Default Host Management Configuration.

upvoted 8 times

cloudenthusiast **Highly Voted** 3 years ago

Selected Answer: A

option A does not involve CloudWatch, while option D does. Therefore, in terms of operational overhead, option A would generally have less complexity and operational overhead compared to option D.

Option A simply enables S3 logging in the Systems Manager console, allowing you to directly send session logs to an S3 bucket. This approach is straightforward and requires minimal configuration.

On the other hand, option D involves installing and configuring the Amazon CloudWatch agent, creating a CloudWatch log group, setting up a CloudWatch Logs subscription, and configuring an Amazon Kinesis Data Firehose delivery stream to store logs in an S3 bucket. This requires additional setup and management compared to option A.

So, if minimizing operational overhead is a priority, option A would be a simpler and more straightforward choice.

upvoted 5 times

pujithacg8 **Most Recent** 1 year, 10 months ago

A, You can choose to store session log data in a specified Amazon Simple Storage Service (Amazon S3) bucket for debugging and troubleshooting purposes.

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager-logging.html#session-manager-logging-s3>

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: A