

Most efficient is A because it is a direct option in SM logging.

B can work but is more operational overhead as you end up using CloudWatch (not sure how but making assumption based on language of option)

C is definitely too much work

D Way too many moving parts

upvoted 3 times

potomac 2 years, 7 months ago

Selected Answer: A

You can choose to store session log data in a specified Amazon Simple Storage Service (Amazon S3) bucket for debugging and troubleshooting purposes.

upvoted 2 times

deechean 2 years, 9 months ago

Selected Answer: A

You can config the log archived to S3 in the Session Manager - > preference tab. Another option is CloudWatch log.

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager-logging.html#session-manager-logging-s3>

upvoted 2 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: A

*Simplicity - Enabling S3 logging requires just a simple configuration in the Systems Manager console to specify the destination S3 bucket. No other services need to be configured.

*Direct integration - Systems Manager has native support to send session logs to S3 through this feature. No need for intermediary services.

*Automated flow - Once S3 logging is enabled, the session logs automatically flow to the S3 bucket without manual intervention.

*Easy management - The S3 bucket can be managed independently for log storage and archival purposes without impacting Systems Manager.

*Cost-effectiveness - No charges for intermediate CloudWatch or Kinesis services. Just basic S3 storage costs.

*Minimal overhead - No ongoing management of complex pipeline of services. Direct logs to S3 minimizes overhead.

upvoted 3 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: A

With the MOST operational efficiency then option A is best.

Otherwise B is also an option with a little bit more ops than option A.

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager-logging.html>

upvoted 2 times

Zox42 2 years, 11 months ago

Selected Answer: A

Answer A. <https://aws-labs.net/winlab5-manageinfra/sessmgrlog.html>

upvoted 2 times

Zuit 2 years, 11 months ago

Selected Answer: A

GPT argued for D.

B could be an option, by installing a logging package on alle managed systems/ECs etc. <https://docs.aws.amazon.com/systems-manager/latest/userguide/distributor-working-with-packages-deploy.html>

However, as it mentions the "Session manager logs" I would tend towards A.

upvoted 2 times

MrAWSAssociate 2 years, 11 months ago

Selected Answer: A

It should be "A".

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager-logging.html>

upvoted 2 times

secdgs 2 years, 12 months ago

Selected Answer: A

It have menu to Enable S3 Logging.

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager-logging.html#session-manager-logging-s3>

upvoted 2 times

Markie999 3 years ago

Selected Answer: B

BBBBBBBBB

upvoted 1 times

pentium75 2 years, 5 months ago
"Install the CloudWatch agent" where?
upvoted 2 times

Bill1000 3 years ago

Selected Answer: B

The option 'A' says "Enable S3 logging in the Systems Manager console." This means that you will enable the logs !! FOR !! S3 events and its is not what the question asks. My vote is for Option B, based on this article: <https://docs.aws.amazon.com/AmazonS3/latest/userguide/logging-with-S3.html>

upvoted 1 times

baba365 2 years, 11 months ago

To log session data using Amazon S3 (console)

Open the AWS Systems Manager console at <https://console.aws.amazon.com/systems-manager/>.

In the navigation pane, choose Session Manager.

Choose the Preferences tab, and then choose Edit.

Select the check box next to Enable under S3 logging.

upvoted 3 times

vrevkov 2 years, 11 months ago

But where do you want to install the Amazon CloudWatch agent in case of B?

upvoted 1 times

omoakin 3 years ago

DDDDDD

upvoted 1 times

Anmol_1010 3 years ago

Option D is definitely not right,

Its option B

upvoted 2 times

omoakin 3 years ago

Chat GPT says option A is incorrect cos it requires enabling S3 logging in the system manager console only logs information about the systems manager service not the session logs

Says correct answer is B

upvoted 1 times

An application uses an Amazon RDS MySQL DB instance. The RDS database is becoming low on disk space. A solutions architect wants to increase the disk space without downtime.

Which solution meets these requirements with the LEAST amount of effort?

- A. Enable storage autoscaling in RDS **Most Voted**
- B. Increase the RDS database instance size
- C. Change the RDS database instance storage type to Provisioned IOPS
- D. Back up the RDS database, increase the storage capacity, restore the database, and stop the previous instance

Correct Answer: A

Community vote distribution

A (100%)

Comments

cloudehustiaft **Highly Voted** 3 years ago

Selected Answer: A

Enabling storage autoscaling allows RDS to automatically adjust the storage capacity based on the application's needs. When the storage usage exceeds a predefined threshold, RDS will automatically increase the allocated storage without requiring manual intervention or causing downtime. This ensures that the RDS database has sufficient disk space to handle the increasing storage requirements.

upvoted 12 times

Gape4 **Most Recent** 1 year, 11 months ago

Selected Answer: A

Autoscaling.... without downtime...

upvoted 2 times

potomac 2 years, 7 months ago

Selected Answer: A

Amazon RDS for MariaDB, Amazon RDS for MySQL, Amazon RDS for PostgreSQL, Amazon RDS for SQL Server and Amazon RDS for Oracle support RDS Storage Auto Scaling. RDS Storage Auto Scaling automatically scales storage capacity in response to growing database workloads, with zero downtime.

upvoted 3 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: A

This question is so obvious

upvoted 4 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: A

RDS Storage Auto Scaling continuously monitors actual storage consumption, and scales capacity up automatically when actual utilization approaches provisioned storage capacity. Auto Scaling works with new and existing database instances. You can enable Auto Scaling with just a few clicks in the AWS Management Console. There is no additional cost for RDS Storage Auto Scaling. You pay only for the RDS resources needed to run your applications.

<https://aws.amazon.com/about-aws/whats-new/2019/06/rds-storage-auto-scaling/#:~:text=of%20the%20rest.,RDS%20Storage%20Auto%20Scaling,-continuously%20monitors%20actual>

upvoted 3 times

james2033 2 years, 10 months ago

Selected Answer: A

Quote "Amazon RDS now supports Storage Auto Scaling" and "... with zero downtime." (Jun 20th 2019) at <https://aws.amazon.com/about-aws/whats-new/2019/06/rds-storage-auto-scaling/>

upvoted 3 times

james2033 2 years, 10 months ago

Hello moderator, please help me delete this discussion, I already add content before this comment.

upvoted 1 times

james2033 2 years, 11 months ago

Selected Answer: A

See "Amazon RDS now supports Storage Auto Scaling. Posted On: Jun 20, 2019. Starting today, Amazon RDS for MariaDB, Amazon RDS for MySQL, Amazon RDS for PostgreSQL, Amazon RDS for SQL Server and Amazon RDS for Oracle support RDS Storage Auto Scaling. RDS Storage Auto Scaling automatically scales storage capacity in response to growing database workloads, with zero downtime." at <https://aws.amazon.com/about-aws/whats-new/2019/06/rds-storage-auto-scaling/>

upvoted 3 times

haoAWS 2 years, 11 months ago

Selected Answer: A

A is the best answer.

B will not work for increasing disk space, it only improve the IO performance.

C will not work because it will cause downtime.

D is too complicated and need much operational effort.

upvoted 3 times

Anmol_1010 3 years ago

The key word is No Down time. A would be bewt option

upvoted 3 times

A consulting company provides professional services to customers worldwide. The company provides solutions and tools for customers to expedite gathering and analyzing data on AWS. The company needs to centrally manage and deploy a common set of solutions and tools for customers to use for self-service purposes.

Which solution will meet these requirements?

- A. Create AWS CloudFormation templates for the customers.
- B. Create AWS Service Catalog products for the customers. **Most Voted**
- C. Create AWS Systems Manager templates for the customers.
- D. Create AWS Config items for the customers.

Correct Answer: B

Community vote distribution

B (100%)

Comments

cloudehustiaast **Highly Voted** 3 years ago

Selected Answer: B

AWS Service Catalog allows you to create and manage catalogs of IT services that can be deployed within your organization. With Service Catalog, you can define a standardized set of products (solutions and tools in this case) that customers can self-service provision. By creating Service Catalog products, you can control and enforce the deployment of approved and validated solutions and tools.

upvoted 10 times

Oblako 2 years, 6 months ago

"within your organization" => not for customers

upvoted 1 times

JA2018 1 year, 6 months ago

for this use case, in order to access the self-service functions, customers must have some form of AWS identities (e.g. accounts, users, user groups or roles) already setup in the service provider's AWS organisation.

upvoted 1 times

Guru4Cloud **Highly Voted** 2 years, 9 months ago

Selected Answer: B

Some key advantages of using Service Catalog:

Centralized management - Products can be maintained in a single catalog for easy discovery and governance.

Self-service access - Customers can deploy the solutions on their own without manual intervention.

Standardization - Products provide pre-defined templates for consistent deployment.

Access control - Granular permissions can be applied to restrict product visibility and access.

Reporting - Service Catalog provides detailed analytics on product usage and deployments.

upvoted 5 times

hsinchang **Most Recent** 2 years, 10 months ago

Selected Answer: B

CloudFormation: a code as infrastructure service

Systems Manager: management solution for resources

Config: assess, audit and evaluate configurations

Other options does not fit this scenario.

upvoted 2 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: B

AWS Service Catalog lets you centrally manage your cloud resources to achieve governance at scale of your infrastructure as code (IaC) templates, written in CloudFormation or Terraform. With AWS Service Catalog, you can meet your compliance requirements while making sure your customers can quickly deploy the cloud resources they need.

<https://aws.amazon.com/servicecatalog/#:~:text=How%20it%20works-,AWS%20Service%20Catalog,-lets%20you%20centrally>
upvoted 2 times

Yadav_Sanjay 3 years ago

Selected Answer: B

<https://docs.aws.amazon.com/servicecatalog/latest/adminguide/introduction.html>
upvoted 3 times


A company is designing a new web application that will run on Amazon EC2 Instances. The application will use Amazon DynamoDB for backend data storage. The application traffic will be unpredictable. The company expects that the application read and write throughput to the database will be moderate to high. The company needs to scale in response to application traffic.

Which DynamoDB table configuration will meet these requirements MOST cost-effectively?

- A. Configure DynamoDB with provisioned read and write by using the DynamoDB Standard table class. Set DynamoDB auto scaling to a maximum defined capacity.
- B. Configure DynamoDB in on-demand mode by using the DynamoDB Standard table class. **Most Voted**
- C. Configure DynamoDB with provisioned read and write by using the DynamoDB Standard Infrequent Access (DynamoDB Standard-IA) table class. Set DynamoDB auto scaling to a maximum defined capacity.
- D. Configure DynamoDB in on-demand mode by using the DynamoDB Standard Infrequent Access (DynamoDB Standard-IA) table class.

Correct Answer: B

Community vote distribution



Comments

Efren **Highly Voted** 3 years ago

B for me. Provisioned if we know how much traffic will come, but its unpredictable, so we have to go for on-demand
upvoted 13 times

VellaDevil 2 years, 11 months ago

Spot On

upvoted 3 times

tch **Most Recent** 1 year, 3 months ago

Selected Answer: B

DynamoDB offers two table classes designed to help you optimize for cost.

The DynamoDB Standard table class is the default, and is recommended for the vast majority of workloads.

The DynamoDB Standard-Infrequent Access (DynamoDB Standard-IA) table class is optimized for tables where storage is the dominant cost. For example, tables that store infrequently accessed data, such as application logs, old social media posts, e-commerce order history, and past gaming achievements, are good candidates for the Standard-IA table class.

upvoted 1 times

fdae619 1 year, 5 months ago

Selected Answer: B

Can someone answer why the standard class is a better choice to standard IA?

upvoted 1 times

tch 1 year, 3 months ago

DynamoDB offers two table classes designed to help you optimize for cost.

The DynamoDB Standard table class is the default, and is recommended for the vast majority of workloads.

The DynamoDB Standard-Infrequent Access (DynamoDB Standard-IA) table class is optimized for tables where storage is the dominant cost. For example, tables that store infrequently accessed data, such as application logs, old social media posts, e-commerce order history, and past gaming achievements, are good candidates for the Standard-IA table class.

upvoted 1 times

emakid 1 year, 11 months ago

Selected Answer: B

Configure DynamoDB in on-demand mode by using the DynamoDB Standard table class.

This option allows DynamoDB to automatically adjust to varying traffic patterns, which is ideal for unpredictable workloads. The Standard table class is suitable for applications with moderate to high read and write throughput, and on-demand mode ensures that you are billed based on the actual usage, providing cost efficiency for variable traffic patterns.

upvoted 3 times

Gape4 1 year, 11 months ago

Selected Answer: B

Key word : On demand. So I think B.

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: B

On demand

<https://docs.aws.amazon.com/wellarchitected/latest/serverless-applications-lens/capacity.html>

"With on-demand capacity mode, DynamoDB charges you for the data reads and writes your application performs on your tables. You do not need to specify how much read and write throughput you expect your application to perform because DynamoDB instantly accommodates your workload as they ramp up or down."

upvoted 3 times

pentium75 2 years, 5 months ago

Selected Answer: B

Not A because of "unpredictable" traffic

Not C and D because we are expecting "moderate to high" traffic

upvoted 4 times

leonliu4 2 years, 6 months ago

Selected Answer: B

Leaning towards B, it's hard to predict the capacity for A, and autoscaling doesn't respond fast

upvoted 3 times

peekingpicker 2 years, 6 months ago

Selected Answer: A

it's A.

remember that :

he company expects that the application read and write throughput to the database will be moderate to high

provisioned throughput is cheaper than ondemand capacity right ?

upvoted 2 times

pentium75 2 years, 5 months ago

but "unpredictable" which usually hints to on-demand

upvoted 2 times

dilaaziz 2 years, 6 months ago

Selected Answer: D

Data storage: <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithTables.tableclasses.html>

upvoted 1 times

Mr_Marcus 2 years ago

Actually, this link confirmed "B" for me...

upvoted 2 times

potomac 2 years, 7 months ago

Selected Answer: B

On-demand mode is great for unpredictable traffic

upvoted 3 times

bsbs1234 2 years, 8 months ago

I choose B

I think the items stored in the table in this question has large size. So each read/write, a big chunk of data pass through. A capacity unit is used to describe data throughput. provision to the high capacity units will be a waste because unpredicted traffic pattern.

upvoted 2 times

Bennyboy789 2 years, 9 months ago

Selected Answer: B

Unpredictable= on demand

upvoted 3 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: B

The key factors are:

With On-Demand mode, you only pay for what you use instead of over-provisioning capacity. This avoids idle capacity costs. DynamoDB Standard provides the fastest performance needed for moderate-high traffic apps vs Standard-IA which is for less frequent access. Auto scaling with provisioned capacity can also work but requires more administrative effort to tune the scaling thresholds.

upvoted 2 times

msdnpro 2 years, 10 months ago

Selected Answer: B

Support for B from AWS:

On-demand mode is a good option if any of the following are true:

- You create new tables with unknown workloads.
- You have unpredictable application traffic.
- You prefer the ease of paying for only what you use.

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ReadWriteCapacityMode.html>

upvoted 2 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: B

Technically both options A and B will work. But this statement 'traffic will be unpredictable' rules out option A, because 'provisioned mode' was made for scenarios where traffic is predictable.

So I will stick with B, because 'on-demand mode' is made for unpredictable traffic and instantly accommodates workloads as they ramp up or down.

upvoted 3 times

0628atv 2 years, 11 months ago

Selected Answer: A

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/AutoScaling.html>

upvoted 4 times

A retail company has several businesses. The IT team for each business manages its own AWS account. Each team account is part of an organization in AWS Organizations. Each team monitors its product inventory levels in an Amazon DynamoDB table in the team's own AWS account.

The company is deploying a central inventory reporting application into a shared AWS account. The application must be able to read items from all the teams' DynamoDB tables.

Which authentication option will meet these requirements MOST securely?

- A. Integrate DynamoDB with AWS Secrets Manager in the inventory application account. Configure the application to use the correct secret from Secrets Manager to authenticate and read the DynamoDB table. Schedule secret rotation for every 30 days.
- B. In every business account, create an IAM user that has programmatic access. Configure the application to use the correct IAM user access key ID and secret access key to authenticate and read the DynamoDB table. Manually rotate IAM access keys every 30 days.
- C. In every business account, create an IAM role named BU_ROLE with a policy that gives the role access to the DynamoDB table and a trust policy to trust a specific role in the inventory application account. In the inventory account, create a role named APP_ROLE that allows access to the STS AssumeRole API operation. Configure the application to use APP_ROLE and assume the crossaccount role BU_ROLE to read the DynamoDB table. **Most Voted**
- D. Integrate DynamoDB with AWS Certificate Manager (ACM). Generate identity certificates to authenticate DynamoDB. Configure the application to use the correct certificate to authenticate and read the DynamoDB table.

Correct Answer: C

Community vote distribution

C (100%)

Comments

cloudenthusiast **Highly Voted** 3 years ago

Selected Answer: C

IAM Roles: IAM roles provide a secure way to grant permissions to entities within AWS. By creating an IAM role in each business account named BU_ROLE with the necessary permissions to access the DynamoDB table, the access can be controlled at the IAM role level.

Cross-Account Access: By configuring a trust policy in the BU_ROLE that trusts a specific role in the inventory application account (APP_ROLE), you establish a trusted relationship between the two accounts.

Least Privilege: By creating a specific IAM role (BU_ROLE) in each business account and granting it access only to the required DynamoDB table, you can ensure that each team's table is accessed with the least privilege principle.

Security Token Service (STS): The use of STS AssumeRole API operation in the inventory application account allows the application to assume the cross-account role (BU_ROLE) in each business account.

upvoted 32 times

TariqKipkemei 2 years, 10 months ago

Well broken down..thank you :)

upvoted 3 times

MandAsh **Highly Voted** 1 year, 12 months ago

C because they have taken effort to explain it in details.. lol

upvoted 9 times

Bennyboy789 **Most Recent** 2 years, 9 months ago

Selected Answer: C

Keyword: IAM ROLES

upvoted 4 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: C

C is the most secure option to meet the requirements.

Using cross-account IAM roles and role chaining allows the inventory application to securely access resources in other accounts. The roles provide temporary credentials and can be permissions controlled.

upvoted 3 times

hsinchang 2 years, 10 months ago

Selected Answer: C

Looks complex, but IAM role seems more probable, I go with C.

upvoted 4 times

mattcl 2 years, 11 months ago

Why not A?

upvoted 3 times

awsgeek75 2 years, 5 months ago

A is wrong because it is incomplete. Just integrating with secrets manager doesn't give any access to DynamoDB.

upvoted 2 times

antropaws 2 years, 11 months ago

Selected Answer: C

It's complex, but looks C.

upvoted 2 times

eehhssaan 3 years ago

i'll go with C .. coming from two minds

upvoted 3 times

nosense 3 years ago

a or c. C looks like a more secure

upvoted 2 times

omoakin 3 years ago

CCCCCCCCCCC

upvoted 2 times

A company runs container applications by using Amazon Elastic Kubernetes Service (Amazon EKS). The company's workload is not consistent throughout the day. The company wants Amazon EKS to scale in and out according to the workload.

Which combination of steps will meet these requirements with the LEAST operational overhead? (Choose two.)

- A. Use an AWS Lambda function to resize the EKS cluster.
- B. Use the Kubernetes Metrics Server to activate horizontal pod autoscaling. **Most Voted**
- C. Use the Kubernetes Cluster Autoscaler to manage the number of nodes in the cluster. **Most Voted**
- D. Use Amazon API Gateway and connect it to Amazon EKS.
- E. Use AWS App Mesh to observe network activity.

Correct Answer: BC

Community vote distribution

BC (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 9 months ago

Selected Answer: BC

B and C are the correct options.

Using the Kubernetes Metrics Server (B) enables horizontal pod autoscaling to dynamically scale pods based on CPU/memory usage. This allows scaling at the application tier level.

The Kubernetes Cluster Autoscaler (C) automatically adjusts the number of nodes in the EKS cluster in response to pod resource requirements and events. This allows scaling at the infrastructure level.

upvoted 8 times

TariqKipkemei **Highly Voted** 1 year, 10 months ago

Selected Answer: BC

This is pretty straight forward.

Use the Kubernetes Metrics Server to activate horizontal pod autoscaling.

Use the Kubernetes Cluster Autoscaler to manage the number of nodes in the cluster.

upvoted 5 times

wsdasdasdqwda **Most Recent** 1 year, 7 months ago

K8S Metrics Server and Autoscaler => B and C

upvoted 4 times

james2033 1 year, 10 months ago

Selected Answer: BC

Kubernetes Metrics Server <https://docs.aws.amazon.com/eks/latest/userguide/metrics-server.html>

AWS Autoscaler <https://docs.aws.amazon.com/eks/latest/userguide/autoscaling.html> and <https://github.com/kubernetes/autoscaler/blob/master/cluster-autoscaler/cloudprovider/aws/README.md>

upvoted 4 times

cloudenthusiast 2 years ago

Selected Answer: BC

By combining the Kubernetes Cluster Autoscaler (option C) to manage the number of nodes in the cluster and enabling horizontal pod autoscaling (option B) with the Kubernetes Metrics Server, you can achieve automatic scaling of your EKS cluster and container applications based on workload demand. This approach minimizes operational overhead as it leverages built-in Kubernetes functionality and automation mechanisms.

upvoted 5 times

nosense 2 years ago

Selected Answer: BC

b and c is right

upvoted 2 times

A company runs a microservice-based serverless web application. The application must be able to retrieve data from multiple Amazon DynamoDB tables. A solutions architect needs to give the application the ability to retrieve the data with no impact on the baseline performance of the application.

Which solution will meet these requirements in the MOST operationally efficient way?

- A. AWS AppSync pipeline resolvers **Most Voted**
- B. Amazon CloudFront with Lambda@Edge functions
- C. Edge-optimized Amazon API Gateway with AWS Lambda functions
- D. Amazon Athena Federated Query with a DynamoDB connector

Correct Answer: A

Community vote distribution



Comments

elmogy **Highly Voted** 3 years ago

just passed yesterday 30-05-23, around 75% of the exam came from here, some with light changes.
upvoted 39 times

omoakin **Highly Voted** 3 years ago

Great work made it to the last question. Goodluck to you all
upvoted 21 times

MostofMichelle 3 years ago

good luck to you as well.
upvoted 7 times

cyber_bedouin 2 years, 7 months ago

Thanks. Do you think the questions after 500 are relevant, they seem to be above associate level (harder)
upvoted 10 times

Drew3000 2 years, 2 months ago

I am jealous that 10 months ago there were only 523 questions :(
upvoted 10 times

Awsbeginner87 2 years, 2 months ago

right...how to go through 840 questions :(
upvoted 3 times

Rhydian25 1 year, 11 months ago

Currently 904 :)
upvoted 5 times

4729e6c 1 year, 4 months ago

1019 now and hard to finish them all in a month :(
upvoted 3 times

37a869e 1 year, 3 months ago

The first 25% is free though. "Only" 765 questions to do in a month. But not sure I'm gonna make it
upvoted 1 times

Awsbeginner87 2 years, 2 months ago

moderators answer is wrong for many questions. which one should we consider:
upvoted 1 times

lalaraj Most Recent 10 months, 3 weeks ago

Selected Answer: A

AWS AppSync (Option A) is the best choice because:

It provides a unified GraphQL interface to query multiple DynamoDB tables efficiently.

Pipeline resolvers enable batching and chaining requests, reducing load on DynamoDB.

Built-in caching reduces repeated queries, preserving baseline performance.

Fully serverless and managed, minimizing operational overhead.

upvoted 2 times

Charlesvg 11 months, 2 weeks ago

Selected Answer: A

App Sync is the best here

upvoted 2 times

CTao 1 year, 4 months ago

Selected Answer: A

AWS AppSync: It provides a managed GraphQL service that simplifies data access across multiple sources, including DynamoDB tables.

upvoted 4 times

4729e6c 1 year, 4 months ago

Selected Answer: A

Deepseek and Qwen are just published and they both say A.

upvoted 2 times

FlyingHawk 1 year, 4 months ago

Selected Answer: A

AWS AppSync is a managed GraphQL service that allows applications to query multiple data sources, such as DynamoDB tables, in a single API call. It is highly efficient and does not require maintaining backend infrastructure.

Pipeline resolvers in AppSync enable the combination of multiple resolvers to fetch data from different sources, aggregate the data, and return it in a single response. This approach is operationally efficient, has minimal impact on baseline performance, and is designed for applications interacting with multiple DynamoDB tables.

<https://docs.aws.amazon.com/appsync/latest/devguide/tutorial-pipeline-resolvers.html>

upvoted 2 times

FlyingHawk 1 year, 6 months ago

Why the suggested answer is D? I read the comments, feel A is more suitable as the requirement is to give the application the ability to retrieve the data with no impact on the baseline performance of the application. D might have affect the baseline performance, A is more effective

upvoted 3 times

LeonSauveterre 1 year, 6 months ago

Selected Answer: A

A - Correct. Retrieves data from multiple data sources.

B - Incorrect. Lambda@Edge extends serverless processing to CloudFront, but it is typically used for content manipulation, caching, or geographic distribution.

C - That works but requires a lot of things to do manually. Why C when you can do with A?

D - Amazon Athena Federated Query is designed for ad hoc querying of data across different sources BUT it has impact on performance, and not operationally efficient.

PS: "ad hoc" means "for this specific purpose". Such requests can lack efficiency for repetitive tasks because it's not optimized or automated, since they're naturally exploratory or tailored to specific immediate needs.

upvoted 4 times

kevindu 1 year, 10 months ago

Selected Answer: A

Is there anyone who has recently passed the exam who can tell me approximately how many of the original questions are in the actual exam?

upvoted 3 times

use4u 2 years, 1 month ago

let me know

for most of questions that system answer are correct or comment answer are correct?

upvoted 2 times

Sergiuss95 2 years, 1 month ago

Comment answer

upvoted 2 times

osmk 2 years, 3 months ago

Selected Answer: D

https://docs.amazonaws.cn/en_us/athena/latest/ug/connect-to-a-data-source.html

upvoted 6 times

FlyingHawk 1 year, 4 months ago

Athena queries are serverless and scalable but are not optimized for low-latency, high-concurrency scenarios typically required for a microservice-based operational web application. Additionally, query execution time and costs can become a concern if frequent access to DynamoDB tables is required.

upvoted 2 times

upliftinghut 2 years, 4 months ago

Selected Answer: D

key word is most operational effective => D requires no coding

upvoted 3 times

awsgeek75 2 years, 5 months ago

Selected Answer: D

I'll go with D as ABC looks too much work or irrelevant. Although not sure how AFQ actually achieves the read without impacting performance.

upvoted 2 times

pentium75 2 years, 5 months ago

Selected Answer: D

Not A - Pipe Resolvers require coding, would not consider that 'operationally efficient'

Not B - CloudFront caches web content at the edge, not DynamoDB query results for apps

Not C - Neither API Gateway or Lambda have anything to do with DynamoDB performance

D - Can do exactly that

upvoted 8 times

aws94 2 years, 6 months ago

Selected Answer: A

I am not an expert but I used Bing+Gemini+Chatgpt=AAA

upvoted 2 times

ekisako 2 years, 6 months ago

Selected Answer: A

multiple database tables = AppSync pipeline resolvers

upvoted 7 times


A company wants to analyze and troubleshoot Access Denied errors and Unauthorized errors that are related to IAM permissions. The company has AWS CloudTrail turned on.

Which solution will meet these requirements with the LEAST effort?

- A. Use AWS Glue and write custom scripts to query CloudTrail logs for the errors.
- B. Use AWS Batch and write custom scripts to query CloudTrail logs for the errors.
- C. Search CloudTrail logs with Amazon Athena queries to identify the errors. **Most Voted**
- D. Search CloudTrail logs with Amazon QuickSight. Create a dashboard to identify the errors.

Correct Answer: C

Community vote distribution



Comments

awsgeek75 **Highly Voted** 1 year, 10 months ago

Selected Answer: C

<https://docs.aws.amazon.com/athena/latest/ug/cloudtrail-logs.html>

When troubleshooting you will want to query specific things in the log and Athena provides query language for that.

Quick Sight is data analytics and visualisation tool. You can use it to aggregate data and maybe make a dashboard for number of errors by type etc but that doesn't help you troubleshoot anything.

C is correct

upvoted 6 times

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: C

Athena allows you to run SQL queries on data in Amazon S3, including CloudTrail logs. It is the easiest way to query the logs and identify specific errors without needing to write any custom code or scripts.

With Athena, you can write simple SQL queries to filter the CloudTrail logs for the "AccessDenied" and "UnauthorizedOperation" error codes. This will return the relevant log entries that you can then analyze.

upvoted 5 times

In2718 **Most Recent** 10 months, 3 weeks ago

Selected Answer: D

Both C and D work. I chose D because the question says LEAST effort. Setting up QuickSight is less effort than Athena queries, and the question does not ask to resolve, but just analyze and troubleshoot. Again, not a well written question.

upvoted 1 times

pentium75 1 year, 11 months ago

Selected Answer: C

"Search CloudTrail logs with Amazon QuickSight", that doesn't work. QuickSight can visualize Athena query results, so "search CloudTrail logs with Amazon Athena, then create a dashboard with Amazon QuickSight" would make sense. But QuickSight without Athena won't work.

upvoted 4 times

bogobob 2 years ago

Selected Answer: D

The question asks specifically to "analyze and troubleshoot". While Athena is easy to get the data, you then just have a list of logs. Not very useful to troubleshoot...

upvoted 1 times

awsgeek75 1 year, 10 months ago

How will pretty pictures in QuickSight help with troubleshooting?

upvoted 2 times

pentium75 1 year, 11 months ago

But without Athena, there is nothing you can visualize in QuickSight.

upvoted 3 times

NickGordon 2 years, 1 month ago

Selected Answer: D

Quick Sight is an analytics tool. Sounds like a LEAST effort option

upvoted 3 times

TariqKipkemei 2 years, 4 months ago

Selected Answer: C

C for me. Using Athena with CloudTrail logs is a powerful way to enhance your analysis of AWS service activity. For example, you can use queries to identify trends and further isolate activity by attributes, such as source IP address or user.

<https://docs.aws.amazon.com/athena/latest/ug/cloudtrail-logs.html#:~:text=CloudTrail%20Lake%20documentation.-,Using%20Athena,-with%20CloudTrail%20logs>

upvoted 2 times

james2033 2 years, 4 months ago

Selected Answer: C

IAM and CloudTrail <https://docs.aws.amazon.com/IAM/latest/UserGuide/cloudtrail-integration.html#stscloudtrailexample-assumerole> .

Query CloudTrail logs by Athena <https://docs.aws.amazon.com/athena/latest/ug/cloudtrail-logs.html#tips-for-querying-cloudtrail-logs#tips-for-querying-cloudtrail-logs>

upvoted 2 times

james2033 2 years, 4 months ago

Choose C, not D, because need "analyze and troubleshoot", not just see on dashboard (in D).

upvoted 2 times

live_reply_developers 2 years, 5 months ago

Selected Answer: C

Amazon Athena is an interactive query service provided by AWS that enables you to analyze data , is a little bit more suitable integrated with cloud trail that permit to verify WHO accessed the service.

upvoted 2 times

manuh 2 years, 5 months ago

Selected Answer: C

Dashboard isnt requires. Also refer to this <https://repost.aws/knowledge-center/troubleshoot-iam-permission-errors>

upvoted 2 times

haoAWS 2 years, 5 months ago

Selected Answer: D

I am struggling for the C and D for a long time, and ask the chatGPT. The chatGPT says D is better, since Athena requires more expertise on SQL.

upvoted 1 times

antropaws 2 years, 5 months ago

Selected Answer: D

Both C and D are feasible. I vote for D:

Amazon QuickSight supports logging the following actions as events in CloudTrail log files:

- Whether the request was made with root or AWS Identity and Access Management user credentials
- Whether the request was made with temporary security credentials for an IAM role or federated user
- Whether the request was made by another AWS service

<https://docs.aws.amazon.com/quicksight/latest/user/logging-using-cloudtrail.html>

upvoted 1 times

PCWu 2 years, 5 months ago

Selected Answer: C

The Answer will be C:

Need to use Athena to query keywords and sort out the error logs.

D: No need to use Amazon QuickSight to create the dashboard.

upvoted 2 times

Axeashes 2 years, 5 months ago

Selected Answer: C

"Using Athena with CloudTrail logs is a powerful way to enhance your analysis of AWS service activity."
<https://docs.aws.amazon.com/athena/latest/ug/cloudtrail-logs.html>

upvoted 2 times

oras2023 2 years, 6 months ago

Selected Answer: C

Analyse and TROUBLESHOOT, look like Athena

upvoted 1 times

oras2023 2 years, 5 months ago

<https://docs.aws.amazon.com/athena/latest/ug/cloudtrail-logs.html>

upvoted 2 times

alexandercamachop 2 years, 6 months ago

Selected Answer: D

It specifies analyze, not query logs.

Which is why option D is the best one as it provides dashboards to analyze the logs.

upvoted 3 times

A company wants to add its existing AWS usage cost to its operation cost dashboard. A solutions architect needs to recommend a solution that will give the company access to its usage cost programmatically. The company must be able to access cost data for the current year and forecast costs for the next 12 months.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Access usage cost-related data by using the AWS Cost Explorer API with pagination. **Most Voted**
- B. Access usage cost-related data by using downloadable AWS Cost Explorer report .csv files.
- C. Configure AWS Budgets actions to send usage cost data to the company through FTP.
- D. Create AWS Budgets reports for usage cost data. Send the data to the company through SMTP.

Correct Answer: A

Community vote distribution

A (100%)

Comments

BrijMohan08 **Highly Voted** 2 years, 3 months ago

Selected Answer: A

Keyword

12 months, API Support

<https://docs.aws.amazon.com/cost-management/latest/userguide/ce-what-is.html>

upvoted 6 times

alexandercamachop **Highly Voted** 2 years, 6 months ago

Selected Answer: A

Answer is: A

says dashboard = Cost Explorer, therefor C & D are eliminated.

also says programmatically, means non manual intervention therefor API.

upvoted 5 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: A

Programmatically + LEAST overhead = API

upvoted 4 times

TariqKipkemei 2 years ago

Selected Answer: A

access to its usage cost programmatically = AWS Cost Explorer API

upvoted 3 times

thanhnv142 2 years, 1 month ago

A: correct

1. programatically = API

2. In the next 12 months = cost explorer

upvoted 4 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: A

Access usage cost-related data by using the AWS Cost Explorer API with pagination

upvoted 3 times

wendaz 2 years, 2 months ago

don't repeat the answer, it is useless... explain , okay? i have seen your replies many time just to copy the options.. it makes no sense...

upvoted 9 times

james2033 2 years, 4 months ago

Selected Answer: A

AWS Cost Explorer API with paginated request: <https://docs.aws.amazon.com/cost-management/latest/userguide/ce-api-best-practices.html#ce-api-best-practices-optimize-costs>

upvoted 3 times

MrAWSAssociate 2 years, 5 months ago

Selected Answer: A

From AWS Documentation*:

"You can view your costs and usage using the Cost Explorer user interface free of charge. You can also access your data programmatically using the Cost Explorer API. Each paginated API request incurs a charge of \$0.01. You can't disable Cost Explorer after you enable it."

* Source:

<https://docs.aws.amazon.com/cost-management/latest/userguide/ce-what-is.html>

<https://docs.aws.amazon.com/AWSJavaScriptSDK/v3/latest/clients/client-cost-explorer/interfaces/costexplorerpaginationconfiguration.html>

upvoted 5 times

oras2023 2 years, 6 months ago

Selected Answer: A

least operational overhead = API access

upvoted 4 times

oras2023 2 years, 6 months ago

least operational overhead = API access

upvoted 2 times


A solutions architect is reviewing the resilience of an application. The solutions architect notices that a database administrator recently failed over the application's Amazon Aurora PostgreSQL database writer instance as part of a scaling exercise. The failover resulted in 3 minutes of downtime for the application.

Which solution will reduce the downtime for scaling exercises with the LEAST operational overhead?

- A. Create more Aurora PostgreSQL read replicas in the cluster to handle the load during failover.
- B. Set up a secondary Aurora PostgreSQL cluster in the same AWS Region. During failover, update the application to use the secondary cluster's writer endpoint.
- C. Create an Amazon ElastiCache for Memcached cluster to handle the load during failover.
- D. Set up an Amazon RDS proxy for the database. Update the application to use the proxy endpoint. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

alexandercamachop **Highly Voted** 3 years ago

Selected Answer: D

D is the correct answer.

It is talking about the write database. Not reader.

Amazon RDS proxy allows you to automatically route write request to the healthy writer, minimizing downtime.

upvoted 13 times

nilandd44gg 2 years, 10 months ago

One of the benefits of Amazon RDS Proxy is that it can improve application recovery time after database failovers. While RDS Proxy supports both MySQL as well as PostgreSQL engines, in this post, we will use a MySQL test workload to demonstrate how RDS Proxy reduces client recovery time after failover by up to 79% for Amazon Aurora MySQL and by up to 32% for Amazon RDS for MySQL.

<https://aws.amazon.com/blogs/database/improving-application-availability-with-amazon-rds-proxy/>

<https://aws.amazon.com/rds/proxy/faqs/>

upvoted 6 times

bora4motion **Most Recent** 1 year, 1 month ago

Selected Answer: D

I'm going with D. It will reduce the failover time between DBs.

upvoted 1 times

ugaga 1 year, 11 months ago

AWS EXAM is also an AWS promotion...

RDS Proxy really does not reduce recovery time if you don't have alternative RDS instance. You will see if you understand what proxy is. But this is exam (= promotion) of AWS, D could be an answer, because some AWS document says so.

upvoted 2 times

pentium75 2 years, 5 months ago

Selected Answer: D

"RDS Proxy reduces client recovery time after failover by up to 79% for Amazon Aurora MySQL "

<https://aws.amazon.com/de/blogs/database/improving-application-availability-with-amazon-rds-proxy/>

upvoted 4 times

ftaws 2 years, 5 months ago

Selected Answer: B

RDS Proxy is used for DB timeout not downtime.

How to reduce downtime with RDS Proxy?

There is no change downtime if we use RDS Proxy.

upvoted 3 times

pentium75 2 years, 5 months ago

"How to reduce downtime with RDS Proxy", by eliminating the need for the application to retrieve the new DNS record after the old one times out.

upvoted 4 times

MatAlves 1 year, 8 months ago

I would never have thought about DNS records in this question. Nice catch!

upvoted 2 times

pentium75 2 years, 5 months ago

"RDS Proxy reduces client recovery time after failover by up to 79% for Amazon Aurora MySQL "

<https://aws.amazon.com/de/blogs/database/improving-application-availability-with-amazon-rds-proxy/>

upvoted 3 times

Cyberkayu 2 years, 5 months ago

Selected Answer: B

they are using Aurora, RDS proxy dont work here

Answer B

upvoted 2 times

pentium75 2 years, 5 months ago

Wrong: <https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/rds-proxy.html>

upvoted 2 times

pentium75 2 years, 5 months ago

"RDS Proxy reduces client recovery time after failover by up to 79% for Amazon Aurora MySQL "

<https://aws.amazon.com/de/blogs/database/improving-application-availability-with-amazon-rds-proxy/>

upvoted 2 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: D

D. Set up an Amazon RDS proxy for the database. Update the application to use the proxy endpoint.

upvoted 3 times

hachiri 2 years, 9 months ago

point is Aurora Multi-Master

Set up a secondary Aurora PostgreSQL cluster in the *same* AWS Region

upvoted 2 times

hachiri 2 years, 9 months ago

I mean correct is B

upvoted 1 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: C

Availability is the main requirement here. Even if RDS proxy is used, it will still find the writer instance unavailable during the scaling exercise. Best option is to create an Amazon ElastiCache for Memcached cluster to handle the load during the scaling operation.

upvoted 1 times

pentium75 2 years, 5 months ago

Failover is faster with RDS proxy

upvoted 2 times

wizcloudifa 2 years, 1 month ago

that will only account for the read operation part of it, we are concerned with the write operations here

upvoted 2 times

pentium75 2 years, 5 months ago

"RDS Proxy reduces client recovery time after failover by up to 79% for Amazon Aurora MySQL "

<https://aws.amazon.com/de/blogs/database/improving-application-availability-with-amazon-rds-proxy/>

upvoted 2 times

AshishRocks 3 years ago

Set up an Amazon RDS proxy for the database. Update the application to use the proxy endpoint.

D is the answer

upvoted 4 times

A company has a regional subscription-based streaming service that runs in a single AWS Region. The architecture consists of web servers and application servers on Amazon EC2 instances. The EC2 instances are in Auto Scaling groups behind Elastic Load Balancers. The architecture includes an Amazon Aurora global database cluster that extends across multiple Availability Zones.


The company wants to expand globally and to ensure that its application has minimal downtime.

Which solution will provide the MOST fault tolerance?

- A. Extend the Auto Scaling groups for the web tier and the application tier to deploy instances in Availability Zones in a second Region. Use an Aurora global database to deploy the database in the primary Region and the second Region. Use Amazon Route 53 health checks with a failover routing policy to the second Region.
- B. Deploy the web tier and the application tier to a second Region. Add an Aurora PostgreSQL cross-Region Aurora Replica in the second Region. Use Amazon Route 53 health checks with a failover routing policy to the second Region. Promote the secondary to primary as needed.
- C. Deploy the web tier and the application tier to a second Region. Create an Aurora PostgreSQL database in the second Region. Use AWS Database Migration Service (AWS DMS) to replicate the primary database to the second Region. Use Amazon Route 53 health checks with a failover routing policy to the second Region.
- D. Deploy the web tier and the application tier to a second Region. Use an Amazon Aurora global database to deploy the database in the primary Region and the second Region. Use Amazon Route 53 health checks with a failover routing policy to the second Region. Promote the secondary to primary as needed. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

TariqKipkemei **Highly Voted** 2 years, 10 months ago

Selected Answer: D

Auto Scaling groups can span Availability Zones, but not AWS regions.

Hence the best option is to deploy the web tier and the application tier to a second Region. Use an Amazon Aurora global database to deploy the database in the primary Region and the second Region. Use Amazon Route 53 health checks with a failover routing policy to the second Region. Promote the secondary to primary as needed.

upvoted 22 times

awsgeek75 **Most Recent** 2 years, 5 months ago

Selected Answer: D

A: Not possible for autoscaling across regions

BC: Using PostgreSQL, not sure why?

D: MOST fault tolerant != MOST scalable. This gives least downtime.

upvoted 4 times

potomac 2 years, 7 months ago

Selected Answer: D

EC2 Auto Scaling groups are regional constructs. They can span Availability Zones, but not AWS regions

upvoted 3 times

thanhvn142 2 years, 7 months ago

527:

D is correct:

- B & C is not correct because it mentions Aurora PostgreSQL which is not mentioned in the question

- A is not correct because Auto scaling group can not span regions

upvoted 4 times

wsdasdasdqwda 2 years, 7 months ago

Simple as that.

upvoted 2 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: D

Using an Aurora global database that spans both the primary and secondary regions provides automatic replication and failover capabilities for the database tier.

Deploying the web and application tiers to a second region provides fault tolerance for those components.

Using Route53 health checks and failover routing will route traffic to the secondary region if the primary region becomes unavailable.

This provides fault tolerance across all tiers of the architecture while minimizing downtime. Promoting the secondary database to primary ensures the second region can continue operating if needed.

A is close, but doesn't provide an automatic database failover capability.

B and C provide database replication, but not automatic failover.

So D is the most comprehensive and fault tolerant architecture.

upvoted 4 times

JA2018 1 year, 6 months ago

PostgreSQL in options b & C could be a distractor.... PostgreSQL not mentioned in stem

upvoted 1 times

Zox42 2 years, 11 months ago

Selected Answer: D

Answer D

upvoted 2 times

Zuit 2 years, 11 months ago

Selected Answer: D

D seems fitting: Global Database and deploying it in the new region

upvoted 2 times

MrAWSAssociate 2 years, 11 months ago

Selected Answer: B

B is correct!

upvoted 1 times

manuh 2 years, 11 months ago

Replicated db doesn't mean they will act as a single db once the transfer is completed. Global db is the correct approach

upvoted 2 times

JA2018 1 year, 6 months ago

PostgreSQL in options b & C could be a distractor.... PostgreSQL not mentioned in stem

upvoted 1 times

r3mo 2 years, 12 months ago

"D" is the answer: because AWS Aurora Global Database allows you to read and write from any region in the global cluster. This enables you to distribute read and write workloads globally, improving performance and reducing latency. Data is replicated synchronously across regions, ensuring strong consistency.

upvoted 4 times

Henrytml 3 years ago

Selected Answer: A

A is the only answer remaining using ELB, both Web/App/DB has been taking care with replicating in 2nd region, lastly route 53 for failover over multiple regions

upvoted 1 times

Henrytml 2 years, 12 months ago

i will revoke my answer to standby web in 2nd region, instead of trigger to scale out

upvoted 1 times

manuh 2 years, 11 months ago

also ASG can't span beyond a region

upvoted 2 times

alexandercamachop 3 years ago

Selected Answer: D

B&C are discarded.

The answer is between A and D.

I would go with D because it explicitly created this web / app tier in second region, instead A just autoscales into a secondary region, rather than always having resources in this second region.

upvoted 4 times

JA2018 1 year, 6 months ago

as mentioned ASG cannot span across different AWS regions

upvoted 1 times

A data analytics company wants to migrate its batch processing system to AWS. The company receives thousands of small data files periodically during the day through FTP. An on-premises batch job processes the data files overnight. However, the batch job takes hours to finish running.

The company wants the AWS solution to process incoming data files as soon as possible with minimal changes to the FTP clients that send the files. The solution must delete the incoming data files after the files have been processed successfully. Processing for each file needs to take 3-8 minutes.

Which solution will meet these requirements in the MOST operationally efficient way?

- A. Use an Amazon EC2 instance that runs an FTP server to store incoming files as objects in Amazon S3 Glacier Flexible Retrieval. Configure a job queue in AWS Batch. Use Amazon EventBridge rules to invoke the job to process the objects nightly from S3 Glacier Flexible Retrieval. Delete the objects after the job has processed the objects.
- B. Use an Amazon EC2 instance that runs an FTP server to store incoming files on an Amazon Elastic Block Store (Amazon EBS) volume. Configure a job queue in AWS Batch. Use Amazon EventBridge rules to invoke the job to process the files nightly from the EBS volume. Delete the files after the job has processed the files.
- C. Use AWS Transfer Family to create an FTP server to store incoming files on an Amazon Elastic Block Store (Amazon EBS) volume. Configure a job queue in AWS Batch. Use an Amazon S3 event notification when each file arrives to invoke the job in AWS Batch. Delete the files after the job has processed the files.
- D. Use AWS Transfer Family to create an FTP server to store incoming files in Amazon S3 Standard. Create an AWS Lambda function to process the files and to delete the files after they are processed. Use an S3 event notification to invoke the Lambda function when the files arrive. **Most Voted**

Correct Answer: D

Community vote distribution

D (96%)

4%

Comments

pentium75 **Highly Voted** 1 year, 11 months ago

Selected Answer: D

Obviously we choose AWS Transfer Family over hosting the FTP server ourselves on an EC2 instance. And "process incoming data files as soon as possible" -> trigger Lambda when files arrive. Lambda functions can run up to 15 minutes, it takes "3-8 minutes" per file -> works.

AWS Batch just schedules jobs, but these still need to run somewhere (Lambda, Fargate, EC2).

upvoted 5 times

wizcloudifa 1 year, 7 months ago

well if you opt for unmanaged compute option in AWS Batch in that case you don't have to worry about the compute, AWS batch automatically provisions it for you

upvoted 2 times

wsdasdasdqwda **Most Recent** 2 years, 1 month ago

FTP => AWS Transfer Family, => C or D, but in C is used EBS not S3 which needs EC2 and in general is more complex => very clear D.

upvoted 2 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: D

The key points:

Use AWS Transfer Family for the FTP server to receive files directly into S3. This avoids managing FTP servers.

Process each file as soon as it arrives using Lambda triggered by S3 events. Lambda provides fast processing time per file. Lambda can also delete files after processing succeeds.

Options A, B, C involve more operational overhead of managing FTP servers and batch jobs. Processing latency would be higher waiting for batch windows.

Storing files in Glacier (Option A) adds latency for retrieving files.

upvoted 2 times

hsinchang 2 years, 4 months ago

Selected Answer: D

Processing for each file needs to take 3-8 minutes clearly indicates Lambda functions.

upvoted 2 times

TariqKipkemei 2 years, 4 months ago

Selected Answer: D

Process incoming data files with minimal changes to the FTP clients that send the files = AWS Transfer Family.

Process incoming data files as soon as possible = S3 event notification.

Processing for each file needs to take 3-8 minutes = AWS Lambda function.

Delete file after processing = AWS Lambda function.

upvoted 4 times

antropaws 2 years, 5 months ago

Selected Answer: D

Most likely D.

upvoted 2 times

r3mo 2 years, 6 months ago

"D" Since each file takes 3-8 minutes to process the lambda function can process the data file without a problem.

upvoted 2 times

maver144 2 years, 6 months ago

Selected Answer: D

You cannot setup AWS Transfer Family to save files into EBS.

upvoted 3 times

oras2023 2 years, 5 months ago

<https://aws.amazon.com/aws-transfer-family/>

upvoted 2 times

secdgs 2 years, 6 months ago

Selected Answer: D

D. Because

1. process immediate when file transfer to S3 not wait for process several file in one time.

2. takes 3-8 can use Lambda.

C. Wrong because AWS Batch is used for running large-scale or large amount of data in one time.

upvoted 2 times

Aymanovitchy 2 years, 6 months ago

To meet the requirements of processing incoming data files as soon as possible with minimal changes to the FTP clients, and deleting the files after successful processing, the most operationally efficient solution would be:

D. Use AWS Transfer Family to create an FTP server to store incoming files in Amazon S3 Standard. Create an AWS Lambda function to process the files and delete them after processing. Use an S3 event notification to invoke the Lambda function when the files arrive.

upvoted 2 times

bajwa360 2 years, 6 months ago

Selected Answer: D

It should be D as lambda is more operationally viable solution given the fact each processing takes 3-8 minutes that lambda can handle

upvoted 2 times

alexandercamachop 2 years, 6 months ago

Selected Answer: C

Answer has to be between C or D.

Because Transfer Family is obvious do to FTP.

Now I would go with C because it uses AWS Batch, which makes more sense for Batch processing rather than AWS Lambda.

upvoted 1 times

pentium75 1 year, 11 months ago

Also Batch just triggers jobs, they still need to run somewhere (like in Lambda).

upvoted 2 times

pentium75 1 year, 11 months ago

Why? "Process incoming data files as soon as possible", by triggering the Lambda function when files arrive. Batch is for scheduled jobs.

upvoted 2 times

Bill1000 2 years, 6 months ago

I am between C and D. My reason is:

"The company wants the AWS solution to process incoming data files as soon as possible with minimal changes to the FTP clients that send the files."

upvoted 3 times

A company is migrating its workloads to AWS. The company has transactional and sensitive data in its databases. The company wants to use AWS Cloud solutions to increase security and reduce operational overhead for the databases.

Which solution will meet these requirements?

- A. Migrate the databases to Amazon EC2. Use an AWS Key Management Service (AWS KMS) AWS managed key for encryption.
- B. Migrate the databases to Amazon RDS Configure encryption at rest. **Most Voted**
- C. Migrate the data to Amazon S3 Use Amazon Macie for data security and protection
- D. Migrate the database to Amazon RDS. Use Amazon CloudWatch Logs for data security and protection.

Correct Answer: B

Community vote distribution

B (100%)

Comments

AshishRocks **Highly Voted** 3 years ago

B is the answer

Why not C - Option C suggests migrating the data to Amazon S3 and using Amazon Macie for data security and protection. While Amazon Macie provides advanced security features for data in S3, it may not be directly applicable or optimized for databases, especially for transactional and sensitive data. Amazon RDS provides a more suitable environment for managing databases.

upvoted 11 times

oras2023 **Highly Voted** 3 years ago

Selected Answer: B

B. Migrate the databases to Amazon RDS Configure encryption at rest.

Looks like best option

upvoted 5 times

venomblade **Most Recent** 1 year, 6 months ago

Selected Answer: B

S3 and Macie works too. However, the question mentioned about transactional data. RDS is much better for running transactional workloads. Thus, RDS is the correct answer here.

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: B

A: Operational overhead of EC2 and whatever DB is running on it

C: Macie is not for data security, it's for identifying PII and sensitive data

D: CloudWatch is for cloud events and does not secure databases

B: RDS is managed so least operational overhead. Encryption at rest means security

upvoted 3 times

allanm 10 months ago

Macie is actually used to protect sensitive data. It's just that C cannot be used because the data resides in databases and unless you have a service to migrate from database to S3, this won't work

upvoted 1 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: B

Migrate the databases to Amazon RDS Configure encryption at rest.

upvoted 4 times

wendaz 2 years, 8 months ago

down voted.

upvoted 1 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: B

Reduce Ops = Migrate the databases to Amazon RDS Configure encryption at rest

upvoted 4 times

alexandercamachop 3 years ago

Selected Answer: B

B for sure.

First the correct is Amazon RDS, then encryption at rest makes the database secure.

upvoted 4 times

A company has an online gaming application that has TCP and UDP multiplayer gaming capabilities. The company uses Amazon Route 53 to point the application traffic to multiple Network Load Balancers (NLBs) in different AWS Regions. The company needs to improve application performance and decrease latency for the online game in preparation for user growth.

Which solution will meet these requirements?

- A. Add an Amazon CloudFront distribution in front of the NLBs. Increase the Cache-Control max-age parameter.
- B. Replace the NLBs with Application Load Balancers (ALBs). Configure Route 53 to use latency-based routing.
- C. Add AWS Global Accelerator in front of the NLBs. Configure a Global Accelerator endpoint to use the correct listener ports.
- D. Add an Amazon API Gateway endpoint behind the NLBs. Enable API caching. Override method caching for the different stages.

Most Voted

Correct Answer: C

Community vote distribution

C (100%)

Comments

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: C

The key considerations are:

The application uses TCP and UDP for multiplayer gaming, so Network Load Balancers (NLBs) are appropriate. AWS Global Accelerator can be added in front of the NLBs to improve performance and reduce latency by intelligently routing traffic across AWS Regions and Availability Zones.

Global Accelerator provides static anycast IP addresses that act as a fixed entry point to application endpoints in the optimal AWS location. This improves availability and reduces latency.

The Global Accelerator endpoint can be configured with the correct NLB listener ports for TCP and UDP.

upvoted 6 times

TariqKipkemei **Highly Voted** 2 years, 4 months ago

Selected Answer: C

TCP, UDP, Gaming = global accelerator and Network Load Balancer

upvoted 5 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: C

A: CloudFront is for caching. Not required

B: ALB is for HTTP layer, won't help with TCP UDP issues

D: API Gateway, API Caching total rubbish, ignore this option

C: Is correct as Global Accelerator uses unicast for reducing latency globally.

upvoted 2 times

Henrytml 2 years, 5 months ago

Selected Answer: C

only b and c handle TCP/UDP, and C comes with accelerator to enhance performance

upvoted 2 times

manuh 2 years, 5 months ago

Does alb handle udp? Can u share a source?

upvoted 1 times

alexandercamachop 2 years, 6 months ago

Selected Answer: C

UDP and TCP is AWS Global accelerator as it works in the Transportation layer.
Now this with NLB is perfect.

upvoted 3 times

oras2023 2 years, 6 months ago

Selected Answer: C

C is helping to reduce latency for end clients

upvoted 3 times

A company needs to integrate with a third-party data feed. The data feed sends a webhook to notify an external service when new data is ready for consumption. A developer wrote an AWS Lambda function to retrieve data when the company receives a webhook callback. The developer must make the Lambda function available for the third party to call.

Which solution will meet these requirements with the MOST operational efficiency?

A. Create a function URL for the Lambda function. Provide the Lambda function URL to the third party for the webhook.

Most Voted

B. Deploy an Application Load Balancer (ALB) in front of the Lambda function. Provide the ALB URL to the third party for the webhook.

C. Create an Amazon Simple Notification Service (Amazon SNS) topic. Attach the topic to the Lambda function. Provide the public hostname of the SNS topic to the third party for the webhook.

D. Create an Amazon Simple Queue Service (Amazon SQS) queue. Attach the queue to the Lambda function. Provide the public hostname of the SQS queue to the third party for the webhook.

Correct Answer: A

Community vote distribution

A (100%)

Comments

TariqKipkemei **Highly Voted** 2 years, 10 months ago

Selected Answer: A

A function URL is a dedicated HTTP(S) endpoint for your Lambda function. When you create a function URL, Lambda automatically generates a unique URL endpoint for you.

upvoted 9 times

james2033 **Highly Voted** 2 years, 10 months ago

Selected Answer: A

Keyword "Lambda function" and "webhook". See <https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-saas-furls.html#create-stripe-cfn-stack>

upvoted 6 times

emakid **Most Recent** 1 year, 11 months ago

Selected Answer: A

AWS Lambda can provide a URL to call using Function URLs. This is a relatively new feature in AWS Lambda that allows you to create HTTPS endpoints for your Lambda functions, making it easy to invoke the function directly over the web.

Key Features of Lambda Function URLs:

Direct Access: Provides a simple and direct way to call a Lambda function via an HTTP(S) request.

Easy Configuration: You can create a function URL for a Lambda function using the AWS Management Console, AWS CLI, or AWS SDKs.

Managed Service: AWS manages the infrastructure for you, handling scaling, patching, and maintenance.

Security: You can configure authentication and authorization using AWS IAM or AWS Lambda function URL settings.

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: A

Apart from simplest and most operational, I think A is the only option that will work!
BCD cannot even be implemented in real world imho. Happy to be corrected

upvoted 3 times

Orit 2 years, 6 months ago

B is the answer The best solution to make the Lambda function available for the third party to call with the MOST operational efficiency is to deploy an Application Load Balancer (ALB) in front of the Lambda function and provide the ALB URL to the third party for the webhook. This solution is the most efficient because it allows the third party to call the Lambda function without having to worry about managing the Lambda function's availability or scaling. The ALB will automatically distribute traffic across multiple Lambda functions, if necessary, and will also provide redundancy in case of a failure.

upvoted 1 times

hro 2 years, 2 months ago

I believe you are correct. Lambda functions as targets - implementing ALBs

<https://docs.aws.amazon.com/elasticloadbalancing/latest/application/lambda-functions.html>

upvoted 1 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: A

The key points:

A Lambda function needs to be invoked by a third party via a webhook.

Using a function URL provides a direct invoke endpoint for the Lambda function. This is simple and efficient.

Options B, C, and D insert unnecessary components like ALB, SNS, SQS between the webhook and the Lambda function. These add complexity without benefit.

A function URL can be generated and provided to the third party quickly without additional infrastructure.

upvoted 5 times

Abrar2022 2 years, 11 months ago

Selected Answer: A

key word: Lambda function URLs

upvoted 3 times

maver144 2 years, 12 months ago

Selected Answer: A

<https://docs.aws.amazon.com/lambda/latest/dg/lambda-urls.html>

upvoted 2 times

jkhan2405 3 years ago

Selected Answer: A

It's A

upvoted 2 times

alexandercamachop 3 years ago

Selected Answer: A

A would seem like the correct one but not sure.

upvoted 2 times

A company has a workload in an AWS Region. Customers connect to and access the workload by using an Amazon API Gateway REST API. The company uses Amazon Route 53 as its DNS provider. The company wants to provide individual and secure URLs for all customers.

Which combination of steps will meet these requirements with the MOST operational efficiency? (Choose three.)

A. Register the required domain in a registrar. Create a wildcard custom domain name in a Route 53 hosted zone and record in the zone that points to the API Gateway endpoint. **Most Voted**

B. Request a wildcard certificate that matches the domains in AWS Certificate Manager (ACM) in a different Region.

C. Create hosted zones for each customer as required in Route 53. Create zone records that point to the API Gateway endpoint.

D. Request a wildcard certificate that matches the custom domain name in AWS Certificate Manager (ACM) in the same Region. **Most Voted**

E. Create multiple API endpoints for each customer in API Gateway.

F. Create a custom domain name in API Gateway for the REST API. Import the certificate from AWS Certificate Manager (ACM). **Most Voted**

Correct Answer: ADF

Community vote distribution

ADF (100%)

Comments

AshishRocks **Highly Voted** 2 years, 12 months ago

Step A involves registering the required domain in a registrar and creating a wildcard custom domain name in a Route 53 hosted zone. This allows you to map individual and secure URLs for all customers to your API Gateway endpoints.

Step D is to request a wildcard certificate from AWS Certificate Manager (ACM) that matches the custom domain name you created in Step A. This wildcard certificate will cover all subdomains and ensure secure HTTPS communication.

Step F is to create a custom domain name in API Gateway for your REST API. This allows you to associate the custom domain name with your API Gateway endpoints and import the certificate from ACM for secure communication.

upvoted 9 times

Guru4Cloud **Highly Voted** 2 years, 9 months ago

Selected Answer: ADF

The key points:

Using a wildcard domain and certificate avoids managing individual domains/certs per customer. This is more efficient.

The domain, hosted zone, and certificate should all be in the same region as the API Gateway REST API for simplicity.

Creating multiple API endpoints per customer (Option E) adds complexity and is not required.

Option B and C add unnecessary complexity by separating domains, certificates, and hosted zones.

upvoted 8 times

awsgeek75 **Most Recent** 2 years, 4 months ago

ADF looks right but not sure why C is wrong:

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-to-api-gateway.html#routing-to-api-gateway-config>

upvoted 3 times

LeonSauveterre 1 year, 6 months ago

My guess on why C is not ideal: Creating individual hosted zones for each customer is too much. It definitely increases operational overhead, while using a wildcard domain eliminates this need (which is more ideal).

upvoted 1 times

ukivanlamipi 2 years, 10 months ago

Selected Answer: ADF

<https://docs.aws.amazon.com/apigateway/latest/developerguide/how-to-custom-domains.html>
<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/AboutHZWorkingWith.html>

upvoted 5 times

jaydesai8 2 years, 11 months ago

Selected Answer: ADF

ADF - makes sense

upvoted 3 times

jkhan2405 3 years ago

Selected Answer: ADF

It's ADF

upvoted 3 times

MAMADOUG 3 years ago

For me AFD

upvoted 3 times

alexandercamachop 3 years ago

Selected Answer: ADF

ADF - One to create the custom domain in Route 53 (Amazon DNS)
Second to request wildcard certificate from ADM
Thirds to import the certificate from ACM.

upvoted 4 times

AncaZalog 3 years ago

is ADF

upvoted 2 times

A company stores data in Amazon S3. According to regulations, the data must not contain personally identifiable information (PII). The company recently discovered that S3 buckets have some objects that contain PII. The company needs to automatically detect PII in S3 buckets and to notify the company's security team.

Which solution will meet these requirements?

- A. Use Amazon Macie. Create an Amazon EventBridge rule to filter the SensitiveData event type from Macie findings and to send an Amazon Simple Notification Service (Amazon SNS) notification to the security team. **Most Voted**
- B. Use Amazon GuardDuty. Create an Amazon EventBridge rule to filter the CRITICAL event type from GuardDuty findings and to send an Amazon Simple Notification Service (Amazon SNS) notification to the security team.
- C. Use Amazon Macie. Create an Amazon EventBridge rule to filter the SensitiveData:S3Object/Personal event type from Macie findings and to send an Amazon Simple Queue Service (Amazon SQS) notification to the security team.
- D. Use Amazon GuardDuty. Create an Amazon EventBridge rule to filter the CRITICAL event type from GuardDuty findings and to send an Amazon Simple Queue Service (Amazon SQS) notification to the security team.

Correct Answer: A

Community vote distribution

A (87%)

13%

Comments

alexandercamachop **Highly Voted** 3 years ago

Selected Answer: A

B and D are discarded as Macie is to identify PII.

Now that we have between A and C.

SNS is more suitable for this option as a pub/sub service, we subscribe the security team and then they will receive the notifications.

upvoted 15 times

LeonSauveterre **Most Recent** 1 year, 6 months ago

Selected Answer: A

Macie + SNS = Identify PII + Notification

upvoted 1 times

awsgeek75 2 years, 5 months ago

Selected Answer: A

BD: Wrong products

AC: Uses Macie which is the right product but C uses SQS to notify security team which is an incomplete solution (what's listening to SQS?)

upvoted 2 times

pentium75 2 years, 5 months ago

Selected Answer: A

Detect PII -> Macie, A or C

Notify security team -> SNS, A or B

upvoted 4 times

potomac 2 years, 7 months ago

Selected Answer: A

C is SQS, not SNS

upvoted 4 times

Wayne23Fang 2 years, 9 months ago

SQS mentioned in C.

upvoted 2 times

Ale1973 2 years, 10 months ago

Selected Answer: A

Amazon SQS is typically used for decoupling and managing messages between distributed application components. It's not typically used for sending notifications directly to humans. On my opinion C isn't a best practice

upvoted 2 times

Kp88 2 years, 10 months ago

Those who say C , please read carefully (I made the same mistake lol). Teams can't be notified with SQS hence A.

upvoted 3 times

ukivanlamipi 2 years, 10 months ago

Selected Answer: C

there are different type of sensitive data: <https://docs.aws.amazon.com/macie/latest/user/findings-types.html>. if the question only focus on PII, then C is the answer. however, in reality, you will use A, because you will not want bank card, credential...etc all sensitive data , not only PII

upvoted 3 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: A

Automatically detect PII in S3 buckets = Amazon Macie

Notify security team = Amazon SNS

Trigger notification based on SensitiveData event type from Macie findings = EventBridge

upvoted 2 times

NASHDBA 2 years, 11 months ago

Selected Answer: C

There are different types of Sensitive Data. Here we are only referring to PII. Hence SensitiveData:S3Object/Personal. to use SNS, the security team must subscribe. SQS sends the information as designed

upvoted 1 times

narddrer 2 years, 11 months ago

Selected Answer: C

SensitiveData:S3Object/Personal

upvoted 1 times

jaydesai8 2 years, 11 months ago

Selected Answer: A

Sensitive = MACIE, and SNS to sent notification to the Security Team

upvoted 3 times

Iragmt 2 years, 11 months ago

C. Because the question mentioned PII only, there are other Sensitive Data aside from PII.

reference: <https://docs.aws.amazon.com/macie/latest/user/findings-publish-event-schemas.html> look for Event example for a sensitive data finding

upvoted 2 times

Ale1973 2 years, 10 months ago

But Amazon SQS is typically used for decoupling and managing messages between distributed application components. It's not typically used for sending notifications directly to humans!

upvoted 3 times

kapit 2 years, 11 months ago

AAAAAAA

upvoted 2 times

jack79 2 years, 12 months ago

C <https://docs.aws.amazon.com/macie/latest/user/findings-types.html>

and notice the ensitiveData:S3Object/Personal

The object contains personally identifiable information (such as mailing addresses or driver's license identification numbers), personal health information (such as health insurance or medical identification numbers), or a combination of the two.

upvoted 3 times

Ale1973 2 years, 10 months ago

But Amazon SQS is typically used for decoupling and managing messages between distributed application components. It's not typically used for sending notifications directly to humans!

upvoted 2 times

MAMADOUG 3 years ago

I vote for A, Sensitive = MACIE, and SNS to prevent Security Team

upvoted 4 times


A company wants to build a logging solution for its multiple AWS accounts. The company currently stores the logs from all accounts in a centralized account. The company has created an Amazon S3 bucket in the centralized account to store the VPC flow logs and AWS CloudTrail logs. All logs must be highly available for 30 days for frequent analysis, retained for an additional 60 days for backup purposes, and deleted 90 days after creation.

Which solution will meet these requirements MOST cost-effectively?

- A. Transition objects to the S3 Standard storage class 30 days after creation. Write an expiration action that directs Amazon S3 to delete objects after 90 days.
- B. Transition objects to the S3 Standard-Infrequent Access (S3 Standard-IA) storage class 30 days after creation. Move all objects to the S3 Glacier Flexible Retrieval storage class after 90 days. Write an expiration action that directs Amazon S3 to delete objects after 90 days.
- C. Transition objects to the S3 Glacier Flexible Retrieval storage class 30 days after creation. Write an expiration action that directs Amazon S3 to delete objects after 90 days. **Most Voted**
- D. Transition objects to the S3 One Zone-Infrequent Access (S3 One Zone-IA) storage class 30 days after creation. Move all objects to the S3 Glacier Flexible Retrieval storage class after 90 days. Write an expiration action that directs Amazon S3 to delete objects after 90 days.

Correct Answer: C

Community vote distribution



Comments

alexandercamachop **Highly Voted** 3 years ago

Selected Answer: C

C seems the most suitable.

Is the lowest cost.

After 30 days is backup only, doesn't specify frequent access.

Therefore we must transition the items after 30 days to Glacier Flexible Retrieval.

Also it says deletion after 90 days, so all answers specifying a transition after 90 days makes no sense.

upvoted 18 times

MAMADOUG 3 years ago

Agree with you

upvoted 3 times

deechan **Highly Voted** 2 years, 9 months ago

Selected Answer: A

The Glacier min storage duration is 90 days. All the options using Glacier are wrong. Only A is feasible.

upvoted 10 times

daniel33 2 years, 8 months ago

S3 Standard is priced at \$0.023 per GB for the first 50 TB stored per month

S3 Glacier Flexible Retrieval costs \$0.0036 per GB stored per month

If you move or delete data in Glacier within 90-days since their creation, you will pay an additional charge, that is called an early deletion fee. In US East you will pay \$0.004/GB if you have deleted 1 GB in 2 months, \$0.008/GB if you have deleted 1 GB in 1 month and \$0.012 if you have deleted 1 GB within 3 months.

Even with the early deletion fee, it appears to me that answer 'A' would still be cheaper.

upvoted 3 times

awsgeek75 2 years, 4 months ago

But the objects are deleted after 90 days so how is this charge applicable?

upvoted 2 times

pentium75 2 years, 5 months ago

But why 'transition to the S3 Standard storage class', aren't they there already by default?

upvoted 4 times

LeonSauveterre Most Recent 1 year, 6 months ago

Selected Answer: C

A: When logs are created, they are initially automatically created to S3 Standard. What are you transitioning anyway?

B & D: You could've just deleted logs after 90 days. No need to move them somewhere else then delete them. That's just money down the drain.

upvoted 1 times

awsgeek75 2 years, 5 months ago

Selected Answer: C

C: Lowest cost

upvoted 2 times

awsgeek75 2 years, 5 months ago

A: Standard storage is default so this is wrong.

B: Looks wrong because it moves object to S3GFR after 90 days when they could just be deleted so extra cost

D: Same problem as B

upvoted 2 times

pentium75 2 years, 5 months ago

Selected Answer: C

Not A: Objects are created in S3 Standard, so it doesn't make sense to 'transition' them there "30 days after creation"

Not B or C: No need to "move all objects to the S3 Glacier Flexible Retrieval storage class after 90 days" because we want to delete, not archive, them. Even if we would delete them right after moving, we would pay 90 days minimum storage duration. Plus, we are using "Infrequent Access" classes here, but we have no access at all.

upvoted 3 times

Rhydian25 1 year, 11 months ago

I guess you wanted to write "Not B or D"

upvoted 3 times

ftaws 2 years, 5 months ago

Selected Answer: A

requirement : frequently analysis

search cost : S3 STD 0.0004 vs IA 0.001

so IA is more expensive than STD(A)

upvoted 1 times

EdenWang 2 years, 6 months ago

Selected Answer: C

C is most cost-effective

upvoted 3 times

Hades2231 2 years, 9 months ago

Selected Answer: C

Things to note are: 30 days frequent access and 90 days after creation, so you only need to do 2 things, not 3. Objects in S3 will be stored by default for 30 days before you can move it to somewhere else, so C is the answer.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html>

upvoted 5 times

rjbihari 2 years, 9 months ago

C is the correct one .

As after 30 days it doesn't say about access / retrieval , only backup so move items after 30 days to Glacier Flexible Retrieval.

And after it says deletion , so expiration action will ensure that the objects are deleted after 90 days, even if they are not accessed

upvoted 2 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: B

I think - it is B

The first 30 days, the logs need to be highly available for frequent analysis. The S3 Standard storage class is the most expensive storage class, but it also provides the highest availability.

After 30 days, the logs still need to be retained for backup purposes, but they do not need to be accessed frequently. The S3 Standard-IA storage

class is a good option for this, as it is less expensive than the S3 Standard storage class.
After 90 days, the logs can be moved to the S3 Glacier Flexible Retrieval storage class. This is the most cost-effective storage class for long-term archiving.
The expiration action will ensure that the objects are deleted after 90 days, even if they are not accessed
upvoted 2 times

pentium75 2 years, 5 months ago

"After 90 days, the logs can be moved to the S3 Glacier Flexible Retrieval storage class. This is the most cost-effective storage class for long-term archiving." yeah but we don't need long-term archiving, we want to delete them after 90 days.
upvoted 3 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: C

C is the most cost effective solution.
upvoted 2 times

antropaws 2 years, 11 months ago

Selected Answer: C

C most likely.
upvoted 2 times

y0eri 2 years, 12 months ago

Selected Answer: A

Question says "All logs must be highly available for 30 days for frequent analysis" I think the answer is A. Glacier is not made for frequent access.
upvoted 2 times

y0eri 2 years, 12 months ago

I take that back. Moderator, please delete my comment.
upvoted 4 times

KMohsoe 2 years, 12 months ago

Selected Answer: B

I think B
upvoted 1 times


A company is building an Amazon Elastic Kubernetes Service (Amazon EKS) cluster for its workloads. All secrets that are stored in Amazon EKS must be encrypted in the Kubernetes etcd key-value store.

Which solution will meet these requirements?

- A. Create a new AWS Key Management Service (AWS KMS) key. Use AWS Secrets Manager to manage, rotate, and store all secrets in Amazon EKS.
- B. Create a new AWS Key Management Service (AWS KMS) key. Enable Amazon EKS KMS secrets encryption on the Amazon EKS cluster. **Most Voted**
- C. Create the Amazon EKS cluster with default options. Use the Amazon Elastic Block Store (Amazon EBS) Container Storage Interface (CSI) driver as an add-on.
- D. Create a new AWS Key Management Service (AWS KMS) key with the alias/aws/ebs alias. Enable default Amazon Elastic Block Store (Amazon EBS) volume encryption for the account.

Correct Answer: B

Community vote distribution



Comments

Guru4Cloud **Highly Voted** 1 year, 9 months ago

Selected Answer: B

B is the correct solution to meet the requirement of encrypting secrets in the etcd store for an Amazon EKS cluster.

The key points:

Create a new KMS key to use for encryption.

Enable EKS secrets encryption using that KMS key on the EKS cluster. This will encrypt secrets in the Kubernetes etcd store.

Option A uses Secrets Manager which does not encrypt the etcd store.

Option C uses EBS CSI which is unrelated to etcd encryption.

Option D enables EBS encryption but does not address etcd encryption.

upvoted 7 times

TariqKipkemei **Highly Voted** 1 year, 10 months ago

Selected Answer: B

EKS supports using AWS KMS keys to provide envelope encryption of Kubernetes secrets stored in EKS. Envelope encryption adds an additional, customer-managed layer of encryption for application secrets or user data that is stored within a Kubernetes cluster.

<https://eksctl.io/usage/kms-encryption/>

upvoted 6 times

FlyingHawk **Most Recent** 1 year, 4 months ago

Selected Answer: B

The problem is to ensure all secrets in Amazon EKS be encrypted, so the correct solution is to enable the secret encryption on the EKS cluster with a new AWS Key:

```
eksctl utils enable-secrets-encryption \
```

```
--cluster my-cluster \
```

```
--key-arn arn:aws:kms:region-code:account:key/key
```

https://docs.aws.amazon.com/en_en/eks/latest/userguide/enable-kms.html

upvoted 1 times

LeonSauveterre 1 year, 6 months ago

Selected Answer: B

"must be encrypted in the Kubernetes etcd key-value store" - that leaves B only.

upvoted 1 times

manuh 1 year, 11 months ago

Selected Answer: A

Why not a

upvoted 1 times

TariqKipkemei 1 year, 10 months ago

option A does not enable Amazon EKS KMS secrets encryption on the Amazon EKS cluster

upvoted 2 times

MrAWSAssociate 1 year, 11 months ago

Selected Answer: B

B is the right option.

<https://docs.aws.amazon.com/eks/latest/userguide/enable-kms.html>

upvoted 5 times

alexandercamachop 2 years ago

Selected Answer: B

It is B, because we need to encrypt inside of the EKS cluster, not outside.

AWS KMS is to encrypt at rest.

upvoted 5 times

AncaZalog 2 years ago

is B, not D

upvoted 2 times


A company wants to provide data scientists with near real-time read-only access to the company's production Amazon RDS for PostgreSQL database. The database is currently configured as a Single-AZ database. The data scientists use complex queries that will not affect the production database. The company needs a solution that is highly available.

Which solution will meet these requirements MOST cost-effectively?

- A. Scale the existing production database in a maintenance window to provide enough power for the data scientists.
- B. Change the setup from a Single-AZ to a Multi-AZ instance deployment with a larger secondary standby instance. Provide the data scientists access to the secondary instance.
- C. Change the setup from a Single-AZ to a Multi-AZ instance deployment. Provide two additional read replicas for the data scientists.
- D. Change the setup from a Single-AZ to a Multi-AZ cluster deployment with two readable standby instances. Provide read endpoints to the data scientists. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

NASHDBA **Highly Voted** 2 years, 11 months ago

Selected Answer: D

Highly Available = Multi-AZ Cluster

Read-only + Near Real time = readable standby.

Read replicas are async whereas readable standby is synchronous.

<https://stackoverflow.com/questions/70663036/differences-b-w-aws-read-replica-and-the-standby-instances>

upvoted 23 times

chickenmf 2 years, 3 months ago

a Multi-AZ instance deployment is also highly available for a lower cost, no?

upvoted 2 times

Smart 2 years, 9 months ago

This^ is the reason.

upvoted 3 times

maver144 **Highly Voted** 2 years, 12 months ago

It's either C or D. To be honest, I find the newest questions to be ridiculously hard (roughly 500+). I agree with @alexandercamachop that Multi Az in Instance mode is cheaper than Cluster. However, with Cluster we have reader endpoint available to use out-of-box, so there is no need to provide read-replicas, which also has its own costs. The ridiculous part is that I'm pretty sure even the AWS support would have troubles to answer which configuration is MOST cost-effective.

upvoted 14 times

manuh 2 years, 11 months ago

Absolutely true that 500+ questions are damn difficult to answer. I still dont know why is B incorrect. Shouldn't 1 extra be better than 2 ?

upvoted 1 times

cyber_bedouin 2 years, 6 months ago

they are not all hard, most are normal. its just this one and that one about EKS encryption control plane (earlier than this page).

upvoted 3 times

maver144 2 years, 12 months ago

Near real-time is clue for C, since read replicas are async, but still its not obvious question.

upvoted 3 times

Onimole Most Recent 5 months, 3 weeks ago

Selected Answer: C

A lot of people skipped the part where it says "The data scientists use complex queries that will not affect the production database"
upvoted 1 times

AdamVigas 1 year, 3 months ago

Selected Answer: C

two readable standby instances...if they are standby they are not supposed to read...am i getting something wrong ? i would go for C
upvoted 3 times

emakid 1 year, 11 months ago

Selected Answer: C

Option D: Multi-AZ cluster deployment with two readable standby instances would be more costly and is not necessary if read replicas are sufficient for the data scientists' needs.

Thus, Option C is the most cost-effective and operationally efficient solution to meet the company's requirements.
upvoted 3 times

osmk 2 years, 3 months ago

Selected Answer: D

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/multi-az-db-clusters-concepts.html>
upvoted 2 times

pentium75 2 years, 5 months ago

Selected Answer: D

Not A - Not "highly available"
Not B - "Access to the secondary instance" is not possible in Multi-AZ
Not C - Multi-AZ + two (!) read replicas is more expensive than cluster
D - Provides "readable standby instances"
upvoted 3 times

SHAAHIBHUSHANAWS 2 years, 6 months ago

D
<https://aws.amazon.com/about-aws/whats-new/2023/01/amazon-rds-multi-az-readable-standbys-rds-postgresql-inbound-replication/>
upvoted 3 times

bogobob 2 years, 6 months ago

Selected Answer: D

<https://aws.amazon.com/blogs/database/choose-the-right-amazon-rds-deployment-option-single-az-instance-multi-az-instance-or-multi-az-database-cluster/>
C would mean you are paying for 4 instances (primary, backup, and 2 read instances). D would be 3 (primary, and 2 backup). Difficult to be sure, pricing calculator doesn't even include clusters yet.
upvoted 2 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: D

Option D is the most cost-effective solution that meets the requirements for this scenario.

The key considerations are:

Data scientists need read-only access to near real-time production data without affecting performance.
High availability is required.
Cost should be minimized.
upvoted 2 times

ukivanlamlpi 2 years, 10 months ago

Selected Answer: D

<https://aws.amazon.com/blogs/database/choose-the-right-amazon-rds-deployment-option-single-az-instance-multi-az-instance-or-multi-az-database-cluster/>

only multi AZ cluster have reader endpoint. multi AZ instance secondary replicate is not allow to access
upvoted 2 times

msdnpro 2 years, 10 months ago

Selected Answer: D

Support for D:

Amazon RDS now offers Multi-AZ deployments with readable standby instances (also called Multi-AZ DB cluster deployments) in preview. You should consider using Multi-AZ DB cluster deployments with two readable DB instances if you need additional read capacity in your Amazon RDS Multi-AZ deployment and if your application workload has strict transaction latency requirements such as single-digit milliseconds transactions.

<https://aws.amazon.com/blogs/database/readable-standby-instances-in-amazon-rds-multi-az-deployments-a-new-high-availability-option/>
upvoted 2 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: D

Unlike Multi-AZ instance deployment, where the secondary instance can't be accessed for read or writes, Multi-AZ DB cluster deployment consists of primary instance running in one AZ serving read-write traffic and two other standby running in two different AZs serving read traffic.

upvoted 4 times

lragmt 2 years, 11 months ago

Selected Answer: D

D. using Multi-AZ DB cluster deployments with two readable DB instances if you need additional read capacity in your Amazon RDS Multi-AZ deployment and if your application workload has strict transaction latency requirements such as single-digit milliseconds transactions.

<https://aws.amazon.com/blogs/database/readable-standby-instances-in-amazon-rds-multi-az-deployments-a-new-high-availability-option/>

while on read replicas, Amazon RDS then uses the asynchronous replication method for the DB engine to update the read replica whenever there is a change to the primary DB instance. https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_ReadRepl.html

upvoted 2 times

manuh 2 years, 11 months ago

Selected Answer: B

Why not b. Shouldn't it have less number of instances than both c and d?

upvoted 2 times

baba365 2 years, 11 months ago

Complex queries on single db will affect performance of db

upvoted 2 times

pentium75 2 years, 5 months ago

You can't 'access the secondary instance' as suggested by B

upvoted 2 times

pentium75 2 years, 5 months ago

"The Multi-AZ instance is suitable for business/mission critical applications that require high availability with low RTO/RPO and resilience to availability zone outage. However, this high availability option isn't a scaling solution for read-only scenarios. You can't use a standby replica to serve read traffic. To serve read-only traffic, use a Multi-AZ DB cluster or a read replica instead."

upvoted 2 times

baba365 2 years, 11 months ago

Multi-AZ is about twice the price of Single-AZ. For example:

db.t2.micro single - \$0.017/hour

db.t2.micro multi - \$0.034/hour

option C: 1 primary + 1 standby + 2 replica = 4Db

option D: 1 primary + 2 standby = 3Db

D. appears to be most cost effective

upvoted 3 times

wsdadsdqwdaw 2 years, 7 months ago

I think the best explanation I've read so far.

upvoted 2 times

0628atv 2 years, 11 months ago

D:

<https://aws.amazon.com/tw/blogs/database/readable-standby-instances-in-amazon-rds-multi-az-deployments-a-new-high-availability-option/>

upvoted 2 times

vrevkov 2 years, 11 months ago

Selected Answer: D

Forgot to vote

upvoted 3 times

A company runs a three-tier web application in the AWS Cloud that operates across three Availability Zones. The application architecture has an Application Load Balancer, an Amazon EC2 web server that hosts user session states, and a MySQL database that runs on an EC2 instance. The company expects sudden increases in application traffic. The company wants to be able to scale to meet future application capacity demands and to ensure high availability across all three Availability Zones.

Which solution will meet these requirements?

A. Migrate the MySQL database to Amazon RDS for MySQL with a Multi-AZ DB cluster deployment. Use Amazon ElastiCache for Redis with high availability to store session data and to cache reads. Migrate the web server to an Auto Scaling group that is in three Availability Zones. **Most Voted**

B. Migrate the MySQL database to Amazon RDS for MySQL with a Multi-AZ DB cluster deployment. Use Amazon ElastiCache for Memcached with high availability to store session data and to cache reads. Migrate the web server to an Auto Scaling group that is in three Availability Zones.

C. Migrate the MySQL database to Amazon DynamoDB. Use DynamoDB Accelerator (DAX) to cache reads. Store the session data in DynamoDB. Migrate the web server to an Auto Scaling group that is in three Availability Zones.

D. Migrate the MySQL database to Amazon RDS for MySQL in a single Availability Zone. Use Amazon ElastiCache for Redis with high availability to store session data and to cache reads. Migrate the web server to an Auto Scaling group that is in three Availability Zones.

Correct Answer: A

Community vote distribution

A (81%)

B (19%)

Comments

alexandercamachop **Highly Voted** 2 years, 6 months ago

Selected Answer: A

Memcached is best suited for caching data, while Redis is better for storing data that needs to be persisted. If you need to store data that needs to be accessed frequently, such as user profiles, session data, and application settings, then Redis is the better choice

upvoted 17 times

nonameforyou 2 years, 5 months ago

and for high availability, it's better than memcached

upvoted 3 times

nonameforyou 2 years, 5 months ago

but does rds multi-az provide the needed scalability?

upvoted 3 times

wsdasdasdqwda 2 years, 1 month ago

it is multi-az cluster deployment, same as B, so yes, it is providing the needed scalability. Great explanation.

upvoted 2 times

AncaZalog **Highly Voted** 2 years, 6 months ago

is A not B

upvoted 5 times

zdi561 **Most Recent** 1 year, 4 months ago

Selected Answer: A

AI answer, For managing user sessions on AWS ElastiCache, Redis is generally the better choice compared to Memcached because Redis offers more advanced data structures, built-in persistence features, and better support for complex session data, while Memcached is simpler but only suitable for basic key-value caching with limited data types

upvoted 1 times

osmk 1 year, 9 months ago

Selected Answer: A

Replication: Redis supports creating multiple replicas for read scalability and high availability.<https://aws.amazon.com/elasticache/redis-vs-memcached/>

upvoted 2 times

awsgeek75 1 year, 10 months ago

Selected Answer: A

A because of "Amazon EC2 web server that hosts user session states"

C: RDS to DynamoDB doesn't make total sense

D: Single zone is not HA

Between A and B, A is suitable because of session state and Elasticache with Redis is more HA than option B

upvoted 2 times

mr123dd 1 year, 11 months ago

Selected Answer: A

B: from what I know, Memcached provide better performance and simplicity but lower availability than redis.

C: mysql is relational database, dynamodb is nosql

D: single AZ

upvoted 2 times

pentium75 1 year, 11 months ago

Selected Answer: A

ElastiCache for Redis supports HA, ElastiCache for Memcached does not:

<https://docs.aws.amazon.com/AmazonElastiCache/latest/mem-ug/SelectEngine.html>

C could in theory work, but session data is typically stored in ElastiCache, not in DynamoDB.

D is not HA.

upvoted 3 times

Cyberkayu 1 year, 11 months ago

Selected Answer: B

'hosts user session states' in question, thus redis

upvoted 2 times

pentium75 1 year, 11 months ago

Right, but Redis is A

upvoted 2 times

potomac 2 years, 1 month ago

Selected Answer: A

Redis is a widely adopted in-memory data store for use as a database, cache, message broker, queue, session store, and leaderboard.

<https://aws.amazon.com/elasticache/redis/>

upvoted 5 times

thanhnv142 2 years, 1 month ago

B is correct.

We are left with 2 options: A and B. But it requires that the system be able to scale to meet future application capacity demands. Redis is very good.

But its drawback is not scalable. That's why they implement memcached.

upvoted 1 times

ErnShm 2 years, 3 months ago

A

Redis as an in-memory data store with high availability and persistence is a popular choice among application developers to store and manage session data for internet-scale applications. Redis provides the sub-millisecond latency, scale, and resiliency required to manage session data such as user profiles, credentials, session state, and user-specific personalization.

upvoted 2 times

Gajendr 1 year, 11 months ago

Redis provides replication while memcached doesn't.

upvoted 2 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: A

The key reasons why option A is preferable:

RDS Multi-AZ provides high availability for MySQL by synchronously replicating data across AZs. Automatic failover handles AZ outages. ElastiCache for Redis is better suited for session data caching than Memcached. Redis offers more advanced data structures and flexibility. Auto scaling across 3 AZs provides high availability for the web tier

upvoted 2 times

ukivanlamipi 2 years, 4 months ago

Selected Answer: B

the different between Redis and Memcache is that Memcache support multithread process to handle the increase of application traffic.
<https://aws.amazon.com/elasticache/redis-vs-memcached/>

upvoted 2 times

pentium75 1 year, 11 months ago

ElastiCache for Memcached says "No" for "High Availability"

<https://docs.aws.amazon.com/AmazonElastiCache/latest/mem-ug/SelectEngine.html>

upvoted 2 times

TariqKipkemei 2 years, 4 months ago

Selected Answer: B

This requirement wins for me: "be able to scale to meet future application capacity demands".
Memcached implements a multi-threaded architecture, it can make use of multiple processing cores. This means that you can handle more operations by scaling up compute capacity.

<https://aws.amazon.com/elasticache/redis-vs-memcached/#:~:text=by%20their%20rank.,Multithreaded%20architecture,-Since%20Memcached%20is>

upvoted 1 times

plndmns 2 years, 5 months ago

cache reads is memcached right?

upvoted 1 times

MrAWSAssociate 2 years, 5 months ago

Selected Answer: B

B is correct!

upvoted 3 times

A global video streaming company uses Amazon CloudFront as a content distribution network (CDN). The company wants to roll out content in a phased manner across multiple countries. The company needs to ensure that viewers who are outside the countries to which the company rolls out content are not able to view the content.

Which solution will meet these requirements?

- A. Add geographic restrictions to the content in CloudFront by using an allow list. Set up a custom error message. **Most Voted**
- B. Set up a new URL for restricted content. Authorize access by using a signed URL and cookies. Set up a custom error message.
- C. Encrypt the data for the content that the company distributes. Set up a custom error message.
- D. Create a new URL for restricted content. Set up a time-restricted access policy for signed URLs.

Correct Answer: A

Community vote distribution

A (100%)

Comments

alexandercamachop **Highly Voted** 2 years, 6 months ago

Selected Answer: A

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/georestrictions.html>
upvoted 5 times

awsgeek75 **Highly Voted** 1 year, 11 months ago

Selected Answer: A

This question asks us to guess Netflix subscription model in 2 mins! lol!

BCD are impractical for geo restrictions as you cannot restrict URL by region and you cannot encrypt by geo region (country etc)
upvoted 5 times

potomac **Most Recent** 2 years, 1 month ago

Selected Answer: A

The CloudFront geographic restrictions feature lets you control distribution of your content at the country level for all files that you're distributing with a given web distribution.
<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/georestrictions.html>
upvoted 5 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: A

Add geographic restrictions to the content in CloudFront by using an allow list. Set up a custom error message
upvoted 3 times

TariqKipkemei 2 years, 4 months ago

Selected Answer: A

Add geographic restrictions to the content in CloudFront by using an allow list. Set up a custom error message.
upvoted 2 times

jaydesai8 2 years, 5 months ago

Selected Answer: A

A makes sense - cloudfront has the capabilities of georestriction
upvoted 2 times

antropaws 2 years, 5 months ago

Selected Answer: A

Pretty sure it's A.

upvoted 2 times

AncaZalog 2 years, 6 months ago

is B not A

upvoted 1 times

awsgeek75 1 year, 10 months ago

How is signed URL going to be geo restricted? Anyone with signed url can access the content on that url regardless of their location so B is wrong.

upvoted 2 times

antropaws 2 years, 5 months ago

Why's that?

upvoted 2 times

manuh 2 years, 5 months ago

Signed url or cookies can be used for the banner country as well?

upvoted 1 times

A company wants to use the AWS Cloud to improve its on-premises disaster recovery (DR) configuration. The company's core production business application uses Microsoft SQL Server Standard, which runs on a virtual machine (VM). The application has a recovery point objective (RPO) of 30 seconds or fewer and a recovery time objective (RTO) of 60 minutes. The DR solution needs to minimize costs wherever possible.

Which solution will meet these requirements?

- A. Configure a multi-site active/active setup between the on-premises server and AWS by using Microsoft SQL Server Enterprise with Always On availability groups.
- B. Configure a warm standby Amazon RDS for SQL Server database on AWS. Configure AWS Database Migration Service (AWS DMS) to use change data capture (CDC). **Most Voted**
- C. Use AWS Elastic Disaster Recovery configured to replicate disk changes to AWS as a pilot light.
- D. Use third-party backup software to capture backups every night. Store a secondary set of backups in Amazon S3.

Correct Answer: B

Community vote distribution

B (57%)

C (43%)

Comments

1Alpha1 **Highly Voted** 2 years, 4 months ago

Selected Answer: B

Backup & Restore (RPO in hours, RTO in 24 hours or less)

Pilot Light (RPO in minutes, RTO in hours)

Warm Standby (RPO in seconds, RTO in minutes) *** Right Answer ***

Active-Active (RPO is none or possibly seconds, RTO in seconds)

upvoted 10 times

1Alpha1 2 years, 4 months ago

[https://disaster-recovery.workshop.aws/en/intro/disaster-recovery.html#:~:text=Pilot%20Light%20\(RPO%20in%20minutes,that%20includes%20that%20critical%20core.](https://disaster-recovery.workshop.aws/en/intro/disaster-recovery.html#:~:text=Pilot%20Light%20(RPO%20in%20minutes,that%20includes%20that%20critical%20core.)

upvoted 4 times

pentium75 **Highly Voted** 2 years, 5 months ago

Selected Answer: C

Not A - too expensive and not using AWS services

Not B - "RDS for SQL Server" does not support everything that "SQL Server Standard which runs on a VM" does; CDC supports even less (https://docs.aws.amazon.com/dms/latest/userguide/CHAP_Source.SQLServer.html). Also it would be more expensive than C.

Not D - "Every night" would not meet the RPO requirement

upvoted 7 times

awsgeek75 2 years, 5 months ago

Thanks I was confused between B and C. This makes perfect sense!

upvoted 1 times

KarlosRC **Most Recent** 3 months ago

Selected Answer: C

Not B: DMS is not designed for DR, failover is not automatic, minimizing costs means not maintaining a fully functioning DR cluster

upvoted 1 times

In2718 10 months, 3 weeks ago

Selected Answer: C

AI follow-up - Yes, a Pilot Light strategy on AWS can achieve an RPO (Recovery Point Objective) of 30 seconds or even lower in specific scenarios, but it typically requires using specific AWS services that specialize in continuous, low-latency data replication.

upvoted 1 times

In2718 10 months, 3 weeks ago

Selected Answer: C

The request for lower costs leans towards C rather than B, assuming the pilot light meets the SLAs as defined.

upvoted 1 times

zdi561 1 year, 4 months ago

Selected Answer: C

AWS Elastic Disaster Recovery (AWS DRS):

A service that can replicate entire virtual machines, including your SQL Server instance, to a secondary region for disaster recovery

upvoted 1 times

FlyingHawk 1 year, 4 months ago

Selected Answer: B

AWS DMS with CDC can replicate changes from the on-premises SQL Server to the RDS instance in near real-time, meeting the 30-second RPO requirement.

A warm standby RDS instance ensures a faster recovery, meeting the 60-minute RTO requirement.

This solution is more cost-effective than upgrading to SQL Server Enterprise and provides a managed, scalable database service.

upvoted 1 times

FlyingHawk 1 year, 4 months ago

C - Not optimized for real-time replication, making it challenging to meet the 30-second RPO requirement.

upvoted 1 times

Saliilgen 1 year, 5 months ago

Selected Answer: B

IMO answer is B.

AWS Elastic Disaster Recovery is cheaper and may meet the RTO/RPO requirements of the question, but option C states that it is configured as a pilot light (RTO/RPO 10s minute)

upvoted 2 times

LeonSauveterre 1 year, 6 months ago

Selected Answer: C

RPO: the maximum acceptable amount of *data loss* during an outage or disaster.

RTO: the maximum acceptable amount of *downtime* after a failure.

Why B (RDS with DMS) might be better:

1. RTO Guarantee: A warm standby database provides faster recovery since the database is preconfigured and ready to accept traffic.
2. Simplicity: For workloads compatible with RDS and DMS, this approach simplifies management.

Why C (Elastic DR) might be better:

1. Workload Compatibility: If the on-premises SQL Server workload uses features that RDS or DMS cannot support, Elastic DR ensures full replication by mirroring the entire VM or disk.
2. Cost: Elastic DR replicates data but does not run a live database, making it much cheaper than a warm standby RDS.

"The DR solution needs to minimize costs wherever possible" - That makes me think C is the answer.

upvoted 2 times

FlyingHawk 1 year, 4 months ago

Elastic DRS is better suited for applications with less stringent RPO requirements. Achieving a 30-second RPO with this solution could be challenging, as it is not optimized for real-time replication.

upvoted 1 times

example_ 1 year, 10 months ago

Selected Answer: B

Pilot light (RPO in minutes, RTO in tens of minutes)

Warm standby (RPO in seconds, RTO in minutes)

https://docs.aws.amazon.com/wellarchitected/latest/reliability-pillar/rel_planning_for_recovery_disaster_recovery.html

upvoted 2 times

MrAliMohsan 1 year, 9 months ago

When cost is a concern, and you wish to achieve a similar RPO and RTO objectives as defined in the warm standby strategy, you could consider cloud native solutions, like AWS Elastic Disaster Recovery, that take the pilot light approach and offer improved RPO and RTO targets.

upvoted 3 times

abhiarns 2 years, 1 month ago

Selected Answer: C

AWS DRS(AWS Elastic Disaster Recovery) enables RPOs of seconds and RTOs of minutes.

upvoted 2 times

osmk 2 years, 3 months ago

Selected Answer: B

<https://docs.aws.amazon.com/whitepapers/latest/disaster-recovery-workloads-on-aws/disaster-recovery-options-in-the-cloud.html#warm-standby>

upvoted 3 times

awsgeek75 2 years, 4 months ago

Selected Answer: C

A: Not possible

B: With RDS it means your failover will launch a different database engine. This is wrong in general

D: No comments

C: It is a disk based replication so it will be similar DB server and this is the product managed by AWS for the DR of on-prem setups.

<https://aws.amazon.com/blogs/modernizing-with-aws/how-to-set-up-disaster-recovery-for-sql-server-always-on-availability-groups-using-aws-elastic-disaster-recovery/>

upvoted 2 times

1rob 2 years, 6 months ago

Selected Answer: C

AWS Elastic Disaster Recovery

If you are considering the pilot light or warm standby strategy for disaster recovery, AWS Elastic Disaster Recovery could provide an alternative approach with improved benefits. Elastic Disaster Recovery can offer an RPO and RTO target similar to warm standby, but maintain the low-cost approach of pilot light

From <https://docs.aws.amazon.com/wellarchitected/latest/reliability-pillar/rel_planning_for_recovery_disaster_recovery.html>

upvoted 3 times

potomac 2 years, 7 months ago

Selected Answer: B

With the pilot light approach, you replicate your data from one environment to another and provision a copy of your core workload infrastructure, not the fully functional copy of your production environment in a recovery environment.

upvoted 2 times

saymolet 2 years, 6 months ago

<https://aws.amazon.com/blogs/architecture/disaster-recovery-dr-architecture-on-aws-part-iii-pilot-light-and-warm-standby/>

upvoted 3 times

pentium75 2 years, 5 months ago

We have no idea if they are using SQL Server features that require OS customization etc., so we can't assume that the app would run on RDS for SQL Server at all. We need a replica of the VM that SQL Server is currently running on, thus C.

upvoted 2 times

thanhnv142 2 years, 7 months ago

C: Pilot light

- In pilot light, databases are always on, thus minimize RPO (can satisfy the 30s requirement)

- Only apps are turn off. But it can satisfy the 60 minutes requirement

- Warm standby, of course, can satisfy all the RPO and RTO requirements, but it is more expensive than pilot light

upvoted 4 times

richguo 2 years, 8 months ago

Selected Answer: C

B(warm standby) is doable, but C (pilot light) is most cost effectively.

<https://aws.amazon.com/tw/blogs/architecture/disaster-recovery-dr-architecture-on-aws-part-iii-pilot-light-and-warm-standby/>

upvoted 3 times

A company has an on-premises server that uses an Oracle database to process and store customer information. The company wants to use an AWS database service to achieve higher availability and to improve application performance. The company also wants to offload reporting from its primary database system.

Which solution will meet these requirements in the MOST operationally efficient way?

- A. Use AWS Database Migration Service (AWS DMS) to create an Amazon RDS DB instance in multiple AWS Regions. Point the reporting functions toward a separate DB instance from the primary DB instance.
- B. Use Amazon RDS in a Single-AZ deployment to create an Oracle database. Create a read replica in the same zone as the primary DB instance. Direct the reporting functions to the read replica.
- C. Use Amazon RDS deployed in a Multi-AZ cluster deployment to create an Oracle database. Direct the reporting functions to use the reader instance in the cluster deployment.
- D. Use Amazon RDS deployed in a Multi-AZ instance deployment to create an Amazon Aurora database. Direct the reporting functions to the reader instances. **Most Voted**

Correct Answer: D

Community vote distribution

D (72%)

C (28%)

Comments

mrsoa **Highly Voted** 2 years, 10 months ago

Selected Answer: D

Its D

Multi-AZ DB clusters aren't available with the following engines:

RDS for MariaDB

RDS for Oracle

RDS for SQL Server

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.RDS_Fea_Regions_DB-eng.Feature.MultiAZDBClusters.html

upvoted 37 times

alexandercamachop **Highly Voted** 3 years ago

Selected Answer: C

C. Use Amazon RDS deployed in a Multi-AZ cluster deployment to create an Oracle database. Direct the reporting functions to use the reader instance in the cluster deployment.

A and B discarded.

The answer is between C and D

D says use an Amazon RDS to build an Amazon Aurora, makes no sense.

C is the correct one, high availability in multi az deployment.

Also point the reporting to the reader replica.

upvoted 14 times

1rob 2 years, 6 months ago

Multi-AZ DB clusters aren't available with the following engines:

RDS for MariaDB

RDS for Oracle

RDS for SQL Server

upvoted 8 times

bogobob 2 years, 6 months ago

using RDS to build Aurora from an Oracle DB <https://aws.amazon.com/tutorials/break-free-from-legacy-databases/migrate-oracle-to-amazon-aurora/>

upvoted 3 times

KarlosRC Most Recent 3 months ago

Selected Answer: D

The correct answer is: RDS Oracle Multi-AZ instance + Read Replica. But that option isn't listed in the answers. Conclusion: the question is either incomplete or written by idiots.

upvoted 1 times

LeonSauveterre 1 year, 6 months ago

Selected Answer: D

In the MOST operationally efficient way, D is better, and A is too much. HOWEVER, Aurora is not Oracle, so this would require migrating from Oracle to Aurora (potentially requiring significant changes to the application, schema, and queries). The question did not say anything about not wanting to migrate.

B - 100% wrong. Single-AZ deployment does not meet the higher availability requirement.

C - 100% wrong. There's no way to create an Oracle db using RDS.

upvoted 1 times

chest_jd 1 year, 7 months ago

Unfortunately ALL ANSWERS ARE NOT CORRECT.

A. Use AWS Database Migration Service (AWS DMS) to create an Amazon RDS DB instance in multiple AWS Regions.
How can I create RDS using DMS ?

B. Single-Zone is not suitable for high available solution.

C. RDS cluster deployment are not available for DBs: RDS for Db2, RDS for MariaDB, RDS for Oracle, RDS for SQL Server
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.RDS_Fea_Regions_DB-eng.Feature.MultiAZDBClusters.html

D. Use Amazon RDS deployed in a Multi-AZ instance deployment to create an Amazon Aurora database.
Aurora DOES NOT support Oracle DB. Only MySQL and PostgreSQL

upvoted 2 times

Ravan 2 years, 3 months ago

Selected Answer: D

Multi-AZ (Availability Zone) deployments are not available for the following Amazon RDS database engines:

1. Amazon Aurora with MySQL compatibility
2. Amazon Aurora with PostgreSQL compatibility
3. Amazon RDS for SQL Server Express Edition
4. Amazon RDS for Oracle Standard Edition One
5. Amazon RDS for Oracle Standard Edition
6. Amazon RDS for Oracle SE2 (Standard Edition 2)

For these database engines, Amazon RDS provides high availability using other mechanisms specific to each engine, such as Read Replicas or different standby configurations. However, Multi-AZ deployments, which automatically provision and maintain a synchronous standby replica in a different Availability Zone for failover support, are not supported for these engines.

upvoted 4 times

noircesar25 2 years, 3 months ago

this link explains why the answer is C and confirms that rds for oracle supports multi-AZ
<https://aws.amazon.com/blogs/aws/multi-az-option-for-amazon-rds-oracle/>

upvoted 2 times

osmk 2 years, 3 months ago

Selected Answer: D

requiring high availability and performance.<https://aws.amazon.com/rds/aurora/>

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: D

Between C&D, D is correct as C is not possible:

<https://aws.amazon.com/blogs/aws/multi-az-option-for-amazon-rds-oracle/>

upvoted 3 times

pentium75 2 years, 5 months ago

Selected Answer: D

Not A - Creating multiple instances and keeping them in sync in DMS is surely not "operationally efficient"

Not B - "replica in the same zone" -> does not provide "higher availability"

Not C - "Multi-AZ cluster" does not support Oracle engine

Thus D. Question does not mention that the app would use Oracle-specific features; we're also not asked to minimize application changes. Ideal solution from AWS point of view is to move from Oracle to Aurora.

upvoted 2 times

XXXXXINN 1 year, 8 months ago

but Aurora doesnt support oracle db though

upvoted 1 times

aws94 2 years, 6 months ago

Selected Answer: C

i am sure just look here

<https://aws.amazon.com/ar/blogs/aws/amazon-rds-multi-az-db-cluster/>

upvoted 1 times

aws94 2 years, 6 months ago

sorry this is the right link:

<https://aws.amazon.com/ar/blogs/aws/multi-az-option-for-amazon-rds-oracle/>

upvoted 2 times

pentium75 2 years, 5 months ago

"Multi-AZ cluster (!)" does not support Oracle. Multi-AZ instance would.

upvoted 2 times

EEK2k 2 years, 7 months ago

Selected Answer: C

It should be C. Oracle DB is supported in RDS Multi-AZ with one standby for HA. <https://aws.amazon.com/rds/features/multi-az/>. Additionally, a reader instance/replica could be added to RDS Multi-AZ with one standby setup to offload the read requests. Aurora is only supported MySQL and Postgres compatible DB so "D" is out.

upvoted 2 times

1rob 2 years, 6 months ago

<https://aws.amazon.com/rds/features/multi-az/> gives: Amazon RDS Multi-AZ is available for RDS for PostgreSQL, RDS for MySQL, RDS for MariaDB, RDS for SQL Server, RDS for Oracle, and RDS for Db2. Amazon RDS Multi-AZ with two readable standbys is available for RDS for PostgreSQL and RDS for MySQL.

So no reader instance.

upvoted 1 times

potomac 2 years, 7 months ago

Selected Answer: D

Multi-AZ DB clusters are NOT available with the following engines:

RDS for MariaDB

RDS for Oracle

RDS for SQL Server

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.RDS_Fea_Regions_DB-eng.Feature.MultiAZDBClusters.html

upvoted 2 times

danielmakita 2 years, 7 months ago

It is C. Aurora database doesn't support Oracle.

upvoted 1 times

JA2018 1 year, 6 months ago

as pentium75 had pointed out, the stem does not mention that the app would use Oracle-specific features; we're also not asked to minimize application changes. Ideal solution from AWS point of view is to move from Oracle to Aurora.

upvoted 1 times

wsdasdasdqwdaw 2 years, 7 months ago

You can use Aurora instead of Oracle. There are tutorials how to migrate Oracle to Aurora. On top C is not supported. The is not Multi-AZ DB CLUSTER for Oracle.

upvoted 3 times

wsdasdasdqwdaw 2 years, 7 months ago

It is D

upvoted 1 times

thanhnv142 2 years, 7 months ago

None options seems valid. Not C because it is not supported. But not D as well. RDS is not Aurora. They are two separate services. Additionally, In multi AZ instance deployment, it only provides fault tolerance, not High avail.

upvoted 2 times

Nikki013 2 years, 9 months ago

Selected Answer: D

Multi-AZ Cluster does not support Oracle as engine:

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.RDS_Fea_Regions_DB-eng.Feature.MultiAZDBClusters.html

upvoted 2 times

Bennyboy789 2 years, 9 months ago

Selected Answer: D

D is my choice.

Multi-AZ DB cluster does not support Oracle DB.

upvoted 3 times

A company wants to build a web application on AWS. Client access requests to the website are not predictable and can be idle for a long time. Only customers who have paid a subscription fee can have the ability to sign in and use the web application.

Which combination of steps will meet these requirements MOST cost-effectively? (Choose three.)

A. Create an AWS Lambda function to retrieve user information from Amazon DynamoDB. Create an Amazon API Gateway endpoint to accept RESTful APIs. Send the API calls to the Lambda function. **Most Voted**

B. Create an Amazon Elastic Container Service (Amazon ECS) service behind an Application Load Balancer to retrieve user information from Amazon RDS. Create an Amazon API Gateway endpoint to accept RESTful APIs. Send the API calls to the Lambda function.

C. Create an Amazon Cognito user pool to authenticate users. **Most Voted**

D. Create an Amazon Cognito identity pool to authenticate users.

E. Use AWS Amplify to serve the frontend web content with HTML, CSS, and JS. Use an integrated Amazon CloudFront configuration. **Most Voted**

F. Use Amazon S3 static web hosting with PHP, CSS, and JS. Use Amazon CloudFront to serve the frontend web content.

Correct Answer: ACE

Community vote distribution

ACE (86%)

9%

Comments

manOfThePeople **Highly Voted** 2 years, 9 months ago

If in doubt between E or F. S3 doesn't support server-side scripts, PHP is a server-side script. The answer is ACE.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/WebsiteHosting.html>

upvoted 22 times

msdnpro **Highly Voted** 2 years, 11 months ago

Selected Answer: ACE

Option B (Amazon ECS) is not the best option since the website "can be idle for a long time", so Lambda (Option A) is a more cost-effective choice.

Option D is incorrect because User pools are for authentication (identity verification) while Identity pools are for authorization (access control).

Option F is wrong because S3 web hosting only supports static web files like HTML/CSS, and does not support PHP or JavaScript.

upvoted 12 times

0628atv 2 years, 11 months ago

https://aws.amazon.com/getting-started/projects/build-serverless-web-app-lambda-apigateway-s3-dynamodb-cognito/module-1/?nc1=h_ls

upvoted 3 times

Gape4 **Most Recent** 1 year, 11 months ago

Selected Answer: ACE

I will go for A C E

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: ACE

A: App may be idle for long time so Lambda is perfect (charge per invocation)

C: Cognito user pool for user auth

E: Amplify is low code web dev tool

B: Wrong, too much cost when idle
D: Identity pool is session management/identification. Does not help with auth.
F: S3 + PHP doesn't work also no security

upvoted 3 times

rcptryk 2 years, 6 months ago

Selected Answer: ACE

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/WebsiteHosting.html>
S3 doesn't support server-side scripting

upvoted 2 times

potomac 2 years, 7 months ago

Selected Answer: ACE

User Pool = authentication
Identity Pool = authorization
upvoted 7 times

thanhnv142 2 years, 7 months ago

A D F:
A: for hosting the dynamic content of the app. Pay as execution
D: for granting temporary privilege access to users who has paid a fee.
F: for hosting the static content of the app
upvoted 1 times

awsgeek75 2 years, 4 months ago

There is no static content in this web application so F is wrong. You cannot host PHP on S3 also so it is just wrong.
upvoted 2 times

kwang312 2 years, 8 months ago

Selected Answer: ACE

ACE is correct answer
upvoted 3 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: CEF

C) Create an Amazon Cognito user pool to authenticate users.

E) Use AWS Amplify to serve the frontend web content with HTML, CSS, and JS. Use an integrated CloudFront configuration.

F) Use Amazon S3 static web hosting with PHP, CSS, and JS. Use Amazon CloudFront to serve the frontend web content.
upvoted 1 times

awsgeek75 2 years, 4 months ago

There is no static content in this web application so F is wrong. You cannot host PHP on S3 also so it is just wrong.
upvoted 2 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: ACE

Build a web application = AWS Amplify
Sign in users = Amazon Cognito user pool
Traffic can be idle for a long time = AWS Lambda

Amazon S3 does not support server-side scripting such as PHP, JSP, or ASP.NET.
https://docs.aws.amazon.com/AmazonS3/latest/userguide/WebsiteHosting.html?icmpid=docs_amazons3_console#:~:text=website%20relies%20on-,server%2Dside,-processing%2C%20including%20server
Traffic can be idle for a long time = AWS Lambda
upvoted 2 times

james2033 2 years, 11 months ago

Selected Answer: ACE

Use exclusion method: No need for Container (no need run all time), remove B. PHP cannot run with static Amazon S3, remove F.
Use selection method: Idle for sometime, choose AWS Lambda, choose A. "Amazon Cognito is an identity platform for web and mobile apps."
(<https://docs.aws.amazon.com/cognito/latest/developerguide/what-is-amazon-cognito.html>), choose C. Create an identity pool
<https://docs.aws.amazon.com/cognito/latest/developerguide/tutorial-create-identity-pool.html> . AWS Amplify <https://aws.amazon.com/amplify/> for build full-stack web-app in hours.
upvoted 6 times

baba365 2 years, 11 months ago

Ans: ACF

use AWS SDK for PHP/JS with S3

https://docs.aws.amazon.com/sdk-for-php/v3/developer-guide/php_s3_code_examples.html

upvoted 1 times

unbendable 2 years, 7 months ago

did you actually read the link or just copy the first link from google here? the sdk is intended for usage in a php application. it does not say anything about php support in a s3 bucket

upvoted 2 times

Zox42 2 years, 11 months ago

Selected Answer: ACE

Answer is ACE

upvoted 2 times

jaydesai8 2 years, 11 months ago

Selected Answer: ACE

Lambda =serverless

User Pool = For user authentication

Amplify = hosting web/mobile apps

upvoted 3 times

live_reply_developers 2 years, 11 months ago

Selected Answer: ACE

S3 doesn't support PHP as stated in answer F.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/WebsiteHosting.html>

upvoted 4 times

wRhIH 2 years, 11 months ago

Selected Answer: ACE

I don't think S3 can handle anything dynamic such as PHP. So I go for ACE

upvoted 2 times

antropaws 2 years, 11 months ago

Selected Answer: ACF

ACF no doubt. Check the difference between user pools and identity pools.

upvoted 2 times

A media company uses an Amazon CloudFront distribution to deliver content over the internet. The company wants only premium customers to have access to the media streams and file content. The company stores all content in an Amazon S3 bucket. The company also delivers content on demand to customers for a specific purpose, such as movie rentals or music downloads.

Which solution will meet these requirements?

- A. Generate and provide S3 signed cookies to premium customers.
- B. Generate and provide CloudFront signed URLs to premium customers. **Most Voted**
- C. Use origin access control (OAC) to limit the access of non-premium customers.
- D. Generate and activate field-level encryption to block non-premium customers.

Correct Answer: B

Community vote distribution

B (100%)

Comments

NayeraB **Highly Voted** 1 year, 9 months ago

This question page is filled with premium customers I just can't upvoted 8 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: B

CloudFront Signed URL with Custom Policy are exactly for this.

A: Nope, cookies don't help as they don't restrict URL

C: Wrong. OAC for non-premium customers, how is that even possible without any details here?

D: Field encryption, while good idea, does not help restricting the content by customer

upvoted 3 times

pentium75 1 year, 11 months ago

Selected Answer: B

Authentication is done by Cloudfront, thus B

upvoted 4 times

ferdzcruz 1 year, 11 months ago

Content on demand = CloudFront. B

upvoted 3 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: B

Generate and provide CloudFront signed URLs to premium customers.

upvoted 3 times

TariqKipkemei 2 years, 4 months ago

Selected Answer: B

Use CloudFront signed URLs or signed cookies to restrict access to documents, business data, media streams, or content that is intended for selectec users, for example, users who have paid a fee.

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/PrivateContent.html#:~:text=CloudFront%20signed%20URLs>

upvoted 4 times

james2033 2 years, 4 months ago

Selected Answer: B

See <https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/private-content-signed-urls.html#private-content-how-signed-urls-work>

upvoted 2 times

haoAWS 2 years, 5 months ago

Selected Answer: B

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/PrivateContent.html>

Notice that A is not correct because it should be CloudFront signed URL, not S3.

upvoted 3 times

antropaws 2 years, 5 months ago

Why not C?

upvoted 1 times

antropaws 2 years, 5 months ago

<https://aws.amazon.com/blogs/networking-and-content-delivery/amazon-cloudfront-introduces-origin-access-control-oac/>

upvoted 2 times

pentium75 1 year, 11 months ago

OAC requires the consumers to have an IAM role with access to the S3 content, this is not what we're after here.

upvoted 2 times

alexandercamachop 2 years, 6 months ago

Selected Answer: B

Signed URLs

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/PrivateContent.html>

upvoted 3 times

haoAWS 2 years, 5 months ago

Then why A is incorrect?

upvoted 1 times

pentium75 1 year, 11 months ago

Because Authentication is done by Cloudfront, not S3.

upvoted 2 times

A company runs Amazon EC2 instances in multiple AWS accounts that are individually billed. The company recently purchased a Savings Plan. Because of changes in the company's business requirements, the company has decommissioned a large number of EC2 instances. The company wants to use its Savings Plan discounts on its other AWS accounts.

Which combination of steps will meet these requirements? (Choose two.)

A. From the AWS Account Management Console of the management account, turn on discount sharing from the billing preferences section. **Most Voted**

B. From the AWS Account Management Console of the account that purchased the existing Savings Plan, turn on discount sharing from the billing preferences section. Include all accounts.

C. From the AWS Organizations management account, use AWS Resource Access Manager (AWS RAM) to share the Savings Plan with other accounts.

D. Create an organization in AWS Organizations in a new payer account. Invite the other AWS accounts to join the organization from the management account. **Most Voted**

E. Create an organization in AWS Organizations in the existing AWS account with the existing EC2 instances and Savings Plan. Invite the other AWS accounts to join the organization from the management account.

Correct Answer: AD

Community vote distribution



Comments

Aigerim2010 **Highly Voted** 2 years, 11 months ago

i had this question today
upvoted 22 times

Nikki013 **Highly Voted** 2 years, 9 months ago

Selected Answer: AD

It is not recommended to have workload on the management account.
upvoted 7 times

JA2018 1 year, 6 months ago

confirm will fail the AWS WAR process
upvoted 2 times

In2718 **Most Recent** 9 months, 3 weeks ago

Selected Answer: AE

Chat says A and E -

A. Correct – You need to enable discount sharing in the management account's billing preferences.

B. Incorrect – Only the management (payer) account controls discount sharing, not member accounts.

C. Incorrect – Savings Plans cannot be shared via AWS RAM.

D. Creating a new organization is unnecessary; you should use the existing account with the Savings Plan.

E. Correct – You must set up (or use) AWS Organizations in the existing account that owns the Savings Plan, then invite the other accounts to join.

upvoted 1 times

zdi561 1 year, 4 months ago

Selected Answer: AD

google 'aws how to share ec2 shaving plans among accounts of an organization'
upvoted 1 times

FlyingHawk 1 year, 4 months ago

Selected Answer: AE

A E, firstly we create an organization in the existing AWS account with savings plan, so this account becomes a management account. Secondly, from the console of this account, turn on discount sharing.

upvoted 1 times

MaxMingxing 1 year, 5 months ago

Selected Answer: AE

Answer from ChatGPT

upvoted 1 times

FlyingHawk 1 year, 6 months ago

Selected Answer: AB

If the organization is already created, I am feeling A and B are the right one as "To share an Reserved Instance or Savings Plans discount with an account, both accounts must have sharing activated." <https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/ri-turn-off.html>. If no AWS organization was there, then you need to create an organization, then invite others join, D. then A and B.

upvoted 1 times

FlyingHawk 1 year, 4 months ago

B is not correct, Discount sharing can only be enabled from the management account of an AWS Organization, not from the account where the Savings Plan was purchased.

The correct answer is A E.

<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/ri-turn-off.html#ri-turn-on-process>

upvoted 1 times

awsgeek75 2 years, 5 months ago

Selected Answer: AD

For me, E makes no sense as the discount is with a new payer and cannot be transferred to an existing account unless customer service is involved.

upvoted 3 times

awsgeek75 2 years, 5 months ago

Also, "A company runs Amazon EC2 instances in multiple AWS accounts that are individually billed"

It's not billed, it is "billed"

upvoted 6 times

JA2018 1 year, 6 months ago

ROTFL :P

upvoted 1 times

pentium75 2 years, 5 months ago

Selected Answer: AD

Organization should be created by a new account that is reserved for management. Thus D, followed by A (discount sharing must be enabled in the management account).

upvoted 5 times

ErnShm 2 years, 9 months ago

AE

<https://repost.aws/questions/QUQoJuQLNOTDiyEuCLARIBFQ/transfer-savings-plan-across-organizations#:~:text=AWS%20Support%20can%20transfer%20Savings%20Plans%20from%20the%20management%20account%20to%20a%20member%20account%20or%20from%20a%20member%20account%20to%20the%20management%20account%20within%20a%20single%20Organization%20with%20an%20AWS%20Support%20Case.>

upvoted 1 times

lemur88 2 years, 9 months ago

Selected Answer: AD

Not E - it mentions using an account with existing EC2s as the management account, which goes against the best practice for a management account

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_best-practices_mgmt-acct.html

upvoted 4 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: AE

AE is best

upvoted 1 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: AE

AE is best

upvoted 1 times

james2033 2 years, 11 months ago

Selected Answer: AE

- B is not accepted, because "include all accounts", remove B.
 - D has "Create an organization in AWS Organization in a new payer account", it is wrong, remove D.
 - at C: AWS Resource Access Manager (AWS RAM) <https://aws.amazon.com/ram/> it is for security, not for billing. Remove C.
- Has A, E remain, and choosed.

A. "turn on discount sharing" is ok. This case: Has discount for many EC2 instances in one account, then want to share with other user. At E, create Organization, then share.

upvoted 2 times

pentium75 2 years, 5 months ago

What is the problem with "include all accounts"?

upvoted 2 times

antropaws 2 years, 11 months ago

Selected Answer: AE

I vote AE.

upvoted 2 times

MrAWSAssociate 2 years, 11 months ago

Selected Answer: AE

AE are correct !

upvoted 2 times

oras2023 2 years, 12 months ago

Selected Answer: CD

It's not good practice to create a payer account with any workload so it must be D.

By the reason that we need Organizations for sharing, then we need to turn on its from our PAYER account. (all sub-accounts start share discounts)

upvoted 1 times

oras2023 2 years, 12 months ago

changed to AD

upvoted 4 times

A retail company uses a regional Amazon API Gateway API for its public REST APIs. The API Gateway endpoint is a custom domain name that points to an Amazon Route 53 alias record. A solutions architect needs to create a solution that has minimal effects on customers and minimal data loss to release the new version of APIs.

Which solution will meet these requirements?

A. Create a canary release deployment stage for API Gateway. Deploy the latest API version. Point an appropriate percentage of traffic to the canary stage. After API verification, promote the canary stage to the production stage. **Most Voted**

B. Create a new API Gateway endpoint with a new version of the API in OpenAPI YAML file format. Use the import-to-update operation in merge mode into the API in API Gateway. Deploy the new version of the API to the production stage.

C. Create a new API Gateway endpoint with a new version of the API in OpenAPI JSON file format. Use the import-to-update operation in overwrite mode into the API in API Gateway. Deploy the new version of the API to the production stage.

D. Create a new API Gateway endpoint with new versions of the API definitions. Create a custom domain name for the new API Gateway API. Point the Route 53 alias record to the new API Gateway API custom domain name.

Correct Answer: A

Community vote distribution

A (100%)

Comments

AudreyNguyenHN **Highly Voted** 2 years, 10 months ago

We made it all the way here. Good luck everyone!

upvoted 20 times

ddddddddddww12 **Highly Voted** 2 years, 11 months ago

what are the total number of questions this package has as on 14 July 2023 , is it 544 or 551 ?

upvoted 8 times

NayeraB 2 years, 3 months ago

It's 20th of Feb 2024, and it's 798 (it says 804 at the top I donno why tho)

upvoted 3 times

DougZ 11 months ago

It's 13th of July 2025, there is 1019 questions.

upvoted 3 times

Faridtnx 2 years, 2 months ago

March 2024 its 825 questions. Constantly adding.

Doe ur question, ExamTopic always shows a few more question in listing compared to actual number

upvoted 5 times

MrAliMohsan 1 year, 9 months ago

August 2024 - 981 questions :(

upvoted 3 times

rosanna **Most Recent** 1 year, 6 months ago

Selected Answer: A

where do you guys get to know the answer to these things? I study from Stephan's content but lots of topics are uncovered, where am I supposed to find all the exam content in one place?

upvoted 1 times

awsgeek75 2 years, 5 months ago

Selected Answer: A

<https://docs.aws.amazon.com/apigateway/latest/developerguide/canary-release.html>

upvoted 4 times

potomac 2 years, 7 months ago

Selected Answer: A

In a canary release deployment, total API traffic is separated at random into a production release and a canary release with a pre-configured ratio. Typically, the canary release receives a small percentage of API traffic and the production release takes up the rest. The updated API features are only visible to API traffic through the canary. You can adjust the canary traffic percentage to optimize test coverage or performance.

<https://docs.aws.amazon.com/apigateway/latest/developerguide/canary-release.html>

upvoted 7 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: A

Using a canary release deployment allows incremental rollout of the new API version to a percentage of traffic. This minimizes impact on customers and potential data loss during the release.

upvoted 3 times

TariqKipkemei 2 years, 10 months ago

Selected Answer: A

Minimal effects on customers and minimal data loss = Canary deployment

upvoted 5 times

james2033 2 years, 11 months ago

Selected Answer: A

Key word "canary release". See this term in See: <https://www.jetbrains.com/teamcity/ci-cd-guide/concepts/canary-release/> and/or <https://martinfowler.com/bliki/CanaryRelease.html>

upvoted 2 times

Abrar2022 2 years, 11 months ago

Selected Answer: A

keyword: "latest versions on an api"

Canary release is a software development strategy in which a "new version of an API" (as well as other software) is deployed for testing purposes.

upvoted 4 times

jkhan2405 3 years ago

Selected Answer: A

It's A

upvoted 2 times

alexandercamachop 3 years ago

Selected Answer: A

A. Create a canary release deployment stage for API Gateway. Deploy the latest API version. Point an appropriate percentage of traffic to the canary stage. After API verification, promote the canary stage to the production stage.

Canary release meaning only certain percentage of the users.

upvoted 4 times

A company wants to direct its users to a backup static error page if the company's primary website is unavailable. The primary website's DNS records are hosted in Amazon Route 53. The domain is pointing to an Application Load Balancer (ALB). The company needs a solution that minimizes changes and infrastructure overhead.

Which solution will meet these requirements?

- A. Update the Route 53 records to use a latency routing policy. Add a static error page that is hosted in an Amazon S3 bucket to the records so that the traffic is sent to the most responsive endpoints.
- B. Set up a Route 53 active-passive failover configuration. Direct traffic to a static error page that is hosted in an Amazon S3 bucket when Route 53 health checks determine that the ALB endpoint is unhealthy. **Most Voted**
- C. Set up a Route 53 active-active configuration with the ALB and an Amazon EC2 instance that hosts a static error page as endpoints. Configure Route 53 to send requests to the instance only if the health checks fail for the ALB.
- D. Update the Route 53 records to use a multivalue answer routing policy. Create a health check. Direct traffic to the website if the health check passes. Direct traffic to a static error page that is hosted in Amazon S3 if the health check does not pass.

Correct Answer: B

Community vote distribution

B (83%)

D (17%)

Comments

TariqKipkemei **Highly Voted** 2 years, 6 months ago

Selected Answer: B

Set up a Route 53 active-passive failover configuration. Direct traffic to a static error page that is hosted in an Amazon S3 bucket when Route 53 health checks determine that the ALB endpoint is unhealthy.

upvoted 6 times

ssa03 **Highly Voted** 2 years, 9 months ago

Selected Answer: B

B is correct

upvoted 5 times

Danilus **Most Recent** 1 year, 7 months ago

key -minimizes changes and infrastructure overhead

key- backup static error page

Option B is the appropriate solution because it sets up automatic failover to an alternative endpoint (the static error page in S3) when the ALB becomes unavailable. Option D is not viable as it does not implement a true failover mechanism and may continue to direct traffic to an unhealthy ALB.

upvoted 1 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: D

Setting up a Route 53 active-passive failover configuration with the ALB as the primary endpoint and an Amazon S3 static website as the passive endpoint meets the requirements with minimal overhead.

Route 53 health checks can monitor the ALB health. If the ALB becomes unhealthy, traffic will automatically failover to the S3 static website. This provides automatic failover with minimal configuration changes

upvoted 3 times

Guru4Cloud 2 years, 9 months ago

Sorry. I mean B

upvoted 6 times

Nirav1112 2 years, 10 months ago

B is correct

upvoted 3 times

mrsoa 2 years, 10 months ago

Selected Answer: B

B seems correct

upvoted 4 times

Bmaster 2 years, 10 months ago

B is correct..

<https://repost.aws/knowledge-center/fail-over-s3-r53>

upvoted 4 times

awsgeek75 2 years, 5 months ago

Nice link find!

upvoted 2 times

A recent analysis of a company's IT expenses highlights the need to reduce backup costs. The company's chief information officer wants to simplify the on-premises backup infrastructure and reduce costs by eliminating the use of physical backup tapes. The company must preserve the existing investment in the on-premises backup applications and workflows.

What should a solutions architect recommend?

- A. Set up AWS Storage Gateway to connect with the backup applications using the NFS interface.
- B. Set up an Amazon EFS file system that connects with the backup applications using the NFS interface.
- C. Set up an Amazon EFS file system that connects with the backup applications using the iSCSI interface.
- D. Set up AWS Storage Gateway to connect with the backup applications using the iSCSI-virtual tape library (VTL) interface.

Most Voted

Correct Answer: D

Community vote distribution

D (100%)

Comments

awsgeek75 **Highly Voted** 1 year, 11 months ago

Selected Answer: D

Tape... lol

The company must preserve its existing investment so they want to keep using existing applications. This means EFS won't work, and NFS may not be compatible. VTL is the only thing that may be compatible with an application workflow that backs up to tapes.

Who the hell comes up with these questions!

upvoted 7 times

TariqKipkemei **Highly Voted** 2 years ago

Selected Answer: D

Use Tape Gateway to replace physical tapes on premises with virtual tapes on AWS—reducing your data storage costs without changing your tape-based backup workflows. Tape Gateway supports all leading backup applications and caches virtual tapes on premises for low-latency data access. It compresses your tape data, encrypts it, and stores it in a virtual tape library in Amazon Simple Storage Service (Amazon S3). From there, you can transfer it to either Amazon S3 Glacier Flexible Retrieval or Amazon S3 Glacier Deep Archive to help minimize your long-term storage costs.

<https://aws.amazon.com/storagegateway/vtl/#:~:text=Use-,Tape%20Gateway,-to%20replace%20physical>

upvoted 5 times

Nisarg2121 **Most Recent** 2 years, 1 month ago

Selected Answer: D

Tape Gateway is used for attaching to app.

upvoted 4 times

gouranga45 2 years, 2 months ago

Selected Answer: D

Option says it all

upvoted 4 times

Po_chih 2 years, 2 months ago

Selected Answer: D

Tape Gateway enables you to replace using physical tapes on premises with virtual tapes in AWS without changing existing backup workflows. Tape Gateway supports all leading backup applications and caches virtual tapes on premises for low-latency data access. Tape Gateway encrypts data between the gateway and AWS for secure data transfer, and compresses data and transitions virtual tapes between Amazon S3 and Amazon S3 Glacier Flexible Retrieval, or Amazon S3 Glacier Deep Archive, to minimize storage costs.

upvoted 3 times

ssa03 2 years, 3 months ago

Selected Answer: D

https://aws.amazon.com/storagegateway/vtl/?nc1=h_ls

upvoted 2 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: D

Set up AWS Storage Gateway to connect with the backup applications using the iSCSI-virtual tape library (VTL) interface.

upvoted 2 times

Bmaster 2 years, 4 months ago

D is correct

https://aws.amazon.com/storagegateway/vtl/?nc1=h_ls

upvoted 2 times

A company has data collection sensors at different locations. The data collection sensors stream a high volume of data to the company. The company wants to design a platform on AWS to ingest and process high-volume streaming data. The solution must be scalable and support data collection in near real time. The company must store the data in Amazon S3 for future reporting.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use Amazon Kinesis Data Firehose to deliver streaming data to Amazon S3. **Most Voted**
- B. Use AWS Glue to deliver streaming data to Amazon S3.
- C. Use AWS Lambda to deliver streaming data and store the data to Amazon S3.
- D. Use AWS Database Migration Service (AWS DMS) to deliver streaming data to Amazon S3.

Correct Answer: A

Community vote distribution

A (83%)

D (17%)

Comments

emakid 1 year, 11 months ago

Selected Answer: A

Amazon Kinesis Data Firehose is a fully managed service for delivering real-time streaming data to destinations such as Amazon S3, Amazon Redshift, Amazon Elasticsearch Service, and Splunk. It requires minimal setup and maintenance, automatically scales to match the throughput of your data, and offers near real-time data delivery with minimal operational overhead.

upvoted 3 times

awsgeek75 2 years, 5 months ago

Selected Answer: A

High volume streaming data = Kinesis
B: Glue is for ETL (to S3 is ok) but not for streaming
C: Lambda more overhead
D: Streaming != Data migration

upvoted 3 times

pentium75 2 years, 5 months ago

Selected Answer: A

sensor data = Kinesis

upvoted 4 times

TariqKipkemei 2 years, 6 months ago

Selected Answer: A

Amazon Kinesis Data Firehose: Capture, transform, and load data streams into AWS data stores (S3) in near real-time.

https://aws.amazon.com/pm/kinesis/?gclid=CjwKCAiAu9yqBhBmEiwAHTx5px9z182o0HBEX0BGXU7VeOCOdNpkJMxgbSfvcHINKN4NHVnbEa0Y1xoCuU0QAvD_BwE&trk=239a97c0-9c5d-42a5-ac65-7381b62f3756&sc_channel=ps&ef_id=CjwKCAiAu9yqBhBmEiwAHTx5px9z182o0HBEX0BGXU7VeOCOdNpkJMxgbSfvcHINKN4NHVnbEa0Y1xoCuU0Q/vD_BwE:G:s&s_kwcid=AL!442213!651612444428!e!!g!!kinesis%20firehose!19836376048!149982297311#:~:text=Kinesis%20Data%20Firehose-,Capture%2C,-transform%2C%20and%20load

upvoted 3 times

potomac 2 years, 7 months ago

Selected Answer: A

A for sure

upvoted 3 times

ssa03 2 years, 9 months ago

Selected Answer: A

Correct Answer: A
upvoted 3 times

manOfThePeople 2 years, 9 months ago

A is the answer, near real-time = Kinesis Data Firehose.
upvoted 4 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: D

Use Amazon Kinesis Data Firehose to deliver streaming data to Amazon S3
upvoted 3 times

pentium75 2 years, 5 months ago

That is A
upvoted 2 times

bjexamprep 2 years, 10 months ago

Selected Answer: D

Kinesis Data Firehose is only real-time answer
upvoted 3 times

pentium75 2 years, 5 months ago

That is A
upvoted 3 times

mrsoa 2 years, 10 months ago

Selected Answer: A

A is the correct answer
upvoted 3 times

Deepakin96 2 years, 10 months ago

Selected Answer: A

Kinesis = Near Real Time
upvoted 4 times

Kaiden123 2 years, 10 months ago

Selected Answer: A

Data collection in near real time = Amazon Kinesis Data Firehose
upvoted 4 times

Bmaster 2 years, 10 months ago

A is correct..
upvoted 2 times

A company has separate AWS accounts for its finance, data analytics, and development departments. Because of costs and security concerns, the company wants to control which services each AWS account can use.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use AWS Systems Manager templates to control which AWS services each department can use.
- B. Create organization units (OUs) for each department in AWS Organizations. Attach service control policies (SCPs) to the OUs.
Most Voted
- C. Use AWS CloudFormation to automatically provision only the AWS services that each department can use.
- D. Set up a list of products in AWS Service Catalog in the AWS accounts to manage and control the usage of specific AWS services.

Correct Answer: B

Community vote distribution

B (89%)

11%

Comments

zdi561 1 year, 4 months ago

Selected Answer: D

SCP set a guardrail (what you can not do, permission on IAM user and role), ASC provides templates services that you can only use.
upvoted 1 times

awsgeek75 1 year, 11 months ago

Selected Answer: B

Departments = Organizational Units
upvoted 3 times

TariqKipkemei 2 years ago

Selected Answer: B

Create organization units (OUs) for each department in AWS Organizations. Attach service control policies (SCPs) to the OUs
upvoted 2 times

ssa03 2 years, 3 months ago

Selected Answer: B

Correct Answer: B
upvoted 2 times

lemur88 2 years, 3 months ago

Selected Answer: B

SCPs to centralize permissioning
upvoted 2 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: B

Create organization units (OUs) for each department in AWS Organizations. Attach service control policies (SCPs) to the OUs.
upvoted 2 times

xyb 2 years, 4 months ago

Selected Answer: B

control services --> SCP
upvoted 2 times

Ale1973 2 years, 4 months ago

Selected Answer: D

My rational: Scenary is "A company has separate AWS accounts", it is not mentioning anything about use of Organizations or needs related to centralized managment of these accounts.
Then, set up a list of products in AWS Service Catalog in the AWS accounts (on each AWS account) is the best way to manage and control the usage of specific AWS services.
upvoted 1 times

pentium75 1 year, 11 months ago

"Separate AWS accounts" just says that it's multiple accounts, it does not indicate that they are NOT connected into a organization.

Service Catalog alone does not restrict anything. You'd need to create a service in Service Catalog for everything you're allowing to use, then grant permissions on those services, and you'd need to remove other permissions from everyone. All of which is not mentioned in D. Just "setting up a lis of products in AWS Service Catalog in the AWS accounts" will not restrict anyone from doing what he could do before.
upvoted 3 times

mrsoa 2 years, 4 months ago

Selected Answer: B

BBBBBBBBB
upvoted 2 times

Deepakin96 2 years, 4 months ago

Selected Answer: B

To control different AWS account you required AWS Organisation
upvoted 2 times

Bmaster 2 years, 4 months ago

B is correct!!!!
upvoted 2 times

A company has created a multi-tier application for its ecommerce website. The website uses an Application Load Balancer that resides in the public subnets, a web tier in the public subnets, and a MySQL cluster hosted on Amazon EC2 instances in the private subnets. The MySQL database needs to retrieve product catalog and pricing information that is hosted on the internet by a third-party provider. A solutions architect must devise a strategy that maximizes security without increasing operational overhead.

What should the solutions architect do to meet these requirements?

- A. Deploy a NAT instance in the VPC. Route all the internet-based traffic through the NAT instance.
- B. Deploy a NAT gateway in the public subnets. Modify the private subnet route table to direct all internet-bound traffic to the NAT gateway. **Most Voted**
- C. Configure an internet gateway and attach it to the VPC. Modify the private subnet route table to direct internet-bound traffic to the internet gateway.
- D. Configure a virtual private gateway and attach it to the VPC. Modify the private subnet route table to direct internet-bound traffic to the virtual private gateway.

Correct Answer: B

Community vote distribution

B (100%)

Comments

awsgeek75 1 year, 11 months ago

Selected Answer: B

A: Probably an old question so this option is here but NAT instance is overhead

C: Not secure as IG opens up a lot of things

D: VPG connects to a service

B: NG is managed solution. Secure by config

upvoted 3 times

TariqKipkemei 2 years ago

Selected Answer: B

Deploy a NAT gateway in the public subnets. Modify the private subnet route table to direct all internet-bound traffic to the NAT gateway

upvoted 4 times

ssa03 2 years, 3 months ago

Selected Answer: B

Correct Answer: B

upvoted 4 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: B

Deploy a NAT gateway in the public subnets. Modify the private subnet route table to direct all internet-bound traffic to the NAT gateway.

upvoted 3 times

Deepakin96 2 years, 4 months ago

Selected Answer: B

NAT Gateway is safe

upvoted 3 times

Bmaster 2 years, 4 months ago

B is correct

upvoted 2 times

A company is using AWS Key Management Service (AWS KMS) keys to encrypt AWS Lambda environment variables. A solutions architect needs to ensure that the required permissions are in place to decrypt and use the environment variables.

Which steps must the solutions architect take to implement the correct permissions? (Choose two.)

- A. Add AWS KMS permissions in the Lambda resource policy.
- B. Add AWS KMS permissions in the Lambda execution role. **Most Voted**
- C. Add AWS KMS permissions in the Lambda function policy.
- D. Allow the Lambda execution role in the AWS KMS key policy. **Most Voted**
- E. Allow the Lambda resource policy in the AWS KMS key policy.

Correct Answer: BD

Community vote distribution

BD (100%)

Comments

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: BD

To decrypt environment variables encrypted with AWS KMS, Lambda needs to be granted permissions to call KMS APIs. This is done in two places:

The Lambda execution role needs kms:Decrypt and kms:GenerateDataKey permissions added. The execution role governs what AWS services the function code can access.

The KMS key policy needs to allow the Lambda execution role to have kms:Decrypt and kms:GenerateDataKey permissions for that specific key. This allows the execution role to use that particular key.

upvoted 8 times

wizcloudifa **Highly Voted** 1 year, 7 months ago

Selected Answer: BD

As per the principle of least privilege, granting permissions = role level

upvoted 5 times

TariqKipkemei **Most Recent** 2 years ago

Selected Answer: BD

Allow the Lambda execution role in the AWS KMS key policy then add AWS KMS permissions in the role.

upvoted 3 times

ssa03 2 years, 3 months ago

Selected Answer: BD

Correct Answer: BD

upvoted 3 times

Nirav1112 2 years, 4 months ago

its B & D

upvoted 2 times

mrsoa 2 years, 4 months ago

Selected Answer: BD

BD BD BD BD

upvoted 2 times

Deepakin96 2 years, 4 months ago

Selected Answer: BD

Its B and D

upvoted 2 times

Bmaster 2 years, 4 months ago

My choice is B,D

upvoted 2 times


A company has a financial application that produces reports. The reports average 50 KB in size and are stored in Amazon S3. The reports are frequently accessed during the first week after production and must be stored for several years. The reports must be retrievable within 6 hours.

Which solution meets these requirements MOST cost-effectively?

- A. Use S3 Standard. Use an S3 Lifecycle rule to transition the reports to S3 Glacier after 7 days. **Most Voted**
- B. Use S3 Standard. Use an S3 Lifecycle rule to transition the reports to S3 Standard-Infrequent Access (S3 Standard-IA) after 7 days.
- C. Use S3 Intelligent-Tiering. Configure S3 Intelligent-Tiering to transition the reports to S3 Standard-Infrequent Access (S3 Standard-IA) and S3 Glacier.
- D. Use S3 Standard. Use an S3 Lifecycle rule to transition the reports to S3 Glacier Deep Archive after 7 days.

Correct Answer: A

Community vote distribution



Comments

zjcorpuz **Highly Voted** 2 years, 10 months ago

Answer is A

Amazon S3 Glacier:

Expedited Retrieval: Provides access to data within 1-5 minutes.

Standard Retrieval: Provides access to data within 3-5 hours.

Bulk Retrieval: Provides access to data within 5-12 hours.

Amazon S3 Glacier Deep Archive:

Standard Retrieval: Provides access to data within 12 hours.

Bulk Retrieval: Provides access to data within 48 hours.

upvoted 27 times

oayoad **Highly Voted** 2 years, 9 months ago

Selected Answer: C

All the "...after 7 days" options are wrong.

Before you transition objects to S3 Standard-IA or S3 One Zone-IA, you must store them for at least 30 days in Amazon S3

[https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-](https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html#:~:text=Minimum%20Days%20for%20Transition%20to%20S3%20Standard%2DIA%20or%20S3%20One%20Zone%2DIA)

[considerations.html#:~:text=Minimum%20Days%20for%20Transition%20to%20S3%20Standard%2DIA%20or%20S3%20One%20Zone%2DIA](https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html#:~:text=Minimum%20Days%20for%20Transition%20to%20S3%20Standard%2DIA%20or%20S3%20One%20Zone%2DIA)

upvoted 11 times

Hades2231 2 years, 9 months ago

This is worth noticing! Glad I came across your comment 1 day before my test.

upvoted 4 times

Marco_St 2 years, 5 months ago

so Could I ask is A or C for this question? I voted for A but it seems you had the same question in the exam and it was C? Thanks! I will attend the exam soon.

upvoted 2 times

MatAlves 1 year, 8 months ago

After reading the document from your link, it's clear that NONE of the restrictions apply to S3 standard, but to S3 Standard IA.

A is the correct answer.

upvoted 2 times

franbarberan 2 years, 8 months ago

the 7 days limitation is only if you want to move from s3 standart to S3 Standard-IA or S3 One Zone-IA, if you move to s3 glacier dont have this limitation, correct answer is A

upvoted 15 times

LeonSauveterre 1 year, 6 months ago

Objects must be stored for at least 30 days before transitioning to S3 Standard-IA or S3 One Zone-IA.

Option A is to S3 Glacier. A is correct.

upvoted 2 times

TEC65 Most Recent 1 year, 1 month ago

Selected Answer: A

a is the answer

upvoted 1 times

rosanna 1 year, 6 months ago

Selected Answer: A

there's a period before you can transition to any storage class (down the hierarchy) "except for glacier of all types" a minimum of 30 days. However, when transitioning to Glacier directly there's no minimum period to wait so you can choose "A" safely

upvoted 1 times

1e22522 1 year, 10 months ago

Selected Answer: A

Its A ya bunch of nerds

upvoted 4 times

Gape4 1 year, 11 months ago

Selected Answer: C

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html>

upvoted 1 times

Linuslin 2 years, 1 month ago

Selected Answer: A

C is incorrect.

Unsupported lifecycle transitions

Amazon S3 does not support any of the following lifecycle transitions.

You can't transition from the following:

Any storage class to the S3 Standard storage class.

Any storage class to the Reduced Redundancy Storage (RRS) class.

The S3 Intelligent-Tiering storage class to the S3 Standard-IA storage class.

The S3 One Zone-IA storage class to the S3 Intelligent-Tiering, S3 Standard-IA, or S3 Glacier Instant Retrieval storage classes.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html>

upvoted 4 times

awsgeek75 2 years, 5 months ago

Selected Answer: A

BC are lifecycle with tiering and infrequent access which are not required here.

D is deep archive and can take hours to retrieve so it is not suitable

A is cheapest workable option

upvoted 4 times

Marco_St 2 years, 5 months ago

Selected Answer: A

frequent access pattern- Standard.

upvoted 2 times

pentium75 2 years, 5 months ago

Selected Answer: A

Not B - More expensive than A

Not C - Intelligent-Tiering moves only objects of at least 128 KB

Not D - Glacier Deep Archive takes more than 6 hours to retrieve

upvoted 4 times

TariqKipkemei 2 years, 6 months ago

Selected Answer: A

Any option with S3 Intelligent-Tiering is out, this is only required when the access patterns are unknown.

From the question the access patterns are well known, enough to tie the frequently accessed reports to S3 standard and transition them to S3 glacier

after 7 days.

upvoted 4 times

iwannabeawsgod 2 years, 7 months ago

Selected Answer: A

its A for me

upvoted 3 times

Carlos_O 2 years, 8 months ago

Selected Answer: A

Tiene mas sentido

upvoted 2 times

sl2man 2 years, 8 months ago

Selected Answer: A

Option A

Amazon S3 Glacier Standard Retrieval: Provides access to data within 3-5 hours.

upvoted 4 times

Ramdi1 2 years, 8 months ago

Selected Answer: A

most cost effective has to be glacier so A

With C it is using intelligence tiering which is 30 days minimum from what I have read, I may be wrong on how I read that.

upvoted 2 times

tabbyDolly 2 years, 8 months ago

answer A

frequent access during the first week -> keeps data in s3 standard for 7 days

stored for several year and retrievable within 6 hours -> can be moved to s3 glacier for data archive purpose

upvoted 2 times

anikety123 2 years, 9 months ago

Selected Answer: A

Its A. Data cannot be transitioned from Intelligent Tiering to Standard IA

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html>

upvoted 4 times

A company needs to optimize the cost of its Amazon EC2 instances. The company also needs to change the type and family of its EC2 instances every 2-3 months.

What should the company do to meet these requirements?

- A. Purchase Partial Upfront Reserved Instances for a 3-year term.
- B. Purchase a No Upfront Compute Savings Plan for a 1-year term. **Most Voted**
- C. Purchase All Upfront Reserved Instances for a 1-year term.
- D. Purchase an All Upfront EC2 Instance Savings Plan for a 1-year term.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 9 months ago

Selected Answer: B

The key considerations are:

The company needs flexibility to change EC2 instance types and families every 2-3 months. This rules out Reserved Instances which lock you into an instance type and family for 1-3 years.

A Compute Savings Plan allows switching instance types and families freely within the term as needed. No Upfront is more flexible than All Upfront.

A 1-year term balances commitment and flexibility better than a 3-year term given the company's changing needs.

With No Upfront, the company only pays for usage monthly without an upfront payment. This optimizes cost.

upvoted 11 times

Kiki_Pass **Highly Voted** 1 year, 10 months ago

Selected Answer: B

"EC2 Instance Savings Plans give you the flexibility to change your usage between instances WITHIN a family in that region."

<https://aws.amazon.com/savingsplans/compute-pricing/>

upvoted 5 times

TariqKipkemei **Most Recent** 1 year, 6 months ago

Selected Answer: B

Only Compute Savings Plan allows you to change instance family.

upvoted 5 times

avkya 1 year, 10 months ago

Selected Answer: B

"needs to change the type and family of its EC2 instances". that means B I think.

upvoted 3 times

mrsoa 1 year, 10 months ago

Selected Answer: B

B is the right answer

upvoted 3 times

Bmaster 1 year, 10 months ago

B is correct..

'EC2 Instance Savings Plans' can't change 'family'.

upvoted 2 times

Josantru 1 year, 10 months ago

Correct B.

To change 'Family' always Compute saving plan, right?

upvoted 4 times

A solutions architect needs to review a company's Amazon S3 buckets to discover personally identifiable information (PII). The company stores the PII data in the us-east-1 Region and us-west-2 Region.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Configure Amazon Macie in each Region. Create a job to analyze the data that is in Amazon S3. **Most Voted**
- B. Configure AWS Security Hub for all Regions. Create an AWS Config rule to analyze the data that is in Amazon S3.
- C. Configure Amazon Inspector to analyze the data that is in Amazon S3.
- D. Configure Amazon GuardDuty to analyze the data that is in Amazon S3.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: A

The key reasons are:

Amazon Macie is designed specifically for discovering and classifying sensitive data like PII in S3. This makes it the optimal service to use. Macie can be enabled directly in the required Regions rather than enabling it across all Regions which is unnecessary. This minimizes overhead. Macie can be set up to automatically scan the specified S3 buckets on a schedule. No need to create separate jobs. Security Hub is for security monitoring across AWS accounts, not specific for PII discovery. More overhead than needed. Inspector and GuardDuty are not built for PII discovery in S3 buckets. They provide broader security capabilities.

upvoted 6 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: A

PII = Macie
Security Hub: Organisation security and logging not for PII
Inspector: Infra vulnerability management
GuardDuty: Network protection
upvoted 4 times

TariqKipkemei 2 years ago

Selected Answer: A

Amazon Macie = PII
upvoted 2 times

mrsoa 2 years, 4 months ago

Selected Answer: A

AWS Macie = PII detection
upvoted 4 times

Deepakin96 2 years, 4 months ago

Selected Answer: A

Amazon Macie will identify all PII
upvoted 3 times

A company's SAP application has a backend SQL Server database in an on-premises environment. The company wants to migrate its on-premises application and database server to AWS. The company needs an instance type that meets the high demands of its SAP database. On-premises performance data shows that both the SAP application and the database have high memory utilization.

Which solution will meet these requirements?

- A. Use the compute optimized instance family for the application. Use the memory optimized instance family for the database.
- B. Use the storage optimized instance family for both the application and the database.
- C. Use the memory optimized instance family for both the application and the database. **Most Voted**
- D. Use the high performance computing (HPC) optimized instance family for the application. Use the memory optimized instance family for the database.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 9 months ago

Selected Answer: C

Since both the app and database have high memory needs, the memory optimized family like R5 instances meet those requirements well. Using the same instance family simplifies management and operations, rather than mixing instance types.

Compute optimized instances may not provide enough memory for the SAP app's needs.

Storage optimized is overkill for the database's compute and memory needs.

HPC is overprovisioned for the SAP app.

upvoted 16 times

manOfThePeople **Highly Voted** 1 year, 9 months ago

High memory utilization = memory optimized.

C is the answer

upvoted 5 times

LeonSauveterre **Most Recent** 1 year, 6 months ago

Selected Answer: C

Why D is incorrect: HPC family is geared toward compute-heavy workloads such as simulations or scientific calculations, not appropriate for a memory-intensive application like SAP.

upvoted 1 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: C

Use the memory optimized instance family for both the application and the database

upvoted 3 times

mrsoa 1 year, 10 months ago

Selected Answer: C

I think its C

upvoted 3 times

A company runs an application in a VPC with public and private subnets. The VPC extends across multiple Availability Zones. The application runs on Amazon EC2 instances in private subnets. The application uses an Amazon Simple Queue Service (Amazon SQS) queue.

A solutions architect needs to design a secure solution to establish a connection between the EC2 instances and the SQS queue.

Which solution will meet these requirements?

- A. Implement an interface VPC endpoint for Amazon SQS. Configure the endpoint to use the private subnets. Add to the endpoint a security group that has an inbound access rule that allows traffic from the EC2 instances that are in the private subnets. **Most Voted**
- B. Implement an interface VPC endpoint for Amazon SQS. Configure the endpoint to use the public subnets. Attach to the interface endpoint a VPC endpoint policy that allows access from the EC2 instances that are in the private subnets.
- C. Implement an interface VPC endpoint for Amazon SQS. Configure the endpoint to use the public subnets. Attach an Amazon SQS access policy to the interface VPC endpoint that allows requests from only a specified VPC endpoint.
- D. Implement a gateway endpoint for Amazon SQS. Add a NAT gateway to the private subnets. Attach an IAM role to the EC2 instances that allows access to the SQS queue.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: A

An interface VPC endpoint is a private way to connect to AWS services without having to expose your VPC to the public internet. This is the most secure way to connect to Amazon SQS from the private subnets.

Configuring the endpoint to use the private subnets ensures that the traffic between the EC2 instances and the SQS queue is only within the VPC. This helps to protect the traffic from being intercepted by a malicious actor.

Adding a security group to the endpoint that has an inbound access rule that allows traffic from the EC2 instances that are in the private subnets further restricts the traffic to only the authorized sources. This helps to prevent unauthorized access to the SQS queue.

upvoted 12 times

Bmaster **Highly Voted** 2 years, 4 months ago

A is correct.

B,C: 'Configuring endpoints to use public subnets' --> Invalid

D: No Gateway Endpoint for SQS.

upvoted 6 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: A

BC are using public subnets so not useful for security

D uses gateway endpoint which is not useful to connect to SQS

A: <https://docs.aws.amazon.com/vpc/latest/privatelink/aws-services-privatelink-support.html>

upvoted 2 times

awsgeek75 1 year, 10 months ago

Sorry, the link is wrong for A. Please ignore it!

upvoted 2 times

ShawnTang 1 year, 12 months ago

A seems the most suitable,
but security group can't add to the endpoint directly, right?

upvoted 2 times

TariqKipkemei 2 years ago

Selected Answer: A

Answer is A

upvoted 2 times

TariqKipkemei 2 years ago

Interface endpoints enable connectivity to services over AWS PrivateLink. It is a collection of one or more elastic network interfaces with a private IP address that serves as an entry point for traffic destined to a supported service.

Implement an interface VPC endpoint for Amazon SQS. Configure the endpoint to use the private subnets. Add to the endpoint a security group that has an inbound access rule that allows traffic from the EC2 instances that are in the private subnets.

upvoted 2 times

potomac 2 years, 1 month ago

A is correct

upvoted 2 times

mrsoa 2 years, 4 months ago

Selected Answer: A

I think its A

upvoted 2 times

A solutions architect is using an AWS CloudFormation template to deploy a three-tier web application. The web application consists of a web tier and an application tier that stores and retrieves user data in Amazon DynamoDB tables. The web and application tiers are hosted on Amazon EC2 instances, and the database tier is not publicly accessible. The application EC2 instances need to access the DynamoDB tables without exposing API credentials in the template.

What should the solutions architect do to meet these requirements?

- A. Create an IAM role to read the DynamoDB tables. Associate the role with the application instances by referencing an instance profile.
- B. Create an IAM role that has the required permissions to read and write from the DynamoDB tables. Add the role to the EC2 instance profile, and associate the instance profile with the application instances. **Most Voted**
- C. Use the parameter section in the AWS CloudFormation template to have the user input access and secret keys from an already-created IAM user that has the required permissions to read and write from the DynamoDB tables.
- D. Create an IAM user in the AWS CloudFormation template that has the required permissions to read and write from the DynamoDB tables. Use the GetAtt function to retrieve the access and secret keys, and pass them to the application instances through the user data.

Correct Answer: B

Community vote distribution

B (87%)

13%

Comments

upliftinghut 1 year, 10 months ago

Selected Answer: B

best practice is using IAM role for database access. From app to DB => need both read & write, only B meets these 2
upvoted 3 times

pentium75 1 year, 11 months ago

Selected Answer: B

Application "stores and retrieves" data in DynamoDB while A grants only access "to read".
upvoted 3 times

Nisarg2121 2 years, 1 month ago

Selected Answer: B

B is correct, A total wrong because "read the DynamoDB tables", so what about write in database.
upvoted 4 times

darekw 2 years, 3 months ago

question says: ...application tier stores and retrieves user data in Amazon DynamoDB tables... so it needs read and write access
A) is only read access
B) seems to be the right answer
upvoted 2 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: B

Option B is the correct approach to meet the requirements:

Create an IAM role with permissions to access DynamoDB
Add the IAM role to an EC2 Instance Profile
Associate the Instance Profile with the application EC2 instances
This allows the instances to assume the IAM role to obtain temporary credentials to access DynamoDB.

upvoted 3 times

anibinaadi 2 years, 3 months ago

Explanation. Both A and B seems suitable. But Option A is incorrect because it says "Associate the role with the application instances by referencing an instance profile". Which just only a Part of the solution.

In API/AWS CLI following steps are required to complete the Role-> instance profile association-> to instance.

1. Create an IAM Role

2. add-role-to-instance-profile (aws iam add-role-to-instance-profile --role-name S3Access --instance-profile-name Webserver)

3. associate-iam-instance-profile (aws ec2 associate-iam-instance-profile --instance-id i-123456789abcde123 --iam-instance-profile Name=admin-role)

hence Option B is correct.

upvoted 3 times

DannyKang5649 2 years, 4 months ago

Selected Answer: B

Why "No read and write" ? The question clearly states that application tier STORE and RETRIEVE the data from DynamoDB. Which means write and read... I think answer should be B

upvoted 3 times

xyb 2 years, 4 months ago

Selected Answer: B

<https://www.examtopycs.com/discussions/amazon/view/80755-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 2 times

Ale1973 2 years, 4 months ago

Selected Answer: B

My rationl: Option A is wrong because the scenario says "stores and retrieves user data in Amazon DynamoDB tables", STORES and RETRIVE, if you set a role to READ, you can write on DinamoDB database

upvoted 2 times

mrsoa 2 years, 4 months ago

Selected Answer: A

AAAAAAAAA

upvoted 2 times

pentium75 1 year, 11 months ago

No because it grants only read access

upvoted 3 times

kangho 2 years, 4 months ago

Selected Answer: A

A is correct

upvoted 1 times

pentium75 1 year, 11 months ago

No because it grants only read access

upvoted 3 times

A solutions architect manages an analytics application. The application stores large amounts of semistructured data in an Amazon S3 bucket. The solutions architect wants to use parallel data processing to process the data more quickly. The solutions architect also wants to use information that is stored in an Amazon Redshift database to enrich the data.

Which solution will meet these requirements?

- A. Use Amazon Athena to process the S3 data. Use AWS Glue with the Amazon Redshift data to enrich the S3 data.
- B. Use Amazon EMR to process the S3 data. Use Amazon EMR with the Amazon Redshift data to enrich the S3 data. **Most Voted**
- C. Use Amazon EMR to process the S3 data. Use Amazon Kinesis Data Streams to move the S3 data into Amazon Redshift so that the data can be enriched.
- D. Use AWS Glue to process the S3 data. Use AWS Lake Formation with the Amazon Redshift data to enrich the S3 data.

Correct Answer: B

Community vote distribution

B (79%)

A (21%)

Comments

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: B

Option B is the correct solution that meets the requirements:

Use Amazon EMR to process the semi-structured data in Amazon S3. EMR provides a managed Hadoop framework optimized for processing large datasets in S3.

EMR supports parallel data processing across multiple nodes to speed up the processing.

EMR can integrate directly with Amazon Redshift using the EMR-Redshift integration. This allows querying the Redshift data from EMR and joining it with the S3 data.

This enables enriching the semi-structured S3 data with the information stored in Redshift

upvoted 18 times

zjcorpuz **Highly Voted** 2 years, 4 months ago

By combining AWS Glue and Amazon Redshift, you can process the semistructured data in parallel using Glue ETL jobs and then store the processed and enriched data in a structured format in Amazon Redshift. This approach allows you to perform complex analytics efficiently and at scale.

upvoted 9 times

upliftinghut **Most Recent** 1 year, 10 months ago

Selected Answer: B

D: not relevant, data is semistructured and Glue is more batch than stream data

A: not correct, Athena is for querying data

B & C look ok but C is out => redundant with Kinesis data stream; EMR already processed data as input into Redshift for parallel processing

Only B is most logical

upvoted 4 times

awsgeek75 1 year, 10 months ago

Selected Answer: B

Key requirement: parallel data processing

parallel data processing is EMR (Kind of Apache Hadoop) so it only leave B and C

C is Kinesis to Redshift which is pointless logic here

B EMR for S3 and EMR for Redshift gives maximum parallel processing here

upvoted 3 times

pentium75 1 year, 11 months ago

Selected Answer: B

A has a pitfall, "use Amazon Athena to PROCESS the data". With Athena you can query, not process, data.
C is wrong because Kinesis has no place here.
D is wrong because it does not process the Redshift data, and Glue does ETL, not analyze

Thus it's B. EMR can use semi-structured data from S3 and structured data from Redshift and is ideal for "parallel data processing" of "large amounts" of data.

upvoted 8 times

aws94 1 year, 12 months ago

Selected Answer: B

large amount of data + parallel data processing = EMR

upvoted 3 times

[Removed] 2 years ago

Selected Answer: A

Amazon Athena is an interactive query service that makes it easy to analyze data directly in Amazon Simple Storage Service (Amazon S3) using standard SQL.

upvoted 1 times

pentium75 1 year, 11 months ago

Y, but A says "process", not "query" data with Athena.

upvoted 2 times

SHAAHIBHUSHANAWS 2 years ago

Selected Answer: D

Glue use apache pyspark cluster for parallel processing. EMR or Glue are possible options. Glue is serverless so better use this plus pyspark is in memory parallel processing.

upvoted 1 times

aragornfsm 2 years ago

i think a is correct

semistructured data ==> Athena

upvoted 1 times

pentium75 1 year, 11 months ago

"Hadoop [as used by EMR] helps you turn petabytes of un-structured or semi-structured data into useful insights"

<https://aws.amazon.com/emr/features/hadoop/>

upvoted 2 times

riyasara 2 years ago

Athena is not designed for parallel data processing. So it's B

upvoted 3 times

TariqKipkemei 2 years ago

Selected Answer: A

Answer is A

upvoted 1 times

TariqKipkemei 2 years ago

Selected Answer: B

From this documentation looks like EMR cannot interface with S3.

<https://aws.amazon.com/emr/>

I will settle with option A.

upvoted 2 times

pentium75 1 year, 11 months ago

Of course EMR can access S3

<https://docs.aws.amazon.com/emr/latest/ManagementGuide/emr-plan-file-systems.html>

upvoted 2 times

bogobob 2 years ago

Selected Answer: B

For those answering A, AWS Glue can directly query S3, it can't use Athena as a source of data. The questions say the Redshift data should be user to "enrich" which means that the redshift data needs to be "added" to the s3 data. A doesn't allow that.

upvoted 2 times

hungta 2 years ago

Selected Answer: B

Choose option B.

Option A is not correct. Amazon Athena is suitable for querying data directly from S3 using SQL and allows parallel processing of S3 data. AWS Glue can be used for data preparation and enrichment but might not directly integrate with Amazon Redshift for enrichment.

upvoted 2 times

potomac 2 years, 1 month ago

Selected Answer: A

Athena and Redshift both do SQL query

upvoted 1 times

Sab123 2 years, 2 months ago

Selected Answer: A

semi-structure supported by Athena not by EMR

upvoted 4 times

pentium75 1 year, 11 months ago

"Hadoop helps you turn petabytes of un-structured or semi-structured data into useful insights about your applications or users."

https://aws.amazon.com/emr/features/hadoop/?nc1=h_ls

upvoted 2 times

JKevin778 2 years, 2 months ago

Selected Answer: A

athena for s3

upvoted 1 times

A company has two VPCs that are located in the us-west-2 Region within the same AWS account. The company needs to allow network traffic between these VPCs. Approximately 500 GB of data transfer will occur between the VPCs each month.

What is the MOST cost-effective solution to connect these VPCs?

- A. Implement AWS Transit Gateway to connect the VPCs. Update the route tables of each VPC to use the transit gateway for inter-VPC communication.
- B. Implement an AWS Site-to-Site VPN tunnel between the VPCs. Update the route tables of each VPC to use the VPN tunnel for inter-VPC communication.
- C. Set up a VPC peering connection between the VPCs. Update the route tables of each VPC to use the VPC peering connection for inter-VPC communication. **Most Voted**
- D. Set up a 1 GB AWS Direct Connect connection between the VPCs. Update the route tables of each VPC to use the Direct Connect connection for inter-VPC communication.

Correct Answer: C

Community vote distribution

C (100%)

Comments

luiscc **Highly Voted** 1 year, 10 months ago

Selected Answer: C

C is the correct answer.

VPC peering is the most cost-effective way to connect two VPCs within the same region and AWS account. There are no additional charges for VPC peering beyond standard data transfer rates.

Transit Gateway and VPN add additional hourly and data processing charges that are not necessary for simple VPC peering.

Direct Connect provides dedicated network connectivity, but is overkill for the relatively low inter-VPC data transfer needs described here. It has high fixed costs plus data transfer rates.

For occasional inter-VPC communication of moderate data volumes within the same region and account, VPC peering is the most cost-effective solution. It provides simple private connectivity without transfer charges or network appliances.

upvoted 6 times

Guru4Cloud **Highly Voted** 1 year, 9 months ago

Selected Answer: C

The key reasons are:

VPC peering provides private connectivity between VPCs without using public IP space.

Data transferred between peered VPCs is free as long as they are in the same region.

500 GB/month inter-VPC data transfer fits within peering free tier.

Transit Gateway (Option A) incurs hourly charges plus data transfer fees. More costly than peering.

Site-to-Site VPN (Option B) incurs hourly charges and data transfer fees. More expensive than peering.

Direct Connect (Option D) has high hourly charges and would be overkill for this use case.

upvoted 5 times

TariqKipkemei **Most Recent** 1 year, 6 months ago

Selected Answer: C

A VPC peering connection is a networking connection between two VPCs that enables you to route traffic between them using private IPv4 addresses or IPv6 addresses. Instances in either VPC can communicate with each other as if they are within the same network. You can create a VPC peering connection between your own VPCs, or with a VPC in another AWS account. The VPCs can be in different Regions (also known as an inter-Region VPC peering connection).

<https://docs.aws.amazon.com/vpc/latest/peering/what-is-vpc-peering.html#:~:text=A-,VPC%20peering,-connection%20is%20a>
upvoted 3 times

BrijMohan08 1 year, 9 months ago

Selected Answer: C

Transit Gateway network peering.
VPC Peering to peer 2 or more VPC in the same region.
upvoted 4 times

mrsoa 1 year, 10 months ago

Selected Answer: C

VPC peering is the most cost-effective solution
upvoted 2 times

Deepakin96 1 year, 10 months ago

Selected Answer: C

Communicating with two VPC in same account = VPC Peering
upvoted 2 times

A company hosts multiple applications on AWS for different product lines. The applications use different compute resources, including Amazon EC2 instances and Application Load Balancers. The applications run in different AWS accounts under the same organization in AWS Organizations across multiple AWS Regions. Teams for each product line have tagged each compute resource in the individual accounts.

The company wants more details about the cost for each product line from the consolidated billing feature in Organizations.

Which combination of steps will meet these requirements? (Choose two.)

A. Select a specific AWS generated tag in the AWS Billing console.

B. Select a specific user-defined tag in the AWS Billing console. **Most Voted**

C. Select a specific user-defined tag in the AWS Resource Groups console.

D. Activate the selected tag from each AWS account.

E. Activate the selected tag from the Organizations management account. **Most Voted**

Correct Answer: BE

Community vote distribution

BE (100%)

Comments

Guru4Cloud **Highly Voted** 2 years, 9 months ago

Selected Answer: BE

The reasons are:

User-defined tags were created by each product team to identify resources. Selecting the relevant tag in the Billing console will group costs. The tag must be activated from the Organizations management account to consolidate billing across all accounts.

AWS generated tags are predefined by AWS and won't align to product lines.

Resource Groups (Option C) helps manage resources but not billing.

Activating the tag from each account (Option D) is not needed since Organizations centralizes billing.

upvoted 12 times

adiorsdior999 1 year, 7 months ago

Resource Groups helps grouping, editing, searching tag for resources, and do not use for viewing cost details

upvoted 2 times

potomac **Most Recent** 2 years, 7 months ago

Selected Answer: BE

Your user-defined cost allocation tags represent the tag key, which you activate in the Billing console.

upvoted 2 times

mrsoa 2 years, 10 months ago

Selected Answer: BE

BE BE BE BE

upvoted 3 times

Kiki_Pass 2 years, 10 months ago

Selected Answer: BE

"Only a management account in an organization and single accounts that aren't members of an organization have access to the cost allocation tags manager in the Billing and Cost Management console."

<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/custom-tags.html>

upvoted 4 times

A company's solutions architect is designing an AWS multi-account solution that uses AWS Organizations. The solutions architect has organized the company's accounts into organizational units (OUs).

The solutions architect needs a solution that will identify any changes to the OU hierarchy. The solution also needs to notify the company's operations team of any changes.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Provision the AWS accounts by using AWS Control Tower. Use account drift notifications to identify the changes to the OU hierarchy. **Most Voted**
- B. Provision the AWS accounts by using AWS Control Tower. Use AWS Config aggregated rules to identify the changes to the OU hierarchy.
- C. Use AWS Service Catalog to create accounts in Organizations. Use an AWS CloudTrail organization trail to identify the changes to the OU hierarchy.
- D. Use AWS CloudFormation templates to create accounts in Organizations. Use the drift detection operation on a stack to identify the changes to the OU hierarchy.

Correct Answer: A

Community vote distribution

A (61%)

B (22%)

C (17%)

Comments

Guru4Cloud **Highly Voted** 2 years, 9 months ago

Selected Answer: A

The key advantages you highlight of Control Tower are convincing:

Fully managed service simplifies multi-account setup.
Built-in account drift notifications detect OU changes automatically.
More scalable and less complex than Config rules or CloudTrail.
Better security and compliance guardrails than custom options.
Lower operational overhead compared to other solution
upvoted 12 times

Bmaster **Highly Voted** 2 years, 10 months ago

A is correct.

<https://docs.aws.amazon.com/controltower/latest/userguide/what-is-control-tower.html>
<https://docs.aws.amazon.com/controltower/latest/userguide/prevention-and-notification.html>
upvoted 8 times

In2718 **Most Recent** 9 months, 3 weeks ago

Selected Answer: C

A. Control Tower drift detection is about accounts or resources drifting from guardrails/baselines, not OU hierarchy changes.
B. Control Tower + AWS Config aggregated rules
AWS Config can track certain resources, but AWS Config does not natively record OU hierarchy changes in AWS Organizations.
upvoted 1 times

In2718 9 months, 3 weeks ago

Selected Answer: C

Chat ---

C. Service Catalog + CloudTrail organization trail

CloudTrail does log AWS Organizations events such as CreateOrganizationalUnit, MoveAccount, UpdateOrganizationalUnit, DeleteOrganizationalUnit

Setting up an organization trail ensures these events are captured across all accounts automatically. Then you can use EventBridge rules to trigger notifications (e.g., to SNS → Ops team). This directly meets the requirement with the least overhead.

upvoted 1 times

111380c 10 months ago

Selected Answer: B

A for checking account drift (account configuration & blueprint)

upvoted 1 times

1166ae3 1 year, 11 months ago

Selected Answer: C

Create Accounts using AWS Service Catalog:

Utilize AWS Service Catalog to provision AWS accounts within AWS Organizations. This ensures standardized account creation and management. Enable AWS CloudTrail Organization Trail:

Set up an AWS CloudTrail organization trail that records all API calls across all accounts in the organization. This trail will capture changes to the OU hierarchy, including any modifications to organizational units.

upvoted 2 times

chickenmf 2 years, 3 months ago

Selected Answer: B

AWS Config helps you maintain a detailed inventory of your resources and their configurations, track changes over time, and ensure compliance with your organization's policies and industry regulations.

upvoted 4 times

chickenmf 2 years, 3 months ago

Furthermore, AWS Config Aggregated Rules are a feature within AWS Config that enables you to evaluate compliance with desired configurations or compliance standards across multiple AWS accounts and regions. They are particularly useful in scenarios where you want to enforce consistent rules and compliance checks across an entire organization with multiple AWS accounts.

upvoted 1 times

chickenmf 2 years, 3 months ago

NVM - This is such a stupid question lol changing my answer to A due to the following:

Account drift notifications in AWS are a feature provided by AWS Control Tower. These notifications help organizations identify and respond to changes made to an AWS account that deviate from the established baseline configuration created during the initial setup by AWS Control Tower. Drift refers to any configuration changes that have been made to an AWS account after it was provisioned by Control Tower.

upvoted 6 times

Avyay 2 years, 3 months ago

This was in my exam today..I selected Answer A

upvoted 3 times

chickenmf 2 years, 3 months ago

what percentage of all these questions would you say were in the exam?

upvoted 2 times

wizcloudifa 2 years, 1 month ago

I read in one of the earlier questions, its around 75%, someone who gave the exam said so

upvoted 2 times

SHAAHIBHUSHANAWS 2 years, 6 months ago

A

<https://docs.aws.amazon.com/controltower/latest/userguide/drift.html>

upvoted 2 times

potomac 2 years, 7 months ago

Selected Answer: A

AWS Control Tower provides passive and active methods of drift monitoring protection for preventive controls.

upvoted 2 times

darekw 2 years, 9 months ago

<https://docs.aws.amazon.com/controltower/latest/userguide/prevention-and-notification.html>

upvoted 2 times

A company's website handles millions of requests each day, and the number of requests continues to increase. A solutions architect needs to improve the response time of the web application. The solutions architect determines that the application needs to decrease latency when retrieving product details from the Amazon DynamoDB table.

Which solution will meet these requirements with the LEAST amount of operational overhead?

- A. Set up a DynamoDB Accelerator (DAX) cluster. Route all read requests through DAX. **Most Voted**
- B. Set up Amazon ElastiCache for Redis between the DynamoDB table and the web application. Route all read requests through Redis.
- C. Set up Amazon ElastiCache for Memcached between the DynamoDB table and the web application. Route all read requests through Memcached.
- D. Set up Amazon DynamoDB Streams on the table, and have AWS Lambda read from the table and populate Amazon ElastiCache. Route all read requests through ElastiCache.

Correct Answer: A

Community vote distribution

A (100%)

Comments

mrsoa **Highly Voted** 2 years, 4 months ago

Selected Answer: A

A, because B,C and D contains ElastiCache which required a heavy code changes, so more operational overhead
upvoted 9 times

TariqKipkemei **Highly Voted** 2 years ago

Selected Answer: A

decrease latency when retrieving product details from the Amazon DynamoDB = Amazon DynamoDB Accelerator (DAX)
upvoted 5 times

sandordini **Most Recent** 1 year, 7 months ago

Selected Answer: A

DAX to reduce latency
upvoted 3 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: A

The key reasons:

DAX provides a DynamoDB-compatible caching layer to reduce read latency. It is purpose-built for accelerating DynamoDB workloads. Using DAX requires minimal application changes - only read requests are routed through it. DAX handles caching logic automatically without needing complex integration code. ElastiCache Redis/Memcached (Options B/C) require more integration work to sync DynamoDB data. Using Lambda and Streams to populate ElastiCache (Option D) is a complex event-driven approach requiring ongoing maintenance. DAX plugs in seamlessly to accelerate DynamoDB with very little operational overhead
upvoted 3 times

Deepakin96 2 years, 4 months ago

Selected Answer: A

DynamoDB = DAX
upvoted 3 times

Bmaster 2 years, 4 months ago

only A
upvoted 3 times


A solutions architect needs to ensure that API calls to Amazon DynamoDB from Amazon EC2 instances in a VPC do not travel across the internet.

Which combination of steps should the solutions architect take to meet this requirement? (Choose two.)

- A. Create a route table entry for the endpoint. **Most Voted**
- B. Create a gateway endpoint for DynamoDB. **Most Voted**
- C. Create an interface endpoint for Amazon EC2.
- D. Create an elastic network interface for the endpoint in each of the subnets of the VPC.
- E. Create a security group entry in the endpoint's security group to provide access.

Correct Answer: AB

Community vote distribution



Comments

ukivanlamipi **Highly Voted** 2 years, 10 months ago

Selected Answer: AB

<https://docs.aws.amazon.com/vpc/latest/privatelink/vpc-endpoints-ddb.html>

upvoted 13 times

Guru4Cloud **Highly Voted** 2 years, 9 months ago

Selected Answer: BE

The reasons are:

A gateway endpoint for DynamoDB enables private connectivity between DynamoDB and the VPC. This allows EC2 instances to access DynamoDB APIs without traversing the internet.

A security group entry is needed to allow the EC2 instances access to the DynamoDB endpoint over the VPC.

An interface endpoint is used for services like S3 and Systems Manager, not DynamoDB.

Route table entries route traffic within a VPC but do not affect external connectivity.

Elastic network interfaces are not needed for gateway endpoints.

upvoted 10 times

allanm 10 months ago

Gateway VPC endpoints provide reliable connectivity to Amazon S3 and DynamoDB without requiring an internet gateway or a NAT device for your VPC. Gateway endpoints do not use AWS PrivateLink, unlike other types of VPC endpoints.

<https://docs.aws.amazon.com/vpc/latest/privatelink/gateway-endpoints.html>

B is correct

upvoted 2 times

unbendable 2 years, 7 months ago

"The outbound rules for the security group for instances that access DynamoDB through the gateway endpoint must allow traffic to DynamoDB",

<https://docs.aws.amazon.com/vpc/latest/privatelink/vpc-endpoints-ddb.html>

The option however is talking about the security group of the endpoint

upvoted 3 times

JoeTromundo **Most Recent** 1 year, 8 months ago

Selected Answer: AB

A & B are correct

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/vpc-endpoints-dynamodb.html>

E is incorrect. There's no need for security group.

From the URL above:

"Once the VPC subnet's gateway endpoint has been granted access to DynamoDB, any AWS account with access to that subnet can use DynamoDB."

upvoted 2 times

a7md0 1 year, 11 months ago

Selected Answer: AB

Creating the gateway endpoint and edit the route table is enough, there are no security group involved

upvoted 2 times

osmk 2 years, 4 months ago

AB <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/vpc-endpoints-dynamodb.html>

upvoted 3 times

upliftinghut 2 years, 4 months ago

Selected Answer: AB

C & D are both not relevant. D looks ok but DynamoDB doesn't go with security group, it only allows route table for VPC endpoint. Link here:

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/vpc-endpoints-dynamodb.html>

upvoted 2 times

upliftinghut 2 years, 4 months ago

Sorry, E not D. E looks ok but DynamoDB doesn't go with security group, it only allows route table for VPC endpoint. Link here:

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/vpc-endpoints-dynamodb.html>

upvoted 2 times

awsgeek75 2 years, 4 months ago

Selected Answer: AB

DynamoDB can only be connected via Gateway endpoint (just like S3)

route table for connecting the VPC to the endpoint

So do B then A

C: interface endpoint for EC2 to what?

D: ENI not applicable here for VPC

E: Incomplete option as to access to what?

upvoted 4 times

theonlyhero 2 years, 5 months ago

go through this video it will show the answer is AB

<https://www.youtube.com/watch?v=8FTnyhklEvU>

upvoted 5 times

pentium75 2 years, 5 months ago

Selected Answer: AB

Gateway Endpoint does not have an ENI, thus it has no security group. Instances have security groups and those must allow access to DynamoDB.

upvoted 6 times

aws94 2 years, 6 months ago

Selected Answer: BE

A. Create a route table entry for the endpoint: This is not necessary, as the gateway endpoint itself automatically creates the required route table entries.

upvoted 2 times

TariqKipkemei 2 years, 6 months ago

Selected Answer: AB

Create a gateway endpoint for DynamoDB then create a route table entry for the endpoint

upvoted 3 times

EdenWang 2 years, 6 months ago

Selected Answer: BE

refer to question 555

upvoted 2 times

cciesam 2 years, 7 months ago

Selected Answer: AB

<https://docs.aws.amazon.com/vpc/latest/privatelink/gateway-endpoints.html#vpc-endpoints-routing>

Traffic from your VPC to Amazon S3 or DynamoDB is routed to the gateway endpoint. Each subnet route table must have a route that sends traffic destined for the service to the gateway endpoint using the prefix list for the service.

upvoted 2 times

potomac 2 years, 7 months ago

Selected Answer: AB

You can access Amazon DynamoDB from your VPC using gateway VPC endpoints. After you create the gateway endpoint, you can add it as a target in your route table for traffic destined from your VPC to DynamoDB.

upvoted 3 times

danielmakita 2 years, 7 months ago

It is A and B. Not E because security group does not span VPCs.

upvoted 3 times

iwannabeawsgod 2 years, 7 months ago

Selected Answer: AB

A and B for sure

upvoted 4 times

loveaws 2 years, 8 months ago

B and D.

upvoted 1 times


A company runs its applications on both Amazon Elastic Kubernetes Service (Amazon EKS) clusters and on-premises Kubernetes clusters. The company wants to view all clusters and workloads from a central location.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use Amazon CloudWatch Container Insights to collect and group the cluster information.
- B. Use Amazon EKS Connector to register and connect all Kubernetes clusters. **Most Voted**
- C. Use AWS Systems Manager to collect and view the cluster information.
- D. Use Amazon EKS Anywhere as the primary cluster to view the other clusters with native Kubernetes commands.

Correct Answer: B

Community vote distribution



Comments

ErnShm **Highly Voted** 2 years, 9 months ago
B

You can use Amazon EKS Connector to register and connect any conformant Kubernetes cluster to AWS and visualize it in the Amazon EKS console. After a cluster is connected, you can see the status, configuration, and workloads for that cluster in the Amazon EKS console. You can use this feature to view connected clusters in Amazon EKS console, but you can't manage them. The Amazon EKS Connector requires an agent that is an open source project on Github. For additional technical content, including frequently asked questions and troubleshooting, see Troubleshooting issues in Amazon EKS Connector

The Amazon EKS Connector can connect the following types of Kubernetes clusters to Amazon EKS.

On-premises Kubernetes clusters

Self-managed clusters that are running on Amazon EC2

Managed clusters from other cloud providers
upvoted 7 times

mrsoa **Highly Voted** 2 years, 10 months ago
seems B

<https://docs.aws.amazon.com/eks/latest/userguide/eks-connector.html>
upvoted 5 times

Danilus **Most Recent** 1 year, 7 months ago

Selected Answer: B

KEY-he company wants to view all clusters and workloads from a central location.

the answer is not D because amazon EKS anywhere allow you to run cluster of kubernetes in on-premises or in other clouds
the answer is B EKS connector helps u to connect kubernetes clusters on premises or other cloud and manage from a single ublication
upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: B

<https://docs.aws.amazon.com/eks/latest/userguide/eks-connector.html>

"You can use Amazon EKS Connector to register and connect any conformant Kubernetes cluster to AWS and visualize it in the Amazon EKS console.

"

B is the right product for this.
upvoted 4 times

pentium75 2 years, 5 months ago

Selected Answer: B

EKS Connector -> 'view clusters and workloads' as requested
EKS Anywhere -> create and manage on-premises EKS clusters
upvoted 4 times

SHAAHIBHUSHANAWS 2 years, 6 months ago

B
EKS connector helps to integrate multiple cluster with EKS console. EKS anywhere is Kubernetes Ditro cluster to be deployed on-prem. It is not for integrating with other cluster.
upvoted 1 times

TariqKipkemei 2 years, 6 months ago

Selected Answer: B

View all clusters and workloads (incl on-prem) from a central location = Amazon EKS Connector
Create and operate Kubernetes clusters on your own infrastructure = Amazon EKS Anywhere

<https://aws.amazon.com/eks/eks-anywhere/#:~:text=Amazon-,EKS%20Anywhere,-lets%20you%20create>

<https://docs.aws.amazon.com/eks/latest/userguide/eks-connector.html#:~:text=You%20can%20use-,Amazon%20EKS%20Connector,-to%20register%20and>
upvoted 2 times

potomac 2 years, 7 months ago

Selected Answer: B

It is B
upvoted 2 times

thainguyensunya 2 years, 9 months ago

Selected Answer: B

Definitely B.
"You can use Amazon EKS Connector to register and connect any conformant Kubernetes cluster to AWS and visualize it in the Amazon EKS console. After a cluster is connected, you can see the status, configuration, and workloads for that cluster in the Amazon EKS console."
<https://docs.aws.amazon.com/eks/latest/userguide/eks-connector.html>
upvoted 3 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: B

The key reasons:

EKS Connector allows registering external Kubernetes clusters (on-premises and otherwise) with Amazon EKS
This provides a unified view and management of all clusters within the EKS console.
EKS Connector handles keeping resources in sync across connected clusters.
This centralized approach minimizes operational overhead compared to using separate tools.
CloudWatch Container Insights (Option A) only provides metrics and logs, not cluster management.
Systems Manager (Option C) is more general purpose and does not natively integrate with EKS.
EKS Anywhere (Option D) would not provide a single pane of glass for external clusters.
upvoted 4 times

RealMarcus 2 years, 10 months ago

Amazon EKS Connector enables you to create and manage a centralized view of all your Kubernetes clusters, regardless of whether they are Amazon EKS clusters or on-premises Kubernetes clusters. It allows you to register these clusters with your Amazon EKS control plane, providing a unified management interface for all clusters.
upvoted 2 times

avkya 2 years, 10 months ago

Selected Answer: B

You can use Amazon EKS Connector to register and connect any conformant Kubernetes cluster to AWS and visualize it in the Amazon EKS console. After a cluster is connected, you can see the status, configuration, and workloads for that cluster in the Amazon EKS console. You can use this feature to view connected clusters in Amazon EKS console, but you can't manage them
upvoted 1 times

ukivanlamlpi 2 years, 10 months ago

Selected Answer: D

only D can connect to on-perm
upvoted 1 times

pentium75 2 years, 5 months ago

No.

"The Amazon EKS Connector can connect the following types of Kubernetes clusters to Amazon EKS.

On-premises Kubernetes clusters"

<https://aws.amazon.com/de/eks/eks-anywhere/>

upvoted 2 times

pentium75 2 years, 5 months ago

Wrong link, statement is from <https://docs.aws.amazon.com/eks/latest/userguide/eks-connector.html>

upvoted 2 times

Bmaster 2 years, 10 months ago

Only B

<https://docs.aws.amazon.com/eks/latest/userguide/eks-connector.html>

upvoted 3 times

A company is building an ecommerce application and needs to store sensitive customer information. The company needs to give customers the ability to complete purchase transactions on the website. The company also needs to ensure that sensitive customer data is protected, even from database administrators.

Which solution meets these requirements?

- A. Store sensitive data in an Amazon Elastic Block Store (Amazon EBS) volume. Use EBS encryption to encrypt the data. Use an IAM instance role to restrict access.
- B. Store sensitive data in Amazon RDS for MySQL. Use AWS Key Management Service (AWS KMS) client-side encryption to encrypt the data. **Most Voted**
- C. Store sensitive data in Amazon S3. Use AWS Key Management Service (AWS KMS) server-side encryption to encrypt the data. Use S3 bucket policies to restrict access.
- D. Store sensitive data in Amazon FSx for Windows Server. Mount the file share on application servers. Use Windows file permissions to restrict access.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: B

The key reasons are:

RDS MySQL provides a fully managed database service well suited for an ecommerce application. AWS KMS client-side encryption allows encrypting sensitive data before it hits the database. The data remains encrypted at rest. This protects sensitive customer data from database admins and privileged users. EBS encryption (Option A) protects data at rest but not in use. IAM roles don't prevent admin access. S3 (Option C) encrypts data at rest on the server side. Bucket policies don't restrict admin access. FSx file permissions (Option D) don't prevent admin access to unencrypted data.

upvoted 12 times

pentium75 **Highly Voted** 1 year, 11 months ago

Selected Answer: B

A, C and D would allow the administrator of the storage to access the data. Besides, it is data about "purchase transactions" which is usually stored in a transactional database (such as RDS for MySQL), not in a file or object storage.

upvoted 8 times

sidharthwader 1 year, 7 months ago

Good thought purchase transactions I had missed that part

upvoted 2 times

SHAAHIBHUSHANAWS **Most Recent** 2 years ago

B

I want to go with B as question is for database administrator. Also client key encryption is possible in code and KMS can be used for encryption but not using KMS keys. Encrypted data available in DB is of no use to DB admin.

upvoted 2 times

riyasara 2 years ago

Answer is option C. option B is not ideal because Amazon RDS for MySQL is a relational database service that is optimized for structured data, not for storing sensitive customer information. Moreover, by using client-side encryption with AWS KMS, you need to encrypt and decrypt the data in your application code, which increases the risk of exposing your data in transit or at rest. You also need to manage the encryption keys yourself, which adds complexity and overhead to your application.

upvoted 2 times

pentium75 1 year, 11 months ago

"optimized for structured data, not for storing sensitive customer information" ... Data related to "purchase transactions" is usually structured; that it contains "sensitive customer information" doesn't change the structured nature.

upvoted 4 times

awsgeek75 1 year, 11 months ago

eCommerce data and transaction data are ideal for RDS which, when encrypted, is secure even from the DBA.

upvoted 2 times

wsdasdasdqwdaw 2 years, 1 month ago

I would go for B, because RDS (database admins), but I would like to see as well encryption at rest as well, not only in transit.

upvoted 2 times

mrsoa 2 years, 4 months ago

Selected Answer: B

Using client-side encryption we can protect specific fields and guarantee only decryption if the client has access to an API key, we can protect specific fields even from database admins

upvoted 3 times

D10SJoker 2 years, 4 months ago

Selected Answer: B

For me it's B because of "client-side encryption to encrypt the data"

upvoted 2 times

h8er 2 years, 4 months ago

keyword - database administrators

upvoted 5 times

allanm 10 months ago

False correlation. Just because it says "database administrators" doesn't mean RDS is the answer to this.

upvoted 1 times

Kiki_Pass 2 years, 4 months ago

Selected Answer: B

"even from database administrators" -> "Client Side encryption"

upvoted 4 times

Bmaster 2 years, 4 months ago

My choice is B

upvoted 4 times

A company has an on-premises MySQL database that handles transactional data. The company is migrating the database to the AWS Cloud. The migrated database must maintain compatibility with the company's applications that use the database. The migrated database also must scale automatically during periods of increased demand.

Which migration solution will meet these requirements?

- A. Use native MySQL tools to migrate the database to Amazon RDS for MySQL. Configure elastic storage scaling.
- B. Migrate the database to Amazon Redshift by using the mysqldump utility. Turn on Auto Scaling for the Amazon Redshift cluster.
- C. Use AWS Database Migration Service (AWS DMS) to migrate the database to Amazon Aurora. Turn on Aurora Auto Scaling. **Most Voted**
- D. Use AWS Database Migration Service (AWS DMS) to migrate the database to Amazon DynamoDB. Configure an Auto Scaling policy.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: C

The key reasons are:

DMS provides an easy migration path from MySQL to Aurora while minimizing downtime.
Aurora is a MySQL-compatible relational database service that will maintain compatibility with the company's applications.
Aurora Auto Scaling allows the database to automatically scale up and down based on demand to handle increased workloads.
RDS MySQL (Option A) does not scale as well as the Aurora architecture.
Redshift (Option B) is for analytics, not transactional data, and may not be compatible.
DynamoDB (Option D) is a NoSQL datastore and lacks MySQL compatibility.

upvoted 9 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: C

A is wrong as you cannot use native MySQL tools for migration. Happy to be corrected though!
B Redshift is not compatible with MySQL
D is DynamoDB
C Aurora MySQL is compatible and supports auto scaling

upvoted 3 times

TariqKipkemei 2 years ago

Selected Answer: C

on-premises MySQL database, transactional data, maintain compatibility, scale automatically = Amazon Aurora
migrating the database to the AWS Cloud = AWS Database Migration Service

upvoted 2 times

potomac 2 years, 1 month ago

Selected Answer: C

Aurora is a MySQL-compatible relational database service

upvoted 2 times

mrsoa 2 years, 4 months ago

Selected Answer: C

Aurora is better in autoscaling then RDS

upvoted 2 times

Bmaster 2 years, 4 months ago

C is correct

A is incorrect. RDS for MySQL does not scale automatically during periods of increased demand.

B is incorrect. Redshift is used for data sharing purposes.

D is incorrect. you muse change application codes.

upvoted 2 times

Eminenza22 2 years, 4 months ago

Amazon RDS now supports Storage Auto Scaling

upvoted 2 times

A company runs multiple Amazon EC2 Linux instances in a VPC across two Availability Zones. The instances host applications that use a hierarchical directory structure. The applications need to read and write rapidly and concurrently to shared storage.

What should a solutions architect do to meet these requirements?

- A. Create an Amazon S3 bucket. Allow access from all the EC2 instances in the VPC.
- B. Create an Amazon Elastic File System (Amazon EFS) file system. Mount the EFS file system from each EC2 instance.
Most Voted
- C. Create a file system on a Provisioned IOPS SSD (io2) Amazon Elastic Block Store (Amazon EBS) volume. Attach the EBS volume to all the EC2 instances.
- D. Create file systems on Amazon Elastic Block Store (Amazon EBS) volumes that are attached to each EC2 instance. Synchronize the EBS volumes across the different EC2 instances.

Correct Answer: B

Community vote distribution

B (96%) 4%

Comments

Josantru **Highly Voted** 2 years, 10 months ago
Correct B.

How is Amazon EFS different than Amazon S3?

Amazon EFS provides shared access to data using a traditional file sharing permissions model and hierarchical directory structure via the NFSv4 protocol. Applications that access data using a standard file system interface provided through the operating system can use Amazon EFS to take advantage of the scalability and reliability of file storage in the cloud without writing any new code or adjusting applications.

Amazon S3 is an object storage platform that uses a simple API for storing and accessing data. Applications that do not require a file system structure and are designed to work with object storage can use Amazon S3 as a massively scalable, durable, low-cost object storage solution.

upvoted 13 times

Gape4 **Most Recent** 1 year, 11 months ago

Selected Answer: B

<https://aws.amazon.com/efs/when-to-choose-efs/>

upvoted 2 times

zinabu 2 years, 2 months ago

hierarchical structure =EFS

upvoted 3 times

pentium75 2 years, 5 months ago

Selected Answer: B

Not A because S3 does not allow a "hierarchical directory structure"

Not C because Multi-attach does not work "across two Availability Zones"

Not D because we need "shared", not synchronized, storage.

upvoted 4 times

TariqKipkemei 2 years, 6 months ago

Selected Answer: B

hierarchical directory structure, read and write rapidly and concurrently to shared storage = Amazon Elastic File System

upvoted 2 times

zits88 1 year, 10 months ago

Best and most concise explanation here. All true.

upvoted 1 times

potomac 2 years, 7 months ago

Selected Answer: B

Amazon EFS simultaneously supports on-premises servers using a traditional file permissions model, file locking, and hierarchical directory structure through the NFS v4 protocol.

upvoted 2 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: B

The key reasons:

EFS provides a scalable, high performance NFS file system that can be concurrently accessed from multiple EC2 instances.

It supports the hierarchical directory structure needed by the applications.

EFS is elastic, growing and shrinking automatically as needed.

It can be accessed from instances across AZs, meeting the shared storage requirement.

S3 object storage (option A) lacks the file system semantics needed by the apps.

EBS volumes (options C and D) are attached to a single instance and would require replication and syncing to share across instances.

EFS is purpose-built for this use case of a shared file system across Linux instances and aligns best with the performance, concurrency, and availability needs.

upvoted 4 times

barracouto 2 years, 10 months ago

Selected Answer: B

Going with b

upvoted 2 times

Bennyboy789 2 years, 10 months ago

Selected Answer: B

C and D involve using Amazon EBS volumes, which are block storage. While they can be attached to EC2 instances, they might not provide the same level of shared concurrent access as Amazon EFS. Additionally, synchronizing EBS volumes across different EC2 instances (as in option D) can be complex and error-prone.

Therefore, for a scenario where multiple EC2 instances need to rapidly and concurrently access shared storage with a hierarchical directory structure, Amazon EFS is the best solution.

upvoted 3 times

ukivanlamlpi 2 years, 10 months ago

Selected Answer: B

s3 is flat structure. EBS multi mount only for same available zone

upvoted 2 times

Dana12345 2 years, 10 months ago

Selected Answer: B

Because Amazon EBS Multi-Attach enables you to attach a single Provisioned IOPS SSD (io1 or io2) volume to multiple instances that are in the same Availability Zone. The infra contains 2 AZ's.

upvoted 2 times

mrsoa 2 years, 10 months ago

Selected Answer: B

B is the correct answer

<https://docs.aws.amazon.com/efs/latest/ug/whatisefs.html>

upvoted 2 times

mrsoa 2 years, 10 months ago

B is the correct answer

<https://docs.aws.amazon.com/efs/latest/ug/whatisefs.html>

upvoted 2 times

RazSteel 2 years, 10 months ago

Selected Answer: C

I think that C is the best option coz io2 can share storage and multi attach.

upvoted 1 times

PLN6302 2 years, 9 months ago

hierarchial directory structure is present in EFS

upvoted 3 times

pentium75 2 years, 5 months ago

Multi-attach does not work "across two Availability Zones".

upvoted 3 times

A solutions architect is designing a workload that will store hourly energy consumption by business tenants in a building. The sensors will feed a database through HTTP requests that will add up usage for each tenant. The solutions architect must use managed services when possible. The workload will receive more features in the future as the solutions architect adds independent components.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use Amazon API Gateway with AWS Lambda functions to receive the data from the sensors, process the data, and store the data in an Amazon DynamoDB table. **Most Voted**
- B. Use an Elastic Load Balancer that is supported by an Auto Scaling group of Amazon EC2 instances to receive and process the data from the sensors. Use an Amazon S3 bucket to store the processed data.
- C. Use Amazon API Gateway with AWS Lambda functions to receive the data from the sensors, process the data, and store the data in a Microsoft SQL Server Express database on an Amazon EC2 instance.
- D. Use an Elastic Load Balancer that is supported by an Auto Scaling group of Amazon EC2 instances to receive and process the data from the sensors. Use an Amazon Elastic File System (Amazon EFS) shared file system to store the processed data.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Guru4Cloud **Highly Voted** 2 years, 9 months ago

Selected Answer: A

The key reasons are:

- ° API Gateway removes the need to manage servers to receive the HTTP requests from sensors
- ° Lambda functions provide a serverless compute layer to process data as needed
- ° DynamoDB is a fully managed NoSQL database that scales automatically
- ° This serverless architecture has minimal operational overhead to manage
- ° Options B, C, and D all require managing EC2 instances which increases ops workload
- ° Option C also adds SQL Server admin tasks and licensing costs
- ° Option D uses EFS file storage which requires capacity planning and management

upvoted 6 times

mrsoa **Highly Voted** 2 years, 10 months ago

Selected Answer: A

A seems to be the right answer

upvoted 5 times

AWSSURI **Most Recent** 1 year, 9 months ago

Options B,C,D involves unwanted operational overheads due to EC2

So A is the right answer

upvoted 2 times

Mikado211 2 years, 6 months ago

Thinking of that, there is not many questions about IoT Core, but this product could be an excellent answer for the need.

upvoted 2 times

TariqKipkemei 2 years, 6 months ago

Selected Answer: A

Workload runs every hour, must use managed services, more features in the future, LEAST operational overhead = AWS Lambda functions.
HTTP requests, must use managed services, more features in the future, LEAST operational overhead = API Gateway.

Must use managed services, more features in the future, LEAST operational overhead =Amazon DynamoDB.

upvoted 4 times

ersin13 2 years, 10 months ago

key word is "must use managed services when possible" api ,lambda dynamodb are serverless. so answer is A

upvoted 2 times

Kiki_Pass 2 years, 10 months ago

Selected Answer: A

"The workload will receive more features in the future ..." -> DynamoDB

upvoted 4 times

Bmaster 2 years, 10 months ago

A is correct.

upvoted 3 times

A solutions architect is designing the storage architecture for a new web application used for storing and viewing engineering drawings. All application components will be deployed on the AWS infrastructure.

The application design must support caching to minimize the amount of time that users wait for the engineering drawings to load. The application must be able to store petabytes of data.

Which combination of storage and caching should the solutions architect use?

- A. Amazon S3 with Amazon CloudFront **Most Voted**
- B. Amazon S3 Glacier with Amazon ElastiCache
- C. Amazon Elastic Block Store (Amazon EBS) volumes with Amazon CloudFront
- D. AWS Storage Gateway with Amazon ElastiCache

Correct Answer: A

Community vote distribution

A (100%)

Comments

Guru4Cloud **Highly Voted** 2 years, 9 months ago

Selected Answer: A

The key reasons are:

S3 provides highly durable and scalable object storage capable of handling petabytes of data cost-effectively. CloudFront can be used to cache S3 content at the edge, minimizing latency for users and speeding up access to the engineering drawings. The global CloudFront edge network is ideal for caching large amounts of static media like drawings. EBS provides block storage but lacks the scale and durability of S3 for large media files. Glacier is cheaper archival storage but has higher latency unsuited for frequent access. Storage Gateway and ElastiCache may play a role but do not align as well to the main requirements.

upvoted 6 times

Bmaster **Highly Voted** 2 years, 10 months ago

A is correct

upvoted 5 times

Gape4 **Most Recent** 1 year, 11 months ago

Selected Answer: A

I think the answer is A.

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: A

Petabyte data on AWS infra with high performance

B is Glacier so slow

C EBS for petabyte data doesn't work

D Storage gateway is for on premise connectivity which is not required

upvoted 4 times

TariqKipkemei 2 years, 6 months ago

Selected Answer: A

Storing and viewing engineering drawings = Amazon S3

Support caching to minimize the amount of time that users wait for the engineering drawings to load = Amazon CloudFront

upvoted 2 times

wsdasdasdqwda 2 years, 7 months ago

CF caching and S3 supports petabytes data

upvoted 3 times

lemur88 2 years, 9 months ago

Selected Answer: A

CF allows caching

upvoted 2 times

mrsoa 2 years, 10 months ago

Selected Answer: A

The answer seems A:

B : Glacier for archiving

C : i dont think EBS scale to petabytes (I am not sure about that)

D : it incorrect becasseAll application components will be deployed on the AWS infrastruc

upvoted 3 times


An Amazon EventBridge rule targets a third-party API. The third-party API has not received any incoming traffic. A solutions architect needs to determine whether the rule conditions are being met and if the rule's target is being invoked.

Which solution will meet these requirements?

- A. Check for metrics in Amazon CloudWatch in the namespace for AWS/Events. **Most Voted**
- B. Review events in the Amazon Simple Queue Service (Amazon SQS) dead-letter queue.
- C. Check for the events in Amazon CloudWatch Logs.
- D. Check the trails in AWS CloudTrail for the EventBridge events.

Correct Answer: A

Community vote distribution



Comments

lemur88 **Highly Voted** 2 years, 3 months ago

Selected Answer: A

<https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-monitoring.html>

upvoted 8 times

awsgeek75 **Highly Voted** 1 year, 11 months ago

Selected Answer: A

"EventBridge sends metrics to Amazon CloudWatch every minute for everything from the number of matched events to the number of times a target is invoked by a rule."

from <https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-monitoring.html>

B: SQS, irrelevant

C: 'Check for events', this wording is confusing but could mean something in wrong context. I would have chosen C if A wasn't an option

D: CloudTrail is for AWS resource monitoring so irrelevant

upvoted 7 times

pentium75 **Most Recent** 1 year, 11 months ago

Selected Answer: A

A per <https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-monitoring.html>

Not B because SQS is not even involved here

Not C because EventBridge sends only metrics, not detailed logs, to CloudWatch

Not D, many fall for CloudTrail supposedly recording "API calls", but this is about calls for the EventBridge API to AWS, not calls to 3rd party APIs by EventBridge.

upvoted 6 times

Min_93 1 year, 11 months ago

Selected Answer: C

Option A, "Check for metrics in Amazon CloudWatch in the namespace for AWS/Events," primarily provides aggregated metrics related to EventBridge, but it may not give detailed information about individual events or their specific content. Metrics in CloudWatch can give you an overview of how many events are being processed, but for detailed inspection of events and their conditions, checking CloudWatch Logs (option C) is more appropriate.

CloudWatch Logs allow you to see the actual event data and details, providing a more granular view that is useful for troubleshooting and understanding the specifics of why a third-party API is not receiving incoming traffic.

upvoted 1 times

SHAAHIBHUSHANAWS 2 years ago

A

Events not generating logs in cloudwatch and cloudtrail. only metric data is available.

upvoted 2 times

TariqKipkemei 2 years ago

Selected Answer: D

CloudWatch is a monitoring service for AWS resources and applications. CloudTrail is a web service that records API activity in your AWS account. CloudWatch monitors applications and infrastructure performance in the AWS environment. CloudTrail monitors actions in the AWS environment.

upvoted 1 times

pentium75 1 year, 11 months ago

"API activity", referring to AWS APIs. This would record if someone modifies the EventBridge configuration.

upvoted 2 times

EdenWang 2 years ago

Selected Answer: C

C should be correct, I check in AWS management console.

upvoted 1 times

potomac 2 years, 1 month ago

Selected Answer: A

should be A

upvoted 2 times

ibu007 2 years, 3 months ago

Selected Answer: D

Check the trails in AWS CloudTrail for the EventBridge events.

upvoted 1 times

pentium75 1 year, 11 months ago

I think CloudTrail captures management events (such as modifying the EventBridge configuration)

upvoted 2 times

Eminenza22 2 years, 3 months ago

Selected Answer: C

Amazon CloudWatch Logs is a service that collects and stores logs from Amazon Web Services (AWS) resources. These logs can be used to troubleshoot problems, monitor performance, and audit activity.
The other options are incorrect:

Option A: CloudWatch metrics are used to track the performance of AWS resources. They are not used to store events.

Option B: Amazon SQS dead-letter queues are used to store messages that cannot be delivered to their intended recipients. They are not used to store events.

Option D: AWS CloudTrail is a service that records AWS API calls. It can be used to track the activity of EventBridge rules, but it does not store the events themselves.

upvoted 2 times

Eminenza22 2 years, 3 months ago

Errata Corrigé

A

EventBridge sends metrics to Amazon CloudWatch every minute for everything from the number of matched events to the number of times a target is invoked by a rule.

<https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-monitoring.html>

upvoted 2 times

Eminenza22 2 years, 3 months ago

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/events/CloudWatch-Events-Monitoring-CloudWatch-Metrics.html>

upvoted 2 times

jayce5 2 years, 3 months ago

Selected Answer: D

The answer is D:

"CloudTrail captures API calls made by or on behalf of your AWS account from the EventBridge console and to EventBridge API operations."

(<https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-logging-monitoring.html>)

upvoted 2 times

pentium75 1 year, 11 months ago

"API calls" to AWS for managing EventBridge. Not "API calls" BY EventBridge to 3rd party APIs.

upvoted 1 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: D

The key reasons:

AWS CloudTrail provides visibility into EventBridge operations by logging API calls made by EventBridge.

Checking the CloudTrail trails will show the PutEvents API calls made when EventBridge rules match an event pattern.

CloudTrail will also log the Invoke API call when the rule target is triggered.

CloudWatch metrics and logs contain runtime performance data but not info on rule evaluation and targeting.

SQS dead letter queues collect failed event deliveries but won't provide insights on successful invocations.

CloudTrail is purpose-built to log operational events and API activity so it can confirm if the EventBridge rule is being evaluated and triggering the target as expected.

upvoted 2 times

Eminenza22 2 years, 3 months ago

Amazon CloudWatch Logs is a service that collects and stores logs from Amazon Web Services (AWS) resources. These logs can be used to troubleshoot problems, monitor performance, and audit activity.

The other options are incorrect:

Option A: CloudWatch metrics are used to track the performance of AWS resources. They are not used to store events.

Option B: Amazon SQS dead-letter queues are used to store messages that cannot be delivered to their intended recipients. They are not used to store events.

Option D: AWS CloudTrail is a service that records AWS API calls. It can be used to track the activity of EventBridge rules, but it does not store the events themselves.

upvoted 1 times

Bennyboy789 2 years, 4 months ago

Selected Answer: A

Option A is the most appropriate solution because Amazon EventBridge publishes metrics to Amazon CloudWatch. You can find relevant metrics in the "AWS/Events" namespace, which allows you to monitor the number of events matched by the rule and the number of invocations to the rule's target.

upvoted 6 times

h8er 2 years, 4 months ago

Selected Answer: A

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/events/CloudWatch-Events-Monitoring-CloudWatch-Metrics.html>

upvoted 2 times

A company has a large workload that runs every Friday evening. The workload runs on Amazon EC2 instances that are in two Availability Zones in the us-east-1 Region. Normally, the company must run no more than two instances at all times. However, the company wants to scale up to six instances each Friday to handle a regularly repeating increased workload.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create a reminder in Amazon EventBridge to scale the instances.
- B. Create an Auto Scaling group that has a scheduled action. **Most Voted**
- C. Create an Auto Scaling group that uses manual scaling.
- D. Create an Auto Scaling group that uses automatic scaling.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Bmaster **Highly Voted** 2 years, 4 months ago

B is correct.

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/ec2-auto-scaling-scheduled-scaling.html>

upvoted 9 times

wizcloudifa **Most Recent** 1 year, 7 months ago

Selected Answer: B

A - too much operation overhead, manually provisioning the instances after you receive the reminder from eventbridge

B - right answer, as you can scale up the EC2 instances and keep them ready before large overload time

C - too much operation overhead in manually scaling

D - automatic scaling will scale up the instances after some duration after it has encountered the heavy workload traffic, not ideal

upvoted 3 times

TariqKipkemei 2 years ago

Selected Answer: B

runs every Friday evening = an Auto Scaling group that has a scheduled action

upvoted 3 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: B

The key reasons:

Auto Scaling scheduled actions allow defining specific dates/times to scale out or in. This can be used to scale to 6 instances every Friday evening automatically.

Scheduled scaling removes the need for manual intervention to scale up/down for the workload.

EventBridge reminders and manual scaling require human involvement each week adding overhead.

Automatic scaling responds to demand and may not align perfectly to scale out every Friday without additional tuning.

Scheduled Auto Scaling actions provide the automation needed to scale for the weekly workload without ongoing operational overhead.

upvoted 4 times

Sat897 2 years, 4 months ago

Selected Answer: B

Predicted period.. So schedule the instance

upvoted 4 times

mrsoa 2 years, 4 months ago

Selected Answer: B

B seems to be correct
upvoted 2 times

Deepakin96 2 years, 4 months ago

Selected Answer: B

When we know the run time is Friday, we can schedule the instance to 6
upvoted 3 times

Josantru 2 years, 4 months ago

Correct B.
upvoted 4 times

A company is creating a REST API. The company has strict requirements for the use of TLS. The company requires TLSv1.3 on the API endpoints. The company also requires a specific public third-party certificate authority (CA) to sign the TLS certificate.

Which solution will meet these requirements?

A. Use a local machine to create a certificate that is signed by the third-party CA. Import the certificate into AWS Certificate Manager (ACM). Create an HTTP API in Amazon API Gateway with a custom domain. Configure the custom domain to use the certificate. **Most Voted**

B. Create a certificate in AWS Certificate Manager (ACM) that is signed by the third-party CA. Create an HTTP API in Amazon API Gateway with a custom domain. Configure the custom domain to use the certificate.

C. Use AWS Certificate Manager (ACM) to create a certificate that is signed by the third-party CA. Import the certificate into AWS Certificate Manager (ACM). Create an AWS Lambda function with a Lambda function URL. Configure the Lambda function URL to use the certificate.

D. Create a certificate in AWS Certificate Manager (ACM) that is signed by the third-party CA. Create an AWS Lambda function with a Lambda function URL. Configure the Lambda function URL to use the certificate.

Correct Answer: A

Community vote distribution

A (78%)

B (22%)

Comments

bjexamprep **Highly Voted** 2 years, 9 months ago

Selected Answer: A

I don't understand why some many people vote B. In ACM, you can either request certificate from Amazon CA or import an existing certificate. There is no option in ACM that allow you to request a certificate that can be signed by third party CA.

upvoted 22 times

markoniz 2 years, 8 months ago

I fully agree

upvoted 6 times

wsdasdasdqwdaw 2 years, 7 months ago

Hmm AWS is saying:

ACM certificates can be used to establish secure communications across the internet or within an internal network. You can request a publicly trusted certificate directly from ACM (an "ACM certificate") or import a publicly trusted certificate issued by a third party. Self-signed certificates are also supported. To provision your organization's internal PKI, you can issue ACM certificates signed by a private certificate authority (CA) created and managed by AWS Private CA. The CA may either reside in your account or be shared with you by a different account.

<https://docs.aws.amazon.com/acm/latest/userguide/gs.html>

upvoted 5 times

pentium75 2 years, 5 months ago

Exactly. You can "import [not create] a publicly trusted certificate issued by a third party".

upvoted 5 times

luiscc **Highly Voted** 2 years, 10 months ago

Selected Answer: B

AWS Certificate Manager (ACM) is a service that lets you easily provision, manage, and deploy SSL/TLS certificates for use with AWS services and your internal resources. By creating a certificate in ACM that is signed by the third-party CA, the company can meet its requirement for a specific public third-party CA to sign the TLS certificate.

upvoted 9 times

pentium75 2 years, 5 months ago

Sounds like ChatGPT answer, "creating a certificate in ACM that is signed by the third-party CA" is not possible.

upvoted 5 times

emakid **Most Recent** 1 year, 11 months ago

Selected Answer: A

A. Use a local machine to create a certificate that is signed by the third-party CA. Import the certificate into AWS Certificate Manager (ACM). Create an HTTP API in Amazon API Gateway with a custom domain. Configure the custom domain to use the certificate.

Reason:

Custom Certificate: Allows you to use a certificate signed by the third-party CA.

TLSv1.3 Support: API Gateway supports TLSv1.3 for custom domains.

Configuration: You can import the third-party CA certificate into ACM and configure API Gateway to use this certificate with a custom domain.

This approach meets all the specified requirements by allowing the use of a third-party CA-signed certificate and ensuring the API endpoints use TLSv1.3.

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: A

A is logical answer.

BCD are either misworded here or intentionally confusing. Regardless, you cannot create a cert in ACM that is signed by 3rd party CA. You can only import these certs to ACM.

upvoted 4 times

Shubhi_08 2 years, 5 months ago

Selected Answer: A

We can't create third party certificates in ACM.

upvoted 2 times

foha2012 2 years, 5 months ago

Is this a question from the associate or professional exam ??

upvoted 2 times

pentium75 2 years, 5 months ago

Selected Answer: A

ACM can import, but not create, 3rd party certificates. Leaves only A.

upvoted 2 times

maged123 2 years, 5 months ago

Selected Answer: A

You have already a publicly trusted certificate issued by a third party and you just need to import it in ACM not to create a new one. So, the correct answer is A which is the only one that importing the certificate in ACM while B, C and D are creating a new one.

upvoted 2 times

sparun1607 2 years, 6 months ago

The answer must be A,

You can't create a certificate in ACM, read the below link

<https://docs.aws.amazon.com/acm/latest/userguide/setup.html>

upvoted 2 times

numark 2 years, 6 months ago

Answer is A: Can I import a third-party certificate and use it with AWS services?

Yes. If you want to use a third-party certificate with Amazon CloudFront, Elastic Load Balancing, or Amazon API Gateway, you may import it into ACM using the AWS Management Console, AWS CLI, or ACM APIs. ACM does not manage the renewal process for imported certificates. You can use the AWS Management Console to monitor the expiration dates of an imported certificates and import a new third-party certificate to replace an expiring one.

upvoted 2 times

TariqKipkemei 2 years, 6 months ago

Selected Answer: A

It's 22/Nov/2023 and from the console you cant create a certificate in AWS Certificate Manager (ACM) that is signed by the third-party CA. But you could obtain it externally then import it into ACM.

upvoted 2 times

Tshring 2 years, 6 months ago

Selected Answer: B

Option B meets these requirements:

- API Gateway HTTP APIs support TLS 1.3
 - ACM can import certificates signed by third-party CAs
 - API Gateway provides REST APIs
- upvoted 1 times

pentium75 2 years, 5 months ago

"ACM can import (!) certificates signed by third-party CA", but not create (!) them as B suggests.

upvoted 2 times

NickGordon 2 years, 7 months ago

Selected Answer: A

In ACM you can't create a cert signed by another CA. Dude, try it by yourself. There is no such option!

upvoted 2 times

chen0305_099 2 years, 9 months ago

WHY NOT A?

upvoted 1 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: B

Use ACM to create a certificate signed by the third-party CA. ACM integrates with external CAs. Create an API Gateway HTTP API with a custom domain name. Configure the custom domain to use the ACM certificate. API Gateway supports configuring custom domains with ACM certificates. This allows serving the API over TLS using the required third-party certificate and TLS 1.3 support.

upvoted 2 times

pentium75 2 years, 5 months ago

"ACM integrates with external CAs." no

upvoted 2 times

taustin2 2 years, 10 months ago

Selected Answer: A

You can provide certificates for your integrated AWS services either by issuing them directly with ACM or by importing third-party certificates into the ACM management system.

upvoted 2 times

vini15 2 years, 10 months ago

Should be A.
We need to import third-party certificate to ACM.

upvoted 5 times

A company runs an application on AWS. The application receives inconsistent amounts of usage. The application uses AWS Direct Connect to connect to an on-premises MySQL-compatible database. The on-premises database consistently uses a minimum of 2 GiB of memory.

The company wants to migrate the on-premises database to a managed AWS service. The company wants to use auto scaling capabilities to manage unexpected workload increases.

Which solution will meet these requirements with the LEAST administrative overhead?

- A. Provision an Amazon DynamoDB database with default read and write capacity settings.
- B. Provision an Amazon Aurora database with a minimum capacity of 1 Aurora capacity unit (ACU).
- C. Provision an Amazon Aurora Serverless v2 database with a minimum capacity of 1 Aurora capacity unit (ACU). **Most Voted**
- D. Provision an Amazon RDS for MySQL database with 2 GiB of memory.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Guru4Cloud **Highly Voted** 2 years, 9 months ago

Selected Answer: C

The key reasons:

Aurora Serverless v2 provides auto-scaling so the database can handle inconsistent workloads and spikes automatically without admin intervention. It can scale down to zero when not in use to minimize costs.

The minimum 1 ACU capacity is sufficient to replace the on-prem 2 GiB database based on the info given.

Serverless capabilities reduce admin overhead for capacity management.

DynamoDB lacks MySQL compatibility and requires more hands-on management.

RDS and provisioned Aurora require manually resizing instances to scale, increasing admin overhead.

upvoted 11 times

dkw2342 2 years, 2 months ago

> It can scale down to zero when not in use to minimize costs.

This part is not correct. Aurora Serverless v1 was able to scale to zero.

upvoted 2 times

kambarami **Highly Voted** 2 years, 8 months ago

the questions are hard from 500 +

upvoted 8 times

foha2012 2 years, 5 months ago

I don't think these are associate exam questions rather are from AWS professional exam

upvoted 3 times

awsgeek75 2 years, 5 months ago

Yes, I agree. I have been reading the pro questions and these are copy paste. On the bright side, it prepares you for the next step!

upvoted 7 times

emakid **Most Recent** 1 year, 11 months ago

Selected Answer: C

C. Provision an Amazon Aurora Serverless v2 database with a minimum capacity of 1 Aurora capacity unit (ACU).

Suitability: Amazon Aurora Serverless v2 is a good option for applications with variable workloads because it automatically adjusts capacity based on

demand. It can handle MySQL-compatible databases and supports auto-scaling. You can set the minimum and maximum capacity based on your needs, making it highly suitable for handling unexpected workload increases with minimal administrative overhead.

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: C

LEAST administrative overhead = Aurora Serverless

upvoted 3 times

TariqKipkemei 2 years, 6 months ago

Selected Answer: C

LEAST administrative overhead = Serverless

upvoted 2 times

ibu007 2 years, 10 months ago

Selected Answer: C

serverless = LEAST overhead

upvoted 3 times

D10SJoker 2 years, 10 months ago

Why not D?

upvoted 2 times

awsgeek75 2 years, 5 months ago

Because "LEAST administrative overhead" is a requirement. RDS configured with mem requirements is an admin overhead

upvoted 2 times

wizcloudifa 2 years, 1 month ago

no autoscaling with RDS

upvoted 2 times

mrsoa 2 years, 10 months ago

Selected Answer: C

C seems to be the right answer

Instead of provisioning and managing database servers, you specify Aurora capacity units (ACUs). Each ACU is a combination of approximately 2 gigabytes (GB) of memory, corresponding CPU, and networking. Database storage automatically scales from 10 gibibytes (GiB) to 128 tebibytes (TiB), the same as storage in a standard Aurora DB cluster

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/aurora-serverless-v1.how-it-works.html>

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/aurora-serverless-v2.html>

upvoted 2 times

Bmaster 2 years, 10 months ago

C is correct.

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/aurora-serverless-v2.how-it-works.html#aurora-serverless-v2.how-it-works.capacity>

upvoted 3 times

A company wants to use an event-driven programming model with AWS Lambda. The company wants to reduce startup latency for Lambda functions that run on Java 11. The company does not have strict latency requirements for the applications. The company wants to reduce cold starts and outlier latencies when a function scales up.

Which solution will meet these requirements MOST cost-effectively?

- A. Configure Lambda provisioned concurrency.
- B. Increase the timeout of the Lambda functions.
- C. Increase the memory of the Lambda functions.
- D. Configure Lambda SnapStart. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: D

The key reasons:

SnapStart keeps functions initialized and ready to respond quickly, eliminating cold starts. SnapStart is optimized for applications without aggressive latency needs, reducing costs. It scales automatically to match traffic spikes, eliminating outliers when scaling up. SnapStart is a native Lambda feature with no additional charges, keeping costs low. Provisioned concurrency incurs charges for always-on capacity reserved. More costly than SnapStart. Increasing timeout and memory do not directly improve startup performance like SnapStart.

upvoted 15 times

anikety123 **Highly Voted** 2 years, 3 months ago

Selected Answer: D

Both Lambda SnapStart and provisioned concurrency can reduce cold starts and outlier latencies when a function scales up. SnapStart helps you improve startup performance by up to 10x at no extra cost. Provisioned concurrency keeps functions initialized and ready to respond in double-digit milliseconds. Configuring provisioned concurrency incurs charges to your AWS account. Use provisioned concurrency if your application has strict cold start latency requirements. You can't use both SnapStart and provisioned concurrency on the same function version.

upvoted 6 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: D

<https://docs.aws.amazon.com/lambda/latest/dg/snapstart.html>

"Lambda SnapStart for Java can improve startup performance for latency-sensitive applications by up to 10x at no extra cost, typically with no changes to your function code."

upvoted 5 times

awsgeek75 1 year, 11 months ago

Also

A: Solves concurrency issues not startup

B is for execution timeout (don't think that possible if I understand the option correctly)

C Memory is not the issue here

upvoted 2 times

TariqKipkemei 2 years ago

Selected Answer: D

Lambda SnapStart it is.

<https://docs.aws.amazon.com/lambda/latest/dg/snapstart.html#:~:text=RSS-,Lambda%20SnapStart,-for%20Java%20can>
upvoted 2 times

TariqKipkemei 2 years ago

only because its a Java 11 app...if it were any other besides Java I believe Provisioned concurrency could help.

upvoted 2 times

potomac 2 years, 1 month ago

Selected Answer: D

Lambda SnapStart for Java can improve startup performance for latency-sensitive applications by up to 10x at no extra cost, typically with no change to your function code.

<https://docs.aws.amazon.com/lambda/latest/dg/snapstart.html>
upvoted 3 times

BrijMohan08 2 years, 3 months ago

Selected Answer: D

<https://docs.aws.amazon.com/lambda/latest/dg/snapstart.html>
upvoted 2 times

skyphilip 2 years, 3 months ago

Selected Answer: D

D is correct

Lambda SnapStart for Java can improve startup performance for latency-sensitive applications by up to 10x at no extra cost, typically with no change to your function code. The largest contributor to startup latency (often referred to as cold start time) is the time that Lambda spends initializing the function, which includes loading the function's code, starting the runtime, and initializing the function code.

With SnapStart, Lambda initializes your function when you publish a function version. Lambda takes a Firecracker microVM snapshot of the memory and disk state of the initialized execution environment, encrypts the snapshot, and caches it for low-latency access. When you invoke the function version for the first time, and as the invocations scale up, Lambda resumes new execution environments from the cached snapshot instead of initializing them from scratch, improving startup latency.

upvoted 2 times

avkya 2 years, 4 months ago

"SnapStart does not support provisioned concurrency, the arm64 architecture, Amazon Elastic File System (Amazon EFS), or ephemeral storage greater than 512 MB." The question says "The company wants to reduce cold starts" This means provisioned concurrency. I'm a little bit confused with D.

upvoted 3 times

Woodlawn5700 2 years, 4 months ago

D

<https://docs.aws.amazon.com/lambda/latest/dg/snapstart.html>
upvoted 2 times

mrsoa 2 years, 4 months ago

Selected Answer: D

D is the answer

Lambda SnapStart for Java can improve startup performance for latency-sensitive applications by up to 10x at no extra cost, typically with no change to your function code. The largest contributor to startup latency (often referred to as cold start time) is the time that Lambda spends initializing the function, which includes loading the function's code, starting the runtime, and initializing the function code.

<https://docs.aws.amazon.com/lambda/latest/dg/snapstart.html>
upvoted 3 times

Bmaster 2 years, 4 months ago

D is best!!

A is not MOST cost effectly.

lambda snapshot is new feature for lambda.

<https://docs.aws.amazon.com/lambda/latest/dg/snapstart.html>
upvoted 4 times

Bmaster 2 years, 4 months ago

misspell.... lambda snapstart

upvoted 2 times

RaksAWS 2 years, 4 months ago

why not D
It should work
upvoted 3 times

A financial services company launched a new application that uses an Amazon RDS for MySQL database. The company uses the application to track stock market trends. The company needs to operate the application for only 2 hours at the end of each week. The company needs to optimize the cost of running the database.

Which solution will meet these requirements MOST cost-effectively?

- A. Migrate the existing RDS for MySQL database to an Aurora Serverless v2 MySQL database cluster. **Most Voted**
- B. Migrate the existing RDS for MySQL database to an Aurora MySQL database cluster.
- C. Migrate the existing RDS for MySQL database to an Amazon EC2 instance that runs MySQL. Purchase an instance reservation for the EC2 instance.
- D. Migrate the existing RDS for MySQL database to an Amazon Elastic Container Service (Amazon ECS) cluster that uses MySQL container images to run tasks.

Correct Answer: A

Community vote distribution

A (92%)

8%

Comments

Guru4Cloud **Highly Voted** 2 years, 9 months ago

Selected Answer: A

The key reasons are:

Aurora Serverless v2 scales compute capacity automatically based on actual usage, down to zero when not in use. This minimizes costs for intermittent usage.

Since it only runs for 2 hours per week, the application is ideal for a serverless architecture like Aurora Serverless.

Aurora Serverless v2 charges per second when the database is active, unlike RDS which charges hourly.

Aurora Serverless provides higher availability than self-managed MySQL on EC2 or ECS.

Using reserved EC2 instances or ECS still incurs charges when not in use versus the fine-grained scaling of serverless.

Standard Aurora clusters have a minimum capacity unlike the auto-scaling serverless architecture.

upvoted 14 times

dkw2342 2 years, 2 months ago

A is correct, but Aurora Serverless v2 only scales down to 0.5 ACU, not to zero.

upvoted 2 times

pentium75 **Highly Voted** 2 years, 5 months ago

Selected Answer: A

2 hours per week = Serverless = A. Recommended for "infrequent, intermittent, or unpredictable workloads"

upvoted 5 times

LeonSauveterre **Most Recent** 1 year, 6 months ago

Selected Answer: A

A: Aurora Serverless v2 offers the most cost-effective and operationally efficient solution for intermittent workloads like this.

B: over-provisioned and too expensive for this workload

C: Also expensive. You gotta pay for the instance 24/7, regardless of usage.

D: Introduces complexity and operational overhead for database management tasks. Overall, ECS is never easy to handle.

upvoted 1 times

awsgeek75 2 years, 5 months ago

Selected Answer: A

B is wrong because Aurora MySQL cluster will just keep on running for the rest of the week and will be costly.

C and D have too much infra bloating so costly

upvoted 2 times

TariqKipkemei 2 years, 6 months ago

Selected Answer: A

Answer is A.

Here are the key distinctions:

Amazon Aurora: provides built-in security, continuous backups, serverless compute, up to 15 read replicas, automated multi-Region replication, and integrations with other AWS services.

Amazon Aurora Serverless: is an on-demand, auto-scaling configuration for Aurora where the database automatically starts up, shuts down, and scales capacity up or down based on your application's needs.

With serverless the db will shut down when not in use.

upvoted 5 times

anikety123 2 years, 9 months ago

Selected Answer: A

Option is A

upvoted 3 times

hachiri 2 years, 9 months ago

Selected Answer: A

Aurora Serverless

- Automated database instantiation and auto-scaling based on actual usage
- Good for infrequent, intermittent or unpredictable workloads
- No capacity planning needed
- Pay per second, can be more cost-effective

upvoted 3 times

vini15 2 years, 10 months ago

will go with A

Amazon Aurora Serverless v2 is suitable for the most demanding, highly variable workloads. For example, your database usage might be heavy for a short period of time, followed by long periods of light activity or no activity at all.

upvoted 3 times

msdnpro 2 years, 10 months ago

Selected Answer: A

"Amazon Aurora Serverless v2 is suitable for the most demanding, highly variable workloads. For example, your database usage might be heavy for a short period of time, followed by long periods of light activity or no activity at all. "

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/aurora-serverless-v2.how-it-works.html>

upvoted 3 times

ersin13 2 years, 10 months ago

A. Migrate the existing RDS for MySQL database to an Aurora Serverless v2 MySQL database cluster.

upvoted 2 times

mrsoa 2 years, 10 months ago

Selected Answer: B

B seems to be the correct answer, because if we have a predictable workload Aurora database seems to be most cost effective however if we have unpredictable workload aurora serverless seems to be more cost effective because our database will scale up and down

for more informations please read this article

<https://medium.com/trackit/aurora-or-aurora-serverless-v2-which-is-more-cost-effective-bcd12e172dcf>

upvoted 3 times

Chef_couincouin 2 years, 7 months ago

according to the link, i understand that Aurora Serverless is ideal for sudden peaks in database usage with moderate or minimal usage during other periods of the day. So Answer is A

upvoted 3 times

Smart 2 years, 9 months ago

True but due to autoscaling - it will be cheaper...check example#1 in the your link.

upvoted 1 times

Smart 2 years, 9 months ago

Correct Answer is A

upvoted 2 times

pentium75 2 years, 5 months ago

Provisioned RDS (as in B) is good for steady (not "predictable") workloads. In this case, the workload is predictable, but the prediction is that it will be used only 2 hours per week.

upvoted 3 times

pentium75 2 years, 5 months ago

Aurora Serverless is for "infrequent, intermittent, OR unpredictable workloads"

upvoted 3 times

A company deploys its applications on Amazon Elastic Kubernetes Service (Amazon EKS) behind an Application Load Balancer in an AWS Region. The application needs to store data in a PostgreSQL database engine. The company wants the data in the database to be highly available. The company also needs increased capacity for read workloads.

Which solution will meet these requirements with the MOST operational efficiency?

- A. Create an Amazon DynamoDB database table configured with global tables.
- B. Create an Amazon RDS database with Multi-AZ deployments.
- C. Create an Amazon RDS database with Multi-AZ DB cluster deployment. **Most Voted**
- D. Create an Amazon RDS database configured with cross-Region read replicas.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Guru4Cloud **Highly Voted** 2 years, 9 months ago

Selected Answer: C

RDS Multi-AZ DB cluster deployments provide high availability, automatic failover, and increased read capacity. A multi-AZ cluster automatically handles replicating data across AZs in a single region. This maintains operational efficiency as it is natively managed by RDS without needing external replication. DynamoDB global tables involve complex provisioning and requires app changes. RDS read replicas require manual setup and management of replication. RDS Multi-AZ clustering is purpose-built by AWS for HA PostgreSQL deployments and balancing read workloads.

upvoted 9 times

luiscc **Highly Voted** 2 years, 10 months ago

Selected Answer: C

DB cluster deployment can scale read workloads by adding read replicas. This provides increased capacity for read workloads without impacting the write workload.

upvoted 5 times

MatAlves **Most Recent** 1 year, 8 months ago

Selected Answer: C

"A Multi-AZ DB cluster deployment is a semisynchronous, high availability deployment mode of Amazon RDS with two readable replica DB instances."

upvoted 2 times

upliftinghut 2 years, 4 months ago

Selected Answer: C

multi-AZ addresses both HA & increased read capacity with synchronous data replication between main DB & standby. Read replica is not enough because only increased read capacity not enabling HA, besides the data replication is async

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: C

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/multi-az-db-clusters-concepts.html>

"A Multi-AZ DB cluster deployment is a semisynchronous, high availability deployment mode of Amazon RDS with two readable standby DB instances"

A: DynamoDB is not Postgres

B: Although HA is achieve but it does not increase the read capacity as much as C without additional operational complexity

D: Cross region is not a requirement and won't solve the same region HA or read issues

upvoted 2 times

aws94 2 years, 6 months ago

Selected Answer: C

Multi-AZ DB Cluster Deployment = Aurora
upvoted 2 times

TariqKipkemei 2 years, 6 months ago

Selected Answer: C

Multi-AZ DB cluster deployments provides two readable DB instances if you need additional read capacity
upvoted 2 times

potomac 2 years, 7 months ago

Selected Answer: C

C is correct
upvoted 2 times

avkya 2 years, 10 months ago

Selected Answer: C

Multi-AZ DB clusters provide high availability, increased capacity for read workloads, and lower write latency when compared to Multi-AZ DB instance deployments.
upvoted 2 times

mrsoa 2 years, 10 months ago

Selected Answer: C

CCCCCCCCCcCCcCCcCCCCccCc
upvoted 2 times

A company is building a RESTful serverless web application on AWS by using Amazon API Gateway and AWS Lambda. The users of this web application will be geographically distributed, and the company wants to reduce the latency of API requests to these users.

Which type of endpoint should a solutions architect use to meet these requirements?

- A. Private endpoint
- B. Regional endpoint
- C. Interface VPC endpoint
- D. Edge-optimized endpoint **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

mrsoa **Highly Voted** 2 years, 4 months ago

Selected Answer: D

The correct answer is D

API Gateway - Endpoint Types

- Edge-Optimized (default): For global clients
- Requests are routed through the CloudFront Edge locations (improves latency)
- The API Gateway still lives in only one region
- Regional:
 - For clients within the same region
 - Could manually combine with CloudFront (more control over the caching strategies and the distribution)
- Private:
 - Can only be accessed from your VPC using an interface VPC endpoint (ENI)
 - Use a resource policy to define access

upvoted 9 times

potomac **Highly Voted** 2 years, 1 month ago

Selected Answer: D

An edge-optimized API endpoint typically routes requests to the nearest CloudFront Point of Presence (POP), which could help in cases where your clients are geographically distributed. This is the default endpoint type for API Gateway REST APIs.

<https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-api-endpoint-types.html>

upvoted 5 times

awsgeek75 **Most Recent** 1 year, 10 months ago

Selected Answer: D

geographically distributed users + low latency = Edge optimized endpoint

upvoted 2 times

TariqKipkemei 2 years ago

Selected Answer: D

An edge-optimized API endpoint typically routes requests to the nearest CloudFront Point of Presence (POP), which could help in cases where your clients are geographically distributed. This is the default endpoint type for API Gateway REST APIs.

<https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-api-endpoint-types.html#:~:text=API%20endpoint%20typically,-routes,-requests%20to%20the>

upvoted 3 times

dilaaziz 2 years, 1 month ago

Selected Answer: D

<https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-api-endpoint-types.html>

upvoted 3 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: D

Edge-optimized endpoint

upvoted 3 times

Josantru 2 years, 4 months ago

Correct D.

Edge-optimized API endpoints

An edge-optimized API endpoint is best for geographically distributed clients. API requests are routed to the nearest CloudFront Point of Presence (POP). This is the default endpoint type for API Gateway REST APIs.

upvoted 3 times

A company uses an Amazon CloudFront distribution to serve content pages for its website. The company needs to ensure that clients use a TLS certificate when accessing the company's website. The company wants to automate the creation and renewal of the TLS certificates.

Which solution will meet these requirements with the MOST operational efficiency?

- A. Use a CloudFront security policy to create a certificate.
- B. Use a CloudFront origin access control (OAC) to create a certificate.
- C. Use AWS Certificate Manager (ACM) to create a certificate. Use DNS validation for the domain. **Most Voted**
- D. Use AWS Certificate Manager (ACM) to create a certificate. Use email validation for the domain.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Bmaster **Highly Voted** 2 years, 4 months ago

C is correct.

"ACM provides managed renewal for your Amazon-issued SSL/TLS certificates. This means that ACM will either renew your certificates automatically (if you are using DNS validation), or it will send you email notices when expiration is approaching. These services are provided for both public and private ACM certificates."

<https://docs.aws.amazon.com/acm/latest/userguide/managed-renewal.html>

upvoted 10 times

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: C

The key reasons are:

AWS Certificate Manager (ACM) provides free public TLS/SSL certificates and handles certificate renewals automatically.

Using DNS validation with ACM is operationally efficient since it automatically makes changes to Route 53 rather than requiring manual validation steps.

ACM integrates natively with CloudFront distributions for delivering HTTPS content.

CloudFront security policies and origin access controls do not issue TLS certificates.

Email validation requires manual steps to approve the domain validation emails for each renewal.

upvoted 7 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: C

For me, C is the only realistic option as I don't think you can do AB without a lot of complexity. D just makes no sense.

upvoted 2 times

ibu007 2 years, 3 months ago

Selected Answer: C

Use AWS Certificate Manager (ACM) to create a certificate. Use DNS validation for the domain

upvoted 4 times

chen0305_099 2 years, 3 months ago

Selected Answer: C

C 似乎是正确的

upvoted 4 times

Kiki_Pass 2 years, 4 months ago

Selected Answer: C

"DNS Validation is preferred for automation purposes" -- Stephane's course on Udemy
upvoted 3 times

mrsoa 2 years, 4 months ago

Selected Answer: C

C seems to be correct
upvoted 2 times

nananashi 2 years, 4 months ago

I think the general product uses DNS rather than email to automate, is the given answer correct?
upvoted 1 times

A company deployed a serverless application that uses Amazon DynamoDB as a database layer. The application has experienced a large increase in users. The company wants to improve database response time from milliseconds to microseconds and to cache requests to the database.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use DynamoDB Accelerator (DAX). **Most Voted**
- B. Migrate the database to Amazon Redshift.
- C. Migrate the database to Amazon RDS.
- D. Use Amazon ElastiCache for Redis.

Correct Answer: A

Community vote distribution

A (91%)

9%

Comments

h8er **Highly Voted** 2 years, 10 months ago

Selected Answer: A

Amazon DynamoDB Accelerator (DAX) is a fully managed, highly available, in-memory cache for Amazon DynamoDB that delivers up to a 10 times performance improvement—from milliseconds to microseconds—even at millions of requests per second.

[https://aws.amazon.com/dynamodb/dax/#:~:text=Amazon%20DynamoDB%20Accelerator%20\(DAX\)%20is,millions%20of%20requests%20per%20second.](https://aws.amazon.com/dynamodb/dax/#:~:text=Amazon%20DynamoDB%20Accelerator%20(DAX)%20is,millions%20of%20requests%20per%20second.)

upvoted 11 times

MatAlves **Most Recent** 1 year, 8 months ago

Amazon ElastiCache for Redis would help with "caching requests", but not "improve database response" itself.

upvoted 2 times

awsgeek75 2 years, 5 months ago

Selected Answer: A

DAX is least operations overhead.

B: Redshift, although powerful, but is for analytics

C: Downgrading to RDS won't help

D: EC for Redis is more for persistent caching so would be good but lot of operational overhead

upvoted 3 times

TariqKipkemei 2 years, 6 months ago

Selected Answer: A

improve DynamoDB response time from milliseconds to microseconds and to cache requests to the database = DynamoDB Accelerator (DAX)

upvoted 3 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: C

Use DynamoDB Accelerator (DAX).

upvoted 2 times

pentium75 2 years, 5 months ago

Which is A, not C.

upvoted 3 times

awsgeek75 2 years, 5 months ago

Quote A but mark C. You need more coffee mate :)

upvoted 5 times

mrsoa 2 years, 10 months ago

Selected Answer: A

A is the right answer

upvoted 3 times

Bmaster 2 years, 10 months ago

Correct A.

upvoted 2 times

A company runs an application that uses Amazon RDS for PostgreSQL. The application receives traffic only on weekdays during business hours. The company wants to optimize costs and reduce operational overhead based on this usage.

Which solution will meet these requirements?

- A. Use the Instance Scheduler on AWS to configure start and stop schedules. **Most Voted**
- B. Turn off automatic backups. Create weekly manual snapshots of the database.
- C. Create a custom AWS Lambda function to start and stop the database based on minimum CPU utilization.
- D. Purchase All Upfront reserved DB instances.

Correct Answer: A

Community vote distribution

A (96%)

4%

Comments

potomac **Highly Voted** 2 years, 1 month ago

Selected Answer: A

The Instance Scheduler on AWS solution automates the starting and stopping of Amazon Elastic Compute Cloud (Amazon EC2) and Amazon Relational Database Service (Amazon RDS) instances.

This solution helps reduce operational costs by stopping resources that are not in use and starting them when they are needed. The cost savings can be significant if you leave all of your instances running at full utilization continuously.

<https://aws.amazon.com/solutions/implementations/instance-scheduler-on-aws/>

upvoted 6 times

pentium75 **Most Recent** 1 year, 11 months ago

Selected Answer: A

B increases operational overhead

C Lambda functions could work but NOT "based on minimum CPU utilization"

D might save cost but not as much as A

upvoted 4 times

ibu007 2 years, 3 months ago

Selected Answer: A

A. Use the Instance Scheduler on AWS to configure start and stop schedules

upvoted 4 times

baba365 2 years, 2 months ago

Why not D?

upvoted 3 times

AndreiWebNet 2 years ago

How do you actually reduce costs enough to buy upfront instances that you pay for them if you use them 1 h or 24 and it is paid to run 24h. It says that you use this instances 8 hours a day 5 days a week, totaling 40h a week.... so is it the difference from 40h to 168 h?

upvoted 3 times

master9 1 year, 11 months ago

When you buy Reserved Instances, the larger the upfront payment, the greater the discount. To maximize your savings, you can pay all up-front and receive the largest discount. Partial up-front RI's offer lower discounts but give you the option to spend less up front. Lastly, you can choose to spend nothing up front and receive a smaller discount, but allowing you to free up capital to spend in other projects.

But you need some mechanism to stop on weekend and in night to save cost.

upvoted 2 times

ErnShm 2 years, 3 months ago

A

<https://docs.aws.amazon.com/solutions/latest/instance-scheduler-on-aws/solution-overview.html>

upvoted 3 times

Guru4Cloud 2 years, 3 months ago

Selected Answer: A

Purpose-built scheduling minimizes operational overhead.

Aligns instance running time precisely with business hour demands.

Maintains backups unlike disabling auto backups.

More cost effective and flexible than reserved instances.

Simpler to implement than a custom Lambda function.

upvoted 4 times

anikety123 2 years, 3 months ago

Selected Answer: B

Its B. Check the AWS link

https://aws.amazon.com/solutions/implementations/instance-scheduler-on-aws/?nc1=h_ls

upvoted 1 times

anikety123 2 years, 3 months ago

Sorry I wanted to select A.

upvoted 5 times

mrsoa 2 years, 4 months ago

Selected Answer: A

A

<https://aws.amazon.com/solutions/implementations/instance-scheduler-on-aws/>

upvoted 2 times

luiscc 2 years, 4 months ago

Selected Answer: A

Scheduler do the job

upvoted 4 times

A company uses locally attached storage to run a latency-sensitive application on premises. The company is using a lift and shift method to move the application to the AWS Cloud. The company does not want to change the application architecture.

Which solution will meet these requirements MOST cost-effectively?

- A. Configure an Auto Scaling group with an Amazon EC2 instance. Use an Amazon FSx for Lustre file system to run the application.
- B. Host the application on an Amazon EC2 instance. Use an Amazon Elastic Block Store (Amazon EBS) GP2 volume to run the application.
- C. Configure an Auto Scaling group with an Amazon EC2 instance. Use an Amazon FSx for OpenZFS file system to run the application.
- D. Host the application on an Amazon EC2 instance. Use an Amazon Elastic Block Store (Amazon EBS) GP3 volume to run the application. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

TariqKipkemei **Highly Voted** 1 year, 6 months ago

Selected Answer: D

MOST cost-effectively =GP3

upvoted 5 times

potomac **Most Recent** 1 year, 7 months ago

Selected Answer: D

gp3 offers SSD-performance at a 20% lower cost per GB than gp2 volumes.

upvoted 3 times

bojila 1 year, 9 months ago

GP3 is the latest version

upvoted 3 times

Hades2231 1 year, 9 months ago

Selected Answer: D

GP3 is the latest version, and it is cost effective

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

GP3 is preferable over GP2, FSx for Lustre, and FSx for OpenZFS is clear and convincing:

GP3 offers identical latency performance to GP2 at a lower price point.

FSx options are higher performance but more expensive and require application changes.

GP3 aligns better with lift and shift needs as a directly attached block storage volume.

upvoted 4 times

taustin2 1 year, 10 months ago

Selected Answer: D

Migrate your Amazon EBS volumes from gp2 to gp3 and save up to 20% on costs.

upvoted 3 times

Vadbro7 1 year, 10 months ago

Y not gp2

upvoted 1 times

Ale1973 1 year, 10 months ago

Selected Answer: D

My rational:

Options A y C are based on autoscaling-group and no make sense for me on this scenary.
Then, use Amazon EBS is the solution and GP2 or GP3 is the question.

Requirement requires the most COST effective solution, then, I choose GP3

upvoted 4 times

A company runs a stateful production application on Amazon EC2 instances. The application requires at least two EC2 instances to always be running.

A solutions architect needs to design a highly available and fault-tolerant architecture for the application. The solutions architect creates an Auto Scaling group of EC2 instances.

Which set of additional steps should the solutions architect take to meet these requirements?

A. Set the Auto Scaling group's minimum capacity to two. Deploy one On-Demand Instance in one Availability Zone and one On-Demand Instance in a second Availability Zone.

B. Set the Auto Scaling group's minimum capacity to four. Deploy two On-Demand Instances in one Availability Zone and two On-Demand Instances in a second Availability Zone. **Most Voted**

C. Set the Auto Scaling group's minimum capacity to two. Deploy four Spot Instances in one Availability Zone.

D. Set the Auto Scaling group's minimum capacity to four. Deploy two On-Demand Instances in one Availability Zone and two Spot Instances in a second Availability Zone.

Correct Answer: B

Community vote distribution

B (75%)

A (25%)

Comments

luiscc **Highly Voted** 2 years, 10 months ago

Selected Answer: B

By setting the Auto Scaling group's minimum capacity to four, the architect ensures that there are always at least two running instances. Deploying two On-Demand Instances in each of two Availability Zones ensures that the application is highly available and fault-tolerant. If one Availability Zone becomes unavailable, the application can still run in the other Availability Zone.

upvoted 19 times

Ale1973 **Highly Voted** 2 years, 10 months ago

Selected Answer: A

My rational is: Highly available = 2 AZ, and then 2 EC2 instances always running is 1 EC2 in each AZ. If an entire AZ fails, SacalinGroup deploy the minimum instances (2) on the running AZ

upvoted 15 times

baba365 2 years, 8 months ago

Ans: A.

The application requires at least two EC2 instances to always be running = 2 minimum capacity... minimum cap of 4 ec2 will work but a waste of resources that doesn't follow well archi. framework.

upvoted 2 times

Ramdi1 2 years, 8 months ago

it says always have to have two running, hence you need 4. two in each AV. it might be a waste of resource but if that what is required by the company then so be it. Also you out the 4 you cannot use spot instances because if the two instances on the on demand go down and you need to use the spot instance they could be called back at any point.

upvoted 5 times

Ramdi1 2 years, 8 months ago

AZ * not AV

upvoted 3 times

bora4motion **Most Recent** 1 year, 1 month ago

Selected Answer: B

A and C are not compatible with the "fault tolerant" condition.

upvoted 1 times

emakid 1 year, 11 months ago

Option A:

Set the Auto Scaling group's minimum capacity to two. Deploy one On-Demand Instance in one Availability Zone and one On-Demand Instance in a second Availability Zone.

This configuration ensures that you have two instances running across two different AZs, which provides high availability. However, it does not take advantage of additional capacity to handle failures or spikes in demand. If either AZ becomes unavailable, you will have one running instance, but this does not meet the requirement of having at least two running instances at all times.

upvoted 1 times

emakid 1 year, 11 months ago

Selected Answer: B

Set the Auto Scaling group's minimum capacity to four. Deploy two On-Demand Instances in one Availability Zone and two On-Demand Instances in a second Availability Zone.

This configuration provides high availability with four instances distributed across two AZs. The minimum capacity of four ensures that even if one instance fails, there are still two instances in each AZ to handle the load. This option is highly available and fault-tolerant but may be more than required if only two instances are needed to be always running.

upvoted 2 times

cjace 1 year, 12 months ago

Answer is A: If one Availability Zone fails, the Auto Scaling group will automatically launch a new instance in a different, healthy Availability Zone to maintain the desired capacity of two instances. This is one of the key benefits of using Auto Scaling groups—they automatically maintain the desired number of instances across multiple Availability Zones, ensuring that your application is highly available and fault-tolerant. So even in the event of a failure in one Availability Zone, your application will continue to run on the required number of instances. This is why it's recommended to distribute instances across multiple Availability Zones when designing architectures for high availability and fault tolerance.

upvoted 1 times

LuongTo 1 year, 7 months ago

that is reasonable. but what do you think about "always"? With A, there's a moment there would be only one EC2 healthy.

upvoted 2 times

Marco_St 2 years, 5 months ago

Selected Answer: B

so indeed ASG can set up a new EC2 instance in another AZ if there is one AZ failed with fault but it failed to meet the need of always having 2 instance running before the new instance replacement is done in the working AZ. so this is why we deploy 2 instances per AZ

upvoted 3 times

pentium75 2 years, 5 months ago

Selected Answer: B

If it would not mention the "stateful" application, and if it would only have to be "highly available" but NOT "fault-tolerant", A would be fine.

upvoted 5 times

1rob 2 years, 6 months ago

Selected Answer: B

From <<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/spot-best-practices.html> : Spot Instances are not suitable for workloads that are inflexible, stateful, fault-intolerant, or tightly coupled between instance nodes. So C and D don't fit.

From <<https://docs.aws.amazon.com/whitepapers/latest/real-time-communication-on-aws/use-multiple-availability-zones.html> : Within the constructs of AWS, customers are encouraged to run their workloads in more than one Availability Zone. This ensures that customer applications can withstand even a complete Availability Zone failure - a very rare event in itself.

So a HA solution in this case implies a total of 4 instances, 2 per AZ.

upvoted 2 times

TariqKipkemei 2 years, 6 months ago

Selected Answer: B

The main requirement here is a 'highly available and fault-tolerant architecture for the application', this covered by option B.

The application requires at least two EC2 instances to always be running, main word here being 'atleast' which means more than two is ok.

upvoted 2 times

Ramdi1 2 years, 8 months ago

Selected Answer: B

B - Need 2 in each AZ and you cant use spot instances as it could be recalled.

upvoted 2 times

Mandar15 2 years, 8 months ago

Selected Answer: B

Stateful is keyword here. 2 is minimum required all time.
upvoted 2 times

JA2018 1 year, 6 months ago

minimum of 2x EC2 instances per AZ at all times.
upvoted 1 times

MI11975 2 years, 9 months ago

Selected Answer: A

If a complete AZ fails, autoscale will launch a second EC2 in the running AZ. If that short period of time is not always, which is not, then the answer is B, but I would take my chances and select A in the exam xD because the application is highly available and fault-tolerant.
upvoted 1 times

Guru4Cloud 2 years, 9 months ago

Selected Answer: B

° Minimum of 4 ensures at least 2 instances are always running in each AZ, meeting the HA requirement.
° On-Demand instances provide consistent performance and availability, unlike Spot.
° Spreading across 2 AZs adds fault tolerance, protecting from AZ failure.
upvoted 4 times

darkknight23 2 years, 10 months ago

Selected Answer: B

While Spot Instances can be used to reduce costs, they might not provide the same level of availability and guaranteed uptime that On-Demand Instances offer. So I will go with B and not D.
upvoted 2 times

Sat897 2 years, 10 months ago

Selected Answer: B

Highly available - 2 AZ and then 2 EC2 instances always running. 2 in each AZ.
upvoted 3 times

Sat897 2 years, 10 months ago

Highly available - 2 AZ and then 2 EC2 instances always running. 2 in each AZ..
upvoted 2 times

An ecommerce company uses Amazon Route 53 as its DNS provider. The company hosts its website on premises and in the AWS Cloud. The company's on-premises data center is near the us-west-1 Region. The company uses the eu-central-1 Region to host the website. The company wants to minimize load time for the website as much as possible.

Which solution will meet these requirements?

- A. Set up a geolocation routing policy. Send the traffic that is near us-west-1 to the on-premises data center. Send the traffic that is near eu-central-1 to eu-central-1. **Most Voted**
- B. Set up a simple routing policy that routes all traffic that is near eu-central-1 to eu-central-1 and routes all traffic that is near the on-premises datacenter to the on-premises data center.
- C. Set up a latency routing policy. Associate the policy with us-west-1.
- D. Set up a weighted routing policy. Split the traffic evenly between eu-central-1 and the on-premises data center.

Correct Answer: A

Community vote distribution

A (89%)

11%

Comments

Guru4Cloud **Highly Voted** 2 years, 3 months ago

Selected Answer: A

The key reasons are:

Geolocation routing allows you to route users to the closest endpoint based on their geographic location. This will provide the lowest latency. Routing us-west-1 traffic to the on-premises data center minimizes latency for those users since it is also located near there. Routing eu-central-1 traffic to the eu-central-1 AWS region minimizes latency for users nearby. This achieves routing users to the closest endpoint on a geographic basis to optimize for low latency.

upvoted 5 times

PLN6302 2 years, 3 months ago

why can't be the option C

upvoted 2 times

lemur88 2 years, 3 months ago

You cannot associate the policy to us-west-1 as the AWS account is in eu-central-1

upvoted 5 times

awsgeek75 **Most Recent** 1 year, 11 months ago

Selected Answer: A

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-geo.html>

B can be done but definition of "near" is ambiguous

C wrong region

D wrong solution as splitting evenly does not reduce latency for on-prem server users

upvoted 2 times

Cyberkayu 1 year, 11 months ago

Selected Answer: A

not C. Client do not have AWS us-west-1 region. Client have a on prem DC near west-1

not D. 2 people visit the site together near eu-central-1, one of the user may be thrown to west-1 due to load balancing on split even weighted policy.

A and B are both valid, latency = how soon user reach the datacenter and received a responses from the DC, round trip. So in short, geolocation or send user to the nearest DC will improve latency.

upvoted 2 times

TariqKipkemei 2 years ago

Selected Answer: A

Geolocation routing policy allows you to route traffic based on the location of your users.
upvoted 4 times

t0nx 2 years ago

Selected Answer: C

C. Set up a latency routing policy. Associate the policy with us-west-1.

Explanation:

A latency routing policy directs traffic based on the lowest network latency to the specified AWS endpoint. Since the on-premises data center is near the us-west-1 Region, associating the policy with us-west-1 ensures that users near that region will be directed to the on-premises data center.

This allows for optimal routing, minimizing the load time for users based on their geographical proximity to the respective hosting locations (us-west-1 and eu-central-1).

Options A, B, and D do not explicitly consider latency or are not optimal for minimizing load time:

Option A (geolocation routing policy) would direct traffic based on the geographic location of the user but may not necessarily optimize for the lowest latency.

upvoted 2 times

awsgeek75 1 year, 10 months ago

There is nothing in us-west-1 as the company's data centre is near us-west-1.

upvoted 1 times

Chiquitabandita 2 years ago

except I don't think that it should be applied to the west region. If Geolocation is applied and the west is closer to the client, but the west is having intermittent issues at the time, they will have a longer latency even though closer to that region. this is why I would apply latency in a real world solution.

upvoted 1 times

Chiquitabandita 2 years ago

in real world I think it should use latency routing if the main concern is to lower the latency but AWS likes to promote geolocation and if that is in the question I think that will be the answer so I choose A.

upvoted 1 times

baba365 2 years, 2 months ago

The company wants to minimize load time for the website as much as possible... between data Centre and website or between users and website?

upvoted 1 times

Hades2231 2 years, 3 months ago

Selected Answer: A

Geolocation is the key word

upvoted 2 times

lemur88 2 years, 3 months ago

Selected Answer: A

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-geo.html>

upvoted 2 times

A company has 5 PB of archived data on physical tapes. The company needs to preserve the data on the tapes for another 10 years for compliance purposes. The company wants to migrate to AWS in the next 6 months. The data center that stores the tapes has a 1 Gbps uplink internet connectivity.

Which solution will meet these requirements MOST cost-effectively?

- A. Read the data from the tapes on premises. Stage the data in a local NFS storage. Use AWS DataSync to migrate the data to Amazon S3 Glacier Flexible Retrieval.
- B. Use an on-premises backup application to read the data from the tapes and to write directly to Amazon S3 Glacier Deep Archive.
- C. Order multiple AWS Snowball devices that have Tape Gateway. Copy the physical tapes to virtual tapes in Snowball. Ship the Snowball devices to AWS. Create a lifecycle policy to move the tapes to Amazon S3 Glacier Deep Archive. **Most Voted**
- D. Configure an on-premises Tape Gateway. Create virtual tapes in the AWS Cloud. Use backup software to copy the physical tape to the virtual tape.

Correct Answer: C

Community vote distribution

C (98%)

2

Comments

adeyinkaamole **Highly Voted** 2 years, 9 months ago

If you have made it to the end of the exam dump, you will definitely pass your exams in Jesus name. After over a year of Procrastination, I am finally ready to write my AWS Solutions Architect Exam. Thank you Exam Topics

upvoted 35 times

Hades2231 **Highly Voted** 2 years, 9 months ago

Selected Answer: C

Ready for the exam tomorrow. Wish you guys all the best. BTW Snowball Device comes in handy when you need to move a huge amount of data but cant afford any bandwidth loss

upvoted 13 times

MatAlves **Most Recent** 1 year, 8 months ago

Oh, to think now we have to study 904 questions instead of just 583 lol

upvoted 3 times

JA2018 1 year, 6 months ago

currently (22 Nov 2024), question count now stands at 1.019

upvoted 3 times

XXXXXINN 1 year, 8 months ago

well, 2 weeks later today, it says 981 questions...

upvoted 3 times

awsgeek75 2 years, 5 months ago

Selected Answer: C

5PB over 1GB connection will take approximately 15 months so anything with "transfer" is invalid. ABD are not practical.

C: Just order snowball

upvoted 5 times

pentium75 2 years, 5 months ago

Selected Answer: C