

A company has multiple AWS accounts with applications deployed in the us-west-2 Region. Application logs are stored within Amazon S3 buckets in each account. The company wants to build a centralized log analysis solution that uses a single S3 bucket. Logs must not leave us-west-2, and the company wants to incur minimal operational overhead.

Which solution meets these requirements and is MOST cost-effective?

- A. Create an S3 Lifecycle policy that copies the objects from one of the application S3 buckets to the centralized S3 bucket.
- B. Use S3 Same-Region Replication to replicate logs from the S3 buckets to another S3 bucket in us-west-2. Use this S3 bucket for log analysis. **Most Voted**
- C. Write a script that uses the PutObject API operation every day to copy the entire contents of the buckets to another S3 bucket in us-west-2. Use this S3 bucket for log analysis.
- D. Write AWS Lambda functions in these accounts that are triggered every time logs are delivered to the S3 buckets (s3:ObjectCreated:* event). Copy the logs to another S3 bucket in us-west-2. Use this S3 bucket for log analysis.

Correct Answer: B

Community vote distribution

B (100%)

Comments

BillaRanga **Highly Voted** 2 years, 3 months ago

Selected Answer: B

The main Use case of S3 same region replication is "log aggregation, live replication between production and test accounts".
upvoted 12 times

Andy_09 **Highly Voted** 2 years, 4 months ago

Option B

upvoted 6 times

zdi561 **Most Recent** 1 year, 4 months ago

Selected Answer: B

A and B can copy, but B is cheaper, it is free.
upvoted 1 times

ckhemani 1 year, 6 months ago

Selected Answer: B

Option C and D is out of question since it required operation overhead.
Between A and B, Lifecycle is used to transfer between S3 storage classes and not within same S3 Buckets. So B is correct answer.
upvoted 2 times

sandordini 2 years, 1 month ago

Selected Answer: B

Needs to be B
upvoted 2 times

A company has an application that delivers on-demand training videos to students around the world. The application also allows authorized content developers to upload videos. The data is stored in an Amazon S3 bucket in the us-east-2 Region.


The company has created an S3 bucket in the eu-west-2 Region and an S3 bucket in the ap-southeast-1 Region. The company wants to replicate the data to the new S3 buckets. The company needs to minimize latency for developers who upload videos and students who stream videos near eu-west-2 and ap-southeast-1.

Which combination of steps will meet these requirements with the FEWEST changes to the application? (Choose two.)

- A. Configure one-way replication from the us-east-2 S3 bucket to the eu-west-2 S3 bucket. Configure one-way replication from the us-east-2 S3 bucket to the ap-southeast-1 S3 bucket.
- B. Configure one-way replication from the us-east-2 S3 bucket to the eu-west-2 S3 bucket. Configure one-way replication from the eu-west-2 S3 bucket to the ap-southeast-1 S3 bucket.
- C. Configure two-way (bidirectional) replication among the S3 buckets that are in all three Regions. **Most Voted**
- D. Create an S3 Multi-Region Access Point. Modify the application to use the Amazon Resource Name (ARN) of the Multi-Region Access Point for video streaming. Do not modify the application for video uploads.
- E. Create an S3 Multi-Region Access Point. Modify the application to use the Amazon Resource Name (ARN) of the Multi-Region Access Point for video streaming and uploads. **Most Voted**

Correct Answer: CE

Community vote distribution



Comments

BillaRanga **Highly Voted** 2 years, 3 months ago

Selected Answer: CE

To keep replication in SYNC across all three regions, we use Bi-directional. Multi-Region Access Point for video streaming and uploads. -> uploads to nearest Low latency region and Bi-directional replication will keep other two regions in SYNC this reducing the upload and streaming latency
upvoted 14 times

bujuman 2 years, 1 month ago

For confirmation purposes: <https://aws.amazon.com/s3/features/multi-region-access-points/>

upvoted 3 times

KarlosRC **Most Recent** 3 months ago

Selected Answer: AE

B: — Chain replication is bad practice. C: Two-way (bidirectional) replication is not required by the condition. D: Where is the latency minimization for upload?

upvoted 1 times

JustEugen 8 months, 2 weeks ago

Selected Answer: CE

If we need to minimize latency, only a combination of C and E can achieve this.

My exam is tomorrow, hope to pass

upvoted 4 times

Onimole 5 months, 3 weeks ago

how did it go?

upvoted 3 times

In2718 10 months, 2 weeks ago

Selected Answer: CE

Because of this line - The company needs to minimize latency for developers who upload videos and students who stream videos near eu-west-2 and ap-southeast-1... there are two locations, so I think C and E may be the best solution, and basically it's the same amount of work as A and D.

upvoted 1 times

Dz1990 1 year ago

Selected Answer: CE

 ****Grammar Breakdown:****

"minimize latency for [developers who upload videos] and [students who stream videos] near eu-west-2 and ap-southeast-1"

The phrase "near eu-west-2 and ap-southeast-1" modifies ****BOTH**** developers and students!

 ****Corrected Answer: CE****

Why CE is Actually Correct:

****Option C - Bidirectional Replication:****

-  Allows developers to upload to their nearest region (low latency)
-  Content automatically syncs to all other regions
-  No single point of failure for uploads

****Option E - MRAP for Everything:****

-  ****Fewest application changes****: Single endpoint for all operations
-  Automatic routing for both uploads AND downloads
-  Future-proof and scalable

upvoted 2 times

srcntpc 1 year ago

Selected Answer: AD

A & D

A: Sets up replication to eu-west-2 and ap-southeast-1.

D: Uses MRAP to optimize streaming latency without modifying upload logic.

upvoted 1 times

sk1974 1 year, 3 months ago

Selected Answer: AD

I will go with A & D too since that option will require least amount of re-work . In reality , I would have preferred B for the copy option since Ap-Sotheast-1 is closer to eu-west , but , that could involve additional work

upvoted 1 times

robotgeek 1 year, 4 months ago

Selected Answer: A

Question does not make a lot of sense with this responses

- 1) If the goal is to modify the app as little as possible, options D and E don't present much of a difference. If you modify the app for uploads, it will cost you the same to modify it for streaming (even more so with an MRAP, since you only need to change the global upload endpoint, which is a trivial change).
- 2) It is not suggested that the need is "to upload from anywhere in the world" but rather to copy to other regions for streaming (download), which an MRAP does not do per-se.
- 3) If the idea of the question is to create an MRAP (D or E) and then set up replication rules in the MRAP (C), it doesn't make sense to replicate from US-West or from Asia-Pacific, but only the other way around.
- 4) Option A stands on its own (but it's a multiple-choice answer), and even though the initial replication latency is slow, uploading video tutorials is not critical.
- 5) Option B involves unnecessary latency.

upvoted 1 times

robotgeek 1 year, 4 months ago

So the idea is you put an MRAP and then replicate... so basically the "FEWEST changes to the application" is either to confuse or poorly constructed

upvoted 1 times

FlyingHawk 1 year, 5 months ago

Selected Answer: AD

The data is stored in an Amazon S3 bucket in the us-east-2 Region implies the content developer is in us-east-2 region. If the content developers are in the us-east-2 Region, then the solution is A and D.

upvoted 3 times

kajiyatta 1 year, 5 months ago

Selected Answer: AD

This is an example of the kind of joke problem AWS provides. The information needed to derive justification is vague and it is not clear which region the uploading developer will be using.

Time spent on these problems is wasted.

upvoted 4 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: AD

I agree with @kbgsgsgs and here's why:

What the scenario is about: Currently, videos are in us, so they're uploaded to us. Then, we want to replicate them to eu and ap. People near eu and ap should be fast delivered with the videos they need.

So with "FEWEST changes", all we need is replicate from us to both eu and ap, which is option B. Option C works but not fewest changes.

The uploads are directly to us, so option E actually induces more latency. You may argue that if we can upload videos directly to eu/ap, it would be faster. Yes, but it requires modifying the application for both streaming and upload workflows, so much more complicated than option D, and definitely not fewest changes.

upvoted 2 times

LeonSauveterre 1 year, 5 months ago

Sorry, typo. All we need is replicate from us to both eu and ap, which is option A, not option B. You can share you ideas here if you disagree

upvoted 1 times

JA2018 1 year, 6 months ago

Selected Answer: CD

My personal take: C & D

C => Although stem did not specify which region S# bucket the developers will upload their contents, this will allow new contents in 1 S3 bucket to replicate over to the other 2 S3 buckets.

D => uploads needs to be performed to the first region only (Closest???) and accessed by remaining two S3 buckets.

upvoted 1 times

JA2018 1 year, 6 months ago

but I am confused by Google AI search's response:

The correct answers to this question are D. Create an S3 Multi-Region Access Point. Modify the application to use the Amazon Resource Name (ARN) of the Multi-Region Access Point for video streaming and E. Create an S3 Multi-Region Access Point. Modify the application to use the Amazon Resource Name (ARN) of the Multi-Region Access Point for video streaming and uploads.

upvoted 1 times

JA2018 1 year, 6 months ago

Explanation:

Minimizing changes to the application:

The question explicitly asks for the solution with the "fewest changes to the application.". Option D and E only require modifying the application to use the Multi-Region Access Point ARN, while options A, B, and C would involve additional replication configurations, making them less ideal.

Multi-Region Access Point benefits:

An S3 Multi-Region Access Point allows users to access data from any region through a single endpoint, effectively minimizing latency for users in different regions without requiring data replication across multiple buckets.

upvoted 1 times

JA2018 1 year, 6 months ago

Why not the other options:

Options A and B:

These involve one-way replication, which would still require users in some regions to access data from a geographically distant bucket, leading to higher latency.

Option C:

Bidirectional replication would create unnecessary data redundancy and additional network traffic, making it less efficient.

Key takeaway: When you need to minimize latency for users in different regions while minimizing application changes, an S3 Multi-Region Access Point is the best solution.

upvoted 1 times

JA2018 1 year, 6 months ago

is it just me or am I missing something in this question?

upvoted 1 times

kbgsgsgs 1 year, 8 months ago

Selected Answer: AD

They are simply trying to replicate to a new S3 bucket. I don't see why it needs to be bidirectional. Also, since the problem assumes that the content developer with permission is the one uploading, it seems like there needs to be a way to centralize the upload without modifying the application.

upvoted 3 times

1166ae3 1 year, 11 months ago

Selected Answer: AE

Since developer upload video to us-east-2, by configuring one-way replication directly from us-east-2 to eu-west-2 and from us-east-2 to ap-southeast-1, you ensure that each region has the latest data without additional replication hops.

upvoted 1 times

Scheldon 1 year, 11 months ago

Selected Answer: CE

AnswerCE

From my understanding Video uploads can happen near new regions hence to speed up that operation we need to upload to nearest region, hence would choose option E, and for the same reason we need to be able to replicate data from any of new region to old one and opposite, hence we need bidirectional (two-way) replication

upvoted 1 times

lenotc 2 years, 2 months ago

Selected Answer: CD

FEWEST changes to the application

D -> MRAP can upload the appropriate S3 bucket

C -> two-way -> to worry about anything

obs: I believe this question dubious, amphibological

upvoted 1 times

67a3f49 2 years, 3 months ago

There is no information where the upload should be performed. If files will be uploaded to first region then:

AD because:

A -> content uploaded to the primary bucket in us-east-2 is automatically replicated to the other regions, minimizing latency for users accessing content near those regions.

D -> uploads needs to be performed to the first region only and accessed to remaining two

Otherwise CE

upvoted 4 times

Salilgen 1 year, 5 months ago

You are right but I think files are only uploaded to first region.

The question states "A company has an application that delivers on-demand training videos to students around the world". It doesn't state anything about developer.

Furthermore, in this case both AD and CE work but AD requires fewest changes to the application

upvoted 2 times

A company has a new mobile app. Anywhere in the world, users can see local news on topics they choose. Users also can post photos and videos from inside the app.


Users access content often in the first minutes after the content is posted. New content quickly replaces older content, and then the older content disappears. The local nature of the news means that users consume 90% of the content within the AWS Region where it is uploaded.

Which solution will optimize the user experience by providing the LOWEST latency for content uploads?

- A. Upload and store content in Amazon S3. Use Amazon CloudFront for the uploads.
- B. Upload and store content in Amazon S3. Use S3 Transfer Acceleration for the uploads. **Most Voted**
- C. Upload content to Amazon EC2 instances in the Region that is closest to the user. Copy the data to Amazon S3.
- D. Upload and store content in Amazon S3 in the Region that is closest to the user. Use multiple distributions of Amazon CloudFront.

Correct Answer: B

Community vote distribution



Comments

Cali182 **Highly Voted** 2 years, 4 months ago

Selected Answer: B

Cloudfront is for reading not for uploading

Option B

upvoted 18 times

Salilgen 1 year, 5 months ago

I don't think so

<https://aws.amazon.com/blogs/aws/amazon-cloudfront-content-uploads-post-put-other-methods/>

upvoted 1 times

BillaRanga **Highly Voted** 2 years, 3 months ago

Selected Answer: B

Question says - " LOWEST latency for content uploads"

Hence Use S3 Transfer Acceleration for the uploads.

upvoted 14 times

flaviobrf **Most Recent** 1 year, 10 months ago

Selected Answer: D

Cloudfront can also upload data, not just for caching content

upvoted 2 times

flaviobrf 1 year, 10 months ago

Sorry, for this scenario, i believe that B its the correct

upvoted 1 times

ike001 1 year, 12 months ago

B is the answer

upvoted 2 times

7fb06b3 2 years ago

Selected Answer: D

CloudFront does support upload acceleration

<https://aws.amazon.com/blogs/aws/amazon-cloudfront-content-uploads-post-put-other-methods/>

upvoted 3 times

TruthWS 2 years, 2 months ago

Option D

upvoted 2 times

alawada 2 years, 2 months ago

Selected Answer: B

Amazon S3 Transfer Acceleration utilizes Amazon CloudFront's globally distributed edge locations to accelerate content uploads to Amazon S3.

upvoted 3 times

xBUGx 2 years, 2 months ago

Selected Answer: B

S3TA is actually using cloudfront's infrastructure.

So, yes B. Which is just an optimized solution for cloudfront itself.

upvoted 2 times

Ipergorta 2 years, 2 months ago

Option D

Regional S3 Buckets: Storing content in S3 buckets located in the same Region as the user minimizes the physical distance the data needs to travel during upload, reducing latency.

CloudFront Distributions: CloudFront is a content delivery network (CDN) that caches content in edge locations around the world. By creating multiple CloudFront distributions with edge locations closest to users, the content can be served with minimal latency for downloads.

upvoted 3 times

Andy_09 2 years, 4 months ago

Option D

upvoted 4 times

jaswantn 2 years, 4 months ago

option B... S3 transfer acceleration for LOWEST latency for content uploads. question is not asking for low latency for content retrieval.

Happy to be corrected

upvoted 4 times

A company is building a new application that uses serverless architecture. The architecture will consist of an Amazon API Gateway REST API and AWS Lambda functions to manage incoming requests.

The company wants to add a service that can send messages received from the API Gateway REST API to multiple target Lambda functions for processing. The service must offer message filtering that gives the target Lambda functions the ability to receive only the messages the functions need.

Which solution will meet these requirements with the LEAST operational overhead?

A. Send the requests from the API Gateway REST API to an Amazon Simple Notification Service (Amazon SNS) topic. Subscribe Amazon Simple Queue Service (Amazon SQS) queues to the SNS topic. Configure the target Lambda functions to poll the different SQS queues. **Most Voted**

B. Send the requests from the API Gateway REST API to Amazon EventBridge. Configure EventBridge to invoke the target Lambda functions.

C. Send the requests from the API Gateway REST API to Amazon Managed Streaming for Apache Kafka (Amazon MSK). Configure Amazon MSK to publish the messages to the target Lambda functions.

D. Send the requests from the API Gateway REST API to multiple Amazon Simple Queue Service (Amazon SQS) queues. Configure the target Lambda functions to poll the different SQS queues.

Correct Answer: A

Community vote distribution



Comments

Kezuko **Highly Voted** 2 years, 2 months ago

Selected Answer: A

"message filtering" = SNS

upvoted 17 times

AMEJack **Highly Voted** 1 year, 6 months ago

Selected Answer: B

The question is asking about one service. If you choose A, you need to build 3 services one per each SNS topic. Also, creating one Event Bridge with rules to forward to the target Lambda is easier than creating 3 topics with 3 SQS queues. The answer is definitely B.

upvoted 6 times

Pucmaster **Most Recent** 3 months, 2 weeks ago

Selected Answer: B

B event bridge works with gateway

upvoted 1 times

In2718 10 months, 2 weeks ago

Selected Answer: B

I vote for answer B from reading through the discussions

upvoted 1 times

GOTJ 1 year, 3 months ago

Selected Answer: A

I'm discarding option "B" for the following reasons

1. You can't send "requests" (messages) to EventBridge, at least not directly as option suggests. There must be an event involved in the process

2. Lack of configuration regarding to discrete Lambda invocations

Option "A", by its side, also lacks of the necessary configuration for discrete lambda invocation (use of subscriber filter), but in my opinion is closer to what company needs:

<https://docs.aws.amazon.com/sns/latest/dg/sns-message-filtering.html>

So my vote goes to "A".

upvoted 2 times

Chen77 1 year, 3 months ago

Selected Answer: A

SNS supports fan-out: API Gateway sends messages to an SNS topic, which then forwards to multiple SQS queues.
Message filtering with SNS: SNS allows each SQS queue to receive only the relevant messages, reducing unnecessary processing.
SQS queues decouple workloads: Lambda functions pull messages only when ready, ensuring efficient and scalable processing.

upvoted 2 times

zdi561 1 year, 4 months ago

Selected Answer: B

A and B both can do message filtering but A only filters message by attributes and B can do filtering by message content too, and B is simpler

upvoted 1 times

hilker1983 1 year, 5 months ago

Selected Answer: B

LEAST operational overhead -> B better than A

upvoted 1 times

Denise123 1 year, 5 months ago

Selected Answer: B

The answer is B. On the other hand, the solution involving Amazon SNS, Amazon SQS, and Lambda polling (Option A) introduces more moving parts and operational overhead:

It requires managing and configuring multiple services (SNS, SQS, and Lambda).

Lambda functions need to continuously poll SQS queues, which can be less efficient and introduce additional latency.

Scaling and managing SQS queues and their associated resources (e.g., dead-letter queues) add operational complexity.

While the SNS-SQS-Lambda solution is still valid and widely used, the EventBridge-Lambda integration offers a more streamlined and serverless-native approach, aligning better with the company's serverless architecture goals and minimizing operational overhead.

upvoted 1 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: B

EventBridge applies filtering rules and invokes the appropriate Lambda functions directly, just like SNS + SQS. Option B is definitely "the LEAST operational overhead".

While there's a 5-target-per-rule limit, you can work around it by creating additional rules. Each rule can directly invoke Lambda functions without the need for SQS queues or polling.

Let's not stick to option A just because it's the textbook exam-standard SNS+SQS combination. There are other options that will help you achieve the goal under different circumstances.

upvoted 2 times

JA2018 1 year, 6 months ago

Selected Answer: B

agreed with AMEJack.... based on the stem, B required the LEAST operational overhead

upvoted 2 times

JA2018 1 year, 6 months ago

From Google AI Search:

The correct answer is B. Send the requests from the API Gateway REST API to Amazon EventBridge. Configure EventBridge to invoke the target Lambda functions.

upvoted 1 times

JA2018 1 year, 6 months ago

Explanation:

Least operational overhead:

EventBridge is designed specifically for event routing and filtering, making it the most efficient option for managing which Lambda functions receive messages based on specific criteria, minimizing operational overhead compared to the other choices.

Message filtering:

EventBridge allows you to define detailed rules with attributes and conditions to filter messages and send them only to the relevant Lambda functions, fulfilling the requirement for message filtering.

Scalability:

Both EventBridge and SNS can handle large volumes of messages, but EventBridge offers better management and control over message routing, especially when dealing with complex event patterns.

upvoted 1 times

JA2018 1 year, 6 months ago

Why the other options are less ideal:

A. SNS with SQS:

While SNS can be used for fan-out messaging, adding SQS queues introduces extra complexity and potential bottlenecks in managing multiple queues and subscriptions.

C. Amazon MSK:

While Kafka is a powerful messaging system, it might be overkill for this scenario due to its complexity and overhead for managing a dedicated Kafka cluster, especially if the message volume is not very high.

D. Multiple SQS queues:

This option would require additional logic to decide which SQS queue to send messages to based on filtering criteria, adding complexity to the architecture.

Key takeaway: When looking for a solution with the least operational overhead for message routing and filtering between an API Gateway and multiple Lambda functions, Amazon EventBridge is the most suitable choice.

upvoted 1 times

1166ae3 1 year, 11 months ago

Selected Answer: B

LEAST operational overhead -> B better than A

upvoted 5 times

Scheldon 1 year, 11 months ago

Selected Answer: A

AnswerA

Hence EventBridge is a solution to handle events and we need to handle messages I believe option A is the best solution here

upvoted 3 times

3bdf1cc 1 year, 12 months ago

<https://aws.amazon.com/blogs/compute/capturing-client-events-using-amazon-api-gateway-and-amazon-eventbridge/>

upvoted 2 times

BBR01 2 years, 1 month ago

Selected Answer: A

The main issue with B is that with Eventbridge, you can only define up to five targets for each rule.

<https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-targets.html>

upvoted 4 times

sandordini 2 years, 1 month ago

B: EventBridge reacts to events, not requests or messages.

C: I don't think so, but I don't know MSK well enough.

D: You can add a filter so that your function only processes Amazon SQS messages containing certain data parameters. but it will still receive, so I assume it's not what the question asks for.

Only A remains... But it still misses steps plus we are looking for the Least ops overhead..

I am confused..

upvoted 2 times

03beafc 2 years, 1 month ago

Selected Answer: B

Eventbridge + lambda is two services, sns + sqs + lambda is 3. Both can filter, but the config involved in eventbridge > lambda is easier

upvoted 3 times

MatAlves 1 year, 8 months ago

SNS provides built-in message filtering.

SNS -> SQS -> Lambda = very common Fanout Scenario

upvoted 2 times

A company migrated millions of archival files to Amazon S3. A solutions architect needs to implement a solution that will encrypt all the archival data by using a customer-provided key. The solution must encrypt existing unencrypted objects and future objects.

Which solution will meet these requirements?

A. Create a list of unencrypted objects by filtering an Amazon S3 Inventory report. Configure an S3 Batch Operations job to encrypt the objects from the list with a server-side encryption with a customer-provided key (SSE-C). Configure the S3 default encryption feature to use a server-side encryption with a customer-provided key (SSE-C). **Most Voted**

B. Use S3 Storage Lens metrics to identify unencrypted S3 buckets. Configure the S3 default encryption feature to use a server-side encryption with AWS KMS keys (SSE-KMS).

C. Create a list of unencrypted objects by filtering the AWS usage report for Amazon S3. Configure an AWS Batch job to encrypt the objects from the list with a server-side encryption with AWS KMS keys (SSE-KMS). Configure the S3 default encryption feature to use a server-side encryption with AWS KMS keys (SSE-KMS).

D. Create a list of unencrypted objects by filtering the AWS usage report for Amazon S3. Configure the S3 default encryption feature to use a server-side encryption with a customer-provided key (SSE-C).

Correct Answer: A

Community vote distribution

A (100%)

Comments

OX_HDR **Highly Voted** 2 years, 4 months ago

Selected Answer: A

A seems correct here.

<https://aws.amazon.com/blogs/storage/encrypting-objects-with-amazon-s3-batch-operations/>

upvoted 9 times

BillaRanga **Highly Voted** 2 years, 3 months ago

Selected Answer: A

S3 inventory list has "Encryption status" field so you can use this to filter the unencrypted objects. and use S3 batch to encrypt it with SSE-C key.

AWS Usage report does not provide details about encryption status of individual objects

upvoted 9 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: A

As stated, "by using a customer-provided key", which rules out B & C.

Option A uses 3 things:

1. S3 Inventory - Provides a native way to generate a list of unencrypted objects, making it scalable for millions of objects.
2. S3 Batch Operations - Encrypts existing objects efficiently.
3. S3 Default Encryption - Ensures future objects are automatically encrypted (with SSE-C).

upvoted 1 times

Scheldon 1 year, 11 months ago

Selected Answer: A

AnswerA

upvoted 2 times

ike001 1 year, 12 months ago

A is the answer

upvoted 2 times

jaswantn 2 years, 4 months ago

option B.... S3 Inventory report to check for unencrypted objects in s3 and then using Batch operation.

upvoted 1 times

mestule 2 years, 4 months ago

Selected Answer: A

The solution must encrypt existing unencrypted objects. Batch will do that.

upvoted 5 times

Andy_09 2 years, 4 months ago

Option B

upvoted 1 times

sk1974 1 year, 3 months ago

Need to use customer provided keys. So, B cannot be the answer

upvoted 1 times

The DNS provider that hosts a company's domain name records is experiencing outages that cause service disruption for a website running on AWS. The company needs to migrate to a more resilient managed DNS service and wants the service to run on AWS.

What should a solutions architect do to rapidly migrate the DNS hosting service?

- A. Create an Amazon Route 53 public hosted zone for the domain name. Import the zone file containing the domain records hosted by the previous provider. **Most Voted**
- B. Create an Amazon Route 53 private hosted zone for the domain name. Import the zone file containing the domain records hosted by the previous provider.
- C. Create a Simple AD directory in AWS. Enable zone transfer between the DNS provider and AWS Directory Service for Microsoft Active Directory for the domain records.
- D. Create an Amazon Route 53 Resolver inbound endpoint in the VPC. Specify the IP addresses that the provider's DNS will forward DNS queries to. Configure the provider's DNS to forward DNS queries for the domain to the IP addresses that are specified in the inbound endpoint.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Andy_09 **Highly Voted** 2 years, 4 months ago

Option A

upvoted 10 times

BillaRanga **Highly Voted** 2 years, 3 months ago

Selected Answer: A

A -> Correct as we need to route to a Company in public network.

B -> No, because it routes only within one or more VPC

C -> Added as a distractor

D -> Inbound resolver is for traffic from On-Prem to VPC

upvoted 8 times

aditianand 2 years ago

Hello BillaRanga, I just bought this examtopics. My exam is on Jun 9th. Did we get questions from exam topics? How was the exam?

upvoted 3 times

ccceb01 **Most Recent** 1 year, 9 months ago

Selected Answer: A

A is the answer

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/migrate-dns-domain-in-use.html>

upvoted 2 times

pawanghujanamazon53 2 years, 1 month ago

Selected Answer: A

option A

upvoted 2 times

A company is building an application on AWS that connects to an Amazon RDS database. The company wants to manage the application configuration and to securely store and retrieve credentials for the database and other services.

Which solution will meet these requirements with the LEAST administrative overhead?

- A. Use AWS AppConfig to store and manage the application configuration. Use AWS Secrets Manager to store and retrieve the credentials. **Most Voted**
- B. Use AWS Lambda to store and manage the application configuration. Use AWS Systems Manager Parameter Store to store and retrieve the credentials.
- C. Use an encrypted application configuration file. Store the file in Amazon S3 for the application configuration. Create another S3 file to store and retrieve the credentials.
- D. Use AWS AppConfig to store and manage the application configuration. Use Amazon RDS to store and retrieve the credentials.

Correct Answer: A

Community vote distribution

A (100%)

Comments

BillaRanga **Highly Voted** 1 year, 9 months ago

Selected Answer: A

AppConfig useCase = You can use AWS AppConfig to deploy configuration data stored in the AWS AppConfig hosted configuration store, AWS Secrets Manager, Systems Manager Parameter Store, or Amazon S3.
So B and C are out.

use RDS to store credentials is not a good practise. So D is out.

Ans is A

upvoted 10 times

Andy_09 **Highly Voted** 1 year, 10 months ago

Option A

upvoted 8 times

Awsbeginner87 **Most Recent** 1 year, 8 months ago

Credentials= secrets Manager

upvoted 3 times

To meet security requirements, a company needs to encrypt all of its application data in transit while communicating with an Amazon RDS MySQL DB instance. A recent security audit revealed that encryption at rest is enabled using AWS Key Management Service (AWS KMS), but data in transit is not enabled.

What should a solutions architect do to satisfy the security requirements?

- A. Enable IAM database authentication on the database.
- B. Provide self-signed certificates. Use the certificates in all connections to the RDS instance.
- C. Take a snapshot of the RDS instance. Restore the snapshot to a new instance with encryption enabled.
- D. Download AWS-provided root certificates. Provide the certificates in all connections to the RDS instance. **Most Voted**

Correct Answer: D

Community vote distribution

D (81%)

A (19%)

Comments

BillaRanga **Highly Voted** 2 years, 3 months ago

Selected Answer: D

Amazon RDS creates an SSL certificate and installs the certificate on the DB instance when the instance is provisioned. So it is AWS provided.
upvoted 13 times

Andy_09 **Highly Voted** 2 years, 4 months ago

Option D

upvoted 5 times

Arad **Most Recent** 1 year, 6 months ago

Selected Answer: A

A is the correct answer
upvoted 1 times

JA2018 1 year, 6 months ago

Selected Answer: D

Why D is correct:

To encrypt data in transit to an RDS MySQL instance, you need to configure the database connection to use TLS/SSL, which requires using AWS-provided root certificates to establish a secure connection.

upvoted 3 times

JA2018 1 year, 6 months ago

Why other options are incorrect:

A (IAM database authentication): While important for user access control, this does not address data encryption in transit.

B (Self-signed certificates): Using self-signed certificates is not recommended for production environments as they cannot be verified by the client and might raise security concerns.

C (Snapshot and restore): Taking a snapshot and restoring to a new instance with encryption enabled only affects data at rest, not data in transit.

upvoted 2 times

JA2018 1 year, 6 months ago

From Google AI Search:

Key point: To encrypt data in transit to an RDS instance, ensure your application uses the appropriate AWS-provided root certificates to establish a secure TLS/SSL connection.

upvoted 1 times

Scheldon 1 year, 11 months ago

Selected Answer: D

AnswerD

upvoted 2 times

DAIYL 2 years, 1 month ago

Selected Answer: D

Even if IAM database authentication is enabled, clients still need to download and configure the AWS-provided root certificate to ensure a secure connection using SSL/TLS encryption. Without configuring the certificate, communication may not be fully encrypted, even with IAM authentication enabled.

https://docs.aws.amazon.com/zh_cn/AmazonRDS/latest/UserGuide/UsingWithRDS.SSL.html

upvoted 4 times

Nm55569 2 years ago

That's not in any of the answers - "Provide the certificates in all connections to the RDS instance." this doesn't make sense with option D - it's not saying configure to trust the CA. Answer can only be option A. Your link includes this "Optionally, your SSL/TLS connection can perform server identity verification by validating the server certificate installed on your database.". This you don't actually need to trust the using CA and can configure the app that way - the traffic is still encrypted though.

upvoted 3 times

Kezuko 2 years, 2 months ago

Selected Answer: A

A

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/UsingWithRDS.IAMDBAuth.html>

upvoted 4 times

Sivaeas 2 years, 3 months ago

Option A:

IAM database authentication provides the following benefits:

Network traffic to and from the database is encrypted using Secure Socket Layer (SSL) or Transport Layer Security (TLS). For more information about using SSL/TLS with Amazon RDS, see Using SSL/TLS to encrypt a connection to a DB instance or cluster.

upvoted 4 times

A company is designing a new web service that will run on Amazon EC2 instances behind an Elastic Load Balancing (ELB) load balancer. However, many of the web service clients can only reach IP addresses authorized on their firewalls.

What should a solutions architect recommend to meet the clients' needs?

- A. A Network Load Balancer with an associated Elastic IP address. **Most Voted**
- B. An Application Load Balancer with an associated Elastic IP address.
- C. An A record in an Amazon Route 53 hosted zone pointing to an Elastic IP address.
- D. An EC2 instance with a public IP address running as a proxy in front of the load balancer.

Correct Answer: A

Community vote distribution

A (95%)

5%

Comments

67a3f49 **Highly Voted** 2 years, 3 months ago

A for sure. The same question was in "AWS Certified Solutions Architect Associate Practice Test 3" on Udemy. There was an explanation that NLB needs to be before ALB because only NLB can have static IP.

upvoted 17 times

alawada **Highly Voted** 2 years, 2 months ago

Selected Answer: A

A - correct (Static ip can thereafter be used for client whitelisting)
Using a Network Load Balancer instead of a Classic Load Balancer has the following benefits:
Support for static IP addresses for the load balancer.
<https://docs.aws.amazon.com/elasticloadbalancing/latest/network/introduction.html>

upvoted 5 times

bora4motion **Most Recent** 1 year ago

Selected Answer: A

Why not B? we're talking about http traffic not tcp / udp. Why not ALB instead of NLB?

upvoted 1 times

FlyingHawk 1 year, 5 months ago

Selected Answer: A

Both A and C are correct:

A Uses an NLB because it supports Elastic IP addresses, which are static IPs that can be whitelisted by clients.

C : Create an Alias record in Route 53 that points to the NLB.

An Alias record is a Route 53-specific feature that allows you to map a domain name (e.g., [webservice.example.com](https://www.example.com)) directly to the NLB without needing to use the NLB's IP address in the DNS record.

Clients can access the web service using the domain name (e.g., [webservice.example.com](https://www.example.com)), which resolves to the NLB's IP address.

However, in the context of exam, I will select A since the question explicitly mentions the use of an Elastic Load Balancer (ELB), and A directly addresses this by recommending a Network Load Balancer (NLB) with an Elastic IP address.

upvoted 2 times

Scheldon 1 year, 11 months ago

Selected Answer: A

AnswerA

upvoted 2 times

Sivaeas 2 years, 3 months ago

Selected Answer: A

Option A

Please look into the below for detailed explanation

<https://www.scalefactory.com/blog/2021/12/13/aws-network-load-balancers-new-features/img/previouslly-firewall-egress.png>

upvoted 3 times

PolarFox 2 years, 3 months ago

Selected Answer: C

Option C

upvoted 1 times

JA2018 1 year, 6 months ago

B. An Application Load Balancer with an associated Elastic IP address:

While an ALB can use an Elastic IP, it operates at the application layer and may not be suitable when the primary concern is client firewall restrictions requiring a specific IP.

C. An A record in an Amazon Route 53 hosted zone pointing to an Elastic IP address:

While Route 53 can manage DNS records, it does not directly address the need for a specific IP address accessible by clients with firewall restrictions.

D. An EC2 instance with a public IP address running as a proxy in front of the load balancer:

This adds unnecessary complexity and introduces a single point of failure, making it less than optimal.

upvoted 1 times

BillaRanga 2 years, 3 months ago

Selected Answer: A

B -> Application Load Balancer cannot be assigned an Elastic IP address (static IP address).

C -> Its DNS after all, "Associated elastic IP" is what IP? Makes no sense

D -> "If you require a persistent public IP address that can be associated to and from instances as you require, use an Elastic IP address instead."

PUBLIC IP of an EC2 is not persistent, although we can give an Elastic IP, Using EC2 in front of a Load Balancer is toooooo much. What if it gets a million request? So to scale that EC2 you use another LB and an ASG? This makes no sense

A is correct because a NLB can have an elastic IP and we can use this in our firewall as per the use case

upvoted 5 times

hajra313 2 years, 4 months ago

Setting up an EC2 instance with a public IP address to act as a proxy in front of the load balancer allows clients with restricted IP access to connect to the web service. The EC2 instance can handle IP address whitelisting and proxy requests to the ELB load balancer, ensuring that only authorized clients can access the service. This solution provides flexibility and control over access while leveraging the scalability and availability benefits of ELB.

upvoted 1 times

BillaRanga 2 years, 3 months ago

Is this ChatGPT answer? Can you provide the AWS documentation link?

upvoted 3 times

Andy_09 2 years, 4 months ago

Option C

upvoted 2 times

jaswantn 2 years, 3 months ago

is there any valid justification for opting C? Glad to be informed, as these questions are tricky to answer.

upvoted 2 times

jaswantn 2 years, 3 months ago

My inclination is for Option D, but not 100 % sure

upvoted 1 times

JA2018 1 year, 6 months ago

D. An EC2 instance with a public IP address running as a proxy in front of the load balancer:

This adds unnecessary complexity and introduces a single point of failure, making it less than optimal.

upvoted 1 times

A company has established a new AWS account. The account is newly provisioned and no changes have been made to the default settings. The company is concerned about the security of the AWS account root user.

What should be done to secure the root user?

- A. Create IAM users for daily administrative tasks. Disable the root user.
- B. Create IAM users for daily administrative tasks. Enable multi-factor authentication on the root user. **Most Voted**
- C. Generate an access key for the root user. Use the access key for daily administration tasks instead of the AWS Management Console.
- D. Provide the root user credentials to the most senior solutions architect. Have the solutions architect use the root user for daily administration tasks.

Correct Answer: B

Community vote distribution

B (100%)

Comments

BillaRanga **Highly Voted** 1 year, 9 months ago

Selected Answer: B

"As a best practice, do not use the AWS account root user for any task where it's not required. Instead, create a new IAM user for each person that requires administrator access."

It's B :)

upvoted 12 times

Andy_09 **Highly Voted** 1 year, 10 months ago

Option B

upvoted 8 times

d401c0d **Most Recent** 1 year, 7 months ago

Selected Answer: B

D is just killing me. If we have reached this far, we all know it is Option B - "As a best practice, do not use the AWS account root user for any task where it's not required. Instead, create a new IAM user for each person that requires administrator access."

upvoted 2 times

Sivaeas 1 year, 9 months ago

Selected Answer: B

its option B

upvoted 3 times

Naveena_Devanga 1 year, 9 months ago

Segregation of roles, also known as separation of duties (SoD), is a business control that helps prevent security or privacy incidents and errors. Therefore, root access must never be used for routine operational activities.

upvoted 2 times

A company is deploying an application that processes streaming data in near-real time. The company plans to use Amazon EC2 instances for the workload. The network architecture must be configurable to provide the lowest possible latency between nodes.

Which combination of network solutions will meet these requirements? (Choose two.)

- A. Enable and configure enhanced networking on each EC2 instance. **Most Voted**
- B. Group the EC2 instances in separate accounts.
- C. Run the EC2 instances in a cluster placement group. **Most Voted**
- D. Attach multiple elastic network interfaces to each EC2 instance.
- E. Use Amazon Elastic Block Store (Amazon EBS) optimized instance types.

Correct Answer: AC

Community vote distribution

AC (100%)

Comments

mestule **Highly Voted** 1 year, 10 months ago

Selected Answer: AC

A. Enable and configure enhanced networking on each EC2 instance. Enhanced networking provides higher bandwidth, higher packet per second (PPS) performance, and consistently lower inter-instance latencies.

C. Run the EC2 instances in a cluster placement group. A cluster placement group is a logical grouping of instances within a single Availability Zone. This configuration is recommended for applications that need low network latency, high network throughput, or both.

upvoted 13 times

Sivaeas **Highly Voted** 1 year, 9 months ago

Selected Answer: AC

To reach speeds up to 10 Gbps between instances, launch your instances in a cluster placement group with the enhanced networking instance type. These instance types are placed physically close to each other. Instance types that are close to each other further reduces latency and improves transfer speeds.

upvoted 5 times

FlyingHawk **Most Recent** 1 year, 5 months ago

Selected Answer: AC

A: <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/enhanced-networking.html>

C: https://aws.amazon.com/ec2/instance-types/#Instance_Features

Select EC2 instances support cluster networking when launched into a common cluster placement group. A cluster placement group provides low-latency networking between all instances in the cluster. The bandwidth an EC2 instance can utilize depends on the instance type and its networking performance specification.

upvoted 1 times

AmirBe 1 year, 8 months ago

AC

Use of Placement Groups: Utilize EC2 Placement Groups to ensure that instances are physically located close to each other within the same Availability Zone. This reduces the latency between instances by minimizing the distance data needs to travel.

Selection of EC2 Instance Types: Choose EC2 instance types optimized for low-latency networking, such as instances with enhanced networking capabilities like Elastic Network Adapter (ENA) or instances that support Amazon EC2 Nitro System. These instances provide high throughput and low latency networking performance.

upvoted 4 times

osmk 1 year, 9 months ago

what's AM?

upvoted 1 times

jaswantn 1 year, 10 months ago

option C & E.

Option A is not viable as ..EC2 provides enhanced networking capabilities using single root I/O virtualization (SR-IOV) only on supported instance types.

upvoted 1 times

jaswantn 1 year, 10 months ago

option E... EBS-optimized instance uses an optimized configuration

upvoted 1 times

sandordini 1 year, 7 months ago

C'mon, EBS is storage. The question does not deal with the storage solutions. Its a distractor...

upvoted 2 times

FlyingHawk 1 year, 5 months ago

Enhanced networking is available on most modern instance families, such as C5, M5, R5, P3, G4, and others.

Always check the AWS documentation <https://aws.amazon.com/ec2/instance-types/> for the latest list of supported instance types:

upvoted 1 times

Andy_09 1 year, 10 months ago

Correct option should be CD

upvoted 1 times

A financial services company wants to shut down two data centers and migrate more than 100 TB of data to AWS. The data has an intricate directory structure with millions of small files stored in deep hierarchies of subfolders. Most of the data is unstructured, and the company's file storage consists of SMB-based storage types from multiple vendors. The company does not want to change its applications to access the data after migration.

What should a solutions architect do to meet these requirements with the LEAST operational overhead?

- A. Use AWS Direct Connect to migrate the data to Amazon S3.
- B. Use AWS DataSync to migrate the data to Amazon FSx for Lustre.
- C. Use AWS DataSync to migrate the data to Amazon FSx for Windows File Server. **Most Voted**
- D. Use AWS Direct Connect to migrate the data on-premises file storage to an AWS Storage Gateway volume gateway.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Sivaeas **Highly Voted** 2 years, 3 months ago

Selected Answer: C

AWS DataSync is a data transfer service that simplifies, automates, and accelerates moving and replicating data between on-premises storage systems and AWS storage services over the internet or AWS Direct Connect. DataSync can transfer your file data, and also file system metadata such as ownership, time stamps, and access permissions.

In DataSync, a location for Amazon FSx for Windows is an endpoint for an FSx for Windows File Server. You can transfer files between a location for Amazon FSx for Windows and a location for other file systems. For information, see Working with Locations in the AWS DataSync User Guide.

DataSync accesses your FSx for Windows File Server using the Server Message Block (SMB) protocol.

upvoted 6 times

aditinand 2 years ago

Does Fsx support SMB? I read prior posts that it only ONTAP supports SMB

upvoted 1 times

Andy_09 **Highly Voted** 2 years, 4 months ago

Option C

upvoted 6 times

Scheldon **Most Recent** 1 year, 11 months ago

Selected Answer: C

AnswerC

upvoted 2 times

Linuslin 2 years ago

Selected Answer: C

AWS Storage Gateway provides a standard set of storage protocols such as iSCSI, SMB, and NFS, which allow you to use AWS storage without rewriting your existing applications.--->A is not complete describe, so A is out.
<https://aws.amazon.com/storagegateway/faqs/?nc=sn&loc=6>

Only FSx for NetApp ONTAP and FSx for Windows File Server support SMB Protocol. --->B is out.
<https://aws.amazon.com/tw/fsx/when-to-choose-fsx/>

AWS Direct Connect is more expensive than AWS DataSync.--->D is out

C is the correct answer.

upvoted 2 times

Naveena_Devanga 2 years, 3 months ago

Correct Answer is C

As most of the data is unstructured, and the company's file storage consists of SMB-based storage types from multiple vendors which is commonly a Windows-Linux file-sharing type so FSx for Windows File Server file systems completely meets the solution.

upvoted 3 times

ogerber 2 years, 3 months ago

Selected Answer: C

Option C since its SMB (windows) , and low operational effort so DataSync over Direct Connect

upvoted 4 times

osmk 2 years, 3 months ago

Selected Answer: C

<https://docs.aws.amazon.com/datasync/latest/userguide/create-fsx-location.html>

upvoted 3 times


A company uses an organization in AWS Organizations to manage AWS accounts that contain applications. The company sets up a dedicated monitoring member account in the organization. The company wants to query and visualize observability data across the accounts by using Amazon CloudWatch.

Which solution will meet these requirements?

- A. Enable CloudWatch cross-account observability for the monitoring account. Deploy an AWS CloudFormation template provided by the monitoring account in each AWS account to share the data with the monitoring account. **Most Voted**
- B. Set up service control policies (SCPs) to provide access to CloudWatch in the monitoring account under the Organizations root organizational unit (OU).
- C. Configure a new IAM user in the monitoring account. In each AWS account, configure an IAM policy to have access to query and visualize the CloudWatch data in the account. Attach the new IAM policy to the new IAM user.
- D. Create a new IAM user in the monitoring account. Create cross-account IAM policies in each AWS account. Attach the IAM policies to the new IAM user.

Correct Answer: A

Community vote distribution



Comments

jaswantn **Highly Voted** 2 years, 3 months ago

option A

below are the links to check both parts of option A.

https://docs.amazonaws.cn/en_us/AmazonCloudWatch/latest/monitoring/cloudwatch_crossaccount_dashboard.html

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/CloudWatch-Unified-Cross-Account-Setup.html#Unified-Cross-Account-SetupSource-SingleTemplate>

upvoted 7 times

MatAlves 1 year, 8 months ago

"If you have multiple Amazon accounts, you can set up CloudWatch cross-account observability and then create rich cross-account dashboards in your monitoring accounts. You can seamlessly search, visualize, and analyze your metrics, logs, and traces without account boundaries."

Nice catch!

upvoted 4 times

Sivaeas **Highly Voted** 2 years, 3 months ago

Selected Answer: A

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/CloudWatch-Unified-Cross-Account.html>

upvoted 6 times

ninasgx **Most Recent** 2 years, 3 months ago

Selected Answer: C

It's C

upvoted 1 times

osmk 2 years, 3 months ago

Selected Answer: A

https://docs.amazonaws.cn/en_us/AmazonCloudWatch/latest/monitoring/cloudwatch_crossaccount_dashboard.html

upvoted 4 times

Andy_09 2 years, 4 months ago

Option A
upvoted 3 times

A company's website is used to sell products to the public. The site runs on Amazon EC2 instances in an Auto Scaling group behind an Application Load Balancer (ALB). There is also an Amazon CloudFront distribution, and AWS WAF is being used to protect against SQL injection attacks. The ALB is the origin for the CloudFront distribution. A recent review of security logs revealed an external malicious IP that needs to be blocked from accessing the website.

What should a solutions architect do to protect the application?

- A. Modify the network ACL on the CloudFront distribution to add a deny rule for the malicious IP address.
- B. Modify the configuration of AWS WAF to add an IP match condition to block the malicious IP address. **Most Voted**
- C. Modify the network ACL for the EC2 instances in the target groups behind the ALB to deny the malicious IP address.
- D. Modify the security groups for the EC2 instances in the target groups behind the ALB to deny the malicious IP address.

Correct Answer: B

Community vote distribution

B (87%) 13%

Comments

Andy_09 **Highly Voted** 2 years, 4 months ago

Option B

upvoted 12 times

mohammadthainat **Highly Voted** 2 years, 2 months ago

Selected Answer: B

in WAF you can define Web ACL (Web Access Control List) Rule:

IP Set: up to 10,000 IP addresses – use multiple Rules for more IPs

upvoted 5 times

FlyingHawk **Most Recent** 1 year, 5 months ago

Selected Answer: B

Network ACL is with VPC at subnet level, A and C are incorrect. security group for EC2 only supports allow rule, not deny rule, D is out. B is correct: <https://docs.aws.amazon.com/waf/latest/developerguide/waf-rule-statement-type-ipset-match.html>

upvoted 2 times

Saliigen 1 year, 5 months ago

Selected Answer: B

<https://aws.amazon.com/blogs/security/how-to-enhance-amazon-cloudfront-origin-security-with-aws-waf-and-aws-secrets-manager/>

upvoted 1 times

bujuman 2 years, 1 month ago

Selected Answer: B

There was option from distribution Security Tab ==> Request logs for the specified time range where someone could target an IP address and block it - which action won't do more than creating a block rule under the associated Web ACL - but function has vanished, i don't ask me why.

So the only feasible option in WEBACLv2 is to go for an Ipset and add a WebACL ip match block condition.

I really liked the option A the first time i experimented it.

upvoted 4 times

xBUGx 2 years, 2 months ago

Selected Answer: A

You only need to block an IP. And Cloudfront is the first layer

upvoted 3 times

JA2018 1 year, 6 months ago

Modifying the CloudFront network ACL would block access at the CDN level. This could impact legitimate users globally, not just the malicious IP.

upvoted 1 times

JA2018 1 year, 6 months ago

When dealing with web application security concerns, leverage the capabilities of AWS WAF to block malicious IPs rather than modifying network access controls at lower levels like EC2 security groups or network ACLs

upvoted 1 times

FlyingHawk 1 year, 5 months ago

CloudFront does not support network ACLs. Network ACLs are used with Amazon VPCs at the subnet level, not with CloudFront distributions.

<https://docs.aws.amazon.com/vpc/latest/userguide/vpc-network-acls.html>

upvoted 1 times

Sivaeas 2 years, 3 months ago

Selected Answer: B

The AWS WAF IP set match statement inspects the IP address of a web request against a set of IP addresses and address ranges. Use this to allow or block web requests based on the IP addresses that the requests originate from

upvoted 4 times

stephensimudemy 2 years, 3 months ago

Selected Answer: B

Option B

upvoted 4 times

A company sets up an organization in AWS Organizations that contains 10 AWS accounts. A solutions architect must design a solution to provide access to the accounts for several thousand employees. The company has an existing identity provider (IdP). The company wants to use the existing IdP for authentication to AWS.

Which solution will meet these requirements?

- A. Create IAM users for the employees in the required AWS accounts. Connect IAM users to the existing IdP. Configure federated authentication for the IAM users.
- B. Set up AWS account root users with user email addresses and passwords that are synchronized from the existing IdP.
- C. Configure AWS IAM Identity Center (AWS Single Sign-On). Connect IAM Identity Center to the existing IdP. Provision users and groups from the existing IdP. **Most Voted**
- D. Use AWS Resource Access Manager (AWS RAM) to share access to the AWS accounts with the users in the existing IdP.

Correct Answer: C

Community vote distribution

C (100%)

Comments

osmk **Highly Voted** 1 year, 9 months ago

c--> Regardless of how you provision users, IAM Identity Center redirects the AWS Management Console, command line interface, and application authentication to your external IdP. IAM Identity Center then grants access to those resources based on policies you create in IAM Identity Center <https://docs.aws.amazon.com/singlesignon/latest/userguide/manage-your-identity-source-idp.html#provisioning-when-external-idp>
upvoted 9 times

ogerber **Most Recent** 1 year, 9 months ago

Selected Answer: C

Option C
<https://docs.aws.amazon.com/singlesignon/latest/userguide/manage-your-identity-source-idp.html>
upvoted 4 times

osmk 1 year, 9 months ago

Selected Answer: C

<https://docs.aws.amazon.com/singlesignon/latest/userguide/manage-your-identity-source-idp.html#provisioning-when-external-idp>
upvoted 3 times

Andy_09 1 year, 10 months ago

Option C
upvoted 3 times

A solutions architect is designing an AWS Identity and Access Management (IAM) authorization model for a company's AWS account. The company has designated five specific employees to have full access to AWS services and resources in the AWS account.

The solutions architect has created an IAM user for each of the five designated employees and has created an IAM user group.

Which solution will meet these requirements?

- A. Attach the AdministratorAccess resource-based policy to the IAM user group. Place each of the five designated employee IAM users in the IAM user group.
- B. Attach the SystemAdministrator identity-based policy to the IAM user group. Place each of the five designated employee IAM users in the IAM user group.
- C. Attach the AdministratorAccess identity-based policy to the IAM user group. Place each of the five designated employee IAM users in the IAM user group. **Most Voted**
- D. Attach the SystemAdministrator resource-based policy to the IAM user group. Place each of the five designated employee IAM users in the IAM user group.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Scheldon **Highly Voted** 1 year, 11 months ago

Selected Answer: C

AnswerC

We need identity-based policy and if we will compare System Admin and Administrator Access policy it clear that SysAdmin have is allowing for limited amount of actions, where Admin Access simple allow for all actions.

https://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies_identity-vs-resource.html

<https://docs.aws.amazon.com/aws-managed-policy/latest/reference/AdministratorAccess.html>

<https://docs.aws.amazon.com/aws-managed-policy/latest/reference/SystemAdministrator.html>

upvoted 7 times

Linuslin **Highly Voted** 2 years ago

Selected Answer: C

The question says "full access to AWS services and resources in the AWS account" and "created an IAM user group." You can see it is identity-based policy, not resource-based.--->A and D are out.

SystemAdministrator: Allow 28 of 412 services.--->B is out.

AdministratorAccess: Allow 412 of 412 services.--->C is the correct answer.

If you are curious about what a policy can allow for, just log in you AWS account and go to IAM-policies to find out.

upvoted 7 times

NSA_Poker **Most Recent** 1 year, 12 months ago

Selected Answer: C

(A & D) eliminated. Resource-based policies are attached to a resource NOT an IAM user, group, or role.

(B) eliminated. SystemAdministrator has fewer permissions than AdministratorAccess.

upvoted 5 times

MattBJ 2 years, 3 months ago

Selected Answer: C

C is the correct answer
upvoted 2 times

osmk 2 years, 3 months ago

Selected Answer: C

C > > > https://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies_manage-attach-detach.html
upvoted 3 times

osmk 2 years, 3 months ago

C > > > https://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies_manage-attach-detach.html
upvoted 3 times

Umuntu 2 years, 4 months ago

C looks correct
upvoted 3 times

Andy_09 2 years, 4 months ago

Option C
upvoted 4 times

A company has a multi-tier payment processing application that is based on virtual machines (VMs). The communication between the tiers occurs asynchronously through a third-party middleware solution that guarantees exactly-once delivery.

The company needs a solution that requires the least amount of infrastructure management. The solution must guarantee exactly-once delivery for application messaging.

Which combination of actions will meet these requirements? (Choose two.)

- A. Use AWS Lambda for the compute layers in the architecture. **Most Voted**
- B. Use Amazon EC2 instances for the compute layers in the architecture.
- C. Use Amazon Simple Notification Service (Amazon SNS) as the messaging component between the compute layers.
- D. Use Amazon Simple Queue Service (Amazon SQS) FIFO queues as the messaging component between the compute layers. **Most Voted**
- E. Use containers that are based on Amazon Elastic Kubernetes Service (Amazon EKS) for the compute layers in the architecture.

Correct Answer: AD

Community vote distribution

AD (100%)

Comments

osmk **Highly Voted** 2 years, 3 months ago

Selected Answer: AD

<https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/FIFO-queues-exactly-once-processing.html>

upvoted 6 times

Scheldon **Highly Voted** 1 year, 11 months ago

Selected Answer: AD

AnswerAD,

SQS FIFO will guarantee exactly one time execution for each operation. The problem is with processing as we do not know if whole process will be closed in 15 min (TTL for Lambda).

I'm choosing Lambda as it is natural thing for Payment procesing in AWS but I'm not 100% sure

upvoted 5 times

tch **Most Recent** 1 year, 5 months ago

Selected Answer: AD

SQS is helpful for Asynchronous processing workflows.

upvoted 1 times

MatAlves 1 year, 8 months ago

Selected Answer: AD

"As its name suggests, exactly-once semantics means that each message is delivered precisely once. The message can neither be lost nor delivered twice (or more times)."

- SNS doesn't provide exactly-once delivery. Thus, we need SQS.

- To achieve "least amount of intra management", we go with Lambda for compute layer.

upvoted 2 times

Sivaeas 2 years, 3 months ago

Selected Answer: AD

Lambda+SQS FIFO

upvoted 4 times

PolarFox 2 years, 3 months ago

someone please explain why the combination of D and E is not the correct?

upvoted 2 times

stephensimudemy 2 years, 3 months ago

because qn says 'least amount of infrastructure management'.

E is not.

upvoted 2 times

JA2018 1 year, 6 months ago

While EKS can provide scalability, it still requires managing container infrastructure. This adds complexity (more overhead in infra mgmt than A) to the solution

upvoted 1 times

jaswantn 2 years, 3 months ago

option A for payment processing.

option D for exactly once delivery.

upvoted 2 times

Umntu 2 years, 4 months ago

CD IS THE BEST ANSWER

upvoted 1 times

hajra313 2 years, 4 months ago

a and d

upvoted 3 times


A company has a nightly batch processing routine that analyzes report files that an on-premises file system receives daily through SFTP. The company wants to move the solution to the AWS Cloud. The solution must be highly available and resilient. The solution also must minimize operational effort.

Which solution meets these requirements?

- A. Deploy AWS Transfer for SFTP and an Amazon Elastic File System (Amazon EFS) file system for storage. Use an Amazon EC2 instance in an Auto Scaling group with a scheduled scaling policy to run the batch operation. **Most Voted**
- B. Deploy an Amazon EC2 instance that runs Linux and an SFTP service. Use an Amazon Elastic Block Store (Amazon EBS) volume for storage. Use an Auto Scaling group with the minimum number of instances and desired number of instances set to 1.
- C. Deploy an Amazon EC2 instance that runs Linux and an SFTP service. Use an Amazon Elastic File System (Amazon EFS) file system for storage. Use an Auto Scaling group with the minimum number of instances and desired number of instances set to 1.
- D. Deploy AWS Transfer for SFTP and an Amazon S3 bucket for storage. Modify the application to pull the batch files from Amazon S3 to an Amazon EC2 instance for processing. Use an EC2 instance in an Auto Scaling group with a scheduled scaling policy to run the batch operation.

Correct Answer: A

Community vote distribution



Comments

Sivaeas **Highly Voted** 2 years, 3 months ago

Selected Answer: A

The Answer should be A not D because ...

Modify the application to pull the batch files from Amazon S3 to an Amazon EC2 instance for processing.--Why we need to do this when we can move the file directly to EFS in EC2 system

AWS Transfer Family now also supports file transfers to Amazon Elastic File System (Amazon EFS) file systems as well as Amazon S3.

upvoted 9 times

BBR01 **Highly Voted** 2 years, 1 month ago

Selected Answer: A

A should be enough. EFS can be mounted to ASG directly, and there is no need to use S3 in the middle.

upvoted 5 times

zdi561 **Most Recent** 1 year, 4 months ago

Selected Answer: A

A and D both work , but EFS is more natural file system for original batch routing, and could be faster than S3

upvoted 1 times

MatAlves 1 year, 8 months ago

Selected Answer: A

"... an on-premises file system receives" = file system > EFS.

SFT in AWS = AWS Transfer family

Even though D works, that would require changes in the current architecture.

upvoted 2 times

Scheldon 1 year, 11 months ago

Selected Answer: A

AnswerA

Option A and D will work, but taking into consideration requirements I would go with A
upvoted 2 times

Franjly 2 years ago

Selected Answer: A

File system = efs, fsx,
upvoted 2 times

7fb06b3 2 years ago

Selected Answer: D

Option D, I'm not 100% sure.. Always prefer S3 over EFS
upvoted 2 times

MatAlves 1 year, 8 months ago

"... an on-premises file system receives" = file system > EFS.
upvoted 2 times

JackyCCK 2 years, 2 months ago

I think the ans is A as well, option D require "Modify the application" which is not "minimize operational effort"
upvoted 4 times

khoantd 2 years, 2 months ago

Selected Answer: C

Option D
upvoted 1 times

Ipergorta 2 years, 2 months ago

Option D
upvoted 1 times

PolarFox 2 years, 3 months ago

Selected Answer: D

transfer + S3 = HA, scheduled scaling = resilient
upvoted 3 times

NayeraB 2 years, 3 months ago

Selected Answer: D

I'm not 100% sure, but D looks like the right flow to me
upvoted 1 times

TwinSpark 2 years ago

I agree on a professional level, that's will make the company save money, that is the only things company care about. But for the exam i will go for A
upvoted 2 times

osmk 2 years, 3 months ago

Selected Answer: A

The service is designed to be highly scalable, highly available, and highly durable. Amazon EFS offers the following file system types to meet your availability and durability needs
-> <https://docs.aws.amazon.com/efs/latest/ug/whatisefs.html>
Amazon S3 achieves high availability by replicating data across multiple servers within AWS data centers-
-> <https://docs.aws.amazon.com/AmazonS3/latest/userguide/Welcome.html>
upvoted 2 times

NayeraB 2 years, 3 months ago

But option A doesn't address the need for the application to pull the batch jobs from the new storage, also is the use of EFS needed here? In terms of it being a shared storage and whatnot..
upvoted 2 times

osmk 2 years, 3 months ago

A>>>>
upvoted 1 times

osmk 2 years, 3 months ago

The service is designed to be highly scalable, highly available, and highly durable. Amazon EFS offers the following file system types to meet your availability and durability needs

-><https://docs.aws.amazon.com/efs/latest/ug/whatisefs.html>

Amazon S3 achieves high availability by replicating data across multiple servers within AWS data centers-

><https://docs.aws.amazon.com/AmazonS3/latest/userguide/Welcome.html>

upvoted 3 times

Andy_09 2 years, 4 months ago

Option D

upvoted 3 times

A company has users all around the world accessing its HTTP-based application deployed on Amazon EC2 instances in multiple AWS Regions. The company wants to improve the availability and performance of the application. The company also wants to protect the application against common web exploits that may affect availability, compromise security, or consume excessive resources. Static IP addresses are required.

What should a solutions architect recommend to accomplish this?

- A. Put the EC2 instances behind Network Load Balancers (NLBs) in each Region. Deploy AWS WAF on the NLBs. Create an accelerator using AWS Global Accelerator and register the NLBs as endpoints.
- B. Put the EC2 instances behind Application Load Balancers (ALBs) in each Region. Deploy AWS WAF on the ALBs. Create an accelerator using AWS Global Accelerator and register the ALBs as endpoints. **Most Voted**
- C. Put the EC2 instances behind Network Load Balancers (NLBs) in each Region. Deploy AWS WAF on the NLBs. Create an Amazon CloudFront distribution with an origin that uses Amazon Route 53 latency-based routing to route requests to the NLBs.
- D. Put the EC2 instances behind Application Load Balancers (ALBs) in each Region. Create an Amazon CloudFront distribution with an origin that uses Amazon Route 53 latency-based routing to route requests to the ALBs. Deploy AWS WAF on the CloudFront distribution.

Correct Answer: B

Community vote distribution

B (85%)

D (15%)

Comments

ogerber **Highly Voted** 1 year, 9 months ago

Selected Answer: B

HTTP based application so ALB is required.

because static IP addresses are required, we should use global accelerator:

"By default, Global Accelerator provides you with static IP addresses that you associate with your accelerator."

upvoted 15 times

Andy_09 **Highly Voted** 1 year, 10 months ago

Option D

upvoted 5 times

Typewriter101 1 year, 9 months ago

Why D cause i think global accelerator will do a better job an cloudfront to increase availability and performance

upvoted 2 times

Typewriter101 1 year, 9 months ago

than cloudfront*

upvoted 1 times

In2718 **Most Recent** 10 months, 2 weeks ago

Selected Answer: D

The best answer is now D - this is an HTTP based app and Cloudfront currently supports static IPs - <https://aws.amazon.com/about-aws/whats-new/2024/11/amazon-cloudfront-anycast-static-ips/> Either this is a trick question - or it is now outdated and needs revision.

upvoted 1 times

f07ed8f 1 year, 6 months ago

Selected Answer: B

CloudFront doesn't support assigning a static IP address to distributions

upvoted 3 times

zinabu 1 year, 7 months ago

Selected Answer: B

Http based app=ALB

static IP= AWS global accelerator

for those who choose "A" NLB doesn't support Http based traffic it is just used for TCP/UDP based traffic.

upvoted 4 times

mohammadthainat 1 year, 8 months ago

Selected Answer: D

Something wrong in the question, here is why:

Static IP addresses are required --> NLB

protect against common web exploits --> WAF (But you can't use WAF directly with NLB)

HTTP-based application --> CloudFront (using CloudFront with NLB is not recommended)

EC2s in multiple AWS Regions --> Route 53 latency-based

upvoted 3 times

mohammadthainat 1 year, 8 months ago

Changing my answer to B

Static IP addresses are required --> We can use Global Accelerator for fixed IP and WAF on the ALB

upvoted 4 times

TruthWS 1 year, 8 months ago

Option A

Static IP --> NLB

against common web exploits --> WAF

performance --> Global Accelerator is best choice in this situation.

upvoted 1 times

dkw2342 1 year, 8 months ago

No, option B is correct.

* WAF (L7) does not work with NLB (L4)

* Traffic enters via the Global Accelerator, so that's the customer-facing (static) IP - <https://docs.aws.amazon.com/global-accelerator/latest/dg/about-accelerators.eip-accelerator.html>

upvoted 3 times

jackky3123213 1 year, 8 months ago

Selected Answer: D

Option D

upvoted 1 times

alawada 1 year, 8 months ago

Selected Answer: B

CloudFront uses multiple sets of dynamically changing IP addresses while Global Accelerator will provide you a set of static IP addresses as a fixed entry point to your applications

upvoted 2 times

lpergorta 1 year, 8 months ago

Option D

upvoted 1 times

Naveena_Devanga 1 year, 9 months ago

Correct Answer is C.

Static IP addresses are required specific to the requirement.

upvoted 1 times

stephensimudemy 1 year, 9 months ago

Selected Answer: B

CloudFront uses multiple sets of dynamically changing IP addresses while Global Accelerator will provide you a set of static IP addresses as a fixed entry point to your applications

upvoted 2 times

osmk 1 year, 9 months ago

Selected Answer: B

Network Load Balancer (NLB): NLB operates at layer 4 and does not support AWS WAF directly

<https://docs.aws.amazon.com/elasticloadbalancing/latest/application/introduction.html>

upvoted 2 times

osmk 1 year, 9 months ago

The company wants to improve the availability and performance of the application

upvoted 1 times

jaswantn 1 year, 9 months ago

Static IP addresses are required, so option B....global accelerator with ALB

upvoted 3 times

Dhokal 1 year, 10 months ago

B is correct

upvoted 3 times

A company's data platform uses an Amazon Aurora MySQL database. The database has multiple read replicas and multiple DB instances across different Availability Zones. Users have recently reported errors from the database that indicate that there are too many connections. The company wants to reduce the failover time by 20% when a read replica is promoted to primary writer.

Which solution will meet this requirement?

- A. Switch from Aurora to Amazon RDS with Multi-AZ cluster deployment.
- B. Use Amazon RDS Proxy in front of the Aurora database. **Most Voted**
- C. Switch to Amazon DynamoDB with DynamoDB Accelerator (DAX) for read connections.
- D. Switch to Amazon Redshift with relocation capability.

Correct Answer: B

Community vote distribution

B (100%)

Comments

osmk **Highly Voted** 1 year, 9 months ago

Selected Answer: B

By using Amazon RDS Proxy, your applications can pool and share database connections. This pooling improves scalability by allowing multiple application instances to reuse existing connections.

It also makes your applications more resilient to database failures. When a primary database instance fails, RDS Proxy automatically connects to a standby DB instance while preserving application connections. => <https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/rds-proxy.htm>

upvoted 10 times

Umuntu **Most Recent** 1 year, 10 months ago

Option B

upvoted 4 times

Andy_09 1 year, 10 months ago

Option B

upvoted 4 times

A company stores text files in Amazon S3. The text files include customer chat messages, date and time information, and customer personally identifiable information (PII).

The company needs a solution to provide samples of the conversations to an external service provider for quality control. The external service provider needs to randomly pick sample conversations up to the most recent conversation. The company must not share the customer PII with the external service provider. The solution must scale when the number of customer conversations increases.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create an Object Lambda Access Point. Create an AWS Lambda function that redacts the PII when the function reads the file. Instruct the external service provider to access the Object Lambda Access Point. **Most Voted**

B. Create a batch process on an Amazon EC2 instance that regularly reads all new files, redacts the PII from the files, and writes the redacted files to a different S3 bucket. Instruct the external service provider to access the bucket that does not contain the PII.

B. Create a web application on an Amazon EC2 instance that presents a list of the files, redacts the PII from the files, and allows the external service provider to download new versions of the files that have the PII redacted.

D. Create an Amazon DynamoDB table. Create an AWS Lambda function that reads only the data in the files that does not contain PII. Configure the Lambda function to store the non-PII data in the DynamoDB table when a new file is written to Amazon S3. Grant the external service provider access to the DynamoDB table.

Correct Answer: A

Community vote distribution

A (94%) 6%

Comments

osmk **Highly Voted** 2 years, 3 months ago

Selected Answer: A

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/tutorial-s3-object-lambda-redact-pii.html>

upvoted 13 times

Sergius95 2 years, 1 month ago

thanks

upvoted 2 times

Andy_09 **Highly Voted** 2 years, 4 months ago

Option A

upvoted 5 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: A

If you want to achieve the goal stipulated in the question, all 4 options work, but only option A minimizes operational overhead and ensures scalability while redacting PII dynamically.

upvoted 1 times

trinh_le 1 year, 11 months ago

Selected Answer: A

Use AWS Lambda functions to change the Object before it is retrieved by the caller application. Only one S3 bucket is needed, on top of which we create S3 Access Point And S3 Object Lambda Access Points

Use case:

1. Redact PII for analytics or non-production environment

2. Convert across data formats ex: XML to Json

3. Resizing and watermarking images on fly using caller-specific details ex: user who requested the object

upvoted 3 times

zinabu 2 years, 1 month ago

Selected Answer: D

Creating a RAID 0 array allows you to achieve a higher level of performance for a file system than you can provision on a single Amazon EBS volume. Use RAID 0 when I/O performance is of the utmost importance. With RAID 0, I/O is distributed across the volumes in a stripe.

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/raid-config.html>

upvoted 1 times

Vlad 2 years, 4 months ago

A is the correct choice.

upvoted 3 times

Umntu 2 years, 4 months ago

A is the best choice

upvoted 3 times

A company is running a legacy system on an Amazon EC2 instance. The application code cannot be modified, and the system cannot run on more than one instance. A solutions architect must design a resilient solution that can improve the recovery time for the system.

What should the solutions architect recommend to meet these requirements?

- A. Enable termination protection for the EC2 instance.
- B. Configure the EC2 instance for Multi-AZ deployment.
- C. Create an Amazon CloudWatch alarm to recover the EC2 instance in case of failure. **Most Voted**
- D. Launch the EC2 instance with two Amazon Elastic Block Store (Amazon EBS) volumes that use RAID configurations for storage redundancy.

Correct Answer: C

Community vote distribution

C (58%)

D (42%)

Comments

Andy_09 **Highly Voted** 2 years, 4 months ago

Option C

upvoted 11 times

Typewriter101 2 years, 3 months ago

i think D is the answer.

Cause the question asks for a resilient solution and EBS with RAID config can balance between the performance and redundancy. EBS can also help with faster launch.

upvoted 2 times

MatAlves 1 year, 8 months ago

But how does that "improve the recovery time for the system"?

upvoted 2 times

mavik 2 years, 3 months ago

Your solution can't resolve the problem

upvoted 3 times

mohammadthainat **Highly Voted** 2 years, 2 months ago

Selected Answer: C

Question about ""improve the recovery time for the system"" RAID improves data resilience, but won't recover the instance if the system itself fails. it's 100% C

upvoted 6 times

jacobtriestech **Most Recent** 8 months, 2 weeks ago

Selected Answer: C

The legacy app must run on a single instance and the code can't change, so you need fast recovery of that single instance if the underlying host or hardware fails. Amazon EC2 instance recovery (triggered by a CloudWatch alarm on the instance status checks) will automatically recover the instance onto healthy hardware while preserving the instance ID, attached EBS volumes, elastic IPs, and network configuration — delivering a much faster, transparent recovery than manual rebuilds.

upvoted 2 times

allanm 10 months ago

Selected Answer: D

Amazon EBS volume data is replicated across multiple servers in an Availability Zone to prevent the loss of data from the failure of any single component. This replication makes Amazon EBS volumes ten times more reliable than typical commodity disk drives.

<https://docs.aws.amazon.com/ebs/latest/userguide/raid-config.html>

upvoted 1 times

ln2718 10 months, 2 weeks ago

Selected Answer: D

Answer D is better because it is the closest thing to first level of prevention. Answer C is after the fact and so, yes you would have Cloudwatch as well but it only ask for one answer and the better one is D for a better first level response, and of course you would have Cloudwatch. Really bad question should have asked for two answers.

upvoted 1 times

buzzinmumbai 2 years, 2 months ago

Option should be B. They are not asking about storage anywhere. In multi-AZ you application runs on the primary and the secondary is kept in sync.

upvoted 1 times

dkw2342 2 years, 2 months ago

Pretty sure option D is NOT correct.

> RAID 5 and RAID 6 are not recommended for Amazon EBS (...).

> RAID 1 is also not recommended for use with Amazon EBS.

<https://docs.aws.amazon.com/ebs/latest/userguide/raid-config.html#raid-config-options>

upvoted 3 times

Awsbeginner87 2 years, 2 months ago

So what is the answer?

upvoted 2 times

sandordini 2 years, 1 month ago

C: You can create an Amazon CloudWatch alarm that monitors an Amazon EC2 instance and automatically recovers the instance if it becomes impaired due to an underlying hardware failure or a problem that requires AWS involvement to repair. Terminated instances cannot be recovered. A recovered instance is identical to the original instance, including the instance ID, private IP addresses, Elastic IP addresses, and all instance metadata.

upvoted 3 times

Salilgen 1 year, 5 months ago

IMO answer is B.

You could set Max desired capacity to 1 in autoscaling group then you always have one instance running

upvoted 1 times

Salilgen 1 year, 5 months ago

Sorry. Answer is C.

In multi-AZ deployment you must have at least one instance per AZ

upvoted 1 times

haci 2 years, 2 months ago

For those who choose C, the question asks that "must design a resilient solution".. C may improve recovery time but it has nothing to do with resiliency.

upvoted 1 times

JackyCCK 2 years, 2 months ago

"resilient solution that can improve the recovery time for the system", resiliency here means only

upvoted 2 times

mavik 2 years, 3 months ago

Option C

upvoted 2 times

stephensimudemy 2 years, 3 months ago

Selected Answer: C

Can only run 1 instance.

improve recovery time.

upvoted 3 times

stephensimudemy 2 years, 3 months ago

Option B.

Question never ask anything about storage.

upvoted 1 times

osmk 2 years, 3 months ago

Selected Answer: D

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/raid-config.html>

upvoted 6 times

A company wants to deploy its containerized application workloads to a VPC across three Availability Zones. The company needs a solution that is highly available across Availability Zones. The solution must require minimal changes to the application.

Which solution will meet these requirements with the LEAST operational overhead?

A. Use Amazon Elastic Container Service (Amazon ECS). Configure Amazon ECS Service Auto Scaling to use target tracking scaling. Set the minimum capacity to 3. Set the task placement strategy type to spread with an Availability Zone attribute.

Most Voted

B. Use Amazon Elastic Kubernetes Service (Amazon EKS) self-managed nodes. Configure Application Auto Scaling to use target tracking scaling. Set the minimum capacity to 3.

C. Use Amazon EC2 Reserved Instances. Launch three EC2 instances in a spread placement group. Configure an Auto Scaling group to use target tracking scaling. Set the minimum capacity to 3.

D. Use an AWS Lambda function. Configure the Lambda function to connect to a VPC. Configure Application Auto Scaling to use Lambda as a scalable target. Set the minimum capacity to 3.

Correct Answer: A

Community vote distribution

A (100%)

Comments

osmk **Highly Voted** 2 years, 3 months ago

Selected Answer: A

Amazon EKS self-managed nodes require you to manually install and configure the Kubernetes node components, such as kubelet, kube-proxy, and Docker, on your Amazon EC2 instances. You also need to manage the security group, IAM role, and subnet for your node group. Amazon ECS handles these tasks for you when you use the Amazon EC2 launch type .

upvoted 9 times

EdricHoang 1 year, 11 months ago

but it requires minimum change in application. I believe when changing to ECS, its a huge change

upvoted 1 times

ajwksldldfgdsg **Most Recent** 1 year, 9 months ago

Selected Answer: A

Containerized.. = ECS

upvoted 1 times

1dd 2 years, 3 months ago

why not lambda?

upvoted 1 times

khoahoang 2 years, 2 months ago

lamda dont have containerized

upvoted 2 times

Sergiuss95 2 years, 1 month ago

Containerized... The solution must require minimal changes to the application.

upvoted 1 times

Andy_09 2 years, 4 months ago

Option A

upvoted 2 times

A media company stores movies in Amazon S3. Each movie is stored in a single video file that ranges from 1 GB to 10 GB in size.


The company must be able to provide the streaming content of a movie within 5 minutes of a user purchase. There is higher demand for movies that are less than 20 years old than for movies that are more than 20 years old. The company wants to minimize hosting service costs based on demand.

Which solution will meet these requirements?

- A. Store all media content in Amazon S3. Use S3 Lifecycle policies to move media data into the Infrequent Access tier when the demand for a movie decreases.
- B. Store newer movie video files in S3 Standard. Store older movie video files in S3 Standard-Infrequent Access (S3 Standard-IA). When a user orders an older movie, retrieve the video file by using standard retrieval. **Most Voted**
- C. Store newer movie video files in S3 Intelligent-Tiering. Store older movie video files in S3 Glacier Flexible Retrieval. When a user orders an older movie, retrieve the video file by using expedited retrieval.
- D. Store newer movie video files in S3 Standard. Store older movie video files in S3 Glacier Flexible Retrieval. When a user orders an older movie, retrieve the video file by using bulk retrieval.

Correct Answer: B

Community vote distribution



Comments

Freddie26 **Highly Voted** 2 years, 3 months ago

Technically, expedited retrieval for files is not guaranteed within 1-5 minutes for files larger than 250 MB+. See <https://docs.aws.amazon.com/AmazonS3/latest/userguide/restoring-objects-retrieval-options.html>.

upvoted 17 times

osmk **Highly Voted** 2 years, 3 months ago

Selected Answer: B

S3 Standard-IA is for data that is accessed less frequently, but requires rapid access when needed. S3 Standard-IA offers the high durability, high throughput, and low latency of S3 Standard, with a low per GB storage price and per GB retrieval charge <https://aws.amazon.com/s3/storage-classes/>

upvoted 12 times

JustEugen **Most Recent** 8 months, 1 week ago

Selected Answer: B

It is B because:

- An AWS Glacier is used for archival purposes, which is not our case.
- We have a known pattern, so there are no reason to use intelligent tearing

upvoted 1 times

tch 1 year, 3 months ago

Selected Answer: C

S3 Standard-Infrequent Access (S3 Standard-IA) do not come with retrieval options??

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/restoring-objects-retrieval-options.html>

upvoted 1 times

zdi561 1 year, 4 months ago

Selected Answer: C

Glacier flex expedited retrieval takes 1-5 min.

upvoted 1 times

FlyingHawk 1 year, 4 months ago

Selected Answer: B

S3 Glacier Flexible Retrieval cannot grant the retrieval within 5 minutes even with expedited retrieval. Intelligent-Tiering is for unknown pattern, however, we already know higher demand for movies that are less than 20 years old than for movies that are more than 20 years old.

upvoted 1 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: B

A - Cost optimization for > 20-yr-old files is insufficient.

B - Good. Although, S3 Standard-IA is still relatively expensive compared to archival classes.

C - Glacier Flexible Retrieval may not consistently meet the 5-minute streaming requirement. Could be minutes, could be hours.

D - Bulk retrieval typically takes hours, let alone 5 minutes.

upvoted 3 times

tch 1 year, 2 months ago

this is the reason we choose C as answer! S3 Glacier Flexible Retrieval has option with 5 minutes time called expedited retrieval.

S3 Standard-IA do not come with option to retrieve data.... please check

<https://aws.amazon.com/s3/storage-classes/glacier/>

upvoted 1 times

Arad 1 year, 6 months ago

Selected Answer: C

C is the correct answer.

upvoted 1 times

Abdullah2004 1 year, 9 months ago

A is most correct ✓

upvoted 2 times

EdricHoang 1 year, 11 months ago

Selected Answer: B

"Expedited Retrieval, you can retrieve small amounts of data (up to 250 MB per request) within 1-5 minutes."

Cannot C

upvoted 8 times

NSA_Poker 1 year, 12 months ago

Selected Answer: B

(A) is eliminated. Demand metric is popularly based & cannot be configured with Lifecycle policies. Ex: an old movie can have resurgent demand 20 years after its sequel is released.

(C) is eliminated. Expedited retrieval is for all but the largest archived objects (250 MB+).

(D) is eliminated. Bulk retrieval takes hours.

(B) is more expensive than S3 Glacier Flexible Retrieval but it's the only one that works.

upvoted 6 times

bujuman 2 years, 1 month ago

Selected Answer: C

AS the pattern is uncertain- Customer could not, in advance, segregate data, the pattern will be determined on the fly - and with regard of the following S3 feature:

S3 Intelligent-Tiering is an additional storage class that provides flexibility for data with unknown or changing access patterns. It automates the movement of your objects between storage classes to optimize cost.

C will be the most cost effective for this use case.

upvoted 1 times

bujuman 2 years, 1 month ago

More insight:

S3 Glacier Flexible Retrieval for the most flexible retrieval options that balance cost with access times ranging from minutes to hours. Your retrieval options permit you to access all the archives you need, when you need them, for one low storage price. This storage class comes with multiple retrieval options:

- Expedited retrievals (restore in 1–5 minutes)

- Standard retrievals (restore in 3–5 hours)

- Bulk retrievals (restore in 5–12 hours). Bulk retrievals are available at no additional charge

upvoted 3 times

Sergiuss95 2 years, 1 month ago

Selected Answer: C

Expedited 1-5min and for new files intelligent tier is a good option

upvoted 2 times

mohammadthainat 2 years, 2 months ago

Selected Answer: C

All old files should be in--> Glacier Flexible Retrieval takes (1-5 minutes) to retrieve the file.
New files should not stay in Standard Storage class forever --> Intelligent-Tiering

upvoted 3 times

JackyCCK 2 years, 2 months ago

I don't think C is an option, S3 Glacier Flexible takes hour to retrieve the data.
Option A is actually valid, but the way the option A describe it does not consider "demand patterns based on time"

So it should be B

upvoted 2 times

JackyCCK 2 years, 2 months ago

expedited retrieval should not be used in that way as well

upvoted 1 times

Drew3000 2 years, 2 months ago

Selected Answer: A

There is something I like about option A. It's the only one that deals with what happens with a movie that goes from "new" to "old". With other options, new movies will be new forever.

upvoted 2 times

dkw2342 2 years, 2 months ago

Option B makes the most sense.

Why not option C:

1. This is not an archival use case, the company runs a video streaming service, so objects are still accessed regularly. Accelerated Retrieval is designed for "occasional urgent requests for a subset of archives".
2. The 5 minute timeframe does not apply to items of 250+ MB.
3. Even if the timeframe were valid, it's not guaranteed ("typically")
4. Expedited retrieval is expensive if used frequently (\$10.00 per 1,000 requests) - depending on access patterns, this may more than offset the savings in storage costs.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/restoring-objects-retrieval-options.html>

upvoted 4 times

A solutions architect needs to design the architecture for an application that a vendor provides as a Docker container image. The container needs 50 GB of storage available for temporary files. The infrastructure must be serverless.

Which solution meets these requirements with the LEAST operational overhead?

- A. Create an AWS Lambda function that uses the Docker container image with an Amazon S3 mounted volume that has more than 50 GB of space.
- B. Create an AWS Lambda function that uses the Docker container image with an Amazon Elastic Block Store (Amazon EBS) volume that has more than 50 GB of space.
- C. Create an Amazon Elastic Container Service (Amazon ECS) cluster that uses the AWS Fargate launch type. Create a task definition for the container image with an Amazon Elastic File System (Amazon EFS) volume. Create a service with that task definition. **Most Voted**
- D. Create an Amazon Elastic Container Service (Amazon ECS) cluster that uses the Amazon EC2 launch type with an Amazon Elastic Block Store (Amazon EBS) volume that has more than 50 GB of space. Create a task definition for the container image. Create a service with that task definition.

Correct Answer: C

Community vote distribution

C (95%)

5%

Comments

stephensimudemey **Highly Voted** 2 years, 3 months ago

Selected Answer: C

Options A and B involve AWS Lambda, which is suitable for event-driven, short-lived compute tasks, but it's NOT ideal for long-running containerized applications and managing large volumes of data.

upvoted 9 times

Andy_09 **Highly Voted** 2 years, 4 months ago

Option C

upvoted 6 times

nj1999 2 years, 4 months ago

Why C and not B?

upvoted 1 times

hajra313 2 years, 4 months ago

the infrastructure must be serverless

upvoted 2 times

Cali182 2 years, 4 months ago

Creating an AWS Lambda function that uses the Docker container image with an Amazon S3 mounted volume might not be suitable because Lambda functions have limitations on execution duration (15 minutes) and storage size (maximum 512 MB in the /tmp directory).

upvoted 5 times

dkw2342 2 years, 2 months ago

There's no indication of runtime, so that's not the reason.

A is wrong because "S3 volumes" do not exist. If the question were about S3 buckets: while it is possible to mount an S3 bucket using FUSE, this is completely unsupported by AWS and definitely won't work in a container running on Lambda (you can't assign SYS_ADMIN cap and mount /dev/fuse).

B is wrong because you can't use EBS volumes with Lambda.

As an aside, Lambda supports up to 10 GB of ephemeral storage (configurable).

upvoted 3 times

03beafc 2 years, 1 month ago

Not B because your lambda container needs the RIC and the image is already provided, presumably without the RIC (or else it would have mentioned it)

upvoted 3 times

sandordini 2 years, 1 month ago

RIC: Runtime interface clients

upvoted 2 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: C

A - Technically, temporary file access in Lambda is limited to the /tmp directory with a 512 MB limit.

B - AWS Lambda does not natively support mounting EBS volumes.

C - Looks good to me.

D - ECS with EC2 is not serverless.

FYI - Serverless infrastructure: Solutions must not involve managing servers or underlying infrastructure. AWS Fargate and AWS Lambda are serverless options, while Amazon ECS with EC2 launch type involves managing EC2 instances.

upvoted 1 times

dragongoseki 1 year, 11 months ago

Selected Answer: C

C is right answer.

upvoted 2 times

DZRomero 1 year, 11 months ago

Selected Answer: C

The combination of ECS with Fargate and EFS (option C) provides a serverless solution that can run Docker containers and meet the storage requirements, all while minimizing operational overhead. You don't need to manage any servers, and the storage will automatically scale as needed. This makes it the best fit for the given requirements.

upvoted 2 times

sandordini 2 years, 1 month ago

Selected Answer: C

Lambda would need Runtime interface clients (RIC) to host a container workload.

Also Lambda storage limit: 10GB

Fargate is Serverless >> C

upvoted 4 times

zinabu 2 years, 1 month ago

Selected Answer: B

the key word here is {" serverless + temporary file"}

A: it uses S3 for storage that is not a temporary file storage system

C: that was good using ECS with fargate for serverless part but it uses EFS file system still it is a durable file system not temporary

D: Using EBS was good to use for temporary file system but it is mounted on EC2 which is not serverless. so that we are left with "B" which uses [lambda(serverless) + EBS(temporary storage)]

upvoted 1 times

zinabu 2 years, 1 month ago

the key word here is {" serverless + temporary file"}

A: it uses S3 for storage that is not a temporary file storage system

C: that was good using ECS with fargate for serverless part but it uses EFS file system still it is a durable file system not temporary

D: Using EBS was good to use for temporary file system but it is mounted on EC2 which is not serverless. so that we are left with "B" which uses [lambda(serverless) + EBS(temporary storage)]

upvoted 1 times

A company needs to use its on-premises LDAP directory service to authenticate its users to the AWS Management Console. The directory service is not compatible with Security Assertion Markup Language (SAML).

Which solution meets these requirements?

- A. Enable AWS IAM Identity Center (AWS Single Sign-On) between AWS and the on-premises LDAP.
- B. Create an IAM policy that uses AWS credentials, and integrate the policy into LDAP.
- C. Set up a process that rotates the IAM credentials whenever LDAP credentials are updated.
- D. Develop an on-premises custom identity broker application or process that uses AWS Security Token Service (AWS STS) to get short-lived credentials. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

kempes **Highly Voted** 2 years, 4 months ago

Selected Answer: D

The solution that best meets the requirements. This approach provides a pathway for authenticating LDAP users to AWS without requiring direct LDAP to AWS IAM Identity Center integration or SAML compatibility, offering a flexible and secure method to extend on-premises authentication mechanisms to AWS services.

upvoted 11 times

aditianand 2 years ago

Why not option A A. Enable AWS IAM Identity Center (AWS Single Sign-On) between AWS and the on-premises LDAP.

upvoted 1 times

NSA_Poker 1 year, 12 months ago

(A) is incorrect bc to use AWS IAM Identity Center (AWS Single Sign-On) with an external IdP, you need SAML.

upvoted 5 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: D

A - AWS IAM Identity Center (formerly AWS SSO) allows integration with SAML-compatible identity providers, which won't work with on-premises LDAP directory (not SAML-compatible).

B - There's no way to "integrate" an IAM policy into LDAP.

C - Too complex. AWS already provides better mechanisms for handling temporary credentials - AWS STS.

D - YES. Actually, this approach is commonly used when SAML is not an option.

upvoted 1 times

Scheldon 1 year, 11 months ago

Selected Answer: D

AnswerD

upvoted 2 times

ike001 1 year, 12 months ago

D is the answer

upvoted 2 times

NSA_Poker 1 year, 12 months ago

Selected Answer: D

Identity federation can be accomplished in one of three ways.

(1) Use a corporate IdP (such as Microsoft Active Directory) or a custom identity broker application. Each option uses AWS STS.

- (2) Create an integration that uses Security Assertion Markup Language (SAML).
(3) Use a web identity provider, such as Amazon Cognito.

upvoted 2 times

1e22522 1 year, 10 months ago

YEA SURE FED

upvoted 1 times

TwinSpark 2 years ago

Selected Answer: D

option D

As per described here:

https://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_common-scenarios_federated-users.html#id_roles_common-scenarios_federated-users-idbroker

option A is wrong because for use SSO need to be compatible with SAML (at least this is what i understand from here:

https://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_common-scenarios_federated-users.html#id_roles_common-scenarios_federated-users-saml20)

upvoted 2 times

Naveena_Devanga 2 years, 3 months ago

Option D

A custom identity broker application can be built to perform a similar function to an identity store that is not compatible with SAML 2.0. The broker application authenticates users, requests temporary credentials from AWS, and provides them to the user to access AWS resources.

upvoted 2 times

aditianand 2 years ago

Why not option A A. Enable AWS IAM Identity Center (AWS Single Sign-On) between AWS and the on-premises LDAP.

upvoted 1 times

allanm 10 months ago

A is not an option because IAM Identity center depends on SAML.

upvoted 1 times

jaswantn 2 years, 3 months ago

If your identity store is not compatible with SAML 2.0, then you can build a custom identity broker application to perform a similar function.option D

upvoted 2 times

A company stores multiple Amazon Machine Images (AMIs) in an AWS account to launch its Amazon EC2 instances. The AMIs contain critical data and configurations that are necessary for the company's operations. The company wants to implement a solution that will recover accidentally deleted AMIs quickly and efficiently.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create Amazon Elastic Block Store (Amazon EBS) snapshots of the AMIs. Store the snapshots in a separate AWS account.
- B. Copy all AMIs to another AWS account periodically.
- C. Create a retention rule in Recycle Bin. **Most Voted**
- D. Upload the AMIs to an Amazon S3 bucket that has Cross-Region Replication.

Correct Answer: C

Community vote distribution

C (100%)

Comments

alawada **Highly Voted** 1 year, 8 months ago

Selected Answer: C

Recycle Bin is a data recovery feature that enables you to restore accidentally deleted Amazon EBS snapshots and EBS-backed AMIs. When using Recycle Bin, if your resources are deleted, they are retained in the Recycle Bin for a time period that you specify before being permanently deleted. You can restore a resource from the Recycle Bin at any time before its retention period expires. This solution has the least operational overhead, as you do not need to create, copy, or upload any additional resources. You can also manage tags and permissions for AMIs in the Recycle Bin. AMIs in the Recycle Bin do not incur any additional charges. Reference:

upvoted 8 times

asdfcdsxdcf **Most Recent** 1 year, 9 months ago

Selected Answer: C

C is correct

upvoted 2 times

Naveena_Devanga 1 year, 9 months ago

Option C

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/recycle-bin-working-with-rules.html>

upvoted 2 times

Freddie26 1 year, 9 months ago

Option C is correct. Recycling bin is a new feature to protect snaps and AMIs from accidental or malicious deleting. Inside the recycling bin, set a retention policy, and then your images or snapshots are protected.

upvoted 4 times

mestule 1 year, 10 months ago

Selected Answer: C

<https://aws.amazon.com/about-aws/whats-new/2022/02/amazon-ec2-recycle-bin-machine-images/>

upvoted 4 times

Andy_09 1 year, 10 months ago

Option C

upvoted 3 times

A company has 150 TB of archived image data stored on-premises that needs to be moved to the AWS Cloud within the next month. The company's current network connection allows up to 100 Mbps uploads for this purpose during the night only.

What is the MOST cost-effective mechanism to move this data and meet the migration deadline?

- A. Use AWS Snowmobile to ship the data to AWS.
- B. Order multiple AWS Snowball devices to ship the data to AWS. **Most Voted**
- C. Enable Amazon S3 Transfer Acceleration and securely upload the data.
- D. Create an Amazon S3 VPC endpoint and establish a VPN to upload the data.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Andy_09 **Highly Voted** 2 years, 4 months ago

Option B

upvoted 13 times

Scheldon **Most Recent** 1 year, 11 months ago

Selected Answer: B

Answer B

SnowBall Storage will give us 80TB. To transfer data we will need 2 devices. Taking into consideration that 1 device = ~300\$ we will spend 600\$. Option B is most Cost effective and will allow us to end operation in less than month.

upvoted 3 times

sandordini 2 years, 1 month ago

Selected Answer: B

Snowball base free from 200USD, Snowmobile base fee from 4100USD (According to AWS)

upvoted 4 times

sandordini 2 years, 1 month ago

Snowmobile advised above 10 Petabytes

Snowball(s) below 10 PB

upvoted 3 times

TruthWS 2 years, 2 months ago

Option B - Snowmobile have higher cost

upvoted 2 times

Mikado211 2 years, 2 months ago

Selected Answer: B

Amazon S3 Transfer Acceleration must be very expensive

Correct in such case : B Snowball

upvoted 3 times

asdfcdsxdcf 2 years, 3 months ago

Selected Answer: B

B is correct

upvoted 2 times

iczcezar 2 years, 3 months ago

Option B

upvoted 2 times

Naveena_Devanga 2 years, 3 months ago

Option B:

1 Snow Ball Max Allowed capacity is 80 TB. Hence, you need to order multiple snowballs to achieve the requirement.

upvoted 3 times

stephensimudemey 2 years, 3 months ago

Selected Answer: B

B. Its only 150TB

upvoted 2 times

A company wants to migrate its three-tier application from on premises to AWS. The web tier and the application tier are running on third-party virtual machines (VMs). The database tier is running on MySQL.

The company needs to migrate the application by making the fewest possible changes to the architecture. The company also needs a database solution that can restore data to a specific point in time.

Which solution will meet these requirements with the LEAST operational overhead?

A. Migrate the web tier and the application tier to Amazon EC2 instances in private subnets. Migrate the database tier to Amazon RDS for MySQL in private subnets.

B. Migrate the web tier to Amazon EC2 instances in public subnets. Migrate the application tier to EC2 instances in private subnets. Migrate the database tier to Amazon Aurora MySQL in private subnets. **Most Voted**

C. Migrate the web tier to Amazon EC2 instances in public subnets. Migrate the application tier to EC2 instances in private subnets. Migrate the database tier to Amazon RDS for MySQL in private subnets.

D. Migrate the web tier and the application tier to Amazon EC2 instances in public subnets. Migrate the database tier to Amazon Aurora MySQL in public subnets.

Correct Answer: B

Community vote distribution

B (82%)

C (18%)

Comments

haci **Highly Voted** 2 years, 3 months ago

Selected Answer: B

I'm between B and C. Since RDS requires an additional configuration for PTR, it adds an operational overhead. So I will go with B.

Aurora provides automated backup and point-in-time recovery, simplifying backup management and data protection. Continuous incremental backups are taken automatically and stored in Amazon S3, and data retention periods can be specified to meet compliance requirements.

RDS provides the same but first, the users should set a retention period for these backups, allowing historical data recovery in case of accidental data loss or corruption, and point-in-time recovery (PITR) allows users to restore the database to any specific moment within the set retention period.

upvoted 13 times

shahrehi **Highly Voted** 2 years, 3 months ago

B: Amazon Aurora is a fully managed relational database engine that's compatible with both MySQL and PostgreSQL

upvoted 5 times

In2718 **Most Recent** 9 months, 3 weeks ago

Selected Answer: C

Chat AI -

Options A/D place tiers in the wrong subnet types;

B adds an engine change to Aurora without need.

upvoted 1 times

In2718 10 months, 2 weeks ago

Selected Answer: C

Also, with RDS you can restore for a specific point in time as well. I would spend the time to migrate all of this away from the current migration, but for now I would just keep apples to apples and stick with answer C

upvoted 1 times

In2718 10 months, 2 weeks ago

Selected Answer: C

I'll go with answer C because moving from MySQL server on prem to a MySQL RDS will involve less schema work and changes.
upvoted 1 times

Bwhizzy 1 year, 8 months ago

Selected Answer: B

Answer is B. please see below article
<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/aurora-pitr.html>
upvoted 2 times

kbgsqgs 1 year, 8 months ago

Selected Answer: C

Auroralo maigeuleisyeonhalyeomyeon MySQLboda akitegcheo byeongyeong-i deo pil-yohal su isseubnida. ganeunghamyeon choesohan-ui byeongyeong-eul yochyeong
67 / 5,000
Migrating to Aurora may require more architectural changes than MySQL. Request minimal changes if possible.
upvoted 1 times

Scheldon 1 year, 11 months ago

Selected Answer: B

AnswerB

The problem with this question is that we do not have enough information. We can execute task with both AURORA or RDS DB. I will go with AURORA as it is Amazon Proprietary and is developed by AWS teams, hence we do not need to think about updates etc. as it is done by AWS teams.
upvoted 2 times

MattBJ 2 years, 2 months ago

Selected Answer: B

B is the correct option.
upvoted 2 times

DEN_ZZ 2 years, 3 months ago

Selected Answer: B

PTR, it's Aurora
upvoted 4 times

stephensimudemy 2 years, 3 months ago

Selected Answer: C

It's C. Strictly speaking, there is no AWS DB call Amazon Aurora "MySQL"
upvoted 1 times

ogerber 2 years, 3 months ago

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/Aurora.AuroraMySQL.html>
upvoted 3 times

hajra313 2 years, 4 months ago

C. This option aligns with the requirements by keeping the web tier in public subnets, migrating the application tier to EC2 instances in private subnets to enhance security, and using Amazon RDS for MySQL in private subnets to meet the database requirements with minimal operational overhead. option A: While migrating the web tier and application tier to EC2 instances in private subnets minimizes exposure to the internet. option B: Migrating the database tier to Amazon Aurora MySQL introduces changes to the database engine, which might require additional testing and adjustments to the application. Additionally, Aurora MySQL does not directly support point-in-time recovery; instead, it uses continuous backups and snapshots for data recovery.
upvoted 3 times

Andy_09 2 years, 4 months ago

Option A works better
upvoted 2 times

Andy_09 2 years, 4 months ago

Option B
upvoted 3 times

A development team is collaborating with another company to create an integrated product. The other company needs to access an Amazon Simple Queue Service (Amazon SQS) queue that is contained in the development team's account. The other company wants to poll the queue without giving up its own account permissions to do so.

How should a solutions architect provide access to the SQS queue?

- A. Create an instance profile that provides the other company access to the SQS queue.
- B. Create an IAM policy that provides the other company access to the SQS queue.
- C. Create an SQS access policy that provides the other company access to the SQS queue. **Most Voted**
- D. Create an Amazon Simple Notification Service (Amazon SNS) access policy that provides the other company access to the SQS queue.

Correct Answer: C

Community vote distribution

C (100%)

Comments

iczcizar **Highly Voted** 2 years, 3 months ago

The correct option to provide access to the SQS queue without giving up the other company's account permissions is:

C. Create an SQS access policy that provides the other company access to the SQS queue.

By creating an SQS access policy, you can define specific permissions for the other company to access the SQS queue without requiring them to modify their own account permissions. This allows for fine-grained control over access to the queue while maintaining security and isolation between accounts. Options A, B, and D are not appropriate for granting access to the SQS queue in this scenario.

upvoted 8 times

7fb06b3 **Highly Voted** 2 years ago

Selected Answer: C

Amazon SQS policy system lets you grant permission to other Amazon Accounts.

https://docs.amazonaws.cn/en_us/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-using-identity-based-policies.html

upvoted 5 times

FlyingHawk **Most Recent** 1 year, 5 months ago

Selected Answer: C

<https://repost.aws/knowledge-center/sqs-queue-access-permissions>

Here is the example:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "AWS": "arn:aws:iam::OtherCompanyAccountID:root"
      },
      "Action": [
        "sqs:ReceiveMessage",
        "sqs:DeleteMessage"
      ],
      "Resource": "arn:aws:sqs:region:YourAccountID:YourQueueName"
    }
  ]
}
```

upvoted 1 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: C

A - Instance profiles are attached to Amazon EC2 instances. This is irrelevant.
B - While you could create an IAM role in your account for the other company to assume, it is explicitly stated that the other company does not want to give up its own permissions, meaning cross-account access must be configured, but IAM policies alone do not provide cross-account access.
C - SQS supports resource-based policies, which allow you to grant access to specific AWS principals (users, roles, or accounts) across accounts.
D - No idea what it would achieve. SNS policies don't even apply to SQS queues.

upvoted 2 times

Scheldon 1 year, 11 months ago

Selected Answer: C

AnswerC

Creating Access policy in SQS which will allow other company to access SQS queue seems to be the only solution which is RIGHT here

upvoted 2 times

sandordini 2 years, 1 month ago

Selected Answer: C

SQS Access Policy for secure, fine-grained Cross-account access

upvoted 3 times

NayeraB 2 years, 3 months ago

Selected Answer: C

<https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-overview-of-managing-access.html>

upvoted 2 times

hajra313 2 years, 4 months ago

option A: Instance profiles are used to grant permissions to EC2 instances, not for granting access to other AWS services like SQS queues. Option B: IAM policies are applied to IAM users, groups, or roles within the same AWS account. They are not directly applicable to granting access to resources in other AWS accounts. option C: SQS access policies allow you to grant cross-account access to SQS resources. You can specify the necessary permissions in the policy and attach it directly to the SQS queue. This way, you can give the other company's AWS account the necessary permissions to poll the queue without compromising their account permissions. option D. Amazon SNS access policies are used to manage access to SNS topics, not SQS queues

upvoted 4 times

kempes 2 years, 4 months ago

Selected Answer: C

Option C

upvoted 4 times

Andy_09 2 years, 4 months ago

Option B

upvoted 1 times

A company's developers want a secure way to gain SSH access on the company's Amazon EC2 instances that run the latest version of Amazon Linux. The developers work remotely and in the corporate office.

The company wants to use AWS services as a part of the solution. The EC2 instances are hosted in a VPC private subnet and access the internet through a NAT gateway that is deployed in a public subnet.

What should a solutions architect do to meet these requirements MOST cost-effectively?

- A. Create a bastion host in the same subnet as the EC2 instances. Grant the `ec2:CreateVpnConnection` IAM permission to the developers. Install EC2 Instance Connect so that the developers can connect to the EC2 instances.
- B. Create an AWS Site-to-Site VPN connection between the corporate network and the VPC. Instruct the developers to use the Site-to-Site VPN connection to access the EC2 instances when the developers are on the corporate network. Instruct the developers to set up another VPN connection for access when they work remotely.
- C. Create a bastion host in the public subnet of the VPC. Configure the security groups and SSH keys of the bastion host to only allow connections and SSH authentication from the developers' corporate and remote networks. Instruct the developers to connect through the bastion host by using SSH to reach the EC2 instances.
- D. Attach the `AmazonSSMMangedInstanceCore` IAM policy to an IAM role that is associated with the EC2 instances. Instruct the developers to use AWS Systems Manager Session Manager to access the EC2 instances. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

alawada **Highly Voted** 1 year, 8 months ago

Selected Answer: D

AWS Systems Manager Session Manager is a service that enables you to securely connect to your EC2 instances without using SSH keys or bastion hosts. You can use Session Manager to access your instances through the AWS Management Console, the AWS CLI, or the AWS SDKs. Session Manager uses IAM policies and roles to control who can access which instances. By attaching the `AmazonSSMMangedInstanceCore` IAM policy to an IAM role that is associated with the EC2 instances, you grant the Session Manager service the necessary permissions to perform actions on your instances. You also need to attach another IAM policy to the developers' IAM users or roles that allows them to start sessions to the instances.

upvoted 7 times

Andy_09 **Highly Voted** 1 year, 10 months ago

Option D

upvoted 5 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: D

A - Bastion hosts are typically placed in public subnets (not private subnets). Also, see option C below.

B - Would work, but this is too complex and too expensive.

C - An additional EC2 instance (bastion host) = instance running costs + endless maintenance.

D - God yes. No need to manage SSH keys, and access is encrypted by default.

upvoted 1 times

TruthWS 1 year, 8 months ago

Option D

upvoted 3 times

Mikado211 1 year, 8 months ago

Selected Answer: D

SSM is always the recommended way of connection for EC2 "using ssh".
It's the most cost effective and the most secure way of doing the job.

upvoted 2 times

iczcezar 1 year, 9 months ago

Why not C?

upvoted 2 times

pila21 1 year, 8 months ago

it doesn't meet requirements MOST cost-effectively

upvoted 3 times

kempes 1 year, 10 months ago

Selected Answer: D

Option D

upvoted 3 times


A pharmaceutical company is developing a new drug. The volume of data that the company generates has grown exponentially over the past few months. The company's researchers regularly require a subset of the entire dataset to be immediately available with minimal lag. However, the entire dataset does not need to be accessed on a daily basis. All the data currently resides in on-premises storage arrays, and the company wants to reduce ongoing capital expenses.

Which storage solution should a solutions architect recommend to meet these requirements?

- A. Run AWS DataSync as a scheduled cron job to migrate the data to an Amazon S3 bucket on an ongoing basis.
- B. Deploy an AWS Storage Gateway file gateway with an Amazon S3 bucket as the target storage. Migrate the data to the Storage Gateway appliance.
- C. Deploy an AWS Storage Gateway volume gateway with cached volumes with an Amazon S3 bucket as the target storage. Migrate the data to the Storage Gateway appliance. **Most Voted**
- D. Configure an AWS Site-to-Site VPN connection from the on-premises environment to AWS. Migrate data to an Amazon Elastic File System (Amazon EFS) file system.

Correct Answer: C

Community vote distribution



Comments

Andy_09 **Highly Voted** 2 years, 4 months ago

Option C

upvoted 11 times

hajra313 **Highly Voted** 2 years, 4 months ago

B. Deploying an AWS Storage Gateway file gateway with an Amazon S3 bucket as the target storage would require the entire dataset to be stored in Amazon S3, which might not be cost-effective considering that only a subset of the data needs to be accessed regularly. Additionally, accessing data directly from S3 might introduce latency. so correct option is C bcz AWS Storage Gateway volume gateway with cached volumes allows the company to keep frequently accessed data locally on-premises while storing the entire dataset in Amazon S3. This solution provides immediate access to the subset of data with minimal lag, as frequently accessed data is cached locally. It also reduces ongoing capital expenses as it leverages Amazon S3 storage, which is cost-effective.

upvoted 8 times

capepenguin 1 year, 11 months ago

Both B and C store the entire data set in Amazon S3 and the AWS Storage Gateway file gateway also supports local caching. I do not know if it is B or C for me.

upvoted 2 times

e10c971 **Most Recent** 6 months ago

Selected Answer: B

Why B fits best from ai

Immediate, low-latency access to a subset: File Gateway provides a local cache for recently accessed files, so researchers can work with hot data with minimal lag via standard NFS/SMB file protocols.

Cost optimization: The entire dataset is stored durably and cost-effectively in Amazon S3, reducing on-premises storage CapEx. Only frequently accessed data stays cached on-prem.

Operational simplicity: It integrates cleanly with existing file-based workflows and supports object-optimized storage in S3 with lifecycle policies (e.g. transitioning cold data to S3 Glacier).

upvoted 1 times

nadeerm 1 year, 3 months ago

Selected Answer: B

Volume Gateway is designed for block storage (EBS volumes), not file storage. It is not the best fit for this use case, which involves file-based data access.

upvoted 3 times

Scheldon 1 year, 11 months ago

Selected Answer: C

AnswerC

Deploying an AWS Storage Gateway volume gateway with cached volumes will allow to store all data in AWS but the most frequently accessed data will be stored/cached locally (on-premises) = low latency for most used data while all data will be stored in the cloud.

<https://docs.aws.amazon.com/storagegateway/latest/vgw/StorageGatewayConcepts.html#storage-gateway-cached-concepts>

upvoted 3 times

BatVanyo 2 years, 1 month ago

A storage guy here.. the question is not clear enough to give a definitive answer between B and C, as both can do the job.

An "on-prem storage array" can be any of the three:

- File array (serving any file protocol, e.g. NFS/SMB) -> requiring a file gateway (supports caching of the most recently used data)
- Block array (iSCSI/Fibre Channel) -> requiring a volume gateway (supports cached volumes most recently used data)
- Combo (providing both File and Block protocols)

Something is clearly missing in the question in order to give a definitive answer between B and C.

upvoted 5 times

mohammadthainat 2 years, 2 months ago

Selected Answer: C

storage arrays = Volume Gateway

upvoted 3 times

lenotc 2 years, 2 months ago

Selected Answer: C

storage array, also known as a disk array so AWS Storage Gateway volume.

its a trap

upvoted 5 times

MattBJ 2 years, 2 months ago

Selected Answer: C

C is correct. Using AWS Storage Gateway volume gateway with cached volumes provide local access to the file.

upvoted 3 times

ninasgx 2 years, 3 months ago

Selected Answer: C

require a subset of the entire dataset => cached volumes

upvoted 4 times

osmk 2 years, 3 months ago

Selected Answer: C

The company's researchers regularly require a subset of the entire dataset to be immediately available with minimal lag <https://docs.aws.amazon.com/storagegateway/latest/vgw/WhatIsStorageGateway.html>

upvoted 2 times

A company has a business-critical application that runs on Amazon EC2 instances. The application stores data in an Amazon DynamoDB table. The company must be able to revert the table to any point within the last 24 hours.

Which solution meets these requirements with the LEAST operational overhead?

- A. Configure point-in-time recovery for the table. **Most Voted**
- B. Use AWS Backup for the table.
- C. Use an AWS Lambda function to make an on-demand backup of the table every hour.
- D. Turn on streams on the table to capture a log of all changes to the table in the last 24 hours. Store a copy of the stream in an Amazon S3 bucket.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Andy_09 **Highly Voted** 2 years, 4 months ago

Option A

upvoted 13 times

MattBJ **Highly Voted** 2 years, 2 months ago

Selected Answer: A

A is correct. One of the highlight features of DynamoDB.

upvoted 5 times

Scheldon **Most Recent** 1 year, 11 months ago

Selected Answer: A

AnswerA

Point-in-time recovery helps protect your DynamoDB tables from accidental write or delete operations. With point-in-time recovery, you don't have to worry about creating, maintaining, or scheduling on-demand backups. For example, suppose that a test script writes accidentally to a production DynamoDB table. With point-in-time recovery, you can restore that table to any point in time during the last 35 days. After you enable point-in-time recovery, you can restore to any point in time from five minutes before the current time until 35 days ago. DynamoDB maintains incremental backup of your table.

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/PointInTimeRecovery.html>

upvoted 3 times

1dd 2 years, 3 months ago

Selected Answer: A

option A

upvoted 3 times

asdfcdsxdcf 2 years, 3 months ago

Selected Answer: A

A looks correct

upvoted 3 times

mavik 2 years, 3 months ago

Selected Answer: A

Option A

upvoted 3 times

A company hosts an application used to upload files to an Amazon S3 bucket. Once uploaded, the files are processed to extract metadata, which takes less than 5 seconds. The volume and frequency of the uploads varies from a few files each hour to hundreds of concurrent uploads. The company has asked a solutions architect to design a cost-effective architecture that will meet these requirements.

What should the solutions architect recommend?

- A. Configure AWS CloudTrail trails to log S3 API calls. Use AWS AppSync to process the files.
- B. Configure an object-created event notification within the S3 bucket to invoke an AWS Lambda function to process the files.
Most Voted
- C. Configure Amazon Kinesis Data Streams to process and send data to Amazon S3. Invoke an AWS Lambda function to process the files.
- D. Configure an Amazon Simple Notification Service (Amazon SNS) topic to process the files uploaded to Amazon S3. Invoke an AWS Lambda function to process the files.

Correct Answer: B

Community vote distribution

B (100%)

Comments

mestule **Highly Voted** 2 years, 4 months ago

Selected Answer: B

B seems to be make most sense to me.

upvoted 5 times

hajra313 **Highly Voted** 2 years, 4 months ago

option b bcz option c is WS AppSync is not the most appropriate solution for file processing.

option d While Amazon Simple Notification Service (SNS) can be used to trigger actions based on S3 events, it's not directly involved in processing files .option c :Kinesis is typically used for real-time data streaming and analytics, which may not be needed for simple file processing tasks such as extracting metadata.

upvoted 5 times

Lin878 **Most Recent** 1 year, 11 months ago

Selected Answer: B

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/EventNotifications.html>

upvoted 3 times

Scheldon 1 year, 11 months ago

Selected Answer: B

AnswerB

upvoted 2 times

alawada 2 years, 2 months ago

Selected Answer: B

The problem with C is how it sends the data to S3, if it was Firehose it would make sense. I waka for B.

upvoted 3 times

MattBJ 2 years, 2 months ago

Selected Answer: B

B is correct. The most cost effective option.

upvoted 3 times

jaswantn 2 years, 3 months ago

option B

upvoted 2 times

kempes 2 years, 4 months ago

Option D

upvoted 2 times

Andy_09 2 years, 4 months ago

Option D

upvoted 1 times

A company's application is deployed on Amazon EC2 instances and uses AWS Lambda functions for an event-driven architecture. The company uses nonproduction development environments in a different AWS account to test new features before the company deploys the features to production.

The production instances show constant usage because of customers in different time zones. The company uses nonproduction instances only during business hours on weekdays. The company does not use the nonproduction instances on the weekends. The company wants to optimize the costs to run its application on AWS.

Which solution will meet these requirements MOST cost-effectively?

- A. Use On-Demand Instances for the production instances. Use Dedicated Hosts for the nonproduction instances on weekends only.
- B. Use Reserved Instances for the production instances and the nonproduction instances. Shut down the nonproduction instances when not in use.
- C. Use Compute Savings Plans for the production instances. Use On-Demand Instances for the nonproduction instances. Shut down the nonproduction instances when not in use. **Most Voted**
- D. Use Dedicated Hosts for the production instances. Use EC2 Instance Savings Plans for the nonproduction instances.

Correct Answer: C

Community vote distribution

C (87%)

13%

Comments

Andy_09 **Highly Voted** 2 years, 4 months ago
Option C
upvoted 13 times

MattBJ **Highly Voted** 2 years, 2 months ago
Selected Answer: C
Definitely C.
upvoted 5 times

Penjerla **Most Recent** 1 year, 5 months ago
Selected Answer: B
A Reserved Instance is preferred if you use the reserve capacity for the full duration of the one- or three-year term. You may also find that a Reserved Instance offers cost savings over On-Demand Instances, even if you are only using the instance at half capacity. Reserved Instances also offer guaranteed availability for stop-and-start scenarios. In addition, by utilizing and continuously monitoring Convertible Reserved Instances, you can save more money.
upvoted 2 times

MatAlves 1 year, 8 months ago

Selected Answer: C

"An Amazon EC2 Dedicated Host is a physical server fully dedicated for your use, so you can help address corporate compliance requirements."

<https://aws.amazon.com/ec2/dedicated-hosts/>

That already eliminates "A" and "D". No need to pay extra since there is no requirement for eligible software licences.

B is INCORRECT = no need to use reserved instances for non-prod, since they will only be active during business hours during weekdays.
upvoted 2 times

Scheldon 1 year, 11 months ago

Selected Answer: C

AnswerC

We need to use somekind of savings plans for production and disable Development servers when we are not using them. Only Option C has both.
upvoted 2 times

Naveena_Devanga 2 years, 3 months ago

Option C

upvoted 2 times

stephensimudemy 2 years, 3 months ago

Selected Answer: C

It's C

upvoted 4 times

A company stores data in an on-premises Oracle relational database. The company needs to make the data available in Amazon Aurora PostgreSQL for analysis. The company uses an AWS Site-to-Site VPN connection to connect its on-premises network to AWS.

The company must capture the changes that occur to the source database during the migration to Aurora PostgreSQL.

Which solution will meet these requirements?

- A. Use the AWS Schema Conversion Tool (AWS SCT) to convert the Oracle schema to Aurora PostgreSQL schema. Use the AWS Database Migration Service (AWS DMS) full-load migration task to migrate the data.
- B. Use AWS DataSync to migrate the data to an Amazon S3 bucket. Import the S3 data to Aurora PostgreSQL by using the Aurora PostgreSQL aws_s3 extension.
- C. Use the AWS Schema Conversion Tool (AWS SCT) to convert the Oracle schema to Aurora PostgreSQL schema. Use AWS Database Migration Service (AWS DMS) to migrate the existing data and replicate the ongoing changes. **Most Voted**
- D. Use an AWS Snowball device to migrate the data to an Amazon S3 bucket. Import the S3 data to Aurora PostgreSQL by using the Aurora PostgreSQL aws_s3 extension.

Correct Answer: C

Community vote distribution

C (100%)

Comments

kempes **Highly Voted** 2 years, 4 months ago

Selected Answer: C

Option C

upvoted 10 times

Andy_09 **Highly Voted** 2 years, 4 months ago

Option C

upvoted 8 times

striker89 **Most Recent** 1 year, 7 months ago

Selected Answer: C

migrate the existing data and replicate the ongoing changes.

upvoted 2 times

Scheldon 1 year, 11 months ago

Selected Answer: C

Answer C

upvoted 1 times

MattBJ 2 years, 2 months ago

Selected Answer: C

C is correct. As we need to capture the change during the migration.

upvoted 2 times

A company built an application with Docker containers and needs to run the application in the AWS Cloud. The company wants to use a managed service to host the application.

The solution must scale in and out appropriately according to demand on the individual container services. The solution also must not result in additional operational overhead or infrastructure to manage.

Which solutions will meet these requirements? (Choose two.)

- A. Use Amazon Elastic Container Service (Amazon ECS) with AWS Fargate. **Most Voted**
- B. Use Amazon Elastic Kubernetes Service (Amazon EKS) with AWS Fargate. **Most Voted**
- C. Provision an Amazon API Gateway API. Connect the API to AWS Lambda to run the containers.
- D. Use Amazon Elastic Container Service (Amazon ECS) with Amazon EC2 worker nodes.
- E. Use Amazon Elastic Kubernetes Service (Amazon EKS) with Amazon EC2 worker nodes.

Correct Answer: AB

Community vote distribution

AB (64%) AC (36%)

Comments

xBUGx **Highly Voted** 2 years, 2 months ago

Selected Answer: AC

I don't want confuse other...

upvoted 10 times

JohnYu 1 year, 7 months ago

AWS Lambda can run Docker containers, but it is more suited for short-duration, event-driven functions rather than long-running services.

While you can use Lambda to run container images, it is typically used for microservices or functions rather than managing scalable containerized applications in the same way that ECS or EKS does.

upvoted 4 times

FlyingHawk **Most Recent** 1 year, 5 months ago

Selected Answer: AB

A or B is okay, why not C?

While AWS Lambda now supports running Linux-based container images, it is primarily designed for microservices or event-driven functions rather than managing scalable containerized applications like Amazon ECS or Amazon EKS

upvoted 2 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: AB

A or B is OK.

About C: Amazon API Gateway and AWS Lambda are not typically used for running Docker containers. Lambda is designed for serverless functions, and while you can use Lambda to invoke containerized tasks, it does not run containers directly.

upvoted 2 times

JA2018 1 year, 6 months ago

Selected Answer: AB

AWS Fargate:

Both ECS and EKS with Fargate provide a fully managed container orchestration service where you don't need to manage underlying infrastructure like EC2 instances, making it ideal for scaling and reducing operational overhead.

Scaling capabilities:

Both options allow for automatic scaling based on demand, meaning the container instances will automatically scale up or down based on application requirements.

upvoted 1 times

striker89 1 year, 7 months ago

Selected Answer: AB

API Gateway + Lambda are not designed for running Docker containers directly

upvoted 3 times

MatAlves 1 year, 8 months ago

Selected Answer: AC

- Container = ECS or EKS
- "managed service to host the application" = Fargate
- "not result in additional operational overhead or infrastructure to manage" = ECS for the win.

The main difference between ECS and EKS = simplicity vs flexibility.

<https://aws.amazon.com/blogs/containers/amazon-ecs-vs-amazon-eks-making-sense-of-aws-container-services/>

upvoted 2 times

MatAlves 1 year, 8 months ago

now, if instead of asking the "best" (combination) solution you want TWO SOLUTIONS, then I see value in A-B.

upvoted 2 times

Salilgen 1 year, 5 months ago

Where do you read "best" solution?

upvoted 1 times

Abdullah2004 1 year, 9 months ago

Selected Answer: AB

For sure

upvoted 3 times

MatAlves 1 year, 8 months ago

Why so many trolls recently...

upvoted 1 times

Abdullah2004 1 year, 7 months ago

Dear MatAlves before you describe my answer as trolls try to study and focus

I will explain to you that C is not correct

Lambda is for running code in response to events and does not natively support running Docker containers in the traditional sense of a containerized application. Lambda can use container image as deployment packages but this setup isn't ideal for managing and calling complex containerized applications

So it's A + B

upvoted 4 times

1e22522 1 year, 10 months ago

Selected Answer: AB

glowies out reeeeeeeeeeeeeeeee

upvoted 2 times

Rhydian25 1 year, 11 months ago

Selected Answer: AB

The question is asking for two alternatives to run Docker containers in a serverless service with minimal effort.

Option C will require a lot of effort to configure the Lambda and the API Gateway to run the Container correctly.

Instead, just use EKS or ECS with Fargate to execute the container image

upvoted 4 times

victor78 1 year, 12 months ago

It should be AB

upvoted 2 times

sheilawu 1 year, 12 months ago

Selected Answer: AB

It should be AB, Container :ECS , EKS

upvoted 2 times

sheilawu 1 year, 12 months ago

No, I decided to change my option since this question is asking "Docker"&"AWS"
it is not asking local Kubernetes, so AC should be the right answer.

upvoted 1 times

jckk202020 2 years, 1 month ago

AB are using AWS Fargate which IS considered a managed service, option C does not run containers, , DE you have to manage your own EC2 instances thus not consider managed

upvoted 2 times

asdfcdsxdcf 2 years, 3 months ago

Selected Answer: AB

Everyone is picking AB too..

upvoted 2 times

NayeraB 2 years, 3 months ago

Are people picking A&B as alternate solutions? Is the question asking for alternates?? Am I missing something? Somebody explain please I'm super confused.

upvoted 2 times

Cali182 2 years, 3 months ago

The question states itself. Which Solutions....?

upvoted 3 times

Drew3000 2 years, 2 months ago

I believe so. Based on other questions, they would have asked "which combination"

upvoted 2 times

kempes 2 years, 4 months ago

Option AB

upvoted 2 times

Andy_09 2 years, 4 months ago

Option AB

upvoted 3 times

An ecommerce company is running a seasonal online sale. The company hosts its website on Amazon EC2 instances spanning multiple Availability Zones. The company wants its website to manage sudden traffic increases during the sale.

Which solution will meet these requirements MOST cost-effectively?

- A. Create an Auto Scaling group that is large enough to handle peak traffic load. Stop half of the Amazon EC2 instances. Configure the Auto Scaling group to use the stopped instances to scale out when traffic increases.
- B. Create an Auto Scaling group for the website. Set the minimum size of the Auto Scaling group so that it can handle high traffic volumes without the need to scale out.
- C. Use Amazon CloudFront and Amazon ElastiCache to cache dynamic content with an Auto Scaling group set as the origin. Configure the Auto Scaling group with the instances necessary to populate CloudFront and ElastiCache. Scale in after the cache is fully populated.
- D. Configure an Auto Scaling group to scale out as traffic increases. Create a launch template to start new instances from a preconfigured Amazon Machine Image (AMI). **Most Voted**

Correct Answer: D

Community vote distribution

D (77%)

C (23%)

Comments

Andy_09 **Highly Voted** 2 years, 4 months ago

Option D

upvoted 8 times

In2718 **Most Recent** 10 months, 2 weeks ago

Selected Answer: C

The key to the question is this - meet these requirements MOST cost-effectively. Using Cloudfront and Elastic Cache would minimize the cost of scaling out EC2 servers. For this I would choose C over D.

upvoted 1 times

JA2018 1 year, 6 months ago

Selected Answer: D

Reasons?

Cost-effective scaling: Allows the company to only pay for the EC2 instances needed during peak traffic periods, automatically scaling up when necessary and scaling down when traffic subsides, which is the most cost-efficient approach.

Launch template and AMI: Using a launch template with a preconfigured AMI simplifies the scaling process by ensuring new instances are quickly provisioned with the required configurations.

upvoted 2 times

viejito 2 years ago

En la respuesta C : Al usar servicios como Amazon CloudFront y Amazon ElastiCache para almacenar en caché el contenido dinámico, reduciendo la carga en las instancias de Amazon EC2 y mejorando la velocidad de entrega del contenido a los usuarios finales. Esto resulta en una solución más rentable y eficiente en comparación con la respuesta D : simplemente escalar instancias de EC2 sin considerar medidas adicionales para optimizar el rendimiento y reducir los costos . Por lo que la opción correcta es la C .

upvoted 1 times

sandordini 2 years, 1 month ago

Selected Answer: D

Cloudfront could be a good idea, but it seems to be a simple scaling scenario.. IMO: D

upvoted 3 times

aditianand 2 years ago

Don't we need a launch template? How can it just launch from AMI image

upvoted 2 times

buzzinmumbai 2 years, 1 month ago

Answer is D .C is not cost effective to use elasticache .Not sure if you can have ASG as the origin.

upvoted 2 times

geraltRebo 2 years, 1 month ago

Selected Answer: D

Sorry D

upvoted 2 times

geraltRebo 2 years, 1 month ago

Selected Answer: C

Option C

upvoted 1 times

TruthWS 2 years, 2 months ago

Option D bring a most cost effective

upvoted 2 times

JCAWS 2 years, 2 months ago

Selected Answer: C

C more suitable

upvoted 1 times

stephensimudemy 2 years, 3 months ago

Selected Answer: D

It's D

upvoted 3 times

A solutions architect must provide an automated solution for a company's compliance policy that states security groups cannot include a rule that allows SSH from 0.0.0.0/0. The company needs to be notified if there is any breach in the policy. A solution is needed as soon as possible.

What should the solutions architect do to meet these requirements with the LEAST operational overhead?

- A. Write an AWS Lambda script that monitors security groups for SSH being open to 0.0.0.0/0 addresses and creates a notification every time it finds one.
- B. Enable the restricted-ssh AWS Config managed rule and generate an Amazon Simple Notification Service (Amazon SNS) notification when a noncompliant rule is created. **Most Voted**
- C. Create an IAM role with permissions to globally open security groups and network ACLs. Create an Amazon Simple Notification Service (Amazon SNS) topic to generate a notification every time the role is assumed by a user.
- D. Configure a service control policy (SCP) that prevents non-administrative users from creating or editing security groups. Create a notification in the ticketing system when a user requests a rule that needs administrator permissions.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Andy_09 **Highly Voted** 1 year, 10 months ago

Option B

upvoted 8 times

kempes **Highly Voted** 1 year, 10 months ago

Selected Answer: B

Option B

upvoted 5 times

sandordini **Most Recent** 1 year, 7 months ago

Selected Answer: B

The others sound 'silly'... to say the least

upvoted 2 times

asdfcdsxdcf 1 year, 9 months ago

Selected Answer: B

B looks correct

upvoted 2 times

Naveena_Devanga 1 year, 9 months ago

Option B

<https://docs.aws.amazon.com/config/latest/developerguide/restricted-ssh.html>

upvoted 4 times

hajra313 1 year, 10 months ago

option b

upvoted 3 times

Use Amazon Elastic Kubernetes Service (Amazon EKS) with Amazon EC2 worker nodes.

A company has deployed an application in an AWS account. The application consists of microservices that run on AWS Lambda and Amazon Elastic Kubernetes Service (Amazon EKS). A separate team supports each microservice. The company has multiple AWS accounts and wants to give each team its own account for its microservices.

A solutions architect needs to design a solution that will provide service-to-service communication over HTTPS (port 443). The solution also must provide a service registry for service discovery.

Which solution will meet these requirements with the LEAST administrative overhead?

A. Create an inspection VPC. Deploy an AWS Network Firewall firewall to the inspection VPC. Attach the inspection VPC to a new transit gateway. Route VPC-to-VPC traffic to the inspection VPC. Apply firewall rules to allow only HTTPS communication.

B. Create a VPC Lattice service network. Associate the microservices with the service network. Define HTTPS listeners for each service. Register microservice compute resources as targets. Identify VPCs that need to communicate with the services. Associate those VPCs with the service network. **Most Voted**

C. Create a Network Load Balancer (NLB) with an HTTPS listener and target groups for each microservice. Create an AWS PrivateLink endpoint service for each microservice. Create an interface VPC endpoint in each VPC that needs to consume that microservice.

D. Create peering connections between VPCs that contain microservices. Create a prefix list for each service that requires a connection to a client. Create route tables to route traffic to the appropriate VPC. Create security groups to allow only HTTPS communication.

Correct Answer: B

Community vote distribution

B (100%)

Comments

1dd **Highly Voted** 1 year, 9 months ago

Selected Answer: B

VPC Lattice is a completely new way to simplify API communication between services or microservices in one or more AWS accounts.
upvoted 7 times

SnaxAttax **Most Recent** 6 months, 2 weeks ago

Selected Answer: B

VPC Lattice is a FMS, is good for service-to-service use cases, and supports HTTPS.
upvoted 1 times

zinabu 1 year, 7 months ago

Selected Answer: B

Amazon VPC Lattice is a new capability of Amazon Virtual Private Cloud (Amazon VPC) designed to simplify networking for service-to-service communication.

link: https://www.bing.com/search?q=what+VPC+Lattice+service+used+for+microservices&cvid=d706d95737274f388660cbda9b7b2c4e&gs_lcrp=EgZjaHJvbWUyBggAEEUYOTIICAQ6QcY_FXSAQkyMTY1N2owajSoAgCwAgE&FORM=ANAB01&PC=U531

upvoted 4 times

aditinand 1 year, 6 months ago

Did you complete the exam recently? Was examtopics useful?

upvoted 2 times

phoenix2023 1 year, 6 months ago

Please keep in mind this is for helpful answers to THIS specific question. Please don't abuse it as a random info pitch for yourself. This is distracting and wasting others' time. Please respect other people's time.

upvoted 7 times

stephensimudemey 1 year, 9 months ago

Selected Answer: B

IT's B. Google VPC Lattice service network

upvoted 3 times

Andy_09 1 year, 10 months ago

Option B

upvoted 4 times

A company has a mobile game that reads most of its metadata from an Amazon RDS DB instance. As the game increased in popularity, developers noticed slowdowns related to the game's metadata load times. Performance metrics indicate that simply scaling the database will not help. A solutions architect must explore all options that include capabilities for snapshots, replication, and sub-millisecond response times.

What should the solutions architect recommend to solve these issues?

- A. Migrate the database to Amazon Aurora with Aurora Replicas.
- B. Migrate the database to Amazon DynamoDB with global tables.
- C. Add an Amazon ElastiCache for Redis layer in front of the database. **Most Voted**
- D. Add an Amazon ElastiCache for Memcached layer in front of the database.

Correct Answer: C

Community vote distribution

C (80%)

B (20%)

Comments

Andy_09 **Highly Voted** 2 years, 4 months ago

Option C is better as we need replication and snapshots

upvoted 20 times

JA2018 1 year, 6 months ago

Not sure if this is applicable to the stem. While Redis can be used as a cache, placing it in front of a relational database like RDS might not be enough to address the performance issues if the primary bottleneck is database access time.

upvoted 1 times

arunkpskpm 2 years, 3 months ago

C is correct as only Redis support snapshot feature :<https://aws.amazon.com/elasticache/redis-vs-memcached/>

upvoted 6 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: C

A - If the root cause is database-related (for example, read-heavy traffic), then OK. However, "Performance metrics indicate that simply scaling the database will not help." means that maybe caching is a more effective option.

B - Migrating from a relational database to DynamoDB requires so much effort. Such a headache.

C - ElastiCache for Redis supports replication and snapshots, and of course, sub-millisecond response times.

D - Memcached lacks built-in persistence, which is why game applications rarely use it.

upvoted 2 times

tch 1 year, 5 months ago

Selected Answer: C

capabilities for snapshots, replication

upvoted 1 times

JA2018 1 year, 6 months ago

Selected Answer: B

Why I choose B?

1 High-performance, low latency: DynamoDB is a NoSQL database designed for fast, sub-millisecond response times, which is crucial for handling the high volume of metadata requests from a popular mobile game.

2 Scalability: DynamoDB scales automatically to handle fluctuating traffic without the need to manually manage database instances, making it ideal for a rapidly growing game.

3 Global tables: This feature allows for geographically distributed data access, ensuring quick response times for players across different regions.

upvoted 2 times

JunsK1e 1 year, 11 months ago

Selected Answer: C

C is correct

upvoted 1 times

f07ed8f 2 years ago

Thinking the capabilities with "snapshots, replication, and sub-millisecond response times" is for the Database or selected solution(ElastiCache).

upvoted 2 times

asdfcdsxdcf 2 years, 3 months ago

Selected Answer: C

C is correct

upvoted 2 times

nbellaiche 2 years, 3 months ago

Selected Answer: C

Réponse C

upvoted 1 times

osmk 2 years, 3 months ago

Selected Answer: C

:<https://aws.amazon.com/elasticache/redis-vs-memcached/>

upvoted 1 times

Andy_09 2 years, 4 months ago

Option D

upvoted 1 times

A company uses AWS Organizations for its multi-account AWS setup. The security organizational unit (OU) of the company needs to share approved Amazon Machine Images (AMIs) with the development OU. The AMIs are created by using AWS Key Management Service (AWS KMS) encrypted snapshots.

Which solution will meet these requirements? (Choose two.)

A. Add the development team's OU Amazon Resource Name (ARN) to the launch permission list for the AMIs. **Most Voted**

B. Add the Organizations root Amazon Resource Name (ARN) to the launch permission list for the AMIs.

C. Update the key policy to allow the development team's OU to use the AWS KMS keys that are used to decrypt the snapshots. **Most Voted**

D. Add the development team's account Amazon Resource Name (ARN) to the launch permission list for the AMIs.

E. Recreate the AWS KMS key. Add a key policy to allow the Organizations root Amazon Resource Name (ARN) to use the AWS KMS key.

Correct Answer: AC

Community vote distribution

AC (100%)

Comments

Andy_09 **Highly Voted** 2 years, 4 months ago
Changing to options AC
upvoted 13 times

osmk **Highly Voted** 2 years, 3 months ago

Selected Answer: AC

c=><https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/share-amis-with-organizations-and-OUs.html#allow-org-ou-to-use-key>

A--><https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/share-amis-with-organizations-and-OUs.html#share-amis-org-ou>

upvoted 5 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: AC

A - The security OU is explicitly allowing the development OU to use the AMI.

B - OK but too broad, and not secure.

C - Now they can use the encrypted snapshots.

D - Would be too much trouble if there are many accounts. A is enough.

E - Recreate the AWS KMS key? Uhm, why?

upvoted 2 times

Scheldon 1 year, 11 months ago

Selected Answer: AC

AnswerAC

Option A will allow to run/lunch AMIs

Option C will allow to decrypt AMIs which is necessary to run AMI.

upvoted 2 times

cjace 1 year, 12 months ago

CD - Solution C: Update the Key Policy

Why: The AMIs are created using KMS-encrypted snapshots, so the KMS keys must allow the development team's accounts to use these keys for decrypting the snapshots.

How: Update the key policy of the KMS key to include permissions for the development OU or specific accounts within that OU. This will enable those accounts to use the KMS key for decrypting the snapshots associated with the AMIs.

Solution D: Add the Development Team's Account ARN to the Launch Permission List

Why: To share the AMIs with the development accounts, you need to grant launch permissions to those accounts. This allows the specified accounts to use the shared AMIs to launch instances.

How: Add the ARNs of the development team's accounts to the launch permission list of the AMIs. This can be done using the modify-image-attribute command in the AWS CLI, specifying the account IDs that should have launch permissions.

upvoted 1 times

Mikado211 2 years, 2 months ago

Selected Answer: AC

A : give users the right to launch

C : give users the right to decrypt

upvoted 4 times

Andy_09 2 years, 4 months ago

Option CD

upvoted 1 times

A data analytics company has 80 offices that are distributed globally. Each office hosts 1 PB of data and has between 1 and 2 Gbps of internet bandwidth.


The company needs to perform a one-time migration of a large amount of data from its offices to Amazon S3. The company must complete the migration within 4 weeks.

Which solution will meet these requirements MOST cost-effectively?

- A. Establish a new 10 Gbps AWS Direct Connect connection to each office. Transfer the data to Amazon S3.
- B. Use multiple AWS Snowball Edge storage-optimized devices to store and transfer the data to Amazon S3. **Most Voted**
- C. Use an AWS Snowmobile to store and transfer the data to Amazon S3.
- D. Set up an AWS Storage Gateway Volume Gateway to transfer the data to Amazon S3.

Correct Answer: B

Community vote distribution



Comments

mestule **Highly Voted** 2 years, 4 months ago

Selected Answer: B

B because too many offices that are geographically separated.

"data analytics company has 80 offices that are distributed globally."

upvoted 16 times

Andy_09 2 years, 4 months ago

Nice spot...completely missed that part !!

upvoted 4 times

MatAlves **Most Recent** 1 year, 8 months ago

Selected Answer: B

AWS Snowmobile: ideal for migrating big datasets containing 10PB or more and stored in one location. Snowmobile can help you migrate all of these large datasets at once, but the process requires a high-speed backbone with hundreds of Gb/s of spare throughput.

AWS Snowball: ideal for datasets storing less than 10PB or datasets distributed across multiple locations. You can use Snowball to migrate data incrementally—this is a good alternative if you do not have enough bandwidth on the network backbone.

<https://bluexp.netapp.com/blog/aws-cvo-blg-aws-snowball-vs-snowmobile-data-migration-options-comparedwork> backbone.

upvoted 3 times

ccceb01 1 year, 9 months ago

Selected Answer: B

10 Snowball price (100TB x 10) still cheaper then 1 Snowmobile price (\$4,100 plus additional cost for data transfer)

upvoted 2 times

aditinand 2 years ago

Why not D? Why not AWS storage gateway? With 1Gbps, they can transfer 1.25 GBPS which translates to 2.8 PB in 4 weeks. They need just 1 PB to be transferred

upvoted 2 times

buzzinmumbai 2 years, 1 month ago

As of March 2024 AWS has stopped offering snowmobile as a service .So B is the right answer.Hopefully they don't ask this question :)

upvoted 3 times

Tanidanindo 2 years, 2 months ago

Selected Answer: C

Too large for snowball devices.

upvoted 2 times

Naveena_Devanga 2 years, 3 months ago

Option C,

An AWS Snowmobile has a maximum storage capacity of 100 petabytes (PB). This is equivalent to the capacity of 1,250 Snowball Edge devices

upvoted 1 times

HarryLopez 2 years, 3 months ago

but there are many offices geographically distributed, so snowmobile for each one of them adds up to a lot of cost as compared to option B).

upvoted 3 times

sandordini 2 years, 1 month ago

Snowmobile advised over 10PB! Definitely snowball

upvoted 1 times

chefKC 2 years, 4 months ago

option B

upvoted 2 times

Andy_09 2 years, 4 months ago

Option C looks good, as option B would lead to usage of too many snowball devices.

upvoted 2 times


A company has an Amazon Elastic File System (Amazon EFS) file system that contains a reference dataset. The company has applications on Amazon EC2 instances that need to read the dataset. However, the applications must not be able to change the dataset. The company wants to use IAM access control to prevent the applications from being able to modify or delete the dataset.

Which solution will meet these requirements?

- A. Mount the EFS file system in read-only mode from within the EC2 instances.
- B. Create a resource policy for the EFS file system that denies the `elasticfilesystem:ClientWrite` action to the IAM roles that are attached to the EC2 instances. **Most Voted**
- C. Create an identity policy for the EFS file system that denies the `elasticfilesystem:ClientWrite` action on the EFS file system.
- D. Create an EFS access point for each application. Use Portable Operating System Interface (POSIX) file permissions to allow read-only access to files in the root directory.

Correct Answer: B

Community vote distribution



Comments

hajra313 **Highly Voted** 2 years, 4 months ago

Create an EFS access point for each application. Use Portable Operating System Interface (POSIX) file permissions to allow read-only access to files in the root directory.

Explanation:

By creating an EFS access point for each application and configuring POSIX file permissions to allow read-only access, you can enforce the desired access control. This approach restricts write and delete actions on the dataset while allowing read access, aligning with the company's requirements.
upvoted 6 times

MatAlves 1 year, 8 months ago

Resource policies are included in "IAM to control":

"Using IAM to control file system data access

NFS clients can identify themselves using an IAM role when connecting to an EFS file system. When a client connects to a file system, Amazon EFS evaluates the file system's IAM resource policy, which is called a file system policy, along with any identity-based IAM policies to determine the appropriate file system access permissions to grant."

<https://docs.aws.amazon.com/efs/latest/ug/iam-access-control-nfs-efs.html>

upvoted 2 times

f07ed8f 2 years ago

Please note that the question is asking "The company wants to use IAM access control to prevent the applications from being able to modify or delete the dataset."

upvoted 5 times

lenotc **Highly Voted** 2 years, 3 months ago

Selected Answer: B

B correct best solution best well architected

C wrong because identity policies are typically associated with users or roles, not directly with the EFS file system

D wrong because POSIX file permissions at the root directory level may not be sufficient to prevent modifications to other directories or files

A is so far away

upvoted 6 times

MatAlves **Most Recent** 1 year, 8 months ago

Selected Answer: B

- Identity-based policies are attached to an IAM user, group, or role.
- Resource-based policies are attached to a resource.
- elasticfilesystem:ClientWrite: Provides write permissions on a file system.

EFS is a RESOURCE, so that excludes "C" (we need a resource policy).

<https://docs.aws.amazon.com/efs/latest/ug/iam-access-control-nfs-efs.html>

https://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies_identity-vs-resource.html

upvoted 2 times

elmyth 1 year, 9 months ago

Selected Answer: B

There is no such thing as an "identity policy" for EFS.

upvoted 2 times

sandordini 2 years, 1 month ago

Selected Answer: B

2 ways to prevent writing to the file system:

1. The mount option in the /etc/fstab file is set to read-only access. > A
2. IAM policy indicates read-only access, or root access disabled. > B

The question clearly states they are looking to use IAM access control

upvoted 3 times

Ansuman_lucky 2 years, 2 months ago

prevent the applications from being able to modify or delete the dataset.-- This means a role would be used. So answer is B

upvoted 4 times

xBUGx 2 years, 2 months ago

IAM policies are used to control access to AWS resources, including Amazon EFS. By default, IAM policies control access to the EFS API actions, such as elasticfilesystem:ClientWrite, which allows clients to write to the file system. However, POSIX file permissions control access to files within the file system itself, which is independent of IAM policies.

While using POSIX file permissions can restrict access to the files within the file system, it doesn't prevent a user or application with the appropriate IAM permissions from modifying or deleting those files directly through the EFS API.

upvoted 4 times

HarryLopez 2 years, 3 months ago

Selected Answer: B

B)

IAM needs to be used, so A) & D) are out.

So b/w B) and C), Resource policies are meant for specific aws service or resource while Identity policies are attached to an identity (user, group or role). C) attached identity policy to EFS, dont know how and why. Hence, B).

upvoted 3 times

osmk 2 years, 3 months ago

Selected Answer: C

company wants to use IAM access control to prevent <https://docs.aws.amazon.com/efs/latest/ug/iam-access-control-nfs-efs.html>

upvoted 3 times

jaswantn 2 years, 3 months ago

Selected Answer: D

option D

upvoted 1 times

Salilgen 1 year, 5 months ago

You use IAM policy to control permissions with access points

<https://docs.aws.amazon.com/efs/latest/ug/access-points-iam-policy.html>

upvoted 1 times

Oo_Cc 2 years, 4 months ago

Selected Answer: C

"The company wasn't to use IAM access control". Yes, it would deny writing action to everything .. but it's still the only one that uses IAM.

upvoted 2 times

MatAlves 1 year, 8 months ago

"Using IAM to control file system data access

NFS clients can identify themselves using an IAM role when connecting to an EFS file system. When a client connects to a file system, Amazon EFS evaluates the file system's IAM resource policy, which is called a file system policy, along with any identity-based IAM policies to determine the appropriate file system access permissions to grant."

<https://docs.aws.amazon.com/efs/latest/ug/iam-access-control-nfs-efs.html>
upvoted 2 times

MatAlves 1 year, 8 months ago

What we need to change is the " IAM resource policy".
upvoted 1 times

Andy_09 2 years, 4 months ago

Option B

upvoted 5 times

A company has hired an external vendor to perform work in the company's AWS account. The vendor uses an automated tool that is hosted in an AWS account that the vendor owns. The vendor does not have IAM access to the company's AWS account. The company needs to grant the vendor access to the company's AWS account.

Which solution will meet these requirements MOST securely?

A. Create an IAM role in the company's account to delegate access to the vendor's IAM role. Attach the appropriate IAM policies to the role for the permissions that the vendor requires. **Most Voted**

B. Create an IAM user in the company's account with a password that meets the password complexity requirements. Attach the appropriate IAM policies to the user for the permissions that the vendor requires.

C. Create an IAM group in the company's account. Add the automated tool's IAM user from the vendor account to the group. Attach the appropriate IAM policies to the group for the permissions that the vendor requires.

D. Create an IAM user in the company's account that has a permission boundary that allows the vendor's account. Attach the appropriate IAM policies to the user for the permissions that the vendor requires.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Mikado211 **Highly Voted** 2 years, 1 month ago

Selected Answer: A

When you have somebody from another account who needs a resource in your account

- create a role to access to this account
- allow the remote account to assume the role.

upvoted 8 times

Andy_09 **Highly Voted** 2 years, 4 months ago

Option A looks ok

upvoted 7 times

Salilgen **Most Recent** 1 year, 5 months ago

Selected Answer: A

https://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_common-scenarios_third-party.html

upvoted 1 times

JA2018 1 year, 6 months ago

Selected Answer: A

Least privilege principle: Using an IAM role instead of a user adheres to the least privilege principle as it only grants the vendor access to specific actions needed for their task, not full account access like a user would have.

No direct access to the company account: By delegating access through a role, the vendor's IAM user in their own account does not need to have direct login credentials to the company's AWS account, enhancing security.

Granular control: IAM policies attached to the role can be carefully crafted to provide only the necessary permissions for the vendor's automated tool, limiting potential damage in case of unauthorized access

upvoted 1 times

Scheldon 1 year, 11 months ago

Selected Answer: A

AnswerA

I would go with option A

upvoted 2 times

osmk 2 years, 3 months ago

Selected Answer: A

Question #222

upvoted 7 times

A company wants to run its experimental workloads in the AWS Cloud. The company has a budget for cloud spending. The company's CFO is concerned about cloud spending accountability for each department. The CFO wants to receive notification when the spending threshold reaches 60% of the budget.

Which solution will meet these requirements?

- A. Use cost allocation tags on AWS resources to label owners. Create usage budgets in AWS Budgets. Add an alert threshold to receive notification when spending exceeds 60% of the budget. **Most Voted**
- B. Use AWS Cost Explorer forecasts to determine resource owners. Use AWS Cost Anomaly Detection to create alert threshold notifications when spending exceeds 60% of the budget.
- C. Use cost allocation tags on AWS resources to label owners. Use AWS Support API on AWS Trusted Advisor to create alert threshold notifications when spending exceeds 60% of the budget.
- D. Use AWS Cost Explorer forecasts to determine resource owners. Create usage budgets in AWS Budgets. Add an alert threshold to receive notification when spending exceeds 60% of the budget.

Correct Answer: A

Community vote distribution

A (100%)

Comments

NayeraB **Highly Voted** 2 years, 3 months ago

Selected Answer: A

Nothing with cost explorer in it, and I don't want to be Captain Obvious but we need to set the budget alerts through AWS Budgets, so A upvoted 9 times

Scheldon **Most Recent** 1 year, 12 months ago

Selected Answer: A

AnswerA

<https://docs.aws.amazon.com/cost-management/latest/userguide/budgets-controls.html>

upvoted 3 times

Andy_09 2 years, 4 months ago

Option A

upvoted 4 times

A company wants to deploy an internal web application on AWS. The web application must be accessible only from the company's office. The company needs to download security patches for the web application from the internet.

The company has created a VPC and has configured an AWS Site-to-Site VPN connection to the company's office. A solutions architect must design a secure architecture for the web application.

Which solution will meet these requirements?

- A. Deploy the web application on Amazon EC2 instances in public subnets behind a public Application Load Balancer (ALB). Attach an internet gateway to the VPC. Set the inbound source of the ALB's security group to 0.0.0.0/0.
- B. Deploy the web application on Amazon EC2 instances in private subnets behind an internal Application Load Balancer (ALB). Deploy NAT gateways in public subnets. Attach an internet gateway to the VPC. Set the inbound source of the ALB's security group to the company's office network CIDR block. **Most Voted**
- C. Deploy the web application on Amazon EC2 instances in public subnets behind an internal Application Load Balancer (ALB). Deploy NAT gateways in private subnets. Attach an internet gateway to the VPC. Set the outbound destination of the ALB's security group to the company's office network CIDR block.
- D. Deploy the web application on Amazon EC2 instances in private subnets behind a public Application Load Balancer (ALB). Attach an internet gateway to the VPC. Set the outbound destination of the ALB's security group to 0.0.0.0/0.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Andy_09 **Highly Voted** 2 years, 4 months ago

Option B

upvoted 7 times

osmk **Highly Voted** 2 years, 3 months ago

Selected Answer: B

none sense why IGW on top of NATGW.

upvoted 6 times

MatAlves 1 year, 8 months ago

<https://docs.aws.amazon.com/network-firewall/latest/developerguide/arch-igw-ngw.html>

Confusing, I agree. But it seems to be recommended in some cases.

upvoted 1 times

striker89 **Most Recent** 1 year, 7 months ago

Selected Answer: B

I Would go for B even if NAT GW allow outbound traffic ONLY. Still wondering how the Company newtwork will access Private Subnet in the VPC.

upvoted 1 times

Scheldon 1 year, 11 months ago

AnswerB

Server and LB in Private will hide WEB application from the word. NAT will allow for server's access to the internet in case of need

upvoted 1 times

NayeraB 2 years, 3 months ago

Selected Answer: B

B is well structured
upvoted 3 times

ogerber 2 years, 3 months ago

To my opinion, with only having inbound of the companys CIDR block, it will not include access for the patches available online.
i would go for D
upvoted 4 times

sandordini 2 years, 1 month ago

Incorrect: B says inbound, D says outbound. Outbound for ALB are the EC2 Instances.
upvoted 1 times

kempes 2 years, 4 months ago

Selected Answer: B

Option B
upvoted 4 times

A company maintains its accounting records in a custom application that runs on Amazon EC2 instances. The company needs to migrate the data to an AWS managed service for development and maintenance of the application data. The solution must require minimal operational support and provide immutable, cryptographically verifiable logs of data changes.

Which solution will meet these requirements MOST cost-effectively?

- A. Copy the records from the application into an Amazon Redshift cluster.
- B. Copy the records from the application into an Amazon Neptune cluster.
- C. Copy the records from the application into an Amazon Timestream database.
- D. Copy the records from the application into an Amazon Quantum Ledger Database (Amazon QLDB) ledger. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Andy_09 **Highly Voted** 2 years, 4 months ago

Option D

upvoted 7 times

asdfcdsxdcf **Highly Voted** 2 years, 3 months ago

Selected Answer: D

Amazon QLDB

- QLDB stands for "Quantum Ledger Database"
- A ledger is a book recording financial transactions
- Fully Managed, Serverless, High available, Replication across 3 AZ
- Used to review history of all the changes made to your application data over time
- Immutable system: no entry can be removed or modified, cryptographically verifiable

upvoted 5 times

Scheldon **Most Recent** 1 year, 11 months ago

Selected Answer: D

AnswerD

Amazon Quantum Ledger Database (Amazon QLDB) is a fully managed ledger database that provides a transparent, immutable, and cryptographically verifiable transaction log owned by a central trusted authority.

<https://docs.aws.amazon.com/qldb/latest/developerguide/what-is.html>

upvoted 4 times

f07ed8f 2 years ago

Selected Answer: D

Amazon Quantum Ledger Database (Amazon QLDB) is a fully managed ledger database that provides a transparent, immutable, and cryptographically verifiable transaction log.

upvoted 3 times

Mikado211 2 years, 2 months ago

Selected Answer: D

immutable, cryptographically verifiable ==> Amazon QLDB

upvoted 3 times

NayeraB 2 years, 3 months ago

Selected Answer: D

D is correct
upvoted 2 times

A company's marketing data is uploaded from multiple sources to an Amazon S3 bucket. A series of data preparation jobs aggregate the data for reporting. The data preparation jobs need to run at regular intervals in parallel. A few jobs need to run in a specific order later.

The company wants to remove the operational overhead of job error handling, retry logic, and state management.

Which solution will meet these requirements?

- A. Use an AWS Lambda function to process the data as soon as the data is uploaded to the S3 bucket. Invoke other Lambda functions at regularly scheduled intervals.
- B. Use Amazon Athena to process the data. Use Amazon EventBridge Scheduler to invoke Athena on a regular interval.
- C. Use AWS Glue DataBrew to process the data. Use an AWS Step Functions state machine to run the DataBrew data preparation jobs. **Most Voted**
- D. Use AWS Data Pipeline to process the data. Schedule Data Pipeline to process the data once at midnight.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Andy_09 **Highly Voted** 2 years, 4 months ago

Option C

upvoted 9 times

FlyingHawk **Most Recent** 1 year, 4 months ago

Selected Answer: C

Both c and d are workable, but AWS Data Pipeline is considered outdated because AWS has introduced newer, more efficient, and serverless alternatives for data orchestration and ETL (Extract, Transform, Load) tasks. While AWS Data Pipeline is still available, it is not actively promoted or updated by AWS, and many companies have shifted to AWS Glue, Step Functions, and EventBridge for data workflow automation.

<https://aws.amazon.com/blogs/big-data/build-event-driven-data-quality-pipelines-with-aws-glue-databrew/>

upvoted 3 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: C

A - Lambda function alone is incompetent for retries or running in a specific order.

B - Athena cannot automatically handle errors or running in a specific order.

C - YES and fully managed. See below.

D - Data Pipeline with retries and error handling is A LOT OF WORK.

Takeaway:

1. Glue DataBrew simplifies the data preparation process with a visual interface, and it integrates with S3 for input/output.

2. Step Functions efficiently manages parallel and sequential workflows while minimizing operational overhead for state management, error handling and retries.

upvoted 4 times

Scheldon 1 year, 11 months ago

Selected Answer: C

AnswerC

With Step Functions' built-in controls, you can examine the state of each step in your workflow to make sure that your application runs in order and as expected.

<https://docs.aws.amazon.com/step-functions/latest/dg/welcome.html>

AWS Glue is a serverless data integration service that makes it easy for analytics users to discover, prepare, move, and integrate data from multiple sources.

<https://docs.aws.amazon.com/glue/latest/dg/what-is-glue.html>

upvoted 3 times

asdfcdsxdcf 2 years, 3 months ago

Selected Answer: C

c looks correct

upvoted 3 times

A solutions architect is designing a payment processing application that runs on AWS Lambda in private subnets across multiple Availability Zones. The application uses multiple Lambda functions and processes millions of transactions each day.

The architecture must ensure that the application does not process duplicate payments.

Which solution will meet these requirements?

- A. Use Lambda to retrieve all due payments. Publish the due payments to an Amazon S3 bucket. Configure the S3 bucket with an event notification to invoke another Lambda function to process the due payments.
- B. Use Lambda to retrieve all due payments. Publish the due payments to an Amazon Simple Queue Service (Amazon SQS) queue. Configure another Lambda function to poll the SQS queue and to process the due payments.
- C. Use Lambda to retrieve all due payments. Publish the due payments to an Amazon Simple Queue Service (Amazon SQS) FIFO queue. Configure another Lambda function to poll the FIFO queue and to process the due payments. **Most Voted**
- D. Use Lambda to retrieve all due payments. Store the due payments in an Amazon DynamoDB table. Configure streams on the DynamoDB table to invoke another Lambda function to process the due payments.

Correct Answer: C

Community vote distribution

C (82%)

D (18%)

Comments

hajra313 **Highly Voted** 2 years, 4 months ago

Standard queues provide at-least-once delivery, which means that each message is delivered at least once.

FIFO queues provide exactly-once processing, which means that each message is delivered once and remains available until a consumer processes it and deletes it. Duplicates are not introduced into the queue. OPTION C

upvoted 18 times

Scheldon **Most Recent** 1 year, 11 months ago

Selected Answer: C

AnswerC

SQS FIFO was created for such tasks

Unlike standard queues, FIFO queues don't introduce duplicate messages. FIFO queues help you avoid sending duplicates to a queue. If you retry the SendMessage action within the 5-minute deduplication interval, Amazon SQS doesn't introduce any duplicates into the queue. <https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/FIFO-queues-exactly-once-processing.html>

upvoted 3 times

escalibran 2 years, 3 months ago

Selected Answer: C

C over D, because

<https://docs.aws.amazon.com/lambda/latest/dg/with-ddb.html> Processing dynamo streams with lambda can cause duplication.

SQS FIFO can be configured for High Throughput to exceed the 3000/s (batched) limit

<https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/high-throughput-fifo.html>

I previously worked with payments and would argue that either option doesn't fully solve duplications. Events might be sent multiple times from source, you definitely want to perform de-duplication and have some sort of idempotent processing for them, instead of just blindly processing each thing you're given.

upvoted 3 times

asdfcdsxdfc 2 years, 3 months ago

Selected Answer: C

c is correct

upvoted 2 times

shahreh1 2 years, 3 months ago

Option C:

FIFO queues

Exactly-Once Processing – A message is delivered once and remains available until a consumer processes and deletes it. Duplicates aren't introduced into the queue.

First-In-First-Out Delivery – The order in which messages are sent and received is strictly preserved.

upvoted 2 times

FZA24 2 years, 3 months ago

Selected Answer: C

Option C Fifo

upvoted 3 times

Mikado211 2 years, 3 months ago

SQS can have duplicate messages in case of problems with the timeout window.

upvoted 2 times

haci 2 years, 3 months ago

Selected Answer: C

"The application does not process duplicate payments" is the key point, which leads us directly to SQS FIFO

upvoted 3 times

Cali182 2 years, 4 months ago

Selected Answer: D

Option D

DynamoDB Streams helps ensure the following:

Each stream record appears exactly once in the stream.

For each item that is modified in a DynamoDB table, the stream records appear in the same sequence as the actual modifications to the item.

DynamoDB Streams writes stream records in near-real time so that you can build applications that consume these streams and take action based on the contents.

upvoted 3 times

jaswantn 2 years, 3 months ago

Option D...If you need to handle millions of transactions each day, you might need to consider other approach instead of SQS FIFO. And amongst the given options, we have DynmamoDB that maintains order in the streams.

upvoted 1 times

NayeraB 2 years, 3 months ago

I'm not sure if the answer is DynamoDB as well, but answering your question, SQS Fifo can handle 300 messages/second without batching, 3,000 messages/second with batching. Assuming we're using the 300/sec option, with 86,400 seconds in a day, that gives you 25,920,000 messages, so in short, yes SQS can handle millions of requests each day.

Not to mention DynamoDB doesn't provide the exactly-once processing the SQS offer and clearly requested in the question. That's just my train o thought, I'm happy to be corrected.

upvoted 3 times

jaswantn 2 years, 3 months ago

Dynamodb streams with partition key can be used to implement exactly once processing. There are many options with dynamodb to check for already processed item, and can be filtered out so that they are processed only once.

upvoted 1 times

jaswantn 2 years, 3 months ago

This calculation limits the number of transactions to 25 million a day. What if there are transactions exceeding this limit? As question say ... millions of transactions a day; that could be 70,80 or 90 millions also. In that case how SQS FIFO would perform?

Happy to be corrected with more convincing facts

upvoted 1 times

c3c7e62 1 year, 1 month ago

It seems that DynamoDB Streams deduplication is not guaranteed -

"A Lambda consumer for a DynamoDB stream doesn't guarantee exactly once delivery and may lead to occasional duplicates. Make sure your Lambda function code is idempotent to prevent unexpected issues from arising because of duplicate processing."

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Streams.Lambda.BestPracticesWithDynamoDB.html>

upvoted 1 times

kempes 2 years, 4 months ago

Option c

upvoted 2 times

Andy_09 2 years, 4 months ago

Option B

upvoted 1 times

A company runs multiple workloads in its on-premises data center. The company's data center cannot scale fast enough to meet the company's expanding business needs. The company wants to collect usage and configuration data about the on-premises servers and workloads to plan a migration to AWS.

Which solution will meet these requirements?

- A. Set the home AWS Region in AWS Migration Hub. Use AWS Systems Manager to collect data about the on-premises servers.
- B. Set the home AWS Region in AWS Migration Hub. Use AWS Application Discovery Service to collect data about the on-premises servers. **Most Voted**
- C. Use the AWS Schema Conversion Tool (AWS SCT) to create the relevant templates. Use AWS Trusted Advisor to collect data about the on-premises servers.
- D. Use the AWS Schema Conversion Tool (AWS SCT) to create the relevant templates. Use AWS Database Migration Service (AWS DMS) to collect data about the on-premises servers.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Kezuko **Highly Voted** 2 years, 2 months ago

Still the planning stage, C and D is out.

upvoted 5 times

Scheldon **Most Recent** 1 year, 11 months ago

Selected Answer: B

AnswerB

AWS Migration Hub delivers a guided end-to-end migration and modernization journey through discovery, assessment, planning, and execution.
<https://aws.amazon.com/migration-hub/>

AWS Application Discovery Service helps you plan your migration to the AWS cloud by collecting usage and configuration data about your on-premises servers and databases. Application Discovery Service is integrated with AWS Migration Hub and AWS Database Migration Service Fleet Advisor.

<https://docs.aws.amazon.com/application-discovery/latest/userguide/what-is-appdiscovery.html>

upvoted 4 times

Ipergorta 2 years, 2 months ago

Option D

upvoted 1 times

Ipergorta 2 years, 2 months ago

Sorry B

upvoted 4 times

asdfcdsxdfc 2 years, 3 months ago

Selected Answer: B

B is correct

upvoted 2 times

Andy_09 2 years, 4 months ago

Option B

upvoted 4 times

A company has an organization in AWS Organizations that has all features enabled. The company requires that all API calls and logins in any existing or new AWS account must be audited. The company needs a managed solution to prevent additional work and to minimize costs. The company also needs to know when any AWS account is not compliant with the AWS Foundational Security Best Practices (FSBP) standard.

Which solution will meet these requirements with the LEAST operational overhead?

A. Deploy an AWS Control Tower environment in the Organizations management account. Enable AWS Security Hub and AWS Control Tower Account Factory in the environment. **Most Voted**

B. Deploy an AWS Control Tower environment in a dedicated Organizations member account. Enable AWS Security Hub and AWS Control Tower Account Factory in the environment.

C. Use AWS Managed Services (AMS) Accelerate to build a multi-account landing zone (MALZ). Submit an RFC to self-service provision Amazon GuardDuty in the MALZ.

D. Use AWS Managed Services (AMS) Accelerate to build a multi-account landing zone (MALZ). Submit an RFC to self-service provision AWS Security Hub in the MALZ.

Correct Answer: A

Community vote distribution

A (100%)

Comments

FlyingHawk **Highly Voted** 1 year, 4 months ago

Selected Answer: A

This question was in my today's exam.

upvoted 7 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: A

A - Control Tower enforces centralized logging, monitoring, and account management. Security Hub monitors compliance with AWS Foundational Security Best Practices (FSBP). Account Factory simplifies the creation of new accounts with preconfigured settings.

B - Control Tower is intended to be deployed in the management account, not a member account, which deviates from AWS-recommended practices.

C - GuardDuty focuses on security threats rather than compliance monitoring for FSBP.

D - AMS Accelerate is more suitable for enterprises with highly complex requirements. Also, more operational overheads and cost more, but the question wants "to prevent additional work and to minimize costs".

upvoted 1 times

LuongTo 1 year, 7 months ago

Why not B?

upvoted 1 times

Kezuko 2 years, 2 months ago

Selected Answer: A

<https://docs.aws.amazon.com/controltower/latest/userguide/security-hub-controls.html>

upvoted 4 times

Ipergorta 2 years, 2 months ago

Option D

upvoted 1 times

asdfcdsxdhc 2 years, 3 months ago

Selected Answer: A

A is correct

upvoted 2 times

kempes 2 years, 4 months ago

Selected Answer: A

Option A

upvoted 4 times

Andy_09 2 years, 4 months ago

Option A

upvoted 2 times

A company has stored 10 TB of log files in Apache Parquet format in an Amazon S3 bucket. The company occasionally needs to use SQL to analyze the log files.

Which solution will meet these requirements MOST cost-effectively?

- A. Create an Amazon Aurora MySQL database. Migrate the data from the S3 bucket into Aurora by using AWS Database Migration Service (AWS DMS). Issue SQL statements to the Aurora database.
- B. Create an Amazon Redshift cluster. Use Redshift Spectrum to run SQL statements directly on the data in the S3 bucket.
- C. Create an AWS Glue crawler to store and retrieve table metadata from the S3 bucket. Use Amazon Athena to run SQL statements directly on the data in the S3 bucket. **Most Voted**
- D. Create an Amazon EMR cluster. Use Apache Spark SQL to run SQL statements directly on the data in the S3 bucket.

Correct Answer: C

Community vote distribution

C (100%)

Comments

LeonSauveterre 1 year, 5 months ago

Selected Answer: C

A - Aurora is cool but migrating 10 TB of data incurs significant costs and operational overhead.
B - Redshift Spectrum allows querying data directly in S3 without loading it into Redshift, but costs are really high especially for infrequent use.
C - Athena is serverless and charges only for the data scanned by queries. Glue Crawler automatically extracts metadata and schema information from the Parquet files. No need to migrate anything.
D - Just by the look of it I know I'll go bankrupt if I choose that.

upvoted 1 times

sandordini 1 year, 7 months ago

Selected Answer: C

S3 + SQL = Athena

upvoted 3 times

Kezuko 1 year, 8 months ago

Selected Answer: C

Apache Parquet => Glue Crawler

upvoted 4 times

asdfcdsxdcf 1 year, 9 months ago

Selected Answer: C

c is correct

upvoted 2 times

kempes 1 year, 10 months ago

Selected Answer: C

Option C

upvoted 4 times

Andy_09 1 year, 10 months ago

Option C

upvoted 4 times

A company needs a solution to prevent AWS CloudFormation stacks from deploying AWS Identity and Access Management (IAM) resources that include an inline policy or "*" in the statement. The solution must also prohibit deployment of Amazon EC2 instances with public IP addresses. The company has AWS Control Tower enabled in its organization in AWS Organizations.

Which solution will meet these requirements?

A. Use AWS Control Tower proactive controls to block deployment of EC2 instances with public IP addresses and inline policies with elevated access or "*". **Most Voted**

B. Use AWS Control Tower detective controls to block deployment of EC2 instances with public IP addresses and inline policies with elevated access or "*".

C. Use AWS Config to create rules for EC2 and IAM compliance. Configure the rules to run an AWS Systems Manager Session Manager automation to delete a resource when it is not compliant.

D. Use a service control policy (SCP) to block actions for the EC2 instances and IAM resources if the actions lead to noncompliance.

Correct Answer: A

Community vote distribution

A (61%)

D (39%)

Comments

jaswantn **Highly Voted** 2 years, 3 months ago

Selected Answer: D

Option D... This is preventive control of Control Tower where we use SCP to disallow actions that lead to policy violation.

upvoted 9 times

osmk **Highly Voted** 2 years, 3 months ago

Selected Answer: A

Proactive controls are implemented using AWS CloudFormation hooks within AWS Control Tower. They operate before resources are deployed to determine compliance with activated controls. SCPs are part of AWS Organizations and are used to manage permissions. vs Define specific purposes for implementing controls. <https://docs.aws.amazon.com/controltower/latest/userguide/proactive-controls.html>

upvoted 6 times

osmk 2 years, 3 months ago

SCPs focus on managing permissions at the OU level, while proactive controls in AWS Control Tower help prevent non-compliance during resource provisioning.

upvoted 3 times

FlyingHawk **Most Recent** 1 year, 4 months ago

Selected Answer: A

Proactive controls check whether resources are compliant with your company policies and objectives, before the resources are provisioned in your accounts. If the resources are out of compliance, they are not provisioned. Proactive controls monitor resources that would be deployed in your accounts by means of AWS CloudFormation templates.

For those who are familiar with AWS: In AWS Control Tower preventive controls are implemented with Service Control Policies (SCPs).

Detective controls are implemented with AWS Config rules.

Proactive controls are implemented with AWS CloudFormation hooks.

upvoted 3 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: A

A - Proactive controls are a feature of AWS Control Tower that prevent noncompliant resources from being deployed by validating configurations before deployment.

B - Detective controls can't block deployment.

C - If you must monitor to get it right, then something's already wrong before you notice that.

D - SCPs can indeed block specific API calls for creating IAM resources with "*" or EC2 instances with public IPs, but can't make the most of AWS CloudFormation stacks. Plus, SCPs apply at the account level and might inadvertently restrict legitimate use cases. Not ideal enough is what I'm saying.

upvoted 1 times

LeonSauveterre 1 year, 5 months ago

And btw, since "The company has AWS Control Tower enabled", option A induces much less overhead than option D.

upvoted 2 times

MatAlves 1 year, 8 months ago

Selected Answer: A

Prevent AWS CloudFormation from deploying IAM resources and EC2 instances based on specific use cases = Control Tower Proactive controls.

"Proactive controls are security controls that are designed to prevent the creation of noncompliant resources.

For example (...), through AWS CloudFormation, the proactive control can prevent the creation of update of any S3 bucket that has public access enabled."

<https://docs.aws.amazon.com/prescriptive-guidance/latest/aws-security-controls/proactive-controls.html>

upvoted 3 times

88f8032 2 years, 1 month ago

Selected Answer: A

this is A

upvoted 2 times

Sergiuss95 2 years, 1 month ago

Selected Answer: D

Is D, the best way to prevent this actions, is deploying SCPs

upvoted 1 times

BBR01 2 years, 1 month ago

Selected Answer: D

It is D. You want to prevent the events from happening.

Proactive controls check whether resources are compliant with your company policies and objectives, before the resources are provisioned in your accounts.

Detective controls detect specific events when they occur and log the action in CloudTrail.

Preventive controls prevent actions from occurring.

Preventive controls are implemented with SCPs. Detective controls are implemented with AWS Config rules. Proactive controls are implemented with AWS CloudFormation hooks.

<https://docs.aws.amazon.com/controltower/latest/userguide/how-control-tower-works.html#how-controls-work>

upvoted 1 times

TwinSpark 2 years ago

as stated by you A is correct, Proactive controls are implemented as cloudformation hooks and the resource will not be deployed if not compliants. It is exactly what is asked in question.

Using an SCP it is actually a valid solution, but it need to be associated to a specific resource that is not specified. If you associated to root account nobody can deploy a public ip ec2, not only cloudformation

upvoted 2 times

NayeraB 2 years, 3 months ago

Selected Answer: A

A would provide a proactive solution, also I'm not sure if SCP are made for granular details like creation of EC2 instances with public IP addresses or IAM resources with certain inline policies.

upvoted 2 times

Andy_09 2 years, 4 months ago

Option D

upvoted 2 times

A company's web application that is hosted in the AWS Cloud recently increased in popularity. The web application currently exists on a single Amazon EC2 instance in a single public subnet. The web application has not been able to meet the demand of the increased web traffic.

The company needs a solution that will provide high availability and scalability to meet the increased user demand without rewriting the web application.

Which combination of steps will meet these requirements? (Choose two.)

- A. Replace the EC2 instance with a larger compute optimized instance.
- B. Configure Amazon EC2 Auto Scaling with multiple Availability Zones in private subnets. **Most Voted**
- C. Configure a NAT gateway in a public subnet to handle web requests.
- D. Replace the EC2 instance with a larger memory optimized instance.
- E. Configure an Application Load Balancer in a public subnet to distribute web traffic. **Most Voted**

Correct Answer: BE

Community vote distribution

BE (100%)

Comments

kempes **Highly Voted** 2 years, 4 months ago

Selected Answer: BE

Option BE

upvoted 6 times

Scheldon **Most Recent** 1 year, 11 months ago

Selected Answer: BE

AnswerBE,

AutoScaling to increase amount of servers per need,
Load Balancer to balance traffic equally to all available servers

upvoted 2 times

sandordini 2 years, 1 month ago

Selected Answer: BE

Only BE makes sense, even though it might require modification of the application
upvoted 4 times

gsgdga 2 years, 2 months ago

Why isn't C the answer?

upvoted 1 times

802c4ff 2 years, 1 month ago

nat gateway is for accessing internet-facing from private subnets, not the other way around
upvoted 3 times

asdfcdsxdcf 2 years, 3 months ago

Selected Answer: BE

be are correct

upvoted 2 times

chefKC 2 years, 4 months ago

Option B & E

upvoted 3 times

Andy_09 2 years, 4 months ago

Option BE

upvoted 4 times

A company has AWS Lambda functions that use environment variables. The company does not want its developers to see environment variables in plaintext.

Which solution will meet these requirements?

- A. Deploy code to Amazon EC2 instances instead of using Lambda functions.
- B. Configure SSL encryption on the Lambda functions to use AWS CloudHSM to store and encrypt the environment variables.
- C. Create a certificate in AWS Certificate Manager (ACM). Configure the Lambda functions to use the certificate to encrypt the environment variables.
- D. Create an AWS Key Management Service (AWS KMS) key. Enable encryption helpers on the Lambda functions to use the KMS key to store and encrypt the environment variables. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Andy_09 **Highly Voted** 2 years, 4 months ago

Option D

upvoted 7 times

osmk **Highly Voted** 2 years, 3 months ago

Selected Answer: D

<https://docs.aws.amazon.com/lambda/latest/dg/configuration-envvars.html#configuration-envvars-encryption>

upvoted 7 times

JA2018 **Most Recent** 1 year, 6 months ago

Selected Answer: D

Reasons for choosing D:

AWS KMS is the primary service for managing encryption keys in AWS. This makes it the most appropriate choice for encrypting sensitive environment variables within Lambda functions.

Encryption helpers in Lambda allow you to easily integrate KMS key usage within your Lambda code, ensuring that environment variables are stored and accessed in an encrypted format.

upvoted 2 times

MatAlves 1 year, 8 months ago

"To configure encryption for your environment variables

Enable console encryption helpers to use client-side encryption to protect your data in transit.

Under Encryption in transit, choose Enable helpers for encryption in transit.

For each environment variable that you want to enable console encryption helpers for, choose Encrypt next to the environment variable.

Under AWS KMS key to encrypt in transit, choose a customer managed key that you created at the beginning of this procedure."

upvoted 3 times

Rhydian25 1 year, 11 months ago

Selected Answer: D

I don't understand why we should use a complex way of encrypting variables instead of using Parameter Store... but in this case the best option is D

upvoted 2 times


An analytics company uses Amazon VPC to run its multi-tier services. The company wants to use RESTful APIs to offer a web analytics service to millions of users. Users must be verified by using an authentication service to access the APIs.

Which solution will meet these requirements with the MOST operational efficiency?

- A. Configure an Amazon Cognito user pool for user authentication. Implement Amazon API Gateway REST APIs with a Cognito authorizer. **Most Voted**
- B. Configure an Amazon Cognito identity pool for user authentication. Implement Amazon API Gateway HTTP APIs with a Cognito authorizer.
- C. Configure an AWS Lambda function to handle user authentication. Implement Amazon API Gateway REST APIs with a Lambda authorizer.
- D. Configure an IAM user to handle user authentication. Implement Amazon API Gateway HTTP APIs with an IAM authorizer.

Correct Answer: A

Community vote distribution



Comments

stephensimudemy **Highly Voted** 1 year, 9 months ago

Selected Answer: A

User pools is for Authentication and user management
upvoted 7 times

sandordini **Highly Voted** 1 year, 7 months ago

Selected Answer: A

User pools are for authentication. Your app users can sign in through the user pool, Identity pools are for authorization, give them access to other AWS services.
upvoted 6 times

zdi561 **Most Recent** 1 year, 3 months ago

Selected Answer: A

When comparing operation efficiency, a Cognito authorizer in AWS API Gateway is generally considered more performant than a Lambda authorizer for basic authentication checks, as Cognito is a managed service optimized for user authentication, while a Lambda authorizer introduces an additional layer of processing with potential cold start overhead, making it better suited for complex, custom authorization logic. But Lambda is more flexible
upvoted 1 times

NayeraB 1 year, 9 months ago

Selected Answer: B

B offers more operational efficiency imo
upvoted 2 times

chefKC 1 year, 10 months ago

Answer is A
upvoted 2 times

Andy_09 1 year, 10 months ago

Option A
upvoted 5 times

A company has a mobile app for customers. The app's data is sensitive and must be encrypted at rest. The company uses AWS Key Management Service (AWS KMS).

The company needs a solution that prevents the accidental deletion of KMS keys. The solution must use Amazon Simple Notification Service (Amazon SNS) to send an email notification to administrators when a user attempts to delete a KMS key.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create an Amazon EventBridge rule that reacts when a user tries to delete a KMS key. Configure an AWS Config rule that cancels any deletion of a KMS key. Add the AWS Config rule as a target of the EventBridge rule. Create an SNS topic that notifies the administrators.
- B. Create an AWS Lambda function that has custom logic to prevent KMS key deletion. Create an Amazon CloudWatch alarm that is activated when a user tries to delete a KMS key. Create an Amazon EventBridge rule that invokes the Lambda function when the DeleteKey operation is performed. Create an SNS topic. Configure the EventBridge rule to publish an SNS message that notifies the administrators.
- C. Create an Amazon EventBridge rule that reacts when the KMS DeleteKey operation is performed. Configure the rule to initiate an AWS Systems Manager Automation runbook. Configure the runbook to cancel the deletion of the KMS key. Create an SNS topic. Configure the EventBridge rule to publish an SNS message that notifies the administrators. **Most Voted**
- D. Create an AWS CloudTrail trail. Configure the trail to deliver logs to a new Amazon CloudWatch log group. Create a CloudWatch alarm based on the metric filter for the CloudWatch log group. Configure the alarm to use Amazon SNS to notify the administrators when the KMS DeleteKey operation is performed.

Correct Answer: C

Community vote distribution

C (86%)

14%

Comments

Andy_09 **Highly Voted** 2 years, 4 months ago

Option C

<https://docs.aws.amazon.com/prescriptive-guidance/latest/patterns/monitor-and-remediate-scheduled-deletion-of-aws-kms-keys.html>

upvoted 15 times

hajra313 **Highly Voted** 2 years, 4 months ago

option c bcz Option C emerges as the clear winner due to its:

Direct event monitoring for the DeleteKey operation

Pre-built automation using Systems Manager Automation runbooks

Efficient notification via Amazon SNS

Minimal code development and operational overhead

Reduced risk of accidental deletion with faster response times

upvoted 10 times

FlyingHawk **Most Recent** 1 year, 4 months ago

Selected Answer: C

A- AWS config is for detection, not for prevention, it cannot cancel the deletion.

B is more effort to implement the custom logic to prevent KMS key deletion.

C is correct to use EventBridge rule to detect it then config the rule to use AWS System Manager Automation rubook to cancel the deletion and

publish the event to SNS.

D - does not have action for prevent the deletion of key.

upvoted 2 times

JA2018 1 year, 6 months ago

Selected Answer: C

Option C provides the most streamlined solution with minimal additional components needed (LEAST operational overhead). It directly leverages EventBridge to monitor for KMS key deletion attempts, triggers a Systems Manager Automation runbook to automatically prevent the deletion, and sends notifications through SNS without requiring additional Lambda functions or complex CloudWatch metric filtering.

upvoted 1 times

MatAlves 1 year, 8 months ago

Selected Answer: C

"Deletion of an AWS KMS key is scheduled.

The scheduled-deletion event is evaluated by an EventBridge rule.

The EventBridge rule engages the Amazon SNS topic.

The EventBridge rule initiates the Systems Manager automation and runbooks.

The runbooks cancel the deletion."

<https://docs.aws.amazon.com/prescriptive-guidance/latest/patterns/monitor-and-remediate-scheduled-deletion-of-aws-kms-keys.html>

upvoted 2 times

JunsK1e 1 year, 11 months ago

Selected Answer: C

I agree with andy_09

upvoted 2 times

Dammy031 1 year, 11 months ago

Selected Answer: D

Cloud trail helps to keep all invoked API calls in the AWS account which can trail back to the delete call made by a user

CloudWatch triggers an alarm when deletion is attempted.

SNS sends a notification to the administration about the attempt made.

All these met the requirement of the question.

upvoted 2 times

sandordini 2 years, 1 month ago

Selected Answer: C

My educated guess was C. Now, reading the comments, from Hajrá313 and knben I feel confident as well :)

upvoted 3 times

camp5 2 years, 2 months ago

It's D <https://docs.aws.amazon.com/kms/latest/developerguide/deleting-keys-creating-cloudwatch-alarm.html>

<https://docs.aws.amazon.com/kms/latest/developerguide/deleting-keys-creating-cloudwatch-alarm.html#cloudwatch-alarm-prerequisites>

upvoted 2 times

Salilgen 1 year, 5 months ago

This procedures create an alarm that notifies you when someone in your account tries to use a KMS key that is pending deletion. It doesn't notify you when a user attempts to delete a KMS key or even prevent deletion of the key

upvoted 1 times

1dd 2 years, 3 months ago

C as it "cancel the deletion of the KMS key"

upvoted 2 times

knben 2 years, 3 months ago

I would go with C

A -> Config is for compliance

B -> No lambda is required, too much complexity

C -> It achieves the goal, since KMS keys are not immediately deleted, which gives time to automation to cancel the action

D -> Cloudtrail is for auditing

upvoted 4 times

NayeraB 2 years, 3 months ago

Selected Answer: C

I agree with hajra313

upvoted 2 times

A company wants to analyze and generate reports to track the usage of its mobile app. The app is popular and has a global user base. The company uses a custom report building program to analyze application usage.

The program generates multiple reports during the last week of each month. The program takes less than 10 minutes to produce each report. The company rarely uses the program to generate reports outside of the last week of each month. The company wants to generate reports in the least amount of time when the reports are requested.

Which solution will meet these requirements MOST cost-effectively?

- A. Run the program by using Amazon EC2 On-Demand Instances. Create an Amazon EventBridge rule to start the EC2 instances when reports are requested. Run the EC2 instances continuously during the last week of each month.
- B. Run the program in AWS Lambda. Create an Amazon EventBridge rule to run a Lambda function when reports are requested. **Most Voted**
- C. Run the program in Amazon Elastic Container Service (Amazon ECS). Schedule Amazon ECS to run the program when reports are requested.
- D. Run the program by using Amazon EC2 Spot Instances. Create an Amazon EventBridge rule to start the EC2 instances when reports are requested. Run the EC2 instances continuously during the last week of each month.

Correct Answer: B

Community vote distribution

B (82%)

C (18%)

Comments

Andy_09 **Highly Voted** 2 years, 4 months ago

Option B

upvoted 6 times

ogerber 2 years, 3 months ago

not sure because it says that the program produces several reports, and each takes less than 10 min. I am voting for option A

upvoted 2 times

1dd 2 years, 3 months ago

Lambda takes duration--> 15 minutes

upvoted 1 times

FZA24 2 years, 3 months ago

each lambda triggering produces a report in less than 10 mins.

upvoted 4 times

LuongTo 1 year, 7 months ago

can lambda host the "custom report building program" ?

upvoted 2 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: B

A - The task only runs for 10 minutes per report. Why would you need an EC2 instance?

B - No infrastructure to manage. Love it.

C - Works, but cost more than B and is more complex to set up than B.

D - More cost-effective than A, but spot instances can be interrupted, probably causing delays in report generation.

upvoted 1 times

agbor_tambe 1 year, 8 months ago

Selected Answer: B

B is correct

upvoted 2 times

EdricHoang 1 year, 11 months ago

Selected Answer: B

"The program takes less than 10 minutes to produce each report. The company rarely uses the program to generate reports outside of the last week of each month The company wants to generate reports in the least amount of time when the reports are requested."

I go for B because of this

upvoted 2 times

Scheldon 1 year, 12 months ago

Selected Answer: C

AnswerC

Option A and D are saying about Running the EC2 instances countinously during the last week of each month which is not necessary from my understanding and will be not so cheap.

Option B - 10min per report and we have couple of reports. So it looks like program is running for at least 20 min so in theory Lambda is not useful here

Option C - ECS is allowing us to run Fargate which will allow to run program for more then 15 min, hence all reports which program is preparing should be created. I'm not sure but I think ECS API is allowing to run task on demand/request.

upvoted 2 times

MatAlves 1 year, 8 months ago

It's "15 minutes per execution", which is enough to produce a report and be ready for another invocation.

upvoted 2 times

asdfcdsxdxc 2 years, 3 months ago

Selected Answer: B

B is corre ct

upvoted 2 times

NayeraB 2 years, 3 months ago

Selected Answer: B

B..maybe?

upvoted 2 times

A company is designing a tightly coupled high performance computing (HPC) environment in the AWS Cloud. The company needs to include features that will optimize the HPC environment for networking and storage.

Which combination of solutions will meet these requirements? (Choose two.)

- A. Create an accelerator in AWS Global Accelerator. Configure custom routing for the accelerator.
- B. Create an Amazon FSx for Lustre file system. Configure the file system with scratch storage. **Most Voted**
- C. Create an Amazon CloudFront distribution. Configure the viewer protocol policy to be HTTP and HTTPS.
- D. Launch Amazon EC2 instances. Attach an Elastic Fabric Adapter (EFA) to the instances. **Most Voted**
- E. Create an AWS Elastic Beanstalk deployment to manage the environment.

Correct Answer: BD

Community vote distribution

BD (100%)

Comments

Andy_09 **Highly Voted** 1 year, 10 months ago

Options BD

upvoted 7 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: BD

A - Global Accelerator is not designed for tightly coupled HPC environments. It focuses on improving performance for geographically distributed users, not for low-latency intra-cluster communication.

B - FSx for Lustre is specifically designed for HPC workloads and integrates well with Amazon S3 for high-throughput and low-latency access to data.

C - CloudFront does not address the low-latency, high-bandwidth networking or storage requirements of tightly coupled HPC environments.

D - EFA is a network interface that provides low-latency, high-throughput connectivity for tightly coupled HPC workloads, such as those using MPI (Message Passing Interface).

E - Elastic Beanstalk is used for deploying and managing web applications and services, which has nothing to do with HPC.

upvoted 2 times

seetpt 1 year, 9 months ago

Selected Answer: BD

BD looks right

upvoted 3 times

asdfcdsxdhc 1 year, 9 months ago

Selected Answer: BD

Elastic Fabric Adapter (EFA)

- Improved ENA for HPC, only works for Linux
- Great for inter-node communications, tightly coupled workloads
- Leverages Message Passing Interface (MPI) standard
- Bypasses the underlying Linux OS to provide low-latency, reliable transport

upvoted 4 times


A company needs a solution to prevent photos with unwanted content from being uploaded to the company's web application. The solution must not involve training a machine learning (ML) model.

Which solution will meet these requirements?

- A. Create and deploy a model by using Amazon SageMaker Autopilot. Create a real-time endpoint that the web application invokes when new photos are uploaded.
- B. Create an AWS Lambda function that uses Amazon Rekognition to detect unwanted content. Create a Lambda function URL that the web application invokes when new photos are uploaded. **Most Voted**
- C. Create an Amazon CloudFront function that uses Amazon Comprehend to detect unwanted content. Associate the function with the web application.
- D. Create an AWS Lambda function that uses Amazon Rekognition Video to detect unwanted content. Create a Lambda function URL that the web application invokes when new photos are uploaded.

Correct Answer: B

Community vote distribution



Comments

LeonSauveterre **Highly Voted** 1 year, 5 months ago

Selected Answer: B

You've made it this far, and you're all gonna make it.

For your quick reference:

Vision and Image Processing

1. Amazon Rekognition: Detects objects, scenes, text, faces, and activities in images/videos.
2. Amazon Textract: Extracts text, tables, and forms from scanned documents.

Speech and Language

1. Amazon Polly: Converts text to lifelike speech.
2. Amazon Transcribe: Converts speech to text for audio files.
3. Amazon Translate: Yeah, translate.
4. Amazon Comprehend: Extracts insights like sentiment, entities, and key phrases from text.

Machine Learning Development

1. Amazon SageMaker: Comprehensive ML platform for building, training, and deploying models.
2. Amazon SageMaker Autopilot: Automatically builds, trains, and tunes ML models, basically simplified model creation for non-experts.
3. Amazon SageMaker JumpStart: Provides pre-built ML models and solutions for quick deployment.

upvoted 10 times

LeonSauveterre 1 year, 5 months ago

Chatbots and Conversations

1. Amazon Lex: Builds conversational AI (chatbots) using speech and text.
2. Amazon Kendra: Intelligent search service for documents and FAQs.

Personalization and Forecasting

1. Amazon Personalize: Provides personalized recommendations using ML.
2. Amazon Forecast: Time-series forecasting for demand, sales, and resource planning.

Data Insights

1. Amazon Fraud Detector: Detects fraudulent activity using ML models.
2. Amazon Lookout for Metrics: Detects anomalies in metrics automatically.
3. Amazon Lookout for Vision: Detects manufacturing defects in images.

Robotics and Edge AI

1. AWS DeepRacer: Provides a platform for learning reinforcement learning using autonomous racing.
2. AWS Panorama: Adds computer vision to on-premises cameras.

upvoted 5 times

Andy_09 **Highly Voted** 1 year, 10 months ago

Option B

upvoted 8 times

zdi561 **Most Recent** 1 year, 3 months ago

Selected Answer: A

Amazon Rekognition is a visual analysis service that includes both Amazon Rekognition Image and Amazon Rekognition Video. Amazon Rekognition Image analyzes images, while Amazon Rekognition Video analyzes videos.

upvoted 1 times

HarryLopez 1 year, 9 months ago

Selected Answer: B

Rekognition: for image and video analysis

Comprehend: natural language processing model for uncovering insights and connections in text

Sagemaker Autopilot: feature set that simplifies and accelerates and automates the various stages of the machine learning workflow

upvoted 6 times

NayeraB 1 year, 9 months ago

Selected Answer: B

B is correct

upvoted 3 times

A company uses AWS to run its ecommerce platform. The platform is critical to the company's operations and has a high volume of traffic and transactions. The company configures a multi-factor authentication (MFA) device to secure its AWS account root user credentials. The company wants to ensure that it will not lose access to the root user account if the MFA device is lost.

Which solution will meet these requirements?

- A. Set up a backup administrator account that the company can use to log in if the company loses the MFA device.
- B. Add multiple MFA devices for the root user account to handle the disaster scenario. **Most Voted**
- C. Create a new administrator account when the company cannot access the root account.
- D. Attach the administrator policy to another IAM user when the company cannot access the root account.

Correct Answer: B

Community vote distribution

B (90%)

10%

Comments

hajra313 **Highly Voted** 2 years, 4 months ago

B. Add multiple MFA devices for the root user account to handle the disaster scenario.

By adding multiple MFA devices for the root user account, the company ensures that it can still access the account even if one MFA device is lost. This approach provides a backup for authentication, addressing the concern of losing access to the root user account if the MFA device is lost.

upvoted 11 times

Tatai2015 2 years ago

<https://docs.aws.amazon.com/IAM/latest/UserGuide/root-user-best-practices.html>

upvoted 2 times

91157fe **Most Recent** 9 months ago

Selected Answer: A

Add multiple MFA devices for the root user account:

AWS does not support multiple MFA devices for the root user. Only one MFA device can be associated at a time.

upvoted 1 times

ad2dj28 1 year, 3 months ago

Selected Answer: B

Your AWS account root user and IAM users can register up to eight MFA devices of any type. Registering multiple MFA devices can provide flexibility and help you reduce the risk of access interruption if a device is lost or broken.

https://docs.aws.amazon.com/IAM/latest/UserGuide/id_credentials_mfa.html

upvoted 1 times

Scheldon 1 year, 12 months ago

Selected Answer: B

Answer B

Because a root user can perform privileged actions, it's crucial to add MFA for the root user as a second authentication factor in addition to the email address and password as sign-in credentials. We strongly recommend enabling multiple MFA for your root user credentials to provide additional flexibility and resiliency in your security strategy. You can register up to eight MFA devices of any combination of the currently supported MFA types with your AWS account root user.

<https://docs.aws.amazon.com/IAM/latest/UserGuide/root-user-best-practices.html#ru-bp-mfa>

upvoted 3 times

asdfcdsxdcf 2 years, 3 months ago

Selected Answer: B

b looks correct
upvoted 2 times

NayeraB 2 years, 3 months ago

Selected Answer: B

I'd go for B
upvoted 3 times

Andy_09 2 years, 4 months ago

Option B
upvoted 4 times

A social media company is creating a rewards program website for its users. The company gives users points when users create and upload videos to the website. Users redeem their points for gifts or discounts from the company's affiliated partners. A unique ID identifies users. The partners refer to this ID to verify user eligibility for rewards.

The partners want to receive notification of user IDs through an HTTP endpoint when the company gives users points. Hundreds of vendors are interested in becoming affiliated partners every day. The company wants to design an architecture that gives the website the ability to add partners rapidly in a scalable way.

Which solution will meet these requirements with the LEAST implementation effort?

- A. Create an Amazon Timestream database to keep a list of affiliated partners. Implement an AWS Lambda function to read the list. Configure the Lambda function to send user IDs to each partner when the company gives users points.
- B. Create an Amazon Simple Notification Service (Amazon SNS) topic. Choose an endpoint protocol. Subscribe the partners to the topic. Publish user IDs to the topic when the company gives users points. **Most Voted**
- C. Create an AWS Step Functions state machine. Create a task for every affiliated partner. Invoke the state machine with user IDs as input when the company gives users points.
- D. Create a data stream in Amazon Kinesis Data Streams. Implement producer and consumer applications. Store a list of affiliated partners in the data stream. Send user IDs when the company gives users points.

Correct Answer: B

Community vote distribution

B (100%)

Comments

kempes **Highly Voted** 2 years, 4 months ago

Selected Answer: B

SNS is designed for precisely this kind of use case. It allows you to publish messages to a topic, which can then be delivered to multiple subscribers. Partners can subscribe to the SNS topic using an HTTP endpoint as the protocol, which meets the requirement to notify partners via an HTTP endpoint. This approach is highly scalable and requires the least implementation effort because it leverages managed services without the need for custom logic to manage subscriptions or deliver notifications.

upvoted 13 times

hajra313 **Highly Voted** 2 years, 4 months ago

Option A involves creating an Amazon Timestream database to store affiliated partners and implementing an AWS Lambda function to read the list and send user IDs to each partner. While this approach can work, it involves more implementation effort than the Amazon SNS solution. It requires setting up and managing a database, as well as configuring the Lambda function to send notifications to partners. The Amazon SNS solution provides a simpler and more scalable approach for rapidly adding partners and notifying them when users receive points. so answer is B

upvoted 9 times

Scheldon **Most Recent** 1 year, 12 months ago

Selected Answer: B

AnswerB

Sending Notification to multiple subscribers = SNS

Amazon Simple Notification Service (Amazon SNS) is a managed service that provides message delivery from publishers to subscribers (also known as producers and consumers). Publishers communicate asynchronously with subscribers by sending messages to a topic, which is a logical access point and communication channel. Clients can subscribe to the SNS topic and receive published messages using a supported endpoint type, such as Amazon Data Firehose, Amazon SQS, AWS Lambda, HTTP, email, mobile push notifications, and mobile text messages (SMS).
<https://docs.aws.amazon.com/sns/latest/dg/welcome.html>

upvoted 1 times

NayeraB 2 years, 3 months ago

Selected Answer: B

This is a perfect SNS use case
upvoted 3 times

jjcode 2 years, 4 months ago

The answer is B, create an SNS topic one subscriptions you can make is HTTP, This completely addresses the question objective.
upvoted 2 times

Andy_09 2 years, 4 months ago

Option A
upvoted 1 times

FlyingHawk 1 year, 4 months ago

Timestream is optimized for time-series data, not for managing partner subscriptions.
upvoted 1 times

A company needs to extract the names of ingredients from recipe records that are stored as text files in an Amazon S3 bucket. A web application will use the ingredient names to query an Amazon DynamoDB table and determine a nutrition score.

The application can handle non-food records and errors. The company does not have any employees who have machine learning knowledge to develop this solution.

Which solution will meet these requirements MOST cost-effectively?

A. Use S3 Event Notifications to invoke an AWS Lambda function when PutObject requests occur. Program the Lambda function to analyze the object and extract the ingredient names by using Amazon Comprehend. Store the Amazon Comprehend output in the DynamoDB table. **Most Voted**

B. Use an Amazon EventBridge rule to invoke an AWS Lambda function when PutObject requests occur. Program the Lambda function to analyze the object by using Amazon Forecast to extract the ingredient names. Store the Forecast output in the DynamoDB table.

C. Use S3 Event Notifications to invoke an AWS Lambda function when PutObject requests occur. Use Amazon Polly to create audio recordings of the recipe records. Save the audio files in the S3 bucket. Use Amazon Simple Notification Service (Amazon SNS) to send a URL as a message to employees. Instruct the employees to listen to the audio files and calculate the nutrition score. Store the ingredient names in the DynamoDB table.

D. Use an Amazon EventBridge rule to invoke an AWS Lambda function when a PutObject request occurs. Program the Lambda function to analyze the object and extract the ingredient names by using Amazon SageMaker. Store the inference output from the SageMaker endpoint in the DynamoDB table.

Correct Answer: A

Community vote distribution

A (100%)

Comments

seetpt **Highly Voted** 2 years, 3 months ago

Selected Answer: A

A correct

upvoted 5 times

Besisco **Most Recent** 1 year, 4 months ago

Selected Answer: A

I'm fascinated by how self-descriptive the names of AWS ML tools are.

upvoted 1 times

Danilus 1 year, 7 months ago

Selected Answer: A

Key: Most cost-effective solution

Key: Extract the names of ingredients from recipe records

Key: The company does not have any employees with machine learning knowledge

Not B because Forecast is used for making predictions in time series, not for extracting text or ingredients.

Not C because it would be too expensive and involve too much operational overhead.

Not D because SageMaker is designed for creating custom machine learning models, and the company lacks employees with machine learning knowledge.

Conclusion: The answer is A because Comprehend is used for natural language processing (NLP), which in this case can extract ingredient names

from the recipes effectively.

so the answer is A because comprehend is used for nlp (natural languages processing) which is in this case use for to extract the names of the recipe
upvoted 3 times

Scheldon 1 year, 12 months ago

Selected Answer: A

AnswerA

Amazon Comprehend uses natural language processing (NLP) to extract insights about the content of documents. It develops insights by recognizing the entities, key phrases, language, sentiments, and other common elements in a document. Use Amazon Comprehend to create new products based on understanding the structure of documents. For example, using Amazon Comprehend you can search social networking feeds for mentions of products or scan an entire document repository for key phrases.

<https://docs.aws.amazon.com/comprehend/latest/dg/what-is.html>

upvoted 3 times

fae3297 2 years ago

Selected Answer: A

A seems right

B. Forecast is time-series

C. Using Polly to create audio recordings just to make your employees listen to them seems inefficient to say the least

D. Question asks for no ML. SageMaker = ML

upvoted 3 times

seetpt 2 years, 3 months ago

Selected Answer: A

A is correct

upvoted 3 times

asdfcdsxdfc 2 years, 3 months ago

shouldn't it be A?

upvoted 3 times

A company needs to create an AWS Lambda function that will run in a VPC in the company's primary AWS account. The Lambda function needs to access files that the company stores in an Amazon Elastic File System (Amazon EFS) file system. The EFS file system is located in a secondary AWS account. As the company adds files to the file system, the solution must scale to meet the demand.

Which solution will meet these requirements MOST cost-effectively?

- A. Create a new EFS file system in the primary account. Use AWS DataSync to copy the contents of the original EFS file system to the new EFS file system.
- B. Create a VPC peering connection between the VPCs that are in the primary account and the secondary account. **Most Voted**
- C. Create a second Lambda function in the secondary account that has a mount that is configured for the file system. Use the primary account's Lambda function to invoke the secondary account's Lambda function.
- D. Move the contents of the file system to a Lambda layer. Configure the Lambda layer's permissions to allow the company's secondary account to use the Lambda layer.

Correct Answer: B

Community vote distribution

B (100%)

Comments

lenotc **Highly Voted** 2 years, 3 months ago

Selected Answer: B

- B -> VPC peering allows the Lambda access secondary account securely and efficiently
- A -> redundancy
- C -> additional complexity
- D -> sharing code libraries

upvoted 9 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: B

- A - Duplicating the EFS file system leads to unnecessary storage costs and additional complexity.
- B - VPC peering connection allows seamless access without additional manual scaling.
- C - Well... I don't know, but you could've directly enabled access to the EFS file system. Why make it even more complex?
- D - Lambda layers are intended for shared libraries or dependencies. You would be crazy to store large amounts of data there.

upvoted 1 times

Scheldon 1 year, 12 months ago

Selected Answer: B

AnswerB

You can configure a function to mount an Amazon EFS file system in another AWS account. Before you mount the file system, you must ensure the following:

VPC peering must be configured, and appropriate routes must be added to the route tables in each VPC.

.
. .
.

<https://docs.aws.amazon.com/lambda/latest/dg/configuration-file-system-cross-account.html>

upvoted 1 times

Nm55569 2 years ago

Selected Answer: B

<https://docs.aws.amazon.com/lambda/latest/dg/configuration-file-system.html#configuration-file-system-cross-account>

upvoted 1 times

osmk 2 years, 3 months ago

Selected Answer: B

<https://docs.aws.amazon.com/efs/latest/ug/efs-different-vpc.html>

upvoted 3 times

asdfcdsxdxc 2 years, 3 months ago

Shouldn't it be B?

upvoted 1 times

1dd 2 years, 3 months ago

I thinks AWS DataSync less costly

upvoted 1 times

[Removed] 2 years, 2 months ago

setting up a peering connection is free. same for data transfer in the same AZ. data sync at the end of the day cost \$\$\$ to move data.

upvoted 3 times

Scheldon 1 year, 12 months ago

When you will SyncData you need to have secondary storage for which you need to pay so it is not cheap solution.

upvoted 1 times

A financial company needs to handle highly sensitive data. The company will store the data in an Amazon S3 bucket. The company needs to ensure that the data is encrypted in transit and at rest. The company must manage the encryption keys outside the AWS Cloud.

Which solution will meet these requirements?

- A. Encrypt the data in the S3 bucket with server-side encryption (SSE) that uses an AWS Key Management Service (AWS KMS) customer managed key.
- B. Encrypt the data in the S3 bucket with server-side encryption (SSE) that uses an AWS Key Management Service (AWS KMS) AWS managed key.
- C. Encrypt the data in the S3 bucket with the default server-side encryption (SSE).
- D. Encrypt the data at the company's data center before storing the data in the S3 bucket. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Johnnopping101 **Highly Voted** 1 year, 10 months ago

You get to do it, keep moving...

upvoted 6 times

srcntpc **Most Recent** 1 year ago

Selected Answer: D

D for sure

upvoted 1 times

Scheldon 1 year, 12 months ago

Selected Answer: D

AnswerD

Hence we need to encrypt data not only during the rest but during the transfer as well, we need execute client-side encryption. SSE will only secure data during rest hence we can eliminate A,B and C.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/UsingClientSideEncryption.html>

upvoted 1 times

chasingsummer 2 years, 2 months ago

Selected Answer: D

Given the requirement to manage encryption keys outside the AWS Cloud, option D is the most suitable solution, despite not directly utilizing AWS's native encryption services like SSE with AWS KMS. Instead, it leverages external encryption mechanisms controlled by the company.

upvoted 4 times

rondelldell 2 years, 2 months ago

A Key is safe but came from the customer

upvoted 2 times

Mikado211 2 years, 2 months ago

Selected Answer: D

A, B and C need to have the key stored in AWS cloud.

D is correct.

upvoted 3 times

osmk 2 years, 3 months ago

Selected Answer: D

Client-side encryption – You encrypt your data client-side and upload the encrypted data to Amazon S3. In this case, you manage the encryption process, encryption keys, and related tools.<https://docs.aws.amazon.com/AmazonS3/latest/userguide/UsingClientSideEncryption.html>

upvoted 3 times

giovanna_mag 2 years, 3 months ago

Selected Answer: D

For me it's D, it's the only one that provides encryption also in transit

upvoted 2 times

asdfcdsxdcf 2 years, 3 months ago

A looks correct

upvoted 3 times

A company wants to run its payment application on AWS. The application receives payment notifications from mobile devices. Payment notifications require a basic validation before they are sent for further processing.

The backend processing application is long running and requires compute and memory to be adjusted. The company does not want to manage the infrastructure.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create an Amazon Simple Queue Service (Amazon SQS) queue. Integrate the queue with an Amazon EventBridge rule to receive payment notifications from mobile devices. Configure the rule to validate payment notifications and send the notifications to the backend application. Deploy the backend application on Amazon Elastic Kubernetes Service (Amazon EKS) Anywhere. Create a standalone cluster.
- B. Create an Amazon API Gateway API. Integrate the API with an AWS Step Functions state machine to receive payment notifications from mobile devices. Invoke the state machine to validate payment notifications and send the notifications to the backend application. Deploy the backend application on Amazon Elastic Kubernetes Service (Amazon EKS). Configure an EKS cluster with self-managed nodes.
- C. Create an Amazon Simple Queue Service (Amazon SQS) queue. Integrate the queue with an Amazon EventBridge rule to receive payment notifications from mobile devices. Configure the rule to validate payment notifications and send the notifications to the backend application. Deploy the backend application on Amazon EC2 Spot Instances. Configure a Spot Fleet with a default allocation strategy.
- D. Create an Amazon API Gateway API. Integrate the API with AWS Lambda to receive payment notifications from mobile devices. Invoke a Lambda function to validate payment notifications and send the notifications to the backend application. Deploy the backend application on Amazon Elastic Container Service (Amazon ECS). Configure Amazon ECS with an AWS Fargate launch type. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Mikado211 **Highly Voted** 2 years, 2 months ago

Selected Answer: D

We want to have least overhead and no infrastructure (aka no server).
So no infrastructure == not C
least overhead == ECS better than EKS == not B and not A

Fargate is serverless so D is still valid.

So the answer is D.

upvoted 5 times

Cpso **Most Recent** 1 year, 6 months ago

Selected Answer: D

D= less operation. But should have SQS to buffer a job.
upvoted 2 times

Scheldon 1 year, 12 months ago

Selected Answer: D

AnswerD

I would go with Lambda function to do basic validation it should not take more then 15min hence Lambda is perfect for that job. Then we have information that backend processing application is long running and we need to make compute and memory adjustment, and everything need to be automatic as company do not want to manage infrastucture. In that situation Fargate with ECS will be ideal as it can run background application for every paymant separetly we need only adjust amunt of resources in use. More payments more application running, mmore resources in use and oposite,

upvoted 1 times

sandordini 2 years, 1 month ago

Selected Answer: D

Lot of grip in this question, but to keep it short:

No infra: B, C Out.

EKS Anywhere: Onprem + AWS: Not needed.

ECS Fargate: Serverless, Least Ops Overhead, SQS fine for the queue, Lambda good for basic validation.

upvoted 3 times

seetpt 2 years, 3 months ago

Selected Answer: D

D is correct

upvoted 4 times

asdfcdsxdfc 2 years, 3 months ago

shouldn't it be D?

upvoted 3 times

A solutions architect is designing a user authentication solution for a company. The solution must invoke two-factor authentication for users that log in from inconsistent geographical locations, IP addresses, or devices. The solution must also be able to scale up to accommodate millions of users.

Which solution will meet these requirements?

- A. Configure Amazon Cognito user pools for user authentication. Enable the risk-based adaptive authentication feature with multifactor authentication (MFA). **Most Voted**
- B. Configure Amazon Cognito identity pools for user authentication. Enable multi-factor authentication (MFA).
- C. Configure AWS Identity and Access Management (IAM) users for user authentication. Attach an IAM policy that allows the AllowManageOwnUserMFA action.
- D. Configure AWS IAM Identity Center (AWS Single Sign-On) authentication for user authentication. Configure the permission sets to require multi-factor authentication (MFA).

Correct Answer: A

Community vote distribution

A (100%)

Comments

mk168898 1 year, 7 months ago

Selected Answer: A

Picked A because cognito user pool => authentication
upvoted 1 times

Scheldon 1 year, 12 months ago

Selected Answer: A

AnswerA

Adaptive authentication

Amazon Cognito can review location and device information from your users' sign-in requests and apply an automatic response to secure the user accounts in your user pool against suspicious activity.

When you activate advanced security, Amazon Cognito assigns a risk score to user activity. You can assign an automatic response to suspicious activity: you can Require MFA, Block sign-in, or just log the activity details and risk score. You can also automatically send email messages that notify your user of the suspicious activity so that they can reset their password or take other self-guided actions.

<https://docs.aws.amazon.com/cognito/latest/developerguide/cognito-user-pool-settings-adaptive-authentication.html>

upvoted 3 times

osmk 2 years, 3 months ago

Selected Answer: A

With adaptive authentication, you can configure your user pool to require second factor authentication in response to an increased risk level. To add adaptive authentication to your user pool, see Adding advanced security to a user pool.<https://docs.aws.amazon.com/cognito/latest/developerguide/cognito-user-pool-settings-advanced-security.html>

upvoted 3 times

lenotc 2 years, 3 months ago

Selected Answer: A

A is correct

B is wrong because is desing for temporary credentials

upvoted 2 times

allanm 10 months ago

B is also wrong because it's identity pool.

User pool = Authentication

Identity pool = Authorisation

upvoted 1 times

giovanna_mag 2 years, 3 months ago

Selected Answer: A

I belive it's A

upvoted 1 times

Tatai2015 2 years ago

A

<https://docs.aws.amazon.com/cognito/latest/developerguide/cognito-user-pool-settings-adaptive-authentication.html>

upvoted 1 times

xBUGx 2 years, 3 months ago

accommodate millions of users and GEO, IP, etc. I think A

upvoted 3 times

A company has an Amazon S3 data lake. The company needs a solution that transforms the data from the data lake and loads the data into a data warehouse every day. The data warehouse must have massively parallel processing (MPP) capabilities.

Data analysts then need to create and train machine learning (ML) models by using SQL commands on the data. The solution must use serverless AWS services wherever possible.

Which solution will meet these requirements?

- A. Run a daily Amazon EMR job to transform the data and load the data into Amazon Redshift. Use Amazon Redshift ML to create and train the ML models.
- B. Run a daily Amazon EMR job to transform the data and load the data into Amazon Aurora Serverless. Use Amazon Aurora ML to create and train the ML models.
- C. Run a daily AWS Glue job to transform the data and load the data into Amazon Redshift Serverless. Use Amazon Redshift ML to create and train the ML models. **Most Voted**
- D. Run a daily AWS Glue job to transform the data and load the data into Amazon Athena tables. Use Amazon Athena ML to create and train the ML models.

Correct Answer: C

Community vote distribution

C (86%)

10%

Comments

Mikado211 **Highly Voted** 2 years, 2 months ago

Selected Answer: C

Data warehouse ==> Redshift

Without additional informations both EMR and Glue Jobs can work.

Since the question asks to use serverless as much as possible, Redshift Serverless is a better solution.

C

upvoted 6 times

BatVanyo **Highly Voted** 2 years, 2 months ago

Selected Answer: C

Neither A, nor B explicitly say "EMR serverless" which is a new AWS offering, so I exclude these two.

MPP goes hand in hand with Redshift, so D is also incorrect.

This leaves C the only possible serverless option here.

upvoted 5 times

In2718 **Most Recent** 10 months, 1 week ago

Selected Answer: A

Gemini AI answer is - A - For massively parallel processing (MPP) capabilities, Amazon EMR generally offers more robust and flexible options than AWS Glue. Then use Redshift serverless for the rest.

upvoted 1 times

sk1974 1 year, 3 months ago

Selected Answer: B

Any reason it cannot be B ?

upvoted 2 times

mk168898 1 year, 7 months ago

Selected Answer: C

Data Warehouse => redshift
Use AWS Services wherever possible => Redshift serverless

upvoted 2 times

rondellidell 2 years, 2 months ago

A

Amazon EMR Serverless is a deployment option for Amazon EMR that provides a serverless runtime environment. This simplifies the operation of analytics applications that use the latest open-source frameworks, such as Apache Spark and Apache Hive. With EMR Serverless, you don't have to configure, optimize, secure, or operate clusters to run applications with these frameworks.

EMR Serverless helps you avoid over- or under-provisioning resources for your data processing jobs. EMR Serverless automatically determines the resources that the application needs, gets these resources to process your jobs, and releases the resources when the jobs finish. For use cases where applications need a response within seconds, such as interactive data analysis, you can pre-initialize the resources that the application needs when you create the application.

upvoted 2 times

1dd 2 years, 3 months ago

Selected Answer: C

Option C

upvoted 2 times

1dd 2 years, 3 months ago

EMR works with big data transfer

upvoted 1 times

1dd 2 years, 3 months ago

MPP --> use Redshift so eliminate B,D

As it required Serverless services --> Glue

upvoted 1 times

1dd 2 years, 3 months ago

A have no serverless

C is the answer

upvoted 1 times

seetpt 2 years, 3 months ago

Selected Answer: C

C is correct

upvoted 3 times

asdfcdsxdfc 2 years, 3 months ago

should be C

upvoted 2 times

A company runs containers in a Kubernetes environment in the company's local data center. The company wants to use Amazon Elastic Kubernetes Service (Amazon EKS) and other AWS managed services. Data must remain locally in the company's data center and cannot be stored in any remote site or cloud to maintain compliance.

Which solution will meet these requirements?

- A. Deploy AWS Local Zones in the company's data center.
- B. Use an AWS Snowmobile in the company's data center.
- C. Install an AWS Outposts rack in the company's data center. **Most Voted**
- D. Install an AWS Snowball Edge Storage Optimized node in the data center.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Mikado211 **Highly Voted** 2 years, 2 months ago

Selected Answer: C

Outpost is a service where AWS has physical servers in your datacenter.

C

upvoted 8 times

asdfcdsxdcf **Highly Voted** 2 years, 3 months ago

Selected Answer: C

C looks correct

upvoted 5 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: C

A - AWS Local Zones provide low-latency access to AWS services by placing AWS infrastructure closer to customers, but the data is still stored in an AWS-managed facility.

B - AWS Snowmobile is used for large-scale data migrations to the cloud, not for running workloads or storing data locally.

C - AWS Outposts allows you to run AWS services, including Amazon EKS, in your own data center while keeping data locally.

D - AWS Snowball Edge is a portable edge computing and data storage device, but it's not suitable for running Kubernetes or AWS-managed services like Amazon EKS.

upvoted 3 times

mk168898 1 year, 7 months ago

Data must remain locally in the company's data center => AWS outpost

upvoted 2 times

JonJon03 1 year, 11 months ago

EKS on SnowBALL could be an option. EKS on SnowMOBILE isn't as it's used for data transfer mostly.

upvoted 2 times

Scheldon 1 year, 12 months ago

Selected Answer: C

AnswerC

AWS Outpost = Bring AWS cloud to your DataCentre, which we need in described scenario

upvoted 4 times

seetpt 2 years, 3 months ago

Selected Answer: C

C is correct

upvoted 2 times

A social media company has workloads that collect and process data. The workloads store the data in on-premises NFS storage. The data store cannot scale fast enough to meet the company's expanding business needs. The company wants to migrate the current data store to AWS.

Which solution will meet these requirements MOST cost-effectively?

- A. Set up an AWS Storage Gateway Volume Gateway. Use an Amazon S3 Lifecycle policy to transition the data to the appropriate storage class.
- B. Set up an AWS Storage Gateway Amazon S3 File Gateway. Use an Amazon S3 Lifecycle policy to transition the data to the appropriate storage class. **Most Voted**
- C. Use the Amazon Elastic File System (Amazon EFS) Standard-Infrequent Access (Standard-IA) storage class. Activate the infrequent access lifecycle policy.
- D. Use the Amazon Elastic File System (Amazon EFS) One Zone-Infrequent Access (One Zone-IA) storage class. Activate the infrequent access lifecycle policy.

Correct Answer: B

Community vote distribution

B (96%) 4%

Comments

alawada **Highly Voted** 2 years, 2 months ago

Selected Answer: B

This solution meets the requirements most cost-effectively because it enables the company to migrate its on-premises NFS data store to AWS without changing the existing applications or workflows. AWS Storage Gateway is a hybrid cloud storage service that provides seamless and secure integration between on-premises and AWS storage. Amazon S3 File Gateway is a type of AWS Storage Gateway that provides a file interface to Amazon S3, with local caching for low-latency access. By setting up an Amazon S3 File Gateway, the company can store and retrieve files as objects in Amazon S3 using standard file protocols such as NFS.

upvoted 10 times

allanm **Most Recent** 10 months ago

Selected Answer: B

Anyone else think this question is weird? It asks for Migration of data but Answers are all about lifecycle policies?

upvoted 1 times

0c21057 1 year, 1 month ago

Selected Answer: B

Why other options are incorrect:

A. Volume Gateway is designed for block storage and backup use cases, not for NFS file storage migration. It would be less cost-effective and more complex to implement.

C. Amazon EFS Standard-IA would be more expensive than S3 for this use case, as it's designed for file systems that need high availability across multiple AZs.

D. Amazon EFS One Zone-IA, while cheaper than Standard-IA, is still more expensive than S3 for this type of data storage requirement. It's also designed for file systems rather than object storage.

The S3 File Gateway solution is most cost-effective because:

S3 storage is generally cheaper than EFS

Lifecycle policies can automatically move data to cheaper storage tiers

It maintains NFS compatibility while providing scalability

It minimizes the need for application changes during migration

upvoted 3 times

Gorkha 1 year, 4 months ago

Selected Answer: B

as their is nfs = file gateway

upvoted 1 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: B

I'm a little confused by option C & D. How does infrequent access have anything to do with this?

upvoted 1 times

mk168898 1 year, 7 months ago

Selected Answer: B

S3 File Gateway => Best for NFS-like file storage workloads

upvoted 2 times

hharbiordun85 1 year, 8 months ago

Answer: C

upvoted 1 times

Scheldon 1 year, 12 months ago

Selected Answer: D

AnswerD

Taking into consideration that we are talking about social media company probably storing a lot of quite small files i would say it cannot be Option A or B.

For example

Amazon S3 File Gateway pricing

Storage Pricing

Storage Pricing Stored and billed as Amazon S3 objects.

Request Pricing

Data written to AWS storage by your gateway \$0.01 per GB+

File storage in S3 Billed as Amazon S3 requests.

It can be quite expensive, especially when we will be working on small files.

I would go with EFS with OneZone-IA (option D) which should be less expensive taking into consideration that we are paying only for Storage and Data Transfer (per GB).

But to be honest we need more information to device which solution will be better.

upvoted 1 times

EdricHoang 1 year, 11 months ago

It say migrating data into cloud, EFS does not satisfy that. You're right, It needs more information to make your choice to be a better answer.

upvoted 2 times

alawada 2 years, 2 months ago

Selected Answer: B

yeah B

upvoted 2 times

seetpt 2 years, 3 months ago

Selected Answer: B

I think B too

upvoted 3 times

asdfcdsxdfc 2 years, 3 months ago

Selected Answer: B

B looks correct

upvoted 2 times


A company uses high concurrency AWS Lambda functions to process a constantly increasing number of messages in a message queue during marketing events. The Lambda functions use CPU intensive code to process the messages. The company wants to reduce the compute costs and to maintain service latency for its customers.

Which solution will meet these requirements?

- A. Configure reserved concurrency for the Lambda functions. Decrease the memory allocated to the Lambda functions.
- B. Configure reserved concurrency for the Lambda functions. Increase the memory according to AWS Compute Optimizer recommendations.
- C. Configure provisioned concurrency for the Lambda functions. Decrease the memory allocated to the Lambda functions.
- D. Configure provisioned concurrency for the Lambda functions. Increase the memory according to AWS Compute Optimizer recommendations. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

JonJon03 **Highly Voted** 1 year, 11 months ago

Selected Answer: B

Reserved concurrency — It guarantees the maximum number of concurrent instances for the function which can be invoked. When a function has been given a reserved concurrency configuration then no other lambda function within the same AWS account and region can use that concurrency. There is no charge for configuring reserved concurrency for a function.
Provisioned concurrency — This concurrency initializes a requested number of execution environments so that they are prepared to respond immediately to your function's invocations. Note that configuring provisioned concurrency incurs charges to your AWS account.
upvoted 11 times

SnaxAttax **Most Recent** 6 months, 2 weeks ago

Selected Answer: D

Provisioned concurrency addresses latency concerns and compute optimizer addresses cost. D is best answer here.
upvoted 1 times

In2718 10 months, 1 week ago

Selected Answer: B

Not a very well written question, but I will go with answer B with this as the deciding factor - The company wants to reduce the compute costs
upvoted 1 times

tch 1 year, 3 months ago

Selected Answer: B

• Reserved concurrency – This represents the maximum number of concurrent instances allocated to your function. When a function has reserved concurrency, no other function can use that concurrency. Reserved concurrency is useful for ensuring that your most critical functions always have enough concurrency to handle incoming requests. Configuring reserved concurrency for a function incurs no additional charges.
upvoted 1 times

Gorkha 1 year, 4 months ago

Selected Answer: D

For reserved we don't know the incoming traffic as traffic is increasing so we choose provisioned and as it related to CPU so compute optimizer.
upvoted 1 times

FlyingHawk 1 year, 4 months ago

Selected Answer: D

I vote D instead of B because of constantly increasing number of messages and to maintain service latency for its customer although provisioned concurrency incurs the charges to your AWS account. If the reduce of comput costs is more important than service latency , B is correct.

upvoted 1 times

Saliilgen 1 year, 5 months ago

Selected Answer: D

IMO answer is D.

The number of messages in message queue is constantly increasing: then, to maintain service latency you need to configure provisioned concurrency.

For cpu-intensive job you can optimize the duration increasing memory

upvoted 2 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: B

Let's think - Should we increase or decrease memory? Know that CPU resources scale proportionally with memory allocation in AWS Lambda, so increasing memory will help reduce the execution time for CPU-intensive tasks. Shorter execution times lower the overall compute cost. Then, A & C are out.

About option D, it ensures performance and prevents cold starts, yes, but it also incurs additional costs to keep Lambda execution environments pre-warmed, even when not in use. Reserved concurrency is more cost-effective for sustained workloads.

upvoted 3 times

dragosky 1 year, 5 months ago

Selected Answer: D

Reserved concurrency – This represents the maximum number of concurrent instances allocated to your function. When a function has reserved concurrency, no other function can use that concurrency. Reserved concurrency is useful for ensuring that your most critical functions always have enough concurrency to handle incoming requests. Configuring reserved concurrency for a function incurs no additional charges.

Provisioned concurrency – This is the number of pre-initialized execution environments allocated to your function. These execution environments are ready to respond immediately to incoming function requests. Provisioned concurrency is useful for reducing cold start latencies for functions. Configuring provisioned concurrency incurs additional charges to your AWS account.

upvoted 2 times

tch 1 year, 3 months ago

then should use reserved concurrency... as no additional charges....

upvoted 1 times

JA2018 1 year, 6 months ago

Selected Answer: D

Provisioned concurrency:

For scenarios with predictable high traffic like marketing events, using provisioned concurrency ensures Lambda functions are always ready to process requests, maintaining low latency.

Increase memory based on Compute Optimizer recommendations:

Since the Lambda functions are CPU intensive, increasing allocated memory based on AWS Compute Optimizer suggestions can often improve performance and reduce costs by allowing the functions to process more data without hitting CPU throttling.

upvoted 1 times

AMEJack 1 year, 6 months ago

Selected Answer: B

Provisioned concurrency will increase cost and we need to maintain latency not to decrease latency. Also, provisioned concurrency cost is very related to memory, so if you use provisioned and increase memory this will affect the cost dramatically.

upvoted 2 times

bujuman 2 years ago

Selected Answer: D

Also Lambda provisioned concurrency incur additional Account charges (<https://docs.aws.amazon.com/lambda/latest/dg/provisioned-concurrency.html>), it's the best option because it is stated:

- The company wants to reduce the compute costs and to maintain service latency for its customers.

So maintaining service latency while reducing compute cost is requested.

That being said, Lambda optimization is not a trivial task, that's why one should rely on AWS Compute Optimizer recommendations to analyze usage and find the best fit.

Please read following for more insights:

<https://aws.amazon.com/blogs/compute/optimizing-aws-lambda-cost-and-performance-using-aws-compute-optimizer/>

upvoted 3 times

JackyCCK 2 years, 2 months ago

Increase the memory according to AWS Compute Optimizer recommendations --> so we can lower the duration of lambda function to reduce the cost.

The ans must be between B & D

upvoted 2 times

alawada 2 years, 2 months ago

Selected Answer: D

Provisioned Concurrency keeps the Lambda functions initialized and ready to process incoming events, reducing the cold start latency associated with spinning up new execution environments.

upvoted 4 times

asdfcdsxdcf 2 years, 3 months ago

Selected Answer: D

D is correct

upvoted 3 times

osmk 2 years, 3 months ago

Selected Answer: A

When a large number of messages are in the SQS queue, Lambda scales out, adding additional functions to process the messages. The scale out can consume the concurrency quota in the account. To prevent this from happening, you can set reserved concurrency for individual Lambda functions. This ensures that the specified Lambda function can always scale to that much concurrency, but it also cannot exceed this number.

<https://docs.aws.amazon.com/lambda/latest/operatorguide/computing-power.html>

upvoted 2 times

osmk 2 years, 3 months ago

When a large number of messages are in the SQS queue, Lambda scales out, adding additional functions to process the messages. The scale out can consume the concurrency quota in the account. To prevent this from happening, you can set reserved concurrency for individual Lambda functions. This ensures that the specified Lambda function can always scale to that much concurrency, but it also cannot exceed this number.

<https://docs.aws.amazon.com/lambda/latest/operatorguide/computing-power.html>

upvoted 1 times

A company runs its workloads on Amazon Elastic Container Service (Amazon ECS). The container images that the ECS task definition uses need to be scanned for Common Vulnerabilities and Exposures (CVEs). New container images that are created also need to be scanned.

Which solution will meet these requirements with the FEWEST changes to the workloads?

A. Use Amazon Elastic Container Registry (Amazon ECR) as a private image repository to store the container images. Specify scan on push filters for the ECR basic scan. **Most Voted**

B. Store the container images in an Amazon S3 bucket. Use Amazon Macie to scan the images. Use an S3 Event Notification to initiate a Macie scan for every event with an s3:ObjectCreated:Put event type.

C. Deploy the workloads to Amazon Elastic Kubernetes Service (Amazon EKS). Use Amazon Elastic Container Registry (Amazon ECR) as a private image repository. Specify scan on push filters for the ECR enhanced scan.

D. Store the container images in an Amazon S3 bucket that has versioning enabled. Configure an S3 Event Notification for s3:ObjectCreated:* events to invoke an AWS Lambda function. Configure the Lambda function to initiate an Amazon Inspector scan.

Correct Answer: A

Community vote distribution

A (96%)

4%

Comments

1dd **Highly Voted** 1 year, 9 months ago

Selected Answer: A

need less workload changes and CVEs

<https://docs.aws.amazon.com/AmazonECR/latest/userguide/image-scanning.html>

upvoted 9 times

xBUGx **Highly Voted** 1 year, 9 months ago

Selected Answer: A

FEWEST changes to the workloads and scan CVE is enough. A looks OK.

upvoted 6 times

zdi561 **Most Recent** 1 year, 3 months ago

Selected Answer: D

A is not right because ECR basic scan only scan OS. in B enhanced scan does scan app packages but it is required to run EKS. D is simpler.

upvoted 1 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: A

Amazon ECR: Designed for storing container images, integrates seamlessly with Amazon ECS.

Scan on Push: Automatically scans images for vulnerabilities (CVEs) when they are uploaded. This is built-in functionality and requires minimal changes to existing workflows.

Fewest Changes: Existing ECS workloads can directly pull images from ECR without modification to the workloads themselves.

B - Amazon Macie focuses on data discovery and classification (such as identifying sensitive data in S3 buckets), not container image vulnerabilities.

C - Result-wise OK, but switching from ECS to EKS is unnecessary and involves a steep learning curve and operational overhead.

D - Result-wise OK. Amazon Inspector can scan for vulnerabilities, but this is too complex and requires a custom setup, including using S3 and Lambda to trigger scans.

upvoted 2 times

sandordini 1 year, 7 months ago

Selected Answer: A

Basic scan looks for Common Vulnerabilities and Exposures (CVEs)
upvoted 5 times

A company uses an AWS Batch job to run its end-of-day sales process. The company needs a serverless solution that will invoke a third-party reporting application when the AWS Batch job is successful. The reporting application has an HTTP API interface that uses username and password authentication.

Which solution will meet these requirements?

A. Configure an Amazon EventBridge rule to match incoming AWS Batch job SUCCEEDED events. Configure the third-party API as an EventBridge API destination with a username and password. Set the API destination as the EventBridge rule target.

Most Voted

B. Configure Amazon EventBridge Scheduler to match incoming AWS Batch job SUCCEEDED events. Configure an AWS Lambda function to invoke the third-party API by using a username and password. Set the Lambda function as the EventBridge rule target.

C. Configure an AWS Batch job to publish job SUCCEEDED events to an Amazon API Gateway REST API. Configure an HTTP proxy integration on the API Gateway REST API to invoke the third-party API by using a username and password.

D. Configure an AWS Batch job to publish job SUCCEEDED events to an Amazon API Gateway REST API. Configure a proxy integration on the API Gateway REST API to an AWS Lambda function. Configure the Lambda function to invoke the third-party API by using a username and password.

Correct Answer: A

Community vote distribution

A (65%)

B (26%)

9%

Comments

venutadi **Highly Voted** 2 years, 1 month ago

Selected Answer: A

<https://aws.amazon.com/blogs/compute/using-api-destinations-with-amazon-eventbridge/>
Amazon EventBridge enables developers to route events between AWS services, integrated software as a service (SaaS) applications, and your own applications. It can help decouple applications and produce more extensible, maintainable architectures. With the new API destinations feature, EventBridge can now integrate with services outside of AWS using REST API calls.

upvoted 13 times

shintaro0914 2 years, 1 month ago

I agree.

upvoted 1 times

sandordini **Highly Voted** 2 years, 1 month ago

I'm confused. Both A and B seem to be viable. There is no requirement of cost, complexity, or overhead. :S

upvoted 5 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: A

A - EventBridge API Destination is specifically designed to send events to HTTP APIs. It supports basic authentication (username and password), which fits the requirement for 3rd-party API authentication.

B - OK, but EventBridge Scheduler is unnecessary because EventBridge rules can already match Batch job events. Introducing Lambda functions adds complexity and costs.

C - An HTTP proxy integration on API Gateway cannot handle authentication directly. We still need to manage username and password authentication separately.

D - OK, but even more complex than option B.

upvoted 3 times

Scheldon 1 year, 12 months ago

Selected Answer: A

AnswerA

EventBridge Schedule will not work as it will allow us to "do something" per schedule.

EventBridge rule will allow us to "do something" when event will occur.

I think there is no possibility to publish/send job "SUCCEEDED" to AMAZON API Gateway REST API or that we can do anykind of integration with AMAZON API Gateway, hence I would choose A

upvoted 1 times

bujuman 2 years ago

Selected Answer: A

Even though option A and B could do the trick and also no statement related to least effort is requested, EventBridge is dedicated for similar use case

<https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-api-destinations.html>

Plus, it can also handle basic authentication

<https://aws.amazon.com/blogs/compute/using-api-destinations-with-amazon-eventbridge/>

upvoted 1 times

Sergiuss95 2 years, 1 month ago

Selected Answer: B

I think is better to programming a lambda and obtain user and password from Secret Manager... So I think the better solution is B

upvoted 3 times

EdricHoang 1 year, 11 months ago

EventBridge also store credentials in Secret manager: <https://aws.amazon.com/blogs/compute/using-api-destinations-with-amazon-eventbridge/>

upvoted 1 times

Oluwatosin09 2 years, 1 month ago

Selected Answer: B

Answer should be B.

upvoted 1 times

boluwatito 2 years, 1 month ago

Selected Answer: B

Configure Amazon EventBridge Scheduler to match incoming AWS Batch job SUCCEEDED events.

Configure an AWS Lambda function to invoke the third-party API using a username and password.

Set the Lambda function as the EventBridge rule target.

upvoted 2 times

AlvinC2024 2 years, 2 months ago

Selected Answer: A

A. Configure an Amazon EventBridge rule to match incoming AWS Batch job SUCCEEDED events. Configure the third-party API as an EventBridge AP destination with a username and password. Set the API destination as the EventBridge rule target.

This option is the most direct and serverless approach to meeting the requirements. Amazon EventBridge can detect the successful completion of the AWS Batch job and trigger actions based on this event. By configuring the third-party API as an API destination with authentication credentials, EventBridge can directly invoke the third-party reporting application without the need for additional services. This approach minimizes complexity and operational overhead.

upvoted 4 times

alawada 2 years, 2 months ago

Selected Answer: D

Create an AWS Lambda function responsible for invoking the third-party reporting application's HTTP API endpoint. The Lambda function will be triggered by the successful completion of the AWS Batch job.

upvoted 3 times

k_k_kkk 2 years, 2 months ago

Selected Answer: B

AWS Batch sends job status change to EventBridge.

https://docs.aws.amazon.com/batch/latest/userguide/batch_cwe_events.html

upvoted 3 times

osmk 2 years, 3 months ago

look like B

upvoted 1 times


A company collects and processes data from a vendor. The vendor stores its data in an Amazon RDS for MySQL database in the vendor's own AWS account. The company's VPC does not have an internet gateway, an AWS Direct Connect connection, or an AWS Site-to-Site VPN connection. The company needs to access the data that is in the vendor database.

Which solution will meet this requirement?

- A. Instruct the vendor to sign up for the AWS Hosted Connection Direct Connect Program. Use VPC peering to connect the company's VPC and the vendor's VPC.
- B. Configure a client VPN connection between the company's VPC and the vendor's VPC. Use VPC peering to connect the company's VPC and the vendor's VPC.
- C. Instruct the vendor to create a Network Load Balancer (NLB). Place the NLB in front of the Amazon RDS for MySQL database. Use AWS PrivateLink to integrate the company's VPC and the vendor's VPC. **Most Voted**
- D. Use AWS Transit Gateway to integrate the company's VPC and the vendor's VPC. Use VPC peering to connect the company's VPC and the vendor's VPC.

Correct Answer: C

Community vote distribution



Comments

Ucy **Highly Voted** 1 year, 10 months ago

Pour yourself a cold beer, when you get to this question, its been a very long run
upvoted 22 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: C

PrivateLink: AWS PrivateLink enables secure and private access to resources (like RDS) across accounts and VPCs using interface endpoints. It avoids exposing the database to the internet.

NLB: A Network Load Balancer acts as the integration point for PrivateLink. The vendor's NLB will route traffic to the RDS database.

VPC Without Internet Gateway: PrivateLink works even if the company's VPC does not have an internet gateway because communication occurs over private AWS infrastructure.

Option A, B, and D all mention VPC Peering, but it doesn't enable direct access to RDS without additional routing or security configurations.

upvoted 4 times

Scheldon 1 year, 12 months ago

Selected Answer: C

AnswerC

AWS PrivateLink enables you to connect to some AWS services, services hosted by other AWS accounts (referred to as endpoint services), and supported AWS Marketplace partner services, via private IP addresses in your VPC. The interface endpoints are created directly inside of your VPC, using elastic network interfaces and IP addresses in your VPC's subnets. That means that VPC Security Groups can be used to manage access to the endpoints.

<https://docs.aws.amazon.com/whitepapers/latest/aws-vpc-connectivity-options/aws-privatelink.html>

upvoted 2 times

Nm55569 2 years ago

Selected Answer: C

<https://aws.amazon.com/blogs/database/access-amazon-rds-across-vpcs-using-aws-privatelink-and-network-load-balancer/>

upvoted 2 times

TwinSpark 2 years ago

I think i go for C, because if you exclude Direct connect, VPN and GW so only C is available. but create an NLB so do not want provision a transit GW sounds weird to me

upvoted 1 times

sandordini 2 years, 1 month ago

Selected Answer: C

Private link:

Does not require VPC linking: NO Internet Gateway, NO NAT Gateway, No Route table

Needs NLB on Service VPC, and ENI on the Customer VPC

upvoted 4 times

rondelldell 2 years, 2 months ago

D

You can peer both intra-Region and inter-Region transit gateways, and route traffic between them, which includes IPv4 and IPv6 traffic. To do this, create a peering attachment on your transit gateway, and specify a transit gateway. The peer transit gateway can be in your account or a different AWS account.

After you create a peering attachment request, the owner of the peer transit gateway (also referred to as the acceptor transit gateway) must accept the request. To route traffic between the transit gateways, add a static route to the transit gateway route table that points to the transit gateway peering attachment.

<https://docs.aws.amazon.com/vpc/latest/tgw/tgw-peering.html>

upvoted 2 times

xBUGx 2 years, 2 months ago

D does not involve internet. But TGW is unnecessary.

A is more simple and clear.

upvoted 1 times

Sivaeas 2 years, 3 months ago

Selected Answer: C

AWS PrivateLink:

AWS PrivateLink enables you to privately access services hosted on AWS in a highly available and scalable manner. With PrivateLink, you can access the vendor's RDS for MySQL instance securely without exposing it to the public internet.

The vendor can create a VPC endpoint for RDS within their own VPC, which acts as an entry point for accessing the RDS instance. This endpoint can then be shared with the company.

The company can create a VPC endpoint service in their VPC and accept the endpoint connection request from the vendor. This allows the company's resources to communicate with the RDS instance securely through PrivateLink.

upvoted 3 times

lenotc 2 years, 3 months ago

Selected Answer: C

C is correct:

<https://aws.amazon.com/blogs/networking-and-content-delivery/how-to-securely-publish-internet-applications-at-scale-using-application-load-balancer-and-aws-privatelink/>

upvoted 1 times

1dd 2 years, 3 months ago

Selected Answer: C

Plz commit the previous comment,

A involve- Direct connect

B involve - peering required same region

D involve - uses internet gateway

upvoted 3 times

1dd 2 years, 3 months ago

Selected Answer: A

No internet gateway XD

No Direct connect XC

No Peering XB

upvoted 1 times

asdfcdsxdcf 2 years, 3 months ago

Shouldn't it be D?

upvoted 3 times

1dd 2 years, 3 months ago

I think it required use of internet gateway .

upvoted 1 times

Jacky_S 1 year, 11 months ago

No, it did not
<https://docs.aws.amazon.com/vpc/latest/tgw/tgw-transit-gateways.html>
upvoted 1 times

rondelldell 2 years, 2 months ago

YES D

transit gateway is like router - u can connect VPCs AND OnPrem. VPCs can be in another account or region or org
upvoted 1 times

Sergantus 1 year, 7 months ago

VPC peering merges two VPCs and exposes all the services across both VPCs, which is more than less desirable
upvoted 1 times

A company wants to set up Amazon Managed Grafana as its visualization tool. The company wants to visualize data from its Amazon RDS database as one data source. The company needs a secure solution that will not expose the data over the internet.

Which solution will meet these requirements?

- A. Create an Amazon Managed Grafana workspace without a VPC. Create a public endpoint for the RDS database. Configure the public endpoint as a data source in Amazon Managed Grafana.
- B. Create an Amazon Managed Grafana workspace in a VPC. Create a private endpoint for the RDS database. Configure the private endpoint as a data source in Amazon Managed Grafana. **Most Voted**
- C. Create an Amazon Managed Grafana workspace without a VPC. Create an AWS PrivateLink endpoint to establish a connection between Amazon Managed Grafana and Amazon RDS. Set up Amazon RDS as a data source in Amazon Managed Grafana.
- D. Create an Amazon Managed Grafana workspace in a VPC. Create a public endpoint for the RDS database. Configure the public endpoint as a data source in Amazon Managed Grafana.

Correct Answer: B

Community vote distribution

B (68%)

C (33%)

Comments

LeonSauveterre **Highly Voted** 1 year, 5 months ago

Selected Answer: B

A & D are out, because they're using public endpoints, which causes exposure over the Internet.

Why C is wrong: If the Grafana workspace is not in a VPC, it cannot use PrivateLink endpoints. PrivateLink endpoints must reside within a VPC. If Grafana is outside the VPC, PrivateLink cannot establish the required secure, private connection.

upvoted 6 times

Sergiuss95 **Highly Voted** 2 years, 1 month ago

Selected Answer: B

I think is b. Private endpoint sounds like private vpc endpoint, that is equals to privatelink

upvoted 6 times

tch **Most Recent** 1 year, 3 months ago

Selected Answer: C

should be C:

<https://docs.aws.amazon.com/grafana/latest/userguide/AMG-configure-nac.html>

Amazon Managed Grafana workspace not created in VPC.....

upvoted 1 times

GOTJ 1 year, 3 months ago

Selected Answer: B

With all due respect, once again options are poorly written: you don't create an Amazon Managed Grafana workspace IN or WITHOUT a VPC. Instead, you create an AMG WITH or WITHOUT a VPC CONNECTION:

<https://docs.aws.amazon.com/grafana/latest/userguide/AMG-create-workspace.html>

(Step 8, Optional): You can choose to connect to an Amazon virtual private cloud (VPC) on this page, or you can connect to a VPC later. To learn more, see Connect to data sources or notification channels in Amazon VPC from Amazon Managed Grafana.

According to the "to learn more" link, if you create the AMG WITHOUT a VPC CONNECTION:

By default, traffic from your Amazon Managed Grafana workspace to data sources or notification channels flows via the public Internet. This limits the connectivity from your Amazon Managed Grafana workspace to services that are publicly accessible.

So option "C" must be discarded as it doesn't comply company requirements of avoid using internet

upvoted 1 times

ARV14 1 year, 6 months ago

Selected Answer: C

<https://docs.aws.amazon.com/grafana/latest/userguide/VPC-endpoints.html> , <https://aws.amazon.com/about-aws/whats-new/2022/11/amazon-managed-grafana-connection-data-sources-hosted-virtual-private-cloud/>

upvoted 1 times

FlyingHawk 1 year, 4 months ago

By default, traffic from your Amazon Managed Grafana workspace to data sources or notification channels flows via the public Internet. This limits the connectivity from your Amazon Managed Grafana workspace to services that are publicly accessible.

If you want to connect to private-facing data sources that are within a VPC, or keep traffic local to a VPC, you can connect your Amazon Managed Grafana workspace to the Amazon Virtual Private Cloud (Amazon VPC) hosting these data sources. After you configure the VPC data source connection, all traffic flows via your VPC.

upvoted 1 times

chest_jd 1 year, 7 months ago

Selected Answer: B

Choice B or C could be resolved in this way:

B. Create an Amazon Managed Grafana workspace in a VPC

C. Create an Amazon Managed Grafana workspace without a VPC

As far as I know we cannot create workspace without VCP

upvoted 3 times

tonybuivannghia 1 year, 7 months ago

Selected Answer: C

After searching effort, I agree C is correct because AMG workspace can't include in VPC. When you have not configured a private VPC, and Amazon Managed Grafana is connecting to publicly accessible data sources, it connects to some AWS services in the same region via AWS PrivateLink. This includes services such as CloudWatch, Amazon Managed Service for Prometheus and AWS X-Ray. Traffic to those services does not flow via the public Internet.

upvoted 1 times

NSA_Poker 1 year, 10 months ago

Selected Answer: C

(B or C)?-1 = Do we create AMG workspace in a VPC OR do we create AMG workspace without a VPC? AMG is NOT created within a VPC; AMG connects to a VPC. "Currently, you can connect one Amazon Managed Grafana workspace to one VPC endpoint in the same region and same account. However, you can use Virtual Private Cloud peering or AWS Transit Gateway to connect the cross-region or cross-account VPCs, then connect the select the VPC endpoint that's in the same account and same region as your Amazon Managed Grafana workspace." -FAQs

(C) is correct.

upvoted 3 times

NSA_Poker 1 year, 10 months ago

(B or C)?-2 = private endpoint OR AWS PrivateLink? The brand-name is more correct.

(B or C)?-3 = Configure the private endpoint as a data source OR Set up Amazon RDS as a data source? In the AMG console, after clicking on Data sources, you'll see a list of AWS services (Athena, Redshift etc) NOT network endpoints. After selecting RDS, you can further specify the Region & Resource ID.

(B) eliminated.

(C) is correct.

upvoted 1 times

LeonSauveterre 1 year, 5 months ago

You're missing out PrivateLink in option C.

upvoted 1 times

EdricHoang 1 year, 11 months ago

Selected Answer: B

Its B.

C is also a valid choice

"Not exposing to the internet" is letting me eliminate C

upvoted 3 times

NSA_Poker 1 year, 11 months ago

(B) "a private endpoint" & (C) "an AWS PrivateLink endpoint" do NOT expose traffic to the internet.

(A & D) eliminated. "a public endpoint for the RDS database" would "expose the data over the internet"

upvoted 1 times

ike001 1 year, 11 months ago

B as you need to create Managed Grafana workspace with a VPC for private access <https://docs.aws.amazon.com/grafana/latest/userguide/AMG-configure-nac.html>

upvoted 2 times

NSA_Poker 1 year, 10 months ago

(B) doesn't say 'with a VPC'; it says "..Grafana workspace IN A VPC."

(B) eliminated.

(C) is correct.

upvoted 2 times

Nm55569 2 years ago

Selected Answer: B

<https://aws.amazon.com/about-aws/whats-new/2022/11/amazon-managed-grafana-connection-data-sources-hosted-virtual-private-cloud/>

upvoted 3 times

sandordini 2 years, 1 month ago

I guess they mean C, But again, it's strange... IMO B would also work... There is no requirement for the least effort... Pls, correct me if I'm wrong...

upvoted 2 times

venutadi 2 years, 1 month ago

Selected Answer: C

Once you configure direct connectivity between a Grafana workspace and a VPC, Amazon Managed Grafana creates and manages an elastic network interface (ENI) per subnet to connect to the VPC. This enables the Grafana workspace to connect to data sources within the VPC, such as OpenSearch domains or RDS databases. Additionally, all traffic is now routed through the configured VPC, including alert destination and data source connectivity.

upvoted 4 times

VortexMD 2 years, 2 months ago

AWS PrivateLink provides private connectivity between virtual private clouds (VPCs), supported AWS services, and your on-premises networks without exposing your traffic to the public internet. Interface VPC endpoints, powered by PrivateLink, connect you to services hosted by AWS Partners and supported solutions available in AWS Marketplace.

upvoted 1 times

VortexMD 2 years, 2 months ago

<https://aws.amazon.com/blogs/mt/announcing-private-vpc-data-source-support-for-amazon-managed-grafana/>

upvoted 1 times

Bazzix 2 years, 2 months ago

Selected Answer: B

B is correct

upvoted 5 times

osmk 2 years, 3 months ago

Selected Answer: C

cccc ccc

upvoted 3 times