

A company is implementing a new application on AWS. The company will run the application on multiple Amazon EC2 instances across multiple Availability Zones within multiple AWS Regions. The application will be available through the internet. Users will access the application from around the world.

The company wants to ensure that each user who accesses the application is sent to the EC2 instances that are closest to the user's location.

Which solution will meet these requirements?

- A. Implement an Amazon Route 53 geolocation routing policy. Use an internet-facing Application Load Balancer to distribute the traffic across all Availability Zones within the same Region.
- B. Implement an Amazon Route 53 geoproximity routing policy. Use an internet-facing Network Load Balancer to distribute the traffic across all Availability Zones within the same Region. **Most Voted**
- C. Implement an Amazon Route 53 multivalue answer routing policy. Use an internet-facing Application Load Balancer to distribute the traffic across all Availability Zones within the same Region.
- D. Implement an Amazon Route 53 weighted routing policy. Use an internet-facing Network Load Balancer to distribute the traffic across all Availability Zones within the same Region.

Correct Answer: B

Community vote distribution

B (62%)

A (38%)

Comments

dhewa **Highly Voted** 1 year, 9 months ago

Selected Answer: B

Keyword closest.

Amazon Route 53 Geoproximity Routing: This routing policy directs traffic based on the geographic location of your users and your resources, ensuring that users are routed to the closest EC2 instances.

upvoted 9 times

nebajp **Highly Voted** 1 year, 10 months ago

Selected Answer: A

Correct answer is A

Geo location is based on users location

GeoProximity is based on the AWS services used by users.

upvoted 8 times

Sergantus 1 year, 6 months ago

The main purpose of geolocation policy is to enforce regional restrictions, provide language-specific content, or balance load across regional endpoints. It doesn't return "closest" records but only relevant (location) records.

upvoted 4 times

In2718 **Most Recent** 10 months ago

Selected Answer: B

AI follow-up --- Geolocation routing directs traffic based on the user's location, while geoproximity routing considers both the user's location and the location of your resources

upvoted 2 times

In2718 10 months ago

Selected Answer: B

Another poorly written question... but I'll go with B for all the reasons stated.

upvoted 1 times

MortisG 1 year ago

Selected Answer: A

Geolocation routing ensures users are routed to their nearest region and is quite simple to set up while geoproximity routing is more suited for custom routing control, requires traffic flow and is more complex to set up.

upvoted 1 times

bora4motion 1 year ago

Selected Answer: A

I am going with A. Its web, which translates to HTTP and therefore ALB.

upvoted 1 times

Salilgen 1 year, 5 months ago

Selected Answer: B

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-geoproximity.html>

upvoted 2 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: B

A - "geolocation" is simple and effective, but it might not always direct users to the physically closest Region. It's more about their geographic "label" (country/region).

B - "geoproximity" is better for ensuring actual closeness but adds a bit more complexity.

Then, I don't think ALB or NLB should be an argument here. NLBs are generally preferred for applications that require high throughput and low latency, especially if they are not using HTTP/HTTPS. NLBs operate at Layer 4 (TCP/UDP), so suitable for a wider range of applications than ALBs (Layer 7, HTTP/HTTPS). Even if the application does use HTTP/HTTPS, NLBs can still be used if you don't need the advanced features of an ALB (like content-based routing or host-based routing).

Then again, the question haven't explicitly stated the protocol so oh well, I'm gonna go with B.

upvoted 4 times

robotgeek 1 year, 4 months ago

Agree, if the question does not state "web", "ALB", or something "layer 7 specific" we should not assume it

upvoted 1 times

deacon967 1 year, 5 months ago

Selected Answer: A

You can argue geolocation vs geoproximity all day but why would one use a NLB for an internet based app? The is the perfect use case for an ALB. A is the correct answer.

upvoted 1 times

LeonSauveterre 1 year, 5 months ago

It's not necessary that an internet based app is bound to use ALBs for traffic balancing. If you immediately go to ALBs when a question mentions web-based apps, you're gonna get so many questions wrong.

upvoted 2 times

dragosky 1 year, 5 months ago

Selected Answer: B

Which AWS Route 53 routing policy should be used to route traffic based on users' location?

Geoproximity routing

Geoproximity routing lets Amazon Route 53 route traffic to your resources based on the geographic location of your users and your resources.

upvoted 1 times

8621a7c 1 year, 7 months ago

Selected Answer: A

Go with A

upvoted 2 times

b3b5fdd 1 year, 7 months ago

Selected Answer: B

Correct answer is B!

upvoted 4 times

Bwhizzy 1 year, 8 months ago

Selected Answer: A

The answer is A. routing traffic to the user region. Also, it can't be B. It says using a Network Load Balancer. you can't use a network load balancer for internet traffic

upvoted 4 times

mk168898 1 year, 7 months ago

TCP UDP traffic?

upvoted 2 times

kbgsqsgs 1 year, 8 months ago

Selected Answer: A

Correct answer is A

upvoted 2 times

blehbleh 1 year, 8 months ago

Selected Answer: B

It is B:

Geolocation:

Primarily focuses on the user's location, allowing you to serve different content based on their geographic region.

Geoproximity:

Considers both the user's location and the location of your resources, dynamically choosing the closest option based on network latency.

upvoted 3 times

komorebi 1 year, 10 months ago

Selected Answer: B

Answer is B

upvoted 5 times

A financial services company plans to launch a new application on AWS to handle sensitive financial transactions. The company will deploy the application on Amazon EC2 instances. The company will use Amazon RDS for MySQL as the database. The company's security policies mandate that data must be encrypted at rest and in transit.

Which solution will meet these requirements with the LEAST operational overhead?

A. Configure encryption at rest for Amazon RDS for MySQL by using AWS KMS managed keys. Configure AWS Certificate Manager (ACM) SSL/TLS certificates for encryption in transit. **Most Voted**

B. Configure encryption at rest for Amazon RDS for MySQL by using AWS KMS managed keys. Configure IPsec tunnels for encryption in transit.

C. Implement third-party application-level data encryption before storing data in Amazon RDS for MySQL. Configure AWS Certificate Manager (ACM) SSL/TLS certificates for encryption in transit.

D. Configure encryption at rest for Amazon RDS for MySQL by using AWS KMS managed keys. Configure a VPN connection to enable private connectivity to encrypt data in transit.

Correct Answer: A

Community vote distribution

A (100%)

Comments

JoeTromundo **Highly Voted** 1 year, 8 months ago

Selected Answer: A

Amazon RDS for MySQL supports encryption at rest using AWS Key Management Service (KMS) managed keys. This encryption is easy to enable during the creation of the RDS instance and requires minimal configuration. AWS KMS provides a fully managed solution for managing encryption keys, and using KMS managed keys reduces operational overhead related to key management and rotation. Encryption in transit ensures that data transmitted between the application and the RDS database is secure. AWS Certificate Manager (ACM) can be used to provide SSL/TLS certificates, which are required to encrypt data in transit. ACM simplifies the management of SSL/TLS certificates by handling certificate renewal and deployment, reducing operational overhead.

upvoted 5 times

dhewa **Most Recent** 1 year, 9 months ago

Selected Answer: A

A is my choice anyday.

upvoted 4 times

A company is migrating its on-premises Oracle database to an Amazon RDS for Oracle database. The company needs to retain data for 90 days to meet regulatory requirements. The company must also be able to restore the database to a specific point in time for up to 14 days.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create Amazon RDS automated backups. Set the retention period to 90 days.
- B. Create an Amazon RDS manual snapshot every day. Delete manual snapshots that are older than 90 days.
- C. Use the Amazon Aurora Clone feature for Oracle to create a point-in-time restore. Delete clones that are older than 90 days.
- D. Create a backup plan that has a retention period of 90 days by using AWS Backup for Amazon RDS. **Most Voted**

Correct Answer: D

Community vote distribution

D (94%)

6%

Comments

nebajp **Highly Voted** 1 year, 10 months ago

Selected Answer: D

Correct Answer is D - Its fulfilling the requirement of point in time.
Automated Backup - Default retention period is 0-35 Days - so option A is wrong.
upvoted 8 times

JoeTromundo **Highly Voted** 1 year, 8 months ago

Selected Answer: D

A: Amazon RDS automated backups support a maximum retention period of 35 days. This option does not meet the requirement to retain backups for 90 days.
B: This approach requires manual snapshot management, including scheduling snapshots and deleting old ones. This increases operational overhead and is prone to human error.
C: This option is not applicable as Aurora Clone is a feature specific to Amazon Aurora and not available for Amazon RDS for Oracle. Additionally, it would require manual management of clones, increasing complexity.
D: AWS Backup supports point-in-time recovery for Amazon RDS, enabling you to restore the database to any specific point within the defined retention period, up to 35 days. For the requirement of 14 days, AWS Backup easily supports this capability.
upvoted 5 times

KarlosRC **Most Recent** 3 months ago

Selected Answer: A

D - is not the best option: AWS Backup stores snapshots but does not manage transaction logs for PITR.
upvoted 1 times

spoved 1 year, 8 months ago

Selected Answer: D

<https://aws.amazon.com/getting-started/hands-on/amazon-rds-backup-restore-using-aws-backup/>
upvoted 3 times

pujithacg8 1 year, 10 months ago

Answer is D
upvoted 2 times

A company is developing a new application that uses a relational database to store user data and application configurations. The company expects the application to have steady user growth. The company expects the database usage to be variable and read-heavy, with occasional writes.

The company wants to cost-optimize the database solution. The company wants to use an AWS managed database solution that will provide the necessary performance.

Which solution will meet these requirements MOST cost-effectively?

A. Deploy the database on Amazon RDS. Use Provisioned IOPS SSD storage to ensure consistent performance for read and write operations.

B. Deploy the database on Amazon Aurora Serverless to automatically scale the database capacity based on actual usage to accommodate the workload. **Most Voted**

C. Deploy the database on Amazon DynamoDB. Use on-demand capacity mode to automatically scale throughput to accommodate the workload.

D. Deploy the database on Amazon RDS. Use magnetic storage and use read replicas to accommodate the workload.

Correct Answer: B

Community vote distribution

B (100%)

Comments

LeonSauveterre 1 year, 5 months ago

Selected Answer: B

- A - You'd pay for high performance even when you don't need it
- B - You only pay for the database capacity you use, making it cost-efficient for unpredictable traffic.
- C - DynamoDB is not even a relational database.
- D - Magnetic disks have significantly slower IOPS compared to SSDs.

P.S. In the AWS context, magnetic storage refers to standard storage volumes (previously called "Magnetic" volumes) in EBS. It offers lower performance compared to modern SSD storage but is cheaper per GB.

upvoted 2 times

AMEJack 1 year, 6 months ago

Selected Answer: B

Option D is out: Aurora data is stored in the cluster volume, which is a single, virtual volume that uses solid state drives (SSDs). So option B is correct

upvoted 2 times

dhewa 1 year, 9 months ago

Selected Answer: B

B is my choice.

upvoted 3 times

A company hosts its application on several Amazon EC2 instances inside a VPC. The company creates a dedicated Amazon S3 bucket for each customer to store their relevant information in Amazon S3.

The company wants to ensure that the application running on EC2 instances can securely access only the S3 buckets that belong to the company's AWS account.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create a gateway endpoint for Amazon S3 that is attached to the VPC. Update the IAM instance profile policy to provide access to only the specific buckets that the application needs.

B. Create a NAT gateway in a public subnet with a security group that allows access to only Amazon S3. Update the route tables to use the NAT Gateway.

C. Create a gateway endpoint for Amazon S3 that is attached to the VPC. Update the IAM instance profile policy with a Deny action and the following condition key:

```
{
  "StringNotEquals" : {
    "s3:ResourceAccount" : [ "CompanyAWSAcctNumber"]
  }
}
```


Most Voted

D. Create a NAT Gateway in a public subnet. Update route tables to use the NAT Gateway. Assign bucket policies for all buckets with a Deny action and the following condition key:

```
{
  "StringNotEquals" : {
    "s3:ResourceAccount" : [ "CompanyAWSAcctNumber"]
  }
}
```

Correct Answer: C

Community vote distribution



Comments

toyaji **Highly Voted** 1 year, 9 months ago

Selected Answer: C

B, C is not secure way because NAT gateway is for internet-facing outbound.

A is not correct because company will create dedicated bucket for each customers it means number of buckets will increase dynamically. so you cant list all on profile.

upvoted 5 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: C

A - It doesn't explicitly ensure that access is limited to buckets only owned by the company. We need further adjustments to enforce this restriction (like adding a Condition key).

B & D - NAT gateways route traffic over the internet, so it's not so secure or cost-effective as S3 Gateway Endpoint.

C - Use an S3 Gateway Endpoint and restrict access using an IAM policy with a condition to deny access to any S3 resource outside the company's account. Simple, secure, and easy to understand.

upvoted 4 times

trinh_le 1 year, 5 months ago

Selected Answer: A

A. Allows option fits to minimum permission policy
C. Deny action is complexity and hard to debug

upvoted 2 times

Denise123 1 year, 5 months ago

Selected Answer: C

The requirement in the question is to ensure that the EC2 instances can securely access only the S3 buckets that belong to (are owned by) the company's AWS account.

In that context, using the "StringNotEquals": {"S3ResourceAccount":["CompanyAWSAcctNumber"]} condition key in the IAM instance profile policy is a valid approach.

This condition key restricts access to S3 buckets that are not owned by the specified AWS account number (the company's account). By setting a Deny with this condition, it effectively allows access only to the S3 buckets owned by the company's AWS account.

upvoted 3 times

Denise123 1 year, 5 months ago

Selected Answer: A

By choosing Option A, the company can leverage the simplicity and security of VPC Gateway Endpoints for Amazon S3, combined with IAM instance profile policies to grant selective access to the required S3 buckets. This solution provides the necessary isolation and access control with minimal operational overhead, making it the most efficient and scalable approach.

Option C and other options have limitations or introduce additional operational overhead Option C (Creating a Gateway Endpoint and using a Deny policy with S3ResourceAccount condition) is not a valid approach. The S3ResourceAccount condition key is used to restrict access based on the AWS account that owns the S3 bucket, not the AWS account that is accessing the bucket.

upvoted 1 times

Denise123 1 year, 5 months ago

Ignore. After re-reading the question carefully, Option C is actually a valid solution.

The right answer is C and I wrote exactly that in my answer myself lol. In the question it says "BELONGS TO" (which I failed to read correctly)

The requirement in the question is to ensure that the EC2 instances can securely access only the S3 buckets that belong to (are owned by) the company's AWS account.

In that context, using the "StringNotEquals": {"S3ResourceAccount":["CompanyAWSAcctNumber"]} condition key in the IAM instance profile policy is a valid approach.

This condition key restricts access to S3 buckets that are not owned by the specified AWS account number (the company's account). By setting a Deny with this condition, it effectively allows access only to the S3 buckets owned by the company's AWS account.

upvoted 2 times

GOTJ 1 year, 4 months ago

Errare humanum est :)

upvoted 1 times

EllenLiu 1 year, 5 months ago

Selected Answer: A

answer C cannot restrict ec2 from accessing s3 from other accounts.

Only by providing access to specific buckets in IAM policy on EC2 are we able to achieve this.

upvoted 1 times

AMEJack 1 year, 6 months ago

Selected Answer: A

Option C will deny the other S3 buckets but will not Allow access to the specified Buckets, thus we should have Allow rules.

Option A will allow the specified Buckets and implicit Deny the other buckets.

upvoted 2 times

ARV14 1 year, 6 months ago

You will keep updating the policy with new s3 buckets as user base grows? Operational overhead?

upvoted 1 times

youkarthik 1 year, 6 months ago

Selected Answer: A

A as per gen Als

upvoted 1 times

Bwhizzy 1 year, 8 months ago

Selected Answer: C

Answer is C. Just specify only the company AWS account number, rather than listing all the Buckets

upvoted 3 times

XXXXXINN 1 year, 8 months ago

Vote A

upvoted 2 times

siheom 1 year, 8 months ago

VOTE A

upvoted 1 times

viejito 1 year, 8 months ago

la respuesta correcta es la A :Los buckets son servicios globales.(o sea no están en una VPC ni Subnet), entonces no hace falta que estén en una subred publica o privada ; los Nat Gateway son para redes publicas o privadas .Entonces ahí descarta B,C y D .Cuando quieres conectar un recurso de Global Service se usa Endpoint Gateway por eso la respuesta es A .

upvoted 1 times

A company is building a cloud-based application on AWS that will handle sensitive customer data. The application uses Amazon RDS for the database, Amazon S3 for object storage, and S3 Event Notifications that invoke AWS Lambda for serverless processing.

The company uses AWS IAM Identity Center to manage user credentials. The development, testing, and operations teams need secure access to Amazon RDS and Amazon S3 while ensuring the confidentiality of sensitive customer data. The solution must comply with the principle of least privilege.

Which solution meets these requirements with the LEAST operational overhead?

- A. Use IAM roles with least privilege to grant all the teams access. Assign IAM roles to each team with customized IAM policies defining specific permission for Amazon RDS and S3 object access based on team responsibilities.
- B. Enable IAM Identity Center with an Identity Center directory. Create and configure permission sets with granular access to Amazon RDS and Amazon S3. Assign all the teams to groups that have specific access with the permission sets. **Most Voted**
- C. Create individual IAM users for each member in all the teams with role-based permissions. Assign the IAM roles with predefined policies for RDS and S3 access to each user based on user needs. Implement IAM Access Analyzer for periodic credential evaluation.
- D. Use AWS Organizations to create separate accounts for each team. Implement cross-account IAM roles with least privilege. Grant specific permission for RDS and S3 access based on team roles and responsibilities.

Correct Answer: B

Community vote distribution

B (100%)

Comments

kevindu **Highly Voted** 1 year, 10 months ago

Selected Answer: B

Is there anyone who has recently passed the exam who can tell me approximately how many of the original questions are in the actual exam?
upvoted 13 times

dhewa **Highly Voted** 1 year, 9 months ago

Selected Answer: B

IAM Identity Center: This service simplifies user management by centralizing credentials and access control.
Permission Sets: You can create granular permission sets that align with the principle of least privilege, ensuring that each team has only the access they need.
Group Assignments: By assigning teams to groups with specific permission sets, you streamline access management and reduce the complexity of individual user permissions.
This approach minimizes operational overhead while maintaining secure and compliant access to sensitive customer data
upvoted 5 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: B

A - Managing and customizing IAM roles for multiple teams can become more and more complex and high-maintenance, especially if team responsibilities evolve.
B - IAM Identity Center (formerly AWS SSO) allows you to manage/update access for multiple teams from one place.
C - TOO MUCH WORK.
D - Cross-account access is harder to manage and introduces additional operational overhead.
upvoted 2 times

JoeTromundo 1 year, 8 months ago

Selected Answer: B

Option A is goo but not the best, which is option B.

upvoted 2 times

aragon_saa 1 year, 9 months ago

Selected Answer: B

Answer is B

upvoted 3 times

A company has an Amazon S3 bucket that contains sensitive data files. The company has an application that runs on virtual machines in an on-premises data center. The company currently uses AWS IAM Identity Center.

The application requires temporary access to files in the S3 bucket. The company wants to grant the application secure access to the files in the S3 bucket.

Which solution will meet these requirements?

- A. Create an S3 bucket policy that permits access to the bucket from the public IP address range of the company's on-premises data center.
- B. Use IAM Roles Anywhere to obtain security credentials in IAM Identity Center that grant access to the S3 bucket. Configure the virtual machines to assume the role by using the AWS CLI. **Most Voted**
- C. Install the AWS CLI on the virtual machine. Configure the AWS CLI with access keys from an IAM user that has access to the bucket.
- D. Create an IAM user and policy that grants access to the bucket. Store the access key and secret key for the IAM user in AWS Secrets Manager. Configure the application to retrieve the access key and secret key at startup.

Correct Answer: B

Community vote distribution

B (89%)

11%

Comments

bora4motion 1 year ago

Selected Answer: D

It does not say there's a VPN to AWS or DX to AWS.
upvoted 1 times

FlyingHawk 1 year, 4 months ago

Selected Answer: B

You can use AWS Identity and Access Management Roles Anywhere to obtain temporary security credentials in IAM for workloads such as servers, containers, and applications that run outside of AWS. Your workloads can use the same IAM policies and IAM roles that you use with AWS applications to access AWS resources. Using IAM Roles Anywhere means you don't need to manage long-term credentials for workloads running outside of AWS.

<https://docs.aws.amazon.com/rolesanywhere/latest/userguide/introduction.html>

upvoted 2 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: B

You have sensitive data in an S3 bucket, and an on-premises application needs temporary access. The solution must be secure and leverage IAM Identity Center (already in use), so now we need to grant temporary and secure credentials to the application while avoiding static credentials (like access keys).

A - IP-based access lacks identity verification, leaving the bucket vulnerable to IP spoofing.

B - This way you're not introducing static credentials or deviating from your existing centralized identity setup.

C - Manually managing credentials for multiple machines or applications is a hassle.

D - Requires additional setup (Secrets Manager integration) and introduces ongoing costs. Also, the underlying issue (static credentials) is still there.

upvoted 1 times

Bwhizzy 1 year, 8 months ago

Selected Answer: B

Answer is B.

AM Roles Anywhere allows on-premises servers and applications to obtain temporary AWS credentials and access AWS resources securely. This

solution allows your on-premises virtual machines to use IAM roles without needing long-term credentials (like access keys). The virtual machines can assume roles and access the S3 bucket temporarily and securely.

Since the company is already using AWS IAM Identity Center, using IAM Roles Anywhere allows the company to leverage its existing Identity Center setup while following AWS best practices for security. This approach ensures the application can securely retrieve credentials without embedding static credentials into the application.

upvoted 3 times

aragon_saa 1 year, 8 months ago

Selected Answer: B

Answer is B

upvoted 2 times

A company hosts its core network services, including directory services and DNS, in its on-premises data center. The data center is connected to the AWS Cloud using AWS Direct Connect (DX). Additional AWS accounts are planned that will require quick, cost-effective, and consistent access to these network services.

What should a solutions architect implement to meet these requirements with the LEAST amount of operational overhead?

- A. Create a DX connection in each new account. Route the network traffic to the on-premises servers.
- B. Configure VPC endpoints in the DX VPC for all required services. Route the network traffic to the on-premises servers.
- C. Create a VPN connection between each new account and the DX VPC route the network traffic to the on-premises servers.
- D. Configure AWS Transit Gateway between the accounts. Assign DX to the transit gateway and route network traffic to the on-premises servers. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

aragon_saa **Highly Voted** 1 year, 8 months ago

Selected Answer: D

Answer is D

upvoted 7 times

LeonSauveterre **Most Recent** 1 year, 5 months ago

Selected Answer: D

New AWS accounts need consistent, quick, and cost-effective access to these services, so we can't do something for each of these accounts. We should find a way to manage these accounts fast.

A - Managing separate DX connections for every account is too much work and DX connections are expensive.

B - VPC endpoints are for accessing AWS services, not on-premises services.

C - Managing individual VPNs for multiple accounts is too much work and VPNs bring more latency.

D - Transit Gateway simplifies the network by acting as a hub for all VPCs and the on-premises network. Instead of multiple DX or VPN connections, you're using one centralized DX connection with Transit Gateway. The best part is that you can easily add new accounts or VPCs without reconfiguring DX or VPN connections.

upvoted 3 times

JA2018 1 year, 5 months ago

Selected Answer: D

I will choose Option D for the following reasons:

#Least operational overhead:

- Using a Transit Gateway centralizes network management, eliminating the need to create separate DX connections or VPNs for each new account, significantly reducing operational overhead.

#Scalability:

- With a Transit Gateway, you can easily add new accounts to the network without modifying existing configurations, making it ideal for scenarios with frequent account creation.

#Cost-effective:

- While creating a dedicated DX connection for each account might seem quick, it can become expensive, especially when dealing with a large number of accounts.

upvoted 1 times

A company hosts its main public web application in one AWS Region across multiple Availability Zones. The application uses an Amazon EC2 Auto Scaling group and an Application Load Balancer (ALB).

A web development team needs a cost-optimized compute solution to improve the company's ability to serve dynamic content globally to millions of customers.

Which solution will meet these requirements?

- A. Create an Amazon CloudFront distribution. Configure the existing ALB as the origin. **Most Voted**
- B. Use Amazon Route 53 to serve traffic to the ALB and EC2 instances based on the geographic location of each customer.
- C. Create an Amazon S3 bucket with public read access enabled. Migrate the web application to the S3 bucket. Configure the S3 bucket for website hosting.
- D. Use AWS Direct Connect to directly serve content from the web application to the location of each customer.

Correct Answer: A

Community vote distribution

A (100%)

Comments

aragon_saa **Highly Voted** 1 year, 8 months ago

Selected Answer: A

Answer is A

upvoted 6 times

GOTJ **Most Recent** 1 year, 4 months ago

Cloudfront to serve dynamic content? Again?

upvoted 1 times

GOTJ 1 year, 4 months ago

I have to correct myself: <https://aws.amazon.com/cloudfront/dynamic-content/>

In this case, you can use cloudfront as an entry-point to AWS networking infrastructure for accessing to the web application, which is wide better than the public network infrastructure

upvoted 2 times

JA2018 1 year, 6 months ago

Selected Answer: A

CloudFront:

is a content delivery network (CDN) service from AWS, designed to distribute content globally at high speeds by caching content in edge locations close to users. This is the most cost-effective way to serve dynamic content globally, which is exactly what the scenario requires.

ALB as the origin:

By setting the existing ALB as the origin for CloudFront, the CDN will fetch content from the web application running on the EC2 instances behind the ALB, effectively caching and delivering content from geographically closer locations to users

upvoted 2 times

Cpso 1 year, 6 months ago

Selected Answer: A

"A. But I'm curious why the exam usually uses CloudFront for dynamic content. It doesn't seem reasonable."

upvoted 4 times

tttttttttttttttttttttt 1 year, 6 months ago

Selected Answer: A

Cloudfront == global

upvoted 3 times

A company stores user data in AWS. The data is used continuously with peak usage during business hours. Access patterns vary, with some data not being used for months at a time. A solutions architect must choose a cost-effective solution that maintains the highest level of durability while maintaining high availability.

Which storage solution meets these requirements?

- A. Amazon S3 Standard
- B. Amazon S3 Intelligent-Tiering **Most Voted**
- C. Amazon S3 Glacier Deep Archive
- D. Amazon S3 One Zone-Infrequent Access (S3 One Zone-IA)

Correct Answer: B

Community vote distribution

B (100%)

Comments

AMEJack 1 year, 6 months ago

Selected Answer: B

Vary then intelligent tier

upvoted 2 times

Rayan999 1 year, 7 months ago

Selected Answer: B

B of course. finally i reached here

upvoted 4 times

viejito 1 year, 8 months ago

Respuesta correcta B . Los patrones de acceso varían = Nivel inteligente de Amazon S3 .

upvoted 2 times

A company is testing an application that runs on an Amazon EC2 Linux instance. A single 500 GB Amazon Elastic Block Store (Amazon EBS) General Purpose SSO (gp2) volume is attached to the EC2 instance.

The company will deploy the application on multiple EC2 instances in an Auto Scaling group. All instances require access to the data that is stored in the EBS volume. The company needs a highly available and resilient solution that does not introduce significant changes to the application's code.

Which solution will meet these requirements?

- A. Provision an EC2 instance that uses NFS server software. Attach a single 500 GB gp2 EBS volume to the instance.
- B. Provision an Amazon FSx for Windows File Server file system. Configure the file system as an SMB file store within a single Availability Zone.
- C. Provision an EC2 instance with two 250 GB Provisioned IOPS SSD EBS volumes.
- D. Provision an Amazon Elastic File System (Amazon EFS) file system. Configure the file system to use General Purpose performance mode. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

GOTJ 1 year, 4 months ago

Selected Answer: D

In my opinion, this question is not written properly, since it doesn't clearly states that the original EBS volume contains files within a filesystem. Changing "All instances require access to the data that is stored in the EBS volume" with "All instances require access to the FILES stored in the EBS volume" would do the trick. Assuming this, "D" should be the right answer.

upvoted 1 times

JA2018 1 year, 6 months ago

Selected Answer: D

3 reasons why I will choose Option D:

#1 - Highly Available and Scalable: EFS is a managed file system designed for high availability and can automatically scale to accommodate the need of multiple EC2 instances in an Auto Scaling group, which perfectly aligns with the company's requirement.

#2 - No code changes: Since EFS is a network file system, the application doesn't need significant code changes to access the data stored on it, as it can mount the file system like a local disk.

#3 - General Purpose performance mode: This mode provides a balance between cost and performance, suitable for most application workloads.

upvoted 3 times

aragon_saa 1 year, 8 months ago

Selected Answer: D

Answer is D

upvoted 4 times

aragon_saa 1 year, 8 months ago

Selected Answer: D

Answer is D

upvoted 3 times

A company recently launched a new application for its customers. The application runs on multiple Amazon EC2 instances across two Availability Zones. End users use TCP to communicate with the application.

The application must be highly available and must automatically scale as the number of users increases.

Which combination of steps will meet these requirements MOST cost-effectively? (Choose two.)

A. Add a Network Load Balancer in front of the EC2 instances. **Most Voted**

B. Configure an Auto Scaling group for the EC2 instances. **Most Voted**

C. Add an Application Load Balancer in front of the EC2 instances.

D. Manually add more EC2 instances for the application.

E. Add a Gateway Load Balancer in front of the EC2 instances.

Correct Answer: AB

Community vote distribution

AB (79%)

11%

11%

Comments

blehbleh **Highly Voted** 1 year, 8 months ago

Selected Answer: AB

This is A and B. TCP best option is NLB. I know tons of you want to use chat GPT and it's going to tell you application load balancer, then ask chatgpt if an application load balancer really is the best option for TCP and it will be like "aw dang dawg, its not, you're right". Then it will switch to A and B. Because chatgpt doesn't know everything. It's a great tool but you still need to research because it doesn't have all the answers.

upvoted 8 times

classic_manda 1 year, 7 months ago

You are wrong! The question says, what will be the most cost effective way. ALB is much cheaper than the NLB. Thus, the best way to use NLB is when your application requires TCP and UDP connection. In this scenario, the users only communicate using TCP, which the ALB supports including variety of protocols such as, HTTP, HTTPS & SSL.

Additionally, ALB's service uptime is up to 99.995%, which falls to the requirement of this question that application must be highly available... and that nullifies your logic!

Correct answer is C

upvoted 2 times

7dcef09 **Most Recent** 1 year, 1 month ago

Selected Answer: AB

I initially selected BC because the question mentions "application" which I immediately assumed Layer 7 communication. However, the question also mentions "End users use TCP to communicate with the application." and looking at the AWS's ELB's comparison, NLB seems like a right fit. Regardless, the question is too vague to fully understand the correct requirements.

<https://aws.amazon.com/elasticloadbalancing/features/#compare>

upvoted 1 times

Dantecito 1 year, 3 months ago

Selected Answer: AB

TCP = NLB

highly available and automatically scale = Auto scaling group

upvoted 1 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: AB

Even for basic TCP, ALBs are still operating at Layer 7 to some extent. They have to do more processing per connection than an NLB, which operates at Layer 4. This translates to higher resource consumption on the ALB side, even for simple TCP forwarding.

Also, both ALBs and NLBs have similar pricing structures based on

1. Hourly usage: You pay for the time the load balancer is running. This is generally comparable between the two.
2. Load Balancer Capacity Units (LCUs)

For TCP traffic, NLBs are significantly more efficient in terms of LCU consumption, especially regarding new connections and bytes processed. Even with low connection rates, the overhead of the ALB's Layer 7 processing will usually result in higher LCU consumption and therefore higher costs.

upvoted 1 times

LeonSauveterre 1 year, 5 months ago

NLBs are designed to handle millions of connections per second with very low latency. Even if the application in the question doesn't require "extreme connection scaling" right now, using an NLB provides headroom for future growth without incurring significant cost increases. The cost difference between a lightly used NLB and an ALB handling the same basic TCP traffic is usually not substantial enough to justify the potential performance bottleneck and lack of scalability of the ALB for TCP. Lastly, NLB configuration for TCP is generally simpler than configuring an ALB for TCP.

upvoted 1 times

Ghoneam 1 year, 5 months ago

Selected Answer: BC

ALB for TCP

upvoted 1 times

siheom 1 year, 7 months ago

Selected Answer: BC

VOTE BC

upvoted 1 times

classic_manda 1 year, 7 months ago

Selected Answer: C

ALB is much cheaper than the NLB. Thus, the best way to use NLB is when your application requires TCP and UDP connection. In this scenario, the users only communicate using TCP, which the ALB supports including variety of protocols such as, HTTP, HTTPS & SSL. Additionally, ALB's service uptime is up to 99.995%, which falls to the requirement of this question that application must be highly available.

upvoted 2 times

classic_manda 1 year, 7 months ago

Answer is B and C

upvoted 1 times

1a0d459 1 year, 8 months ago

Selected Answer: AB

TCP - NLB

upvoted 4 times

A company is designing the architecture for a new mobile app that uses the AWS Cloud. The company uses organizational units (OUs) in AWS Organizations to manage its accounts. The company wants to tag Amazon EC2 instances with data sensitivity by using values of sensitive and nonsensitive. IAM identities must not be able to delete a tag or create instances without a tag.

Which combination of steps will meet these requirements? (Choose two.)

A. In Organizations, create a new tag policy that specifies the data sensitivity tag key and the required values. Enforce the tag values for the EC2 instances. Attach the tag policy to the appropriate OU. **Most Voted**

B. In Organizations, create a new service control policy (SCP) that specifies the data sensitivity tag key and the required tag values. Enforce the tag values for the EC2 instances. Attach the SCP to the appropriate OU.

C. Create a tag policy to deny running instances when a tag key is not specified. Create another tag policy that prevents identities from deleting tags. Attach the tag policies to the appropriate OU.

D. Create a service control policy (SCP) to deny creating instances when a tag key is not specified. Create another SCP that prevents identities from deleting tags. Attach the SCPs to the appropriate OU. **Most Voted**

E. Create an AWS Config rule to check if EC2 instances use the data sensitivity tag and the specified values. Configure an AWS Lambda function to delete the resource if a noncompliant resource is found.

Correct Answer: AD

Community vote distribution

AD (100%)

Comments

FlyingHawk 1 year, 4 months ago

Selected Answer: AD

As depicted in this aws blog

<https://aws.amazon.com/blogs/mt/implement-aws-resource-tagging-strategy-using-aws-tag-policies-and-service-control-policies-scps/>

A - Tag Policy only enforces the accepted value of a tag, and not its presence. (step 1)

D - SCP Policies enforce the presence of tag. (step 2 and 3)

upvoted 1 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: AD

B - SCPs cannot directly enforce tag values; they can only allow or deny actions based on tags.

C - Tag policies in AWS Organizations can't deny actions, such as launching untagged instances or deleting tags.

E - AWS Lambda is definitely irrelevant.

upvoted 2 times

Rayan999 1 year, 7 months ago

Selected Answer: AD

Ofcourse A and D. B is little bit close to the answer but its not

upvoted 3 times

jingen11 1 year, 7 months ago

Selected Answer: AD

A: https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scps_examples_tagging.html

D: https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scps_examples_tagging.html

upvoted 4 times

viejito 1 year, 8 months ago

Respuesta A D .

upvoted 3 times

A company runs database workloads on AWS that are the backend for the company's customer portals. The company runs a Multi-AZ database cluster on Amazon RDS for PostgreSQL.

The company needs to implement a 30-day backup retention policy. The company currently has both automated RDS backups and manual RDS backups. The company wants to maintain both types of existing RDS backups that are less than 30 days old.

Which solution will meet these requirements MOST cost-effectively?

- A. Configure the RDS backup retention policy to 30 days for automated backups by using AWS Backup. Manually delete manual backups that are older than 30 days.
- B. Disable RDS automated backups. Delete automated backups and manual backups that are older than 30 days. Configure the RDS backup retention policy to 30 days for automated backups.
- C. Configure the RDS backup retention policy to 30 days for automated backups. Manually delete manual backups that are older than 30 days. **Most Voted**
- D. Disable RDS automated backups. Delete automated backups and manual backups that are older than 30 days automatically by using AWS CloudFormation. Configure the RDS backup retention policy to 30 days for automated backups.

Correct Answer: C

Community vote distribution

C (89%)

11%

Comments

blehbleh **Highly Voted** 1 year, 8 months ago

Selected Answer: C

This is C, you are looking for the most cost effective solution. Again if any of you use ChatGPT it will say A because it makes the most sense for automation and less management. But it is not the most cost effective. AWS Backup costs extra money. So A, is not correct. It is C.

upvoted 11 times

XXXXXINN 1 year, 8 months ago

Agree!

Additionally we can configure RDS backup easily for a built-in 30 days retention, which meets the requirement and less cost than using AWS backup service

upvoted 2 times

LeonSauveterre **Highly Voted** 1 year, 5 months ago

Selected Answer: C

A - AWS Backup is not required for configuring automated backup retention on RDS. It's native to the RDS service. So not using that will save us a fortune.

B - Disabling backups and then re-enabling a 30-day retention policy? I don't get it

C - Manual deletion will add operational tasks for sure, but in this case, it's actually reasonable and acceptable.

D - Even sounds more absurd than B, because using AWS CloudFormation to automate deletion adds extra complexity.

upvoted 5 times

jacinml **Most Recent** 1 year, 8 months ago

Selected Answer: D

D makes more sense, since AWS CF will keep last 30 days from now on.

upvoted 2 times

viejito 1 year, 8 months ago

.respuesta correcta A-D : Para cumplir con los requisitos de etiquetado de instancias de Amazon EC2 con sensibilidad de datos y garantizar que las identidades de IAM no puedan eliminar una etiqueta o crear instancias sin una etiqueta, la combinación de pasos más adecuada es:

A. En Organizaciones, cree una nueva política de etiqueta que especifique la clave de etiqueta de sensibilidad de datos y los valores requeridos. Aplica los valores de etiqueta para las instancias de EC2. Adjunte la política de etiquetas a la OU apropiada.

D. Cree una política de control de servicio (SCP) para denegar la creación de instancias cuando no se especifica una clave de etiqueta. Crea otro SCP que evite que las identidades eliminen las etiquetas. Adjunte los SCP a la OU apropiada.

upvoted 1 times

A company is planning to migrate a legacy application to AWS. The application currently uses NFS to communicate to an on-premises storage solution to store application data. The application cannot be modified to use any other communication protocols other than NFS for this purpose.

Which storage solution should a solutions architect recommend for use after the migration?

- A. AWS DataSync
- B. Amazon Elastic Block Store (Amazon EBS)
- C. Amazon Elastic File System (Amazon EFS) **Most Voted**
- D. Amazon EMR File System (Amazon EMRFS)

Correct Answer: C

Community vote distribution

C (100%)

Comments

78b9037 1 year, 6 months ago

Selected Answer: C

Amazon EFS is the only one which provides a fully managed NFS file system
upvoted 2 times

aragon_saa 1 year, 8 months ago

Selected Answer: C

Answer is C
upvoted 4 times

A company uses GPS trackers to document the migration patterns of thousands of sea turtles. The trackers check every 5 minutes to see if a turtle has moved more than 100 yards (91.4 meters). If a turtle has moved, its tracker sends the new coordinates to a web application running on three Amazon EC2 instances that are in multiple Availability Zones in one AWS Region.

Recently, the web application was overwhelmed while processing an unexpected volume of tracker data. Data was lost with no way to replay the events. A solutions architect must prevent this problem from happening again and needs a solution with the least operational overhead.

What should the solutions architect do to meet these requirements?

- A. Create an Amazon S3 bucket to store the data. Configure the application to scan for new data in the bucket for processing.
- B. Create an Amazon API Gateway endpoint to handle transmitted location coordinates. Use an AWS Lambda function to process each item concurrently.
- C. Create an Amazon Simple Queue Service (Amazon SQS) queue to store the incoming data. Configure the application to poll for new messages for processing.
- D. Create an Amazon DynamoDB table to store transmitted location coordinates. Configure the application to query the table for new data for processing. Use TTL to remove data that has been processed.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Bwhizzy 1 year, 7 months ago

Selected Answer: C

The correct Answer is C. Amazon SQS is a fully managed message queuing service that allows you to decouple and scale applications by buffering the incoming data. This would ensure that data is stored in a reliable, scalable queue until the web application is ready to process it. SQS provides high availability and fault tolerance, and guarantees that no data will be lost, as the messages remain in the queue until they are processed.

The web application can poll the SQS queue for new messages and process them at its own pace, preventing the application from being overwhelmed by large volumes of tracker data.

Operational overhead is minimal because SQS is a fully managed service. The application only needs to poll the queue for messages, and there is no need to manage infrastructure.

upvoted 3 times

hharbiordun85 1 year, 8 months ago

Answer: C

Amazon SQS provides a reliable and scalable way to decouple the components of the application. By using a queue, the incoming tracker data can be buffered, ensuring that if the web application is overwhelmed with data, the messages will still be retained in the queue until they can be processed.

upvoted 3 times

A company's software development team needs an Amazon RDS Multi-AZ cluster. The RDS cluster will serve as a backend for a desktop client that is deployed on premises. The desktop client requires direct connectivity to the RDS cluster.

The company must give the development team the ability to connect to the cluster by using the client when the team is in the office.

Which solution provides the required connectivity MOST securely?

- A. Create a VPC and two public subnets. Create the RDS cluster in the public subnets. Use AWS Site-to-Site VPN with a customer gateway in the company's office.
- B. Create a VPC and two private subnets. Create the RDS cluster in the private subnets. Use AWS Site-to-Site VPN with a customer gateway in the company's office. **Most Voted**
- C. Create a VPC and two private subnets. Create the RDS cluster in the private subnets. Use RDS security groups to allow the company's office IP ranges to access the cluster.
- D. Create a VPC and two public subnets. Create the RDS cluster in the public subnets. Create a cluster user for each developer. Use RDS security groups to allow the users to access the cluster.

Correct Answer: B

Community vote distribution

B (83%)

C (17%)

Comments

Bwhizzy **Highly Voted** 1 year, 7 months ago

Selected Answer: B

The Correct Answer is B.

Explanation:

VPC and Private Subnets: By placing the RDS cluster in private subnets, you ensure that the RDS cluster is not publicly accessible from the internet. This significantly improves security as the database is only accessible through secure channels, not directly from the public internet.

AWS Site-to-Site VPN: Using a Site-to-Site VPN establishes a secure, encrypted connection between the on-premises office and the AWS environment. This provides secure access to the RDS cluster without exposing it to the internet, ensuring that the developers can only access the cluster when connected to the office network.

Customer Gateway: The customer gateway is configured in the company's office to handle the VPN connection, providing secure connectivity for the desktop client to the RDS cluster when the development team is in the office.

upvoted 9 times

blehbleh **Highly Voted** 1 year, 8 months ago

Selected Answer: B

This is B site to site von adds additional security. We are going for more secure.

upvoted 6 times

kbgsgsgs **Most Recent** 1 year, 8 months ago

Selected Answer: C

The goal is to limit the team to only being in the office to be in the RDS cluster, so wouldn't checking IP ranges based on the office network rather than bringing up the internet be better suited to what you really need?

upvoted 3 times

trongod05 1 year, 8 months ago

But if they are in private subnets, how do they connect? Can't over public internet. And there's no connection between their office and the VPC. Needs more info I think.

upvoted 2 times

A solutions architect is creating an application that will handle batch processing of large amounts of data. The input data will be held in Amazon S3 and the output data will be stored in a different S3 bucket. For processing, the application will transfer the data over the network between multiple Amazon EC2 instances.

What should the solutions architect do to reduce the overall data transfer costs?

- A. Place all the EC2 instances in an Auto Scaling group.
- B. Place all the EC2 instances in the same AWS Region.
- C. Place all the EC2 instances in the same Availability Zone. **Most Voted**
- D. Place all the EC2 instances in private subnets in multiple Availability Zones.

Correct Answer: C

Community vote distribution

C (83%)

B (17%)

Comments

FlyingHawk 1 year, 4 months ago

Selected Answer: C

Data transfer within the zone is free, but cross-zone is not free:

<https://aws.amazon.com/blogs/architecture/overview-of-data-transfer-costs-for-common-architectures/>

upvoted 1 times

ckhemani 1 year, 6 months ago

Selected Answer: C

K. After research it seems C make sense since it is saying data transfer between EC2 and data transfer between two EC2 in AZ is free but not between AZ's.

No cost within AZ: Data transfer between EC2 instances in the same AZ is free.

Cross-AZ charges: If you transfer data between instances in different Availability Zones, you will be charged for the inter-zone data transfer.

Region-based pricing: While data transfer within an AZ is free, data transfer across different regions may also incur charges

upvoted 3 times

ckhemani 1 year, 6 months ago

Selected Answer: B

S3 is replicated only to 3 AZ's and you can't define AZ while creating the Bucket. If this Question would have said about Directory which is One Zone express where you can define the AZ then it makes sense to select option C.

Option B makes sense to configure EC2 in same region since EC2/S3 transfer data between same region is free too.

upvoted 2 times

78b9037 1 year, 6 months ago

Selected Answer: C

Two reasons:

1. EC2 instances in the same AZ can communicate with each other at no cost.

2. Data transfer between EC2 instances and S3 within the same AZ is free

<https://aws.amazon.com/blogs/architecture/overview-of-data-transfer-costs-for-common-architectures/>

upvoted 2 times

sOI852POL 1 year, 7 months ago

Selected Answer: C

C

upvoted 2 times

viejito 1 year, 8 months ago

Respuesta correcta C :

upvoted 2 times

aragon_saa 1 year, 8 months ago

Selected Answer: C

Answer is C

upvoted 2 times

A company hosts a multi-tier web application that uses an Amazon Aurora MySQL DB cluster for storage. The application tier is hosted on Amazon EC2 instances. The company's IT security guidelines mandate that the database credentials be encrypted and rotated every 14 days.

What should a solutions architect do to meet this requirement with the LEAST operational effort?

A. Create a new AWS Key Management Service (AWS KMS) encryption key. Use AWS Secrets Manager to create a new secret that uses the KMS key with the appropriate credentials. Associate the secret with the Aurora DB cluster. Configure a custom rotation period of 14 days. **Most Voted**

B. Create two parameters in AWS Systems Manager Parameter Store: one for the user name as a string parameter and one that uses the SecureString type for the password. Select AWS Key Management Service (AWS KMS) encryption for the password parameter, and load these parameters in the application tier. Implement an AWS Lambda function that rotates the password every 14 days.

C. Store a file that contains the credentials in an AWS Key Management Service (AWS KMS) encrypted Amazon Elastic File System (Amazon EFS) file system. Mount the EFS file system in all EC2 instances of the application tier. Restrict the access to the file on the file system so that the application can read the file and that only super users can modify the file. Implement an AWS Lambda function that rotates the key in Aurora every 14 days and writes new credentials into the file.

D. Store a file that contains the credentials in an AWS Key Management Service (AWS KMS) encrypted Amazon S3 bucket that the application uses to load the credentials. Download the file to the application regularly to ensure that the correct credentials are used. Implement an AWS Lambda function that rotates the Aurora credentials every 14 days and uploads these credentials to the file in the S3 bucket.

Correct Answer: A

Community vote distribution

A (100%)

Comments

LeonSauveterre 1 year, 5 months ago

Selected Answer: A

A - AWS KMS + Secrets Manager support custom rotation intervals, such as 14 days, with built-in integration for Aurora DB clusters.
B, C, D - Manual implementation of rotation? This is really hard.

upvoted 1 times

viejito 1 year, 8 months ago

respuesta correcta A . Credenciales , rotación automática =AWS Secrets Manager .

upvoted 2 times

aragon_saa 1 year, 8 months ago

Selected Answer: A

Answer is A

upvoted 2 times

A streaming media company is rebuilding its infrastructure to accommodate increasing demand for video content that users consume daily.

The company needs to process terabyte-sized videos to block some content in the videos. Video processing can take up to 20 minutes.

The company needs a solution that will scale with demand and remain cost-effective.

Which solution will meet these requirements?

A. Use AWS Lambda functions to process videos. Store video metadata in Amazon DynamoDB. Store video content in Amazon S3 Intelligent-Tiering.

B. Use Amazon Elastic Container Service (Amazon ECS) and AWS Fargate to implement microservices to process videos. Store video metadata in Amazon Aurora. Store video content in Amazon S3 Intelligent-Tiering. **Most Voted**

C. Use Amazon EC2 instances in an Auto Scaling group behind an Application Load Balancer (ALB) to process videos. Store video content in Amazon S3 Standard. Use Amazon Simple Queue Service (Amazon SQS) for queuing and to decouple processing tasks.

D. Deploy a containerized video processing application on Amazon Elastic Kubernetes Service (Amazon EKS) on Amazon EC2. Store video metadata in Amazon RDS in a single Availability Zone. Store video content in Amazon S3 Glacier Deep Archive.

Correct Answer: B

Community vote distribution

B (85%)

C (15%)

Comments

kbgsqsgs **Highly Voted** 1 year, 8 months ago

Selected Answer: B

S3 Intelligent-Tiering is cost-effective for storing large amounts of video content, and since Lambda doesn't work, shouldn't we consider serverless?
upvoted 8 times

LeonSauveterre **Highly Voted** 1 year, 5 months ago

Selected Answer: B

A - Video processing can take up to 20 minutes. So lambda is not suitable here.
B - Fargate allows running containers without managing EC2 instances, scaling automatically based on demand. Besides, S3 Intelligent-Tiering is cost-effective for storing large video files.
C - I gotta say, the SQS is a really good choice here! It's just that S3 Standard is more expensive than S3 Intelligent-Tiering for large files, and that EC2 (even with Auto Scaling) requires managing the underlying infrastructure.
D - EKS is more complex to set up but the real issue is that RDS in a single Availability Zone introduces a single point of failure, so scalability and reliability can't be met. Also, S3 Glacier Deep Archive is designed for long-term archival storage with high latency and not suitable for frequently accessed video content.
upvoted 6 times

EllenLiu **Most Recent** 1 year, 5 months ago

Selected Answer: B

it is a tough question.
terabyte-sized videos are difficult to process. we need to split the videos into chunks either with fargate or EC2.
seems DB is essential in this scenario as I have no idea how to get the videos out from s3 without metadata which should be stored into DB. So I prefer B. SQS is used for peak clipping during burst periods.
upvoted 1 times

EllenLiu 1 year, 5 months ago

this the whole work flow for a ideal solution:

- 1.Upload: Users upload videos to CloudFront with origin S3.
 - 2.Trigger: S3 sends an event notification to SQS.
 - 3.EC2 instances poll the SQS queue and download and process Video files from S3
 - 4.The processed video is uploaded to another S3 bucket, metadata is saved into Aurora
- upvoted 1 times

78b9037 1 year, 6 months ago

Selected Answer: B

B and C are both possible but i decided for B because:

- Fargate provides serverless containers that scale automatically
- No infrastructure management is required
- Aurora provides high availability for metadata
- S3 Intelligent-Tiering optimizes storage costs automatically
- It is not tile limited like Lambda

On the other hand C requires more management overhead it is less cost-effective than containerized solutions, and need to manage EC2 instances manually

upvoted 2 times

XXXXXINN 1 year, 8 months ago

This is a tough one I would say. I wish B and C can combine. but I vote for B.

For B, it is more cost efficiency focus rather than decoupling the whole process for improving overall reliability.

For C, the use of SQS is perfect solution for the downside of option B. But EC2 comes in picture which increases the cost and operational complexity

How to pick then? lets go back to the question and see what it focuses - cost or operational complexity and stability? It looks like it leans more on focusing scalability and cost-efficiency. In that case, I would go for B because fargate provides cost-efficiency and store just metadata in DB and the rest data in S3 also provides a lower cost and improves its performance.

upvoted 2 times

XXXXXINN 1 year, 8 months ago

Additionally, C also introduced ALB service in the picture, that increases the cost as well.

upvoted 2 times

blehbleh 1 year, 8 months ago

Selected Answer: C

C makes the most sense out of the options and given requirements.

upvoted 1 times

hharbiordun85 1 year, 8 months ago

Answer: C

The question did not mention the application need to be containerized , i will choose C

upvoted 1 times

GOTJ 1 year, 4 months ago

But it mentioned "rebuilding its infrastructure", opening the door to a re-engineering process in which micro services can be applied

upvoted 2 times

blehbleh 1 year, 8 months ago

Selected Answer: C

I personally think C, I could be wrong.

upvoted 2 times

A company runs an on-premises application on a Kubernetes cluster. The company recently added millions of new customers. The company's existing on-premises infrastructure is unable to handle the large number of new customers. The company needs to migrate the on-premises application to the AWS Cloud.

The company will migrate to an Amazon Elastic Kubernetes Service (Amazon EKS) cluster. The company does not want to manage the underlying compute infrastructure for the new architecture on AWS.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use a self-managed node to supply compute capacity. Deploy the application to the new EKS cluster.
- B. Use managed node groups to supply compute capacity. Deploy the application to the new EKS cluster.
- C. Use AWS Fargate to supply compute capacity. Create a Fargate profile. Use the Fargate profile to deploy the application.
- D. Use managed node groups with Karpenter to supply compute capacity. Deploy the application to the new EKS cluster.

Most Voted

Correct Answer: C

Community vote distribution

C (100%)

Comments

Razvan_C **Highly Voted** 1 year, 7 months ago

Selected Answer: C

"The company does not want to manage the underlying compute infrastructure for the new architecture on AWS" => Serverless => Fargate
upvoted 7 times

FlyingHawk **Most Recent** 1 year, 4 months ago

Selected Answer: C

With Fargate, you don't have to provision, configure, or scale groups of virtual machines on your own to run containers. You also don't need to choose server types, decide when to scale your node groups, or optimize cluster packing.
A Fargate profile is used to specify which Kubernetes pods should run on Fargate. By creating a Fargate profile, you can ensure that your application runs on Fargate without needing to manage node groups or EC2 instances.
upvoted 2 times

elmyth 1 year, 7 months ago

Selected Answer: C

Fargate
upvoted 3 times

viejito 1 year, 8 months ago

Respuesta correcta C = AWS Fargate = menor sobrecarga operativa. A, B y D tienen gestión de infraestructura.
upvoted 2 times

viejito 1 year, 8 months ago

Respuesta correcta C = AWS Fargate = menor sobrecarga operativa. A, B y C tienen gestión de infraestructura.
upvoted 2 times

A company is launching a new application that requires a structured database to store user profiles, application settings, and transactional data. The database must be scalable with application traffic and must offer backups.

Which solution will meet these requirements MOST cost-effectively?

- A. Deploy a self-managed database on Amazon EC2 instances by using open source software. Use Spot Instances for cost optimization. Configure automated backups to Amazon S3.
- B. Use Amazon RDS. Use on-demand capacity mode for the database with General Purpose SSD storage. Configure automatic backups with a retention period of 7 days.
- C. Use Amazon Aurora Serverless for the database. Use serverless capacity scaling. Configure automated backups to Amazon S3. **Most Voted**
- D. Deploy a self-managed NoSQL database on Amazon EC2 instances. Use Reserved Instances for cost optimization. Configure automated backups directly to Amazon S3 Glacier Flexible Retrieval.

Correct Answer: C

Community vote distribution

C (100%)

Comments

aragon_saa **Highly Voted** 1 year, 8 months ago

Selected Answer: C

Answer is C

upvoted 5 times

Itachi28 **Most Recent** 1 year, 2 months ago

Selected Answer: C

answer is C

upvoted 1 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: C

A - Spot Instance is never a good idea for databases. These instances can be terminated with little notice, while databases are almost always persistent workloads.

B - On-Demand Capacity is suitable for steady or predictable workloads, but not as cost-efficient as Aurora Serverless for spiky or unpredictable workloads.

C - Aurora Serverless is a fully managed relational database that scales automatically with traffic. Automated Backups is also integrated with Aurora. Good!

D - We need a structured database, not NoSQL database.

upvoted 2 times

A company runs its legacy web application on AWS. The web application server runs on an Amazon EC2 instance in the public subnet of a VPC. The web application server collects images from customers and stores the image files in a locally attached Amazon Elastic Block Store (Amazon EBS) volume. The image files are uploaded every night to an Amazon S3 bucket for backup.

A solutions architect discovers that the image files are being uploaded to Amazon S3 through the public endpoint. The solutions architect needs to ensure that traffic to Amazon S3 does not use the public endpoint.

Which solution will meet these requirements?

- A. Create a gateway VPC endpoint for the S3 bucket that has the necessary permissions for the VPC. Configure the subnet route table to use the gateway VPC endpoint. **Most Voted**
- B. Move the S3 bucket inside the VPC. Configure the subnet route table to access the S3 bucket through private IP addresses.
- C. Create an Amazon S3 access point for the Amazon EC2 instance inside the VPC. Configure the web application to upload by using the Amazon S3 access point.
- D. Configure an AWS Direct Connect connection between the VPC that has the Amazon EC2 instance and Amazon S3 to provide a dedicated network path.

Correct Answer: A

Community vote distribution

A (73%)

C (27%)

Comments

FlyingHawk 1 year, 4 months ago

Selected Answer: A

A gateway VPC endpoint allows private connectivity between your VPC and Amazon S3. It ensures that traffic between the VPC and S3 stays within the AWS network and does not traverse the public internet.

B - S3 buckets cannot be "moved inside the VPC." S3 is a regional service and does not reside within a VPC. This option is not valid.

C - S3 access points are used to simplify access control for S3 buckets, but they do not provide private connectivity. Traffic to S3 access points still uses the public endpoint unless combined with a gateway VPC endpoint.

upvoted 1 times

FlyingHawk 1 year, 4 months ago

D - Direct Connect provides a dedicated network connection between on-premises environments and AWS. It is not necessary for connecting resources within AWS (e.g., EC2 instances in a VPC to S3). This option is overkill and adds unnecessary cost and complexity.

upvoted 1 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: A

A - Gateway VPC Endpoint provides private connectivity between resources in a VPC and AWS services like S3 without requiring an internet gateway, NAT gateway, or public IP address.

B - This is just bull. Amazon S3 is a global service that cannot "move" into a VPC.

C - This doesn't ensure private traffic between the EC2 instance and S3 at all. You still need VPC Endpoint.

D - This may actually work but more expensive and complex compared to option A.

upvoted 4 times

LeonSauveterre 1 year, 5 months ago

About connections:

1. VPN Connection: For secure, encrypted traffic between on-premises and AWS over the public internet.
2. Direct Connect: For dedicated private network connectivity between on-premises and AWS (like data centers with large workloads).
3. PrivateLink: For secure private connectivity to AWS services or third-party services (like Salesforce) within AWS.
4. VPC Peering: For direct connectivity between two AWS VPCs (like multi-account or region scenarios).

upvoted 4 times

LeonSauveterre 1 year, 5 months ago

5. VPC Endpoint (Gateway): Provides private connectivity between VPC and services like S3 and DynamoDB.
6. VPC Endpoint (Interface): Connects to services over private IP using an ENI in your VPC. Suitable for services like SNS, SQS, or KMS.
7. S3 Access Point: Simplifies access to S3 buckets with distinct policies for specific use cases.
8. Internet Gateway: Allows resources in a VPC to connect to the internet.
9. NAT Gateway/Instance: Allows private instances to connect to the internet or public AWS services. Aka, enables outbound internet access for instances in private subnets.

upvoted 5 times

dfgdsfgfdgreg 1 year, 4 months ago

Correct. The S3 access point provides access management, but not a private connection. A VPC endpoint allows an instance to directly upload to the S3 whereas the S3 access point still requires the EC2 to travel over the internet.

upvoted 2 times

JA2018 1 year, 5 months ago

Selected Answer: C

I will choose Option C for the following reasons:

Private access:

S3 access points allow you to create a private endpoint within your VPC that can be used to access the S3 bucket without going through the public endpoint.

Security best practice:

Using an S3 access point is considered a secure way to manage access to your S3 buckets from within your VPC.

Configuration simplicity:

You only need to configure the web application to use the S3 access point, making it a relatively straightforward implementation.

upvoted 3 times

aragon_saa 1 year, 8 months ago

Selected Answer: A

Answer is A

upvoted 3 times

A company is creating a prototype of an ecommerce website on AWS. The website consists of an Application Load Balancer, an Auto Scaling group of Amazon EC2 instances for web servers, and an Amazon RDS for MySQL DB instance that runs with the Single-AZ configuration.

The website is slow to respond during searches of the product catalog. The product catalog is a group of tables in the MySQL database that the company does not update frequently. A solutions architect has determined that the CPU utilization on the DB instance is high when product catalog searches occur.

What should the solutions architect recommend to improve the performance of the website during searches of the product catalog?

- A. Migrate the product catalog to an Amazon Redshift database. Use the COPY command to load the product catalog tables.
- B. Implement an Amazon ElastiCache for Redis cluster to cache the product catalog. Use lazy loading to populate the cache. **Most Voted**
- C. Add an additional scaling policy to the Auto Scaling group to launch additional EC2 instances when database response is slow.
- D. Turn on the Multi-AZ configuration for the DB instance. Configure the EC2 instances to throttle the product catalog queries that are sent to the database.

Correct Answer: B

Community vote distribution

B (100%)

Comments

aragon_saa **Highly Voted** 1 year, 8 months ago

Selected Answer: B

Answer is B

upvoted 7 times

LeonSauveterre **Highly Voted** 1 year, 5 months ago

Selected Answer: B

A - For transactional or low-latency queries like product catalog searches in an ecommerce website, Redshift is not suitable because it's a data warehouse solution optimized for analytical queries.

B - ElastiCache for Redis is a highly performant, in-memory caching service that can significantly reduce database load by caching frequent queries.

C - Adding more web servers can't help alleviate the load on the database. The database remains the bottleneck for product catalog queries.

D - If you turn on throttle queries, customers may gradually ask for a refund :) Let's get rid of the bad things in database instead of user experience.

BTW, what Lazy Loading means: Data is added to the cache only when requested. If not found in the cache, it is fetched from the database and added to the cache for subsequent requests.

upvoted 6 times

A company currently stores 5 TB of data in on-premises block storage systems. The company's current storage solution provides limited space for additional data. The company runs applications on premises that must be able to retrieve frequently accessed data with low latency. The company requires a cloud-based storage solution.

Which solution will meet these requirements with the MOST operational efficiency?

- A. Use Amazon S3 File Gateway. Integrate S3 File Gateway with the on-premises applications to store and directly retrieve files by using the SMB file system.
- B. Use an AWS Storage Gateway Volume Gateway with cached volumes as iSCSI targets. **Most Voted**
- C. Use an AWS Storage Gateway Volume Gateway with stored volumes as iSCSI targets.
- D. Use an AWS Storage Gateway Tape Gateway. Integrate Tape Gateway with the on-premises applications to store virtual tapes in Amazon S3.

Correct Answer: B

Community vote distribution

B (82%)

C (18%)

Comments

sOI852POL **Highly Voted** 1 year, 7 months ago

Selected Answer: B

Question says, "Apps must be able to retrieve frequently accessed data with low latency" so we go with cached volumes. We'd have chosen the stored volumes if question was about low-latency access to the entire dataset.

upvoted 12 times

FlyingHawk **Most Recent** 1 year, 4 months ago

Selected Answer: B

B - Cached volumes store only the most frequently accessed data on-premises, while the full dataset is stored in Amazon S3. This ensures low-latency access to frequently accessed data while leveraging the scalability and durability of S3 for the full dataset.

C - Stored volumes store the entire dataset on-premises and asynchronously back it up to AWS. This option does not address the company's requirement for limited on-premises storage capacity.

upvoted 2 times

FlyingHawk 1 year, 4 months ago

A - S3 File Gateway provides file-based storage (SMB or NFS) and is not suitable for block storage requirements.

D - Tape Gateway is designed for backup and archival use cases, not for low-latency access to frequently accessed data. It uses virtual tapes stored in S3 and is not suitable for block storage or real-time application access.

upvoted 2 times

EllenLiu 1 year, 5 months ago

Selected Answer: C

choose store volume mode

upvoted 1 times

FlyingHawk 1 year, 4 months ago

Stored volumes store the entire dataset on-premises and asynchronously back it up to AWS. This option does not address the company's requirement for limited on-premises storage capacity, as it requires storing the full dataset locally.

upvoted 1 times

8621a7c 1 year, 7 months ago

Selected Answer: C

Block storage use volume gateway, low latency use store volumes.

upvoted 2 times

viejito 1 year, 8 months ago

Respuesta correcta B . Baja latencia = Elastic Cache . A , C y D no cumplen un acceso con baja latencia (descartados) .

upvoted 3 times

A company operates a food delivery service. Because of recent growth, the company's order processing system is experiencing scaling problems during peak traffic hours. The current architecture includes Amazon EC2 instances in an Auto Scaling group that collect orders from an application. A second group of EC2 instances in an Auto Scaling group fulfills the orders.

The order collection process occurs quickly, but the order fulfillment process can take longer. Data must not be lost because of a scaling event.

A solutions architect must ensure that the order collection process and the order fulfillment process can both scale adequately during peak traffic hours.

Which solution will meet these requirements?

- A. Use Amazon CloudWatch to monitor the CPUUtilization metric for each instance in both Auto Scaling groups. Configure each Auto Scaling group's minimum capacity to meet its peak workload value.
- B. Use Amazon CloudWatch to monitor the CPUUtilization metric for each instance in both Auto Scaling groups. Configure a CloudWatch alarm to invoke an Amazon Simple Notification Service (Amazon SNS) topic to create additional Auto Scaling groups on demand.
- C. Provision two Amazon Simple Queue Service (Amazon SQS) queues. Use one SQS queue for order collection. Use the second SQS queue for order fulfillment. Configure the EC2 instances to poll their respective queues. Scale the Auto Scaling groups based on notifications that the queues send.
- D. Provision two Amazon Simple Queue Service (Amazon SQS) queues. Use one SQS queue for order collection. Use the second SQS queue for order fulfillment. Configure the EC2 instances to poll their respective queues. Scale the Auto Scaling groups based on the number of messages in each queue. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

sOI852POL **Highly Voted** 1 year, 7 months ago

Selected Answer: D

SQS to ensure data is not lost because of a scaling event, scale based on number of messages on the queue.

upvoted 5 times

8621a7c **Most Recent** 1 year, 7 months ago

Selected Answer: D

Go with D

upvoted 2 times

aragon_saa 1 year, 8 months ago

Selected Answer: D

Answer is D

upvoted 2 times

An online gaming company is transitioning user data storage to Amazon DynamoDB to support the company's growing user base. The current architecture includes DynamoDB tables that contain user profiles, achievements, and in-game transactions.

The company needs to design a robust, continuously available, and resilient DynamoDB architecture to maintain a seamless gaming experience for users.

Which solution will meet these requirements MOST cost-effectively?

- A. Create DynamoDB tables in a single AWS Region. Use on-demand capacity mode. Use global tables to replicate data across multiple Regions.
- B. Use DynamoDB Accelerator (DAX) to cache frequently accessed data. Deploy tables in a single AWS Region and enable auto scaling. Configure Cross-Region Replication manually to additional Regions.
- C. Create DynamoDB tables in multiple AWS Regions. Use on-demand capacity mode. Use DynamoDB Streams for Cross-Region Replication between Regions.
- D. Use DynamoDB global tables for automatic multi-Region replication. Deploy tables in multiple AWS Regions. Use provisioned capacity mode. Enable auto scaling. **Most Voted**

Correct Answer: D

Community vote distribution

D (92%)

8%

Comments

trinh_le 1 year, 6 months ago

Selected Answer: D

Following keywords

robust: provisioned capacity reduces cost

continuously available: global tables

resilient: multiple regions enhances available

upvoted 3 times

Cpso 1 year, 6 months ago

Selected Answer: D

c & d can do 2 ways rep. But C need user logic to handle conflict, so it not robust

upvoted 3 times

78b9037 1 year, 6 months ago

Selected Answer: D

<https://docs.aws.amazon.com/wellarchitected/latest/serverless-applications-lens/capacity.html>

upvoted 2 times

striker89 1 year, 7 months ago

Selected Answer: C

DynamoDB Streams for replication are less expensive than Global tables.

On-demand capacity mode can be less expensive than provisioned mode.

Multi region deployment ensure high availability.

upvoted 1 times

Bwhizzy 1 year, 7 months ago

Selected Answer: D

D. Use DynamoDB global tables for automatic multi-Region replication. Deploy tables in multiple AWS Regions. Use provisioned capacity mode. Enable auto scaling.

Explanation:

DynamoDB global tables automatically replicate data across multiple Regions, ensuring that the data is available and consistent across all Regions. This provides resilience and high availability by allowing users in different geographical locations to access data from the closest Region.

Provisioned capacity mode allows you to pre-allocate read and write capacity units, which can result in cost savings over on-demand capacity mode if the traffic is predictable. Additionally, auto scaling can be enabled to dynamically adjust the capacity based on the actual traffic, ensuring that you only pay for the capacity that you need.

Multi-Region deployment improves the resilience of the system. If a failure occurs in one Region, another Region can seamlessly take over, ensuring an uninterrupted gaming experience.

upvoted 4 times

GOTJ 1 year, 4 months ago

"if the traffic is predictable". Good point, but how do you know this from the question?

upvoted 1 times

dce7e4b 1 year ago

« Most cost efficiently »

« Growing user base »

The provisionned capacity is more cost effective but as the user base is growing you should enable auto scaling. Like this, predictable or not, you are « most cost efficient »

upvoted 1 times

A company runs its media rendering application on premises. The company wants to reduce storage costs and has moved all data to Amazon S3. The on-premises rendering application needs low-latency access to storage.

The company needs to design a storage solution for the application. The storage solution must maintain the desired application performance.

Which storage solution will meet these requirements in the MOST cost-effective way?

- A. Use Mountpoint for Amazon S3 to access the data in Amazon S3 for the on-premises application.
- B. Configure an Amazon S3 File Gateway to provide storage for the on-premises application. **Most Voted**
- C. Copy the data from Amazon S3 to Amazon FSx for Windows File Server. Configure an Amazon FSx File Gateway to provide storage for the on-premises application.
- D. Configure an on-premises file server. Use the Amazon S3 API to connect to S3 storage. Configure the application to access the storage from the on-premises file server.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Clouddon 11 months ago

Selected Answer: B

Explanation:

The scenario requires:

Low-latency access for an on-premises rendering application

All data is stored in Amazon S3

The solution should be cost-effective and maintain application performance

✅ Why Option B is the best choice:

Amazon S3 File Gateway is part of AWS Storage Gateway.

It provides a local cache of frequently accessed S3 objects via a standard file system interface (NFS/SMB).

This allows low-latency access to cloud-stored data from on-prem applications.

It's cost-effective because data is stored in Amazon S3 (cheap, durable), and only frequently used files are cached locally, reducing storage costs.
upvoted 1 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: B

A - Mountpoint is not designed for on-premises environments or low-latency access; it works well for cloud-based applications.

B - Amazon S3 File Gateway provides low-latency, on-premises access to data stored in Amazon S3 through a file-based interface. It caches frequently accessed data locally to reduce latency.

C - Copying data to FSx increases costs, as FSx storage is more expensive than S3. Let alone unnecessary complexity.

D - Direct interaction with S3 introduces higher latency for on-premises applications and lacks local caching.

upvoted 4 times

Americanman 1 year, 6 months ago

Amazon S3 File Gateway enables your existing file-based applications, devices, and workflows to use Amazon S3, without modification. Amazon S3 File Gateway securely and durably stores both file contents and metadata as objects, while providing your on-premises applications low-latency access to cached data.

upvoted 2 times

Bwhizzy 1 year, 7 months ago

Selected Answer: B

B. Configure an Amazon S3 File Gateway to provide storage for the on-premises application.

Explanation:

Amazon S3 File Gateway provides a way for on-premises applications to access objects stored in Amazon S3 as files. It caches frequently accessed data locally, which ensures low-latency access to the data. This is crucial for maintaining the performance of the rendering application.

By keeping the data in Amazon S3, the company benefits from lower storage costs compared to using other storage services like Amazon FSx, while still providing the necessary performance for the on-premises application through the local caching capabilities of the File Gateway.

The File Gateway seamlessly integrates with Amazon S3, allowing the application to access data using standard file protocols like NFS or SMB, which simplifies the setup.

upvoted 3 times

A company hosts its enterprise resource planning (ERP) system in the us-east-1 Region. The system runs on Amazon EC2 instances. Customers use a public API that is hosted on the EC2 instances to exchange information with the ERP system. International customers report slow API response times from their data centers.

Which solution will improve response times for the international customers MOST cost-effectively?

- A. Create an AWS Direct Connect connection that has a public virtual interface (VIF) to provide connectivity from each customer's data center to us-east-1. Route customer API requests by using a Direct Connect gateway to the ERP system API.
- B. Set up an Amazon CloudFront distribution in front of the API. Configure the CachingOptimized managed cache policy to provide improved cache efficiency.
- C. Set up AWS Global Accelerator. Configure listeners for the necessary ports. Configure endpoint groups for the appropriate Regions to distribute traffic. Create an endpoint in the group for the API. **Most Voted**
- D. Use AWS Site-to-Site VPN to establish dedicated VPN tunnels between Regions and customer networks. Route traffic to the API over the VPN connections.

Correct Answer: C

Community vote distribution

C (60%)

B (40%)

Comments

TewatiaAmit **Highly Voted** 1 year, 8 months ago

Selected Answer: C

CloudFront can reduce response times by caching API responses, but if the API is dynamic and not cacheable, it may not be as effective. Global Accelerator is better for improving latency when caching is not an option.

upvoted 12 times

Sergantus 1 year, 6 months ago

While dynamic content typically has low caching potential, CloudFront reduces latency by routing requests to the nearest edge location. There is also TCP Connection Reuse, which is also beneficial for low latency.

upvoted 4 times

8621a7c **Highly Voted** 1 year, 7 months ago

Selected Answer: B

CloudFront for Dynamic content (such as API acceleration and dynamic site delivery)

upvoted 10 times

SBoksh **Most Recent** 10 months, 3 weeks ago

Selected Answer: B

<https://aws.amazon.com/cloudfront/dynamic-content/>

upvoted 1 times

Clouddon 11 months ago

Selected Answer: C

Why other options are less ideal:

A. AWS Direct Connect

Very expensive and complex to set up for each international customer.

Designed for private, dedicated connections, typically for a single customer, not multiple public API users.

Not scalable or cost-effective in this scenario.

B. Amazon CloudFront

CloudFront is ideal for static content (e.g., images, scripts, web pages), not dynamic APIs.

CachingOptimized policy works best for cacheable responses, not for dynamic ERP APIs, which likely serve real-time data.

D. AWS Site-to-Site VPN

Requires customer coordination to set up VPN tunnels.

Adds network complexity and operational overhead.

Not cost-effective or scalable for public API access by many international users.

upvoted 1 times

eduardi 1 year ago

Selected Answer: B

CloudFront is better for https

upvoted 1 times

tch 1 year, 3 months ago

Selected Answer: B

How is AWS Global Accelerator different from Amazon CloudFront?

AWS Global Accelerator and Amazon CloudFront are separate services that use the AWS global network and its edge locations around the world. CloudFront improves performance for both cacheable content (such as images and videos) and dynamic content (such as API acceleration and dynamic site delivery). Global Accelerator improves performance for a wide range of applications over TCP or UDP by proxying packets at the edge to applications running in one or more AWS Regions. Global Accelerator is a good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP, as well as for HTTP use cases that specifically require static IP addresses or deterministic, fast regional failover. Both services integrate with AWS Shield for DDoS protection

upvoted 1 times

tch 1 year, 3 months ago

MOST cost-effectively.. this is another reason to use CloudFront...

upvoted 1 times

FlyingHawk 1 year, 4 months ago

Selected Answer: B

Given that the ERP system is hosted in us-east-1 and the API is likely HTTP/HTTPS-based, Option B (Amazon CloudFront) is the best solution. CloudFront's global edge network and caching capabilities will significantly improve response times for international customers while reducing the load on the backend EC2 instances.

upvoted 1 times

joanna91 1 year, 4 months ago

Selected Answer: C

C. AWS global accelerator is more suitable for cost-effective solutions to improve API response time for int'l users by using AWS global network to optimize traffic routing. Whereas, D. AWS Cloudfront is less effective for dynamic, real-time API interactions like this ERP system.

upvoted 2 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: C

A - Direct Connect requires dedicated connections for each customer, making it expensive and impractical for a public API.
B - CloudFront improves performance but caching efficiency depends on the API use case. Dynamic & personalized API responses may not benefit significantly from caching. CloudFront is best suited for static content or APIs with predictable responses.
C - Global Accelerator provides a global static IP address and uses the AWS global network to route traffic optimally, reducing latency for geographically distant users. Unlike CloudFront, it works well in this case because it doesn't rely on caching contents.
D - VPNs are not ideal for public APIs because they require specific setup for each customer, so high operational overhead.

upvoted 2 times

GOTJ 1 year, 5 months ago

Selected Answer: C

I think that the following quote "to exchange information with the ERP system" suggest that the clients use the api not only to get, but to add and/or update ERP data. If this is the case, using a Cloudfront distribution for caching should be useless as the add/update operations must be performed by the ERP system itself.

Option "C" approach is different: it's "forcing" users to access API ERP via AWS routing optimised networks by configuring the current API as an endpoint for the globally distributed endpoints groups (no need to duplicate current infrastructure or setup replication)

upvoted 3 times

Anyio 1 year, 5 months ago

Selected Answer: C

The correct answer is C.

Explanation:

Option C: Correct. AWS Global Accelerator provides access via the AWS global network, reducing latency and improving performance for globally distributed users. By using Global Accelerator, you can direct global traffic to optimal AWS endpoints, offering lower latency and better connectivity compared to standard internet routing. It achieves this without modifying the API infrastructure, making it a highly cost-effective solution for improving API response times for international users.

upvoted 1 times

tch 1 year, 5 months ago

Selected Answer: B

MOST cost-effectively -
(ERP) system in the us-east-1 Region (one region only)
will be B

AWS Global Accelerator use case for one or multiple AWS regions!

upvoted 1 times

dragosky 1 year, 5 months ago

Selected Answer: C

C. Set up AWS Global Accelerator. Configure listeners for the necessary ports. Configure endpoint groups for the appropriate Regions to distribute traffic. Create an endpoint in the group for the API.

upvoted 3 times

DhirajBansal 1 year, 6 months ago

Selected Answer: B

Cloudfront used for API acceleration and Dynamic content delivery

upvoted 1 times

A company tracks customer satisfaction by using surveys that the company hosts on its website. The surveys sometimes reach thousands of customers every hour. Survey results are currently sent in email messages to the company so company employees can manually review results and assess customer sentiment.

The company wants to automate the customer survey process. Survey results must be available for the previous 12 months.

Which solution will meet these requirements in the MOST scalable way?

A. Send the survey results data to an Amazon API Gateway endpoint that is connected to an Amazon Simple Queue Service (Amazon SQS) queue. Create an AWS Lambda function to poll the SQS queue, call Amazon Comprehend for sentiment analysis, and save the results to an Amazon DynamoDB table. Set the TTL for all records to 365 days in the future. **Most Voted**

B. Send the survey results data to an API that is running on an Amazon EC2 instance. Configure the API to store the survey results as a new record in an Amazon DynamoDB table, call Amazon Comprehend for sentiment analysis, and save the results in a second DynamoDB table. Set the TTL for all records to 365 days in the future.

C. Write the survey results data to an Amazon S3 bucket. Use S3 Event Notifications to invoke an AWS Lambda function to read the data and call Amazon Rekognition for sentiment analysis. Store the sentiment analysis results in a second S3 bucket. Use S3 lifecycle policies on each bucket to expire objects after 365 days.

D. Send the survey results data to an Amazon API Gateway endpoint that is connected to an Amazon Simple Queue Service (Amazon SQS) queue. Configure the SQS queue to invoke an AWS Lambda function that calls Amazon Lex for sentiment analysis and saves the results to an Amazon DynamoDB table. Set the TTL for all records to 365 days in the future.

Correct Answer: A

Community vote distribution

A (100%)

Comments

EllenLiu 1 year, 5 months ago

Selected Answer: A

<https://docs.aws.amazon.com/comprehend/latest/dg/what-is.html>

upvoted 1 times

trinh_le 1 year, 6 months ago

Selected Answer: A

Following keywords

reach thousands of customers every hour: use SQS

Automate handle survey: aws comprehen

upvoted 1 times

siheom 1 year, 6 months ago

Selected Answer: A

VOTE A

upvoted 2 times

Changwha 1 year, 6 months ago

A. This solution is the most scalable, cost-effective, and tailored to the requirements. It uses Amazon Comprehend for sentiment analysis, leverages serverless architecture, and retains data for 12 months using DynamoDB's TTL feature.

upvoted 2 times

A company uses AWS Systems Manager for routine management and patching of Amazon EC2 instances. The EC2 instances are in an IP address type target group behind an Application Load Balancer (ALB).


New security protocols require the company to remove EC2 instances from service during a patch. When the company attempts to follow the security protocol during the next patch, the company receives errors during the patching window.

Which combination of solutions will resolve the errors? (Choose two.)

- A. Change the target type of the target group from IP address type to instance type.
- B. Continue to use the existing Systems Manager document without changes because it is already optimized to handle instances that are in an IP address type target group behind an ALB.
- C. Implement the `AWSEC2-PatchLoadBalancerInstance` Systems Manager Automation document to manage the patching process. **Most Voted**
- D. Use Systems Manager Maintenance Windows to automatically remove the instances from service to patch the instances. **Most Voted**
- E. Configure Systems Manager State Manager to remove the instances from service and manage the patching schedule. Use ALB health checks to re-route traffic.

Correct Answer: CD

Community vote distribution



Comments

LeonSauveterre **Highly Voted** 1 year, 5 months ago

Selected Answer: CD

YOU SHOULD KNOW: The errors occur because the company is trying to patch instances that are still actively serving traffic.

A - Unnecessary when existing tools can handle the issue.

B - The existing document is not optimized for instances in an IP address type target group, so errors would persist.

C - This document automates the process of deregistering instances from the ALB during patching and re-registering them afterward, ensuring no traffic is routed to the instances while they are being patched. It directly addresses the issue and adheres to best practices.

D - Maintenance Windows allow you to define a schedule for patching and ensure instances are temporarily removed from service during the process.

E - This alone can't de-register and re-register instances.

upvoted 6 times

JoeAWS **Most Recent** 1 year, 3 months ago

Selected Answer: AC

The document deregisters instances from target group. Since they are using IP address type, it won't prevent traffic from reaching the server. Need to work choice A into the answer.

upvoted 1 times

EllenLiu 1 year, 5 months ago

Selected Answer: CD

<https://docs.aws.amazon.com/systems-manager/latest/userguide/state-manager-vs-maintenance-windows.html>

upvoted 1 times

bujuman 1 year, 6 months ago

Selected Answer: CD

<https://docs.aws.amazon.com/systems-manager-automation-runbooks/latest/userguide/automation-awsec2-patch-load-balancer-instance.html>

upvoted 3 times

aragon_saa 1 year, 6 months ago

Selected Answer: CD

Answer is CD

upvoted 3 times

A medical company wants to perform transformations on a large amount of clinical trial data that comes from several customers. The company must extract the data from a relational database that contains the customer data. Then the company will transform the data by using a series of complex rules. The company will load the data to Amazon S3 when the transformations are complete.


All data must be encrypted where it is processed before the company stores the data in Amazon S3. All data must be encrypted by using customer-specific keys.

Which solution will meet these requirements with the LEAST amount of operational effort?

- A. Create one AWS Glue job for each customer. Attach a security configuration to each job that uses server-side encryption with Amazon S3 managed keys (SSE-S3) to encrypt the data.
- B. Create one Amazon EMR cluster for each customer. Attach a security configuration to each cluster that uses client-side encryption with a custom client-side root key (CSE-Custom) to encrypt the data.
- C. Create one AWS Glue job for each customer. Attach a security configuration to each job that uses client-side encryption with AWS KMS managed keys (CSE-KMS) to encrypt the data. **Most Voted**
- D. Create one Amazon EMR cluster for each customer. Attach a security configuration to each cluster that uses server-side encryption with AWS KMS keys (SSE-KMS) to encrypt the data.

Correct Answer: C

Community vote distribution



Comments

SBoksh 10 months, 3 weeks ago

Selected Answer: B

<https://docs.aws.amazon.com/emr/latest/ReleaseGuide/emr-s3a.html>

upvoted 1 times

bora4motion 1 year, 1 month ago

Selected Answer: D

I am tempted to go with C but does it support CSE-KMS ?

upvoted 1 times

FlyingHawk 1 year, 4 months ago

Selected Answer: C

Option C is the best solution because it uses AWS Glue for serverless ETL processing and CSE-KMS for end-to-end encryption using customer-specific keys

upvoted 3 times

sOI852POL 1 year, 7 months ago

Selected Answer: C

AWS Glue for ETL, then AWS KMS (CSE-KMS)

upvoted 4 times

TewatiaAmit 1 year, 8 months ago

Selected Answer: C

AWS Glue with client-side encryption using AWS KMS (CSE-KMS) provides the required pre-processing encryption with minimal operational effort.

upvoted 2 times

blehbleh 1 year, 8 months ago

Selected Answer: C

It's C gotta use glue and since it's before the company stores the data in Amazon S3 gotta be client side.
upvoted 4 times

blehbleh 1 year, 8 months ago

Selected Answer: C

I say C, aws glue reduces the operational management uses server side encryption and ins which allows for user specific keys.
upvoted 4 times

A company hosts a website analytics application on a single Amazon EC2 On-Demand Instance. The analytics application is highly resilient and is designed to run in stateless mode.

The company notices that the application is showing signs of performance degradation during busy times and is presenting 5xx errors. The company needs to make the application scale seamlessly.

Which solution will meet these requirements MOST cost-effectively?

- A. Create an Amazon Machine Image (AMI) of the web application. Use the AMI to launch a second EC2 On-Demand Instance. Use an Application Load Balancer to distribute the load across the two EC2 instances.
- B. Create an Amazon Machine Image (AMI) of the web application. Use the AMI to launch a second EC2 On-Demand Instance. Use Amazon Route 53 weighted routing to distribute the load across the two EC2 instances.
- C. Create an AWS Lambda function to stop the EC2 instance and change the instance type. Create an Amazon CloudWatch alarm to invoke the Lambda function when CPU utilization is more than 75%.
- D. Create an Amazon Machine Image (AMI) of the web application. Apply the AMI to a launch template. Create an Auto Scaling group that includes the launch template. Configure the launch template to use a Spot Fleet. Attach an Application Load Balancer to the Auto Scaling group. **Most Voted**

Correct Answer: D

Community vote distribution

D (75%)

A (25%)

Comments

In2718 9 months, 2 weeks ago

Selected Answer: A

So, for the test the answer is D, but for real-life there's a possibility that the spot instances could all be terminated and although the app is stateless and resilient the application is down and you could be looking for another job.

As such, I worked through the other options with AI and the next best answer was A although it cannot scale, but at least it's closer to real world.

Though A costs a bit more, it's relative to the loss in revenue from the fallout of having the application periodically down and losing customers. Just an FYI, you would never do answer D in real life, but would opt for an on-demand or reserved instances instead with similar configuration.

It's funny that the answers are loaded so you can pick the "obvious" one, but to make them such, a comprehensive analysis cannot always be applied of course.

upvoted 3 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: D

A - This can't scale automatically during peak times.

B - Weighted routing is not ideal for seamless scaling because it requires manual intervention to adjust weights or add new instances.

C - It introduces downtime during instance type changes, let alone scaling.

D - (1) ASG dynamically adjusts the number of instances based on demand. (2) Spot Instances reduces costs significantly compared to On-Demand Instances. (3) ALB + ASG ensures efficient traffic routing to healthy instances.

upvoted 3 times

JA2018 1 year, 6 months ago

Selected Answer: D

3 reasons why I will choose Option D:

1. Scalability: Auto Scaling groups are designed to automatically scale up or down the number of instances based on demand, which is ideal for handling fluctuating traffic during busy periods.

2. Cost-effectiveness: Using a Spot Fleet allows you to leverage unused EC2 capacity at significantly lower costs compared to On-Demand instances.

3. Stateless Application: Since the application is stateless, it can easily be scaled horizontally with multiple instances behind an Application Load Balancer without any data consistency issues

upvoted 2 times

Razvan_C 1 year, 7 months ago

Selected Answer: D

With auto scaling group it can also scale down on low demand (costs saved)

upvoted 4 times

A company runs an environment where data is stored in an Amazon S3 bucket. The objects are accessed frequently throughout the day. The company has strict data encryption requirements for data that is stored in the S3 bucket. The company currently uses AWS Key Management Service (AWS KMS) for encryption.

The company wants to optimize costs associated with encrypting S3 objects without making additional calls to AWS KMS.

Which solution will meet these requirements?

- A. Use server-side encryption with Amazon S3 managed keys (SSE-S3).
- B. Use an S3 Bucket Key for server-side encryption with AWS KMS keys (SSE-KMS) on the new objects. **Most Voted**
- C. Use client-side encryption with AWS KMS customer managed keys.
- D. Use server-side encryption with customer-provided keys (SSE-C) stored in AWS KMS.

Correct Answer: B

Community vote distribution

B (74%)

A (26%)

Comments

8621a7c **Highly Voted** 1 year, 7 months ago

Selected Answer: B

reduce cost <https://docs.aws.amazon.com/AmazonS3/latest/userguide/bucket-key.html>

upvoted 8 times

In2718 **Most Recent** 9 months, 2 weeks ago

Selected Answer: B

Problem: The company uses SSE-KMS, and every object read/write normally generates a call to AWS KMS cost scales with object access frequency.

S3 Bucket Keys: Introduced to reduce the number of direct KMS calls by generating and caching a data key at the bucket level. Objects still benefit from SSE-KMS, but the cost of KMS requests drops significantly.

Encryption requirement: They want to keep AWS KMS-based encryption (not switch to SSE-S3). Optimization goal: Cut down KMS API calls, exactly what S3 Bucket Keys do.

upvoted 1 times

yangbo 1 year, 3 months ago

Selected Answer: A

without making additional calls to AWS KMS

upvoted 1 times

FlyingHawk 1 year, 4 months ago

Selected Answer: B

When you configure your bucket to use an S3 Bucket Key for SSE-KMS, AWS generates a short-lived bucket-level key from AWS KMS, then temporarily keeps it in S3. This bucket-level key will create data keys for new objects during its lifecycle. S3 Bucket Keys are used for a limited time period within Amazon S3, reducing the need for S3 to make requests to AWS KMS to complete encryption operations. This reduces traffic from S3 to AWS KMS, allowing you to access AWS KMS-encrypted objects in Amazon S3 at a fraction of the previous cost.

upvoted 2 times

FlyingHawk 1 year, 4 months ago

Selected Answer: B

Using S3 Bucket Keys allows you to save on AWS KMS request costs by decreasing your requests to AWS KMS for Encrypt, GenerateDataKey, and Decrypt operations through the use of a bucket-level key. By design, subsequent requests that take advantage of this bucket-level key do not result in AWS KMS API requests or validate access against the AWS KMS key policy.

upvoted 1 times

dragossky 1 year, 5 months ago

Selected Answer: B

Amazon S3 Bucket Keys reduce the cost of Amazon S3 server-side encryption with AWS Key Management Service (AWS KMS) keys (SSE-KMS). Using a bucket-level key for SSE-KMS can reduce AWS KMS request costs by up to 99 percent by decreasing the request traffic from Amazon S3 to AWS KMS. With a few clicks in the AWS Management Console, and without any changes to your client applications, you can configure your bucket to use an S3 Bucket Key for SSE-KMS encryption on new objects.

upvoted 2 times

djhtoon 1 year, 7 months ago

optimize costs associated with encrypting S3 objects without making additional calls to AWS KMS. The answer should be A and it also optimize costs- free.

upvoted 2 times

elmyth 1 year, 7 months ago

Selected Answer: A

All S3 are now encrypted by default and free with SSE-S3, but B is also possible.

upvoted 4 times

Sergantus 1 year, 6 months ago

Strict data encryption requirements wouldn't be met with SSE-S3

upvoted 2 times

A company runs multiple workloads on virtual machines (VMs) in an on-premises data center. The company is expanding rapidly. The on-premises data center is not able to scale fast enough to meet business needs. The company wants to migrate the workloads to AWS.

The migration is time sensitive. The company wants to use a lift-and-shift strategy for non-critical workloads.

Which combination of steps will meet these requirements? (Choose three.)

A. Use the AWS Schema Conversion Tool (AWS SCT) to collect data about the VMs.

B. Use AWS Application Migration Service. Install the AWS Replication Agent on the VMs. **Most Voted**

C. Complete the initial replication of the VMs. Launch test instances to perform acceptance tests on the VMs. **Most Voted**

D. Stop all operations on the VMs. Launch a cutover instance. **Most Voted**

E. Use AWS App2Container (A2C) to collect data about the VMs.

F. Use AWS Database Migration Service (AWS DMS) to migrate the VMs.

Correct Answer: BCD

Community vote distribution

BCD (100%)

Comments

LeonSauveterre **Highly Voted** 1 year, 5 months ago

Selected Answer: BCD

Well this is new. Never seen a question like that before.

WHAT YOU SHOULD KNOW:

1. Lift-and-shift means replicating on-premises workloads to AWS with minimal modifications.
 2. A cutover instance is used to direct traffic away from the source region and to the workload in the target region.
- upvoted 6 times

LeonSauveterre 1 year, 5 months ago

- A - SCT is used for database schema conversions when migrating databases, not for migrating VMs.
 - B - Application Migration Service is designed for lift-and-shift migrations. It uses the AWS Replication Agent to replicate VMs to AWS.
 - C - After replication, testing ensures that the workloads function correctly before completing the migration. (If you're wondering why, see this: <https://docs.aws.amazon.com/mgn/latest/ug/launch-cutover-gs.html>)
 - D - This is a critical step in the migration process. After testing, cutover instances are launched, and the source VMs are stopped to finalize the migration.
 - E - AWS App2Container is used for containerizing applications, not for migrating VMs.
- upvoted 2 times

JA2018 **Most Recent** 1 year, 6 months ago

Selected Answer: BCD

Why I will choose Options B, C & D:

#B: AWS Application Migration Service (AWS Migration Service) is the primary tool for a lift-and-shift migration strategy, allowing you to easily replicate on-premises VMs to AWS by installing the Replication Agent on the VMs.

#C: Complete initial replication and launch test instances:

After the initial replication of VMs to AWS, it's crucial to perform acceptance tests on the newly migrated instances to ensure functionality before full transitioning.

#D: Stop operations and launch a cutover instance:

For a cutover migration, you would stop operations on the on-premises VMs and then launch the migrated instances on AWS to switch over completely

upvoted 3 times

JA2018 1 year, 6 months ago

Why the other options are not suitable:

A and E: AWS Schema Conversion Tool (AWS SCT) and AWS App2Container (A2C):

These tools are primarily used for database migration and containerization respectively, which are not the best fit for a simple lift-and-shift migration of VMs.

F: AWS Database Migration Service (AWS DMS):

While AWS DMS is used for database migration, the question specifies migrating entire VMs, not just databases.

upvoted 1 times

JA2018 1 year, 6 months ago

ey points to remember:

For a lift-and-shift migration, AWS Application Migration Service is the recommended tool.

Always perform testing on migrated instances before fully transitioning to AWS.

A cutover migration involves stopping operations on the source and switching to the new environment.

upvoted 4 times

aragon_saa 1 year, 6 months ago

Selected Answer: BCD

Answer is BCD

upvoted 2 times

A company hosts an application in a private subnet. The company has already integrated the application with Amazon Cognito. The company uses an Amazon Cognito user pool to authenticate users.

The company needs to modify the application so the application can securely store user documents in an Amazon S3 bucket.

Which combination of steps will securely integrate Amazon S3 with the application? (Choose two.)

- A. Create an Amazon Cognito identity pool to generate secure Amazon S3 access tokens for users when they successfully log in. **Most Voted**
- B. Use the existing Amazon Cognito user pool to generate Amazon S3 access tokens for users when they successfully log in.
- C. Create an Amazon S3 VPC endpoint in the same VPC where the company hosts the application. **Most Voted**
- D. Create a NAT gateway in the VPC where the company hosts the application. Assign a policy to the S3 bucket to deny any request that is not initiated from Amazon Cognito.
- E. Attach a policy to the S3 bucket that allows access only from the users' IP addresses.

Correct Answer: AC

Community vote distribution

AC (100%)

Comments

bujuman **Highly Voted** 1 year, 6 months ago

Selected Answer: AC

securely integrate Amazon S3 with the application:

<https://docs.aws.amazon.com/cognito/latest/developerguide/identity-pools.html>

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/privatelink-interface-endpoints.html#types-of-vpc-endpoints-for-s3>

upvoted 7 times

FlyingHawk **Most Recent** 1 year, 4 months ago

Selected Answer: AC

An Amazon Cognito identity pool allows you to generate temporary AWS credentials (e.g., access tokens) for authenticated users. These credentials can be used to securely access AWS services like Amazon S3.

By creating an identity pool, the application can generate secure Amazon S3 access tokens for users after they successfully log in, ensuring that only authenticated users can upload or access documents in the S3 bucket.

upvoted 1 times

FlyingHawk 1 year, 4 months ago

C_ A VPC endpoint for Amazon S3 allows private connectivity between the application (hosted in a private subnet) and the S3 bucket. This ensures that data transfer between the application and S3 stays within the AWS network, improving security and performance.

upvoted 1 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: AC

A - Amazon Cognito identity pools provide temporary AWS credentials for authenticated users.

B - User pools are for authentication (who the user is). Identity pools are for authorization (what the user can do).

C - Traffic between your VPC and S3 stays within the AWS network, good. This also removes the need for a NAT Gateway for S3 access.

D - Unnecessary and not secure.

E - Users' IP addresses can change frequently (especially mobile users).

upvoted 3 times

viejito 1 year, 7 months ago

respuesta correcta : A - B

upvoted 2 times

dragosky 1 year, 5 months ago

identity pool - authorization

user pool - authentication, in this case something needs to modify docs, so authorization

upvoted 1 times

luther77 1 year, 6 months ago

B doesnt make sense here. because user pools are used for authentication, not authorization

upvoted 2 times

A company has a three-tier web application that processes orders from customers. The web tier consists of Amazon EC2 instances behind an Application Load Balancer. The processing tier consists of EC2 instances. The company decoupled the web tier and processing tier by using Amazon Simple Queue Service (Amazon SQS). The storage layer uses Amazon DynamoDB.

At peak times, some users report order processing delays and hangs. The company has noticed that during these delays, the EC2 instances are running at 100% CPU usage, and the SQS queue fills up. The peak times are variable and unpredictable.

The company needs to improve the performance of the application.

Which solution will meet these requirements?

- A. Use scheduled scaling for Amazon EC2 Auto Scaling to scale out the processing tier instances for the duration of peak usage times. Use the CPU Utilization metric to determine when to scale.
- B. Use Amazon ElastiCache for Redis in front of the DynamoDB backend tier. Use target utilization as a metric to determine when to scale.
- C. Add an Amazon CloudFront distribution to cache the responses for the web tier. Use HTTP latency as a metric to determine when to scale.
- D. Use an Amazon EC2 Auto Scaling target tracking policy to scale out the processing tier instances. Use the ApproximateNumberOfMessages attribute to determine when to scale. **Most Voted**

Correct Answer: D

Community vote distribution

D (85%)

B (15%)

Comments

In2718 10 months ago

Selected Answer: B

In real life I would look at where the real bottleneck is, which as stated below may be the database that is causing the EC2 CPU thrashing, in which case an ELB may be the better solution. Again, a poorly written question, which I believe they want you to choose answer D, and would only do after a thorough investigation if the extra overhead and potential cost would justify answer D.

upvoted 1 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: D

Just scale base on how many tasks left in the queue. CPU utilization is not an accurate metric. There could be just one task but it takes quite a while to process, in which case scaling is unnecessary.

upvoted 2 times

FlyingHawk 1 year, 5 months ago

Selected Answer: D

The bottleneck is at the processing tier, B and C are incorrect, processing tier instances are running at 100% CPU usage, and the SQS queue fills up. Because the peak times are variable and unpredictable, we should use the EC2 Auto Scaling target tracking policy to scale out the processing tier instances based on the size of the queue, A uses scheduled scaling for Amazon EC2 Auto Scaling is incorrect.

upvoted 3 times

GOTJ 1 year, 5 months ago

Selected Answer: D

Check this out:

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/as-using-sqs-queue.html>

upvoted 1 times

EllenLiu 1 year, 5 months ago

Selected Answer: D

The bottleneck is from the processing tier.

upvoted 3 times

dragossky 1 year, 5 months ago

Selected Answer: B

Yes, ElastiCache can be used with DynamoDB to improve the performance and scalability of read-heavy or frequently accessed workloads. Amazon ElastiCache is a fully managed in-memory data store that supports popular caching engines such as Memcached and Redis.

upvoted 1 times

GOTJ 1 year, 5 months ago

Agree. But this "SQS queue fill up" (remember, SQS is meant for decouple web and processing tiers) suggest that the bottleneck is produced by the processing CPU bottleneck because web servers are sending more requests to the queue than the processing servers are able to process from the queue. As the first approach, I would try option D

upvoted 2 times

Salilgen 1 year, 5 months ago

The problem is from application tier

upvoted 1 times

EllenLiu 1 year, 5 months ago

DynamoDB Accelerator is cache for dynamoDB

upvoted 1 times

aragon_saa 1 year, 7 months ago

Selected Answer: D

Answer is D

upvoted 2 times

A company's production environment consists of Amazon EC2 On-Demand Instances that run constantly between Monday and Saturday. The instances must run for only 12 hours on Sunday and cannot tolerate interruptions. The company wants to cost-optimize the production environment.

Which solution will meet these requirements MOST cost-effectively?

A. Purchase Scheduled Reserved Instances for the EC2 instances that run for only 12 hours on Sunday. Purchase Standard Reserved Instances for the EC2 instances that run constantly between Monday and Saturday. **Most Voted**

B. Purchase Convertible Reserved Instances for the EC2 instances that run for only 12 hours on Sunday. Purchase Standard Reserved Instances for the EC2 instances that run constantly between Monday and Saturday.

C. Use Spot Instances for the EC2 instances that run for only 12 hours on Sunday. Purchase Standard Reserved Instances for the EC2 instances that run constantly between Monday and Saturday.

D. Use Spot Instances for the EC2 instances that run for only 12 hours on Sunday. Purchase Convertible Reserved Instances for the EC2 instances that run constantly between Monday and Saturday.

Correct Answer: A

Community vote distribution

A (100%)

Comments

LeonSauveterre 1 year, 5 months ago

Selected Answer: A

Scheduled Reserved Instance: Perfect for workloads with predictable and recurring schedules (like 12 hours on Sunday). This is cost-efficient and ensures uninterrupted instances.

Convertible Reserved Instance: Allows flexibility to change instance types, sizes, or OS during the term. However, they are not ideal for short, predictable schedules like 12 hours on Sunday. Scheduled Reserved Instances are a better fit.

Spot Instances can be interrupted at any time. They do not guarantee uninterrupted operations.

upvoted 3 times

FlyingHawk 1 year, 5 months ago

Selected Answer: A

as the instance cannot tolerate interruptions, so spot instances are out, C and D are incorrect, leave only A and B, Scheduled RI or Convertible RI, based on the AWS doc, Purchasing Reserved Instances on a recurring schedule lets you pay for compute capacity by the hour and obtain a capacity reservation ahead of time for only the time periods you will need it. as it only 12 hours on Sunday, so A will be cheaper than B <https://aws.amazon.com/ec2/pricing/reserved-instances/pricing/>.

upvoted 1 times

FlyingHawk 1 year, 5 months ago

Not the minimum commit for RI is one year, if you want the flexibility of change instance type within one year, B might be a better choice, however, considering the flexibility, saving plan for computing is a recommended solution over RIs, I guess this question should not appear in the exam if AWS wants to advocate the saving plan.

upvoted 1 times

GOTJ 1 year, 5 months ago

Selected Answer: A

"A" is the most straightforward option. However, I'm afraid this feature could be eventually removed and substituted by On-Demand Capacity Reservations and Savings Plans. If you try to use scheduled reservation in the EC2 console, you'll get the following banner:

"We recommend Savings Plans over Reserved Instances. Savings Plans are the easiest and most flexible way to save money on your AWS compute costs and offer lower prices (up to 72% off) just like Reserved Instances. To learn more and get started with Savings Plans click here"

upvoted 1 times

EllenLiu 1 year, 5 months ago

Selected Answer: A

Scheduled reserved instances are purchase of reserved Instances on a recurring schedule
<https://aws.amazon.com/ec2/pricing/reserved-instances/buyer/>

upvoted 1 times

bujuman 1 year, 6 months ago

Selected Answer: A

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-reserved-instances.html>

upvoted 3 times

8621a7c 1 year, 7 months ago

Selected Answer: A

standard reserved instance is cheaper than convertible reserved instance

upvoted 4 times

A digital image processing company wants to migrate its on-premises monolithic application to the AWS Cloud. The company processes thousands of images and generates large files as part of the processing workflow.

The company needs a solution to manage the growing number of image processing jobs. The solution must also reduce the manual tasks in the image processing workflow. The company does not want to manage the underlying infrastructure of the solution.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use Amazon Elastic Container Service (Amazon ECS) with Amazon EC2 Spot Instances to process the images. Configure Amazon Simple Queue Service (Amazon SQS) to orchestrate the workflow. Store the processed files in Amazon Elastic File System (Amazon EFS).
- B. Use AWS Batch jobs to process the images. Use AWS Step Functions to orchestrate the workflow. Store the processed files in an Amazon S3 bucket. **Most Voted**
- C. Use AWS Lambda functions and Amazon EC2 Spot Instances to process the images. Store the processed files in Amazon FSx.
- D. Deploy a group of Amazon EC2 instances to process the images. Use AWS Step Functions to orchestrate the workflow. Store the processed files in an Amazon Elastic Block Store (Amazon EBS) volume.

Correct Answer: B

Community vote distribution

B (100%)

Comments

MortisG 1 year ago

Selected Answer: B

All other options require managing infrastructure. AWS Batch handles the processing and Step Functions the workflow. Best solutions for what's solicited.

upvoted 1 times

ElI89 1 year, 2 months ago

Selected Answer: B

B- because all other answers mention EC2 which goes against the company not wanting to manage the underlying infrastructure.

upvoted 2 times

LeonSauveterre 1 year, 5 months ago

Selected Answer: B

A & C - Spot Instances are cost-effective but are subject to interruptions. "The solution must also reduce the manual tasks."

B - AWS Batch is a fully managed service for batch processing, ideal for workloads like image processing. It scales automatically, handles job scheduling, and eliminates the need to manage infrastructure. Step Functions provides robust orchestration with built-in support for retries, error handling, and dependencies, reducing manual tasks in the workflow.

D - Using EC2 instances requires managing the underlying infrastructure. Also, EBS is not as scalable or cost-effective as S3 when you want to store large, growing numbers of files.

upvoted 3 times

trinh_le 1 year, 6 months ago

Selected Answer: B

A. Use Amazon Elastic Container Service (Amazon ECS) complexity

C. Amazon FSx is not an ideal compared to s3

D. Use a group of Amazon EC2 instances do not scalable

upvoted 4 times

aragon_saa 1 year, 7 months ago

Selected Answer: B

Answer is B

upvoted 4 times

A company's image-hosting website gives users around the world the ability to upload, view, and download images from their mobile devices. The company currently hosts the static website in an Amazon S3 bucket.

Because of the website's growing popularity, the website's performance has decreased. Users have reported latency issues when they upload and download images.

The company must improve the performance of the website.

Which solution will meet these requirements with the LEAST implementation effort?

A. Configure an Amazon CloudFront distribution for the S3 bucket to improve the download performance. Enable S3 Transfer Acceleration to improve the upload performance. **Most Voted**

B. Configure Amazon EC2 instances of the right sizes in multiple AWS Regions. Migrate the application to the EC2 instances. Use an Application Load Balancer to distribute the website traffic equally among the EC2 instances. Configure AWS Global Accelerator to address global demand with low latency.

C. Configure an Amazon CloudFront distribution that uses the S3 bucket as an origin to improve the download performance. Configure the application to use CloudFront to upload images to improve the upload performance. Create S3 buckets in multiple AWS Regions. Configure replication rules for the buckets to replicate users' data based on the users' location. Redirect downloads to the S3 bucket that is closest to each user's location.

D. Configure AWS Global Accelerator for the S3 bucket to improve network performance. Create an endpoint for the application to use Global Accelerator instead of the S3 bucket.

Correct Answer: A

Community vote distribution

A (100%)

Comments

FlyingHawk 1 year, 4 months ago

Selected Answer: A

A - By configuring a CloudFront distribution for the S3 bucket, the company can significantly improve download performance for users by serving images from the nearest edge location. S3 Transfer Acceleration uses CloudFront's globally distributed edge locations to accelerate upload performance. It optimizes the path between the user's device and the S3 bucket, reducing latency and improving upload speeds.

upvoted 2 times

FlyingHawk 1 year, 4 months ago

Option C (Configure CloudFront for downloads, use CloudFront for uploads, create S3 buckets in multiple Regions, and configure replication rules): While this option improves performance, it requires significant effort to configure multiple S3 buckets, replication rules, and redirection logic. It is more complex than necessary for the company's requirements.

upvoted 2 times

Denise123 1 year, 5 months ago

Selected Answer: A

Obvious.

upvoted 1 times

aragon_saa 1 year, 7 months ago

Selected Answer: A

Answer is A

upvoted 2 times

A company runs an application in a private subnet behind an Application Load Balancer (ALB) in a VPC. The VPC has a NAT gateway and an internet gateway. The application calls the Amazon S3 API to store objects.

According to the company's security policy, traffic from the application must not travel across the internet.

Which solution will meet these requirements MOST cost-effectively?

- A. Configure an S3 interface endpoint. Create a security group that allows outbound traffic to Amazon S3.
- B. Configure an S3 gateway endpoint. Update the VPC route table to use the endpoint. **Most Voted**
- C. Configure an S3 bucket policy to allow traffic from the Elastic IP address that is assigned to the NAT gateway.
- D. Create a second NAT gateway in the same subnet where the legacy application is deployed. Update the VPC route table to use the second NAT gateway.

Correct Answer: B

Community vote distribution

B (100%)

Comments

sOI852POL **Highly Voted** 1 year, 7 months ago

Selected Answer: B

S3 gateway endpoint

<https://docs.aws.amazon.com/vpc/latest/privatelink/vpc-endpoints-s3.html>

upvoted 7 times

GOTJ **Most Recent** 1 year, 5 months ago

Selected Answer: B

Here's another document that support option "B" as the most cost-effectively:

<https://aws.amazon.com/blogs/architecture/choosing-your-vpc-endpoint-strategy-for-amazon-s3/>

upvoted 3 times

trinh_le 1 year, 6 months ago

Selected Answer: B

Keyword: ec2 connect s3 -> vpc endpoint or s3 endpoint

upvoted 4 times

aragon_saa 1 year, 7 months ago

Selected Answer: B

Answer is B

upvoted 2 times

A company has an application that runs on an Amazon Elastic Kubernetes Service (Amazon EKS) cluster on Amazon EC2 instances. The application has a UI that uses Amazon DynamoDB and data services that use Amazon S3 as part of the application deployment.

The company must ensure that the EKS Pods for the UI can access only Amazon DynamoDB and that the EKS Pods for the data services can access only Amazon S3. The company uses AWS Identity and Access Management (IAM).

Which solution meets these requirements?

- A. Create separate IAM policies for Amazon S3 and DynamoDB access with the required permissions. Attach both IAM policies to the EC2 instance profile. Use role-based access control (RBAC) to control access to Amazon S3 or DynamoDB for the respective EKS Pods.
- B. Create separate IAM policies for Amazon S3 and DynamoDB access with the required permissions. Attach the Amazon S3 IAM policy directly to the EKS Pods for the data services and the DynamoDB policy to the EKS Pods for the UI.
- C. Create separate Kubernetes service accounts for the UI and data services to assume an IAM role. Attach the AmazonS3FullAccess policy to the data services account and the AmazonDynamoDBFullAccess policy to the UI service account. **Most Voted**
- D. Create separate Kubernetes service accounts for the UI and data services to assume an IAM role. Use IAM Role for Service Accounts (IRSA) to provide access to the EKS Pods for the UI to Amazon S3 and the EKS Pods for the data services to DynamoDB.

Correct Answer: C

Community vote distribution

C (69%)

D (31%)

Comments

jingen11 **Highly Voted** 1 year, 7 months ago

Selected Answer: D

<https://docs.aws.amazon.com/eks/latest/userguide/service-accounts.html#service-accounts-iam>
upvoted 9 times

LeonSauveterre **Highly Voted** 1 year, 5 months ago

Selected Answer: C

A - This is a major security flaw. All Pods running on the EC2 instances would inherit both the S3 and DynamoDB permissions.
B - Directly attaching IAM policies to Pods is not a valid AWS mechanism. You should use roles.
C - Not recommended unless we throw least privilege principles to the wind. BUT this is the only feasible option here, so.
D - This option mixes up the permissions. It gives UI Pods access to S3 and data service Pods access to DynamoDB, which is the OPPOSITE of the requirement.
upvoted 7 times

je33 **Most Recent** 8 months, 1 week ago

Selected Answer: D

Option C's FullAccess policy violates the principle of least privilege due to excessive permissions.
upvoted 2 times

je33 8 months, 1 week ago

I was mistaken. D has the wrong permissions assigned.
upvoted 1 times

FlyingHawk 1 year, 4 months ago

Selected Answer: C

IRSA allows Kubernetes service accounts to assume IAM roles, enabling fine-grained access control for EKS Pods. This ensures that each Pod can access only the AWS resources it needs.

Attach the AmazonDynamoDBFullAccess policy to the service account for the UI Pods, allowing them to access only DynamoDB.

Attach the AmazonS3FullAccess policy to the service account for the data services Pods, allowing them to access only S3.

upvoted 4 times

BugsyWarribwoy 1 year, 5 months ago

Selected Answer: C

D is sneakily WRONG because it swaps the services incorrectly. Pay attention. UWC.

upvoted 6 times

GOTJ 1 year, 5 months ago

Selected Answer: C

Even though there are a couple of comments rightfully discarding "D" because the UI --> S3/DataServices --> DynamoDB swap, I found an AWS document claiming that Kubernetes Service accounts and IAM Role for Service Accounts combination should be the right answer:

<https://docs.aws.amazon.com/eks/latest/userguide/iam-roles-for-service-accounts.html>

Having this in mind, I'm also discarding "D" (right reasoning, wrong scenario), as well "A" and "B". I didn't like the "full access" policy, but is technically correct, so my vote goes to "C"

upvoted 3 times

trinh_le 1 year, 6 months ago

Selected Answer: C

* A. Attach both IAM policies to the EC2 instance profile. Does not separate each EKS Pod

* B. Aws does not support Attach the Amazon S3 IAM policy directly to the EKS Pods

*

* D. provide access "UI to Amazon S3" and "data services to DynamoDB" => it does not meet requirements

upvoted 3 times

rosanna 1 year, 6 months ago

Selected Answer: C

The answer is C as they're switching data pods with DynamoDB service and vice versa (configure the wrong resources)

upvoted 2 times

tm1000000 1 year, 7 months ago

answer is D

upvoted 1 times

A company needs to give a globally distributed development team secure access to the company's AWS resources in a way that complies with security policies.

The company currently uses an on-premises Active Directory for internal authentication. The company uses AWS Organizations to manage multiple AWS accounts that support multiple projects.

The company needs a solution to integrate with the existing infrastructure to provide centralized identity management and access control.

Which solution will meet these requirements with the LEAST operational overhead?

A. Set up AWS Directory Service to create an AWS managed Microsoft Active Directory on AWS. Establish a trust relationship with the on-premises Active Directory. Use IAM roles that are assigned to Active Directory groups to access AWS resources within the company's AWS accounts.

B. Create an IAM user for each developer. Manually manage permissions for each IAM user based on each user's involvement with each project. Enforce multi-factor authentication (MFA) as an additional layer of security.

C. Use AD Connector in AWS Directory Service to connect to the on-premises Active Directory. Integrate AD Connector with AWS IAM Identity Center. Configure permissions sets to give each AD group access to specific AWS accounts and resources.

Most Voted

D. Use Amazon Cognito to deploy an identity federation solution. Integrate the identity federation solution with the on-premises Active Directory. Use Amazon Cognito to provide access tokens for developers to access AWS accounts and resources.

Correct Answer: C

Community vote distribution

C (80%)

A (20%)

Comments

LeonSauveterre 1 year, 5 months ago

Selected Answer: C

A - On-premises Active Directory already exists.

B - "Manually" ? Oh hell no. But MFA is something I guess.

C - AD Connector seamlessly connects AWS to the on-premises Active Directory without the need to synchronize or replicate the directory. Identity Center (formerly AWS SSO) allows centralized access management across AWS accounts in an AWS Organizations setup. Permissions sets can be configured to map Active Directory groups to specific AWS accounts and resources, making access control easy and secure.

D - Amazon Cognito is better suited for application-level identity management (like customer-facing apps), not for internal teams working across multiple AWS accounts.

upvoted 3 times

GOTJ 1 year, 5 months ago

Selected Answer: A

Why not "A"? Check out the note of this link: https://docs.aws.amazon.com/directoryservice/latest/admin-guide/directory_ad_connector.html

AD Connector cannot be shared with other AWS accounts. If this is a requirement, consider using AWS Managed Microsoft AD to Share your AWS Managed Microsoft AD. AD Connector is also not multi-VPC aware, which means that AWS applications like WorkSpaces are required to be provisioned into the same VPC as your AD Connector.

And I think managing multiple aws accounts is, indeed, a requirement

upvoted 3 times

FlyingHawk 1 year, 4 months ago

While this option provides centralized identity management, it requires setting up a separate AWS managed Microsoft AD, which increases operational overhead. A trust relationship is also more complex to configure than AD Connector.

upvoted 1 times

GOTJ 1 year, 4 months ago

I guess so, but... is the alternative configure one AD connector per account more operational friendly?

upvoted 1 times

trinh_le 1 year, 6 months ago

Selected Answer: C

C: AD Connector allows AWS to use the on-premises Active Directory for authentication without replicating directory data to AWS.

- Maintains centralized identity management in the on-premises directory, adhering to the company's security policies.

A: Higher Overhead: Requires creating and maintaining a separate managed Active Directory instance in AWS.

B: Scalability Issues: Manually creating and managing IAM users for a globally distributed team is cumbersome.

D: Unnecessary Federation Layer: Cognito is more suited for customer identity use cases rather than managing internal developer access to AWS resources.

upvoted 3 times

78b9037 1 year, 6 months ago

Selected Answer: C

Once set up, it requires minimal ongoing management. User provisioning and deprovisioning are handled through the existing Active Directory processes.

upvoted 2 times

xekiva3329 1 year, 7 months ago

Selected Answer: C

answer C

upvoted 2 times

aragon_saa 1 year, 8 months ago

Selected Answer: C

Answer is C

upvoted 2 times

A company is developing an application in the AWS Cloud. The application's HTTP API contains critical information that is published in Amazon API Gateway. The critical information must be accessible from only a limited set of trusted IP addresses that belong to the company's internal network.

Which solution will meet these requirements?

- A. Set up an API Gateway private integration to restrict access to a predefined set of IP addresses.
- B. Create a resource policy for the API that denies access to any IP address that is not specifically allowed. **Most Voted**
- C. Directly deploy the API in a private subnet. Create a network ACL. Set up rules to allow the traffic from specific IP addresses.
- D. Modify the security group that is attached to API Gateway to allow inbound traffic from only the trusted IP addresses.

Correct Answer: B

Community vote distribution

B (92%)

8%

Comments

Jassn3 **Highly Voted** 1 year, 8 months ago
OMG finally I have reached here
upvoted 38 times

Bwhizzy **Highly Voted** 1 year, 7 months ago
Finally made it. Congratulations to everyone who got here. I know the journey has been hard and long. best of luck
upvoted 22 times

DRODRO **Most Recent** 6 months, 3 weeks ago
Selected Answer: B
Good luck to everyone.
upvoted 5 times

whoeevd2eww3322d 1 year, 4 months ago
🇺🇸 I cleared the AWS Solutions Architect – Associate exam!

For those on the same journey, many exam questions are from or related to these 1019 questions. Go through them twice for a solid prep. Also, check out Shaping Pixel's YouTube channel—focus only on the latest videos as most questions overlap.

You've got this—good luck! 🚀
upvoted 13 times

techghost 1 year, 4 months ago
Selected Answer: B
Got to the end, worth it and great contributions
upvoted 2 times

LeonSauveterre 1 year, 5 months ago
Selected Answer: B

A - A private integration is for connecting API Gateway to private resources within VPCs, such as EC2 instances or Lambda functions. It doesn't restrict access based on IP addresses. It controls how the API interacts with private resources.
B - This is the most direct way to restrict access at the API Gateway level.
C - Deploying the API in a private subnet would require users to connect to your VPC (e.g., via VPN or Direct Connect). This is not what the question implies (the API should be publicly accessible, but restricted by IP).
D - API Gateway doesn't even use security groups because it's a regional managed service, not directly tied to VPCs like EC2 instances.
upvoted 5 times

LeonSauveterre 1 year, 5 months ago

Now people let's go back to square one and nail this frickin' exam!

upvoted 6 times

jfedotov 1 year, 5 months ago

Selected Answer: B

The correct answer should be B, but the question is misleading. This is because HTTP APIs do not support resource policies; only REST APIs do.

upvoted 2 times

Denise123 1 year, 5 months ago

Selected Answer: B

I can't believe that I made it! Congrats everyone! Time to restart from the Q1 to polish, and good luck for the exam!

upvoted 3 times

EllenLiu 1 year, 5 months ago

Selected Answer: B

<https://docs.aws.amazon.com/apigateway/latest/developerguide/apigateway-resource-policies-examples.html#apigateway-resource-policies-source-ip-address-example>

upvoted 4 times

JA2018 1 year, 6 months ago

Selected Answer: D

yes, finally, the finishing point.....

Why not D?

By modifying the security group attached to API Gateway, you can explicitly define which IP addresses are allowed to access the API, effectively restricting access to only the trusted internal network IPs

upvoted 3 times

trinh_le 1 year, 6 months ago

Selected Answer: B

B: api gateway allows configure resource policy to restricted IP

A: private integration-> private api

C: private subnet -> private api

D: security group does not support api gateway

upvoted 6 times

Sergantus 1 year, 6 months ago

Congrats to everyone who made it here. On to the next step, gents.

upvoted 7 times

Ben_88 1 year, 7 months ago

And now it's time to start all over again to see if you can do better this time . good luck lads !

upvoted 7 times

xekiva3329 1 year, 7 months ago

Selected Answer: B

answer B

upvoted 6 times